

Command Control Interface

01-83-03/03

Command Reference

This document describes and provides instructions for using the Command Control Interface (CCI) software to configure and perform operations on the Hitachi RAID storage systems.

© 2010, 2025 Hitachi, Ltd. All rights reserved.

No part of this publication may be reproduced or transmitted in any form or by any means, electronic or mechanical, including copying and recording, or stored in a database or retrieval system for commercial purposes without the express written permission of Hitachi, Ltd., Hitachi Vantara, Ltd., or Hitachi Vantara LLC (collectively "Hitachi"). Licensee may make copies of the Materials provided that any such copy is: (i) created as an essential step in utilization of the Software as licensed and is used in no other manner; or (ii) used for archival purposes. Licensee may not make any other copies of the Materials. "Materials" mean text, data, photographs, graphics, audio, video and documents.

Hitachi reserves the right to make changes to this Material at any time without notice and assumes no responsibility for its use. The Materials contain the most current information available at the time of publication.

Some of the features described in the Materials might not be currently available. Refer to the most recent product announcement for information about feature and product availability, or contact Hitachi Vantara LLC at https://support.hitachivantara.com/en_us/contact-us.html.

Notice: Hitachi products and services can be ordered only under the terms and conditions of the applicable Hitachi agreements. The use of Hitachi products is governed by the terms of your agreements with Hitachi Vantara LLC.

By using this software, you agree that you are responsible for:

1. Acquiring the relevant consents as may be required under local privacy laws or otherwise from authorized employees and other individuals; and
2. Verifying that your data continues to be held, retrieved, deleted, or otherwise processed in accordance with relevant laws.

Notice on Export Controls. The technical data and technology inherent in this Document may be subject to U.S. export control laws, including the U.S. Export Administration Act and its associated regulations, and may be subject to export or import regulations in other countries. Reader agrees to comply strictly with all such regulations and acknowledges that Reader has the responsibility to obtain licenses to export, re-export, or import the Document and any Compliant Products.

Hitachi and Lumada are trademarks or registered trademarks of Hitachi, Ltd., in the United States and other countries.

AlX, DB2, DS6000, DS8000, Enterprise Storage Server, eServer, FICON, FlashCopy, GDPS, HyperSwap, IBM, OS/390, PowerHA, PowerPC, S/390, System z9, System z10, Tivoli, z/OS, z9, z10, z13, z14, z15, z16, z/VM, and z/VSE are registered trademarks or trademarks of International Business Machines Corporation.

Active Directory, ActiveX, Bing, Excel, Hyper-V, Internet Explorer, the Internet Explorer logo, Microsoft, Microsoft Edge, the Microsoft corporate logo, the Microsoft Edge logo, MS-DOS, Outlook, PowerPoint, SharePoint, Silverlight, SmartScreen, SQL Server, Visual Basic, Visual C++, Visual Studio, Windows, the Windows logo, Windows Azure, Windows PowerShell, Windows Server, the Windows start button, and Windows Vista are registered trademarks or trademarks of Microsoft Corporation. Microsoft product screen shots are reprinted with permission from Microsoft Corporation.

All other trademarks, service marks, and company names in this document or website are properties of their respective owners.

The open source content used in Hitachi Vantara products may be found within the Product documentation or you may request a copy of such information (including source code and/or modifications to the extent the license for any open source requires Hitachi make it available) by sending an email to OSS_licensing@hitachivantara.com.

Contents

Preface.....	10
Product version.....	10
Changes in this revision.....	10
Release notes.....	10
Accessing product documentation.....	11
Getting help.....	11
Comments.....	11
Chapter 1: Differences between CCI and Device Manager - Storage Navigator.....	12
Supported characters.....	12
Maximum number of characters.....	17
Operational differences.....	18
Methods for specifying IP addresses.....	20
Displaying NAA identifiers (LUN WWNs) and namespace identifiers (NGU IDs).....	21
Specifying a mirror unit (MU) number by using the -d option.....	22
Chapter 2: Data management commands.....	23
paircreate*	23
pairsplit*	42
pairresync*	56
pairevtwait*	71
pairmon*	81
pairvolchk*	84
pairdisplay*	96
paircurchk* (for TrueCopy/global-active device).....	129
pairsyncwait*	135
horctakeover*	143
raidscan*	148
raidar*	168
raidqry*	171
raidvchkset*	176
raidvchkdsp*	180
raidvchkscan*	187

raidvchkscan for Universal Replicator*	192
raidvchkscan for Thin Image, Copy-on-Write Snapshot, and DP pools*	198
horcmstart	207
horcmshutdown	209
horcctl*	210
horctakeoff	213
Chapter 3: Subcommands	219
Windows subcommands	219
findcmddev	219
drivescan	220
portscan	222
sync, syncd	223
mount	225
umount, umountd	227
Environment variable subcommands	229
Chapter 4: Command tools	231
inqraid	231
mkconf	248
rmawk	252
Chapter 5: Configuration setting commands	262
raidcom	262
Method for specifying LDEV number	266
Methods for specifying multiple LDEVs	266
Operations where multiple LDEVs can be specified	267
Specifying and displaying storage system serial numbers	269
Maximum number of acceptable asynchronous commands	269
Command execution with MP blade or controller failure	269
Resource group operation	273
Resource lock operation	274
Resource locking and CCI commands	275
Request ID function	293
raidcom add clpr	294
raidcom delete clpr	295
raidcom get clpr	296
raidcom modify clpr	297
raidcom get system_opt	299
raidcom modify user_system_opt	301
raidcom get user_system_opt	302
raidcom get command_status	303

raidcom reset command_status.....	307
raidcom add copy_grp.....	309
raidcom delete copy_grp.....	310
raidcom get copy_grp.....	311
raidcom add device_grp.....	312
raidcom delete device_grp.....	313
raidcom get device_grp.....	313
raidcom get drive	315
raidcom modify drive	316
raidcom get error_message.....	317
raidcom add external_grp.....	318
raidcom check_ext_storage external_grp.....	323
raidcom delete external_grp.....	324
raidcom disconnect external_grp.....	325
raidcom get external_grp.....	327
raidcom modify external_grp.....	330
raidcom discover external_storage.....	331
raidcom add host_grp.....	335
raidcom delete host_grp.....	337
raidcom get host_grp.....	338
raidcom modify host_grp.....	344
raidcom add host_nqn.....	346
raidcom get host_nqn.....	347
raidcom modify host_nqn.....	348
raidcom delete host_nqn.....	349
raidcom add chap_user.....	350
raidcom delete chap_user.....	351
raidcom set chap_user.....	352
raidcom reset chap_user.....	353
raidcom get chap_user.....	354
raidcom add hba_wwn.....	355
raidcom delete hba_wwn.....	356
raidcom get hba_wwn.....	357
raidcom reset hba_wwn.....	358
raidcom set hba_wwn.....	359
raidcom add hba_iscsi.....	359
raidcom delete hba_iscsi.....	360
raidcom set hba_iscsi.....	361
raidcom reset hba_iscsi.....	362
raidcom get hba_iscsi.....	363
raidcom add external_iscsi_name.....	364

raidcom delete external_iscsi_name.....	368
raidcom get external_iscsi_name.....	370
raidcom get initiator_iscsi_name.....	374
raidcom discover external_iscsi_name.....	376
raidcom check external_iscsi_name.....	379
raidcom modify external_chap_user.....	382
raidcom modify initiator_chap_user.....	385
raidcom add journal.....	386
raidcom delete journal.....	388
raidcom get journal.....	389
raidcom modify journal.....	393
raidcom add ldev.....	398
raidcom delete ldev.....	408
raidcom extend ldev.....	411
raidcom get ldev.....	413
Internal volume examples.....	419
External volume examples.....	428
Dynamic Provisioning V-VOL examples.....	431
Dynamic Provisioning V-VOL namespace examples.....	437
Deduplication system data volume example.....	442
Volume deletion examples.....	443
Dynamic Tiering V-VOL examples.....	444
Thin Image primary volume example.....	447
Pool volume example.....	449
rmawk command examples.....	450
Capacity saving examples.....	452
QoS examples.....	453
raidcom initialize ldev.....	457
raidcom modify ldev.....	461
raidcom add license.....	473
raidcom delete license.....	473
raidcom modify license.....	474
raidcom get license.....	475
raidcom add quorum.....	477
raidcom delete quorum.....	478
raidcom modify quorum.....	479
raidcom get quorum.....	480
raidcom replace quorum.....	482
raidcom modify local_replica_opt.....	482
raidcom get local_replica_opt.....	483
raidcom get remote_replica_opt.....	484

raidcom modify remote_replica_opt.....	487
raidcom add lun.....	488
raidcom delete lun.....	492
raidcom discover lun.....	494
Getting the external storage system's iSCSI target information corresponding to the pseudo WWN	499
raidcom get lun.....	500
raidcom modify lun.....	503
raidcom add path.....	506
raidcom get path.....	509
raidcom check_ext_storage path.....	516
raidcom delete path.....	519
raidcom disconnect path.....	521
raidcom modify path.....	524
raidcom delete pool.....	526
raidcom get pool.....	528
raidcom modify pool.....	549
raidcom monitor pool.....	557
raidcom reallocate pool.....	558
raidcom rename pool.....	559
raidcom initialize pool.....	560
raidcom modify system.....	561
raidcom get system.....	563
raidcom get port.....	572
raidcom modify port.....	596
raidcom add parity_grp	615
raidcom delete parity_grp.....	617
raidcom get parity_grp.....	618
raidcom initialize parity_grp	624
raidcom modify parity_grp.....	625
raidcom extend parity_grp.....	627
raidcom add rcu.....	628
raidcom delete rcu.....	630
raidcom get rcu.....	632
raidcom modify rcu.....	637
raidcom add rcu_iscsi_port.....	639
raidcom delete rcu_iscsi_port.....	641
raidcom get rcu_iscsi_port.....	642
raidcom add rcu_path.....	644
raidcom delete rcu_path.....	645
raidcom add ssid.....	647
raidcom delete ssid.....	648

raidcom get ssid.....	649
raidcom add resource.....	650
raidcom modify resource.....	657
raidcom delete resource.....	658
raidcom get resource.....	660
raidcom lock resource.....	663
raidcom unlock resource.....	664
raidcom map resource.....	665
raidcom unmap resource.....	667
raidcom add snap_pool.....	669
raidcom get snap_pool.....	673
raidcom add snapshot.....	674
raidcom map snapshot.....	676
raidcom unmap snapshot.....	677
raidcom delete snapshot.....	678
raidcom modify snapshot.....	680
raidcom get snapshot.....	688
raidcom replace snapshot.....	698
raidcom add spm_wwn.....	699
raidcom delete spm_wwn.....	700
raidcom modify spm_wwn.....	701
raidcom get spm_wwn.....	702
raidcom monitor spm_wwn.....	704
raidcom add spm_group.....	706
raidcom delete spm_group.....	708
raidcom modify spm_group.....	709
raidcom get spm_group.....	710
raidcom monitor spm_group.....	712
raidcom modify spm_ldev.....	713
raidcom delete spm_ldev.....	715
raidcom monitor spm_ldev.....	716
raidcom get spm_ldev.....	718
raidcom add dp_pool.....	720
raidcom get dp_pool.....	726
raidcom send ping.....	733
raidcom add server.....	734
raidcom delete server.....	736
raidcom modify server.....	737
raidcom get server.....	740
raidcom get apn.....	743
raidcom add qos_grp.....	744

raidcom delete qos_grp.....	745
raidcom modify qos_grp.....	746
raidcom get qos_grp.....	748
raidcom monitor resource.....	751
raidcom add nvm_subsystem.....	754
raidcom modify nvm_subsystem.....	755
raidcom delete nvm_subsystem.....	757
raidcom get nvm_subsystem.....	758
raidcom add nvm_subsystem_port.....	762
raidcom delete nvm_subsystem_port.....	762
raidcom get nvm_subsystem_port.....	763
raidcom add namespace_path.....	764
raidcom delete namespace_path.....	765
raidcom get namespace_path.....	766
raidcom add namespace.....	767
raidcom delete namespace.....	768
raidcom modify namespace.....	769
raidcom get namespace.....	770
raidcom get ctg.....	772

Appendix A: Valid parameter ranges for CCI options (VSP One Block)	774
--	-----

Appendix B: Valid parameter ranges for CCI options (VSP 5000 series)	776
--	-----

Appendix C: Valid parameter ranges for CCI options (VSP G130, G/F350, G/F370, G/F700, G/F900)	778
---	-----

Preface

This document describes and provides instructions for using Command Control Interface (CCI) software to configure and perform operations on Hitachi RAID storage systems.

Please read this document carefully to understand how to use these products, and maintain a copy for your reference.

Product version

This document revision applies to CCI software version 01-83-03/03 or later.

Changes in this revision

- Added information about specifying local and remote replication pairs ([paircreate*](#) (on page 23)).
- Updated the information about the swaps(p) option ([pairresync*](#) (on page 56)).
- Added support for 25-Gbps Ethernet connections ([raidcom add external_iscsi_name](#) (on page 364), [raidcom add rcu_iscsi_port](#) (on page 639), [raidcom get rcu_iscsi_port](#) (on page 642)).
- Added support for the Thin Image Advanced virtual clone (vClone) function ([raidcom get ldev](#) (on page 413), [Thin Image primary volume example](#) (on page 447), [Capacity saving examples](#) (on page 452), [raidcom modify snapshot](#) (on page 680), [raidcom get snapshot](#) (on page 688)).
- Updated the descriptions of capacity saving data displayed by the **raidcom get pool** command ([raidcom get pool](#) (on page 528)).
- Added display of the path blocked SIM monitoring ([raidcom get rcu](#) (on page 632), [raidcom modify rcu](#) (on page 637)).
- Added an example of displaying consistency group (CTG) information for remote replication ([raidcom get ctg](#) (on page 772)).

Release notes

Read the release notes before installing and using this product. They may contain requirements or restrictions that are not fully described in this document or updates or corrections to this document. Release notes are available on the Hitachi Vantara documentation website: <https://docs.hitachivantara.com>.

Accessing product documentation

Product user documentation is available on: <https://docs.hitachivantara.com>. Check this site for the most current documentation, including important updates that may have been made after the release of the product.

Getting help

The [Hitachi Vantara Support Website](#) is the destination for technical support of products and solutions sold by Hitachi Vantara. To contact technical support, log on to the Hitachi Vantara Support Website. For additional contact methods, go to <https://support.hitachivantara.com/en/contact-support.html>.

[Hitachi Vantara Community](#) is a global online community for Hitachi Vantara customers, partners, independent software vendors, employees, and prospects. It is the destination to get answers, discover insights, and make connections. **Join the conversation today!** Go to community.hitachivantara.com, register, and complete your profile.

Comments

Please send comments to doc.feedback@hitachivantara.com. Include the document title and number, including the revision level (for example, -07), and refer to specific sections and paragraphs whenever possible. All comments become the property of Hitachi Vantara LLC.

Thank you!

Chapter 1: Differences between CCI and Device Manager - Storage Navigator

This information describes the differences between CCI and Storage Navigator/Device Manager - Storage Navigator.

Supported characters

Supported characters for CCI

There are some differences between the supported characters for CCI and the supported characters for Device Manager - Storage Navigator. Because of this, you should always use only characters that are supported by both CCI and Device Manager - Storage Navigator.

The following figure shows the characters that can be used in CCI commands and their ASCII codes. Note that "\" can be used only on Windows, whereas "/" can be used only on UNIX. For CCI, do not use characters that are not supported by the operating system in which the command is executed.



Caution: When you execute a command to change the configuration, make sure you specify the parameters correctly. If you specify a parameter that is not contained in the syntax of the command, or you do not specify necessary parameters, the result might be different from the one you expected.

Low 4 bits	High 3 bits							
	0	1	2	3	4	5	6	7
0	NUL	DLE	SP	0	@	P	.	p
1	SOH	DC1	!	1	A	Q	a	q
2	STX	DC2	"	2	B	R	b	r
3	ETX	DC3	#	3	C	S	c	s
4	EOT	DC4	\$	4	D	T	d	t
5	ENQ	NAC	%	5	E	U	e	u
6	ACK	SYN	&	6	F	V	f	v
7	BEL	ETB	'	7	G	W	g	w
8	BS	CAN	(8	H	X	h	x
9	HT	EM)	9	I	Y	i	y
A	LF/NL	SUB	*	:	J	Z	j	z
B	VT	ESC	+	;	K	[k	{
C	FF	FS	,	<	L	\	l	
D	CR	GS	-	=	M]	m	}
E	SO	RS	.	>	N	^	n	~
F	SI	US	/	?	O	_	o	DEL

: usable

: unusable



Caution: Observe the following restrictions for using characters:

- A hyphen can be used in a name, but it cannot be used at the beginning of the name. You can neither use some character codes according to the OS on which you execute the command.
- A space (SP) can be used in a name, but if you specify a space for a user name, reference commands will fail, because a space is used as a delimiter and the commands cannot distinguish whether the space is for user name or delimiter.
- A space can only be used between other characters. If you use a space at the beginning or end of a name, the space is omitted automatically. Also, you cannot use a name consisting only of one or more space characters.

Examples of using space characters (SP) in a name:

- "group name" (different from "groupname")
- "g r o u p n a m e" (different from "groupname" and "group name")

Examples in which you cannot use a space:

- " name" (same as "name")
- "name " (same as "name")
- " " (invalid)
- The characters that can be used in Storage Advisor Embedded or VSP One Block Administrator are the same as the ones that can be used in CCI.

Supported characters for Device Manager - Storage Navigator

The following figure shows the characters that can be used in Device Manager - Storage Navigator commands and their ASCII codes.



Note: Characters that can be used vary depending on names. For details, see the System Administrator Guide of your system.

Low 4 bits	High 3 bits							
	0	1	2	3	4	5	6	7
0	NUL	DLE	SP	0	@	P	`	p
1	SOH	DC1	!	1	A	Q	a	q
2	STX	DC2	"	2	B	R	b	r
3	ETX	DC3	#	3	C	S	c	s
4	EOT	DC4	\$	4	D	T	d	t
5	ENQ	NAC	%	5	E	U	e	u
6	ACK	SYN	&	6	F	V	f	v
7	BEL	ETB	'	7	G	W	g	w
8	BS	CAN	(8	H	X	h	x
9	HT	EM)	9	I	Y	i	y
A	LF/NL	SUB	*	:	J	Z	j	z
B	VT	ESC	+	;	K	[k	{
C	FF	FS	,	<	L	\	l	
D	CR	GS	-	=	M]	m	}
E	SO	RS	.	>	N	^	n	~
F	SI	US	/	?	O	_	o	DEL

: usable

: unusable

When using CCI and Device Manager - Storage Navigator together, and setting the name from CCI, see the notes in the following table for characters that can be used in each name.

Name	Parameter to be set	Characters that can be used in CCI	Note when using Device Manager - Storage Navigator together with CCI
Host group name (For FC)	-host_grp_name	See figure in Supported	Do not use the following characters:*

Name	Parameter to be set	Characters that can be used in CCI	Note when using Device Manager - Storage Navigator together with CCI
		characters for CCI (on page 12)	\\ / : ,
LDEV nickname	-ldev_name	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\ / : ,
WWN nickname	-wwn_nickname	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\ / : ,
Pool name	-pool_name	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\ / : ,
Resource group name	-resource_name	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\ / : ,
SPM name	-spm_name	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\ / : ,
iSCSI name	-iscsi_nickname	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\ / : ,
CHAP user name	-target_chap_user	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\
	-initiator_chap_user	See figure in Supported characters for CCI (on page 12)	Do not use the following characters: * \\

Name	Parameter to be set	Characters that can be used in CCI	Note when using Device Manager - Storage Navigator together with CCI
* If you use these characters, you might not be able to use some Device Manager - Storage Navigator functions such as configuration reports.			

Maximum number of characters

There are some differences between the maximum number of characters for names in CCI commands and the maximum number of characters for names in Device Manager - Storage Navigator. Because of this, you should always try to use the number of characters that can be used in both CCI and Device Manager - Storage Navigator.

If you enter more than the maximum number of characters for a name in a CCI command, the name is truncated and only the allowable number of characters is used. For example, if the maximum number of characters is 32 and you enter 35 characters, only the first 32 characters are used.

The following table specifies the maximum number of characters for names in CCI and Device Manager - Storage Navigator. To ensure that names are usable in both CCI and Device Manager - Storage Navigator, observe the following requirements:

- WWN nickname: Do not use more than 64 characters.
- User ID: Do not use more than 63 characters.
- Password: Do not use more than 63 characters.

Name	CCI command	Device Manager - Storage Navigator
Host group name	64 characters	64 characters
Device group name	32 characters	-
Device name	32 characters	-
LDEV nickname	32 characters	32 characters
WWN nickname	64 characters	64 characters
Copy group name	32 characters	32 characters
Pool name	32 characters	32 characters
Resource group name	32 characters	32 characters
User name (User ID)	63 characters	256 characters

Name	CCI command	Device Manager - Storage Navigator
User name (Password)	63 characters	256 characters
Raw device file name	63 characters* * If you want to specify a name that consists of 64 characters or more for such as a UNIX device path, create a symbolic link to the name and make sure that the name that you specify consists of 63 characters or less.	-

Operational differences

The following table lists the operational differences between CCI and Device Manager - Storage Navigator.

Operation	CCI	Device Manager - Storage Navigator
Adding or changing path for an external volume (UVM)	The path definition is required for each external volume in the path group.	You can specify the path groups on the screen and execute the add/change path at one time.
Operations when the software product is not installed	You can change and delete existing resources for this software product, but you cannot add new resources.	You cannot change or delete existing resources or add new resources.
Displaying WWN	The WWN is displayed only when LUN security is enabled.	The WWN is displayed when LUN security is enabled or disabled.
Setting Pool ID	Optional	Required

Operation	CCI	Device Manager - Storage Navigator
Creating LDEV	You can specify the LDEV size in GB, LBA, or cylinders. When you specify GB, CCI can perform with or without size correction. If the capacity of LDEVs that are created by each GUI and CLI is the same, a copy pair might not be created. To create a pair with the LDEV that is created by GUI, create an LDEV by specifying LBA.	You can specify the LDEV size in GB, LBA, or cylinders. When you specify GB, size correction is performed.
Expanding the capacity of Dynamic Provisioning or Dynamic Provisioning for Mainframe virtual volume	You need to specify the capacity to be added to the volume.	You need to specify the total capacity after the volume is expanded.
Moving the CLPR assigned to the LUSE configuration volumes or the CLPR of parity groups containing LUSE configuration volumes	CLPRs can be moved, but it is not recommended.	CLPRs cannot be moved.
Moving the CLPR assigned to the journal volumes	CLPRs cannot be moved.	CLPRs can be moved if you specify all LDEVs in the journal.
Deleting SPM name	The SPM name of WWN is deleted, and the registration of the SPM name in the specified port is released.	The SPM name of WWN is deleted, but the SPM registration is maintained.
Deleting SPM group	WWN is deleted from the group, and the registration of SPM in the group of the specified port is released.	WWN is deleted from the group, but the SPM registration is maintained.
Moving the parity groups which configure the distributed parity group between the CLPRs	Parity groups cannot be moved.	All parity groups which are concatenated are moved.

Operations that cannot be used with Device Manager - Storage Navigator

You cannot use CCI together with Device Manager - Storage Navigator for some operations.

The following table shows operations that cannot be used with Device Manager - Storage Navigator.

Function	Operation
ShadowImage	Device Manager - Storage Navigator cannot split a pair created as a consistency group by CCI.
Universal Replicator	<ul style="list-style-type: none"> Device Manager - Storage Navigator cannot split, resynchronize, or delete a pair created by CCI. CCI cannot split, resynchronize, or delete a pair created by Device Manager - Storage Navigator.

Methods for specifying IP addresses

CCI supports the following methods for specifying IP addresses.

For IPv4

- Method 1: a.b.c.d

IPv4 dotted-decimal notation that separates an IPv4 (4-byte) IP address in units of one byte using dots. This is the most common method. CCI uses this format when it displays IPv4 IP addresses.

Example

Decimal representation of a class C IP address using method 1

192.168.1.1

- Method 2: a.b.c

This notation separates an IPv4 (4-byte) IP address into two sets of eight bits (one byte) and one set of 16 bits (two bytes) using dots.

Example

Decimal representation of a class C IP address using method 2

192.168.257 (192.168.1.1 when expressed using method 1)

- Method 3: a.b

This notation separates an IPv4 (4-byte) IP address into one set of eight bits (one byte) and one set of 24 bits (three bytes).

Example

Decimal representation of a class C IP address using method 3

192.11010305 (192.168.1.1 when expressed using method 1)

- Method 4: a

This notation represents an IPv4 (4-byte) address as a 4-byte value.

Example

Decimal representation of a class C IP address using method 4

3232235777 (192.168.1.1 when expressed using method 1)

When prefixed with "0", a, b, c, and d above are treated as octal numbers. When prefixed with "0x", a, b, c, and d above are treated as hexadecimal numbers.

Example

Hexadecimal representation of a class C IP address using method 1

0xc0.0xa8.0x01.0x01 (192.168.1.1 when expressed in decimal notation)

For IPv6

The specification methods defined in RFC 2373 can be used.

Displaying NAA identifiers (LUN WWNs) and namespace identifiers (NGU IDs)

Some CCI commands are used to display the NAA identifiers and namespace identifiers of volumes. However, the CCI commands cannot display the NAA identifiers and namespace identifiers if any of the following conditions applies:

NAA identifiers cannot be displayed if:

- Volumes communicate using a protocol other than SCSI.
- Volumes do not have LU paths.
- Virtual LDEV IDs are not assigned to volumes.
- The GAD reserve attribute is set for volumes.

Namespace identifiers cannot be displayed if:

- Volumes communicate using a protocol other than NVMe-oF or NVMe/TCP.
- Virtual LDEV IDs are not assigned to volumes.
- The GAD reserve attribute is set for volumes.

Specifying a mirror unit (MU) number by using the -d option

Certain commands specify the LDEV identification information and a mirror unit (MU) number by using the -d option to identify the target pair defined in the configuration definition file. If the pair identified by the specified LDEV identification information and the MU number is not defined in the configuration definition file, the command returns an error. If the MU number is not specified, the MU number is assumed to be zero (MU#0).



Note: The LDEV identification information cannot be omitted.

CCI interprets the specified MU number as the MU number of the target pair type on which the command operates so that the pair registered in the configuration definition file as the pair type is the operation target. For example, if the target pair type of the command is TrueCopy, the pair with the MU number specified in the configuration definition file is operated as a TrueCopy pair.

For details about a target pair type on which the command operates, see the `-I [H] [M] [instance#]` or `-I [TC] [SI] [instance#]` option for each command.

For more information about how to register a pair and the settings that depend on a pair type (ShadowImage, Copy-on-Write Snapshot, TrueCopy, Universal Replicator, or global-active device) in the configuration definition file, see the *Command Control Interface User and Reference Guide*.

Chapter 2: Data management commands

This information provides specifications for the CCI data replication and protection management commands.



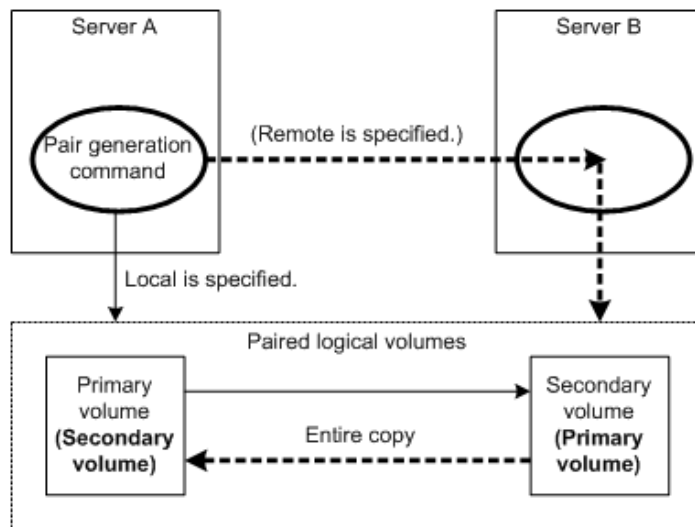
Note: For the following commands marked with an asterisk (*), if the `HORCC_CHECK_CHARACTER_CODE` environment variable is set when executing the command, the command checks the codes of the characters used for the command options and parameters. If the `HORCC_CHECK_CHARACTER_CODE` environment variable is set, when you specify a character code of 0x1F or less, or 0x7F or more as a command option or parameter, then the command stops the processing and responds `EX_INVARG`.

paircreate*



Warning: Use the `paircreate` command with caution. The command starts the initial copy operation, which overwrites all data on the secondary (target) volume of the pair. If the primary and secondary volumes are not identified correctly, or if the wrong options are specified (for example, `vl` instead of `vr`), data is copied to the wrong volume overwriting and the data in the target of transferring is overwritten.

The `paircreate` command is used to create a new volume pair from two unpaired volumes. The `paircreate` command can create either a paired logical volume or a group of paired volumes. The `paircreate` command allows you to specify the direction (local or remote) of the pair generation (see the following figure for pair creation). If local (`-vl` option) is specified, the server issuing the `paircreate` command has the primary volume. If remote (`-vr` option) is specified, the remote server has the primary volume. The `-split` option of the `paircreate` command (ShadowImage and Copy-on-Write Snapshot only) allows you to simultaneously create and split pairs using the `paircreate` command only. When `-split` is used, the pair status changes from `COPY` to `PSUS` (instead of `PAIR`) when the initial copy operation is complete.



Before issuing the `paircreate` command, make sure that the secondary volume is not mounted on any system. If the secondary volume is found to be mounted after `paircreate`, delete the pair (`pairsplit -S`), unmount the secondary volume, and then reissue the `paircreate` command.

When you use TrueCopy, TrueCopy for Mainframe, TrueCopy Async, TrueCopy Asynchronous for Mainframe, Universal Replicator, Universal Replicator for Mainframe, or global-active device, if a failure that requires maintenance operations, such as blocking the controller board or the cache memory, occurs on the primary storage system, the `paircreate` command cannot be executed.

When you use global-active device, if a failure that requires maintenance operations occurs on the secondary storage system, the `paircreate` command cannot be executed as well.

The `paircreate` command terminates before the initial copy operation is complete (except when the `nocopy` option is specified). Use the `pairevtwait` or `pairdisplay` command to verify that the initial copy operation completed successfully (status changes from COPY to PAIR, or from COPY to PSUS if the `-split` option was specified).

If the ROW column in the `pairdisplay` command displays "Y", a Thin Image Advanced pair can be created. To create a Thin Image Advanced pair, specify a DRS-VOL with an LU path configured as the P-VOL and the S-VOL, and then specify the pool ID of the P-VOL and the S-VOL for the `-pid` option.

A Thin Image pair created by using this command do not belong to a snapshot group. Perform the pair operations by using the `raidcom` command only when the created Thin Image pair cannot be deleted by using the `pairsplit -S` command.

When the Thin Image Advanced pair is created by using this command, the available range of CTG IDs 0 through 255. However, if the maximum of the CTG ID range supported by the storage system is less than 255, the maximum CTG ID supported by the storage system is the available maximum CTG ID used by this command.

If the Thin Image Advanced pair is created by using this command, an audit log of "+Command=Add Snapshot" is output. If the Thin Image Advanced pair cannot be created, an SSB1 error code starting with 2E might be output. For details about the actions to be taken if this SSB code is output, see the *Command Control Interface User and Reference Guide*. The SSB code description is not displayed in the command output.

The operations can be complicated in a cascade configuration in which the S-VOL of the Thin Image Advanced pair is used as the P-VOL of the Thin Image Advanced pair, therefore do not create this cascade configuration by using the **paircreate** command.

If a volume belonging to a virtual storage machine created by a user is specified for HORCM_CMD in the configuration definition file, this command cannot be run for the Thin Image Advanced pair. If this command is run, an EX_CMDRJE (SSB1=2EF3, SSB2=9300) error is displayed.

Specify the remote replication volume pairs by combining the following options. For details about the option, see [Options and parameters \(on page 27\)](#). There are no restrictions on the order in which the options are specified.

- To create a volume pair of Universal Replicator or Universal Replicator for Mainframe:
 - Specify `-I[H]` or `-I[TC]`.
It can be omitted if the HORCC_MRCF environment variable is not set in the command execution environment.
 - Specify `-jp -js -f async [| -startctg]`.
- To create a volume pair of TrueCopy or TrueCopy for Mainframe:
 - Specify `-I[H]` or `-I[TC]`.
It can be omitted if the HORCC_MRCF environment variable is not set in the command execution environment.
 - Specify `-f[g] {data|status|never} [| -startctg]`.
Do not specify `-jp` (or `-jq`), and `-js`.
- To create a volume pair of global-active device or HAM:
 - Specify `-I[H]` or `-I[TC]`.
It can be omitted if the HORCC_MRCF environment variable is not set in the command execution environment.
 - Specify `-jp (or -jq) -f[g] never [| -startctg]`.
Do not specify `-js`.

Specify the local replication volume pairs by combining the following options. For details about the option, see [Options and parameters \(on page 27\)](#). There are no restrictions on the order in which the options are specified.

- To create a volume pair of Volume Migration:
 - Specify `-I [M]` or `-I [SI]`.
It can be omitted if the `HORCC_MRCF` environment variable is set in the command execution environment.
 - Specify `-m cc`. Do not specify `-pid`.
- To create a volume pair of Thin Image (TI only) or Thin Image Advanced:
 - Specify `-I [M]` or `-I [SI]`.
It can be omitted if the `HORCC_MRCF` environment variable is set in the command execution environment.
 - Specify `-pid`. Do not specify `-m cc`.
- To create a volume pair of ShadowImage or ShadowImage for Mainframe:
 - Specify `-I [M]` or `-I [SI]`.
It can be omitted if the `HORCC_MRCF` environment variable is set in the command execution environment.
 - Do not specify `-m cc` and `-pid`.



Note: While the `paircreate -m grp` command is run, if the command is suspended because the host on which CCI is running is down or the storage has lost communication with CCI, a consistency group with no pairs registered might remain on the storage system. In that case, use the following procedure to delete the consistency group with no pairs registered, and then retry the command:

1. Identify a consistency group in which no pairs are registered.
 - (VSP One B20) Run the `raidcom get ctg` command to find a consistency group in which the STS is in the CHG status. The consistency group in which the STS is in the CHG status is the one in which no pairs are registered.
 - (All other models): On the **Consistency Group** tab in the **Local Replication** window in HDvM - SN, the consistency group with a status of Changing is the one in which no pairs are registered.
2. Specify the ID of the consistency group identified in step 1 for the `-m grp` option in the `paircreate` command, and then create a ShadowImage pair or a Thin Image pair. If the consistency group ID is 128 or more, create a Thin Image pair.
3. Delete the ShadowImage pair or the Thin Image pair created in step 2.



Note: When a consistency group is created with TC, GAD, and UR by using the CCI **paircreate** or **pairresync** command, the same CTG ID cannot be used between different software. For example, if a consistency group including a GAD pair whose CTG ID is 0 exists, a consistency group including a UR pair whose CTG ID is 0 cannot be created. Therefore, if a consistency group is created with multiple software products by using CCI, the maximum number of consistency groups supported by the storage system might not be created. The following example shows the consistency groups created with multiple software products by using CCI. Make sure to use Device Manager - Storage Navigator to create the maximum number of consistency groups supported by the storage system.

The following example lists the allowed maximum numbers and the range of assigned CTG IDs of consistency groups if they are created in 3DC configurations with GAD and UR for VSP G1x00 and VSP F1500.

- GAD: 85 (CTG ID 0 through CTG ID 84)
- UR: 85 (CTG ID 85 through CTG ID 169)
- UR delta resync: 85 (CTG ID 170 through CTG ID 254)

Syntax

```
paircreate { -h | -q | -z[x] | -I[H][M][instance#] or
-I[TC][SI][instance#] | -g <group> | -d <pair Vol> | -d[g] <raw_device>
[MU#] [-n] | -FHORC [MU#] or -FCA[MU#] | -d[g] <seq#> <ldev#>
[MU#] | -f[g] <fence> [<CTG ID> | -startctg <START CTG ID>] | -v | -c <size> | -
nocopy | -nomsg |
-split | [-m <mode>] | -jp <id> | -jq <id> | -js <id> | -pid <PID> |
-fq <mode> | -cto <o-time> [c-time] [r-time] | -pvol(svol)[ldevgrp] | -nocsus | -pr
<io preference> | -el <error level> }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **paircreate** command enter interactive mode. The **-zx** option monitors if the HORCM is operating in interactive mode. When this option detects a HORCM shutdown, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies a group name defined in the configuration definition file. The command is executed for the specified group unless the -d <pair Vol> option is specified.

-d <pair Vol>

Specifies paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volume.

-d[g] <raw_device> [MU#] [-n]

Searches whether the specified `raw_device` is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of `-g <group>` option. If the specified `raw_device` is contained in two or more groups, the command is executed for the first group.

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <ldev#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of `'-g <group>'` option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The `<seq#>` `<LDEV#>` values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-f[g] <fence> [<CTG ID> | -startctg <START CTG ID>]

TrueCopy/TrueCopy Async/Universal Replicator/High Availability Manager/global-active device only.

Specifies the fence level for assuring the consistency of paired volume data. A fence level of 'data', 'status', 'never', or 'async' must be specified. This option must always be specified. Fence level '-f async' can be specified only for TrueCopy Async/Universal Replicator. The '-fg' option is used to create a TrueCopy Sync CTG volume, and fence level must be specified as '-fg data', '-fg status', or '-fg never'.

Specifies the fence level of '-f never' for HAM.

Specifies the fence level of '-f never' or '-fg never' for GAD. The '-fg never' option is used to create a CTG volume.

'CTG ID' is assigned as follows:

- When no CTG ID is assigned to the other device in the specified volume group:
The CTG ID specified by the "CTGID" option is assigned.
If the "CTGID" option is omitted, a new CTG ID is assigned.
- When the CTG ID is already assigned to another device in the specified volume group:
The CTG ID which has been assigned to the other device in the volume group is assigned. Use the `raidcom get ctg` command to verify how the CTG is being used. For details about the storage system models that support the command and how to use the command, see the command summary section.

Note that the CTG ID specified by "CTGID" option is invalid.

If 'CTGID' is not specified (with '-f async' or '-fg' option) and the maximum number of consistency groups already exists, an EX_ENOCTG error is returned. Therefore, the 'CTG ID' option can forcibly assign an existing CTG ID to a volume group on the RAID storage systems only when no CTG ID has been assigned to the volume group. The CTGID option is ignored unless you specify the '-f async' or '-fg' option.

The "-startctg" option is valid only when "-fg" is specified. START CTG ID is a value indicating the search start position of CTG automatic allocation. Using "START CTG ID", an unassigned CTG ID to another device is searched from the value specified by the "START CTG ID" option to the maximum value of the consistency group in ascending order. If the unassigned CTG ID is found, the search ends and the CTG ID is assigned to CTG. If the unassigned CTG ID is not found, an EX_ENOCTG error is returned. Use the `raidcom get ctg` command to verify how the CTG is being used.

-vl or -vr ; -pvol [ldevgrp] or -svol [ldevgrp]

Specifies the direction in which the data is transferred. Make sure to specify this option.

If the `-vl` option is specified, the volume managed by the instance from which the command is issued becomes a P-VOL. If `-pvol [ldevgrp]` is specified, the device group specified in [ldevgrp] becomes the P-VOL. If the `-vr` option is specified, the volume managed by the remote instance of the instance from which the command is issued becomes the P-VOL. If `-svol [ldevgrp]` is specified, the device group specified in [ldevgrp] becomes the S-VOL.

-c <size>

TrueCopy/TrueCopy for Mainframe/TrueCopy Async/TrueCopy Asynchronous for Mainframe/ShadowImage/global-active device only.

Specifies the track size of extents (1 to 15) to be used for the copy operation. If you specify a large number, the time for copy operation will be shortened. When you want to copy in a short time by stopping the writing of data to the P-VOL, specify the maximum value 15. If this option is not specified, the default value (3) is used.

The relationship between the track size and the copy pace is shown as follows. If the copy pace is fast, the host I/O performance might be degraded. To reduce the effect on host I/O performance, lower the copy pace.

- TrueCopy/ Async/global-active device
 - When the track size is 1, the copy pace is slow.
 - When the track size is 2 or 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.
- ShadowImage
 - When the track size is 1 or 2, the copy pace is slow.
 - When the track size is 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.
- TrueCopy for Mainframe/TrueCopy Asynchronous for Mainframe
 - When the track size is 1 to 3, the number of the tracks copied simultaneously is 3 and the copy pace is slow.
 - When the track size is 4 or more, the number of the tracks copied simultaneously is 15 and the copy pace is fast.



Note: This option is disabled in Universal Replicator, Universal Replicator for Mainframe, or ShadowImage for Mainframe. To change the copy pace in Universal Replicator or Universal Replicator for Mainframe, use the `raidcom modify journal` command for JNL option changes, or Storage Navigator. ShadowImage for Mainframe works only at medium pace.

-nocopy

Creates paired volumes without the initial copy operation when the data consistency of simplex volumes is assured by the user.



Note:

This option cannot be specified for ShadowImage or ShadowImage for Mainframe.

-nomsg

Suppresses messages to be displayed when this command is executed. It is used to execute this command from a user program. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

-split

ShadowImage/Copy-on-Write Snapshot only.

Splits the paired volume after the initial copy operation is complete. This option will return after changed the state in P-VOL_PSUS & S-VOL_COPY immediately, and S-VOL state is changed to 'SVOL_SSUS' after all data is copied.

-m <mode>

Specifies the mode.

mode = noread (ShadowImage only): Specifies the noread mode for hiding the secondary volume. The secondary volume cannot be read when this mode option is specified. The secondary volume can be read when this mode option is omitted.

**Note:**

The primary volume becomes read-disabled during a reverse resync operation (restore option of **pairresync** command).

mode = cyl (TrueCopy/TrueCopy Async/Universal Replicator/High Availability Manager/global-active device only): Set this value when the difference of bitmap is managed by a cylinder unit.

mode = trk (TrueCopy/TrueCopy Async/Universal Replicator/High Availability Manager/global-active device only): Set this value when the difference of bitmap is managed by a track unit.

**Note:**

When this mode (cyl or trk) is not specified, the default value is used. About the default value, see the user guide of the storage system or software product.

**Note:**

When the storage system does not support the bitmap management by a unit of cylinder, the bitmap difference is managed by track even if you specify "cyl". For details about the bitmap management which is supported by the storage system, see the user guide of the storage system or software product.

**Note:**

When the connection destination storage system does not support the bitmap management by a unit of cylinder, the bitmap difference is managed by track even if you specify "cyl". For details about the bitmap management which is supported by the storage system, see the user guide of the storage system or software product.

mode = grp [CTG ID] (ShadowImage, Thin Image, or Copy-on-Write Snapshot only). Creates a ShadowImage pair or a Thin Image pair, and then register the pair in a consistency group. Like a TrueCopy Async or Universal Replicator consistency group, ShadowImage guarantees data consistency in a group at a single point in time when splitting a pair using the 'pairsplit -g <group>' command (except '-S' or '-E' option).

'CTG ID' is assigned as follows:

- When no CTG ID is assigned to other devices in the specified volume group:
The CTG ID specified by "CTG ID" option is assigned.
If the "CTG ID" option is omitted, new CTG ID is assigned.
- When the CTG ID is already assigned to another device in the specified volume group:
The CTG ID which has been assigned to the other device in the volume group is assigned. Use the `raidcom get ctg` command to verify how the CTG is being used. For details about the storage system models that support the command and how to use the command, see the command summary section.

Note that the CTG ID specified by "CTG ID" option is invalid.

If 'CTG ID' is not specified and the number of consistency groups reaches the maximum, an EX_ENOCTG error is returned. Therefore, the 'CTG ID' option can forcibly assign an existing CTG ID to a volume group (for example, 0 to 127 on 9900V) only when no CTG ID has been assigned to the volume group.

For detail, please refer to Restrictions on specified volumes with `-m grp` option.



Note:

This option cannot be specified with `-split` option in the same command.

`mode = cc` (Volume Migration only): Specifies Volume Migration.

The `-vl` option specifies 'local', and copies data from the local instance LU (P-VOL) to the remote instance LU (S-VOL). The original volume as the local instance LU is migrated from P-VOL to S-VOL, and the physical volume mapping between P-VOL and S-VOL is switched after copied.

The `-vr` option specifies 'remote', and copies data from the remote instance LU (P-VOL) to the local instance LU (S-VOL). The original volume as the remote instance LU is migrated from P-VOL to S-VOL, and the physical volume mapping between P-VOL and S-VOL is switched after copied. During maintenance work on the storage system (Device Manager - Storage Navigator or the maintenance utility is in modify mode), this operation cannot be completed.



Note:

This option cannot be specified with the `-split` option in the same command.

This option ignores the '`-c <size>`' option.

-jp <id> or -jq <id>

Universal Replicator, the HAM configuration or the GAD configuration only.

- For Universal Replicator:

You can use -jp <id> option when specifying a journal ID for P-VOL.

The -jp <id> option is valid when the fence level is set to 'ASYNC', and a journal ID is automatically bound to the CTG ID.

- For the HAM configuration or the GAD configuration:

You can create a HAM/GAD pair by using either -jp <id> or -jq <id> option.

Specify the quorum ID with -f fence (never) option and quorum ID when creating a HAM pair or a GAD pair. You do not need to check the resource group of the quorum volume.

The following conditions must be met for the HAM configuration.

- It is not a consistency group.
- The fence level is set to 'Never'.
- The RCU path is set to 'CU Free'.

The following conditions must be met for the GAD configuration.

- The fence level is set to 'Never'.
- The RCU path is set to 'CU Free'.

-js <id>

Universal Replicator only.

This option is used when specifying a journal ID for S-VOL.

Both the -jp <id> and -js <id> options are valid when the fence level is set to 'ASYNC', and each journal ID is automatically bound to the CTG ID.

-pid <PID>

Copy-on-Write Snapshot and Thin Image only.

This option is used to specify the ID of the pool in which snapshot data to be created is stored.

If a V-VOL in Copy-on-Write Snapshot or Thin Image (TI only) is specified as an S-VOL, and a pool for Thin Image (TI only) or Dynamic Provisioning is specified with this option, a Thin Image (TI only) pair is created. If this option is omitted, the default ID of the pool containing the snapshot data is set to zero.

If a V-VOL in Copy-on-Write Snapshot or Thin Image (TI only) is specified as an S-VOL, and a pool for Copy-on-Write Snapshot is specified with this option, a Copy-on-Write Snapshot pair is created. If this option is omitted, the default ID of the pool containing the snapshot data is set to zero.

If DRS-VOLs are specified as a P-VOL and an S-VOL, the ROW row in the **pairdisplay** command displays Y, and a pool for Dynamic Provisioning is specified with this option, a Thin Image Advanced pair is created. If this option is omitted, a ShadowImage pair or a Volume Migration pair is created.

If a volume with the DP attribute is used as the S-VOL with this option, a Thin Image (TI only) pair cannot be created. Use the **raidcom add snapshot** command.

-fq <mode>

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

ShadowImage only.

This option is used when specifying the mode whether `-split` is performed or not as 'QUICK'.

- **mode = normal:** The **paircreate -split** command is performed as non-quick mode regardless of setting of \$HORCC_SPLT environment variable or the system option mode 122.
- **mode = quick:** The **paircreate -split** command is performed as Quick Split regardless of setting of \$HORCC_SPLT environment variable or the system option mode 122.

If this option is not specified, then the performing of the 'Split' depends on \$HORCC_SPLT environment variable or the system option mode setting through the SVP, whether the paircreate operation is Quick Split or not.

This `-fq` option is also valid for TrueCopy-TrueCopy/ShadowImage cascading operations using the `-FBC [MU#]` option.

The relationship between the `-fq` option and \$HORCC_SPLT is as shown below:

-fq option	\$HORCC_SPLT	Behavior
quick	Invalid	Quick Split
normal	Invalid	Normal Split
Omitted	QUICK	Quick Split
Omitted	NORMAL	Normal Split

-fq option	\$HORCC_SPLT	Behavior
Omitted	Omitted	Depends on system option mode (SOM)122. If SOM 122 is disabled, Normal Split is performed, whereas if SOM 122 is enabled, Quick Split is performed. This option is disabled by default.

**Note:**

- This `-fq` option is ignored to maintain the compatibility on 9900V so that you can add this option to the same script.

-FHORC [MU#] or -FCA [MU#]

This option is used to create the cascading configuration with `-g <group>` and `-gs <group>` options from the local node (takeover node).

`-g <group>` is used when specifying the cascaded P-VOL, and `-gs <group>` option is used when specifying the cascaded S-VOL. This operation ignores the `-vl` or `vr` option, because S-VOL is specified with `-gs <group>` option.

`-gs <group>`: This 's' option is used when specifying a group name for cascading S-VOL (defined in the configuration definition file). The command is executed for the specified group unless the `-ds <pair Vol>` option shown below is specified.

`-ds <pair Vol>`: The following options can be specified for cascading S-VOL.

`-d[g]s <raw_device> [MU#] ...`

`-d[g]s <seq#> <LDEV#> [MU#]`

-cto <o-time> [c-time] [r-time]

TrueCopy Async and Universal Replicator only.

If you specify `-cto <o-time> [c-time] [r-time]` option on TrueCopy Sync, it is ignored.

`o-time`:

This option is used when setting offloading timer for controlling inflow of write I/Os to the specified consistency group.

For TrueCopy Async, an integer from 1 through 255 (seconds) must be specified. If this option is not specified, 90 seconds is set by default.

For Universal Replicator, an integer from 1 through 255 (seconds) must be specified. If this option is not specified, the value set for the journal on the P-VOL side is applied. Use the `raidcom modify journal` command to change the value range to 256 to 600.

If `o-time=0` is specified, the inflow of write I/O becomes invalid.

When the sidefile capacity exceeds the limit of the sidefile area, write I/Os from the host wait, within the specified timeout period, until the space which is large enough to store next new data becomes available. As the timeout period, you can specify a value from 1 through 255 (seconds) for TrueCopy Async, 1 through 600 (seconds) for Universal Replicator. The default timeout values are 90 seconds for TrueCopy Async, and 60 seconds for Universal Replicator. If the timeout happens occurs during this waiting state, then the pair status changes from PAIR to PSUS of sidefile (Journal) Full, and its host side Write I/Os continue and data in waiting state is managed by BITMAP mode.

Therefore the o-time timeout value must be less than the I/O timeout value of the host system

[c-time](TrueCopy Async only): This option is used when setting Copy Pending timer to the specified consistency group. Specify an integer from 1 to 15 (minutes) for c-time. If this option is not specified, then this value is set as follows:

- If a consistency group is created, 5 (minutes) is set by default.
- If a consistency group is not created, the setting is not changed.

[r-time] (TrueCopy Async only): This option is used when specifying RCU Ready timer for the specified consistency group. r-time can be set from 1 to 10 minutes. If this option is not specified, this value is set as follows:

- If a consistency group is created, 5 (minutes) is set by default.
- If a consistency group is not created, the setting is not changed.



Note:

For TrueCopy Async, settings changed by these options are invalid if a consistency group already exists. These parameters are also forwarded to the S-VOL side with the **paircreate** command, and are used when S-VOL is changed to P-VOL. These parameters are maintained and become valid until and when the pair-volumes are changed to SMPL.

For Universal Replicator, these parameters can be set and changed when a pair is created in a journal volume in the P(S)JSN status. The parameters are set for the journal on the P-VOL side when the pair is created. To set parameters for the journals on both P-VOL and S-VOL sides, specify as follows:

1. `paircreate -g <group> -vr -f async -nocopy -jp <id> -js <id> -cto <o-time>`
2. `pairsplit -g <group> -S`
3. `paircreate -g <group> -vl -f async -jp <id> -js <id> -cto <o-time>`

These parameters are maintained on each journal. Therefore, if you set the value for offloading timer, execute the **raidcom modify journal** command on both P-VOL and S-VOL sides.

-nocsus

Universal Replicator only.

This option is used to create the suspended journal volumes without copying data in order to make the delta-resync between DC2 (Sync-S-VOL) and DC3 (Universal Replicator-S-VOL).

-pr <io preference>

This option is only available for global-active device.

Sets the I/O preference mode (I/O preference mode when remote path failed) when a failure occurs in the remote path between the storage systems of the primary site and the secondary site and communication becomes impossible. For details of the I/O preference mode when remote path failed, see the Global-Active Device User Guide. If this option is specified for any pair that is not a global-active device pair, the command returns an error.

The following value can be specified for `<io preference>`:

`pvol`: Primary Volume. When remote path failed, the primary volume continues host I/O preferentially.

If this option is omitted, the I/O Preference Mode When Remote Path Failed is disabled.

-el <error level>

This option is only available for Universal Replicator and specified only for VSP One B20.

Specify this option with the `-f async` option at the same time.

This option specifies an error level in Universal Replicator. If specifying this option is omitted, the error level is group. If this option is specified for any pair that is not a Universal Replicator pair, the command returns an error.

The following value can be specified as `-el <error level>`:

`volume`: The error level is volume.

This option cannot be used with the `-FHORC` option at the same time.

Returned values

The `paircreate` command sets either of the following returned values in `exit ()`, which allows users to check the execution results using a user program.

- **Normal termination:**

- **0**: When creating groups, 0 = normal termination for all pairs.

- **Abnormal termination:**

- **other than 0**: Refer to the error code for error details.

Error codes

Unrecoverable errors are not resolved, even after re-executing the command. If the command fails, the detailed status is logged in the CCI command log (`$HORCC_LOG`).



Note: When an option that is valid for only specific volumes is specified for other volumes, the error code might be EX_UNWOPT or EX_UNWCMD. Check whether the specified option is valid for the volume.

Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	Confirm status using the pairdisplay command. Make sure all volumes in the group have the same fence level and volume attributes.	236
	EX_INCSTG	Inconsistent status in group	Confirm pair status using pairdisplay .	229
	EX_INVVOL	Invalid volume status	Confirm volume status/attribute using pairdisplay (-l option) or 'raidvchkdsp -v aou'.	222
	EX_INVSTP	Invalid pair status	Confirm pair status using pairdisplay .	228
	EX_ENQSIZ	Unmatched volume size for pairing	Confirm volume size or number of LUSE volume using raidscan -f, and make sure volume sizes are identical.	212
Resource (Unrecoverable)	EX_ENOC TG	Not enough consistency groups in the RAID	Choose an existing CTG ID (pairvolchk displays CTG IDs). Use '-f async <CTG ID>' or '-m grp <CTG ID>' option of paircreate to force the pair into a pre-existing CTG ID.	217
	EX_ENXC TG	No consistency groups left for OPEN Vol use.	Confirm whether all consistency groups are already used by TrueCopy/TrueCopy Async/GAD or ShadowImage.	215
	EX_ENOPOL	Not enough Pool in RAID	Unable to retain the pool for executing a command because the threshold rate has been exceeded. Delete unnecessary/earlier generations paired volume, or re-synchronize unnecessary/earlier generations split volume.	206

Restrictions on specified volumes with `-m grp` option

- Volume group definition
 - Volume groups specified with `-m grp` cannot be defined across the storage systems.
 - When multiple groups of CCI are contained within the same consistency group (CTG ID), pairs with the specified group are operated for the entire consistency group.
 - When ShadowImage, Thin Image, or Copy-on-Write Snapshot volumes are cascaded by TrueCopy, TrueCopy Async, Universal Replicator, or GAD volumes, the data consistency is not guaranteed by the `pairsplit` command (including `pairsplit -FMRCF`) for which I/Os are being processed.

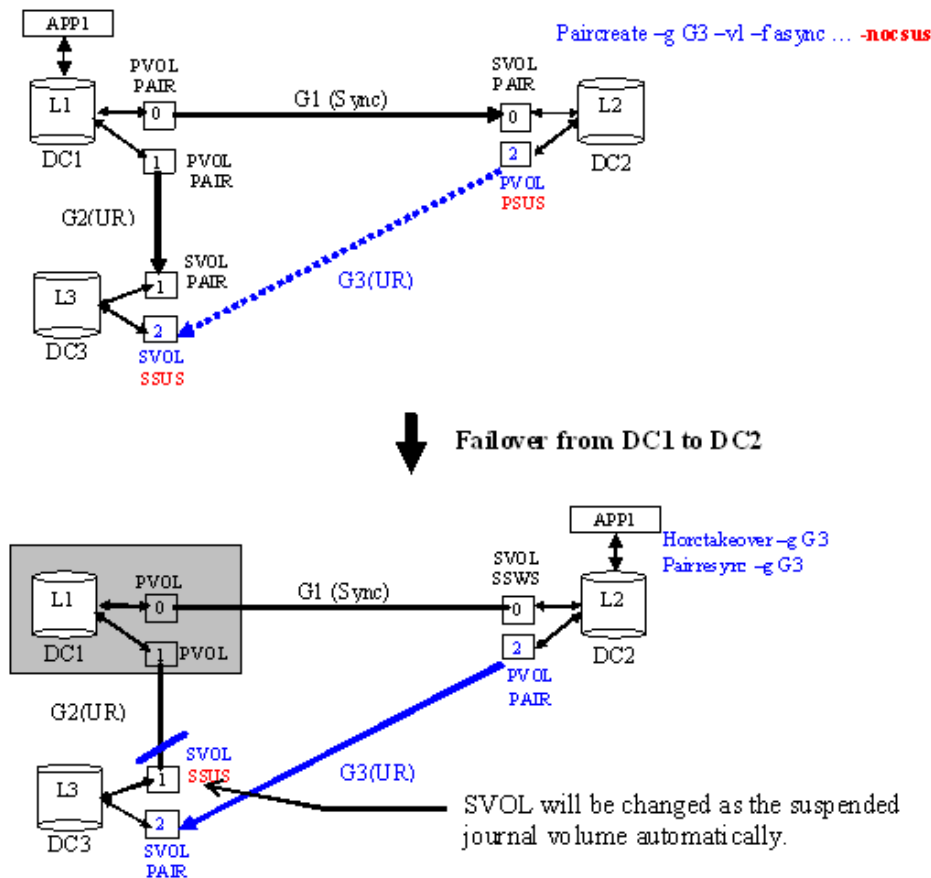
- Registration and limitations of the number of CTG IDs

When you create a pair, CCI maps and assigns the configuration definition file group to a CTG ID managed by the storage system. The maximum number of consistency groups that can be registered to a storage system is shown below. If you register over the maximum number of consistency groups, pair creation terminates with EX_ENOCTG error.

- VSP G1x00 and VSP F1500: 256 (CTG ID 0 - CTG ID 255)
- VSP G800, VSP F800, 9900V: 128 (CTG ID 0 - CTG ID 127)
- VSP G400, VSP G600, VSP F400, VSP F600: 64 (CTG ID 0 - CTG ID 63)
- VSP G200: 16 (CTG ID 0 - CTG ID 15)
- VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900 (ShadowImage): 128 (CTG ID 0 to CTG ID 127)
- VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900 (Thin Image): 2048 (CTG ID 0 to CTG ID 2047)
- VSP G/F350, G/F370, G/F700, VSP E590, VSP E790: (TrueCopy/ Universal Replicator/global-active device): 128 (CTG ID 0 to CTG ID 127)
- VSP G/F900, VSP E990, VSP E1090: (TrueCopy/ Universal Replicator/global-active device): 256 (CTG ID 0 to CTG ID 255)
- VSP 5000 series: (except global-active device): 256 (CTG ID 0 - CTG ID 255)
- VSP 5000 series: (global-active device): 1024 (CTG ID 0 - CTG ID 1023)

Example 1


The following figure shows an example of creating a cascading configuration with `-g <group>` and `-gs <group>` option from the local node (takeover node).



pairsplit*

The **pairsplit** command is used to split or delete volume pairs. This command stops updates to the secondary volume of a pair and can either maintain (status = PSUS) or delete (status = SMPL) the pairing status of the volumes. The **pairsplit** command can be applied to a paired logical volume or a group of paired volumes. The **pairsplit** command allows read access or read/write access to the secondary volume, depending on the selected options. When the **pairsplit** command is specified, acceptance of write requests to the primary volume depends on the fence level of the pair (data, status, never, or async). For Volume Migration, only the -S option is accepted.

The following figure illustrates the pair splitting process.



The primary volume's server is automatically detected by the **pairsplit** command, so the server does not need to be specified in the **pairsplit** command parameters. If the **-s** option (simplex) is used, the volume pair is deleted, the volumes are returned to the simplex state, and the primary and secondary volume status is lost. Paired volumes are split as soon as the **pairsplit** command is issued. If you want to synchronize the volumes, the **pairsplit** command must be issued after write I/Os to the paired volume have completed.

You can create and split ShadowImage pairs simultaneously using the **-split** option of the **paircreate** command.

You can delete pairs by using the **-s** option of the **pairsplit** command. When the **pairsplit -s** command is issued, the specified pair is deleted, and each volume is changed to SMPL (simplex) mode. If you want to re-establish a pair that has been deleted, you must use the **paircreate** command (not **pairresync**).

For Volume Migration, options other than **-S** are rejected.

Note on Quick Split:

If '\$HORCC_SPLT=QUICK' environment variable is set for any of the following storage systems, the **pairsplit** and **paircreate -split** operations are performed as Quick Split regardless of the system option mode 122 setting:

- VSP 5000 series
- VSP One B20
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

The \$HORCC_SPLT=QUICK environment variable is ignored for 9900V.

The Thin Image Advanced pair operations are available if both of the following conditions are met:

- Both the P-VOL and the S-VOL are DRS-VOLs with LU paths configured.
- The ROW column in the **pairdisplay** command displays "Y".

If the Thin Image Advanced pair operations are performed by using this command, an audit log of "+Command=Modify Snapshot" or "+Command=Delete Snapshot" is output. If the Thin Image Advanced pair cannot be created, an SSB1 error code starting with 2E might be output. For details about the actions to be taken if this SSB code is output, see the *Command Control Interface User and Reference Guide*. The SSB code description is not displayed in the command output.

If a volume belonging to a virtual storage machine created by a user is specified for HORCM_CMD in the configuration definition file, this command cannot be run for the Thin Image Advanced pair. If this command is run, an EX_CMDRJE (SSB1=2EF3, SSB2=9300) error is displayed.

Syntax

```
pairsplit {-h | -q | -z[x] | -I[H][M][instance#] or -I[TC][SI]
[instance#] | -g <group> | -d <pair Vol> | -d[g] <raw_device>
[MU#][~n] | -FHORC [MU#] | -FMRCF [MU#] | -d[g] <seq#> <LDEV#> [MU#]
| -r | -rw | -S | -SF[V] | -R[S|B] | -RF[V] | -P | -l |
-t <timeout> | -nomsg | -C <size> | -E | -fq <mode>} | -iomd <mode>
```

Options and parameters

Only one **pairsplit** option (-r, -rw, -S, -R, -P, -C, or -E) can be specified unless stated in the description of each option. If more than one option is specified, only the last specified option will be executed, unless stated in the description of each option.

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits this command.

-z or -zx

Makes the **pairsplit** command enter the interactive mode. The -zx option guards performing of the HORCM in the interactive mode. When this option is specified and a HORCM shutdown is detected, interactive mode terminates.

OpenVMS cannot use the -zx option

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies a group name defined in the configuration definition file. The command is executed for the specified group unless the -d <pair Vol> option is specified.

If the HORCC_CHECK_PAIRSPLIT_GOPTION environment variable is set, when you specify more than one parameter in the -g option, then the command suspends the processing and respond with EX_INVARG.

-d <pair Vol>

Specifies the paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volumes. This option is valid when the -g <group> option is specified.

-d[g] <raw_device> [MU#] [-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-r or -rw

TrueCopy, TrueCopy Async, and Universal Replicator only.

Specifies a mode of access to the S-VOL after paired volumes are split. If this option is omitted, -r is set. The -r option enables read-only access to the S-VOL. For mainframe volumes, both read and write are disabled.

The -rw option enables read and write access to the S-VOL.

-S

Specify this option when deleting a pair to change the P-VOL and S-VOL back to the simplex status.

**Note:**

- Either volume of a pair might be unpaired due to a failure. To recover from this failure, delete the pair by using the **pairsplit -s** command to cancel the P-VOL and S-VOL relation, and then create the pair again by using the **paircreate** command.
- Due to a path failure, the S-VOL might not be unpaired even if you execute the **pairsplit -s** command. If the S-VOL cannot be unpaired, CCI unpairs the P-VOL.

-SF[V]

Specify the **-SF[V]** option only when you must delete a TrueCopy, a Universal Replicator, or a GAD pair forcibly.

- TrueCopy pair

This option is specified only for VSP One B20.

This option changes the pair status of a TrueCopy pair to the unpaired status forcibly. Specifying **[V]** is ignored. Use this option only if the P-VOL status cannot be restored to Simplex by using the **-s** option.

- Universal Replicator pair

This option is specified only for VSP One B20.

This option changes only the P-VOL status of a Universal Replicator pair to the unpaired status forcibly. All CTG pairs including the operation targets are deleted regardless of whether **-d** is specified. Specifying **[V]** is ignored. Use this option only if the P-VOL status cannot be restored to Simplex by using the **-s** option.

- global-active device pair

This option changes only the P-VOL status of a GAD pair to the unpaired status forcibly, but does not change the S-VOL pair status.

- **-SFV** option: Specify this option to access the P-VOL from the host after deleting a GAD pair. This option leaves the virtual LDEV ID of the P-VOL.
- **-SF** option: Specify this option if you do not access the P-VOL from the host after deleting a GAD pair. This option deletes the virtual LDEV ID of the P-VOL.

**Caution:**

Use the **-SF[V]** option only if you cannot change the P-VOL status to Simplex by using the **-s** option. Specify this option only when the I/O mode of both P-VOL and S-VOL is blocked. Contact technical support if you want to split GAD pair forcibly when the I/O mode is other than blocked. To confirm the I/O mode of the P-VOL and S-VOL, check the R/W row output with **pairdisplay -fe** command. B/B in the R/W row indicates the I/O mode of the volume is blocked.

**Note:**

The data of P-VOL and S-VOL becomes inconsistent and it might cause errors if the volume status is changed to Simplex by using `-SFV` or `-RFV` option when the server can access both P-VOL and S-VOL. Follow the steps below when you delete GAD pair forcibly.

1. Stop the access to either the P-VOL or the S-VOL from the server.
2. Change the volume to which you stopped access to Simplex status. For example, if you stopped the access to P-VOL, change the P-VOL status to Simplex by specifying the `-SF` option. If you stopped the access to S-VOL, change the S-VOL status to Simplex by specifying the `-RF` option. When the volume status is changed to Simplex by specifying the `-SF` or `-RF` option, the virtual LDEV ID is deleted and the GAD reserve attribute is set. The volume becomes inaccessible from the server when the reserve attribute is set.
3. Change the volume to which you did not stop access to Simplex status. For example, if you did not stop access to the P-VOL, change the P-VOL status to Simplex by specifying the `-SFV` option. If you did not stop access to the S-VOL, change the S-VOL status to Simplex by specifying the `-RFV` option.

**Note:**

In some cases (incorrect configuration definition file or other reasons), the paired volumes that are the operation target of the `pairsplit` command might both become P-VOLs. If the `pairsplit -SF[V]` command is issued under these conditions, the CCI instance that executes the command directs the storage system to delete the TrueCopy, the Universal Replicator, or the GAD pair of only the volume managed by that CCI instance.

-R[S|B]

TrueCopy, TrueCopy Async, Universal Replicator, and GAD only.

This option issues a command from the host at the secondary site if a failure or downtime occurs on the host accessing the P-VOL.

The `-R` option changes the S-VOL to the unpaired status (simplex volume).

- For TrueCopy, TrueCopy Async, and Universal Replicator:

The `-R` option changes the S-VOL status, but it does not change the P-VOL pair status. However, if communication between MCU and RCU is established, the status of the P-VOL is PSUS, and the S-VOL is SMPL.

- For global-active device:

The `-R` option deletes a GAD pair, and it changes the pair status of both the S-VOL and the P-VOL to SMPL. If you use this option to delete a GAD pair, the virtual LDEV ID of the S-VOL remains, but the virtual LDEV ID of the P-VOL is deleted.

This option is available only when the pair status of the S-VOL is SSWS, and the I/O mode is local (L/L).

When a temporary failure such as a link failure occurs and you switch the control from the primary site host to the secondary site host, if you specify the `-RS` option, you can change the S-VOL pair status to SSWS.

- `-R`: Unpairs the S-VOL.
- `-RS`: Changes the S-VOL status to SSWS.
- `-RB`: Changes the S-VOL status from SSWS back to PSUS (PSUE) (SSUS).

Use the `-RB` option when the S-VOL status is SSWS and you want to resynchronize the pair from the P-VOL to the S-VOL. This option cannot be used for GAD.

-RF[V]

Specify the `-RF[V]` option only when you must delete a TrueCopy, a Universal Replicator, or a GAD pair forcibly. This option unpairs only the S-VOL of the GAD pair. It does not change the P-VOL pair status.

- TrueCopy pair

This option is specified only for VSP One B20.

This option changes only the S-VOL status of a TrueCopy pair to the unpaired status forcibly. Specifying `[V]` is ignored. Use this option if the S-VOL status cannot be restored to Simplex by using the `-R` option.

- Universal Replicator pair

This option is specified only for VSP One B20 and VSP One SDS Block.

This option changes only the S-VOL status of a Universal Replicator pair to the unpaired status forcibly. All CTG pairs including the operation targets are deleted regardless of whether the `-d` option is specified. Specifying `[V]` is ignored. Use this option if the S-VOL status cannot be restored to Simplex by using the `-R` option.

- global-active device pair

This option changes only the S-VOL status of a GAD pair to the unpaired status forcibly, but does not change the P-VOL pair status.

- `-RFV`: Specify this option to access the S-VOL from the host after deleting a GAD pair. This option leaves the virtual LDEV ID of the S-VOL.
- `-RF`: Specify this option if you do not access the S-VOL from the host after deleting a GAD pair. This option deletes the virtual LDEV ID of the S-VOL.



Caution: Use the option only if you cannot change the S-VOL status to Simplex by using the `-R` option. Specify this option only when the I/O mode of both P-VOL and S-Vol is blocked. Contact technical support if you want to split GAD pair forcibly when the I/O mode is other than blocked. To confirm the I/O mode of the P-VOL and S-Vol, check the R/W row output with `pairdisplay -fe` command. B/B in the R/W row indicates the I/O mode of the volume is blocked.



Note: The data of P-VOL and S-VOL becomes inconsistent and it might cause errors if the volume status is changed to Simplex by using `-SFV` or `-RFV` option when the server can access both P-VOL and S-VOL. Follow the steps below when you delete GAD pair forcibly.

1. Stop the access to either the P-VOL or the S-VOL from the server.
2. Change the volume to which you stopped access to Simplex status. For example, if you stopped the access to P-VOL, change the P-VOL status to Simplex by specifying the `-SF` option. If you stopped the access to S-VOL, change the S-VOL status to Simplex by specifying the `-RF` option. When the volume status is changed to Simplex by specifying the `-SF` or `-RF` option, the virtual LDEV ID is deleted and the GAD reserve attribute is set. The volume becomes inaccessible from the server when the reserve attribute is set.
3. Change the volume to which you did not stop access to Simplex status. For example, if you did not stop access to the P-VOL, change the P-VOL status to Simplex by specifying the `-SFV` option. If you did not stop access to the S-VOL, change the S-VOL status to Simplex by specifying the `-RFV` option.



Note: In some cases (incorrect configuration definition file or other reasons), the paired volumes that are the operation target of the **pairsplit** command might both become S-VOLs. If the **pairsplit -RF[V]** command is issued under these conditions, the CCI instance that executes the command directs the storage system to delete the TrueCopy, the Universal Replicator, or the GAD pair of only the volume managed by that CCI instance.

-P

TrueCopy, TrueCopy Async, and Universal Replicator only.

For TrueCopy Sync, this option is used to bring the primary volume forcibly into write disabled mode like PSUE with 'fence=data'. It is issued by the secondary host to disable P-VOL data changes by the host possessing the primary volume.

For TrueCopy Async and Universal Replicator, this option is used to suspend and purge data remaining in the sidefile or journal without updating S-VOL like a link failure (PSUE). This option stops journal operations forcibly when the journal utilization traffic becomes high. This is the same as the failure that S-VOL data is not updated, but it allows write access if you specify `-rw -P`. In that situation, if you are using the S-VOL as a file system (that is, UFS, NTFS, HANFS), then an FSCK(CHKDSK) is necessary before the volume is mounted even after the P-VOL is unmounted.

-I

When the remote host cannot be used due to host down, this option enables a **pairsplit** operation by a local host only. Except the `-R` option, the target volume of a local host must be P-VOL. The pair split operation in ShadowImage, Copy-on-Write Snapshot, or Thin Image pairs are available by specifying the S-VOL.

-t <timeout>

This option is for Universal Replicator. Specifies the maximum time to wait in seconds (1 to 21474836) for splitting or deleting a volume pair. If timeout occurs, the pairsplit command fails with EX_EWSTOT. To avoid the timeout, specify the time required to split or delete a volume pair for this value. If this option is omitted, the value is set to the default value of 7,200 seconds (2 hours).

-nomsg

Suppresses messages to be displayed when this command is executed. It is used to execute a command from a user program. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

-C <size>

ShadowImage only. Specifies the track size for copying in the range of 1 to 15.

Copies differential data retained in the primary volume into the secondary volume, then enables reading and writing from/to the secondary volume after completion of the copying (default). If not specified, the value used for **paircreate** or **pairresync** command is used. In ShadowImage, specify 1 or 2 for slow copy pace, specify 3 for medium copy pace, and specify 4 for fast copy pace.

-E

ShadowImage only.

Specify this option when suspending a paired volume forcibly. Not normally used.

-FHORC [MU#] or -FCA [MU#]

Forcibly specifies a cascading remote copy volume for specified volume pair in a local copy environment. If the **-l** option is specified, a cascading remote copy volume is split on a local host (campus distance). If the **-l** option is not specified, a cascading remote copy volume is split on a remote host (metro distance). The target HORC volume must be a P-VOL, or '**-R[S][B]**' option can be specified for S-VOL.

-FMRCF [MU#] or -FBC [MU#]

Forcibly specifies a cascading local copy volume for specified volume pair in a remote copy environment. If the **-l** option is specified, a cascading local copy volume is split on a local host (near site). If **-l** option is not specified, a cascading local copy volume is split on a remote host (far site). The target local copy volume must be a P-VOL, and the **-E** option cannot be specified.

-fq <mode>

ShadowImage only.

This option is used to specify the mode whether **pairsplit** is performed as 'QUICK'.

mode = normal: **pairsplit** is performed as Non quick mode regardless of setting of \$HORCC_SPLT environment variable and/or the system option mode 122 via SVP.

mode = quick: **pairsplit** is performed as Quick Split regardless of setting of \$HORCC_SPLT environment variable and/or the system option mode 122 via SVP.

If this option is not specified, then performing 'Quick Split' depends on \$HORCC_SPLT environment variable and/or the system option mode setting through the SVP.

The `-fq` option is also valid for TrueCopy-TrueCopy/ShadowImage cascading operation using '`-FBC [MU#]`' option.

The relationship between '`-fq`' option and `$HORCC_SPLT` is as shown below.

-fq option	\$HORCC_SPLT	Behavior
quick	Invalid	Quick Split
normal	Invalid	Normal Split
Omitted	QUICK	Quick Split
Omitted	NORMAL	Normal Split
Omitted	Omitted	Depends on system option mode (SOM)122. If SOM 122 is disabled, Normal Split is performed, whereas if SOM 122 is enabled, Quick Split is performed. This option is disabled by default.



Note:

This `-fq` option is ignored for 9900V to maintain compatibility on 9900V so that you can add this option to the same script.

-iomd <mode>

This option is supported by global-active device only.

Specify this option to change the I/O mode of a pair volume.

You can specify the following values for <mode>.

- local: Change the I/O mode to local.
- block: Change the I/O mode to block.

The pair volume whose I/O mode is to be changed is the pair volume described in the configuration definition file of the instance used to issue the command, among the pair volumes of the specified pair.

In the following example, change the I/O mode of the pair volume described in the configuration definition file of "instance: 100" to Local, among the pairs in "device group: mygroup".

Example:

```
pairsplit -g mygroup -iomd local -IH100
```

Returned values

The **pairsplit** command sets the following returned values during exit allowing you to check the execution results.

- **Normal termination:**

- **0:** When splitting groups, 0 = normal termination for all pairs.

- **Abnormal termination:**

- **other than 0:** Refer to the error codes for error details.

Error codes

The following table lists and describes the error codes for the **pairsplit** command. Unrecoverable errors are not resolved, even after re-executing the command. If the command failed, the detailed status is logged in the CCI command log (\$HORCC_LOG).



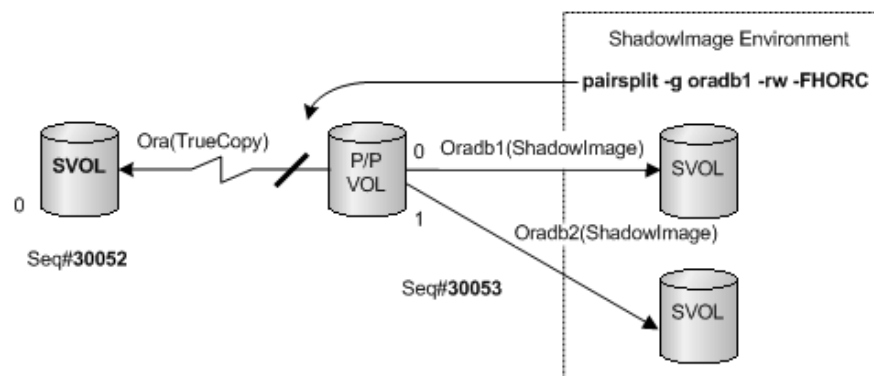
Note: When an option that is valid for only specific volumes is specified for other volumes, the error code might be EX_UNWOPT or EX_UNWCMD. Check whether the specified option is valid for the volume.

Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	Confirm status using the pairdisplay command. Make sure all volumes in the group have the same fence level and volume attributes.	236
	EX_INCSTG	Inconsistent status in group	Confirm pair status using pairdisplay .	229
	EX_INVVOL	Invalid volume status	Confirm volume status/attribute using pairdisplay -l or ' raidvchkdsp -v aou '. "Aou" (allocation on use) refers to dynamic provisioning.	222
	EX_EVOLCE	Pair volume combination error	Confirm pair status using pairdisplay , and change combination of volumes.	235
	EX_INVSTP	Invalid pair status	Confirm pair status using pairdisplay .	228


Category	Error Code	Error Message	Recommended Action	Value
	EX_EWSUSE	Pair suspended at WAIT state	Issue pairresync manually to the identified failed paired volume to try to recover it. If the trouble persists, call the customer support.	234
Timer (Recoverable)	EX_EWSTOT	Timeout waiting for specified status	<p>If an error occurs when the horctakeover or pairsplit command is run for a consistency group, verify that the pair configuration and the pair status are in the status in which the performed pair operation can be accepted by referring to the user guide for each software.</p> <p>If the pair operation is acceptable, increase the timeout value specified by the -t option, and then retry the operation.</p> <p>If this error occurs when a different command is run, verify that the operation and the timeout value are valid.</p>	233

Examples

See the following figure for example of **-FHORC** option for **pairsplit** command.



See the following figure for example of **-FMRCF** option for **pairsplit** command.



pairresync*

The **pairresync** command re-establishes a split pair volume and then restarts the update copy operations to the secondary volume. The **pairresync** command can resynchronize either a paired logical volume or a group of paired volumes. The normal direction of resynchronization is from the primary volume to the secondary volume. If the **-restore** option is specified (ShadowImage/Copy-on-Write Snapshot/Thin Image only), the pair is resynchronized in the reverse direction (that is, secondary volume to primary volume). The primary volume remains accessible during **pairresync**, except when the **-restore** option is specified. The secondary volume becomes write-disabled when the **pairresync** command is issued.

The **pairresync** command cannot be used for Volume Migration.

The **pairresync** command terminates before resynchronization of the secondary (or primary) volume is complete. Use the pair event waiting (**pairevtwait**) or **pairdisplay** command to verify that the resync operation completed successfully (status changes from COPY to PAIR). The execution log file also shows completion of the resync operation. The status transition of the paired volume is judged by the status of the primary volume. The fence level is not changed (only for TrueCopy, TrueCopy Async, Universal Replicator, or global-active device).

If no data was written to the secondary volume while the pair was split, the differential data on the primary volume is copied. If data was written to the secondary volume, the differential data on the primary volume and secondary volume is copied. This process is reversed when the ShadowImage **-restore** option is specified.

Before issuing the **pairresync** command (normal or reverse direction), make sure that the secondary volume is not mounted on any UNIX system. Before issuing a reverse **pairresync** command, make sure that the primary volume is not mounted on any UNIX system.

When you use TrueCopy, TrueCopy for Mainframe, TrueCopy Async, TrueCopy Asynchronous for Mainframe, Universal Replicator, Universal Replicator for Mainframe, or global-active device, if a failure that requires maintenance operations, such as blocking the controller board or the cache memory, occurs on the primary storage system, the **paircreate** command cannot be executed.

When you use global-active device, if a failure that requires maintenance operations occurs on the secondary storage system, the **paircreate** command cannot be executed as well.

Note on Quick Resync/Restore:

If the '\$HORCC_RSYN=QUICK' or the '\$HORCC_REST=QUICK' environment variable is set for any of the following systems, the **pairresync** operation is performed as Quick Resync regardless of the system option mode 87/80 setting via SVP:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

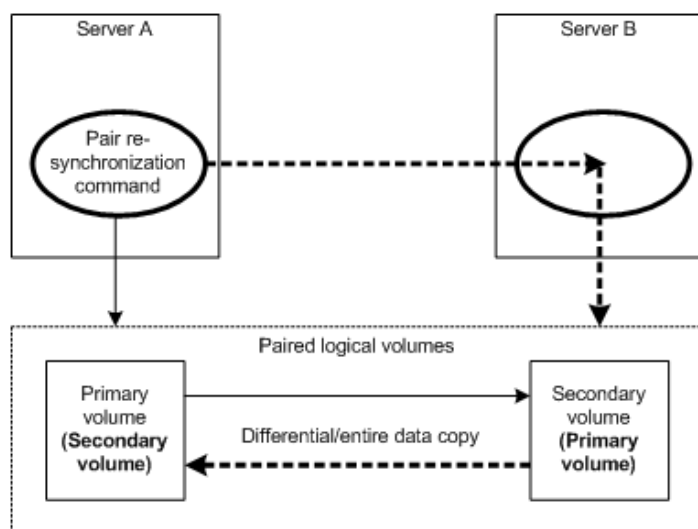
The \$HORCC_RSYN and \$HORCC_REST environment variables are ignored for 9900V.

TrueCopy, TrueCopy Async, Universal Replicator, and GAD only: The **swaps (p)** option copies the differential data from an S-VOL to a P-VOL (pair resynchronization) after setting the P-VOL status to PSUS and the S-VOL status to SSWS. After the copy of the differential data is complete, the volume attributes of the P-VOL and the S-VOL are swapped. The **swaps** option is used to run the command from an instance that manages the S-VOL before swapping the attributes. The **swapp** option is used to run the command from an instance that manages the P-VOL before swapping the attributes. This option is used for operations, such as failback.

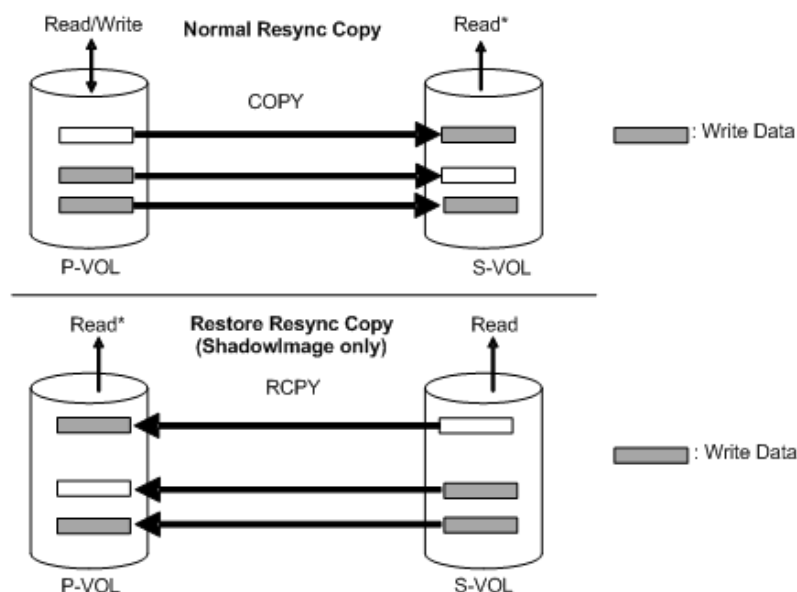
The **swaps (p)** option:

- Ignores the **-l** option.
- If **-c size** option is omitted, use 3 (default value of the number of copy tracks (-c size)).
- Executes the command when the pair status is PAIR including PSUS or PSUE (not applicable to COPY and SMPL).
- If the target volume (remote volume for **-swapp**) is already the P-VOL, the pair operation is skipped.
- Copies the differential data from an S-VOL to a P-VOL (pair resynchronization) after setting the P-VOL status to PSUS and the S-VOL status to SSWS. After the copy of the differential data is complete, the volume attributes of the P-VOL and the S-VOL are swapped.

The following figure shows the pair resynchronization operation.



The following figure shows the Normal Resync and ShadowImage Restore Resync.



*An option for creating a pair. When you specify `-m noread`, you cannot read data in the volume.

The Thin Image Advanced pair operations are available if both of the following conditions are met:

- Both the P-VOL and the S-VOL are DRS-VOLs with LU paths configured.
- The ROW column in the pairdisplay command displays "Y".

If the Thin Image Advanced pair operations are performed by using this command, an audit log of "+Command=Modify Snapshot" or "+Command=Delete Snapshot" is output. If the Thin Image Advanced pair cannot be created, an SSB1 error code starting with 2E might be output. For details about the actions to be taken if this SSB code is output, see the *Command Control Interface User and Reference Guide*. The SSB code description is not displayed in the command output.

If a volume belonging to a virtual storage machine created by a user is specified for HORCM_CMD in the configuration definition file, this command cannot be run for the Thin Image Advanced pair. If this command is run, an EX_CMDRJE (SSB1=2EF3, SSB2=9300) error is displayed.

Syntax

```
pairresync { -h | -q | -z[x] | -I[H][M][instance#] or -I[TC][SI][instance#] | -g
<group> | -d <pair Vol> | -d[g] <raw_device> [MU#][-n] | -FHORC [MU#] | -FMRCF [MU#] | -
d[g] <seq#> <LDEV#> [MU#] | -c <size> | -nomsg | -l | -restore | -swaps | -swapp | -fq
<mode> | -cto <o-time> [c-time] [r-time] | -f[g] <fence> [CTG ID] | -pr <io
preference> | -el <error level> | -hold y }
```

Options and parameters

The primary volume's server is automatically detected by the **pairresync** command, so the server does not need to be specified in the **pairresync** command options.

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits this command.

-z or -zx

Makes the **pairresync** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Used to specify a group name defined in the configuration definition file. The command is executed for the specified group unless the -d <pair Vol> option is specified.

-d <pair Vol>

Specifies a paired logical volume name defined in the configuration definition file.

When this option is specified, the command is executed for the specified paired logical volumes. This option is valid when the -g <group> option is specified.

-d[g] <raw_device> [MU#][-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-FHORC [MU#] or -FCA [MU#]

Forcibly specifies a cascaded remote copy volume for specified pair logical volumes on local copy environment. If the -l option is specified, this option resyncs a cascaded remote copy volume on a local host (near site). If no -l option is specified, this option resyncs a cascaded remote copy volume on a remote host (far site). The target remote copy volume must be a P-VOL. The -swapp option cannot be specified.

-FMRCF [MU#] or -FBC [MU#]

Forcibly specifies a cascaded local copy volume for specified volume pair on remote copy environment. If the -l option is specified, a cascaded local copy volume is split on a local host (near site). If -l option is not specified, a cascaded local copy volume is split on a remote host (far site). The target local copy volume must be a P-VOL.

-swaps -FHORC [MU#]

Swaps the cascaded TrueCopy or Universal Replicator volume from the primary node for failback.

In failback operation from a 3DC cascaded site failure, if you want to failback to DC1 from DC3 directly, it is necessary to operate all cascaded volume from DC1.

In order to make this operation possible, CCI supports the '**pairresync -swaps -FHORC**' option that swaps Universal Replicator volumes on the cascaded TrueCopy Sync/Universal Replicator volume.

-c <size>

TrueCopy, TrueCopy for Mainframe, TrueCopy Async, TrueCopy Async for Mainframe, ShadowImage, and global-active device only.

Specifies the copy pace for the resync operation (range = 1 to 15 tracks). If not specified, the value used for paircreate or **pairsplit** command is used. However, when you use the **-swaps** or **-swapp** option, the default value is 3.

The relationship between the track size and the copy pace is shown as follows. If the copy pace is fast, the host I/O performance might be degraded. To reduce the effect on host I/O performance, lower the copy pace.

- TrueCopy, TrueCopy Async, and global-active device
 - When the track size is 1, the copy pace is slow.
 - When the track size is 2 or 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.
- ShadowImage
 - When the track size is 1 or 2, the copy pace is slow.
 - When the track size is 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.
- TrueCopy for Mainframe, TrueCopy Async for Mainframe
 - When the track size is 1 to 3, the number of the tracks copied simultaneously is 3 and the copy pace is slow.
 - When the track size is 4 or more, the number of the tracks copied simultaneously is 15 and the copy pace is fast.



Note: This option is disabled in Universal Replicator, Universal Replicator for Mainframe, or ShadowImage for Mainframe. To change the copy pace in Universal Replicator or Universal Replicator for Mainframe, use the **raidcom modify journal** command for JNL option changes, or Device Manager - Storage Navigator. ShadowImage for Mainframe works only at medium pace.

-nomsg

Suppresses messages to be displayed when this command is executed. It is used to execute this command from a user program. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

-l

When the remote host cannot be used for host down, this option enables the **pairresync** operation on the local host only. The target volume of the local host must be P-VOL. (For ShadowImage, Copy-on-Write Snapshot, or Thin Image, only S-VOLs can be resynchronized.)

If you specify the **-swapp** option, this option becomes unsupported and cannot be specified at the same time.

-restore

ShadowImage, Copy-on-Write Snapshot, or Thin Image only.

Performs reverse resync (from secondary volume to primary volume).

-swaps

TrueCopy, TrueCopy Async, Universal Replicator, and global-active device only.

Executed from the S-VOL side when there is no host on the P-VOL side to support. Typically executed in PSUS that facilitates 'fast failback' without requiring a full copy. Both **-swaps** and **-swapp** are copied to the original P-VOL based on the delta data from the original S-VOL, and then the S-VOL and the P-VOL are swapped.

-swapp

TrueCopy, TrueCopy Async, Universal Replicator, and global-active device only.

Executes the the same operation as **-swaps** from the original P-VOL, however, **-swapp** does require the cooperation of hosts at both sides.

-fq <mode>

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

ShadowImage and Copy-on-Write Snapshot only.

Specifies the mode whether **pairresync** or the **-restore** option is performed as 'QUICK'.

mode = normal: **pairresync** is performed as Non quick mode regardless of setting of \$HORCC_RSYN environment variable and/or the system option mode 87 via SVP.

mode = quick: **pairresync** is performed as Quick Resync regardless of setting of \$HORCC_RSYN environment variable and/or the system option mode 87 via SVP.

If this option is not specified, then performing **pairresync** is determined by the \$HORCC_RSYN environment variable and/or the system option mode setting through the SVP regardless of whether the **pairresync** operation is Quick Split or not.

The following table shows the relationship between the **-fq** option and the \$HORCC_RSYN environment variable.

-fq option	\$HORCC_RSYN	Behavior
quick	Invalid	Quick Resync
normal	Invalid	Normal Copy
Omitted	QUICK	Quick Resync
Omitted	NORMAL	Normal Copy
Omitted	Omitted	Depends on system option mode (SOM) 87. If SOM 87 is disabled, the option is performed in the non quick mode, whereas if SOM 87 is enabled, the option is performed in the quick mode. If SOMs 87 and 122 are enabled at the same time, the option is performed in the non quick mode. This option is disabled by default.

In the case of Restore (-restore is specified):

mode = normal

The **pairresync -restore** command is performed as Non quick mode regardless of the setting of the \$HORCC_REST environment variable and/or the system option mode 80 via SVP.

mode = quick

The **pairresync -restore** command is performed as Quick Restore regardless of the setting of the \$HORCC_REST environment variable and/or the system option mode 80 via SVP.

If this option is not specified, then the performing of the **pairresync** is determined by the \$HORCC_REST environment variable and/or the system option mode setting through the SVP regardless of whether the **pairresync** operation is Quick Restore or not.

This `-fq` option is also valid for TrueCopy-TrueCopy or ShadowImage cascaded operations using '-F BC [MU#]' option.

The relationship between the `-fq` option and `$HORCC_REST` are shown as below.

-fq option	\$HORCC_REST	Behavior
quick	Invalid	Quick Restore
normal	Invalid	Normal Restore
Omitted	QUICK	Quick Restore
Omitted	NORMAL	Normal Restore
Omitted	Omitted	Depends on system option mode (SOM) 80. If SOM 80 is disabled, Quick Restore is performed, whereas if SOM 80 is enabled, Normal Restore is performed. This option is disabled by default.



Note:

- This `-fq` option is ignored for 9900V to maintain the compatibility on 9900V so that you can add this option to the same script.
- If this `-fq` option is combined with the `-restore` option during the maintenance work in the storage system (SVP is in modify mode), this operation cannot be completed.

-cto <o-time> [c-time] [r-time]

TrueCopy Async and Universal Replicator only.

If you specify '-cto <o-time> [c-time] [r-time]' option on TrueCopy sync, then it is ignored.

- **o-time**: Sets the offloading timer for controlling inflow of write I/O to the specified consistency group. For TrueCopy Async, specify an integer from 1 to 255 (seconds). If this option is not specified, 90 seconds is set by default. For Universal Replicator, specify an integer from 1 to 255 (seconds). If this option is not specified, the value set for the journal on the P-VOL side is applied. Use the **raidcom modify journal** command to change the value range to 256 to 600. If o-time=0 is specified, the inflow control of write I/Os is disabled. When the sidefile capacity exceeds the limit of the sidefile area, write I/Os from the host wait, within the specified timeout period, until the space which is large enough to store next new data becomes available. As the timeout period, you can specify a value from 1 through 255 (seconds) for TrueCopy Async, 1 through 600 (seconds) for Universal Replicator. The default timeout values are 90 seconds for TrueCopy Async, and 60 seconds for Universal Replicator. If the timeout occurs during this waiting state then pair status changes from PAIR to PSUS state of sidefile (Journal) Full, and its host side Write I/O is continued and data is managed by BITMAP mode. Therefore the o-time timeout value should be less than the I/O timeout value of the host system.
- **[c-time]**: (TrueCopy Async only) Sets the Copy Pending timer for the specified consistency group. c-time can be specified from 1 to 15 minutes in increments of 1 minute. If this option is not specified, then this value is set as below.
 - If a consistency group is created, then 5 minutes is set as the default. If not, it is not changed.
- **[r-time]**: (TrueCopy Async only) Sets the RCU Ready timer for the specified consistency group. r-time can be specified from 1 to 10 minutes in increments of 1 minute. If this option is not specified, then the value is not changed.

**Note:**

For TrueCopy Async, settings changed by these options are invalid if a consistency group already exists. These parameters are also forwarded to S-VOL journal with **pairresync** command, and are used when S-VOL is changed to P-VOL. These parameters are maintained and become valid until and when the pair-volumes are changed to SMPL.

For Universal Replicator, these parameters can be set and changed when a pair is resynchronized in a journal volume in the P(S)JSN status. The parameters are forwarded to the S-VOL side and set when the pair is resynchronized. To change settings of journals on the S-VOL side, specify as follows:

1. `pairsplit -g <group>`
2. `pairresync -g <group> -cto <o-time>`

To change settings of journals on the P-VOL side, swap them from the P-VOL to the S-VOL before specifying the above. These parameters are maintained on each journal. Therefore, if you set the value for offloading timer, execute the **raidcom modify journal** command on both P-VOL and S-VOL sides.

-f[g] <fence> [CTG ID]

TrueCopy, TrueCopyAsync, Universal Replicator, and global-active device only.

`-f[g] <fence> [CTG ID]` (TrueCopy only): This option is used to change from existing TrueCopy Sync volumes to TrueCopy Sync CTG without deleting paired-volume. It is possible to change the option (fence level) and the volume attribute as shown below. This option is valid without '-swaps' or '-swapp' option.

Changing the option (fence level) and the volume attribute:

Attribute	Options		
	-f fence	-fg fence	-fg fence CTG ID
Sync	Updates fence	Cmd rejected Abnormal term	Changes to Sync CTG Updates fence*
Sync CTG	Changes to Sync Updates fence	Updates fence	Cmd rejected Abnormal term
Async	Cmd rejected Abnormal term	Cmd rejected Abnormal term	Cmd rejected Abnormal term
Explanation of terms: fence: data, status, or never ('async' cannot be specified) Cmd rejected/Abnormal term.: Command execution is rejected, and the operation terminates abnormally.			

Attribute	Options		
	-f fence	-fg fence	-fg fence CTG ID
*If the identical CTG ID is specified with the different group name of CCI, then the command is rejected. So different CTG ID must be specified for the different group name. Use the raidcom get ctg command to verify how the CTG is being used.			

Changing the volume attribute for global-active device:

Attribute	Options		
	-f never	-fg never	-fg never CTG ID
GAD	No Updates	Cmd rejected Abnormal term	Changes to GAD CTG
GAD CTG	Cmd rejected Abnormal term	No Updates	Cmd rejected Abnormal term
Explanation of terms: Cmd rejected/Abnormal term.: Command execution is rejected, and the operation terminates abnormally.			

-pr <io preference>

This option is only available for global-active device. If this option is specified for any pair that is not a global-active device pair, this option is ignored. If this option is specified for a pair that is already in the pair status, this option is also ignored.

Sets the I/O preference mode (I/O preference mode when remote path failed) when a failure occurs in the remote path between the storage systems of the primary site and the secondary site and communication becomes impossible. For details of the I/O preference mode when remote path failed, see the Global-Active Device User Guide.

The following values can be specified for <io preference>.

pvol: Primary Volume. When remote path failed, the primary volume continues host I/O preferentially.

disable: I/O Preference Mode When Remote Path Failed is disabled. Specify this value to disable I/O Preference Mode When Remote Path Failed.

If this option is omitted, the I/O Preference Mode When Remote Path Failed retains the current setting.

Pairs with different settings of I/O preference mode when remote path failed cannot be mixed in the same consistency group. Therefore, if you want to change the I/O preference mode when remote path failed of a pair that already belongs to a consistency group, make the change for the entire consistency group. If you do not specify the `-d` option, you can change the I/O preference mode when remote path failed for the entire consistency group.

-el <error level>

This option is only available for Universal Replicator and specified only for VSP One B20. If this option is specified for any pair that is not a Universal Replicator pair, this option is ignored. If this option is specified for a pair that is already in the pair status, this option is also ignored.

Specifies an error level in Universal Replicator. To specify this option, make sure to specify the `-d` option at the same time. If the `-d` option is specified but the `-el` option is not specified, the error level is group.

The following value can be specified as `-el <error level>`:

volume: The error level is volume.

This option cannot be used with the `-FHORC` option at the same time.

| -hold y

Changes the HOLD ERROR status to the HOLD status.

This option cannot be used with the `-FHORC` option at the same time. This option is used for troubleshooting. Use this option by following the instructions in the Troubleshooting section in the *Hitachi Universal Replicator User Guide*.

Returned values

The `pairresync` command sets the following returned values during exit allowing you to check the execution results.

- **Normal termination:**
 - **0:** When resynchronizing groups, 0 = normal termination for all pairs.
- **Abnormal termination:**
 - **other than 0:** Refer to the error code for error details.

Error codes

Unrecoverable errors are fixed and not resolved, even after re-executing the command. If the command failed, the detailed status is logged in the CCI command log (\$HORCC_LOG), even if the user script has no error handling.

The `pairresync` command updates both the P-VOL and the S-VOL. When you execute this command, make sure that neither the P-VOL nor the S-VOL is mounted. See the following table for specific error codes for `pairresync`.

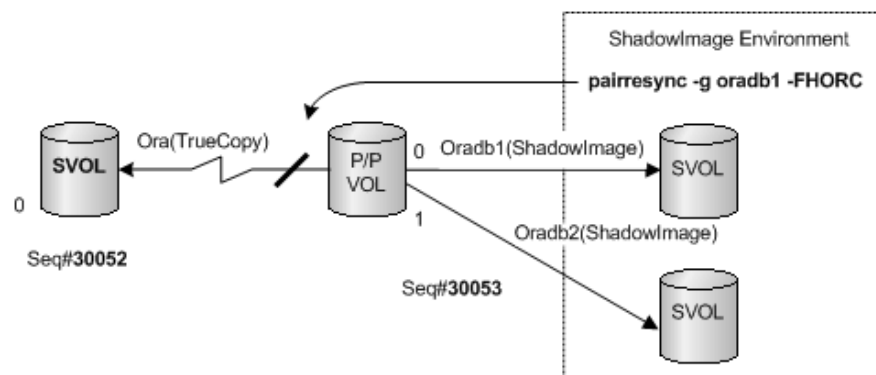


Note: When an option that is valid for only specific volumes is specified for other volumes, the error code might be EX_UNWOPT or EX_UNWCMD. Check whether the specified option is valid for the volume.

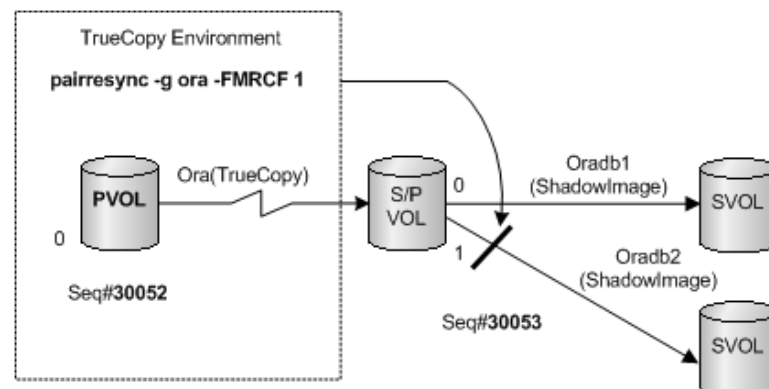
Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	Confirm status using the pairdisplay command. Make sure all volumes in the group have the same fence level and volume attributes.	236
	EX_INCSTG	Inconsistent status in group	Confirm pair status using pairdisplay .	229
	EX_INVVOL	Invalid volume status	Confirm pair status using pairdisplay -l .	222
	EX_INVSTP	Invalid pair status	Confirm pair status using pairdisplay .	228

Examples


See the following figure as an example of **-FHORC** option for **pairresync**



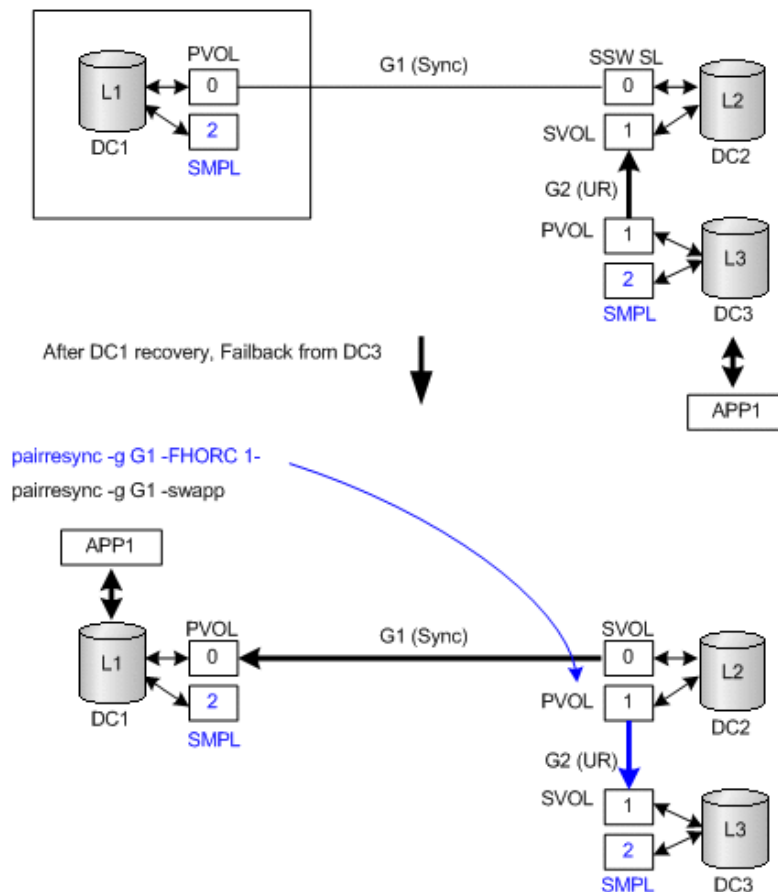
See the following figure as an example of **-FMRCF** option for **pairresync**



See the following figure for swap operation.



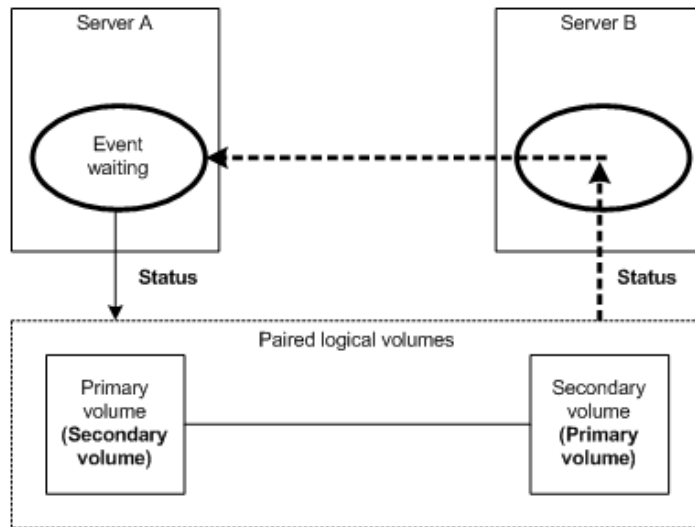
See the following figure as an example of -swaps option with -FHORC [MU#]



pairevtwait*

The **pairevtwait** (pair event waiting) command is used to wait for another command to complete pair operations and to confirm the result of pair operations (see the following figure for pair even waiting operation.). It waits ('sleeps') until the paired volume status becomes identical to a specified status, and terminates abnormally if the PSUE status or any other abnormal status is detected. The transition of the paired volume status is judged by the status of the volume, which is searched automatically.

The **pairevtwait** command can be executed by specifying a paired logical volume or a group of a paired volume. If the **pairevtwait** command is issued by specifying a group and the status, the command waits until the status of each volume in the group changes to the specified status. When the **pairevtwait** command with the **-nowait** or **-nowaits** option is issued for a group, the **pairevtwait** command returns the response immediately according to the status of each volume in the group. The primary and secondary volume's servers are automatically detected by the **pairevtwait** command, so the server does not need to be specified in the **pairevtwait** command parameters. For ShadowImage pairs, this command must be used to confirm a pair status transition.



The operations in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair are available if both of the following conditions are met. This command is used to verify the pair status of the Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair that are displayed in Status in the **pairdisplay** command.

- Both the P-VOL and the S-VOL are used in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute in which an LU path is configured, or a Thin Image Advanced pair in which an LU path is configured.
- The ROW column in the **pairdisplay** command displays "Y".

Syntax

```

pairevtwait{ -h | -q | -z[x] | -I[H][M][instance#] or
             -I[TC][SI][instance#] | -g <group> | -d <pair Vol>
             | -d[g] <raw_device> [MU#][~n] | -FHORC [MU#] | -FMRCF [MU#]
             | -d[g] <seq#> <LDEV#> [MU#] | -s [s] <status> ...
             | -t <timeout>[interval] | -nowait[s] | -l | -nomsg }

```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits this command.

-z or -zx

Makes the `pairevtwait` command enter the interactive mode. The `-zx` option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the `-zx` option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the `HORCC_MRCF` environment variable is set in the command execution environment.

- If the `HORCC_MRCF` environment variable is not set:
a TrueCopy, Universal Replicator, or global-active device pair
- If the `HORCC_MRCF` environment variable is set:
a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the `HORCMINST` environment variable is set in the command execution environment.

- If the `HORCMINST` environment variable is not set:
an instance with no instance number is used.
- If the `HORCMINST` environment variable is set:
an instance whose instance number is set as the `HORCMINST` environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: `-IH0`

-g <group>

Specifies a group name defined in the configuration definition file. The command is executed for the specified group unless the `-d <pair Vol>` option is specified.

-d <pair Vol>

Specifies a paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volumes. This option is valid when the `-g <group>` option is specified.

-d[g] <raw_device> [MU#][-n]

Searches whether the specified `raw_device` is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (`-d`) or group (`-dg`). This option is effective without specification of `-g <group>` option. If the specified `raw_device` is contained in two or more groups, the command is executed for the first group.

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-FHORC [MU#] or -FCA [MU#]

Forcibly specifies a cascaded remote copy volume for specified pair logical volumes on local copy environment. If the `-l` option is specified, this option tests status of a cascaded remote copy volume on a local host (near site). If no `-l` option is specified, this option tests status of a cascaded remote copy volume on a remote host (far site).

-FMRCF [MU#] or -FBC [MU#]

Forcibly specifies a cascaded local copy volume for specified pair logical volumes on remote copy environment. If the `-l` option is specified, this option tests status of a cascaded local copy volume on a local host (near site). If no `-l` option is specified, this option tests status of a cascaded local copy volume on a remote host (far site).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (`-d`) or group (`-dg`). This option is effective without specification of `-g <group>` option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The `<seq#>` `<LDEV#>` values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.



Note:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-s <status>

Specifies the waiting status, which is 'smpl', 'copy'(including 'RCPY'), 'pair', 'psus', or 'psue'. If two or more statuses are specified following -s, waiting is done according to the logical OR of the specified statuses. This option is valid when the `-nowait` option is not specified.

-ss <status>

Specifies the waiting status, which is 'smpl', 'copy'('RCPY' is included), 'pair', 'ssus', 'psue' on S-VOL. If two or more statuses are specified following -ss, waiting is done according to the logical OR of the specified statuses. This option is valid when the `-nowaits` option is not specified.

-t <timeout> [interval]

Specifies the interval (1 to 21474836) of monitoring a status specified using the `-s` and `-ss` option and the time-out period (0 to 21474836) in increments of 1 second. Unless [interval] is specified, the default value (3 seconds) is used. This option is valid when the `-nowait` option is not specified. If <timeout> is specified more than 1999999, then 'WARNING' message appears.

If you execute the command in the Out-of-Band method, specify 3 seconds (default value) or more to <timeout>.

This option is valid only when the `-nowait` or `-nowaits` option is not specified.

-nowait

When this option is specified, the pair status at that time is reported without waiting. The pair status is set as a returned value for this command. When this option is specified, the `-t` and `-s` options are not needed.

-nowaits

When this option is specified, the pairing status on S-VOL at that time is reported without waiting. The pairing status is set as a returned value for this command. When this option is specified, the `-t` and `-ss` options are not needed.

-l

When this command cannot utilize a remote host for host down, this option executes this command by a local host only. When a paired volume used by this command meets any of the following conditions, if this option is not specified, the command fails with EX_EVOLCE:

- Both paired volumes are P-VOLs
- Both paired volumes are S-VOLs
- Either paired volume is SMPL, and the other volume is S-VOL.

-nomsg

Suppresses messages to be displayed when this command is executed. It is used to execute a command from a user program. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

Returned values

The `pairevtwait` command sets the following returned values during exit allowing you to check the execution results. See the table below for more details.

- When the `-nowait` option is specified:
 - **Normal termination:**
 - **1:** The status is SMPL.
 - **2:** The status is COPY or RCPY.
 - **3:** The status is PAIR.
 - **4:** The status is PSUS.
 - **5:** The status is PSUE.
 - When monitoring groups, 1/2/3/4/5 = normal termination for all pairs.
 - **Abnormal termination:**
 - other than 0 to 127: refer to the error codes for error details.
 - The returned value of SSUS is set for PFUS.
 - The returned value of PAIR is set for PFUL.
- When the `-nowaits` option is specified:
 - **Normal termination:**
 - **1:** The status is SMPL.
 - **2:** The status is COPY or RCPY.
 - **3:** The status is PAIR.
 - **4:** The status is SSUS (Note that SVOL_PSUS appears as SSUS).
 - **5:** The status is PSUE.
 - **Abnormal termination:**
 - other than 0 to 127: refer to the error codes for error details.
 - The returned value of SSUS is set for SSWS/PFUS.
 - The returned value of PAIR is set for PFUL.
- When the `-nowait` and/or `nowaits` option is not specified:
 - **Normal termination:**
 - **0:** When monitoring groups, 0 = normal termination for all pairs.
 - **Abnormal termination:**
 - **other than 0 to 127:** refer to the error codes for error details.

Error codes

Unrecoverable errors are fixed and not resolved, even after re-executing the command. If the command failed, the detailed status is logged in the CCI command log (\$HORCC_LOG), even if the user script has no error handling.

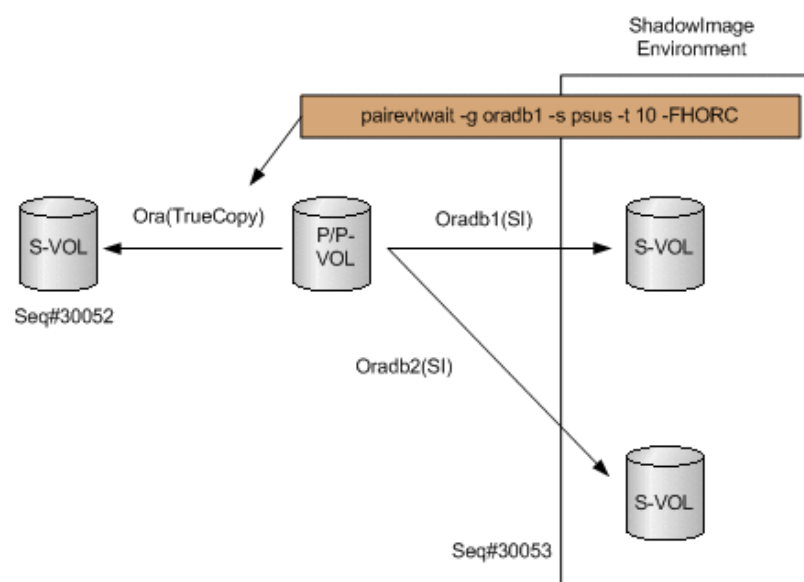
See the table below for specific error codes for **pairevtwait**.

Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	Confirm status using the pairdisplay command. Make sure all volumes in the group have the same fence level and volume attributes.	236
	EX_INCSTG	Inconsistent status in group	Confirm pair status using pairdisplay .	229
	EX_INVVOL	Invalid volume status	Confirm pair status using pairdisplay -l .	222
	EX_EVOLCE	Pair volume combination error	Confirm pair status using pairdisplay , and change combination of volumes.	235
	EX_EWSUSE	Pair suspended at WAIT state	If the PSUE status or any other abnormal status is detected, this error message is output. Issue pairresync manually to the identified failed paired volume to try to recover it. If the trouble persists, contact customer support	234
Timer (Recoverable)	EX_EWSTOT	Timeout waiting for specified status	If an error occurs when the horctakeover or pairsplit command is run for a consistency group, verify that the pair configuration and the pair status are in the status in which the performed pair operation can be accepted by referring to the user guide for each software. If the pair operation is acceptable, increase the timeout value specified by the -t option, and then retry the operation. If this error occurs when a different command is run, verify that the operation and the timeout value are valid.	233

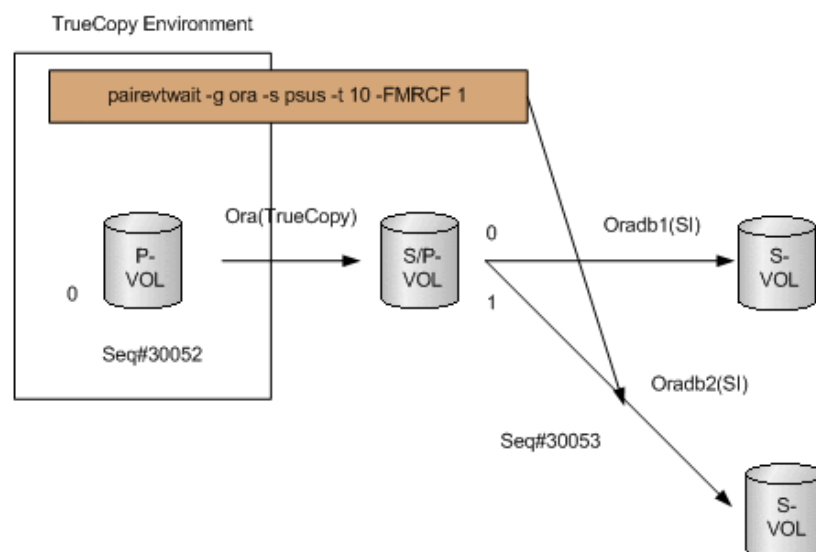
Category	Error Code	Error Message	Recommended Action	Value
	EX_EWSLT O	Timeout waiting for specified status on the local host	Confirm that CCI (HORCM) on the remote host is running.	232

Examples

See the following figure as an example of `-FHORC` option for `pairevtwait`.



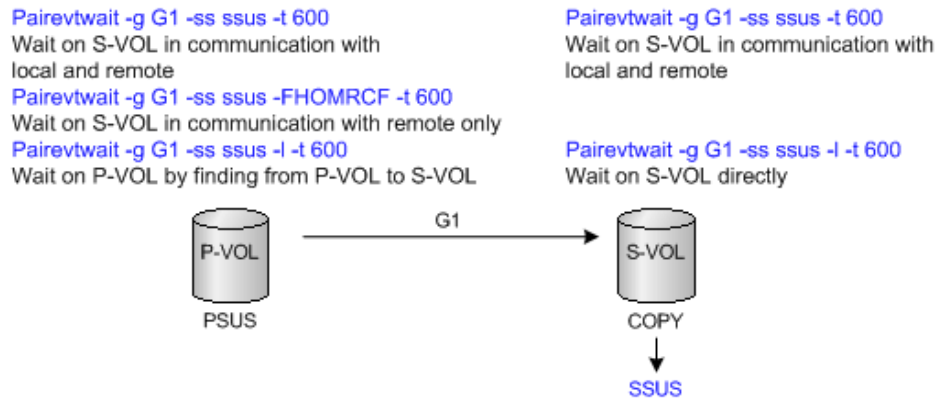
See the following figure as an example of `-FMRCF` option for `pairevtwait`.



Using the -ss <status> ... and -nowaits options.

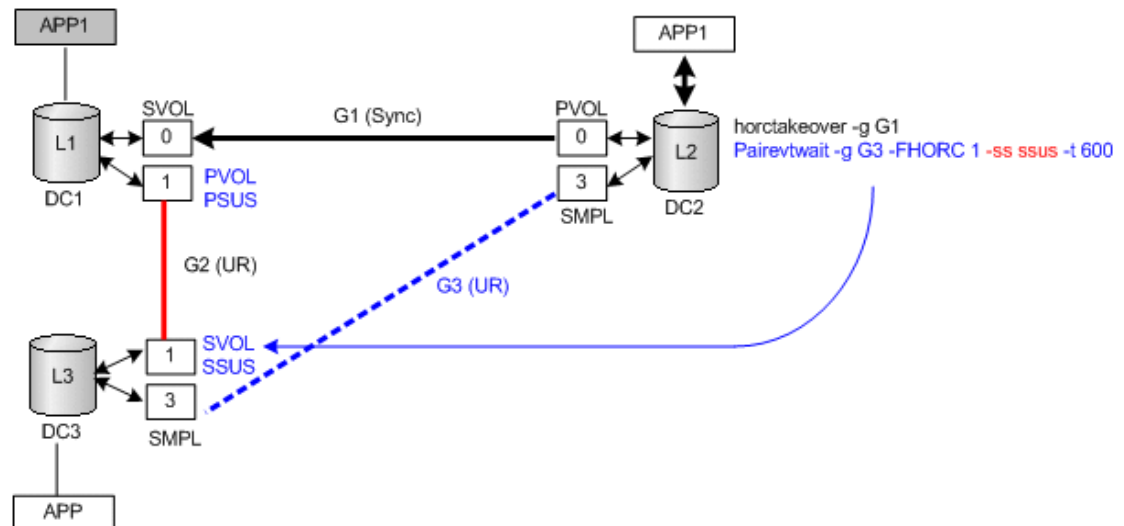
In PVOL_PSUS & SVOL_COPY state of ShadowImage quick mode, **pairevtwait** returns immediately even if the S-VOL is still in SVOL_COPY state because P-VOL is already in PVOL_PSUS state. If you want to wait the SVOL_SSUS state, then use the **-ss <status>** and **-nowaits** options to wait for the pair status on S-VOL side. This is needed for operating **pairresync -restore** or **pairsplit -S**.

The figure below shows an example of waiting on ShadowImage. The figure depicts five examples of waiting until 'PVOL_PSUS' & 'SVOL_COPY' state is changed to SVOL_SSUS.



The **horctakeover** command suspends G2(CA-Jnl) automatically if **horctakeover** returns 'Swap-takeover' as an exit code. In DC1 host failure, if APP1 wants to wait until DC3 becomes the suspend state, then they can verify the 'SSUS' state by using the **pairevtwait** command as shown below.

The following figure shows an example for waiting 'SSUS' on 3DC using TC/UR



pairmon*

The **pairmon** (pair monitor) command, which is connected to the HORCM daemon, obtains the pair status transition of each volume pair and reports the status change. If the pair status changes (due to an error or a user-specified command), the **pairmon** command displays a message.

The pair status transition events exist in the HORCM pair state transfer queue. The **-resevt** option (reset event) deletes one/all events from the HORCM pair state transfer queue. If reset event is not specified, the pair state transfer queue is maintained. If the **-s** option is not specified, **pairmon** displays all events for which it receives information from HORCM. If the **-s** option is specified, only the specified status transitions are displayed.

The CCI software supports the error monitoring and configuration confirmation commands for linkage with the system operation management of the UNIX server.

The operations in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair are available if both of the following conditions are met. This command is used to verify the pair status of the Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair that are displayed in Status in the **pairdisplay** command.

- Both the P-VOL and the S-VOL are used in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute in which an LU path is configured, or a Thin Image Advanced pair in which an LU path is configured.
- The ROW column in the **pairdisplay** command displays "Y".

Syntax

```
pairmon { -h | -q | -z[x] | -I[H][M][instance#] or
          -I[TC][SI][instance#] | -D | -allsnd | -resevt | -nowait
          | -s <status> ... }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits this command.

-z or -zx

Makes the **pairmon** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-D

Selects the default report mode. In the default mode, if there is pair status transition information to be reported, one event is reported and the event is reset. If there is no pair status transition information to be reported, the command waits. The report mode consists of the three flags: `-allsnd`, `-resevt`, and `-nowait` options.

-allsnd

Reports all events if there is pair status transition information.

-resevt

Reports events if there is pair status transition information, and then resets all events.

-nowait

When this option is specified, the command does not wait when there is no pair status transition information.

-s <status>...

Specifies the pair status transition to be reported: smpl, copy (includes rcpy), pair, psus, psue. If two or more statuses are specified following -s, masking is done according to the logical OR of the specified statuses. If this option is not specified, **pairmon** displays all events which received information from HORCM.

Examples

The following shows an example of the **pairmon** command and its output.

```
# pairmon -allsnd -nowait
```

```
Group Pair vol Port targ#lun#LDEV#...Oldstat code
oradb oradb1 CL1-A 1 5 145...SMPL 0x00
oradb oradb2 CL1-A 1 6 146...PAIR 0x02
```

Description of the pairmon command output:**Group**

The group name (dev_group) described in the configuration definition file.

Pair vol

The paired volume name (dev_name) in the specified group that is described in the configuration definition file.

Port targ# lun#

The port ID, TID, and LUN which is described in the configuration definition file. For further information on fibre-to-SCSI address conversion, see the *Command Control Interface Installation and Configuration Guide*.

LDEV#

The LDEV ID for the specified device.

Oldstat

The pair status before changing the volume status.

Newstat

The pair status after changing the volume status.

code

The storage system-internal code for the specified status.

The following table specifies the results of the **pairmon** command options.

-D	-nowait	-resevt	-allsnd	Actions
-D	-	-	-	When HORCM does not have an event, this option waits until an event occurs. If one or more events exist, then it reports one event and resets the event which it reported.
Invalid	-	-	-allsnd	When HORCM does not have an event, this option waits until an event occurs. If one or more events exist, then it reports all events.
Invalid	-	-resevt	-	When HORCM does not have an event, this option waits until an event occurs. If one or more events exist, then it reports one event and resets all events.
Invalid	-	-resevt	-allsnd	When HORCM does not have an event, this option waits until an event occurs. If one or more events exist, then it reports all events and resets all events.
Invalid	-nowait	-	-	When HORCM does not have an event, this option reports event nothing. If one or more events exist, then it reports one event and resets the event which it reported.
Invalid	-nowait	-	-allsnd	When HORCM does not have an event, this option reports event nothing. If one or more events exist, then it reports all events.
Invalid	-nowait	-resevt	-	When HORCM does not have an event, this option reports event nothing. If one or more events exist, then it reports one event and resets all events.
Invalid	-nowait	-resevt	-allsnd	When HORCM does not have an event, this option reports event nothing. If one or more events exist, then it reports all events and resets all events.

pairvolchk*

The **pairvolchk** command checks the attributes and status of a pair volume. It acquires and reports the attribute of a volume or group connected to the local host (issuing the command) or remote host. The volume attribute is SMPL (simplex), P-VOL (primary volume), or S-VOL (secondary volume). The **-s[s]** option reports the pair status in addition to the attribute.

The operations in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair are available if both of the following conditions are met. This command is used to verify the pair status of the Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair that are displayed in Status in the **pairdisplay** command.

- Both the P-VOL and the S-VOL are used in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute in which an LU path is configured, or a Thin Image Advanced pair in which an LU path is configured.
- The ROW column in the **pairdisplay** command displays "Y".



Note: If a snapshot is created by using the **raidcom add snapshot** command, the displayed consistency group ID is not correct. To verify the status of the snapshot created by using the **raidcom add snapshot** command, use the **raidcom get snapshot** command.

Syntax

```
pairvolchk { -h | -q | -z[x] | -I[H][M][instance#] or
-I[TC][SI][instance#] | -g <group> | -d <pair Vol> |
-d[g] <raw_device> [MU#] [-n] | -FHORC [MU#] | -FMRCF [MU#] |
-d[g] <seq#> <LDEV#> [MU#] | -c | -ss | -nomsg }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the pair volume check command.

-z or -zx

Makes the **pairvolchk** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies the group name defined in the configuration definition file. The command is executed for the specified group unless the `-d <pair Vol>` option is specified.

-d <pair Vol>

Specifies the paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volumes. This option is valid when the `-g <group>` option is specified.

-d[g] <raw_device> [MU#] [-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-c

Checks the conform ability of the paired volumes of the local and remote hosts and reports the volume attribute of the remote host. If this option is not specified, the volume attribute of the local host is reported.

-ss

Used when acquiring the attribute of a volume and the pair status of a volume. If this option is not specified, the volume attribute of the local host is reported.

For details on the pair status to be displayed, see the table for **pairvolchk** group status.

-nomsg

Suppresses messages to be displayed when this command is executed. It is used to execute a command from a user program. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

-FHORC [MU#] or -FCA [MU#]

Forcibly specifies a cascaded remote copy volume for specified pair logical volumes on local copy environment. If no -c option is specified, this option acquires the attributes of a cascaded remote copy volume on a local host (near site). If the -c option is specified, this option acquires the attributes of a cascaded remote copy volume on a remote host (far site).

-FMRCF [MU#] or -FBC [MU#]:

Forcibly specifies a cascaded local copy volume for specified pair logical volumes on remote copy environment. If no -c option is specified, acquires the attributes of a cascaded local copy volume on a local host (near site). If the -c option is specified, acquires the attributes of a cascaded local copy volume on a remote host (far site).

Returned values

The `pairvolchk` command sets the following returned values during exit allowing you to check the execution results. See the table below for more details.

- When the `-ss` option is not specified
 - **Normal termination:**
 - **1:** The volume attribute is SMPL.
 - **2:** The volume attribute is P-VOL.
 - **3:** The volume attribute is S-VOL.
 - **Abnormal termination:**
 - **other than 0 to 127:** refer to the error codes for error details.

- When the `-ss` option is specified
 - **Normal termination:**
 - **11:** The status is SMPL.
 - **Abnormal termination:**
 - specific error codes (see Error Codes) and generic error codes.
 - For TrueCopy, ShadowImage, Volume Migration, and global-active device
 - **22:** The status is P-VOL_COPY or P-VOL_RCPY.
 - **23:** The status is P-VOL_PAIR.
 - **24:** The status is P-VOL_PSUS.
 - **25:** The status is P-VOL_PSUE.
 - **26:** The status is P-VOL_PDUB (LUSE volume only).
 - **32:** The status is S-VOL_COPY or S-VOL_RCPY.
 - **33:** The status is S-VOL_PAIR.
 - **34:** The status is S-VOL_PSUS.
 - **35:** The status is S-VOL_PSUE.
 - **36:** The status is S-VOL_PDUB (LUSE volume only).
 - For TrueCopy Async and Universal Replicator

The following values (<status-code> + 20) are returned:

 - **42:** The status is P-VOL_COPY.
 - **43:** The status is P-VOL_PAIR.
 - **44:** The status is P-VOL_PSUS.
 - **45:** The status is P-VOL_PSUE.
 - **46:** The status is P-VOL_PDUB. (LUSE volume only)
 - **47:** The status is P-VOL_PFUL.
 - **48:** The status is P-VOL_PFUS.
 - **52:** The status is S-VOL_COPY or S-VOL_RCPY.
 - **53:** The status is S-VOL_PAIR.
 - **54:** The status is S-VOL_PSUS.
 - **55:** The status is S-VOL_PSUE.
 - **56:** The status is S-VOL_PDUB. (LUSE volume only)
 - **57:** The status is S-VOL_PFUL.
 - **58:** The status is S-VOL_PFUS.
 - For Thin Image and Copy-on-Write Snapshot

Snapshot must display the status when the snapshot pool is full as the snapshot status.

The following values are returned:

- **22:** The status is P-VOL_COPY or P-VOL_RCPY.
- **23:** The status is P-VOL_PAIR.
- **24:** The status is P-VOL_PSUS.
- **25:** The status is P-VOL_PSUE.
- **26:** The status is P-VOL_PDUB. (LUSE volumes only)
- **27:** The status is P-VOL_PFUL.
- **28:** The status is P-VOL_PFUS.
- **32:** The status is S-VOL_COPY or S-VOL_RCPY.
- **33:** The status is S-VOL_PAIR.
- **34:** The status is S-VOL_PSUS.
- **35:** The status is S-VOL_PSUE.
- **36:** The status is S-VOL_PDUB. (LUSE volumes only)
- **37:** The status is S-VOL_PFUL.
- **38:** The status is S-VOL_PFUS.



Note: For pairs in the SSWS status, the command returns the same value as the S-VOL_PSUS status.

Error codes

Category	Error code	Error message	Recommended action	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	Confirm status using the pairdisplay command. Make sure all volumes in the group have the same fence level and volume attributes.	236
	EX_EVOLCE	Pair volume combination error	Confirm pair status using the pairdisplay command, and change combination of volumes.	235

Example 1

Display example for ShadowImage/Copy-on-Write Snapshot:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR ]
```

Display example for ShadowImage (specified with '-m grp' option):

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR, CTGID = 1 ]
```

Display example for TrueCopy:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR, fence = DATA,
MINAP = 2 ]
```

Display example for TrueCopy Sync CTG:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR, fence = DATA,
CTGID = 2 MINAP = 2 ]
```

Display example for TrueCopy Async:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR fence = ASYNC
CTGID = 2 MINAP = 2 ]
```


Description of command output

MINAP

Displays the following two conditions (status) according to the pair status:

PVOL

This shows the minimum in Active Paths on specified group in TrueCopy/TrueCopy Async.



SVOL_SSUS(SSWS)

MINAP shows the result of the suspend operation that indicates whether the remaining data on P-VOL was completely passed (synchronized) to S-VOL. If MINAP is 'one', all data is passed. Otherwise, all data is not passed from P-VOL.



Note: If the microcode on the storage system does not support the active paths, then the 'MINAP' item is not displayed as follows:

```
pairvolchk : Volstat is P-VOL.[status = PAIR fence = ASYNC CTGID = 2]
```

Example 2

Display example for Universal Replicator:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR fence = ASYNC
CTGID = 2 MINAP = 2 ]
```


Description of command output

MINAP

Displays the following two conditions (status) according to the pair status:

PVOL or SVOL_PAIR (except SSUS(SSWS))

This shows the minimum in Active Paths on specified group in Universal Replicator.



SVOL_SSUS(SSWS)

MINAP shows the result of the suspend operation that indicates whether or not the remaining All data on P-VOL were Passed (synchronized) to S-VOL completely. If MINAP is 'one', All data were passed. If not, all data were not passed from P-VOL.



Note: If the microcode on the storage system does not support the active paths, then the 'MINAP' item is not displayed as follows:

```
pairvolchk : Volstat is P-VOL.[status = PAIR fence = ASYNC
CTGID = 2]
```

Example 3

Display example for LDEV blockading:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR fence = DATA
MINAP = 2 LDEV = BLOCKED]
```

Description of command output

LDEV = BLOCKED

Displays the status of LDEV blockading in order to detect a link failure of E-LUN.

Example 4

The following shows examples of the **pairvolchk** command and its output.

TrueCopy Async:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR fence = ASYNC
CTGID = 2]
```

TrueCopy:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR fence = DATA ]
```

ShadowImage:

```
pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR ]
```

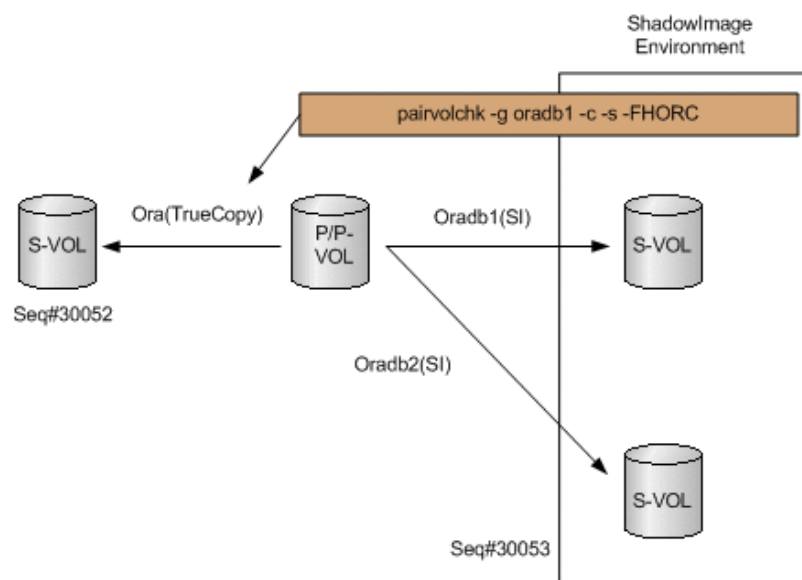
ShadowImage pair splitting with specifying the consistency group:

```
# pairvolchk -g oradb
pairvolchk : Volstat is P-VOL.[status = PAIR CTGID = 1]
```


Example 5

The following shows a pairvolchk example that acquires the status (PVOL_PSUS) of the intermediate P/Pvol through specified pair group on ShadowImage environment. The following figure shows a pairvolchk example that acquires the status (PVOL_PSUS) of the intermediate S/Pvol (MU#1) through specified pair group on TrueCopy environment.

See the following figure as an example of **-FHORC** option for pairvolchk.



See the following figure as an example of **-FMRCF** option for **pairvolchk**



The following table shows the **pairvolchk** group status.

Option	COPY/ RCPY	PSUE	PDU B	PFUS	PSUS	PFUL	PAIR	Group status
	Status of each volume in the group							-
-s*	exist	ignore	ignore	ignore	ignore	ignore	ignore	COPY/ RCPY
	not-exist	exist	ignore	ignore	ignore	ignore	ignore	PSUE
	not-exist	not-exist	exist	ignore	ignore	ignore	ignore	PDUB
	not-exist	not-exist	not-exist	exist	ignore	ignore	ignore	PFUS
	not-exist	not-exist	not-exist	not-exist	exist	ignore	ignore	PSUS
	not-exist	not-exist	not-exist	not-exist	not-exist	exist	ignore	PFUL
	not-exist	not-exist	not-exist	not-exist	not-exist	not-exist	exist	PAIR
-ss	exist	ignore	ignore	ignore	ignore	ignore	ignore	COPY/ RCPY
	not-exist	exist	ignore	ignore	ignore	ignore	ignore	PSUE
	not-exist	not-exist	exist	ignore	ignore	ignore	ignore	PDUB
	not-exist	not-exist	not-exist	ignore	ignore	exist	ignore	PFUL
	not-exist	not-exist	not-exist	ignore	ignore	not-exist	exist	PAIR
	not-exist	not-exist	not-exist	exist	not-exist	not-exist	not-exist	PFUS

Option	COPY/ RCPY	PSUE	PDU B	PFUS	PSUS	PFUL	PAIR	Group status
	not-exist	not- exist	not- exist	not- exist	exist	not- exist	not- exist	PSUS
<p>*This option can be used only when pairvolchk -s has the variable <i>USE_OLD_VCHK</i>.</p> <p>Legend:</p> <ul style="list-style-type: none"> ▪ exist: The target status exists in the group. ▪ not-exist: The target status does not exist in the group. 								

The PFUL state appears as PAIR by all commands (except the -fc option of the **pairedisplay** command), since PFUL indicates PAIR state with sidefile at the HWM.

The PFUS state appears as PSUS by all commands (except the -fc option of the **pairedisplay** command), since PFUS indicates SUSPENDED state due to sidefile full.

The SVOL_PSUS state appears as SSUS by the **pairedisplay** command and other commands.

pairedisplay*

The **pairedisplay** command displays the pair status allowing you to verify completion of pair operations (for example, **paircreate**, **pairresync**). The **pairedisplay** command is also used to confirm the configuration of the pair volume connection path (the physical link of paired volumes and servers). The **pairedisplay** command can be used for a paired volume or a group of paired volumes.

If you want to view a Volume Migration pair created on HDvM - SN or VSP One Block Administrator, use HDvM - SN or VSP One Block Administrator. If you use the **pairedisplay** command, the SMPL status is displayed.

The operations in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair are available if both of the following conditions are met:

- Both the P-VOL and the S-VOL are used in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute in which an LU path is configured, or a Thin Image Advanced pair in which an LU path is configured.
- The ROW column in the **pairedisplay** command displays "Y".

Syntax

```
pairedisplay{ -h | -q | -z[x] | -I[H][M][instance#] or
-I[TC][SI][instance#]
| -g <group> | -d <pair Vol> | -d[g] <raw_device> [MU#][-n]
| -FHORC [MU#]| -FMRCF [MU#] | -d[g] <seq#> <LDEV#> [MU#]
```



```
| -c | -l | -f[xcdmnew] | -CLI | -m <mode> | -v jnl[t]  
| -v ctg | -v pid | -v pidb [-fp]}
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the pair volume check command.

-z or -zx

Makes the **pairdisplay** command enter interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies the group name defined in the configuration definition file. The command is executed for the specified group unless the -d <pair Vol> option is specified.

-d <pair Vol>

This option is used to specify the paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volumes. This option is valid when the -g <group> option is specified.

-d[g] <raw_device> [MU#] [-n]

Searches whether the specified `raw_device` is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified `raw_device` is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-FHORC [MU#] or -FCA [MU#]

Forcibly specifies a cascaded remote copy volume for specified pair logical volumes on local copy environment. If the -l option is specified, this option displays status of a cascaded remote copy volume on a local host (near site). If no -l option is specified, this option displays status of a cascaded remote copy volume on a remote host (far site). This option cannot be specified with -m <mode> option on the same command line.

-FMRCF [MU#] or -FBC [MU#]

Forcibly specifies a cascaded local copy volume for specified pair logical volumes on remote copy environment. If the -l option is specified, this option displays status of a cascaded local copy volume on a local host (near site). If no -l option is specified, this option displays status of a cascaded local copy volume on a remote host (far site). This option cannot be specified with -m <mode> option on the same command line.

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-c

Checks the configuration of the paired volume connection path (physical link of paired volume among the servers) and displays illegal pair configurations. If this option is not specified, the status of the specified paired volume is displayed without checking the path configuration.

-l

Displays the paired volume status of the local host (which issues this command).

-fx

Displays the LDEV ID as a hexadecimal number.

-fc

Displays copy operation progress, sidefile percentage, bitmap percentage, or Universal Replicator journal percentage. The Status item displays the following status in addition to the status displayed when the command is run without the option specified:

- PFUL/PFUS in TrueCopy Async and Universal Replicator
- SSWS displayed after S-VOL_SSUS-takeover

-fd

Displays the relationship between the Device_File and the paired volumes, based on the group (as defined in the local instance configuration definition file). If Device_File column shows 'Unknown' as shown in the Display example, it means that the volume cannot be recognized by the host because the volume has not been registered when you start CCI, and pair operation are rejected (except the local option such as '-l') in protection mode.

Display example:

```
# pairdisplay -g oradb -fd
Group PairVol(L/R) Device_File M ,Seq#,LDEV#.P/S,Status, Seq#,P-LDEV# M
oradb oradev1(L) c0t3d0 0 35013 17..P-VOL COPY, 35013 18 -
oradb oradev1(R) Unknown 0 35013 ****..----- ----, ----- -- -
```

Specify ENABLE for HORCM_NVME in the configuration file. If ENABLE is not specified, a device file name for NVMe is not displayed.

-fm

Displays the management units for the differential data of TC,UR,GAD, or SI pairs. If this option is specified, the management units are displayed in M column.

-fe

Displays the serial# and LDEV# of the external LUNs mapped to the LDEV and additional information for the pair volume. This option displays the information above by adding to last column, and then ignores the format of 80 column. This option is not valid if the cascade options (-m all,-m cas) are specified.

Display example for TC/TC Async/UR/GAD:

```
#pairedisplay.exe -g gad008 -I100 -fe
Group PairVol(L/R) (Port#,TID, LU),Seq#,LDEV#.P/
S,Status,Fence,Seq#,P-LDEV# M CTG JID AP EM E-Seq# E-LDEV# R/W QM DM P PR CS
D_Status ST ELV PGID CT(s) LUT
gad008 gad008dev001(L) (CL1-A-1, 1, 78)800089 78.P-VOL PAIR NEVER ,800013 78 -
- 0 2 - - - L/M - S N D 3 PAIR 02 - 255 0 2023-11-02T19:42:15
gad008 gad008dev001(R) (CL1-A-1, 1, 78)800013 78.S-VOL PAIR NEVER ,----- 78 -
- 0 2 - - - L/M - S N D 3 PAIR 02 - 255 - 2023-11-02T19:42:01
gad008 gad008dev002(L) (CL1-A-1, 1, 79)800089 79.P-VOL PAIR NEVER ,800013 79 -
- 0 2 - - - L/M - S N D 3 PAIR 02 - 255 0 2023-11-02T19:42:15
gad008 gad008dev002(R) (CL1-A-1, 1, 79)800013 79.S-VOL PAIR NEVER ,----- 79 -
- 0 2 - - - L/M - S N D 3 PAIR 02 - 255 - 2023-11-02T19:42:01
```

Display example for ShadowImage/Copy-on-Write Snapshot/Thin Image/Volume Migration:

```
#pairedisplay.exe -g si -IM12 -fe
Group PairVol(L/R) (Port#, TID, LU-M) ,Seq#,LDEV#.P/S,Status,Seq#,P-LDEV# M
CTG CM EM E-Seq# E-LDEV# DM P CS D_Status ROW SNAP_MODE
si sidev1(L) (CL7-H-0,16, 18-0 )500054 30020.P-VOL PAIR,500054 30022 - - N - -
- S N 3 - NI -----
si sidev1(R) (CL7-H-0,16, 20-0 )500054 30022.S-VOL PAIR,----- 30020 - - N - -
- S N - - NI -----
si sidev2(L) (CL7-H-0,16, 19-0 )500054 30021.P-VOL PAIR,500054 30023 - - N - -
- S N 3 - NI -----
si sidev2(R) (CL7-H-0,16, 21-0 )500054 30023.S-VOL PAIR,----- 30021 - - N - -
- S N - - NI -----
```

- CTG: For TC, TC Async, UR, and GAD, it displays the consistency group ID when the fence level is ASYNC. For TC Sync CTG, it displays the consistency group ID when the fence level is DATA, STATUS, or NEVER.

**Note:**

If the snapshot is made by `raidcom add snapshot`, the displayed consistency group ID is not correct. To confirm the status of a snapshot that was made by `raidcom add snapshot`, use the **`raidcom get snapshot`** command.

- JID: The journal ID for P-VOL or S-VOL. In a HAM or GAD configuration, it shows the quorum ID and the fence level is set to 'Never'. When it is not the UR or HAM configuration, '-' is displayed.
- AP: The number of active paths for UR links on P-VOL, and it displays the number of active paths for UR links on P-VOL and S-VOL.'Unknown'is shown as '-'. Refer to 'MINAP' information that is displayed by running the `pairvolchk` command.

- CM: The Copy mode
 - N: Non SnapShot
 - S: Thin Image (TI only) snapshot with the cascade attribute disabled. In the SMPL status, the volume is a V-VOL used in Thin Image.
 - C: Volume Migration
 - R: Thin Image Advanced snapshot. This information is displayed only when the ROW column displays "Y".
 - A: Thin Image (TI only) snapshot with the cascade attribute enabled. This information is displayed only when the ROW column displays "Y".
- EM: The external connection mode
 - H: Mapped E-lun as hidden from the host.
 - V: Mapped E-lun as visible to the host
 - ' - ': Unmapped to the E-lun
 - BH: Mapped E-lun as hidden from the host, but LDEV blockading.
 - BV: Mapped E-lun as visible to the host, but LDEV blockading
 - B: Unmapped to the E-lun, but LDEV blockading
- E-Seq#: The production (serial) number of the external LUN. 'Unknown' is shown as '-'.
 - The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
 - The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).
- E-LDEV#: The LDEV# of the external LUN, 'Unknown' is shown as '-'.
 - L/L: Read/Write is performed on both the primary and secondary volumes.
 - L/M: Read is performed on both the primary and secondary volumes. Write is performed on the primary volume first, and then done on the secondary volume.
 - B/B: Read/Write is rejected (Illegal Request is replied). LU-undefined is returned as a response of the Inquiry command on this volume.
- R/W: Displays the I/O mode when Read/Write is performed. For GAD configuration, the following modes are displayed:
 - L/L: Read/Write is performed on both the primary and secondary volumes.
 - L/M: Read is performed on both the primary and secondary volumes. Write is performed on the primary volume first, and then done on the secondary volume.
 - B/B: Read/Write is rejected (Illegal Request is replied). LU-undefined is returned as a response of the Inquiry command on this volume.

- QM: For GAD pairs, displays the pair operation mode when the quorum disk is blocked.
 - SP: Pair suspended.
 - AS: Pair retained (P-VOL accessible, S-VOL inaccessible).
 - AA: Pair retained (Both P-VOL and S-VOL accessible).
 - - (hyphen): Displayed for any of the following:
 - A non-GAD pair
 - A suspended GAD pair
 - A GAD pair without operation mode display support
 - A GAD pair, but the operation mode that is not supported by CCI is set
 - No volume is set for the quorum disk.
- DM: Displays the differential data management method
Supported storage systems:
 - VSP 5000 series
 - VSP G1x00 and VSP F1500
 - VSP One B20
 - VSP E series
 - VSP G130, G/F350, G/F370, G/F700, G/F900
 - S: The differential data is managed using a shared memory.
 - D: The differential data is managed using a drive (tier difference method).
 - - (hyphen): This information is not available.

- P: Displays the status of expansion processing of DP-VOL

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- E: Expansion is in progress.
- N: Not processed.
- - (hyphen): This information is not available.

A hyphen (-) is displayed if the microcode version is not supported.

- PR: Displays the I/O preference mode when remote path failed of the GAD pair. This information is valid only for GAD pairs.

- P: Primary Volume.
- D: I/O Preference Mode When Remote Path Failed is disabled.
- - (hyphen): This information is not available.

A hyphen (-) is displayed if the DKCMAIN microcode version is not supported, or the pair is not a GAD pair.

-fe

- CS: Displays the copy pace. A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid.
 - TrueCopy, TrueCopy Async, global-active device:
 - When the track size is 1 or 2, the copy pace is slow.
 - When the track size is 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.
 - ShadowImage
 - When the track size is 1 or 2, the copy pace is slow.
 - When the track size is 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.

- **D_Status:** Displays the detailed pair status. This status information might be changed or added in the future for further details.

The following items are displayed in **D_Status**. For the other status, the same Status information as that displayed when the **-fc** option is specified is displayed.

A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid. This information is only supported in Universal Replicator, global-active device, ShadowImage, and Thin Image.

- **CPPD:** A ShadowImage pair or a Thin Image pair is in an initial copy process. For the Thin Image pair, this information is displayed only when the ROW column displays "Y".
- **INIT:** global-active device is about to start an initial copy operation.
- **CPRS:** A ShadowImage pair or a Thin Image pair is being resynchronized. For the Thin Image pair, this information is displayed only when the ROW column displays "Y".
- **SMPP:** A ShadowImage pair or a Thin Image pair is being deleted. For the Thin Image pair, this information is displayed only when the ROW column displays "Y".
- **DELI:** A Universal Replicator pair is being deleted.
- **SUSI:** A Universal Replicator pair is being suspended.
- **HOLD:** A Universal Replicator pair for delta resync is ready for delta resync.
- **HLDI:** A Universal Replicator pair for delta resync is being transitioned to the HOLD status.
- **HLDE:** An error occurred in a Universal Replicator pair for delta resync in the HOLD status.
- **PSUP:** A Thin Image pair is being split. This information is displayed only when the ROW column displays "Y".
- **CPYP:** The snapshot data of a Thin Image pair is being relocated. This information is displayed only when the ROW column displays "Y".
- **CPYD:** The differential data of a Thin Image pair is being copied. This information is displayed only when the ROW column displays "Y".

- ROW

Displays whether the Thin Image Advanced pair operation is available by using the **pairXX** command.

- NI: The Thin Image Advanced pair operation cannot be performed because the instance started when HORCM_SUPROW is not listed in the configuration definition file or HORCM_SUPROW environment variables are not set.
- NM: The instance started when HORCM_SUPROW is listed in the configuration definition file or HORCM_SUPROW environment variables are set, however, the DKCMAIN microcode does not support the Thin Image Advanced pair operations by using the **pairXX** command.
- Y: The Thin Image Advanced pair operations are available by using the **pairXX** command. To view a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, or the Thin Image Advanced pair, use the **pairXX** command.

The operations above are available if all the following conditions are met:

- Make sure to start the instances for both the P-VOL and the S-VOL with the same settings when HORCM_SUPROW is listed in the configuration definition file or HORCM_SUPROW environment variables are set.
- The DKCMAIN microcode supports the Thin Image Advanced pair operations by using the **pairXX** command.
- (hyphen): This information is not valid.

- SNAP_MODE

Displays the same information as "MODE" in the **raidcom get snapshot -key details** command. This information is available only when the ROW column displays "Y". This information is not valid if the pair to be displayed is not a Thin Image pair, or if the ROW column displays a value that is not "Y".

- ST: Displays the pair status in TrueCopy, Universal Replicator, or global-active device in hexadecimal.
 - 0x01: COPY status
 - 0x02: PAIR status
 - Status other than those listed above: The same status as the pair suspend type

For details about the pair suspend type, see the *Hitachi TrueCopy® User Guide*, the *Hitachi Universal Replicator User Guide*, or the *Global-Active Device User Guide*. A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid.

- ELV: Displays Y if the error level of a Universal Replicator pair is volume. Otherwise, displays N. A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid.
- PGID: Displays the path group ID in decimal. A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid.

- CT: Displays the pair copy time in decimal seconds. A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid.
- LUT: Displays the date and time of the last update in the following format. A hyphen (-) is displayed if displaying this information is not supported on the storage system or the information is not valid.

Format: YYYY-MM-DDThh:mm:ss

YYYY is the 4-digit year, MM is the month, and DD is the day of the month. hh is the hour in 24-hour format and the range is 0-23. mm is the minutes, and ss is the seconds.

Example: The following displays January 2, 1970 at 1:00 PM.

```
1970-01-02T13:00:00
```



Note: T is the symbol that separates the date and the time.

-fw

Used when displaying the WWN setting (defined by '50060E80+RAID Type,Serial#,Port#)) to the port instead of a port name. If this option is specified with the '-fe' option at the same line, then 'LUN-WWN' appears as shown below. If values cannot be displayed for LUN-WWN, then 'Unknown' appears (for example, SCSI). For details about the cases where LUN WWNs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#).

Display example:

```
# pairedisplay -g PG01 -fw
Group PairVol(L/R) (WWN, LU-M),Seq#,LDEV#.P/S,Status, Seq#,P-LDEV# M
PG01 ora01(L) (500060e804f42001, 3-0) 62496 67. P-VOL COPY, 62496 69 -
PG01 ora01(R) (500060e804f42001, 5-0) 62496 69. S-VOL COPY, ----- 67 -
PG01 ora02(L) (500060e804f42001, 4-0) 62496 68. P-VOL COPY, 62496 64 -
PG01 ora02(R) (500060e804f42001, 6-0) 62496 64. S-VOL COPY, ----- 68 -
```

Display example:

```
# pairedisplay -IH -g PG01 -CLI -l -fwe
Group PairVol L/R WWN LU Seq# LDEV# P/S Status Fence Seq#
P-LDEV# M CTG JID AP EM E-Seq# E-LDEV# LUN-WWN R/W QM DM P PR
PG01 oradb01 L 500060e804fa0f01 1 64015 11 P-VOL PAIR ASYNC 62496 11 - 5 - 1
- - - 60060e8005fa0f000000fa0f0000000b -/- - - - -
PG01 oradb02 L 500060e804fa0f01 2 64015 12 P-VOL PAIR ASYNC 62496 12 - 5 - 1
- - - 60060e8005fa0f000000fa0f0000000c -/- - - - -
```

- NGUID: Displays the namespace identifier (NGU ID). If NGUID is not assigned, 'Unknown' is displayed.

-fne

Used when displaying the namespace identifier (NGU ID).

Display example:

```
#pairdisplay.exe -g ur008 -I100 -fne
Group PairVol(L/R) (Port#,TID, LU),Seq#,LDEV#.P/S,Status,Fence,Seq#,P-LDEV# M
CTG JID AP EM E-Seq# E-LDEV# NGUID R/W QM DM P PR CS D_Status ST ELV PGID
CT(s) LUT
ur008 ur008dev001(L) (CL1-A-1, 1, 62)800089 62.P-VOL PAIR ASYNC ,800013 62 -
31 7 2 - - - 00500036000000000060e80800360101 -/- - S N - 3 PAIR 02 N 254 1
2023-11-02T19:51:46
ur008 ur008dev001(R) (CL1-A-1, 1, 62)800013 62.S-VOL PAIR ASYNC ,----- 62 - 31
7 2 - - - 00500036000000000060e80800360102 -/- - S N - 3 PAIR 02 N 254 - 2023-
11-02T19:51:32
```

- NGUID: Displays the namespace identifier (NGU ID). If values cannot be displayed for NGUID, 'Unknown' is displayed. For details about the cases where NGU IDs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#).

-fwne

Used when displaying the namespace identifier (NGU ID) and LUNWWN.

Display example:

```
#pairdisplay.exe -g ur008 -I100 -fwne -CLI
Group PairVol L/R WWN LU Seq# LDEV# P/S Status Fence Seq# P-LDEV# M CTG JID AP
EM E-Seq# E-LDEV# LUNWWN NGUID R/W QM DM P PR CS D_Status ST ELV PGID CT(s) LUT
ur008 ur008dev001 L 50060e8028005900 62 800089 62 P-VOL PAIR ASYNC 800013 62 -
31 7 2 - - - Unknown 00500036000000000060e80800360101 -/- - S N - 3 PAIR 02 N
254 1 2023-11-02T19:51:46
ur008 ur008dev001 R 50060e8028000d00 62 800013 62 S-VOL PAIR ASYNC - 62 - 31 7
2 - - - Unknown 00500036000000000060e80800360102 -/- - S N - 3 PAIR 02 N 254 -
2023-11-02T19:51:32
```

- NGUID: Displays the namespace identifier (NGU ID). If values cannot be displayed for NGUID, 'Unknown' is displayed. For details about the cases where NGU IDs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#).

-CLI

Used when specifying display for command line interface (CLI). This option displays to the same position that defined number of columns, and displays one header. The delimiters between columns are displayed as spaces or hyphens (-).

Display example:

```
Group PairVol L/R Port# TID LU-M Seq# LDEV# P/S Status Seq#
P-LDEV# M
homrcf1 deva0 L CL1-D 3 5 0 30053 271 P-VOL PAIR 30053
```

```

263 -
homrcf1 deva1 L CL1-D 3 5 1 30053 271 SMPL - -
- -
homrcf1 deva2 L CL1-D 3 5 2 30053 271 SMPL - -
- -

```

-m <mode>

Used when displaying a paired status of each mirror descriptors for the specified pair logical volume, and used for paired status indication of the cascaded volume. The <mode> option can be designated 'cas' or 'all':

- The 'cas' option is used to verify whether a LDEV in the pair has a different mirror descriptor (MU#) in the configuration definition file in addition to the pair information specified with the `-g` or `-d` option. If the LDEV has a different mirror descriptor (MU#), the pair information is displayed.
- The 'all' option is used when displaying a paired status of all mirror descriptors (MU#).

This option (`-m <mode>`) is not affected with command execution environment (TrueCopy/TrueCopy Async/Universal Replicator/GAD and ShadowImage/Copy-on-Write Snapshot/Thin Image/Volume Migration), and displays the paired status. This option cannot be specified with the `-FHORC/-FMRCF` option on the same command line.

-v jnl[t]

Displays the JNL status for the local and remote interconnected to the group. Also finds the journal ID for each local and remote interconnected to the group via the specified group or <raw_device>, and displays information of each journal ID corresponding the local and remote. The first line shows the journal information for the local host, second line shows the journal information for the remote host. The item for displaying is the same as `raidvchkscan -v jnl[t]`. The serial number (Seq#) display varies depending on storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Example:

```
# pairedisplay -g VG01 -v jnl
```

JID	MU	CTG	JNLS	AP	U(%)	Q-Marker	Q-CNT	D-SZ (BLK)	Seq#	Nnm	LDEV#
-----	----	-----	------	----	------	----------	-------	------------	------	-----	-------

001	0	2	PJNN	4	21	43216fde	30	512345	62500	2	265
-----	---	---	------	---	----	----------	----	--------	-------	---	-----

002	0	2	SJNN	4	95	3459fd43	52000	512345	62538	3	270
-----	---	---	------	---	----	----------	-------	--------	-------	---	-----

Example:

```
# pairedisplay -g VG01 -v jnlt
```

JID	MU	CTG	JNLS	AP	U(%)	Q-Marker	Q-CNT	D-SZ (BLK)	Seq#	DOW	PBW	APW
-----	----	-----	------	----	------	----------	-------	------------	------	-----	-----	-----

001	1	2	PJNN	4	21	43216fde	30	512345	62500	20	300	40
-----	---	---	------	---	----	----------	----	--------	-------	----	-----	----

002	1	2	SJNN	4	95	3459fd43	52000	512345	62538	20	300	40
-----	---	---	------	---	----	----------	-------	--------	-------	----	-----	----

Example:

```
# pairedisplay -g VG01 -v jnl -FCA 1
```

JID	MU	CTG	JNLS	AP	U(%)	Q-Marker	Q-CNT	D-SZ (BLK)	Seq#	Nnm	LDEV#
-----	----	-----	------	----	------	----------	-------	------------	------	-----	-------

003	1	2	PJNN	4	21	43216fde	30	512345	62500	2	265
-----	---	---	------	---	----	----------	----	--------	-------	---	-----

**Note:**

- This option can be specified with following options on the same command line:

```
{-g<group> | -d <pair Vol> | -d[g] <raw_device> [MU#] | -FCA [MU#] |  
-d[g] <seq#> <LDEV#> [MU#] | -l | -f[x] }
```

- The '-FHORC [MU#]' or '-FCA [MU#]' option is used when displaying the journal information of cascaded Universal Replicator volume, and then it shows the journal information for the remote host only.
- This option displays nothing if the target volume is NOT a Universal Replicator volume.

-v ctg

This option finds the consistency group for each local and remote interconnected to the group via the specified group or <raw_device>, and displays any information of each consistency group corresponding the local and remote. The first line shows the consistency group information for the local host, second line shows the consistency group information for the remote host.

Example:

```
# pairedisplay -g VG01-v ctg
```

```
CTG P/S Status AP U(%) Q-Marker QM-Cnt SF(%) Seq# IFC OT/s  
CT/m RT/m
```

```
001 P-VOL PAIR 2 0 00000080 3 50 63528 ON 90 5 5
```

```
001 S-VOL PAIR - 0 0000007d - 50 63528 - - - -
```

CTG: Displays the consistency group ID.

P/S: The attribute of a volume in first LDEV of the specified group.

Status: The status of the paired volume in first LDEV of the specified group.

AP: Displays the number of Active Path in Universal Replicator links on P-VOL, also displays the number of active path in Universal Replicator links on P-VOL and S-VOL, 'Unknown' is shown as '-'.
'Unknown' is shown as '-'.

U(%): The usage sidefile/journal data, it is valid at PAIR state. For VSP G350, G370, G700, G900, VSP F350, F370, F700, F900, this indicates the actual usage rate when the -fp option is specified.

For TrueCopy Async: The sidefile percentage for consistency group in relationship to a 100% full sidefile in cache.

For Universal Replicator: The usage rate of the current journal data as 100% of the journal data space.

Q-Marker: In P-VOL, the latest sequence # of the MCU P-VOL when the write command was received. In S-VOL, the latest sequence # of the DFW on RCU. This item is valid at PAIR state.

QM-Cnt: The number of remaining Q-Marker within consistency group of the Unit. TrueCopy Async sends a token called 'dummy recordset' at regular interval time, therefore QM-Cnt always shows '2' or '3' even if host has NO writing. This item is valid at PAIR state.

SF(%): The usage of cache setting as the sidefile regardless of Universal Replicator and TrueCopy Async.

Seq#: The serial number of the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

IFC: Shows 'ON'/'OFF' for controlling Inflow.

OT/s: Displays the value of the 'offloading timer' (in seconds) for TrueCopy Async or the value of the Data Overflow Watch timer (in seconds) for Universal Replicator. Either value is displayed only for the consistency group containing the P-VOL when control of write I/O inflow (IFC) is enabled (ON). In other cases, a hyphen (-) is displayed.

CT/m: The 'Copy Pending timer' (specified in minutes) setting to consistency group for only TrueCopy Async.

RT/m: The 'RCU Ready timer' (specified in minutes) setting to consistency group for only TrueCopy Async.



Note:

This option displays nothing if the target volume is NOT a TrueCopy Async or Universal Replicator volume. The '-FHORC [MU#]' or '-FCA [MU#]' option is used when displaying the information of cascaded TrueCopy Async/Universal Replicator volume, and then it shows the consistency group information for the remote host only.

-v pid

Finds the pool ID for each local and remote interconnected to the group via the specified group or <raw_device>, and displays any information of each pool ID corresponding the local and remote. The first line shows the pool information for the local host, second line shows the pool information for the remote host. The displayed items are the same as for raidvchkscan -v pid.



Note:

- This option displays nothing if the target volume is not a V-VOL used in Copy-on-Write Snapshot or Thin Image.

Example:

```
# pairedisplay -g VG01 -v pid
```

PID	POLS	U(%)	SSCNT	Available(MB)	Capacity(MB)	Seq#	Num	LDEV#	H(%)
-----	------	------	-------	---------------	--------------	------	-----	-------	------

127	POLN	0	6	3000	3000	63528	2		
864		80							

127	POLN	0	6	3000	3000	63528	2		
864		80							

```
# pairedisplay -g VG01 -v pid -l
```

PID	POLS	U(%)	SSCNT	Available(MB)	Capacity(MB)	Seq#	Num	LDEV#	H(%)
-----	------	------	-------	---------------	--------------	------	-----	-------	------

127	POLN	0	6	3000	3000	63528	2		
864		80							

H(%): Displays the threshold rate being set to the Copy-on-Write Snapshot pool as the high water mark. 'Unknown' is shown as '-'.

-v pid -FMRCF

The '-FBC [MU#]' option is used for displaying the pool information of cascaded snapshot volumes, so that you can monitor the pool status on remote host connected to cascaded TrueCopy P-VOL to TrueCopy S-VOL/P-VOL. It shows the pool information for the remote host only.

**Note:**

- This option displays nothing if the target volume is not a V-VOL used in Copy-on-Write Snapshot or Thin Image.

Example:

```
# pairedisplay -g VG01 -v pid -FMRCF
```

PID	POLS	U(%)	SSCNT	Available(MB)	Capacity(MB)	Seq#	Num	LDEV#	H(%)
-----	------	------	-------	---------------	--------------	------	-----	-------	------

127	POLN	0	6	3000	3000	63528	2		
864		80							

-v pidb

Displays basic information of the pool.



Note:

- This option displays nothing if the target volume is not a snapshot volume.

-v pidb -fp

Displays the actual capacity of the pool.

If the target volume is not a snapshot volume, nothing is displayed.

Returned values

None

Error codes

None

Examples

pairedisplay command example for TrueCopy, TrueCopy Async, Universal Replicator, or global-active device

```
# pairedisplay -g oradb -fcx
```


```
Group  Pair Vol (L/R) (P,T#,L#), Seq#, LDEV#..P/S, Status,
Fence, Copy%, P-LDEV# M
oradb  oradb1 (L) (CL1-B, 1,0) 1234 64..P-VOL PAIR
Never, 75 C8 -
oradb  oradb1 (R) (CL1-A, 1,0) 5678 C8..S-VOL PAIR
Never, ---- 64 -
```

pairedisplay command example for ShadowImage, Copy-on-Write Snapshot, or Thin Image

```
# pairedisplay -g oradb
```

```
Group  Pair Vol (L/R) (Port#,TID,LU-M), Seq#, LDEV#..P/S, Status,
Fence, Seq#, P-LDEV# M
oradb  oradb1 (L) (CL1-A, 1,0) 30053 18..P-VOL PAIR
Never, 30053 19 -
oradb  oradb1 (R) (CL1-D, 1,0) 30053 19..S-VOL PAIR
Never, ---- 18 -
```

Examples of pairedisplay -m



Display example for -m cas

```
# pairedisplay -g oradb -m cas
```

Group	PairVol(L/R)	(Port#,TID,LU-M)	Seq#	LDEV#..P/S	Status
oradb	oradev1(L)	(CL1-D , 3, 0-0)	30052	266....SMPL	----
oradb	oradev1(L)	(CL1-D , 3, 0)	30052	266....P-VOL COPY,	
	30053 268				-
oradb1	oradev11(R)	(CL1-D , 3, 2-0)	30053	268....P-VOL COPY,	
	30053 270				-
oradb2	oradev21(R)	(CL1-D , 3, 2-1)	30053	268....P-VOL PSUS,	
	30053 272				W
oradb	oradev1(R)	(CL1-D , 3, 2)	30053	268....S-VOL COPY,	
	----- 266				-

Display examples for -m all

```
# pairedisplay -g oradb -m all
```

Group	PairVol(L/R)	(Port#,TID,LU-M)	Seq#	LDEV#..P/S	Status
oradb	oradev1(L)	(CL1-D , 3, 0-0)	30052	266....SMPL	----
	----- (L)	(CL1-D , 3, 0-1)	30052	266....SMPL	----
	----- (L)	(CL1-D , 3, 0-2)	30052	266....SMPL	----
oradb	oradev1(L)	(CL1-D , 3, 0)	30052	266....P-VOL PAIR,	
	30053 268				-
oradb1	oradev11(R)	(CL1-D , 3, 2-0)	30053	268....P-VOL COPY,	
	30053 270				-
oradb2	oradev21(R)	(CL1-D , 3, 2-1)	30053	268....P-VOL PSUS,	
	30053 272				W
	----- (R)	(CL1-D , 3, 2-1)	30053	268....SMPL	----
	-----				-

```
oradb   oradev1(R)   (CL1-D , 3, 2)   30053  268....S-VOL COPY,
----- 266         -
```

```
# pairedisplay -d /dev/rdisk/c0t3d0 -l -m all
```

```
Group   PairVol(L/R) (Port#,TID,LU-M), Seq#, LDEV#..P/S, Status,
Seq#, P-LDEV# M
oradb   oradev1(L)   (CL1-D , 3, 0-0) 30052  266....SMPL ----,
----- ---- -
----- ---- (L)     (CL1-D , 3, 0-1) 30052  266....SMPL ----,
----- ---- -
----- ---- (L)     (CL1-D , 3, 0-2) 30052  266....SMPL ----,
----- ---- -
oradb   oradev1(L)   (CL1-D , 3, 0)   30052  266....P-VOL PAIR,
30053   268         -
```

Display examples for -fne

```
# pairedisplay -g nvme -IM10 -fne
Group PairVol(L/R) (Port#,TID, LU-M) ,Seq#,LDEV#.P/S,Status, Seq#,P-LDEV# M CTG CM EM
E-Seq# E-LDEV# NGUID
DM P CS D_Status ROW SNAP_MODE
nvme ndev001(L) (CL7-D-0,17,0-0 )500054 257.SMPL ----,----- - - - - -
00500036000000000060e80800360101 - - 4 - NI -----
nvme ndev001(R) CL7-D-0,17,1-0 )500054 258.SMPL ----,----- - - - - -
00500036000000000060e80800360102 - - 4 - NI -----
```

Description of the pairedisplay command output:

Group

Group name (dev_group) as described in the configuration definition file

Pair Vol(L/R)

Paired volume name (dev_name) as described in the configuration definition file. (L) = local host; (R) = remote host

(P,T#,L#) (TrueCopy, TrueCopy Async, Universal Replicator, global-active device)

Port, TID, and LUN as described in the configuration definition file.

(Port#,TID,LU-M) (ShadowImage, Copy-on-Write Snapshot, Thin Image)

Port number, TID, LUN, and MU number as described in the configuration definition file.

Seq#

Serial number of the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

LDEV#

Logical device number

P/S

Volume attribute

Status

Status of the paired volume

For details about the status displayed, see the manual for each software product.

Available pair status of a Thin Image pair with the cascade attribute, the clone attribute, or the diff clone attribute when the Thin Image pair and the ROW column display "Y". If the ROW column does not display "Y", SMPL is always displayed as the pair status.

The following describes the pair status displayed when the Thin Image pair and the ROW column display "Y".

Thin Image pair status	Pair status displayed in Pairedisplay	
	Status	D_Status
SMPL	SMPL	SMPL
COPY(PD)	COPY	CPPD
PAIR	PAIR	PAIR
PFUL	PFUL (PAIR is displayed if the <code>-fc</code> option is not specified.)	PFUL
PSUS(SP)	COPY	PSUP
PSUS	P-VOL: PSUS	P-VOL: PSUS

Thin Image pair status	Pair status displayed in Pairedisplay	
	Status	D_Status
	S-VOL: SSUS	S-VOL: SSUS
PFUS	PFUS (PSUS or SSUS is displayed if the <code>-fc</code> option is not specified.)	PFUS
COPY(RS)	COPY	CPRS
RCPY	RCPY	RCPY
COPY(SP)	COPY	CPPS
SMPL(PD)	SMPL	SMPP
PSUE	PSUE	PSUE

Fence (TrueCopy, TrueCopy Async, Universal Replicator, global-active device)

Fence level


%

Progress rate of the copy operation or the pair synchronization.

The following tables describe the meanings of the values displayed in %.


Vol.	TrueCopy Async status			TrueCopy Sync/ GAD status			ShadowImage, Copy-on-Write Snapshot, or Volume Migration status			
	Copy	Pair	Other	Copy	Pair	Other	Copy	Pair	Pvol_psu Svol_cop y	Other
P-VOL	CR	SF	BMP	CR	BMP	BMP	CR	CR	BMP	CR
S-VOL	-	SF	BMP	-	BMP	BMP	CR	CR	CR	CR

Vol.	TrueCopy Async status			TrueCopy Sync/ GAD status			ShadowImage, Copy-on-Write Snapshot, or Volume Migration status			
	Copy	Pair	Other	Copy	Pair	Other	Copy	Pair	Pvol_psu Svol_cop y	Other


Note:

- CR: Copy operation progress rate (pair synchronization rate)
- SF: Sidefile usage rate for each consistency group in the caches of both the P-VOL and the S-VOL when the entire sidefile is 100%.
The following is a mathematical expression using high water mark (HWM) when the sidefile space is 100 %:
$$\text{HWM}(\%) = \text{HWM}(\%)/\text{Sidefile space (30 through 70)} \times 100$$
- BMP : BITMAP synchronization rate on both the P-VOL and the S-VOL.
- When Thin Image Advanced is used, if D_STAT is COPY, CPYP, RCPY, or PSUP, it indicates the operation progress rate. If D_STAT is a different status, this value is not valid and zero or a hyphen (-) is displayed.

Universal Replicator				
Volume	Copy	Pair	PSUS/SSUS (PJNS/SJNS)	Other
P-VOL	CR	JF	BMP	BMP
S-VOL	-	JF	BMP	BMP


Note:

- CR: Copy operation progress rate (pair synchronization rate).
- JF: Usage rate of the current journal data when the journal data space is 100%.
- BMP : BITMAP synchronization rate on both the P-VOL and the S-VOL.

P-LDEV#

LDEV number of the partner volume of the pair

M

- When the `-fm` option is not specified:

If a P-VOL is in the PSUS status:

- M='W' shows that S-VOL is in the suspended status with read/write enabled through the pairsplit.
- M='-' (hyphen) shows that S-VOL is in the suspended status with read only through the pairsplit. However, a hyphen (-) is displayed only when a pair operation is performed by setting the S-VOL to read-only, such as with the `pairsplit -r` command. If the S-VOL is set to read-only by an operation other than the pair operation, such as with Data Retention Utility, the results of such an operation are not reflected in this column.

If an S-VOL is in the SSUS status:

- M='W' shows that S-VOL has been altered since entering SSUS state.
- M='-' (hyphen) shows that S-VOL has NOT been altered since entering SSUS state.

If the pair status is COPY, RCPY, PAIR, or PSUE:

- M='N' shows that its volume is read-disabled through the `paircreate -m noread`.
- M=other than 'N' shows that a hyphen (-) is displayed.

- When the `-fm` option is specified:

- M='T' shows that the pairs are TC/UR/GAD pairs, and the management unit for their differential data is "track".
- M='C' shows that the pairs are TC/UR/GAD pairs, and the management unit for their differential data is "cylinder".
- M='U' shows that the system is unable to display differential data, or the pairs are SI/TI/Copy-on-Write Snapshot pairs.
- M='-' shows that there are no pairs.

CTG

Displays the consistency group ID.

If the snapshot is created by using the `raidcom add snapshot` command, the displayed consistency group ID is not correct. To confirm the status of the snapshot that is created by using the `raidcom add snapshot` command, use the `raidcom get snapshot` command.

CM

Displays the copy mode.

- N: Non snapshot
- S: Thin Image (TI only) snapshot with the cascade attribute disabled. In the SMPL status, the volume is a V-VOL used in Thin Image.
- C: Volume Migration
- R: Thin Image Advanced snapshot. This information is displayed only when the ROW column displays "Y".
- A: Thin Image (TI only) snapshot with the cascade attribute enabled. This information is displayed only when the ROW column displays "Y".

EM

Displays the external connection mode.

- H: Mapped E-lun that is hidden from the host
- V: Mapped E-lun that is visible to the host
- ' - ': Unmapped to the E-lun
- BH: Mapped E-lun that is hidden from the host, however the LDEV is blocked.
- BV: Mapped E-lun that is visible to the host, however the LDEV is blocked.
- B: Unmapped to the E-lun, however the LDEV is blocked.

E-Seq#

Displays the production (serial) number of the external LUN. 'Unknown' is displayed as a hyphen (-).

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

E-LDEV#

The LDEV# of the external LUN. 'Unknown' is displayed as a hyphen (-).

NGUID

Displays the namespace identifier (NGU ID). If NGUID cannot be calculated, 'Unknown' is displayed.

DM

Displays the differential data management method.

- S: The differential data is managed using a shared memory.
- D: The differential data is managed using a drive.
- - (hyphen): This information is not available.

CS

Displays the copy pace. A hyphen (-) is displayed if the storage system does not support displaying information or the information is not valid.

- TrueCopy, Universal Replicator, global-active device:
 - When the track size is 1, the copy pace is slow.
 - When the track size is 2 or 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.
- ShadowImage
 - When the track size is 1 or 2, the copy pace is slow.
 - When the track size is 3, the copy pace is medium.
 - When the track size is 4 or more, the copy pace is fast.

D_Status

Displays the detailed pair status. This status information might be changed or added in the future for further details.

The following items are displayed in D_Status. For other statuses, D_Status displays the same status information as that displayed when the `-fc` option is specified.

A hyphen (-) is displayed if the storage system does not support displaying information or the information is not valid. This information is only supported in Universal Replicator, global-active device, ShadowImage, and Thin Image.

- CPPD: A ShadowImage pair or a Thin Image pair is being split. This information is displayed only when the ROW column displays "Y".
- INIT: global-active device is about to start an initial copy operation.
- CPRS: A ShadowImage pair or a Thin Image pair is being resynchronized. This information is displayed only when the ROW column displays "Y".
- SMPP: A ShadowImage pair or a Thin Image pair is being deleted. This information is displayed only when the ROW column displays "Y".
- DELI: A Universal Replicator pair is being deleted.
- SUSI: A Universal Replicator pair is being suspended.
- HOLD: A Universal Replicator pair for delta resync is ready for delta resync.
- HLDI: A Universal Replicator pair for delta resync is being transitioned to the HOLD status.
- HLDE: An error occurred in a Universal Replicator pair for delta resync in the HOLD status.
- PSUP: A Thin Image pair is being split. This information is displayed only when the ROW column displays "Y".
- CPYP: The snapshot data of a Thin Image pair is being relocated. This information is displayed only when the ROW column displays "Y".
- CPYD: The differential data of a Thin Image pair is being copied. This information is displayed only when the ROW column displays "Y".

ROW

Displays whether the Thin Image Advanced pair operations are available by using the **pairXX** command.

- NI: The Thin Image Advanced pair operation cannot be performed because the instance started when HORCM_SUPROW is not listed in the configuration definition file or HORCM_SUPROW environment variables are not set.
- NM: The instance started when HORCM_SUPROW is listed in the configuration definition file or HORCM_SUPROW environment variables are set, however, the DKCMAIN microcode does not support the Thin Image Advanced pair operations by using the **pairXX** command.
- Y: The Thin Image Advanced pair operations are available by using the **pairXX** command. To view a Thin Image (TI only) pair or the Thin Image Advanced pair with the cascade attribute, the clone attribute, or the diff clone attribute, use the **pairXX** command.

The operations above are available if all the following conditions are met:

- Make sure to start the instances for both the P-VOL and the S-VOL with the same settings when HORCM_SUPROW is listed in the configuration definition file or HORCM_SUPROW environment variables are set.
- The DKCMAIN microcode supports the Thin Image Advanced pair operations by using the **pairXX** command.
- - (hyphen): This information is not valid.

SNAP_MODE

Displays the same information as "MODE" in the **raidcom get snapshot -key details** command. This information is available only when the ROW column displays "Y". This information is not valid if the pair to be displayed is not a Thin Image pair, or if the ROW column displays a value that is not "Y".

Example of pairdisplay -v pidb

```
# pairdisplay -g VG01 -v pidb
PID POLS U(%) LCNT SSCNT Available(MB) Capacity(MB) Snap_Used(MB)
TL_CAP(MB) BM TR_CAP(MB) RCNT Seq# Num LDEV# W(%) H(%) STIP
VCAP(%)
TYPE PM PT POOL_NAME D_PM D_PM(%) EFFECTIVE(MB)

001 POLN 0 11001 11001 46998 46998 0
2432398 NB 0 0 300050 1 0 70 80 YES - OPEN
N HDP dp_ti_pool N - 46998
001 POLN 0 11001 11001 46998 46998 0
2432398 NB 0 0 300050 1 0 70 80 YES - OPEN
N HDP dp_ti_pool N - 46998
```

Example of pairdisplay -v pidb -fp

```
# pairdisplay -g VG01 -v pidb -fp
PID POLS U(%) LCNT SSCNT Available(MB) Capacity(MB) Snap_Used(MB)
```

```

TL_CAP(MB) BM TR_CAP(MB) RCNT Seq# Num LDEV# W(%) H(%) STIP VCAP(%)
TYPE PM PT POOL_NAME D_PM D_PM(%) EFFECTIVE(MB)

001 POLN 0 11001 11001 46998 46998 -
2432398 NB 0 0 300050 1 0 70 80 YES -
OPEN N HDP dp_ti_pool N - 46998
001 POLN 0 11001 11001 46998 46998 -
2432398 NB 0 0 300050 1 0 70 80 YES -
OPEN N HDP dp_ti_pool N - 46998

```

Output of the pairedisplay command:

PID

Pool ID

POLS

Status of the pool

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is blocked due to a failure.) In this status, the pool information cannot be displayed.

U(%)

Usage rate of the pool

Displays the actual usage rate of the pool when the -fp option is specified.

LCNT

Total number of Dynamic Provisioning virtual volumes mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

SSCNT

Total number of snapshot data items mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

Available(MB)

Available capacity (megabytes) for volume data in the pool

Displays the actual pool capacity when the -fp option is specified.

Capacity(MB)

Total capacity (megabytes) of the pool

Displays the actual pool capacity when the -fp option is specified.

Snap_Used(MB)

Capacity used for Thin Image data (megabytes). If the value is less than 1 MB, it is rounded up. A hyphen (-) is displayed if the information is not available for this pool.

If you entered the -fp option, a hyphen (-) is displayed.

TL_CAP(MB)

Total capacity (megabytes) of all DP-VOLs and Thin Image pairs mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

BM

I/O Blocking Mode of the pool:

- PF (Pool Full): If the pool is full, you cannot read from or write to the target DP-VOL. If the pool VOL is blocked, you can read from or write to the target DP-VOL.
- PB (Pool vol Blockade): If the pool VOL is blocked, you cannot read from or write to the target DP-VOL. If the pool is full, you can read from or write to the target DP-VOL.
- FB (Full or Blockade): If the pool is full or pool VOL is blocked, you cannot read from or write to the target DP-VOL.
- NB (No Blocking): If the pool is full or pool VOL is blocked, you can read from or write to the target DP-VOL.
- - (Not supported): The configuration does not support the I/O Blocking Mode.

TR_CAP(MB)

Sum of the pool capacities reserved for the volumes for which Full Allocation is enabled or reserved for Proprietary Anchor. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

RCNT

Number of volumes for which Full Allocation is enabled and mapped to the pool. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

Seq#

Serial number of the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Num

Number of LDEVs belonging to the pool

LDEV#

First number of LDEVs in the pool. "65535 (ffff)" is displayed while the pool is being created.

W(%)

Threshold value of WARNING set for the pool. A hyphen (-) is displayed if the information is not available for this pool.

H(%)

Threshold value set for the pool as the high water mark

STIP

Setting for suspending Thin Image pairs when the high water mark threshold is exceeded:

- YES: Thin Image pairs are suspended.
- NO: Thin Image pairs are not suspended.
- - (hyphen): The information is not available for this pool.

VCAP(%)

The percentage of the subscription limit of V-VOL and Thin Image pair to the pool capacity:

- UNLIMITED: Unlimited.
- - (hyphen): The information is not available for this pool.

For VSP G130, G/F350, G/F370, G/F700, G/F900, a hyphen (-) indicating invalid is displayed.

TYPE

Platform type of the pool:

- OPEN: pool for open systems
- M/F: pool for mainframe systems

PM

Pool status:

- N: Normal status
- S: Shrinking or rebalancing
- NT: The pool for Thin Image is in the normal status.
- ST: The pool for Thin Image is shrinking or rebalancing.

Check **D_PM** to view the detailed status of the pool.

PT

Pool type:

- HDP: Pool for Dynamic Provisioning
- HDT: Pool for Dynamic Tiering
- RT: Pool for active flash
- TI: Pool for Thin Image
- CW: Pool for Copy-on-Write Snapshot
- DM: Pool for Dynamic Provisioning with the data direct mapping attribute

POOL_NAME

Pool name

D_PM

Processing status of the pool

If multiple statuses coexist, the statuses are displayed separated by commas (,). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

- N: Processing of the pool is not running.
- WS: The pool is waiting to shrink.
- S: The pool is shrinking.
- WRB: The pool is waiting to rebalance.

- RB: The pool is rebalancing.
- E: The pool is being expanded due to DDP expansion.
- PE: The pool expansion due to DDP expansion has been suspended.
- - (hyphen): The information is not available for this pool.

D_PM(%)

Percentage of progress of the processing status of the pool

If multiple statuses coexist, the percentage of progress for each status is displayed separated by a comma (.). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported or if the percentage of progress is not valid.

EFFECTIVE(MB)

Effective capacity of the pool

A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

paircurchk* (for TrueCopy/global-active device)

The **paircurchk** command is used to check the currency of the TrueCopy secondary volume(s) by evaluating the data consistency based on pair status and fence level.

The following table specifies the data consistency for each possible state of a TrueCopy volume. A paired volume or group can be specified as the target of the **paircurchk** command. The **paircurchk** command assumes that the target is an S-VOL. If the **paircurchk** command is specified for a group, the data consistency of each volume in the group is checked, and all inconsistent volumes are found in the execution log file and displayed. The **paircurchk** command is also executed as part of the **horctakeover** command .

Object Volume			Currency	
Attribute	Status	Fence	paircurchk	SVOL_takeover
SMPL	-	-	To be confirmed	-
P-VOL	-	-	To be confirmed	-
S-VOL	COPY	Data	Inconsistent	Inconsistent
		Status		
		Never		
		Async	Inconsistent	Inconsistent
	PAIR	Data	OK	OK
		Status	OK	OK

Object Volume			Currency	
Attribute	Status	Fence	paircurchk	SVOL_takeover
		Never	To be analyzed	To be analyzed
	PAIR	Async	To be analyzed	OK (assumption)
	PFUL		To be analyzed	OK (assumption)
	PSUS	Data	Suspected	Suspected
		Status	Suspected	Suspected
		Never	Suspected	Suspected
	PSUS	Async	Suspected	Suspected
	PFUS		Suspected	Suspected
	PSUE PDUB	Data	OK	OK
		Status	Suspected	Suspected
		Never	Suspected	Suspected
		Async	Suspected	OK (assumption)
	SSWS	Data	Suspected	-
		Status	Suspected	
		Never	Suspected	
		Async	Suspected	

Legend:

- To be confirmed = It is necessary to check the object volume, since it is not the secondary volume.
- Inconsistent = Data in the volume is inconsistent because it was being copied.
- To be analyzed = It cannot be judged from the status of the secondary volume whether data is consistent or not. It is OK if the status of the primary volume is PAIR. It is Suspected if the status is PSUS or PSUE.
- Suspected = The primary volume data and secondary volume data are not consistent.
- OK (assumption) = Mirroring consistency is not assured, but as S-VOL of TrueCopy Async or Universal Replicator, the sequence of write data is ensured.

Syntax

```
paircurchk { -h | -q | -z[x] | -I[H][M][instance#] or
             -I[TC][SI][instance#]
```

```
| -g <group> | -d <pair Vol> | -d[g] <raw_device> [MU#] [-n]  
| -d[g] <seq#> <LDEV#> [MU#] | -nomsg }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **paircurchk** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies a group name defined in the configuration definition file. The command is executed for the specified group unless the `-d <pair Vol>` option is specified.

-d <pair Vol>

Specifies paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volume.

-d[g] <raw_device> [MU#][-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-nomsg

Suppresses messages to be displayed when this command is executed. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

Returned values

The paircurchk command sets the following returned values during exit allowing you to check the execution results.

- **Normal termination:**
 - **0:** Data is consistent
- **Abnormal termination:**
 - **other than 0:** Refer to the error code for error details.

Error codes

The following table lists and describes the error codes for the **paircurchk** command. Unrecoverable errors are fixed and is not resolved, even after re-executing the command. If the command fails, the detailed status is logged in the CCI command log (\$HORCC_LOG), even if the user script has no error handling.

Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_VOLCUR	S-VOL currency error	Check volume list to see if an operation was directed to the wrong S-VOL.	225

Examples

The following shows an example of the **paircurchk** command for a group and the resulting display of inconsistent volumes in the specified group.

```
# paircurchk -g oradb
```

```
Group  Pair vol  Port  targ# lun# LDEV#  Volstatus  Status  Fence
To be...
oradb  oradb1    CL1-A  1     5     145    S-VOL      PAIR     NEVER   Analyzed
oradb  oradb2    CL1-A  1     6     146    S-VOL      PSUS     STATUS  Suspected
```

Output of the **paircurchk** command:

- **Group:** Displays the group name (dev_group) defined in the configuration definition file.
- **Pair vol:** Displays the pair volume name (dev_name) for the group defined in the configuration definition file.
- **Port targ# lun#:** Displays the port ID, TID, and LUN defined in the configuration definition file.
- **LDEV#:** Displays the LDEV number from the storage system.
- **Volstat:** Displays the volume attributes (P-VOL, S-VOL, SMPL).
- **Status:** Displays the status of a pair volume. Displays the same status as the status displayed when the **-fc** option is not specified in the **pairdisplay** command.
- **Fence:** Displays the fence level of a pair volume.
- **To be...:** Displays S-VOL's currency which was confirmed by mirror consistency check for S-VOL.

pairsyncwait*

The **pairsyncwait** command is used to confirm data consistency between the TrueCopy Async/Universal Replicator P-VOL and S-VOL by confirming that required writing was stored in the DFW area of RCU, and confirming whether the last writing just before this command reached the RCU DFW area.


This command obtains the MCU sidefile sequence number (the sequence number of the P-VOL within the consistency group ID), and the sequence number of RCU DFW in the CTG ID for <group> or <raw_device> specified by **pairsyncwait**. Both sequence numbers are the latest when the command is received. After that, this command compares the MCU and RCU sequence numbers which are the latest when the command is received at a certain interval.

If the RCU sequence number is over the value of the MCU sequence number within the term specified by **pairsyncwait**, this command displays the return code 0, indicating completion of synchronization. The **-nowait** option shows the latest sequence numbers (Q-marker) of the MCU P-VOL and CTG ID that are obtained when the command is received. The marker is a 10-digit hexadecimal number.

When a client issues the **pairsyncwait** command, this command is placed on the queue buffer for waiting in the HORCM daemon as a command request. HORCM gets the latest sequence number obtained from the MCU sidefile when the command is received and the sequence number whose block was transferred and stored in the DFW area of RCU with data consistency, and compares the latest sequence number of MCU sidefile with the sequence number of the RCU DFW area within the term. HORCM replies return codes to this command, when the write of MCU sidefile was stored in RCU DFW area.

Using this function, a client can confirm that a commit() has reached the remote site, and also the backup utility on a remote site can split the cascaded ShadowImage volumes (TrueCopy Async/Universal Replicator to TrueCopy Async/ShadowImage/Universal Replicator) without splitting TrueCopy Async/Universal Replicator.

More robust systems need to confirm the data consistency between the TrueCopy Async/Universal Replicator P-VOL and S-VOL. In DB operations (for example, Oracle), the commit() of DB transaction (see the following figure showing the synchronization for TrueCopy Async/Universal Replicator) is needed to confirm that a last writing for the commit() on a local site reached the remote site by using CCI-unique API command.



Caution: If an extended consistency group is used in Universal Replicator, the **pairsyncwait** command cannot be used.

Syntax

```
pairsyncwait{ -h | -q | -z[x] | -I[H][M][instance#] or -I[TC][SI][instance#] | -g
<group> | -d <pair Vol> | -d[g] <raw_device> [MU#] [-n] | -d[g] <seq#> <LDEV#> [MU#] | -
m <marker> | -t <timeout> | -nowait | -nomsg | -fq }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **pairsyncwait** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies a group name defined in the configuration definition file. The command is executed for the specified group unless the `-d <pair Vol>` option is specified.

-d <pair Vol>

Specifies paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volume.

-d[g] <raw_device> [MU#][-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-m <marker>

Specifies the sequence number of the MCU P-VOL in 10-digit hexadecimal, called the Q-marker. If the application gets Q-marker as the result of execution of **pairsyncwait** because of timeout or '**-nowait**', the application can reconfirm the completion of Async transfer by using **pairsyncwait** with Q-marker. If the application does not specify Q-marker, CCI uses the latest sequence number obtained when CCI receives **pairsyncwait**. It is also possible to wait for the completion from S-VOL side with this option.

Q-Marker format: iissssssss, where ii = regeneration number of pair volume, and ssssssss = latest sequence number obtained when the command is received on the side of the P-VOL.

-t <timeout>

Specifies the timeout value (0 to 21474836) to wait for the completion of RCU DFW area. The unit is 100 ms. MCU gets the latest sequence number obtained from RCU when the command is received at regular interval.

-nowait

Gets the latest sequence numbers of the MCU P-VOL and CTG ID obtained when the command is received without waiting. When this option is specified, the latest sequence number of the MCU P-VOL obtained when the command is received is reported immediately, and `-t <timeout>option` is ignored.

-nomsg

Suppresses messages to be displayed when this command is executed from a user program. This option must be specified at the beginning of the command arguments.

-fq

Displays the number of remaining Q-Markers within the consistency group by adding 'QM-Cnt' to the last column. 'QM-Cnt' is shown as follows:

- When specifying '-nowait -fq'

'QM-Cnt' is shown as the number of remaining Q-Marker at this time within consistency group.
- When specifying '-nowait -m <marker> -fq'

'QM-Cnt' is shown as the number of remaining Q-Marker from the specified <marker> within consistency group.
- When specifying 'TIMEOUT' without '-nowait'

'QM-Cnt' is shown as the number of remaining Q-Marker at this timeout within consistency group.
- 'QM-Cnt' is shown as '-', if the status for Q-Marker is invalid (that is, status is 'BROKEN' or 'CHANGED').

Example:

```
# pairsyncwait -g oradb -nowait -fq
```

UnitID	CTGID	Q-Marker	Status	Q-Num	QM-Cnt
--------	-------	----------	--------	-------	--------

0	3	01003408ef	NOWAIT	2	120
---	---	------------	--------	---	-----

```
# pairsyncwait -g oradb -nowait -m 01003408e0 -fq
```

UnitID	CTGID	Q-Marker	Status	Q-Num	QM-Cnt
--------	-------	----------	--------	-------	--------

0	3	01003408e0	NOWAIT	2	105
---	---	------------	--------	---	-----

```
# pairsyncwait -g oradb -t 50 -fq
```

UnitID	CTGID	Q-Marker	Status	Q-Num	QM-Cnt
--------	-------	----------	--------	-------	--------

0	3	01003408ef	TIMEOUT	2	5
---	---	------------	---------	---	---

Restriction

Specified <group> volume must be P-VOL with status PAIR. Other cases reply with error (EX_INVVOL). It is possible to issue **pairsyncwait** from S-VOL side with -m <marker>.

Returned values

The `pairsyncwait` command sets the following returned values during exit allowing you to check the execution results.

- When the `-nowait` option is specified
 - **Normal termination:**
 - **0:** The status is NOWAIT.
 - **Abnormal termination:**
 - **other than 0 to 127:** Refer to the error code for error details.
- When the `-nowait` option is not specified
 - **Normal termination:**
 - **0:** The status is DONE (completion of synchronization).
 - **1:** The status is TIMEOUT (timeout).
 - **2:** The status is BROKEN (Q-marker synchronized process is rejected).
 - **3:** The status is CHANGED (Q-marker is invalid due to resynchronize).
 - **Abnormal termination:**
 - **other than 0 to 127:** Refer to the error code for error details.

Error codes

Unrecoverable errors are fixed and not resolved, even after re-executing the command. If the command failed, the detailed status is logged in the CCI command log (\$HORCC_LOG), even if the user script has no error handling.

Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_INVVOL	Invalid volume status	Confirm pair status using <code>pairdisplay -l</code> .	222

Examples

If the latest sequence number is not synchronized within 10 seconds when the command is received:

```
# pairsyncwait -g oradb -t 100
UnitID      CTGID      Q-Marker      Status      Q-Num
0           3         01003408ef    TIMEOUT     3
```

If the latest sequence number is synchronized within 10 seconds when the command is received (restored in S-VOLs):

```
# pairsyncwait -g oradb -t 100
```

UnitID	CTGID	Q-Marker	Status	Q-Num
0	3	01003408ef	DONE	2

When you check the synchronization status for a specific checkpoint, specify the Q-Maker value acquired by **-nowait** with the **-m** option, and then check the Status.

If you obtain the latest sequence number when the command is received:

```
# pairsyncwait -g oradb -nowait
```

UnitID	CTGID	Q-Marker	Status	Q-Num
0	3	02006811DE	NOWAIT	2

If the specified sequence number (02006811DE) is not synchronized within 10 seconds:

```
# pairsyncwait -g oradb -t 100 -m 02006811DE
```

UnitID	CTGID	Q-Marker	Status	Q-Num
0	3	0200682000	TIMEOUT	3

If the specified sequence number (02006811DE) is synchronized within 10 seconds:

```
# pairsyncwait -g oradb -t 100 -m 02006811DE
```

UnitID	CTGID	Q-Marker	Status	Q-Num
0	3	0200682F8D	DONE	0

Description of the **pairsyncwait** command output:

UnitID

Unit ID in the case of multiple storage system connection

CTGID

Consistency group ID within Unit ID

Q-Marker

The latest sequence # of MCU PVol (Marker) when the command is received.

Status

The status after the execution of command.

Displays the same status as Status displayed when the **-fc** option is not specified in the **pairedisplay** command.

Q-Num

The number of process queues to wait for synchronization within the consistency group (CTG).

QM-Cnt

The number of remaining Q-Markers within consistency group of the Unit. TrueCopy Async/Universal Replicator sends a token called 'dummy recordset' at regular intervals, therefore QM-Cnt always shows '2' or '3' even if host has NO writing.

Arithmetic expression for determining the remaining data in a consistency group (CTG):

$$\text{Remaining data in CTG} = \text{sidefile capacity} \times \text{sidefile percentage} \div 100$$

Sidefile percentage is the rate showed to '%' column with 'PAIR' state by the **pairedisplay** command. Sidefile capacity is the capacity within 30% to 70% of the cache setting as the sidefile.

Arithmetic expression for determining the average data per Q-Marker in a consistency group (CTG):

$$\text{Data per Q-Marker} = \text{Remaining data in CTG} \div \text{QM-Cnt}$$

horctakeover*

The **horctakeover** command is a scripted command for executing several takeover operations. The **horctakeover** command checks the specified volume's or group's attributes (**paircurchk**), decides the takeover function based on the attributes, executes the chosen takeover function, and returns the result. The four takeover functions designed for HA software operation are (see **horctakeover** command functions in the *User and Reference Guide*): takeover-switch, swap-takeover, PVOL-takeover, and SVOL-takeover. A paired volume or a group can be specified as the target of the TrueCopy takeover command. If SVOL-takeover is specified for a group, the data consistency check is executed for all volumes in the group, and all inconsistent volumes are found in the execution log file and displayed (same as **paircurchk** command). You can execute the **horctakeover** command only for remote copy pairs.

Syntax

```
horctakeover { -h | -q | -z[x] | -I[H][M][instance#] or
               -I[TC][SI][instance#] | -g <group> | -d <pair Vol>
               | -d[g] <raw_device> [MU#] [-n] | -d[g] <seq#> <LDEV#> [MU#]
               | -S | -l | -t <timeout> | -nomsg }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **horctakeover** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies a group name defined in the configuration definition file. The command is executed for the specified group unless the `-d <pair Vol>` option is specified.

-d <pair Vol>

Specifies paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volume.

-d[g] <raw_device> [MU#][-n]

Searches whether the specified `raw_device` is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of `-g <group>` option. If the specified `raw_device` is contained in two or more groups, the command is executed for the first group.

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of `'-g <group>'` option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The `<seq#> <LDEV#>` values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-S

Selects and executes SVOL-takeover. The target volume of the local host must be an S-VOL. If this option is specified, then the following '-l' option is invalid.

-l

Enables read and write to the primary volume(s) by a local host only without a remote host, and executes PVOL-takeover when the primary volume cannot be used because it is fenced (fence = DATA or STATUS, state = PSUE or PDUB, or PSUE or PDUB volume is contained in the group). If the primary volume can be accessed, nop-takeover is executed. The target volume of the local host must be a P-VOL.

-t <timeout>

Can be specified for asynchronous pairs only, ignored for synchronous pairs. Specifies the maximum time to wait in seconds (0 to 21474836) for swap-takeover and SVOL-takeover operation to synchronize the P-VOL and S-VOL. If this timeout occurs, the **horctakeover** command fails with EX_EWSTOT. To avoid timeout, set this value less than or equal to the start-up timeout value of the HA control script. This option is required for asynchronous pairs.

-nomsg

Suppresses messages to be displayed when this command is executed. This option must be specified at beginning of a command argument. The command execution log is not affected by this option.

Returned values

The **horctakeover** command sets the following returned values during exit allowing you to check the execution results.

- **Normal termination:**

- **0:** Nop-takeover (no operation).
- **1:** Swap-takeover was successfully executed.
- **2:** SVOL-takeover was successfully executed.
- **3:** PVOL-SMPL-takeover was successfully executed.
- **4:** PVOL-PSUE-takeover was successfully executed. (This value depends on the microcode level.)
- **5:** SVOL-SSUS-takeover was successfully executed. (This value depends on the microcode level.)

- **Abnormal termination:**

- **other than 0-5:** Refer to the error code for error details.

Error codes

The following table lists and describes the error codes for the **horctakeover** command. Unrecoverable errors are fixed and not resolved, even after re-executing the command. If the command fails, the detailed status is logged in the CCI command log (\$HORCC_LOG), even if the user script has no error handling.

Category	Error Code	Error Message	Recommended Action	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	Confirm status using pairdisplay command. Make sure all volumes in the group have the same fence level and volume attributes.	236

Category	Error Code	Error Message	Recommended Action	Value
	EX_INCSTG	Inconsistent status in group	Confirm pair status using pairdisplay .	229
	EX_EVOLCE	Pair volume combination error	Confirm pair status using pairdisplay , and change combination of volumes.	235
	EX_VOLCUR	S-VOL currency error	A currency check error on the S-VOL occurred. Verify the volume list that the command outputs and the table for the data consistency displayed by using the paircurchk command, and then determine whether to resume the processing from the S-VOL.	225
	EX_VOLCUE	Local volume currency error	Confirm pair status of the local volume.	224
	EX_VOLCRE	Local and remote volume currency error	Confirm pair status of remote and local volumes using pairdisplay command.	223
Timer (Recoverable)	EX_EWSTOT	Timeout waiting for specified status	If an error occurs when the horctakeover or pairsplit command is run for a consistency group, verify that the pair configuration and the pair status are in the status in which the performed pair operation can be accepted by referring to the user guide for each software. If the pair operation is acceptable, increase the timeout value specified by the -t option, and then retry the operation. If this error occurs when a different command is run, verify that the operation and the timeout value are valid.	233

raidscan*

The **raidscan** command displays configuration and status information for the specified port/TID(s)/LUN#/MU#. The information is acquired directly from the storage system (not the configuration definition file).

The operations in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair are available if all the following conditions are met. This command recognizes the pair state of the Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, and the Thin Image Advanced pair from the pair status displayed in Status in the **pairdisplay** command.

- Both the P-VOL and the S-VOL are used in a Thin Image (TI only) pair with the cascade attribute, the clone attribute, or the diff clone attribute, or a Thin Image Advanced pair.
- The ROW column in the **pairdisplay** command displays "Y".

Syntax

```
raidscan { -h | -q | -z[x] | -I[H][M][instance#] or
-I[TC][SI][instance#] | -p <port#> [hgrp]|-pd[g] <raw_device>[-n]
| -s <Seq#> | -t <targ> | -l <lun> | [ -f[xfgde] ] | -CLI
| -find[g] [op] [MU#] | [-g <group>][-n] | -pi <strings>
| -m <MU#> | -fw | -fn | -nguid[v] <nguid>}
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidscan** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:
a TrueCopy, Universal Replicator, or global-active device pair
- If the HORCC_MRCF environment variable is set:
a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:
an instance with no instance number is used.
- If the HORCMINST environment variable is set:
an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-p <port#> [hgrp]

Specifies the port ID of the port to be scanned.

- Valid ports are CL1-A to CL1-R and CL2-A to CL2-R (excluding CL1-I, CL1-O, CL2-I, CL2-O).
- Valid expanded ports are CL3-A to CL3-R (excluding CL3-I and CL3-O) and CLG-A to CLG-R (excluding CLG-I and CLG-O).

The port is not case sensitive (for example, CL1-A = cl1-a = CL1-a = cl1-A).

This option must be specified if '-find' or '-pd <raw_device>' option is not specified.

[hgrp] is specified to display only the LDEVs mapped to a host group on a port.

If only the <port> option is specified, the absolute LUN of CCI is displayed for LU#.

-pd[g] <raw_device> [-n]

Specifies the raw device name. This option is used to detect the serial number and the port name of the specified device, and display the LU information of the port for which the device serial number and whose port name are the same as those detected by the command device listed in the configuration definition file. This option must be specified if the '-find' option is not specified. If this option is specified, the following -s <seq#> option is not valid.

The -pdg option is used when displaying a LUN on the host view by finding a host group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-s <Seq#>

Specifies the serial number of the target storage system if the command devices on the multiple storage systems are listed in the configuration definition file and a unit ID is not specified for '-p <port>'. This option is used to specify the Seq# (serial#) of the storage system. This option scans the port specified by '-p <port>' option of the storage system which corresponds with the unit ID that searches the unit ID from seq#. If this option is specified, then the unit ID that is contained in '-p <port>' option is invalid.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-t <targ>

Specifies the target ID of the specified port. If this option is not specified, the command applies to all target IDs.

-l <lun>

Specifies the LUN of the specified target ID. If this option is not specified, the command applies to all LUNs. If this option is specified, the TID must also be specified.

-f or -ff

Specifies display of volume-type for a display column. If this is specified, -f[g] [d] option is invalid.

-fx

Displays the LDEV number in hexadecimal notation.

-fg

Specifies display of group_name for a display column. This option searches a group on the configuration definition file (local CCI instance) from the scanned LDEV, and displays a group_name when the scanned LDEV is contained in the group. If this option is specified, the `-f[f]` option is not allowed and the `-f[d]` option is invalid.

-fd

Displays the Device_File that was registered to the group of the HORCM in the output, based on the LDEV (as defined in local instance config. def. file). If this option is specified, `-f[f][g]` option is not allowed.

Specify ENABLE for HORCM_NVME in the configuration file. If ENABLE is not specified, a device file name for NVMe is not displayed.

-fe

Displays the serial# (E-Seq#) and LDEV# (E-LDEV#) of the external LUNs only mapped to the LDEV. If the external LUN mapped to the LDEV on a specified port does not exist, then this option does nothing. Also if this option is specified, `-f[f][g][d]` option is not allowed.

Display example:

```
# raidscan -p cl1-a-0 -fe -CLI
```

```
PORT# /ALPA/C TID# LU# Seq# Num LDEV# P/S Status Fence
E-Seq# E-LDEV#
```

```
CL1-A-0 ef 0 0 48 62468 2 256 SMPL - -
30053 17
```

```
CL1-A-0 ef 0 0 49 62468 2 272 SMPL - -
30053 23
```

```
CL1-A-0 ef 0 0 50 62468 1 288 SMPL - -
30053 28
```

-CLI

Specifies display for command line interface (CLI). This option displays to the same position that defined number of columns, and displays one header. The delimiters between columns are displayed as spaces or hyphens (-).

Display example:

Port#	TargetID#	Lun#	Seq#	Num	LDEV#	P/S	Status	Fence	P-
Seq#	P-LDEV#								

CL1-C		1	0	30053	1	274	SMPL	-	-
-		-							

CL1-C		2	2	30053	1	260	P-VOL	PAIR	NEVER
30053		268							

CL1-C		2	3	30053	1	261	P-VOL	PAIR	NEVER
30053		269							

-pi <strings>

Changes a strings via STDIN for -find option to '<strings>'. If this option is specified, the -find option is ignored a raw device file provided via STDIN, and <strings> is used as input. A <strings> must be specified within 255 characters.

-m <MU#>

This option is used for displaying only the specified mirror descriptor. If you want to display all mirror descriptor, specify '-m all' for displaying all MUs.

If this option is omitted and the command execution environment is HORC, the MU#0 information is displayed.

If this option is omitted and the command execution environment is HOMRCF, the MU#0-2 information is displayed.

-fw

Specify this option to display the NAA identifier (LUN WWN) for the LU.

-find [op] [MU#] [-n]

Executes the specified [op] using a raw device file provided via STDIN. If the -pi <strings> option is specified, this option does not use a strings via STDIN, and -pi <strings> is used as input.

Restrictions: Special files via STDIN specify the following values:

- HP-UX:
 - /dev/rdisk/*
 - /dev/rdisk/disk*
- Solaris:
 - /dev/rdisk/*s2
 - c*s2

- Linux:
 - /dev/sd...
 - /dev/rd...
 - /dev/raw/raw*
 - /dev/nvme...
- zLinux:
 - /dev/sd...
 - /dev/dasd...
 - /dev/rd...
 - /dev/raw/raw*
- AIX:
 - /dev/rhdisk*
 - /dev/hdisk*
 - hdisk*
- DIGITAL or Tru64:
 - /dev/rrz*c
 - /dev/rdisk/dsk*c
 - /dev/cport/scp*
- DYNIX:
 - /dev/rdsk/sd*
 - sd* for only unpartitioned raw device

- IRIX64:
 - /dev/rdisk/*vol
 - /dev/rdisk/node_wwn/*vol/*
 - /dev/dsk/*vol
 - /dev/dsk/node_wwn/*vol/*
- Windows:
 - hdX-Y
 - \$LETALL
 - \$Volume
 - \$Phys
 - D:\Vol(Dms,Dmt,Dmr)X\DskY
 - \Vol(Dms,Dmt,Dmr)X\DskY
 - D:\DskX\pY
 - \DskX\pY

For further information on LDM volumes for Windows systems, see **Volume Discovery Function** in the *User and Reference Guide*.

Lines starting with '#' via STDIN are interpreted as comments.

Lines starting with 'quit' via STDIN are interpreted as exit.

This `-n` option is specified to use a device file for NVMe-oF. Specifying this option enables the use of both device files for NVMe-oF and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe-oF and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-find[g] [-n]

Displays the port, target ID, LUN (RAID storage system notation) and so on that was mapped for LDEV using a special file (raw device file) provided via STDIN. If target ID and LUN are Unknown for the target device file, you should start CCI without any description for `HORCM_DEV` and `HORCM_INST`, and should describe the shown port, target ID, and LUN for `HORCM_DEV`. This option also uses the `-fx` option to display the LDEV numbers in hexadecimal.

The `-findg` option is used to show a LUN on the host view by finding a host group.

This `-n` option is specified to use a device file for NVMe-oF. Specifying this option enables the use of both device files for NVMe-oF and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe-oF and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-find inst [-n]

Registers the device file name (raqw device file via STDIN) to the instance. The registered device becomes an allowed volume when the command device security is enabled. This command is run automatically when the instance starts, therefore there is usually no need to run this command. This option is also terminated to avoid wasteful scanning when the registration has been finished with based on HORCM. Therefore if HORCM does not need the registration any more, then nothing is done and it exits. This option can be used with '-fx' option to display LDEV numbers in hexadecimal.

Example for HP-UX:

```
# ioscan -fun | grep rdsd | raidscan -find inst
```

```
DEVICE_FILE Group PairVol PORT TARG LUN M SERIAL LDEV
```

```
/dev/rdsd/c0t3d0 oradb oradev1 CL1-D 3 0 - 35013 17
```

```
/dev/rdsd/c0t3d0 oradb oradev1 CL1-D 3 0 0 35013 17
```

```
/dev/rdsd/c0t3d0 oradb1 oradev2 CL1-D 3 0 1 35013 17
```

**Note:**

When multiple device files share the same LDEV, the first device file is registered to the LDEV map table.

- **Group:** Displays the group name (dev_group) defined in the configuration definition file.
- **PairVol:** Displays the paired volume name (dev_name) within the group defined in the configuration definition file.
- **PORT:** Displays the port number (port#) defined in the configuration definition file.
- **TARG:** Displays the target ID (TargetID) defined in the configuration definition file.
- **LUN:** Displays the LUN (LU#) defined in the configuration definition file.
- **M:** Displays the MUN (MU#) defined in the configuration definition file. MU# for TrueCopy/GAD are shown as '-'. MU# for ShadowImage/Copy-on-Write Snapshot/Thin Image are shown as '0', '1', '2'.
- **SERIAL:** Displays the production (serial#) number of the RAID storage system.
 - The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
 - The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).
- **LDEV:** Displays the LDEV# within the RAID storage system.

This `-n` option is specified to use a device file for NVMe-oF. Specifying this option enables the use of both device files for NVMe-oF and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe-oF and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-find verify [MU#] [-n]

Displays the relationship between group on the configuration definition file and Device_File registered to the LDEV map tables, based on the Device File name (raw

device file provided via STDIN). This option can be used with '-fx' and '-fd' options. This option is affected by the command execution environment (HORCC_MRCF).

```
# ioscan -fun | grep rdsd | raidscan -find verify
```

DEVICE_FILE	Group	PairVol	PORT	TARG	LUN	M	SERIAL	LDEV
-------------	-------	---------	------	------	-----	---	--------	------

/dev/rdsd/c0t3d0	oradb	oradev1	CL1-D	3	0	0	35013	17
------------------	-------	---------	-------	---	---	---	-------	----

/dev/rdsd/c0t3d1	oradb	oradev2	CL1-D	3	1	0	35013	18
------------------	-------	---------	-------	---	---	---	-------	----

/dev/rdsd/c0t3d2	-	-	-	-	-	0	35013	19
------------------	---	---	---	---	---	---	-------	----

```
# ioscan -fun | grep rdsd | raidscan -find verify 1 -fd
```

DEVICE_FILE	Group	PairVol	Device_File	M	SERIAL	LDEV
-------------	-------	---------	-------------	---	--------	------

/dev/rdsd/c0t3d0	oradb	oradev1	c0t3d0	1	35013	17
------------------	-------	---------	--------	---	-------	----

/dev/rdsd/c0t3d1	oradb	oradev2	Unknown	1	35013	18
------------------	-------	---------	---------	---	-------	----

/dev/rdsd/c0t3d2	-	-	-	1	35013	19
------------------	---	---	---	---	-------	----

**Note:**

If the device name is different between DEVICE_FILE and Device_File, then it shows shared LDEV among multiple device files.

- Group: Displays the group name (dev_group) defined in the configuration definition file.
- PairVol: Displays the paired volume name (dev_name) within the group defined in the configuration definition file.
- PORT: Displays the port number (port#) defined in the configuration definition file.
- TARG: Displays the target ID (TargetID) defined in the configuration definition file.
- LUN: Displays the LUN (LU#) defined in the configuration definition file.
- M: Displays the MUN (MU#) defined in the configuration definition file. MU# for TrueCopy/GAD are shown as '-'. MU# for ShadowImage/Copy-on-Write Snapshot/Thin Image are shown as '0', '1', '2'.
- Device_File: Displays Device_File which is registered to the LDEV map tables into the CCI.
- SERIAL: Displays the production (serial#) number of the RAID storage system.
 - The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
 - The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).
- LDEV: Displays the LDEV# within the RAID storage system.

This `-n` option is specified to use a device file for NVMe-oF. Specifying this option enables the use of both device files for NVMe-oF and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe-oF and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-find[g] conf [MU#] [-g <group>]

Displays the port, target ID, LUN associated with a special file (raw device file) provided via STDIN. The displayed port, target ID, and LUN are used to list in HORCM_DEV in the configuration definition file. If target ID & LUN are Unknown for the target device file, you must start CCI without any description for HORCM_DEV and HORCM_INST, and should be described the shown port, target ID, LUN for HORCM_DEV. This option can be used with the '-fx' option.

`[-g <group>]` specifies the group for 'dev_group' on horcm.conf. If omitted, the group applies 'VG' as default.

The `-findg` option is used when displaying a LUN on the host view by finding a host group.

```
# cat /etc/horcmperm.conf | raidscan -find conf 0 -g ORA
```

```
HORCM_DEV
```

#dev_group	dev_name	port#	TargetID	LU#	MU#
------------	----------	-------	----------	-----	-----

```
# /dev/rdisk/c23t0d0 SER = 61456 LDEV = 192 [ FIBRE FCTBL = 4 ]
```

ORA	ORA_000	CL2-J	0	0	0
-----	---------	-------	---	---	---

```
# /dev/rdisk/c23t0d1 SER = 61456 LDEV = 193 [ FIBRE FCTBL = 4 ]
```

ORA	ORA_001	CL2-J	0	1	0
-----	---------	-------	---	---	---

```
# /dev/rdisk/c23t0d2 SER = 61456 LDEV = 194 [ FIBRE FCTBL = 4 ]
```

ORA	ORA_002	CL2-J	0	2	0
-----	---------	-------	---	---	---

```
# /dev/rdisk/c23t0d3 SER = 61456 LDEV = 195 [ FIBRE FCTBL = 4 ]
```

ORA	ORA_003	CL2-J	0	3	0
-----	---------	-------	---	---	---

```
# ERROR [CMDDEV] /dev/rdisk/c23t0d7 SER = 61456 LDEV = 259 [ OPEN-3-  
CM ]
```

If the target device has shared an LDEV among multiple device files and an LDEV is displayed by another target device already, then its target device is suppressed as a comment as shown below:

```
# ERROR [LDEV LINK] /dev/rdisk/c24t0d3 SER = 61456 LDEV = 195 [FIBRE  
FCTBL = 4]
```

If the target device does not have a valid MU#, then its target device is suppressed as a comment as shown below:

```
# ERROR [INVALID MUN (2 < 1)] /dev/rdisk/c24t0d3 SER = 61456 LDEV = 195
[ OPEN-3 ]
```

If the target device is mixed with a different RAID TYPE, then its target device is suppressed as a comment as shown below:

```
# ERROR [MIXING RAID TYPE] /dev/rdisk/c24t0d3 SER = 61456 LDEV = 195
[ OPEN-3 ]
```

-find sync[d] [MU#] [-g <group>]

Flushes the system buffer associated to a logical drive which corresponds to a [-g <group>] through the KEY WORD (**\$Volume**, **\$LETALL**, **\$Physical**) provided via STDIN.

[-g <group>] specifies the group for 'dev_group' on horcm.conf. If this option is not specified, then flushes the system buffer associated to all groups for the local instance.

Example of flushing the system buffer associated to ORB group through \$Volume (Windows):

```
echo $Volume | raidscan -find sync -g ORB or
```

```
raidscan -pi $Volume -find sync -g ORB
```

```
[SYNC] : ORB ORB_000[-] -> \Dmt1\Dsk1 : Volume{bf48a395-0ef6-11d5-8d69-00c00d003b1e}
```

```
[SYNC] : ORB ORB_001[-] -> \Dmt1\Dsk2 : Volume{bf48a395-0ef6-11d5-8d69-00c00d003b1e}
```

```
[SYNC] : ORB ORB_002[-] -> \Dmt1\Dsk3 : Volume{bf48a395-0ef6-11d5-8d69-00c00d003b1e}
```


Example of flushing the system buffer associated to all groups for the local instance (Windows):

```
echo $Volume | raidscan -find sync          or
```

```
raidscan -pi $Volume -find sync
```

```
[SYNC] : ORA ORA_000[-] -> \Vol144\Dsk0 : Volume{56e4954a-28d5-4824-a408-3ff9a6521e5d}
```

```
[SYNC] : ORA ORA_000[-] -> \Vol145\Dsk0 : Volume{56e4954a-28d5-4824-a408-3ff9a6521e5e}
```

```
[SYNC] : ORB ORB_000[-] -> \Dmt1\Dsk1 : Volume{bf48a395-0ef6-11d5-8d69-00c00d003b1e}
```

```
[SYNC] : ORB ORB_001[-] -> \Dmt1\Dsk2 : Volume{bf48a395-0ef6-11d5-8d69-00c00d003b1e}
```

```
[SYNC] : ORB ORB_002[-] -> \Dmt1\Dsk3 : Volume{bf48a395-0ef6-11d5-8d69-00c00d003b1e}
```



Note:

1. The option cannot be specified the device object name as follows:
D:\Vol(Dms,Dmt,Dmr)X\DskY,\Vol(Dms,Dmt,Dmr)X\DskY
2. Sync executes the following behavior under any conditions:
 - If the logical drive which corresponds to a [-g <group>] is not open by any applications, then sync flushes the system buffer to a drive and makes the dismount state for this drive.
 - If the logical drive which corresponds to a [-g <group>] is already opened by any applications, then sync flushes only the system buffer to a drive. This is allowed to flush the system buffer before pairsplit with mounting the P-VOL (opening state), and indicates the behavior as [FLUSH] below: [FLUSH] : ORA ORA_000[-] -> \Vol144\Dsk0 : Volume{56e4954a-28d5-4824-a408-3ff9a6521e5d}

-fn

Specifies the namespace identifier (NGU ID).

-nguid <nguid>

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays the LDEV number and the serial number for the entered NGUID.

If you specify this option, you cannot use the `-s` option to specify the serial number. If you specify the NGUID of either of the following namespaces, the `-nguid` option might not be able to display the appropriate information. If so, the `-nguid` option might display a hyphen (-) for Seq# and LDEV#.

- NGUID of the namespace with the actual LDEV ID and the virtual LDEV ID that differ in the storage system
- NGUID of the namespace that belongs to the virtual storage machine created by a user

Use the `-nguidv` option to use namespaces for GAD pairs.

-nguidv <nguid>

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays the LDEV number and the serial number for the entered NGUID.

Specify this option instead of the `-nguid` option when you use namespaces for GAD pairs. When multiple LDEVs have the specified NGUID, the `-nguidv` option displays the information for all of them. When you specify this option, you cannot use the `-s` option to specify the serial number.

Returned values

None

Error codes

None

raidscan command examples for Fibre Channel ports (when the command execution environment is HORC)

```
# raidscan -p cll-r
```

```
PORT#/ALPA/C,TID#,LU#Num(LDEV#...)P/S,Status,Fence,LDEV#,P-Seq#P-LDEV#
CLl-R/ ce/15,15, 7 5(100,101..)P-VOL PAIR NEVER 100,5678 200
CLl-R/ ce/15,15, 6 5(200,201..)SMPL ---- ---- ---- ---- ----
```

```
# raidscan -p cll-r -f
```

```
PORT#/ALPA/C,TID#,LU#Num(LDEV#...)P/S,Status,Fence,LDEV#,Vol.Type
CLl-R/ ce/15,15, 7 5(100,101..)P-VOL PAIR NEVER 100,OPEN-3
CLl-R/ ce/15,15, 6 5(200,201..)SMPL ---- ---- ----OPEN-3
```

raidscan command examples for Fibre Channel ports (when the command environment is HOMRCF)

```
#raidscan -p CLl-A
```

```
PORT# /ALPA/C,TID#, LU#..Num(LDEV#....)...P/S, Status, LDEV#,P-Seq#,P-LDEV#
CLl-A /ef/ 5, 1, 0-0 .1(1).....SMPL ---- ----, ----
-----
CLl-A /ef/ 5, 1, 0-1 .1(1).....SMPL ---- ----, ----
-----
```

```
#raidscan -p CLl-A -CLI
```

```
PORT# /ALPA/C TID# LU# Seq# Num LDEV# P/S Status P-Seq# P-LDEV#
CLl-A ef 5 1 0 0 800013 1 1 SMPL - - -
CLl-A ef 5 1 0 1 800013 1 1 SMPL - - -
```

Example of the -find option for raidscan

```
# ls /dev/rdsd/* | raidscan -find
```

```
DEVICE_FILE      UID S/F  PORT TARG LUN SERIAL LDEV      PRODUCT_ID
/dev/rdsd/c0t0d4  0   S CLl-M    0   4  31168  216 OPEN-3-CVS-CM
/dev/rdsd/c0t0d1  0   S CLl-M    0   1  31168  117 OPEN-3-CVS
/dev/rdsd/clt0d1  -   - CLl-M    -   -  31170  121 OPEN-3-CVS
```

Example of the -fw option for raidscan

```
#raidscan -p cll-e-0 -l 0 -CLI -fw -IH
```

```
PORT#   /ALPA/C TID# LU#   Seq# Num LDEV# P/S   Status Fence LUN-WWN
CLL-E-0 cd 4    16    0   493017 1    768 P-VOL   PAIR NEVER
60060e80160164000001016400000893
#raidscan -p cll-e-0 -l 0 -fw -IM -CLI
PORT#   /ALPA/C TID# LU#   Seq# Num LDEV# P/S   Status LUN-WWN
CLL-E-0 cd 4    16    0 0 493017 1    768 SMPL   - 60060e80160164000001016400000893
CLL-E-0 cd 4    16    0 1 493017 1    768 SMPL   - 60060e80160164000001016400000893
CLL-E-0 cd 4    16    0 2 493017 1    768 SMPL   - 60060e80160164000001016400000893
```

Description of the raidscan command output:**Port#, ALPA/C, TID#, LU#**

Port ID, arbitrated loop physical address, target ID, LUN.

MU#

For ShadowImage, **raidscan** displays the MU# for each LUN (for example, LUN 7-0, 7-1, 7-2).

Num(LDEV#...):

Number of LDEVs and LDEV ID for a LUSE volume

P/S

Volume attribute

Status

Status of the paired volume. Displays the same status as the status displayed when the -fc option is not specified in the **pairedisplay** command.

Fence

Fence level (TrueCopy/global-active device only)

P-Seq#

Serial # of the storage system which contains the partner volume of the pair.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

P-LDEV#

LDEV number of the partner volume of the pair

Vol.Type

Logical unit (LU) type (for example, OPEN-V, OPEN-9)

Group

Group name (dev_group) as described in the configuration definition file

UID

Displays the unit ID for multiple storage system configuration. If UID appears as '-', the command device for HORCM_CMD is not found.

S/F

Displays "S" for a SCSI port or "F" for a port other than SCSI. Confirm the port type with the TYPE field which is the execution result of the **raidcom get port** command.

PORT

Displays the RAID storage system port number

TARG

Displays the target ID (that was converted by the fibre conversion table)

LUN

Displays the LUN (that was converted by the fibre conversion table). When the command execution environment is HOMRCF, MU# is displayed after LUN.

SERIAL

Displays the production (serial#) number of the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

LDEV

Displays the LDEV# within the RAID storage system

PRODUCT_ID

Displays product-id field in the STD inquiry page

LUN-WWN

NAA identifier (LUN WWN) of the LU. If values cannot be displayed for LUN-WWN, a hyphen (-) is displayed. For details about the cases where LUN WWNs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#).

Example of the -fn option for raidscan

```
#raidscan -p c11-e-0 -l 0 -CLI -fn -IH
PORT# /ALPA/C TID# LU# Seq# Num LDEV# P/S Status Fence NGUID
CL1-E-0 cd 4 16 0 493017 1 768 P-VOL PAIR NEVER 1234...
```

Description of each column in output example:**NGUID**

Displays the namespace identifier (NGU ID). If values cannot be displayed for NGUID, 'Unknown' is displayed. For details about the cases where NGU IDs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#).

Example of the LDEV number and the serial number when using the -nguid option to specify NGUID

The following output example shows the case when the **raidscan** command is run with the **-nguid** option specifying NGUID: 12345678901234567890123456789012. If the specified NGUID is not found, the command displays a hyphen (-) for Seq# and LDEV#.

```
#raidscan -nguid 12345678901234567890123456789012
NGUID                               Seq# LDEV# K
12345678901234567890123456789012 12345 2345 P
```

Description of each column in output example:**NGUID**

Displays the namespace identifier (NGU ID).

Seq#

Displays the serial number corresponding to the specified NGUID.

- If the LDEV corresponding to the specified NGUID does not belong to the virtual storage machine created by the user, the actual serial number of the storage system is displayed. If the target storage system is VSP 5000 series, the serial number added with 500,000 is displayed.
- If the LDEV corresponding to the specified NGUID belongs to the virtual storage machine created by the user, the virtual serial number of the virtual storage machine is displayed. If the model of the virtual storage machine is VSP 5000 series, the serial number added with 500,000 is displayed.
- If the LDEV corresponding to the specified NGUID is not found, a hyphen (-) is displayed.

LDEV#

Displays the LDEV number corresponding to the specified NGUID.

- If the LDEV corresponding to the specified NGUID does not belong to the virtual storage machine created by the user, the actual LDEV number in the storage system is displayed.
- If the LDEV corresponding to the specified NGUID belongs to the virtual storage machine created by the user, the virtual LDEV number is displayed.
- If the LDEV corresponding to the specified NGUID is not found, a hyphen (-) is displayed.

K

Indicates the type of the information displayed for Seq# and LDEV#.

P: The displayed information is the actual serial number of a storage system and an actual LDEV number.

V: The displayed information is the serial number of a virtual storage machine and a virtual LDEV number.

- (hyphen): The information is invalid.

Example of the LDEV number and the serial number when using the `-nguidv` option to specify NGUID

The following output example shows the case when the `raidscan` command is run with the `-nguidv` option specifying NGUID: 005000010000000000060e80800360064. If the specified NGUID is not found, no information is displayed.

```
#raidscan -nguidv 005000010000000000060e80800360064
NGUID                               Seq#  LDEV#  PHY_Seq#  PHY_LDEV#
005000010000000000060e80800360064 500001   100    500054    150
005000010000000000060e80800360064 500001   100    500079    200
```

Description of each column in output example:

NGUID

Displays the namespace identifier (NGU ID).

Seq#

Displays the serial number corresponding to the specified NGUID.

- If the LDEV corresponding to the specified NGUID does not belong to the virtual storage machine created by the user, the actual serial number of the storage system is displayed. If the target storage system is VSP 5000 series, the serial number added with 500,000 is displayed. In this case, the displayed serial number is the same as the number displayed for PHY_Seq#.
- If the LDEV corresponding to the specified NGUID belongs to the virtual storage machine created by the user, the virtual serial number of the virtual storage machine is displayed. If the model of the virtual storage machine is VSP 5000 series, the serial number added with 500,000 is displayed.

LDEV#

Displays the LDEV number corresponding to the specified NGUID.

- If the LDEV corresponding to the specified NGUID does not belong to the virtual storage machine created by the user, the actual LDEV number in the storage system is displayed. In this case, the displayed LDEV number is the same as the number displayed for PHY_LDEV#.
- If the LDEV corresponding to the specified NGUID belongs to the virtual storage machine created by the user, the virtual LDEV number is displayed.

PHY_Seq#

Displays the actual serial number of the storage system corresponding to the specified NGUID. If the target storage system is VSP 5000 series, the serial number added with 500,000 is displayed.

PHY_LDEV#

Displays the actual LDEV number in the storage system corresponding to the specified NGUID.

raidar*

The **raidar** command displays configuration, status, and I/O activity information for the specified port/TID(s)/LUN at the specified time interval. The configuration information is acquired directly from the storage system (not from the configuration definition file).

The I/O activity of a TrueCopy, TrueCopyAsync, Universal Replicator, or global-active device S-VOL in the COPY or PAIR state includes TC/TC Async/UR/GAD remote I/Os (update copy operations) in addition to host-requested I/Os.

The I/O activity of a ShadowImage, Copy-on-Write Snapshot, Thin Image, or Volume Migration S-VOL in the COPY or PAIR state includes only host-requested I/Os (ShadowImage, Copy-on-Write Snapshot, Thin Image, and Volume Migration update copy operations are excluded).

The I/O activity of a P-VOL or simplex volume includes only host-requested I/Os. If the status changed into SMPL in S-VOL (COPY, PAIR) I/O activity, I/O activity is reported as the SMPL status, until the pair status is changed.

Syntax

```
raidar { -h | -q | -z[x] | -I[H][M][instance#] or -I[TC][SI][instance#] |
        -p <port#> <targ> <lun> [mun] | -pd[g] <raw_device> [mun] [-n] |
        -s [interval] [count] }
```

Options and parameters**-h**

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidar** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:
a TrueCopy, Universal Replicator, or global-active device pair
- If the HORCC_MRCF environment variable is set:
a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:
an instance with no instance number is used.
- If the HORCMINST environment variable is set:
an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-p <port#> <targ> <lun> [mun]....

Monitors one or more (up to 16) devices at a time.

- <port#>: Specifies the port to be reported: CL1-A to CL1-R and CL2-A to CL2-R (excluding CL1-I, CL1-O, CL2-I, CL2-O). In case of the expanded port, specify from following: CL3-A to CL3-R (excluding CL3-I and CL3-O), or CLG-A to CLG-R (excluding CLG-I and CLG-O).

The port is not case sensitive (for example, CL1-A= cl1-a= CL1-a= cl1-A, CL3-a= CL3-A= cl3-a= cl3-A).

- <targ>: Specifies the SCSI TID of the specified port.

- `<lun>`: Specifies the LUN on the specified TID.
- `[mun]`: Specifies the MU number of the specified LUN within the range of 0 to 63 (ShadowImage, Thin Image, or Copy-on-Write Snapshot only).

-pd[g] <raw_device> [mun][-n]

Allows designation of an LDEV by raw device file name. The `-pdg` option is used to show a LUN on the host view by finding a host group.

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-s [interval[count]] or -sm [interval[count]]

Designates the time interval in seconds.

- `-s`: Interprets the time interval as seconds.
- `-sm`: Interprets the time interval as minutes.
- `[interval]`: Designates the time interval value (1 to 60). If not specified, the default interval (3) is used.
- `[count]`: Designates the number of repeats. When this parameter is omitted, this command repeats until CNTL-C. If you want to specify a value for this option, specify in the range of 0 to 2147483647.

Returned values

None

Error codes

None

Examples

The following shows an example of the **raidar** command and its output.

```
# raidar -p cl1-a 15 6 -p cl1-b 14 5 -p cl1-a 12 3 -s 3
```

```
TIME[03] PORT   T L VOL   STATUS IOPS HIT(%)W(%) IOCNT
13:45:25 -      - - -      -      -      -      -
13:45:28 CL1-A 15 6 SMPL -      200.0 80.0  40.0  600
          CL1-B 14 5 P-VOL PAIR 133.3 35.0  13.4  400
          CL1-A 12 3 P-VOL PSUS 200.0 35.0  40.6  600
```

Description of the `raidar` command output:**IOPS**

of I/Os (read/write) per second (total I/O rate).

HIT(%)

Hit rate for read I/Os (read hit rate).

W(%)

Ratio of write I/Os to total I/Os (percent writes).

IOCNT

number of write and read I/Os.

raidqry*

The **raidqry** command (RAID query) displays the configuration of the connected host and RAID storage system.

Syntax

```
raidqry { -h | -q | -z[x] | -I[H][M][instance#] or
        -I[TC][SI][instance#] | -l[m] | -r <group> | [ -f ] | -g | -c }
```

Options and parameters**-h**

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidqry** command enter the interactive mode. The `-zx` option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the `-zx` option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-I[m]

Displays the configuration information for the local host and the local RAID storage system.

- -l option: Display DKCMAIN microcode version in CCI format.
- -lm option: Display DKCMAIN microcode version in the same format as used in Device Manager - Storage Navigator or maintenance utility.

-r <group>

Displays the configuration information for the remote host and the remote storage system which contains the specified group.

-f

Displays the host name (ip_address) as specified in the configuration definition file. Use this option if 'floatable IP address' is used for the host name (ip_address) in the configuration file.

-g

This option is used when displaying the lists of group name (dev_group) which described in the configuration file of a local host (instance).

-c

Displays the information about the copy group in an MxN configuration managed by CCI instances and the instances that perform the processing to maintain the data consistency in the copy group. In a Universal Replicator MxN configurations with redundant CCI, one of the CCI instances performs the processing, such as issuing a time stamp (CTQ-Marker) to maintain the data consistency in the consistency group. The other instance is in the standby status in case of failure. If you issue the **pairsplit** command from this standby instance, an error (EX_INVVOL) might occur. Use this option to check which instance can issue the **pairsplit** command in an MxN configuration with redundant CCI.

Returned values

None

Error codes

None

Example 1

The following shows the example of the **raidqry** command and its output.

```
# raidqry -l
```

No	Group	Hostname	HORCM_ver	Uid	Serial#	Micro_ver	Cache(MB)
1	---	HOSTA	01-22-03/06	0	30053	50-04-00/00	256
1	---	HOSTA	01-22-03/06	1	30054	50-04-00/00	256

```
# raidqry -lm
```

No	Group	Hostname	HORCM_ver	Uid	Serial#	Micro_version
1	---	JSSA1012	01-37-03/04	0	302614	80-04-20-00/00

```
# raidqry -r oradb
```

No	Group	Hostname	HORCM_ver	Uid	Serial#	Micro_ver	Cache(MB)
1	oradb	HOSTA	01-22-03/06	0	30053	50-04-00/00	256
2	oradb	HOSTB	01-22-03/06	0	30053	50-04-00/00	256

```

1 oradb   HOSTA  01-22-03/06   1   30054 50-04-00/00   256
2 oradb   HOSTB  01-22-03/06   1   30054 50-04-00/00   256

```

```
# raidqry -l -f
```

```

No Group Floatable Host  HORCM_ver Uid Serial# Micro_ver Cache(MB)
1  ---          FH001 01-22-03/06   0   30053 50-04-00/00 256

```

Output of the **raidqry** command (Example 1):

No

This column shows the order when the group name (dev_group) which is described in the configuration definition file has multiple remote hosts.

Group

When the `-r` option is used, this column shows the group name (dev_group) which is described in the configuration definition file.

Hostname

When using `-l[m]` option, this column shows the host name of local host. When using `-r` option, this column shows the host name of remote host which is included the group name (dev_group) described in a configuration definition file. Over 30 characters long of the host name is not displayed.

Floatable Host

When the `-f` option is used, this column shows the host name (ip_address) which is described in the configuration definition file. Up to 30 host names can be displayed. The `-f` option interprets the host name as utilizing floatable IP for a host.

HORCM_ver

This column shows the version of CCI on the local or remote host. The `-l [m]` option specifies local host. The `-r` option specifies remote host.

HORCM_ver: This column shows the version of CCI on the local or remote host. The `-l [m]` option specifies local host. The `-r` option specifies remote host.

Uid Serial# Micro_ver (Micro_version)

This column shows unitID, serial number, and (DKCMAIN) microcode version of the storage system that is connected to the local or remote host. The `-l [m]` option specifies local host. The `-lm` option displays Micro_version. The `-r` option specifies remote host.

Cache(MB)

Shows the logical cache capacity (megabytes) of the storage system connected to the local or remote host. The `-l` option specifies local host, and `-r` specifies remote host.

Example 2

```
# raidqry -g
```

GNo	Group	RAID_type	IV/H	IV/M	MUN/H	MUN/M
1	ora	HTC_RAID	12	9	4	64
2	orb	XP_RAID	12	9	4	64
3	orc	HTC_DF	8	6	1	1

Output of the raidqry command (Example 2):**GNo**

The order of the group name (dev_group) described in the configuration definition file.

Group

The group name (dev_group) described in the configuration definition file.

RAID_type

The type of RAID configured in the group.

IV/H

The interface version for TrueCopy/TrueCopy Async/Universal Replicator/global-active device in a group, and this is used for maintenance.

IV/M

The interface version for ShadowImage/Copy-on-Write Snapshot/Thin Image/Volume Migration in a group, and this is used for the maintenance.

MUN/H

The number of maximum MUs for Universal Replicator in a group.

MUN/M

The number of maximum MUs for ShadowImage/Copy-on-Write Snapshot/Thin Image in a group.

Example 3

```
# raidqry -c
GNo  Group  S
1    ora   N
2    orb   N
3    orc   Y
```

Output of the raidqry command (Example 3):**GNo**

The order of the group name (dev_group) described in the configuration definition file.

Group

The group name (dev_group) described in the configuration definition file.

S

Displays the information about the copy group in an MxN configuration managed by CCI instances and the instances that perform the processing to maintain the data consistency in the copy group.

- **Y:** This instance performs the processing to maintain the data consistency in the copy group, or it does not support a redundant configuration. You can run the **pairsplit** command from this instance. For more information about the requirements for the redundancy support, see the Command Control Interface User and Reference Guide.
- **N:** Instances other than this instance perform the processing to maintain the data consistency in the copy group, or the instances are in the pair status that cannot run the **pairsplit** command. You cannot run the **pairsplit** command from this instance.

raidvchkset*

The **raidvchkset** command sets the parameters for validation checking of the specified volumes, and can also be used to turn off all validation checking without specifying [type]. Unit of checking for the validation is based on the group of CCI configuration definition file.

This command is controlled as protection facility. This command is rejected with EX_ERPERM by connectivity checking between CCI and the RAID storage system.

Syntax

```
raidvchkset { -h | -q | -z[x] | -I[H][M][instance#] or
              -I[TC][SI][instance#]
              | -g <group> | -d <pair Vol> | -d[g] <raw_device> [MU#] [-n]
              | -d[g] <seq#> <LDEV#> [MU#] | -nomsg | -vg [type] [rtime]
              | -vext <size> }
```

Options and parameters**-h**

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidvchkset** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]
a TrueCopy, Universal Replicator, or global-active device pair
- [M] or [SI]
a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:
a TrueCopy, Universal Replicator, or global-active device pair
- If the HORCC_MRCF environment variable is set:
a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:
an instance with no instance number is used.
- If the HORCMINST environment variable is set:
an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g <group>

Specifies a group name written in the configuration definition file.

-d <pair Vol>

Specifies paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volume.

-d[g] <raw_device> [MU#] [-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the `-d` option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (`-d`) or group (`-dg`). This option is effective without specification of '`-g <group>`' option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The `<seq#> <LDEV#>` values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the `-d` option \(on page 22\)](#).

-nomsg

Suppresses messages to be displayed when this command is executed. It is used to execute this command from a user program. This option must be specified at the beginning of a command argument. The command execution log is not affected by this option.

-vg [type] [rtime]

Specifies the following guard type to the target volumes for Data Retention Utility. If `[type]` is not specified, this option releases all of the guarding.

The following values are available to specify on `[type]`.

- `inv`: The target volumes are concealed from SCSI Inquiry command by responding 'unpopulated volume'.
- `sz0`: The target volumes replies with 'SIZE 0' through SCSI Read capacity command.
- `rwd`: The target volumes are prohibited from reading and writing.

- wtd: The target volumes are prohibited from writing.
- svd: If the target volume is SMPL, it is protected from paircreate (from becoming an S-VOL). If the target volume is P-VOL, it is protected from pairresync restore or pairresync swaps(p). If the target volume is SVOL_PSUS(SSUS), it is protected from pairresync synchronous copy.

svd option can be used with the other option (inv, sz0, rwd, wtd) in parallel. For example, if you want to protect the absolute volume from the writing executed by the copy series software product and the host access, set the both wtd and svd options. The only setting of wtd option cannot protect the absolute volume from the writing by the copy processing of the copy series software product.

[*ptime*]: Specifies the retention time in days. If [*ptime*] is not specified, the default value 0 (zero) defined by the storage system is used. This option is ignored in 9900V microcode versions 21-06-xx or 21-07-xx (default = infinite). To change a flag that displays E in PI-C-R-W-S by using the **raidvchkdsp** command for a volume for which the retention time is set, specify a value that is not less than the retention time displayed by using the **raidvchkdsp** command for [*ptime*].

- This option sets each four flags for guarding type as follows:

```
typeINQRCAPREADWRITE inv1111 Sz00111 rwd0011 wtd0001
```

-vext <size>

Used when extending the LUN capacity of a Dynamic Provisioning volume.

The increment size of capacity can be specified in bytes or blocks. When specifying in bytes, the unit is t/T (terabyte), g/G (gigabyte), m/M (megabyte), or k/K (kilobyte). If this unit is omitted, block (512 bytes) is used.

Examples for extending 1GB (gigabyte) are: -vext 1G, -vext 1g, -vext 1024M, -vext 1024m, -vext 1048576K, -vext 1048576k, -vext 2097152

Examples for extending 1KB (kilobyte) are: -vext 1K, -vext 1k, -vext 2

LUN capacity and usage rate for Dynamic Provisioning volume can be verified by referring 'LU_CAP' of the 'raidvchkdsp -v aou' or 'raidvchkdsp -v aoub' command. "Aou" (allocation on use) refers to Dynamic Provisioning.



Note:

When a group operation is specified, a warning message appears, and this command enters the interactive mode.

Returned values

The **raidvchkset** command sets the following returned values during exit allowing you to check the execution results.

- Normal termination: 0
- Abnormal termination: Specific error codes (see Error codes below) and generic error codes.

Error codes

The **raidvchkset -vg** option command returns the following error code as well as generic error codes. See the table below.

Category	Error Code	Error Message	Recommended Action	Value
Volume Status (Unrecoverable)	EX_EPRORT	Mode changes denied due to retention time	Confirm the retention time for a target volume by using raidvchkscan -v gflag command.	208

Examples

Disables writing to volumes in oralog group:

```
raidvchkset -g oralog -vg wtd
```

Disables writing and sets retention time to volumes in oralog group:

```
raidvchkset -g oralog -vg wtd 365
```

Releases all guarding to volumes in oralog group:

```
raidvchkset -g oralog -vg
```

raidvchkdsp*

The **raidvchkdsp** command displays the parameters for validation checking of the specified volumes. Unit of checking for the validation is based on the group of CCI configuration definition file.

Syntax

```
raidvchkdsp { -h | -q | -z[x] | -I[H][M][instance#] or
              -I[TC][SI][instance#] | -g <group> | -d <pair Vol>
              | -d[g] <raw_device> [MU#] [-n] | -d[g] <seq#> <LDEV#> [MU#]
              | -f[xde] | -v <op> | -c }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidvchkdsp** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- **-I[H][M] or -I[TC][SI]**

Specifies a target pair type on which the command operates.

- **[H] or [TC]**

a TrueCopy, Universal Replicator, or global-active device pair

- **[M] or [SI]**

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- **[instance#]**

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: **-IH0**

-g <group>

Specifies a group name written in the configuration definition file.

-d <pair Vol>

Specifies paired logical volume name defined in the configuration definition file. When this option is specified, the command is executed for the specified paired logical volume.

-d[g] <raw_device> [MU#] [-n]

Searches whether the specified raw_device is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified raw_device is contained in two or more groups, the command is executed for the first group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-d[g] <seq#> <LDEV#> [MU#]

Searches whether the specified LDEV is included in a group on the configuration definition file (local instance). If it is in the group, the target volume is executed as the paired logical volume (-d) or group (-dg). This option is effective without specification of -g <group> option. If the specified LDEV is contained in two or more groups, the command is executed on the first group. The <seq#> <LDEV#> values can be specified in hexadecimal (add the '0x' prefix) or decimal notation.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

For details about MU#, see [Specifying a mirror unit \(MU\) number by using the -d option \(on page 22\)](#).

-fx

Displays the LDEV number in hexadecimal.

-fd

Displays the relationship between the Device_File and the paired volumes, based on the group (as defined in the local instance configuration definition file). If Device_File column shows 'Unknown' to host (instance), then the volume is not recognized on own host, and raidvchkdsp command is rejected in protection mode. Non-permitted volume is shown without LDEV# information (LDEV# '-').

Specify ENABLE for HORCM_NVME in the configuration file. If ENABLE is not specified, a device file name for NVMe is not displayed.

-fe

Displays the serial# and LDEV# of the external LUNs mapped to the LDEV for the target volume by adding to last column (ignores the format of 80 column).

-c

When CCI starts, HORCM_DEV in horcm.conf is translated from port/target/lun numbers to the CU:Ldev information, on one hand HORCM_LDEV in horcm.conf is translated from the CU:Ldev information to port/target/lun numbers, because RAID needs to specify 'Port#, Targ#, Lun#' and 'LDEV' for specifying the target device, and then HORCM keeps this information as internal database for the configuration.

If a storage administrator changes the LDEV to LUN/port mapping, such as

- a new/different LDEV is mapped to a previously used port/LUN, or
- an LDEV is mapped to a different/new port

then pair operations might be rejected because the new mapping is different from the mapping information of the database in the running CCI instance. A **pairdisplay** command shows the real LDEV mapping at the time of the command execution and hence shows different information than what is stored in the internal database of the CCI instance.

The **-c** option for **raidvchkdsp** allows you to see if there is a difference between the current running CCI instance information and the real mapping. This indication should be used to find such issues which indicate that:

- the CCI instance should be restarted to discover and use the new mapping information, or
- a configuration change occurred without changing the affected configuration files of the CCI instance.

Example change from LDEV#785 to LDEV#786:

```
# raidvchkdsp -g VG000 -c
```

```
Group PairVol Port# TID LU Seq# LDEV# LDEV#(conf) -change-> LDEV#
VG000 vg0001 CL4-E-0 0 17 63528 786 785(conf) -change-> 786
```

```
# raidvchkdsp -g VG000 -c -fx
```

```
Group PairVol Port# TID LU Seq# LDEV# LDEV#(conf) -change-> LDEV#
VG000 vg0001 CL4-E-0 0 17 63528 312 311(conf) -change-> 312
```

Example remove LDEV#785 from a port:

```
# raidvchkdsp -g VG000 -c
```

```
Group PairVol Port# TID LU Seq# LDEV# LDEV#(conf) -change-> LDEV#
VG000 vg0001 CL4-E-0 0 17 63528 - 785(conf) -change-> NO LDEV
```

```
# raidvchkdsp -g VG000 -c -fx
```

```
Group PairVol Port# TID LU Seq# LDEV# LDEV#(conf) -change-> LDEV#
VG000 vg0001 CL4-E-0 0 17 63528 - 311(conf) -change-> NO LDEV
```



Note:

If there have not been any changes to the specified volumes, this option displays nothing.

-v [op]

Specifies the following operation that displays each parameter for validation checking:

- gflag: Displays the parameter for guarding on the specified target volumes.
- pool : This option displays the information about the Dynamic Provisioning pool to which the LDEV belongs. If the LDEV does not belong to a Dynamic Provisioning pool, information about the Thin Image or Copy-on-Write Snapshot pool is displayed. See Example 2 below.
- aou[b] : Displays the LUN capacity and usage rate for Dynamic Provisioning volume corresponding to the group of RM configuration file, and displays the ID of the pool to which LDEV belongs.
- aoub : Displays the LUN capacity in blocks (1 block = 512 bytes).

Returned values

None

Error codes

None

Example 1

Example of the **raidvchkdsp** command with the **-v gflag** option and its output:

```
# raidvchkdsp -g vg01 -fd -v gflag
```

Group	PairVol	Device_File	Seq#	LDEV#	GI-C-R-W-S	PI-C-R-W-S	R-Time
vg01	oradb1	c4t0d2	2332	2	E E D D E	E E D D E	365
vg01	oradb2	c4t0d3	2332	3	E E D D E	E E D D E	-

Description of the **raidvchkdsp command output with the **-v gflag** option:****GI-C-R-W-S**

Displays the flags for guarding as for the target volume.

- I=E: Enabled for Inquiry command.
- I=D: Disabled for Inquiry command.
- C=E: Enabled for Read Capacity command.
- C=D: Disabled for Read Capacity command.
- R=E: Enabled for Read command.
- R=D: Disabled for Read command.
- W=E: Enabled for Write command.
- W=D: Disabled for Write command.

- S=E: Enabled for becoming the S-VOL.
- S=D: Disabled for becoming the S-VOL.

PI-C-R-W-S

Displays the permission flags that show whether each mode flag can be changed to enable or not.

- I=E: 'I' flag can be changed to enable.
- I=D: 'I' flag cannot be changed to enable.
- C=E: 'C' flag can be changed to enable.
- C=D: 'C' flag cannot be changed to enable.
- R=E: 'R' flag can be changed to enable.
- R=D: 'R' flag cannot be changed to enable.
- W=E: 'W' flag can be changed to enable.
- W=D: 'W' flag cannot be changed to enable.
- S=E: 'S' flag can be changed to enable.
- S=D: 'S' flag cannot be changed to enable.

R-Time

Displays the retention time for write protect in days. The hyphen (-) shows that the retention time is infinite. The application knows whether the target volume is denied to change to writing enable by referring 'R-Time'.

R-time (Rtime) is identical to rtime and both of them indicate Retention Time. This setting value can normally be identified as a value of R-time that is output by **raidychkdsp** (the logging format is Rtime=xxxx). However, R-time (Rtime) is indicated as the value of 'Retention Time + 1000000' when the expiration lock is enabled. The setting of **raidvchkset** command in this status is denied.

Audit lock status is shown as the retention time plus 1000000. 'R-Time + 1000000' shows the retention time with Audit lock status.

Example 2

Example of the **raidvchkdsp** command with the **-v pool** option and its output:

```
raidvchkdsp -g vg01 -v pool
```

Group	PairVol	Port#	TID	LU	Seq#	LDEV#	Bsize	Available	Capacity
Vg01	oradb1	CL2-D	2	7	62500	167	2048	100000	1000000000
Vg01	oradb2	CL2-D	2	10	62500	170	2048	100000	1000000000

Description of the `raidvchkdsp` command output with the `-v pool` option:**Bsize**

Displays the data block size of the pool in blocks (512 bytes).

Available(Bsize)

Displays the available capacity for the volume data on the Copy-on-Write Snapshot pool in units of Bsize.

Capacity(Bsize)

Displays the total capacity in the Copy-on-Write Snapshot pool in units of Bsize.

Example 3

Example of the `raidvchkdsp` command with the `-v aou` option and its output ("aou" (allocation on use) refers to Dynamic Provisioning):

```
# raidvchkdsp -v aou -g AOU
```

```
Group  PairVol Port# TID LU Seq# LDEV# Used(MB) LU_CAP(MB) U(%)
T(%) PID
AOU    AOU_001 CL2-D  2   7 62500  167   20050  1100000   10
70     1
AOU    AOU_002 CL2-D  2  10 62500  170  110000  1100000   10
70     1
```

Description of the `raidvchkdsp` command output with the `-v aou` option:**Used(MB)**

Displays the usage size (megabytes) of the allocated block on this LUN. Range: $0 \leq \text{Used (MB)} < \text{LU_CAP(MB)} + 42 \text{ MB}$

LU_CAP(MB)

Displays the LUN capacity(megabytes) responded to the 'Readcapacity' command as SCSI interface.

LU_CAP(BLK)

Displays the LUN capacity (in block/512 bytes) responded to the 'Readcapacity' command as SCSI interface.

U(%)

Displays the usage rate of the allocated block on the Dynamic Provisioning pool containing this LU.

T(%)

Displays the threshold rate being set to the Dynamic Provisioning pool as the high water mark.

PID

Displays the Dynamic Provisioning pool ID assigned to this Dynamic Provisioning volume.

raidvchkscan*

The **raidvchkscan** command displays the port of the storage system (9900V or later), target ID, LDEV mapped for LUN# and MU#, and status of LDEV, regardless of the configuration definition file.

**Note:**

This command is rejected with EX_ERPERM by connectivity checking between CCI and the RAID storage system.

Syntax

```
raidvchkscan { -h | -q | -z[x] | -I[H][M][instance#] or
               -I[TC][SI][instance#] | -p <port#> [hgrp] | -pd[g] <raw_device>[-n]
               | -s <seq#> | -t <target> | -l <lun> | [ -f[x] ] | -v <op> }
```

Options and parameters**-h**

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidvchkscan** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-p <port#> [hgrp]

Specifies the port ID of the port to be scanned. Valid ports are CL1-A to CL1-R and CL2-A to CL2-R (excluding CL1-I, CL1-O, CL2-I, CL2-O). In addition, it is able to specify from CL3-a to CL3-r (except CL3-i, CL3-o), or CL4-a to CL4-r (except CL4-i, CL4-o) for the expanded port.

The port is not case sensitive (for example, CL1-A= cl1-a= CL1-a= cl1-A, CL3-a= CL3-A= cl3-a= cl3-A). This option must be specified if '-find' or '-pd <raw_device>' option is not specified.

Specify [hgrp] to display only the LDEVs mapped to a host group on a port.

-pd[g] <raw_device> [-n]

Specifies the raw device name. This option finds Seq# and port_name of the storage system to which the specified device can be connected, and scans the port of the storage system which corresponds with the unit ID that searches the unit ID from Seq#. This option must be specified if the '-find' option is not specified. If this option is specified, the following -s <seq#> option is invalid.

-pdg: Shows a LUN on the host view by finding a host group.

This -n option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the HORCC_NVME environment variable is specified, both device files for NVMe and SCSI can be used without the -n option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

-s <seq#>

Used to specify the seq# (serial#) of the storage system when this option cannot specify the unit ID which is contained for '-p <port>' option. This option scans the port specified by '-p <port>' option of the storage system which corresponds with the unit ID that searches the unit ID from seq#. If this option is specified, then the unit ID which is contained in '-p <port>' option is invalid.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-t <target>

Specifies a target ID (0 to 15) of the specified port. If this option is not specified, the command applies to all target IDs.

-l <lun>

Specifies a LUN (0 to 2047) of the specified target ID. For VSP 5000 series, the LUN that can be specified is (0 to 4095). If this option is not specified, the command applies to all LUNs. If this option is specified, the TID must also be specified.

-fx

Displays the LDEV number in hexadecimal notation.

-v [op]

Specifies the following operation that displays each parameter for validation checking:

gflag: Displays the parameter for guarding on the specified target volumes.

pool: This option displays the information about the Dynamic Provisioning pool to which the LDEV belongs. If the LDEV does not belong to a Dynamic Provisioning pool, information about the Thin Image or Copy-on-Write Snapshot pool is displayed. See Example 2 below.

aou[b]: Displays the LUN capacity and usage rate for only Dynamic Provisioning volume mapped to the specified port, and displays the ID of the pool to which LDEV belongs.

aoub: Displays the LUN capacity in blocks (512 bytes).

Returned values

None

Error codes

None

Example 1

Example of the **raidvchkscan** command with the **-v gflag** option

```
# raidvchkscan -p CL1-A -v gflag
```

PORT#	/ALPA/C	TID#	LU#	Seq#	Num	LDEV#	GI-C-R-W-S	PI-C-R-W-S	R-Time
CL1-A	/ ef/	0	0	2332	1	0	E E D D E	E E D D E	365
CL1-A	/ ef/	0	1	2332	1	1	E E D D E	E E D D E	-
CL1-A	/ ef/	0	2	2332	1	2	E E D D E	E E D D E	0

Description of the **raidvchkscan command output with the **-v gflag** option:**

GI-C-R-W-S

Displays the flags for guarding as for the target volume.

- I=E: Enabled for Inquiry command.
- I=D: Disabled for Inquiry command.
- C=E: Enabled for Read Capacity command.
- C=D: Disabled for Read Capacity command.
- R=E: Enabled for Read command.
- R=D: Disabled for Read command.
- W=E: Enabled for Write command.
- W=D: Disabled for Write command.
- S=E: Enabled for becoming the S-VOL.
- S=D: Disabled for becoming the S-VOL.

PI-C-R-W-S

Displays the permission flags that show whether each mode flag can be changed to enable or not.

- I=E: 'I' flag can be changed to enable.
- I=D: 'I' flag cannot be changed to enable.
- C=E: 'C' flag can be changed to enable.
- C=D: 'C' flag cannot be changed to enable.
- R=E: 'R' flag can be changed to enable.
- R=D: 'R' flag cannot be changed to enable.
- W=E: 'W' flag can be changed to enable.
- W=D: 'W' flag cannot be changed to enable.
- S=E: 'S' flag can be changed to enable.
- S=D: 'S' flag cannot be changed to enable.

R-Time

Displays the retention time for write protect in days. The hyphen (-) shows that the retention time is infinite. The application knows whether the target volume is denied to change to writing enable by referring 'R-Time'.



Note: Audit lock status is shown as the retention time plus 1000000. 'R-Time + 1000000' shows the retention time with Audit lock status.

Example 2

Example of the **raidvchkscan** command with the **-v pool** option:

```
# raidvchkscan -v pool -p CL2-d-0
```

PORT#	/ALPA/C	TID#	LU#	Seq#	Num	LDEV#	Bsize	Available	Capacity	
CL2-D-0	/e4/	0	2	0	62500	1	160	2048	100000	1000000000
CL2-D-0	/e4/	0	2	1	62500	1	161	2048	100000	1000000000

Description of the **raidvchkscan command output with the **-v pool** option:****Bsize**

Displays the data block size of the pool in blocks (512 bytes).

Available(Bsize)

Displays the available capacity for the volume data on the Copy-on-Write Snapshot pool in units of Bsize.

Capacity(Bsize)

Displays the total capacity in the Copy-on-Write Snapshot pool in units of Bsize.

Example 3

Example of the **raidvchkscan** command with the **-v aou** option ("aou" (allocation on use) refers to Dynamic Provisioning)

```
# raidvchkscan -v aou -p CL2-d-0
```

PORT#	/ALPA/C	TID#	LU#	Seq#	Num	LDEV#	Used(MB)	LU_CAP(MB)	U(%)	T(%)	PID	
CL2-D-0	/e4/	0	2	0	62500	1	160	20050	1100000	1	60	1
CL2-D-0	/e4/	0	2	1	62500	1	161	200500	1100000	18	60	2

Description of the raidvchkscan command output with the -v aou option:**Used(MB)**

Displays the usage size (megabytes) the allocated block on this LUN. Range: $0 \leq \text{Used (MB)} < \text{LU_CAP(MB)} + 42\text{MB}$

LU_CAP(MB)

Displays the LUN capacity (megabytes) responded to the 'Readcapacity' command as SCSI interface.

LU_CAP(BLK)

Displays the LUN capacity (in block/512 bytes) responded to the 'Readcapacity' command as SCSI interface.

U(%)

Displays the usage rate of the allocated block on the Dynamic Provisioning pool containing this LU.

T(%)

Displays the threshold rate being set to the Dynamic Provisioning pool as the high water mark.

PID

Displays the Dynamic Provisioning pool ID assigned to this Dynamic Provisioning volume.

raidvchkscan for Universal Replicator*

The **raidvchkscan** command supports the **(-v jnl [t] [unit#])** option to find the journal volume list. It also displays any information for the journal volume.

Syntax

```
raidvchkscan { -h | -q | -z[x] | -I[H][M][instance#] or -I[TC][SI][instance#] | -v jnl  
[t] [unit#] | [ -s <seq#> ] | [ -f[x] ] }
```


Options and parameters**-h**

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidvchkscan** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- **-I[H][M] or -I[TC][SI]**

Specifies a target pair type on which the command operates.

- **[H] or [TC]**

a TrueCopy, Universal Replicator, or global-active device pair

- **[M] or [SI]**

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the **HORCC_MRCF** environment variable is set in the command execution environment.

- If the **HORCC_MRCF** environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the **HORCC_MRCF** environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- **[instance#]**

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the **HORCMINST** environment variable is set in the command execution environment.

- If the **HORCMINST** environment variable is not set:

an instance with no instance number is used.

- If the **HORCMINST** environment variable is set:

an instance whose instance number is set as the **HORCMINST** environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: **-IH0**

-s <seq#>

Used to specify the seq# (serial#) of the storage system when this option cannot specify unitID which is contained for '-v jnl' option. If this option is specified, the unitID which is contained in '-v jnl' is invalid.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-fx

Displays the LDEV number in hexadecimal notation.

-v jn

Displays information for the journal volume.

-v jnlt

Displays the DOW , DPW , and APW time-out values for controlling the journal.

Returned values

None

Error codes

Non

Example 1

Example of the **raidvchkscan** command with the **-v jnl 0** option

```
# raidvchkscan -v jnl 0
```

JID	MU	CTG	JNLS	AP	U(%)	Q-Marker	Q-CNT	D-SZ (BLK)	Seq#	Nnm	LDEV#
001	0	1	PJNN	4	21	43216fde	30	512345	62500	2	265
002	1	2	PJNF	4	95	3459fd43	52000	512345	62500	3	270
002	2	2	SJNS	4	95	3459fd43	52000	512345	62500	3	270
003	0	3	PJSN	4	0	-	-	512345	62500	1	275
004	0	4	PJSF	4	45	1234f432	78	512345	62500	1	276
005	0	5	PJSE	0	0	-	-	512345	62500	1	277
006	-	-	SMPL	-	-	-	-	512345	62500	1	278
007	0	6	SMPL	4	5	345678ef	66	512345	62500	1	278

Description of the **raidvchkscan command output with the **-v jnl 0** option:**

JID

Journal ID

MU

Mirror descriptions on Universal Replicator.

CTG

Consistency group ID

JNLS


Status of the journal:

- SMPL: This means the journal volume which does not have a pair, or in the state of deleting.
- P(S)JNN: This means 'P(S) Journal Volume is in Normal status'.
- P(S)JNS: This means 'P(S) Journal Volume is suspended in Normal status' created with `-nocsus` option.
- P(S)JSN: This means 'P(S) Journal Volume is suspended in Normal status'.
- P(S)JNF: This means 'P(S) Journal Volume is in Full status'.
- P(S)JSF: This means 'P(S) Journal Volume is suspended in Full status'.
- P(S)JSE: This means 'P(S) Journal Volume is suspended by an error (including link failures)'.
- P(S)JES: This means 'P(S) Journal Volume is suspended by an error' created with `-nocsus` option.

AP

Displays the following two conditions (status) according to the pair status.

Shows the number of active paths on the initiator port in Universal Replicator links. 'Unknown' is shown as '-'.

**AM**

The activity monitor that detects whether or not there is a request for data from the initiator at regular intervals. If AM detects a time-out, the P-JNL state is changed from P-JNN to PJSE.



Note: The same path information is used for AP for three commands (`pairvolchk`, `pairdisplay`, `raidvchkscan`). The differential is that `pairvolchk` and `pairdisplay` are to show a special meaning with SSUS(SSWS) state.


Q-Marker

Displays the sequence # of the journal ID, called the Q-marker. For P-JNL, Q-Marker shows the latest sequence # on the P-JNL volume. For S-JNL, the Q-Marker shows the latest sequence # of the cache(DFW).

Q-CNT

Displays the number of remaining Q-Markers within each journal volume.

The following figure shows an example of Q-Marker and Q-CNT

**U(%)**

Displays the usage rate of the journal data.

D-SZ

Displays the capacity for the journal data on the journal volume. For details about the displayed capacity, see the *Hitachi Universal Replicator User Guide*.

Seq#

Displays the serial number of the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Num

Displays the number of LDEVs composing the journal volume.

LDEV#

Displays the following LDEV numbers of journal volumes for storage systems that support reserve journal volumes:

- LDEV number of the master journal volume for the master journal
- LDEV number of the restore journal volume for the restore journal. Whether to support reserve journal volumes depends on the model of your storage system. For the support status, see the *Hitachi Universal Replicator User Guide* for your model.

The following table lists information about the different journal volume statuses. QCNT=0 indicates that the number of remaining Q-Markers is '0'. The letter 'N' indicates a non-zero.

JNLS		Other Information		Description
P-JNL	S-JNL	QCNT	AP	
SMPL		0	-	Configured as journal volume, but NOT pair

JNLS		Other Information		Description
P-JNL	S-JNL	QCNT	AP	
		N	-	Deleting the journal volume
PJNN (PJNS)	SJNN (SJNS)	0	-	Normal state of the journal volume without data
PJNN (PJNS)	-	N	-	Normal state of the journal volume with data
-	SJNN (SJNS)	N	N	Normal state of the journal volume with data
			0	Still normal state of the journal volume at Link failure
PJSN	SJSN	0	-	Suspended journal volume via operation
		N	-	Suspending the journal volume
PJNF	-	N	-	High water mark state
PJSF	SJSF	0	-	Suspended journal volume due to full journal
		N	-	Suspending the journal volume due to full journal
PJSE	-	0	-	Suspended journal volume due to failure/Link failure
		N	-	Suspending the journal volume due to failure/Link failure
-	SJSE	0	N	Suspended journal volume due to failure
			0	Suspended journal volume due to Link failure
		N	N	Suspending the journal volume due to failure
			0	Suspending the journal volume due to Link failure

Example 2

Example of the **raidvchkscan** command with the **-v jnl** option

```
# raidvchkscan -v jnl
```

```
JID MU CTG JNLS AP U(%) Q-Marker Q-CNT D-SZ (BLK) Seq# DOW PBW APW
001 0 1 PJNN 4 21 43216fde 30 512345 63528 20 300 40
```

002	1	2	PJNF	4	95	3459fd43	52000	512345	63528	20	300	40
003	0	3	PJSN	4	0	-	-	512345	63528	20	300	40

Description of the `raidvchkscan` command output with the `-v jnl` option:

DOW

Data Overflow Watch' timer (in seconds) setting per the Journal.

PBW:

Path Blockade Watch timer setting (in seconds) per the Journal. If the setting is more than 3600 seconds, it displays 6000 seconds.

APW

This shows 'Active Path Watch' timer (in seconds) for detecting Link failure.

raidvchkscan for Thin Image, Copy-on-Write Snapshot, and DP pools*

The `raidvchkscan` command supports the option (`-v pid[a] [unit#]`) to find the Thin Image, Copy-on-Write Snapshot, or Dynamic Provisioning pool settings via SVP, and displays information for the Thin Image, Copy-on-Write Snapshot, or Dynamic Provisioning pool.

Syntax

```
raidvchkscan { -h | -q | -z[x] | -I[H][M][instance#] or
               -I[TC][SI][instance#] | -v pid[a][s][b][-fp] [unit#]
               | [ -s <seq#> ] | [ -f[x] | ] }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the `raidvchkdsp` command enter the interactive mode. The `-zx` option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shutdown, interactive mode terminates.

OpenVMS cannot use the `-zx` option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-s <seq#>

Used to specify the Seq# (serial#) of the storage system when this option cannot specify unitID which is contained for '-v jnl[a]' option. If this option is specified, the unitID which is contained in '-v jnl[a]' is invalid.

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-fx

Displays the LDEV number in hexadecimal notation.

-v pidb -fp

Displays the actual capacity of the pool.

The following GUI displays the actual capacity information as the capacity of the accelerated compression enabled pool. Specify this option to confirm the capacity equivalent to the capacity displayed in the GUI.

- Device Manager - Storage Navigator (VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900)
- Storage Advisor Embedded or VSP One Block Administrator

If the microcode version does not support this option, nothing is displayed when this option is specified.

-v pid[s]

Displays information for the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

-v pid[a]

Displays information for the Dynamic Provisioning pool.

-v pidb

Displays basic information for the pool.

If the microcode version does not support this option, nothing is displayed when this option is specified.

Returned values

None

Error codes

None

Example 1

Example of the **raidvchkscan** command with the **-v pid** option:

```
# raidvchkscan -v pid 0
```

PID	POLS	U (%)	SSCNT	Available (MB)	Capacity (MB)	Seq#	Num	LDEV#	H (%)
FMT_CAP (MB)									
001	POLN	10	330	10000000	1000000000	62500	2	265	80 100
002	POLF	95	9900	100000	1000000000	62500	3	270	70 100
003	POLS	100	10000	100	1000000000	62500	1	275	70 100
004	POLE	0	0	0	0	62500	0	0	80 100

Description of the **raidvchkscan command output with the **-v pid** option:**

PID

Displays the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool ID.

POLS

Displays the status of the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is suspended in the failure status.) In this status, the pool information cannot be displayed.

U(%)

Displays the usage rate of the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

SSCNT

Displays the number of Thin Image/Copy-on-Write Snapshot volumes in the Thin Image/Copy-on-Write Snapshot pool or the total number of Dynamic Provisioning volumes mapped in this Dynamic Provisioning pool.

Available(MB)

Displays the available capacity (megabytes) for the volume data on the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

Capacity(MB)

Displays the total capacity (megabytes) in the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

Seq#

Displays the serial number of the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Num

Displays the number of LDEVs composing the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

LDEV#

Displays the first number of LDEV composing the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool.

H(%)

Displays the threshold rate being set to the Thin Image/Copy-on-Write Snapshot/Dynamic Provisioning pool as the high water mark. 'Unknown' is shown as '-'.

FMT_CAP(MB)

Displays the formatted pool capacity (megabytes). If there is no valid information for the pool, a hyphen (-) is displayed.

Example 2

Example of the **raidvchkscan** command with the **-v pida** option

```
# raidvchkscan -v pida 0
```

PID	POLS	U(%)	AV_CAP(MB)	TP_CAP(MB)	W(%)	H(%)	Num	LDEV#	LCNT	TL_CAP(MB)
001	POLN	10	45000000	50000000	50	80	2	265	33	65000000
002	POLF	95	10000	100000000	50	80	3	270	900	100000000
004	POLN	0	10000000	100000000	80	90	2	280	0	0

Description of the output of raidvchkscan command with the -v pida option:**PID**

Displays the Dynamic Provisioning pool ID.

POLS

Displays the status of the Dynamic Provisioning pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is suspended in the failure status.) In this status, the pool information cannot be displayed.

U(%)

Displays the usage rate of the Dynamic Provisioning pool.

AV_CAP(MB)

Displays the available capacity (megabytes) for the Dynamic Provisioning volumes mapped to this pool.

TP_CAP(MB)

Displays the total capacity (megabytes) of the Dynamic Provisioning pool.

W(%)

Displays the threshold value for 'WARNING' set for this Dynamic Provisioning pool.

H(%)

Displays the threshold rate set for the Dynamic Provisioning pool as the high water mark.

Num

Displays the number of LDEVs composing the Dynamic Provisioning pool.

LDEV#

Displays the first number of LDEV composing the Dynamic Provisioning pool.

LCNT

Displays the total number of Dynamic Provisioning volumes mapped to this Dynamic Provisioning pool.

TL_CAP(MB)

Displays the total capacity (megabytes) of all Dynamic Provisioning volumes mapped to this Dynamic Provisioning pool.

Example 3

Example of the **raidvchkscan** command with the **-v pidb** option:

```
# raidvchkscan -v pidb 0
```

```
PID POLS U(%) LCNT SSCNT Available(MB) Capacity(MB) Snap_Used(MB)
TL_CAP(MB) BM TR_CAP(MB) RCNT Seq# Num LDEV# W(%) H(%) STIP VCAP(%)
TYPE PM PT POOL_NAME D_PM D_PM(%) EFFECTIVE(MB)

001 POLN 0 11001 11001 46998 46998 0
2432398 NB 0 0 300050 1 0 70 80 YES - OPEN
N HDP dp_ti_pool N - 46998
001 POLN 0 11001 11001 46998 46998 0
2432398 NB 0 0 300050 1 0 70 80 YES - OPEN
N HDP dp_ti_pool N - 46998
```

Example of the **raidvchkscan** command with the **-v pidb -fp** option:

```
# raidvchkscan -v pidb 0 -fp
PID POLS U(%) LCNT SSCNT Available(MB) Capacity(MB) Snap_Used(MB)
TL_CAP(MB) BM TR_CAP(MB) RCNT Seq# Num LDEV# W(%) H(%) STIP VCAP(%)
TYPE PM PT POOL_NAME D_PM D_PM(%) EFFECTIVE(MB)

001 POLN 0 11001 11001 46998 46998 -
2432398 NB 0 0 300050 1 0 70 80 YES -
OPEN N HDP dp_ti_pool N - 46998

001 POLN 0 11001 11001 46998 46998 -
2432398 NB 0 0 300050 1 0 70 80 YES -
OPEN N HDP dp_ti_pool N - 46998
```

Description of the output of the **raidvchkscan command with the **-v pidb** or **-v pidb -fp** option:**

PID

Displays the pool ID.

POLS

Displays the status of the pool:

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is blocked due to a failure.) In this status, the pool information cannot be displayed.

U(%)

Displays the usage rate of the pool. The actual usage rate is displayed if the `-fp` option is specified.

LCNT

Displays the total number of Dynamic Provisioning virtual volumes mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

SSCNT

Displays the total number of snapshot data items mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

Available(MB)

Displays the available capacity (megabytes) for the volume data in the pool. Displays the actual pool capacity when the `-fp` option is specified.

Capacity(MB)

Displays the total capacity (megabytes) of the pool. Displays the actual pool capacity when the `-fp` option is specified.

Snap_Used(MB)

Displays the capacity (megabytes) used for Thin Image data. If the used capacity is less than 1 MB, the value is rounded up. A hyphen (-) is displayed if the information is not available for this pool.

If you entered the `-fp` option, a hyphen (-) is displayed.

TL_CAP(MB)

Displays the total capacity (megabytes) of all Dynamic Provisioning virtual volumes and Thin Image pairs mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

BM

Displays the I/O Blocking Mode of the pool:

- PF (Pool Full): If the pool is full, you cannot read from or write to the target DP-VOL. If the pool VOL is blocked, you can read from or write to the target DP-VOL.
- PB (Pool vol Blockade): If the pool VOL is blocked, you cannot read from or write to the target DP-VOL. If the pool is full, you can read from or write to the target DP-VOL.

- FB (Full or Blockade): If the pool is full or pool VOL is blocked, you cannot read from or write to the target DP-VOL.
- NB (No Blocking): If the pool is full or pool VOL is blocked, you can read from or write to the target DP-VOL.
- - (Not supported): The configuration does not support the I/O Blocking Mode.

TR_CAP(MB)

Displays the sum of the pool capacities (megabytes) reserved for the volumes for which Full Allocation or Proprietary Anchor is enabled. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

RCNT

Displays the number of volumes for which Full Allocation is enabled that are mapped to a pool. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

Seq#

Displays the serial number.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Num

Displays the number of LDEVs belonging to the pool.

LDEV#

Displays the number of the first LDEV in the pool. "65535 (ffff)" is shown while the pool is being created.

W(%)

Displays the threshold value for WARNING set for the pool. A hyphen (-) is displayed if the information is not available for this pool.

H(%)

Displays the threshold value set for the pool as the high water mark.

STIP

Displays the setting for Thin Image pairs when the high water mark threshold is exceeded.

- YES: Thin Image pairs are suspended.
- NO: Thin Image pairs are not suspended.
- - (hyphen): The information is not available for this pool.

VCAP(%)

Displays the percentage of the subscription limit of V-VOL and Thin Image pair to the pool capacity.

- UNLIMITED: Unlimited.
- - (hyphen): The information is not available for this pool.

For VSP G130, G/F350, G/F370, G/F700, G/F900, a hyphen (-) indicating invalid is displayed.

TYPE

Displays the platform type of the pool:

- OPEN: Pool for open systems
- M/F: Pool for mainframe systems

PM

Displays the pool status:

- N: Normal status.
- S: Shrinking or rebalancing.
- NT: The pool for Thin Image is in the normal status.
- ST: The pool for Thin Image is shrinking or rebalancing.

Check **D_PM** to view the detailed status of the pool.

PT

Displays the pool type:

- HDP: Pool for Dynamic Provisioning
- HDT: Pool for Dynamic Tiering
- RT: Pool for active flash
- TI: Pool for Thin Image
- CW: Pool for Copy-on-Write Snapshot
- DM: Pool for Dynamic Provisioning that has the data direct mapping attribute

POOL_NAME

Displays the pool name.

D_PM

Processing status of the pool

If multiple statuses coexist, the statuses are displayed separated by commas (.). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

- N: Processing of the pool is not running.
- WS: The pool is waiting to shrink.
- S: The pool is shrinking.
- WRB: The pool is waiting to rebalance.
- RB: The pool is rebalancing.
- E: The pool is being expanded due to DDP expansion.
- PE: The pool expansion due to DDP expansion has been suspended.
- - (hyphen): The information is not available for this pool.

D_PM(%)

Percentage of progress of the processing status of the pool

If multiple statuses coexist, the percentage of progress for each status is displayed separated by a comma (.). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported or if the percentage of progress is not valid.

EFFECTIVE(MB)

Effective capacity of the pool

A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

horcmstart

The **horcmstart** command is a script that starts HORCM. This script also sets the environment variables for HORCM as needed (for example, HORCM_CONF, HORCM_LOG, HORCM_LOGS).

When setting the instance number by specifying the HORCMINST environment variable, make sure that the setting value of the environment variable is in the range of 0 to 2047, and then execute **horcmstart**.

Syntax

```
horcmstart.sh { inst ... }    (UNIX systems)
horcmstart.exe { inst ... }   (Windows systems)
```

Options and parameters

inst

Specifies the HORCM instance number (range= from 0 to 2047). When this option is specified, the `horcmstart` shell script sets the environment variables (HORCMINST, HORCM_CONF, HORCM_LOG, HORCM_LOGS) corresponding to the instance number, and starts the specified HORCM instance. (Environment variables set by the user become invalid.) When this option is not specified, the `horcmstart` shell script starts 1 HORCM and uses the environment variables set by the user. If you have designated full environment variables, use `horcmstart.sh` without any arguments. If you did not designate environment variables (HORCM_CONF, HORCM_LOG, HORCM_LOGS), then this shell script sets the environment variables as follows:

For UNIX-based platforms:

- If HORCMINST is specified:
 - HORCM_CONF = /etc/horcm*.conf (* is instance number)
 - HORCM_LOG = /HORCM/log*/curlog
 - HORCM_LOGS = /HORCM/log*/tmplog
- If HORCMINST is not specified:
 - HORCM_CONF = /etc/horcm.conf
 - HORCM_LOG = /HORCM/log/curlog
 - HORCM_LOGS = /HORCM/log/tmplog

For Windows platform:

- If HORCMINST is specified:
 - HORCM_CONF = %windir%\horcm*.conf (* is the instance number)
 - HORCM_LOG = \HORCM\log*\curlog
 - HORCM_LOGS = \HORCM\log*\tmplog
- If HORCMINST is not specified:
 - HORCM_CONF = %windir%\horcm.conf
 - HORCM_LOG = \HORCM\log\curlog
 - HORCM_LOGS = \HORCM\log\tmplog

[environmental variable]: The HORCM_LOGS environment variable is used when specifying the log file directory for automatic storing. When HORCM starts up, the log files created in the operation are stored automatically in the HORCM_LOGS directory. This log directory must give an equality class with HORCM_LOG

HORCMSTART_WAIT (for waiting the CCI instance with start-up). Horcmgr does fork/exec() horcmd_XX as daemon process, and verifies/waits until HORCM become ready state. The timeout is used for only avoiding infinite loop, currently the default time is 200 sec in consideration of maximum LDEV. However, it may be needed to change the default timeout value for starting HORCM under high-loading of the server, or the remote command device. In such a case, this environmental variable is used when changing a timeout value (in seconds) from the current default value (200 sec), this value must be specified more than 5 seconds and multiple of 5 seconds. For example, setting 500 sec:

```
HORCMSTART_WAIT=500 Export HORCMSTART_WAIT
```

For OpenVMS platform: OpenVMS needs to make the Detached LOGINOUT.EXE Process as a JOB in the background by using the 'RUN /DETACHED' command. For details see **Requirements and restrictions for OpenVMS** (item 4) in the *Installation and Configuration Guide*.

horcmshutdown

The **horcmshutdown** command is a script for stopping HORCM.

Syntax

```
horcmshutdown.sh {inst...}      (UNIX systems)
horcmshutdown.exe {inst...}     (Windows systems)
```

Options and parameters

inst

Specifies the HORCM (CCI) instance number (range= from 0 to 2047). When this option is specified, the command stops the specified HORCM instance. When this option is not specified, the command refers to the instance (environment variable HORCMINST) of the execution environment of this shell script and stops the following the HORCM instance.

- When HORCMINST is specified, this command stops the HORCM instance of the execution environment of this shell script.
- When HORCMINST is not specified, this command stops the HORCM having no instance setting.



Caution: After direction of stopping HORCM instance, this command returns a response just before stopping HORCM instance. Thus to return the response of this command, it does not mean HORCM instance disappeared.

horcctl*

The HORCM software have logs that identify the cause of software and/or hardware errors as well as a tracing function for investigating such errors. The location of the log files depends on the user's command execution environment and the HORCM execution environment. The command trace file and core file reside together under the directory specified in the HORCM execution environment.

The **horcctl** command can be used for both maintenance and troubleshooting. The **horcctl** command allows you to change and display the internal trace control parameters (for example, level, type, buffer size) of the HORC Manager (CCI) software. If a new value for a parameter is not specified, the current trace control parameter is displayed.



Caution:

Do not change the trace level unless directed to do so by a Hitachi Vantara representative. Level 4 is the normal trace level setting. Levels 0-3 are for troubleshooting. Setting a trace level other than 4 may impact problem resolution. If you request a change of the trace level using the **horcctl -l <level>** command, a warning message appears, and this command enters interactive mode.

Syntax

```
horcctl { -h | -q | -z[x] | -I[H][M][instance#] or
        -I[TC][SI][instance#] | -d | -c | -l <level> | -b <y/n>
        | -s <size(KB)> | -t <type> | -S | -D[I] | -C
        | [-u <-unitid> | -ND | -NC | -g <group> ] }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **horcctl** command enter the interactive mode. The **-zx** option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shut down, interactive mode terminates.

OpenVMS cannot use the **-zx** option.

-l[H][M] [instance#] or -l[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -l[H][M] or -l[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -lH0

-d

Interprets the control options following this option (-l <level>, -b <y/n>, -s <size(KB)>, and -t <type>) as the parameters of the CCI commands.

-c

Interprets the control options following this option (-l <level>, -b <y/n> and -t <type>) as the parameters of the HORC Manager (HORCM).

-l <level>

Sets the trace level (range = 0 to 15). If a negative value is specified, the trace mode is canceled. A negative value 'n' must be specified as '--n'.

**Caution:**

Do not change the trace level unless directed to do so by a Hitachi Vantara representative. Level 4 is the normal trace level setting. Levels 0-3 are for troubleshooting. Setting a trace level other than 4 may impact problem resolution. If you request a change of the trace level using the **horcctl -l <level>** command, a warning message appears, and this command enters interactive mode.

-b <y/n>

Sets the trace writing mode: Y = buffer mode, N = synchronous mode.

-t <type>

Sets the trace type (range = 0 to 511). When this option is used, only traces of the specified type are output. One or more values can be specified.

-s <size(KB)>

Sets the trace buffer size in increments of 1024 bytes (default = 1 MB).

-S

Shuts down HORCM.

-D

Displays the command device name currently used by HORCM. If the command device is blocked due to online maintenance (microcode replacement) of the storage system, you can check the command device name in advance using this option.

-C

Changes the command device name being used by HORCM and displays the new command device name. If the command device is blocked due to online maintenance (microcode replacement) of the storage system, you can change the command device in advance using this option.

The **horcctl -D -C** command designates a protection mode command device by adding an asterisk (*) to the device file name as shown in the following example.

HP-UX example with command device security:

```
# horcctl -D
```

Current control device = /dev/rdisk/c0t0d0*

The **horcctl -DI** command shows the number of CCI instances of when HORCM has been started as shown in the following example.:

HP-UX example without command device security:

```
# horcctl -DI
```

Current control device = /dev/rdisk/c0t0d0 AI = 14 TI = 0 CI = 1, where

- AI: NUM of actual instances in use
- TI: NUM of temporary instances in the storage system
- CI: NUM of instances using current (own) instance

-u <unitid>

Used when specifying the unit ID of a command device as the target. This option is effective when the -D or -C option is specified. If this option is not specified, the unit ID is 0.

-ND -g <group>

Displays the network address and port name being used by HORCM. The -g <group> option is used when specifying the group name defined in the configuration definition file.

-NC -g <group>

Changes the network address and port name being used by HORCM and displays the new network address name. The -g <group> option specifies the group name defined in the configuration definition file.

horctakeoff

This is a scripted command for executing several HORC operation commands combined. It checks the volume attribute (optionally specified) and decides a takeover action. The **horctakeoff** operation is defined to change from 3DC multi-target to 3DC multi-hop with the state of running APP, after that the **horctakeover** command is able to configure 3DC multi-target on the remote site without stopping the APP. The granularity of either a logical volume or volume group can be specified with this command.

Syntax

```
horctakeoff | -h | -q | -z[x] | -I[H][M][instance#] or
-I[TC][SI][instance#] | -g[s] <group> | -d[s] <pair Vol>
| -d[g][s] <raw_device> [MU#] | -d[g][s] <seq#> <LDEV#> [MU#]
| -jp <id> | -js <id> | [-t <timeout> ] | -nomsg }
```

Options and parameters

-h

Displays help/usage and version information.

-q

Terminates the interactive mode and exits the command.

-z or -zx

Makes the **raidvchkdsp** command enter the interactive mode. The -zx option guards performing of the HORCM in the interactive mode. When this option detects a HORCM shutdown, interactive mode terminates.

OpenVMS cannot use the -zx option.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

-g[s] <group>

Specifies a group name (defined in the configuration definition file). The command is executed for the specified group unless the `-d <pair Vol>` option shown below is specified.

-d[s] <pair Vol>

Specifies a logical (named) volume (defined in the configuration definition file). When this option is specified, the command is executed for the specified paired logical volume.

-d[g][s] <raw_device> [MU#]

Searches the configuration definition file (local instance) for a volume that matches the specified raw device. If a volume is found, the command is executed on the paired volume (-d) or group (-dg). This option is effective without specification of the -g <group> option. If the specified raw_device is listed in multiple device groups, this applies to the first one encountered.

-d[g][s] <seq#> <LDEV#> [MU#]

Searches the configuration definition file (local instance) for a volume that matches the specified sequence # and LDEV. If a volume is found, the command is executed on the paired logical volume (-d) or group (-dg). This option is effective without specification of the -g <group> option. If the specified LDEV is listed in multiple device groups, this applies to the first one encountered. <seq#> <LDEV#> can be specified in a hexadecimal (by addition of '0x') or decimal.

**Note:**

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-jp <id> (HORC/UR only)

The **horctakeoff** command can change 3DC configuration from 3DC multi-target to 3DC multi-hop. To create a 3DC multi-hop (TC_Sync to TC_Sync/UR P-VOL to UR), you must specify a journal ID for UR P-VOL. So this option is used for that purpose. If this option is not specified, a journal ID for UR P-VOL used for 3DC multi-target is inherited automatically.

-js <id>(HORC/UR only)

The **horctakeoff** command can be changed 3DC configuration from 3DC multi-target to 3DC multi-hop. To create a 3DC multi-hop (TC_Sync to TC_Sync/UR to UR S-VOL), you must specify a journal ID for UR S-VOL. So this option is used for that purpose. If this option is not specified, a journal ID for UR S-VOL used with 3DC multi-target is inherited automatically. The CTG ID is also inherited automatically for the internal **paircreate** command.

-t <timeout>

Specifies the maximum time to wait for the Sync P-VOL to Sync S-VOL delta data resynchronizing operation. Used for the internal **pairresync** command with the timeout period in seconds. If this option is not specified, the default timeout value (7200 sec) is used.

-noms

Suppresses messages when this command is executed from a user program. This option must be specified at the beginning of the command arguments.

Returned values

The **horctakeoff** command sets the following returned values during exit allowing you to check the execution results.

- **Normal termination:**

- **0:**

- **Abnormal termination:**


- **other than 0:** The **horctakeoff** command returns the following error codes as well as generic error.

Error codes

Unrecoverable error should have been done without re-execute by handling of an error code. If the command failed, then the detailed status is logged in the CCI command log (\$HORCC_LOG), even though the user script has no error handling.

Category	Error Code	Error Message	Value
Volume status (Unrecoverable)	EX_ENQVOL	Unmatched volume status within the group	236
	EX_INCSTG	Inconsistent status in group	229
	EX_EVOLCE	Pair Volume combination error	235
	EX_VOLCRE	Local and Remote Volume currency error	223
Timer (Recoverable)	EX_EWSTOT	Timeout waiting for specified status	233

Example 1



horctakeoff command on L1 local site

```
# horctakeoff -g G1 -gs G2
```

```
horctakeoff : 'pairsplit -g G1 -S -FHORC 2' is in progress
horctakeoff : 'pairsplit -g G1' is in progress
horctakeoff : 'pairsplit -g G2 -S' is in progress
```




```

horctakeoff : 'paircreate -g G1 -gs G2 -FHORC 2 -nocopy -f async
-jp 0 -js 1' is in progress
horctakeoff : 'pairsplit -g G1 -FHORC 2' is in progress
horctakeoff : 'pairresync -g G1' is in progress
horctakeoff : 'pairresync -g G1 -FHORC 2' is in progress
horctakeoff : horctakeoff done

```

Example 2



horctakeoff command on L2 local site


```
# horctakeoff -g G1 -gs G3
```

```

horctakeoff : 'pairsplit -g G1 -S -FHORC 1' is in progress.
horctakeoff : 'pairsplit -g G1' is in progress.
horctakeoff : 'pairsplit -g G3 -S' is in progress.
horctakeoff : 'paircreate -g G1 -gs G3 -FHORC 1 -nocopy -f async
-jp 0 -js 1' is in progress.
horctakeoff : 'pairsplit -g G1 -FHORC 1' is in progress.
horctakeoff : 'pairresync -g G1' is in progress.
horctakeoff : 'pairresync -g G1 -FHORC 1' is in progress.
horctakeoff : horctakeoff done.

```


Example 3



horctakeoff command on L1 remote site

```
# horctakeoff -g G1 -gs G2
```

```
horctakeoff : 'pairsplit -g G2 -S' is in progress.
horctakeoff : 'pairsplit -g G1' is in progress.
horctakeoff : 'pairsplit -g G1 -FHORC 2 -S' is in progress.
horctakeoff : 'paircreate -g G2 -vl -nocopy -f async -jp 0
-js 1' is in progress.
horctakeoff : 'pairsplit -g G2' is in progress.
horctakeoff : 'pairresync -g G1' is in progress.
horctakeoff : 'pairresync -g G2' is in progress.
horctakeoff : horctakeoff done.
```

Example 4**horctakeoff** command on L2 remote site

```
# horctakeoff -g G1 -gs G3
```

```
horctakeoff : 'pairsplit -g G3 -S' is in progress.
horctakeoff : 'pairsplit -g G1' is in progress.
horctakeoff : 'pairsplit -g G1 -FHORC 1 -S' is in progress.
horctakeoff : 'paircreate -g G3 -vl -nocopy -f async -jp 0
-js 1' is in progress.
horctakeoff : 'pairsplit -g G3' is in progress.
horctakeoff : 'pairresync -g G1' is in progress.
horctakeoff : 'pairresync -g G3' is in progress.
horctakeoff : horctakeoff done.
```

Chapter 3: Subcommands

CCI provides subcommands for the Windows platform and environment variable subcommands.

Windows subcommands

The CCI subcommands for Windows are executed as an option of another command. When you specify a subcommand as the only option of a command, you do not need to start HORCM. If another option of the command and the subcommand are specified on the same command line, place the other option after the subcommand.

The Windows subcommands are **findcmddev**, **drivescan**, **portscan**, **sync(d)**, **mount**, and **umount(d)**.

findcmddev

The **findcmddev** subcommand searches for command devices within the specified range of disk drive numbers. If it is found, the command device appears in the same format as in the configuration definition file. This subcommand is used when the command device name is not known and when the HORCM is not started.



Caution: The **findcmddev** subcommand must be used when HORCM is running.



Note: The **findcmddev** subcommand searches for the physical and logical drives associated with the command device. If the command device is indicated as a logical drive in addition to a physical drive, then a drive letter is assigned to the command device. You must delete the drive letter assigned to the command device to prevent utilization by general users.

The 'Volume{GUID}' must be made by setting a partition using the disk management. Do not format it (no file system). In a SAN environment, the physical drive number might be changed on every reboot. For this case, use the Volume (GUID) that keeps as the same name.

Syntax

```
-x findcmddev drive#(0-N)
```

Argument**drive#(0-N)**

Specifies the range of disk drive numbers on the Windows system.

Example

The following shows an example of the **findcmddev** subcommand used as an option of the **raidscan** command and its output. This example searches for command devices in the range of disk drive numbers 0 through 20.

```
D:\HORCM\etc> raidscan -x findcmddev hdisk0, 20
cmddev of Ser# 62496 = \\.\PhysicalDrive0
cmddev of Ser# 62496 = \\.\E:
cmddev of Ser# 62496 = \\.\Volume{b9b31c79-240a-11d5-a37f-00c00d003b1e}
```

drivescan

The **drivescan** subcommand displays the relationship between the disk numbers assigned by the Windows system and the LDEVs on the RAID storage system, and also displays attribute and status information for each LDEV.

Syntax

```
-x drivescan drive#(0-N)
```

Argument**drive#(0-N)**

Specifies the range of data drive numbers on the Windows system.

Example

The following shows an example of the **drivescan** subcommand used as an option of the **raidscan** command and its output. This example displays the devices for the range of disk drive numbers from 0 to 20.

```
raidscan -x drivescan harddisk0,20
```

```
Harddisk 0..Port[1] PhId[0] TId[0] Lun[0] [HITACHI] [DK328H-43WS]
Harddisk 1..Port[2] PhId[4] TId[29] Lun[0] [HITACHI] [OPEN-3]
      Port[CL1-J] Ser#[ 30053] LDEV#[ 9(0x009)]
      HORC=P-VOL HOMRCF[MU#0 = SMPL MU#1 = SMPL MU#2 = SMPL]
      RAID5[Group 2- 1] SSID = 0x0008 CTGID = 3
Harddisk 2..Port[ 2] PhId[ 4] TId[29] Lun[ 1] [HITACHI] [OPEN-3]
      Port[CL1-J] Ser#[ 30053] LDEV#[ 10(0x00A)]
      HORC=S-VOL HOMRCF[MU#0 = SMPL MU#1 = SMPL MU#2 = SMPL]
      RAID5[Group 2- 1] SSID = 0x0004 CTGID = 3
```

```
Harddisk 3..Port[2] PhId[4] TId[29] Lun[ 6] [HITACHI] [OPEN-3-CM]
      Port[CL1-J] Ser#[ 30053] LDEV#[ 15(0x00F)]
```

Description of the **drivescan** subcommand output:

Harddisk #

Shows the data drive recognized by the Windows system.

Port

Shows the port number on the device adapter recognized by the Windows system.

PhId

Shows the bus number on the device adapter port recognized by Windows system.

TId

Shows the target ID of the data drive(s) connected to device adapter port. For the detail of fibre-to-SCSI address conversion, see the *Command Control Interface Installation and Configuration Guide*.

Lun

Shows the LU number of the data drive connected to device adapter port.

Port[CLX-Y]

Shows the port number on the storage system.

Ser#

Shows the production number (serial number) of the storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

LDEV#

Shows the LDEV ID (hexadecimal) of the specified volume.

HORC

Shows the TrueCopy/global-active device attribute (P-VOL, S-VOL, SMPL) of the specified volume.

HOMRCF

Shows the ShadowImage, Copy-on-Write Snapshot, or Thin Image attribute (P-VOL, S-VOL, or SMPL) and MU number (0-2) of the specified volume.

RAIDX[Group]

Shows the physical location (frame number-parity group number) of the specified volume and the RAID level of this parity group.

SSID

Shows the SSID (hexadecimal) of the specified volume.



Note: SSID is a parameter used by enterprise storage systems. Although SSID is not used by VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, or VSP G200, G400, G600, G800 and VSP F400, F600, F800, the set value is displayed.

CTGID (TrueCopy Async, Universal Replicator, global-active device only)

Shows the consistency group ID of specified volume.

portscan

The **portscan** subcommand displays the devices on the specified port(s).

Syntax

```
-x portscan port#(0-N)
```

Argument

port#(0-N)

Specifies the range of port numbers on the Windows system.

Example

The following shows an example of the **portscan** subcommand used as an option of the **raidscan** command and its output. This example displays the devices for the range of ports from 0 to 20.

```
raidscan -x portscan port0,20
```

```
PORT[ 0] IID [ 7] SCSI Devices
  PhId[ 0] TId[ 3] Lun[ 0] [MATSHIT] [CD-ROM CR-508]...Claimed
  PhId[ 0] TId[ 4] Lun[ 0] [HP      ] [C1537A      ]...Claimed
PORT[ 1] IID [ 7] SCSI Devices
  PhId[ 0] TId[ 0] Lun[ 0] [HITACHI] [DK328H-43WS ]...Claimed
PORT[ 2] IID [ 7] SCSI Devices
  PhId[ 0] TId[ 5] Lun[ 0] [HITACHI] [OPEN-3      ]...Claimed
  PhId[ 0] TId[ 5] Lun[ 1] [HITACHI] [OPEN-3      ]...Claimed
  PhId[ 0] TId[ 5] Lun[ 2] [HITACHI] [OPEN-3      ]...Claimed
  PhId[ 0] TId[ 6] Lun[ 0] [HITACHI] [3390-3A     ]...Claimed
```



Note: This example displays the devices for the range of ports from 0 to 20.

Description of the portscan subcommand output:

Port

Shows the port number on the device adapter recognized by the Windows system

IID

Shows the initiator ID on the specified device adapter port

Phid

Shows the BUS number on the specified device adapter port

Tid

Shows the target ID of the data drive(s) connected to device adapter port.

LUN

Shows the LU number of each data drive connected to device adapter port. This item shows LDEV# of the partner who becomes a pair in or among the RAID storage systems.

sync, syncd

The **sync** (synchronization) subcommand sends unwritten data remaining on the Windows server to the specified device(s) to synchronize the pair(s) before the CCI command is executed. The **syncd** (synchronization delay) subcommand waits for the delayed IO for dismount after issued '**sync**'.

Syntax

```
-x sync[d] A: B: C: ...
-x sync[d] all
-x sync[d] drive#(0-N)
-x sync[d] Volume#(0-N) ...
-x sync[d] D:\Directory or \Directory pattern...
```

Arguments

A: B: C: [\directory or \Directory pattern]

Specifies the logical drive that you want to synchronize. Data is flushed into the specified logical drive and the physical drive corresponding to the logical drive. If the specified logical drive has the directory mount volumes then SYNC is executed to all of the volumes on the logical drive as shown below:

```
pairsplit -x sync D:
[SYNC] D: HarddiskVolume2
[SYNC] D:\hd1 HarddiskVolume8
[SYNC] D:\hd2 HarddiskVolume9
```

[\directory or \Directory pattern] is used to find the directory mount point on the logical drive. If the directory is specified, then SYNC does execute to a directory mounted volume only.

```
pairsplit -x sync D:\hd1
[SYNC] D:\hd1 HarddiskVolume8
```

If the directory pattern is specified, then SYNC does execute to any directory mounted volumes identified to 'directory pattern'.

```
pairsplit -x sync D:\h
[SYNC] D:\hd1 HarddiskVolume8
[SYNC] D:\hd2 HarddiskVolume9
```

all

Synchronizes all logical drives and the physical drives corresponding to the logical drives assuming that they are on the data drives. The logical drive on which the CCI software is installed and the logical drive containing the Windows directory are excluded. If the logical drive has the directory mount volumes then SYNC is executed to all volumes on the logical drive as shown below:

```
pairsplit -x sync all
[SYNC] C: HarddiskVolume1
[SYNC] D:\hd1 HarddiskVolume8
[SYNC] D:\hd2 HarddiskVolume9
[SYNC] G: HarddiskVolume10
```

drive#(0-N)

Specifies the physical drives to be flushed.

Volume#(0-N)

Specifies the LDM volumes to be flushed. Volume# must be specified '\Vol#,\Dms#,\Dmt#,\Dmr# or Volume{...}' as LDM volume for Windows systems.

To flush HarddiskVolumeX:

```
-x sync \VolX
```

For information on '\Vol#,\Dms#,\Dmt#,\Dmr# or Volume{...}' for LDM volumes, see **Volume Discovery Function** in the *User and Reference Guide*.

Examples

The following examples show the **sync** subcommand used as an option of the **pairsplit** command.

sync subcommand with pairsplit

For the following example, the data remaining on logical drives C: and D: is written to disk, all pairs in the specified group are split (status = PSUS), and read/write access is enabled for all S-VOLs in the specified group.

```
pairsplit -x sync C: D: -g oradb -rw
```

sync subcommand with pairsplit -S

For the following example, the data remaining on physical devices `hddisk2` and `hddisk3` is written to disk, all pairs in the specified group are deleted (status = SMPL), which enables read/write access for all secondary volumes.

```
pairsplit -x sync hddisk2 hddisk3 -g oradb -S
```



Note: The **sync** subcommand has the following behavior on any conditions:

- If the logical drives designated as the objects of the **sync** subcommand are not opened to any applications, then **sync** flushes the system buffer to a drive and makes the dismount state for this drive.
- If the logical drives designated as the objects of the **sync** subcommand are already opened to any applications, then **sync** only flushes the system buffer to a drive. This flushes the system buffer before **pairsplit** without unmounting the P-VOL (opening state), and indicates as [WARNING] below:

```
pairsplit -x sync C:
WARNING: Only flushed to [\\.\C:] drive due to be opening
[SYNC] C: HarddiskVolume3
```



Note: The **syncd** subcommand flushes the system buffer to a drive and waits for 30 seconds.



Note: If **sync** has failed, you need to confirm the following conditions:

- The logical and physical drives designated as the objects of the **sync** subcommand are not opened to any applications. For example, confirm that Explore is not pointed on the target drive. If Explore is pointed on the target drive, the target drive is opening.
- The **sync** subcommand does not ignore the detected error on the NT file system, so **sync** executes successfully in normal case (NO ERROR case) only on NT file system. For example, confirm the target drive has no failure on the system for Event Viewer. In this case, you must reboot the system or delete the partition and reconfigure the target drive.

mount

The **mount** subcommand mounts the specified drive (indicates HDD, SSD, SCM, and FMD) to the specified partition on the specified data drive using the drive letter. When the **mount** subcommand is executed without an argument, all currently mounted drives (including directory mounted volumes) are displayed, and logical drive has been mounting an LDM volume then displays `Harddisk#[n]` configured an LDM volume.

Syntax

```
-x mount
-x mount drive: hddisk# [partition#] (for Windows)
```

```
-x mount drive: Volume#(0-N) (for Windows)
-x mount drive: [\directory] Volume#(0-N) (for Windows)
```

Arguments

drive: **hdisk#** [**partition #**]

Specifies the logical drive, data drive (number), and partition to be mounted.

drive: [**directory**] **Volume#**

Specifies the logical drive and LDM volume name and number to be mounted.

Volume# must be specified '\Vol#' or '\Dms#' or '\Dmt#' or '\Dmr#' as LDM volume for Windows.

To mount HarddiskVolumeX: -x mount C: hdX or -x mount C: \VolX

For information on '\Vol#' or '\Dms#' or '\Dmt#' or '\Dmr#' for LDM volumes, see **Volume Discovery Function** in the *User and Reference Guide*.

[directory]: Specifies the directory mount point on the logical drive.

pairsplit -x mount D:\hd1 \Vol8 D:\hd1 <+> HarddiskVolume8 pairsplit -x mount D:\hd2 \Vol9 D:\hd2 <+> HarddiskVolume9



Caution: The partition on the specified data drive must be recognized on the Windows system.

[directory] for the mount must be specified a mount point without embedded space character.

If [directory] is detected as mount point with embedded space (that is, aaabbb), then the directory is shown by adding '...' to first strings as below.

pairsplit -x mount

Drive FS_name VOL_name Device Partition ... Port PathID Targ Lun

D: NTFS Null Harddiskvolume3 ... Harddisk2

D:\aaa... NTFS Null Harddiskvolume4 ... Harddisk3

The same method is used for the 'in RAID \$LETALL' and 'raidscan -pi \$LETALL -find' commands.

Example

The following example shows the **mount** subcommand used as an option of the **pairsplit** command and its outputs.

Examples for Windows

This example executes the **mount** subcommand from a subcommand option of **pairsplit**. It mounts the "F:" drive to the **harddiskvolume2**, and then displays mounted devices.

```
pairsplit -x mount F: hdisk2
pairsplit -x mount
```

Drive	FS_name	VOL_name	Device	Partition	...	Port	PathID	Targ	Lun
C:	NTFS	Null	Harddiskvolume1	...	Harddisk0				
F:	NTFS	Null	Harddiskvolume2	...	Harddisk1				
D:	NTFS	Null	Harddiskvolume3	...	Harddisk2				
D:\hd1	NTFS	Null	Harddiskvolume4	...	Harddisk3				
D:\hd2	NTFS	Null	Harddiskvolume5	...	Harddisk4				
G:	NTFS	Null	HarddiskDmVolumes\...\Volume1...		Harddisk5[3]				

This example executes **mount** from command option of the **pairsplit** command and then displays the mounted devices:

Description of the mount subcommand output:

Drive

Shows the logical drive recognized by the Windows system

FS_name

Shows the name of the file system formatted on the logical drive

VOL_name

Shows the volume label name for the logical drive

Device, Partition

Shows the device name and partition for the mounted logical drive

Port,PathID,Targ,Lun

Shows the port number, path group ID (bus), target ID, and LUN for the device adapter of the mounted logical drive.

umount, umountd

The **umount** subcommand unmounts the specified logical drive and deletes the drive letter. Before deleting the drive letter, this subcommand executes sync internally for the specified logical drive and flushes unwritten data. The **umountd** subcommand unmounts the logical drive after waiting the delayed IO for dismount.

Syntax

```
-x umount[d] drive: [time]

-x umount[d] drive:[\directory] [time] (for Windows)
```

Argument**drive**

Specifies the mounted logical drive.

[directory] [time]

Specifies the directory mount point on the logical drive. This command option calls 'mountvol /P' internally, if 'USE_MOUNTVOL_P' environment variable is specified. In case of Windows, it is required to specify 'USE_MOUNTVOL_P' variable to avoid a problem of mount.

```
pairsplit -x umount D:\hd1 D:\hd1 <-> HarddiskVolume8 set USE_MOUNTVOL_P=1
pairsplit -x umount D:\hd2 D:\hd2 <-> HarddiskVolume9
```

Example for waiting 45 sec:

```
pairsplit -x umount D: 45
```

```
D: <-> HarddiskVolume8
```



Caution: **umountd** has above restriction and it will prompt the delayed IO for dismount. Wait 30 seconds until the completion and release the mount point after making dismount status of the logical drive. This avoids a problem (Windows 2003 only) that the delayed writing for dismount as Event ID51, 57.

Example

The following example shows the **umount** subcommand used as an option of the **pairsplit** command.

```
pairsplit -x umount F: -x umount G: -g oradb -rw
pairsplit -x mount
```

Drive	FS_name	VOL_name	Device	Partition...	Port	PathID	Targ	Lun
C:	FAT	Null	Harddisk0	Partition1...	1	0	0	0
Z:	Unknown	Unknown	CdRom0	...	Unknown			

This example unmounts the F: and G: drives, splits all pairs in the specified group (status = PSUS), enables read/write access to all secondary volumes in the specified group, and then displays all mounted drives.

Description of the umount subcommand output:**Drive:**

Shows the logical drive recognized by the Windows system

FS_name

Shows the name of the file system formatted on the logical drive

VOL_name

Shows the volume label name for the logical drive

Device, Partition

Shows the device name and partition for the mounted logical drive

Port,Phid,Targ,Lun

Shows the port number, path group ID (bus), target ID, and LUN for the device adapter of mounted logical drive.

**Note:**

The **umount** subcommand flushes (sync) the system buffer of the associated drive before deleting the drive letter. If **umount** fails, confirm the following conditions:

- The logical and physical drives designated as the objects of the **umount** subcommand are not opened to any applications. For example, confirm that Explore is not pointed on the target drive. If it is, then the target drive is opening.
- The **umount** subcommand does not ignore the detected error on the NT file system, so that **umount** is successful in a normal case (NO ERROR case) only on NT file system. For example, confirm the target drive has no failure on the system for Event Viewer. If so, you must reboot the system or delete the partition and reconfigure the target drive.

**Note:**

The **umountd** subcommand has the following behavior as well.

Unmount the logical drive after waiting (30 sec) the delayed (paging) IO for dismount after flushed the system buffer to a drive.

This avoids a problem (Windows 2003 only) that NTFS on P-VOL is split on inconsistent state because Windows 2003 (SP1) delays the IO for dismounting. This also avoids a problem that the delayed (paging) IO for dismounting is written on SVOL_PAIR(Write Disable) state by rescan, and logged as Windows event (that is, ID51,57).

These problems do not occur on Windows 2008 and later systems.

Environment variable subcommands

If no environment variables are set in the execution environment, the following environment variable subcommands set or cancel an environment variable within the CCI command.

- **setenv**: The **setenv** subcommand sets the specified environment variable(s).
- **unsetenv**: The **unsetenv** subcommand deletes the specified environment variable(s).
- **env**: The **env** subcommand displays the environment variable(s).
- **sleep**: The **sleep** subcommand causes CCI to wait for the specified time.

Syntax

```
-x setenv varname value
-x unsetenv varname
-x env
-x sleep time
```

Arguments

Varname

Specifies the environment variable to be set or canceled.

Value

Specifies the value or character string of the environment variable to be set.

Time

Specifies the sleep time in seconds.



Caution: The environment variables must be set before connecting to HORCM. And it must be specified during interactive mode (`-z [x]` option). If specified with other than interactive mode, all specified environment variables are not enabled. Changing an environment variable after a CCI command execution error is invalid.

Example

The following examples show the **setenv** and **unsetenv** subcommands used as an option of the **raidscan** command. This example changes from 'HORC' to 'HOMRCF' an execution environment of the **raidscan** command that makes a dialog mode, because of establishing 'HORCC_MRCF' as an environment variable.

```
raidscan[HORC]: -x setenv HORCC_MRCF 1
raidscan[HOMRCF]:
```

```
raidscan[HOMRCF]: -x unsetenv HORCC_MRCF
raidscan[HORC]:
```

Chapter 4: Command tools

CCI provides the following command tools:

- inqraid
- mkconf
- rmawk

inqraid

The **inqraid** command is a tool used to confirm the drive connection between the storage system and host system. The **inqraid** command displays the relationship between special file(s) on the host system and actual physical drive of the RAID storage system.

Syntax

```
/HORCM/usr/bin/inqraid [-h | quit | -inqdump  
| -fx[p][l][g][c][h][n] | -find[c] | <special file>  
| -CLI[W][P][N][B] | -sort [-CM ] | -export][-n]  
  
/HORCM/etc/inqraid [-h | quit | -inqdump | -fx[p][l][g][c][h][n]  
| -find[c] | <special file> | -CLI[W][P][N][B] | -sort [-CM ]  
| -gvinf | -svinf | -gplba | -fv| -export]
```

Options and parameters

-h

Displays Help/Usage.

quit

Terminates from waiting STDIN and exits this command.

-inqdump

Displays information for standard inquiry with Dump Image of hexadecimal.

-fx

Displays the LDEV number (hexadecimal).

-find[c]

Finds the appropriate group within the configuration file using a special file provided by STDIN.

- **-find:** Searches a group on the configuration definition file (local instance) from <special file> of STDIN by using **pairedisplay** command, and uses the following options of the **pairedisplay** command to display its state. This option must be specified HORCMINST as command execution environment.

For ShadowImage/Copy-on-Write Snapshot/Thin Image:

```
pairedisplay -d <Seq#> <LDEV#> 0 1 2 -l [-fx] [-CLI] 2>/dev/null
```

For TrueCopy/Universal Replicator/global-active device:

```
pairedisplay -d <Seq#> <LDEV#> -l [-fx] [-CLI] 2>/dev/null
```

**Note:**

<Seq#> and <LDEV#> are included using SCSI Inquiry command.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
 - When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.
- **<special file>:** This option is used to specify the special file name as argument of command. If no argument, this command makes mode that waits for STDIN. The input from the special file has to be waited.

- **-findc:** Uses the following options of the **pairstdisplay** command, and displays with CLI format by editing an output of **pairstdisplay** command. This option must be specified HORCMINST as command execution environment.

For ShadowImage/Copy-on-Write Snapshot/Thin Image:

```
pairstdisplay -d <Seq#> <LDEV#> <MU#> -fd -CLI 2>/dev/null
```

For TrueCopy/Universal Replicator/global-active device:

```
pairstdisplay -d <Seq#> <LDEV#> -fd -CLI 2>/dev/null
```

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
 - When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.
- **<special file>:** Specifies a special file name as the argument of a command. If no argument, this command makes mode that waits for STDIN. The input from the special file has to be waited.

-CLI

Specifies the display of structured column output for command line interface (CLI) parsing. Also used for '-find' option. The delimiters between columns can be spaces and/or dashes (-).

-CLIWP, -CLIWN

Displays the WWN (world wide name for host adapter) with command line interface (CLI) format, also used for '-find' option.

-sort[CM]

Sorts the target devices by Serial#,LDEV# order.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

[CM] Searches command device from the specified special file (STDIN or argument) and displays the command device only in structure definition file image. This option is valid within '-sort' option.

-gvinf or -gvinfex

Windows systems only. The -gvinfex option is for GPT disk only.

Gets the signature and volume layout information of a raw device file provided via STDIN or arguments, and saves this information to the system disk with the following format:

\\WindowsDirectory\\VOLssss_IIII.ini

where ssss = serial#, llll = LDEV#

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Normally this option is used by the Windows Disk Management after setting the signature and volume layout information for S-VOL. You do not need to be aware of this file.

-svinf[=PTN] or -svinfex[=PTN]

Windows systems only. The -svinfex[=PTN] option is for GPT disk only.

Reconfigure the signature and volume layout information that was saved to the system disk to a device provided by STDIN or arguments. Gets the serial# and LDEV# for the target device issuing SCSI Inquiry, and sets the signature and volume layout information into VOLssss_ llll.ini file to the readout device. This option will set correctly because the signature and volume layout information is managed by the serial# and LDEV# without depend on Harddisk#, even if Harddisk# is changed by the configuration changes.

ssss = serial#

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

[=PTN]: Specifies a strings pattern to interpret the strings provided by STDIN or argument as a raw device.

\Device\HarddiskVolume#(number) is made in a sequential order executed -svinf to Harddisk, and its number will remain the same as long as the system configuration is not changed. If you want to make \Device\HarddiskVolume#(number) more absolutely, then make \Device\HarddiskVolume# in serial# and LDEV# order by using the '-sort' option as shown below:

[VOL61459_451_5296A763] -> Harddisk3 [OPEN-3]

[VOL61459_452_5296A760] -> Harddisk4 [OPEN-3]

[VOL61459_453_5296A761] -> Harddisk5 [OPEN-3]

-gplba or -gplbaex

Windows systems only. The -gplbaex option is for GPT disk only.

Displays usable LBA on a physical drive in units of 512 bytes, and specifies [slba] [elba] options for **raidvchkset** command.

Example:

```
C:\HORCM\etc>inqraid $Phys -CLI -gplba -sort
Harddisk11 : SLBA = 0x00003f00 ELBA = 0x000620d9 PCNT = 7 [OPEN-3-CVS ]
Harddisk12 : SLBA = 0x00003f00 ELBA = 0x00042ad1 PCNT = 4 [OPEN-3-CVS ]
Harddisk13 : SLBA = 0x0000003f ELBA = 0x000620d9 PCNT = 1 [OPEN-3-CVS ]
```

SLBA: Displays usable starting LBA in units of 512 bytes. ELBA: Displays usable ending LBA (ELBA -1) in units of 512 bytes. PCNT: Displays the number of partitions.

Example for setting of Harddisk11: C:\HORCM\etc>raidvchkset -d hd11 -vs 16 0x00003f00 0x000620d9

-fv

Windows systems only.

Displays the Volume{GUID} via \$Volume with wide format. The serial number (serial#) is SERIAL.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Example:

```
C:\HORCM\etc>inqraid -CLI $Vol -fv
```

DEVICE_FILE	PORT	SERIAL	LDEV CTG	H/M/12	SSID R:Group	PRODUCT_ID
Volume{cec25efe-d3b8-11d4-aead-00c00d003b1e}					CL2-D	62496 256
-	-	-	-		OPEN-3-CVS-CM	

-fl

If the target device file is set as a protection volume, the device file name is shown by appending '*'. It is valid when specified with "-CLI" option. If the -fl option is specified, the data protection volume is a Data Retention Utility volume.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

-fg

Shows a LUN on the host view by finding a host group.

-CLIB -sort [-fh | -fc]

Displays the number of the tables of the differential bitmap which are required on the shared memory. This option is used to find the number of pairs that can be created to a paired volume in the storage system. Specify this option with the -sort option to sort the specified special files (the standard input or the argument) in order of the following priority, Serial#, and then LDEV#.

- When you specify `-fh` option: the number of tables of the differential bitmap for TC, TCz, UR, and GAD is displayed.
- When you specify `-fc` option: the number of tables of the differential bitmap in a cylinder size for TC, TCz, UR, and GAD is displayed.
- When you omit the `-fh` option and `-fc` option: the number of tables of the differential bitmap for SI and Slz is displayed.

Example to display the number of tables of the differential bitmap for SI and Slz:

```
# ls /dev/rdisk/* | inqraid -CLIB -sort
DEVICE_FILE PORT  SERIAL LDEV  SL CL +SI/SI UNUSED PRODUCT_ID
c1t0d0      CL1-E 63516  0      0 0 - - OPEN-9-CM
c1t0d1      CL1-E 63516 12288  0 0 1 30718 OPEN-3
c1t0d2      CL1-E 63516 12403  0 0 4 30718 OPEN-9
c1t0d3      CL1-E 63516 12405  0 0 9 30718 OPEN-E
c1t0d4      CL1-E 63516 12800  0 0 12 30718 OPEN-8
c1t0d5      CL1-E 63516 12801  0 0 18 30718 OPEN-8*2
c1t0d6      CL1-E 63516 13057  0 0 31 30718 OPEN-L
c2t0d6      CL2-E 63516 13057  0 0 31 30718 OPEN-L
```

Example to display the number of tables of the differential bitmap for TC, TCz, UR, and GAD:

```
# ls /dev/rdisk/* | inqraid -CLIB -sort -fh
DEVICE_FILE PORT  SERIAL LDEV  SL CL +TC/UR UNUSED PRODUCT_ID
c1t0d0      CL1-E 63516  0      0 0 - - OPEN-9-CM
c1t0d1      CL1-E 63516 12288  0 0 1 11605 OPEN-3
c1t0d2      CL1-E 63516 12403  0 0 3 11605 OPEN-9
c1t0d3      CL1-E 63516 12405  0 0 10 11605 OPEN-E
c1t0d4      CL1-E 63516 12800  0 0 11 11605 OPEN-8
c1t0d5      CL1-E 63516 12801  0 0 13 11605 OPEN-8*2
c1t0d6      CL1-E 63516 13057  0 0 21 11605 OPEN-L
c2t0d6      CL2-E 63516 13057  0 0 21 11605 OPEN-L
```

- SERIAL: Serial number
- SL: The SLPR number of LDEV.
- CL: The CLPR ID of LDEV.

- **+SI/Sl:** Shows the accumulated total number of tables of the differential bitmap for SI and Slz. The increased number shows necessary tables of the differential bitmap for one SI or Slz pair .

**Note:**

The number of tables of the differential bitmap for the following LDEV are excepted from the accumulated total.

- LDEVs whose Serial# and LDEV# are same (they are calculated as one LDEV if the multiple ports share the LDEV).
- An LDEV which does not manage differential bitmap on a shared memory.
- An LDEV which is used as a command device.

- **+TC/UR:** Shows the accumulated total number of tables of the differential bitmap for TC, TCz, UR, and GAD. The increased number shows necessary tables of the differential bitmap for TC, TCz, UR, and GAD. The '-fc' option displays the number of tables of differential bitmap in a cylinder size.

**Note:**

The number of tables of the differential bitmap for the following LDEVs are not included in the accumulated total:

- LDEVs whose Serial# and LDEV# are same (they are calculated as one LDEV if the multiple ports share the LDEV).
- An LDEV which does not manage differential bitmap on a shared memory.
- An LDEV which is used as a command device.

- **UNUSED:** Shows the number of tables of unused differential bitmap for SI, Slz, TC, TCz, UR, and GAD.

-CLI -fn

Displays the LDEV nickname in the PRODUCT ID. If the storage system does not support LDEV nickname, '-' is displayed. The serial number (serial#) is SERIAL.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Example:

```
# ls /dev/rdisk/c57t4* | ./inqraid -CLI -fn
DEVICE_FILE PORT  SERIAL LDEV CTG C/B/12 SSID R:Group LDEV_NIC_NAME
c57t4d0      CL1-D 62496 32   -   s/P/ss 0004 5:01-03 my_volume_1
c57t4d3      CL1-D 62496 35   -   s/P/ss 0004 5:01-03 my_volume_2
```

```
c57t4d4    CL1-D 62496 36 - s/P/ss 0004 5:01-01 -
c57t4d5    CL1-D 62496 37 - s/P/ss 0004 5:01-02 -
```

-CLI -export

This option outputs as export format after obtaining the storage system device information from the specified special file (STDIN or argument). The device information includes 'Keyword, Serial#, Ldev#, Device file name, ...'. This export formatted file can import as STDIN of the raidscan -find verify or the raidscan -find inst. If the application server and CCI server are operated on different hosts, volume discovery can be operated with using this option among the hosts.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Example:

```
# ls /dev/rdisk/clt* | inqraid -CLI -export
INQRAID:@CL4-G@64015@0@124@OPEN-V-CM@/dev/rdisk/clt0d0s2
INQRAID:@CL4-G@64015@1@124@OPEN-V-CM@/dev/rdisk/clt0d1s2
INQRAID:@CL4-G@64015@2@95@OPEN-V@/dev/rdisk/clt0d2s2
INQRAID:@CL4-G@64015@3@95@OPEN-V@/dev/rdisk/clt0d3s2
INQRAID:@CL4-G@64015@4@95@OPEN-V@/dev/rdisk/clt0d4s2
INQRAID:@CL4-G@64015@5@95@OPEN-V@/dev/rdisk/clt0d5s2
INQRAID:@CL4-G@64015@7@95@OPEN-V@/dev/rdisk/clt0d7s2
```

Restrictions on device naming

STDIN or special files are specified as follows:

- HP-UX:
 - /dev/rdisk/*
 - /dev/rdisk/disk*
- Solaris:
 - /dev/rdisk/*s2
 - c*s2
- Linux:
 - /dev/sd...
 - /dev/rd...
 - /dev/raw/raw*

- zLinux:
 - /dev/sd...
 - /dev/dasd...
 - /dev/rd...
 - /dev/raw/raw*
- AIX:
 - /dev/rhdisk*
 - /dev/hdisk*
 - hdisk*
- DIGITAL or Tru64:
 - /dev/rrz*c
 - /dev/rdisk/dsk*c
 - /dev/cport/scp*
- DYNIX:
 - /dev/rdisk/sd*
 - sd* for only unpartitioned raw device
- IRIX64:
 - /dev/rdisk/*vol
 - /dev/rdisk/node_wwn/*vol/*
 - /dev/dsk/*vol
 - /dev/dsk/node_wwn/*vol/*
- Windows: For information about LDM volumes for Windows systems, see **Remote Volume Discovery** in the *Command Control Interface User and Reference Guide*.
 - hdX-Y
 - \$LETALL
 - \$Volume
 - \$Phys
 - D:\Vol(Dms,Dmt,Dmr)X\DskY
 - \Vol(Dms,Dmt,Dmr)X\DskY
 - D:\DskX\pY
 - \DskX\pY

Lines starting with '#' via STDIN are interpreted as comments.

[-n]

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

Example 1: using `inqraid` and `system` command to display the connection between STDIN special file and actual physical drive of storage system

HP-UX system:

```
# ioscan -fun | grep rdsd | ./inqraid
/dev/rdsd/c0t2d1 -> [HP] CL2-D Ser = 30053 LDEV = 9 [HP ]
      [ OPEN-3 ] HORC = P-VOL  HOMRCF[MU#0 = SMPL MU#1 = SMPL
      MU#2 = SMPL] RAID5[Group 2- 1] SSID = 0x0008 CTGID = 3
/dev/rdsd/c0t4d0 -> [HP] CL2-D Ser = 30053 LDEV = 14 [HP ]
      [OPEN-3-CM ] RAID5[Group 2- 1] SSID = 0x0008
```

Linux and zLinux system:

```
# ls /dev/sd* | ./inqraid
/dev/sdh -> CHNO = 0 TID = 1 LUN = 7
      [HP] CL2-B Ser = 30053 LDEV = 23 [HP ]
      [OPEN-3 ] HORC = P-VOL  HOMRCF[MU#0 = SMPL
      MU#1 = SMPL MU#2 = SMPL] RAID5[Group 1- 2]
      SSID = 0x0004 CTGID = 2
/dev/sdi -> CHNO = 0 TID = 4 LUN = 0
      [HP] CL2-B Ser = 30053 LDEV = 14 [HP ]
      [OPEN-3-CM ] RAID5[Group 1- 2] SSID = 0x0004
```

Linux system (namespace for NVMe-oF):

Namespace = Y is displayed if a namespace is used.

```
[root@localhost ~]# ls /dev/nvme4n1 | /HORCM/usr/bin/inqraid -n
/dev/nvme4n1 -> [ST] CL1-D Ser = 500054 LDEV =1536 [HITACHI ] [OPEN-V-CM ]
      A-LUN[PoolID 0002] SSID = 0x0011 Namespace = Y
```

Solaris system:

```
# ls /dev/rdsd/* | ./inqraid
/dev/rdsd/c0t2d1 -> [HP] CL2-D Ser = 30053 LDEV = 9 [HP ]
      [OPEN-3 ] CA = P-VOL  BC[MU#0 = SMPL MU#1 = SMPL
      MU#2 = SMPL] RAID5[Group 2- 1] SSID = 0x0008
      CTGID = 3
/dev/rdsd/c0t4d0 -> [HP] CL2-D Ser = 30053 LDEV = 14 [HP ]
      [OPEN-3-CM ] RAID5[Group 2- 1] SSID = 0x0008
```


AIX system:

```
# lsdev -C -c disk | grep hdisk | ./inqraid
hdisk1 -> [SQ] CL2-D Ser = 30053 LDEV = 9 [HITACHI ] [OPEN-3 ]
          HORC = P-VOL HOMRCF[MU#0 = SMPL MU#1 = SMPL
          MU#2 = SMPL] RAID5[Group 2- 1] SSID = 0x0008 CTGID = 3
hdisk2 -> [SQ] CL2-D Ser = 30053 LDEV = 14 [HITACHI ] [OPEN-3-CM ]
          RAID5[Group 2- 1] SSID = 0x0008
```

Windows system:

```
C:\HORCM\etc> echo hd1-2 | inqraid ( or inqraid hd1-2 )
Harddisk 1 -> [SQ] CL2-D Ser = 30053 LDEV = 9 [HITACHI ] [OPEN-3 ]
              HORC = P-VOL HOMRCF[MU#0 = SMPL MU#1 = SMPL
              MU#2 = SMPL] RAID5[Group 2- 1] SSID = 0x0008
              CTGID = 3
Harddisk 2 -> [SQ] CL2-D Ser = 30053
              LDEV = 14 [HITACHI ] [OPEN-3-CM ] RAID5[Group 2- 1]
              SSID = 0x0008
```

Tru64 UNIX system:

```
# ls /dev/rdisk/dsk* | ./inqraid
/dev/rdisk/dsk10c -> [SQ] CL2-D Ser = 30053 LDEV = 9 [HITACHI ]
                    [OPEN-3 ] HORC = P-VOL HOMRCF[MU#0 = SMPL
                    MU#1 = SMPL MU#2 = SMPL] RAID5[Group 2- 1]
                    SSID = 0x0008 CTGID = 3
/dev/rdisk/dsk11c -> [SQ] CL2-D Ser = 30053 LDEV = 14 [HITACHI ]
                    [OPEN-3-CM ] RAID5[Group 2- 1] SSID = 0x0008
```

IRIX system with FC_AL:

```
# ls /dev/rdisk/*vol | ./inqraid
/dev/rdisk/dks1d6vol -> [SQ] CL2-D Ser = 30053 LDEV = 9 [HITACHI ]
                        [OPEN-3 ] HORC = P-VOL
                        HOMRCF[MU#0 = SMPL MU#1 = SMPL MU#2 = SMPL]
                        RAID5[Group 2- 1] SSID = 0x0008 CTGID = 3
/dev/rdisk/dks1d7vol -> [SQ] CL2-D Ser = 30053 LDEV = 14 [HITACHI ]
                        [OPEN-3-CM ]
                        RAID5[Group 2- 1] SSID = 0x0008
```

IRIX system with fabric:

```
# ls /dev/rdisk/*/*vol/* | ./inqraid
/dev/rdisk/50060e8000100262/lun3vol/c8p0 -> [SQ] CL2-D Ser = 30053
          LDEV = 9 [HITACHI] [OPEN-3 ] HORC = P-VOL
          HOMRCF[MU#0 = SMPL MU#1 = SMPL MU#2 = SMPL]
          RAID5[Group 2- 1] SSID = 0x0008 CTGID = 3
/dev/rdisk/50060e8000100262/lun4vol/c8p0 -> [SQ] CL2-D Ser = 30053
```

```
LDEV = 14 [HITACHI] [OPEN-3-CM]
RAID5[Group 2- 1] SSID = 0x0008
```

OpenVMS system:

```
$ inqraid dka145-146
DKA145 -> [ST] CL2-D Ser = 30053 LDEV = 9 [HITACHI] [OPEN-3 ]
          HORC = P-VOL HOMRCF[MU#0 = SMPL MU#1 = SMPL
          MU#2 = SMPL] RAID5[Group 2- 1] SSID = 0x0008 CTGID = 3
DKA146 -> [ST] CL2-D Ser = 30053 LDEV = 14 [HITACHI] [OPEN-3-CM ]
          RAID5[Group 2- 1] SSID = 0x0008
```

Description of the inqraid command tool output for example 1:

CLX -Y:

Displays the port number on the RAID storage system.

Ser

Displays the production (serial#) number on the RAID storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

LDEV

Displays the LDEV# in the RAID storage system.

HORC

Displays the attribute ('P-VOL/S-VOL/SMPL') of a volume in the RAID storage system.

HOMRCF

Displays the attribute ('P-VOL/S-VOL/SMPL') of a volume as MU#0-2 of ShadowImage, Copy-on-Write Snapshot, or Thin Image in the RAID storage system.

Group

Displays the relationship of the physical volume mapped to LDEV.

LDEV Mapping	Display Formats
RAID Group	RAID1[Group Group number - Sub number] RAID5[Group Group number - Sub number] RAID6[Group Group number - Sub number]
Copy-on-Write Snapshot S-VOL	SNAPS[PoolID poolID number]
Unmapped	UNMAP[Group 00000]
External LUN	E-LUN[Group External Group number]

LDEV Mapping	Display Formats
Dynamic Provisioning volume	A-LUN[PoolID poolID number]

SSID

Displays the storage subsystem ID (hexadecimal) of the LDEV in the RAID storage system.



Note: SSID is a parameter used by enterprise storage systems. Although SSID is not used by VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, or VSP G200, G400, G600, G800 and VSP F400, F600, F800, the set value is displayed.

CTGID

Displays the consistency group ID of TrueCopy Async and Universal Replicator when the LDEV has been defined as the P-VOL or S-VOL of the TrueCopy Async or Universal Replicator.

CHNO

Displays the channel number on the device adapter that recognizes on the Linux host. Displayed only for Linux systems.

TID

Displays target ID of the data drive that connects on the device adapter port. Displayed only for Linux systems.

LUN

Displays logical unit number of the data drive that connects on the device adapter port. Displayed only for Linux systems.

**Note:**

The display of group, SSID, and CTG ID depends on the storage system microcode level. The **CHNO**, **TID**, and **LUN** items are displayed only for Linux systems.

Example 2: inqraid with -find option (Linux shown)

```
ls /dev/sd* | inqraid -find
/dev/sdb -> No such on the group
Group PairVol (L/R) (Port#,TID,LU),Seq#,LDEV#.P/S, Status,Fence,
Seq#,P-LDEV# M
oradb oradev2 (L) (CL2-N , 3, 2) 8071 22..SMPL -----,
----- ---- -
->/dev/sdc
```

Example 3: inqraid with -find option (HP-UX shown)

```
# echo /dev/rdisk/c23t0d0 /dev/rdisk/c23t2d3 | ./inqraid -find
Group   PairVol (L/R) (Port#,TID,LU-M),Seq#,LDEV#.P/S,Status,
Seq#,P-LDEV# M
horc1   dev00 (L)      (CL2-J, 0, 0-0) 61456 192..S-VOL SSUS,
----- 193 -
->/dev/rdisk/c23t0d0
Group   PairVol (L/R) (Port#,TID,LU-M),Seq#,LDEV#.P/S,Status,
Seq#,P-LDEV# M
horc1   dev10 (L)      (CL2-J , 2, 3-0) 61456 209..S-VOL SSUS,
----- 206 -
->/dev/rdisk/c23t2d3
```

Example 4: inqraid with -findc option (HP-UX shown)

```
# echo /dev/rdisk/c23t0d0 /dev/rdisk/c23t2d3 | ./inqraid -findc
DEVICE_FILE M Group PairVol P/S   Stat R_DEVICE M P/S   Stat LK
c23t0d0      0 horc1   dev00 S-VOL SSUS  c23t0d1 0 P-VOL PSUS OK
/dev/rdisk/c23t0d0[1] -> No such on the group
/dev/rdisk/c23t0d0[2] -> No such on the group
DEVICE_FILE M Group PairVol P/S   Stat R_DEVICE M P/S   Stat LK
c23t2d3      0 horc1   dev10 S-VOL SSUS  c23t2d2 0 P-VOL PSUS OK
/dev/rdisk/c23t2d3[1] -> No such on the group
/dev/rdisk/c23t2d3[2] -> No such on the group

# echo /dev/rdisk/c23t0d0 /dev/rdisk/c23t2d3 | ./inqraid -findc -CLI
DEVICE_FILE M Group PairVol P/S   Stat R_DEVICE M P/S   Stat LK
c23t0d0      0 horc1   dev00 S-VOL SSUS  c23t0d1 0 P-VOL PSUS OK
c23t2d3      0 horc1   dev10 S-VOL SSUS  c23t2d2 0 P-VOL PSUS OK
```

Description of the inqraid tool output for example 2 to 4:**DEVICE_FILE**

Device file name.

M:

MU# of local and remote.

Group

Group name (dev_group) defined in the configuration file.

PairVol

Paired vol. name (dev_name) within the group defined in the configuration file.

P/S

Volume attribute (P-VOL or S-VOL or simplex).

Stat

Status of the paired volume.

R_DEVICE

Device file name of remote site.

LK

Check result of the paired volume connection path.

Example 5: inqraid with -cli option (Linux shown)

```
# ls /dev/sd* | ./inqraid -CLI
DEVICE_FILE PORT  SERIAL LDEV CTG  H/M/12 SSID R:Group PRODUCT_ID
sdh          CL2-B  30053   23   2   S/P/ss 0004 5:02-01 OPEN-3
sdi          CL1-A  64015   14   -   -       0004 E:00002 OPEN-3-CM
sdj          -      -        -   -   -       -      -      -
```

Description of the inqraid with -cli option (Linux)**DEVICE_FILE:**

Displays the device file name only.

PORT

Displays the RAID storage system port number.

SERIAL

Displays the production (serial#) number of the storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

LDEV

Displays the LDEV# within the storage system.

CTG

Displays the consistency group ID of TrueCopy Async and Universal Replicator when the LDEV has been defined as a TrueCopy Async or Universal Replicator P-VOL or S-VOL.

H/M/12

Displays the attribute (P = P-VOL, S = S-VOL, s = SMPL) of a remote copy volume (TC/UR/GAD), local copy volume (SI/TI/Copy-on-Write Snapshot), or local copy MU#1,2 volume.

- H: Status of MU#0 for remote copy
- M: Status of MU#0 for local copy
- 1: Status of MU#1 for local copy
- 2: Status of MU#2 for local copy

SSID

Displays the storage subsystem ID of an LDEV in the storage system.



Note: SSID is a parameter used by enterprise storage systems. Although SSID is not used by VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, or VSP G200, G400, G600, G800 and VSP F400, F600, F800 , the set value is displayed.

R:Group

Displays the physical position of an LDEV according to mapping of LDEV in the storage system.

LDEV Mapping	R:	Group
RAID Group	RAID Level 1: RAID1 5: RAID5 6: RAID6	RAID Group number - Sub number
Copy-on-Write Snapshot S-VOL	S	Pool ID number
Unmapped	U	00000
External LUN	E	External group number
Dynamic Provisioning volume	A	Pool ID number

PRODUCT_ID

Displays product-id field in the STD inquiry page.



Note:

For a command device, PORT/SERIAL/LDEV/PRODUCT_ID is the SCSI Inquiry information for the external command device, if the command device is mapped as ELUN(**R:** =E).

Example 6: inqraid with -cliwp and -cliwn options (HP-UX shown)

```
# echo /dev/rdisk/c23t0d0 /dev/rdisk/c23t0d1 | ./inqraid -CLIWP
DEVICE_FILE PWNN          AL PORT  LUN  SERIAL LDEV PRODUCT_ID
c23t0d0     500060e802f01018 - CL2-J -   61456 192   OPEN-3
c23t0d1     500060e802f01018 - CL2-J -   61456 193   OPEN-3

# echo /dev/rdisk/c0t2d3 | ./inqraid -CLIWN
DEVICE_FILE NWWN          AL PORT  LUN  SERIAL LDEV PRODUCT_ID
c0t2d3      5000E000E0005000 - CL1-A -   30015 2054 OPEN3-CVS
```

Description of the inqraid output with -cliwp and -cliwn options (HP-UX shown)**DEVICE_FILE**

Displays the device file name only.

WWN

CLIWP option displays Port_WWN of the host adapter included in the STD inquiry page. CLIWN option displays Node_WWN of host adapter included in STD inquiry page.

AL

Always displays '-'.

PORT

Displays the RAID storage system port number.

LUN

Always displays '-'.

SERIAL

Displays the production (serial#) number of the storage system.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

LDEV

Displays the LDEV# within the storage system.

PRODUCT_ID

Displays product-id field in the STD inquiry page.

Example 7: inqraid with -sort[cm] option (HP-UX shown)

```
#ioscan -fun | grep rdsd | ./inqraid -sort -CM -CLI
HORCM_CMD
#dev_name          dev_name          dev_name
#UnitID 0 (Serial# 30012)
/dev/rdsd/c0t3d0    /dev/rdsd/clt2d1
#UnitID 1 (Serial# 30013)
/dev/rdsd/c2t3d0
```

Example 8: inqraid with -gvinf option

```
D:\HORCM\etc>inqraid $Phys -gvinf -CLI
\\.\PhysicalDrive0:
# Harddisk0      -> [VOL61459_448_DA7C0D91] [OPEN-3    ]
\\.\PhysicalDrive1:
# Harddisk1      -> [VOL61459_449_DA7C0D92] [OPEN-3    ]
\\.\PhysicalDrive2:
```

```
# Harddisk2          -> [VOL61459_450_DA7C0D93] [OPEN-3      ]
                        -> S/N  LDEV  Signature
```

Example 9: inqraid with -svinf[=PTN] option

```
D:\HORCM\etc>pairdisplay -l -fd -g URA
Group PairVol(L/R) Device_File M ,Seq#,LDEV#.P/S,Status, Seq#,
P-LDEV# M
URA   URA_000(L)   Harddisk3   0 61459 451..S-VOL SSUS,-----
448    -
URA   URA_001(L)   Harddisk4   0 61459 452..S-VOL SSUS,-----
449    -
URA   URA_002(L)   Harddisk5   0 61459 453..S-VOL SSUS,-----
450    -

D:\HORCM\etc>pairdisplay -l -fd -g URA | inqraid -svinf=Harddisk
[VOL61459_451_5296A763] -> Harddisk3      [OPEN-3      ]
[VOL61459_452_5296A760] -> Harddisk4      [OPEN-3      ]
[VOL61459_453_5296A761] -> Harddisk5      [OPEN-3      ]
```



Caution: If the S-VOL is created with 'Noread' option (ShadowImage only) and the system is rebooted, the system cannot create a Device object (\Device \HarddiskVolume#) and Volume{guid} for S-VOL, but a Device object (\Device \HarddiskVolume#) and Volume{guid} is created by using -svinf option after splits the S-VOL.

mkconf

The **mkconf** command is a CCI command tool used to make a configuration file from a special file list (raw device file list) provided via STDIN.

Syntax

```
/HORCM/usr/bin/mkconf.sh      (UNIX systems)
\HORCM\Tool\mkconf.exe       (Windows and OpenVMS systems)
mkconf.sh  [ -g[g] <group> [-m <mu#>] [-i <inst#>] [-s <service>]
[-a] ] [-n]
mkconf.exe [ -g[g] <group> [-m <mu#>] [-i <inst#>] [-s <service>]
[-a] [-c <drive#>] ]
```

Options and parameters

No argument

No option displays help/usage and version information.

-g <group>

Specifies the 'dev_group' name for a configuration file. If not specified, 'VG' is used as default.

-gg

Shows a LUN on the host view by finding a host group.

-m <mu#>

Specifies the mirror descriptor for ShadowImage, Copy-on-Write Snapshot, or Thin Image volumes. TrueCopy volume does not specify the mirror descriptor.

-i <inst#>

Specifies the instance number for HORCM.

-s <service>

Specifies the service name (port name) for a configuration file. If not specified, '52323' is used as default.

-a

Specifies an addition of the group to a configuration file.

-c <drive#>

Windows systems only.

Specifies the range of drive numbers that should be examined in order to discover the command devices. If not specified, '\$PhysicalDrive' is used as default.

-c <DKA#-#>

OpenVMS systems only.

Specifies the range of drive numbers that should be examined to discover the command devices. If not specified, '\$1\$DGA0-10000 DKA0-10000 DGA0-10000' is used as default.

[-n]

This `-n` option is specified to use a device file for NVMe. Specifying this option enables the use of both device files for NVMe and SCSI.

If the `HORCC_NVME` environment variable is specified, both device files for NVMe and SCSI can be used without the `-n` option specified.

To specify this option, a command device must be enabled in a namespace on the storage system. For more information about whether a command device can be enabled in a namespace on the storage system, see the *Provisioning Guide* for your storage system.

Example 1: command example for special file list creation

The following are command examples that output the special file list provided via STDIN by using the `mkconf` command.

The special file list must contain the special file name and the hint path.

- **Windows:** `c:\HORCM\etc\inqraid $Phys -CLI`
- **LINUX:** `ls /dev/sd*`
- **unix:** `ls /<device-file-storage-directory-path>/*`

Example 2: mkconf command tool (HP-UX shown)

In this example, the configuration file is created as 'horcm*.conf' in the current directory. The log directory of HORCM is specified as 'log*' in the current directory. You must modify the 'ip_address & service' of an existing configuration file as needed.

```
# cd /tmp/test
# ls /dev/rdisk/* > /etc/devicelist.txt
# cat /etc/devicelist.txt | /HORCM/usr/bin/mkconf.sh -g ORA -i 9
-m 0
starting HORCM inst 9
HORCM inst 9 starts successfully.
HORCM Shutdown inst 9 !!!
A CONFIG file was successfully completed.
starting HORCM inst 9
HORCM inst 9 starts successfully.
DEVICE_FILE      Group PairVol  PORT  TARG  LUN M  SERIAL  LDEV
/dev/rdisk/c23t0d0  ORA  ORA_000  CL2-J   0    0 0   61456   192
/dev/rdisk/c23t0d1  ORA  ORA_001  CL2-J   0    1 0   61456   193
/dev/rdisk/c23t0d2  ORA  ORA_002  CL2-J   0    2 0   61456   194
/dev/rdisk/c23t0d3  ORA  ORA_003  CL2-J   0    3 0   61456   195
/dev/rdisk/c23t0d4  ORA  ORA_004  CL2-J   0    4 0   61456   256
/dev/rdisk/c23t0d5  ORA  ORA_005  CL2-J   0    5 0   61456   257
/dev/rdisk/c23t0d6  ORA  ORA_006  CL2-J   0    6 0   61456   258
/dev/rdisk/c23t0d7  -    -        -    -    - 0   61456   259
HORCM Shutdown inst 9 !!!
Please check '/tmp/test/horcm9.conf','/tmp/test/log9/curlog/
horcm_*.log', and modify 'ip_address & service'.

# ls                                <=Verify configuration and log files.
horcm9.conf  log9
# vi *.conf    <=Verify config file, check ip address & service.

# Created by mkconf.sh on Mon Jan 22 17:59:11 JST 2001

HORCM_MON
#ip_address      service          poll(10ms)      timeout(10ms)
127.0.0.1        52323                1000            3000

HORCM_CMD
#dev_name        dev_name          dev_name
#UnitID 0 (Serial# 61456)
/dev/rdisk/c23t3d0
```

Example 3: mkconf command tool (HP-UX shown)

```

HORCM_DEV
#dev_group      dev_name      port#      TargetID      LU#      MU#
# /dev/rdisk/c23t0d0 SER = 61456 LDEV = 192 [FIBRE FCTBL = 4] ORA
ORA_000 CL2-J 0      0      0
# /dev/rdisk/c23t0d1 SER = 61456 LDEV = 193 [FIBRE FCTBL = 4] ORA
ORA_001 CL2-J 0      1      0
# /dev/rdisk/c23t0d2 SER = 61456 LDEV = 194 [FIBRE FCTBL = 4] ORA
ORA_002 CL2-J 0      2      0
# /dev/rdisk/c23t0d3 SER = 61456 LDEV = 195 [FIBRE FCTBL = 4] ORA
ORA_003 CL2-J 0      3      0
# /dev/rdisk/c23t0d4 SER = 61456 LDEV = 256 [FIBRE FCTBL = 4] ORA
ORA_004 CL2-J 0      4      0
# /dev/rdisk/c23t0d5 SER = 61456 LDEV = 257 [FIBRE FCTBL = 4] ORA
ORA_005 CL2-J 0      5      0
# /dev/rdisk/c23t0d6 SER = 61456 LDEV = 258 [FIBRE FCTBL = 4] ORA
ORA_006 CL2-J 0      6      0
# ERROR [CMDDEV] /dev/rdisk/c23t0d7 SER = 61456 LDEV = 259
[OPEN-3-CM]                                     <=See Notes.

HORCM_INST
#dev_group      ip_address      service
ORA             127.0.0.1      52323      <=Check and update as needed.

```

**Note:**

- A unitID is added to the Serial# order. If two or more command devices exist in the storage system, select the device file that is shared among the storage system ports on a priority basis, and treat as an alternate command device.
- If the STDIN device includes the command device, the target device is commented out as shown below:

```
# ERROR [CMDDEV] /dev/rdisk/c23t0d7 SER = 61456 LDEV = 259 [ OPEN-3-
CM ]
```

- If the STDIN device is shared among multiple device files and already displayed as a target device, the target device is commented out as shown below:

```
# ERROR [LDEV LINK] /dev/rdisk/c24t0d3 SER = 61456 LDEV = 195
[FIBREFCTBL = 4]
```

- If the STDIN device does not have appropriate mirror description (MU#), the target device is commented out as shown below:

```
# ERROR [INVALID MUN (2 < 1)] /dev/rdisk/c24t0d3 SER = 61456 LDEV = 195
[ OPEN-3 ]
```

- If the STDIN device is mixed among the storage systems of differential mirror control, the target device is commented out as shown below:

```
# ERROR [MIXING RAID TYPE] /dev/rdisk/c24t0d3 SER = 61456 LDEV = 195
[ OPEN-3]
```

rmawk

The **rmawk** command is a scriptable command for associating with the pair operation commands and raidcom commands.

This command provides basically the following three functions:

- Outputs by filtering STDIN with the specified conditions.

```
Command | rmawk @1-eq:PAIR -a @2-ne:0
```

- Performs the command with the variable parameter specified by interpreting STDIN with the specified conditions.

```
Command | rmawk @3-eq:TAR exe="Command line @1"
```

- Tests by interpreting the output command with the specified conditions or waits until the conditions become TRUE, by performing the specified command.

```
rmawk exe="Command line" @2-eq:PAIR -o @2-eq:PSUS
rmawk exe="Command line" @2-eq:COPY timeout=300
```

Syntax

```
/HORCM/usr/bin/rmawk [-h | exe="..." | sys="..." | timeout=value | interval=value | -BL
| -AT | -EC[VAL] | @variable | operators ]
```

```
\HORCM\etc\rmawk [-h | exe="..." | sys="..." | timeout=value | interval=value | -BL | -
AT | -EC[VAL] | @variable | operators ]
```

Options and parameters

No argument or -h

No option displays help/usage and version information.

exe="command_line"

Specifies the command line to be performed when the result of the specified formula and the conditions are TRUE. The testing/waiting specifies the command to be its target.

"@variable" with the field variables of STDIN can be included in the command line.

Example:

```
exe="raidcfg -a gry -o hgrp -pport @1"
exe="raidcfg -a reg -o hgrp -pport @1-@2 -pname @3"
exe="type map.txt | rmawk @@L-eq:@L exe=\"raidcfg -a map -o
snap -pname @1 -pldev @5 @@1\""
```

exe="print "

Specifies printing (with Line Feed) when the result of the specified formula and the conditions are TRUE.

"@variable" with the field variables of STDIN can be included in the command line. When "@variable" is specified as "@C#/C#*", the variable is interpreted as character string. When "@variable" is specified as "@C#?/C#+/L/C/R", the variable is interpreted as integer.

The printing range is from "print" to the end including spaces.

Example:

```
exe="print PORT=@1 LDEV=@3"
exe="print PORT=@1 LDEV=@3 Total CAP = @4+"
```

exe="printn.... " or exe="prints.... "

Specifies printing (without Line Feed) when the result of the specified formula and the conditions are TRUE.

"@variable" with the field variables of STDIN can be included in the command line. When "@variable" is specified as "@C#/C#*", the variable is interpreted as character string. When "@variable" is specified as "@C#?/C#+/L/C/R", the variable is interpreted as integer.

The printing range is from "print" to the end including spaces.

Example:

```
exe="printn PORT=@1 LDEV=@3"
```



Note: "prints..." is used to print to the strings buffer of '@0*'.

exe=exit

Specifies to exit this command when the result of the specified formula and the conditions are TRUE. The return value of exit is "@R" variable.

sys="command_line"

Specifies the command line to be performed in the transparent mode when the result of the specified formula and the conditions are TRUE.

"@variable" and "@expression" in the command line are not interpreted as variables but are filtered out as variables to path them to the specified command.

Example for UNIX:

```
sys='rmawk exe="pairedisplay -g G1 -CLI -l" @6-eq:PAIR'
```

Example for Windows:

```
sys="rmawk exe=\"pairedisplay -g G1 -CLI -l\" @6-eq:PAIR"
```

Example for openVMS:

```
sys="rmawk exe=""pairedisplay -g G1 -CLI -l"" @6-eq:PAIR"
```

Exception: When "exe=..." is specified in the command line and there is "@variable" as an argument in this command, it is interpreted as "@variable" that is including the field variables of STDIN.

Example for UNIX:

```
sys='rmawk exe="pairdisplay -d @3 @5 -CLI -l" @6-eq:PAIR'
```

Example for Windows:

```
sys="rmawk exe=\"pairdisplay -d @3 @5 -CLI -l\" @6-eq:PAIR"
```

Example for OpenVMS:

```
sys="rmawk exe=""pairdisplay -d @3 @5 -CLI -l"" @6-eq:PAIR"
```

timeout=value

Specifies the timeout time in seconds (1 to 2147483647) for waiting by interpreting this output command with the specified conditions when the specified command is performed.

interval=value

Specifies the interval time in seconds (1 to 2147483647) for waiting by interpreting this output command with the specified conditions when the specified command is performed.

-BL

Specifies to include blank lines from STDIN. This is used to detect the blank lines when there are in the STDIN.

-AT

- When "-AT" is specified (the default setting does not specify this option)
 Specifies to testing or waiting until the conditions become TRUE in all the command lines, by interpreting the output of the specified command with the specified conditions.
 Testing ends as an ERROR (1) if there is even one FALSE in the specified condition. And waiting waits until all the lines become TRUE by retrying the specified command until becoming timeout if there is even one FALSE in the lines of the specified conditions. If it becomes timeout, the command ends as an ERROR (1).
- When "-AT" is not specified
 It tests or waits until at least one line becomes TRUE, by interpreting the output of the specified command with the specified conditions.
 If there is one TRUE at least in the specified conditions, the testing ends as NORMAL (0). If there is one TRUE at least in the specified conditions, the waiting ends as NORMAL (0). If all the lines are FALSE, it waits until at least one line become TRUE by retrying the specified command until becoming timeout. If it becomes timeout, the command ends as an ERROR (1).

-EC[VAL]

Performs only once with ending of lines or "exe=exit" as Line#0. If you need to perform something special at end of lines, specify this option that is called as Line#0. This is enabled in the run mode.

If VAL(integer) is specified, @C#? variable for all are initialized to the specified VAL before starting. If VAL is invalid such as -ECM, then @C#? variable for all are initialized to the MAX value with 64-bit integer.

@variable

Specifies the following variables. The variables can be specified in both the left and right side of the formula. The variables can also be specified as arguments of the command to be performed or of printing.

- @C#
Field strings corresponding to the column number (Column#) from STDIN. If "C#" is specified in hexadecimal ("0x..."), then its field is treated as hexadecimal strings.
- @C#*
Strings from the field to the end corresponding to the column number from STDIN.
- @C#?
64 bit variables corresponding to the column number from STDIN (initial value 0).
- @C#+
64 bit variables adding field variables corresponding to the column number from STDIN (initial value 0).
- @0
A variable that indicates one line (This is used when printing or searching within the line.).
- @L
A variable that indicates the line number from STDIN.
- @C
A variable that indicates the number of column in each line of STDIN.
- @R
A return value of performing command.
- @0*
A string variable that can be memorized one line.
- null
A special variable that tests if the character string is null or not.
- @C#?t
A variable that can be specified within exe="..." to print the C#? variable as TOD (time of day). If the C#? variable is zero, the current time is printed.

- **@C#?x**
A variable that can be specified within `exe="..."` to print the `C#?` variable as hexadecimal strings by adding `0x`.
- **@@**
A variable that can be specified within `exe="..."` to print the `@` character.

operators (-operator:)

Specifies the following comparison operation. Variables in the field are operated in 64-bit integers.

The numerical values in the field strings are "1234..." or "0x1234..." format without octal code.

- **-eq:**
Treats the target as character strings or integers by comparing with `"=="`. When `"@C#/C#"` is specified in the left side of an equation, this is treated as character strings. When `"@C#?/C#+/L/C/R"` is specified in it, this is treated as integers.
- **-ne:**
Treats the target as character strings or integers by comparing with `"!="`. When `"@C#/C#"` is specified in the left side of an equation, this is treated as character strings. When `"@C#?/C#+/L/C/R"` is specified in it, this is treated as integers.
- **-gt:**
Treats the target as integers by comparing with `">"`.
- **-ge:**
Treats the target as integers by comparing with `">="`.
- **-lt:**
Treats the target as integers by comparing with `"<"`.
- **-le:**
Treats the target as integers by comparing with `"<="`.

operators (=operator:)

Specifies the following arithmetical operation. Variables in the field are operated in 64-bit integers.

The numerical values in the field strings are "1234..." or "0x1234..." format without octal code.

- **=ad:**
Treats the target as integers by operating with `"+" (Addition)"`.
- **=sb:**
Treats the target as integers by operating with `"-" (Subtraction)"`.
- **=ml:**
Treats the target as integers by operating with `"*" (Multiplication)"`.

- **=dv:**
Treats the target as integers by operating with "/" (Division)".
- **=st:**
Treats the target as integers by substituting with "=" (Setting)".

operators (-operator)

Specifies the following logical operation for the result of the comparison operation.

- **-o**
Judges the result of the comparison operation as the "Logical OR".
- **-a**
Judges the result of the comparison operation as the "Logical AND".
- **-n**
Inverts the result of the comparison operation. (TRUE becomes FALSE, FALSE becomes TRUE.)

Returned values

The **rmawk** command sets the following returned values during exit allowing you to check the execution results.

- **Normal termination:**
 - **0:** The command ends normally with the specified condition.
 - **1:** Testing: The specified condition is FALSE ; Waiting: The specified condition is Timeout.
- **Abnormal termination:**
 - **125:** The command ends with a syntax error.
 - **126:** The command ends with a system error.

Examples of comparison expression

- **@20-eq:PAIR:** Compares if the character string in Column #20 from STDIN matches "PAIR".
- **@20-eq:PSU*:** Compares if "PSU" is included in the character string in Column #20 from STDIN.
- **@0-eq:PSU*:** Compares if "PSU" is included in one line from STDIN. This is equivalent to "grep PSU".
- **@20-eq:@21:** Compares if the character strings in Column #20 and Column #21 match.
- **@20-ge:50:** Compares if it is "value >= 50" of Column #20 from STDIN.
- **@L-ge:20:** Compares if it is the current "number of lines >= 20" from STDIN.
- **@C-ge:20:** Compares if it is "number of columns >= 20" of the current lines from the STDIN.
- **@R-gt:0:** Compares if it is "return value > 0" of the specified command.

Examples of arithmetic expression

- `@8?=ad:@8`: Performs "`@8? = @8? + @8`". This is equivalent to "`@8+`".
- `@8=ad:@5`: Performs "`@8? = @8 + @5`".
- `@8=ad:@5?`: Performs "`@8? = @8 + @5?`".
- `@8=ad:30`: Performs "`@8? = @8 + 30?`".
- `@8=st:30`: Performs "`@8? = 30`".
- `@5=st:@5`: Performs "`@5? = @5`" for converting to the integer.
- `:`
- `:`
- `:`

Example:

```
Command | rmawk @8?=ad:@8 exe="print Total = @8?"
Command | rmawk exe="print Total = @8+"

```

Example of the command option format

Performs the operation by interpreting the specified option format and classifying them in the following three function formats.

- Performs the operation as a filter if there is no "exe=..." in the specified option.

```
Command | rmawk [@expression1] Logical operator [@expression2] ...
```

Example:

```
Command | rmawk @1-eq:PAIR -a @2-ne:0
```

- Performs the operation as a testing if there is one "exe=..." but no "@variable" in the specified option.

```
rmawk exe="Command line" [@expression1] Logical operator [@expression2] ...
```

Example:

```
rmawk exe="Command line" @2-eq:PAIR -o @2-eq:PSUS
```

And performs as an waiting command if "timeout=value" is specified as an option.

```
rmawk exe="Command line" [@expression1] Logical operator [@expression2] timeout=6
```

Example:

```
rmawk exe="Command line" @2-eq:COPY timeout=300
```

- Other than the above, it performs the multiple "exe="Command line"" commands that include contexts written in multiple formulas and the variable parameters by interpreting the STDIN.

```
Command | rmawk [@expression1] Logical operator [@expression2] exe="Command line"
```

Example:

```
Command | rmawk @3-eq:TAR exe="Command line @1"
Command | rmawk @3-eq:TAR exe="Command line(true) @1" -n exe="Command line(false) @1"
```

Because the command runs with the result of TRUE, multiple command can be performed if you keep writing lines ("exe="command line2"", "exe="command line3"", and so on).

```
Command | rmawk [@expression1] Logical operator [@expression2] exe="Command line1"
exe="Command line2" exe="Command line3"
```

Example:

```
Command | rmawk @3-eq:TAR exe="Command line @1" exe="Command line2 @1"
Command | rmawk @3-eq:TAR exe="Command line(true) @1" exe="Command line2(true) @1"
-n exe="Command line(false) @1"
```

Chapter 5: Configuration setting commands

The CCI **raidcom** command enables you to perform configuration and provisioning operations on the Hitachi RAID storage systems.



Note: If a failure occurs in a processor on an MP blade or controller, some commands cannot be executed on any MP blades or controllers (even if no failure occurs on those MP blades or controllers). For details, see [Command execution with MP blade or controller failure \(on page 269\)](#).

raidcom

Specifies a configuration change.

Syntax

Displaying help

```
raidcom { [-h]
```

Log-in and log-out

```
raidcom {-login [<user_name> <password>] | -logout} [-I[H][M] <instance#>] [-I[TC]  
[SI] <instance#>]
```

Configuration change: line-by-line mode

```
raidcom <action> <object> [<param> <value>...] [-fx] [-checkmode <check mode>] [-store  
<filename>] [-nomsg] [-login [<user_name> <password>] [-resource <resource_grp_id>...]  
| -logout] [-s <seq#> | -u <unit#>] [-I[H][M] <instance#>] [-I[TC][SI] <instance#>]
```

Configuration change: transaction mode

```
raidcom -zt <filename> [-load <work_filename>]  
[-checkmode <check mode>]
```

Options and parameters

[-h]

Displays help of raidcom.

[-login [<user_name> <password>]]

Specifies user authentication for the storage system. Specifies user name and password.

If no user authentication is done yet and the `-login` option is omitted, the input of the user name and the password is required. And, in case of omitting the user name and the password by specifying `-login` option, the input of the user name and the password is required.

The maximum number of the user who can login at same time is 512.

If you specify the user name and the password with the `-login` option, user authentication is performed for all storage systems that require user authentication described in the configuration definition file.

If the configuration definition file contains command devices for multiple storage systems and the user name and the password are different for each storage system, specify the `-login` option without the user name and the password, and then run the command by manually entering the user name and the password for each storage system. If the user name and the password is not omitted, the specified user name and password are used to log in to all storage systems in order. For details about user authentication, see the *Command Control Interface User and Reference Guide*.

[-logout]

Logs out from the storage system. For details about user authentication, see the *Command Control Interface User and Reference Guide*.

[-s <seq#>]

Specifies the serial number.

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

[-u <Unit#>]

Specifies the unit ID of command device as HORCM_CMD of configuration file.

-I[H][M] [instance#] or -I[TC][SI] [instance#]

Specifies a target pair type and instance number on which the command operates.

- -I[H][M] or -I[TC][SI]

Specifies a target pair type on which the command operates.

- [H] or [TC]

a TrueCopy, Universal Replicator, or global-active device pair

- [M] or [SI]

a ShadowImage, Volume Migration, or Thin Image pair

If not specified, the target pair type is determined by whether the HORCC_MRCF environment variable is set in the command execution environment.

- If the HORCC_MRCF environment variable is not set:

a TrueCopy, Universal Replicator, or global-active device pair

- If the HORCC_MRCF environment variable is set:

a ShadowImage, Volume Migration, or Thin Image pair

- [instance#]

Specifies a number of an instance that is used when running the command.

If not specified, the instance number is determined by whether the HORCMINST environment variable is set in the command execution environment.

- If the HORCMINST environment variable is not set:

an instance with no instance number is used.

- If the HORCMINST environment variable is set:

an instance whose instance number is set as the HORCMINST environment variable is used.

Example

Specify TrueCopy, Universal Replicator, or global-active device as a target pair type on which the command operates, and an instance number: -IH0

<action> <object> [<param> <value>...]

Specifies each command name that is expressed in **raidcom add copy_grp** and later section and the parameters and values that are required at each operation.

[-fx]

Displays the LDEV number in hexadecimal notation.

[-checkmode <check mode>]

Specifies when executing Precheck function (execute checking commands only). It specifies the following value.

Precheck: When this option is specified, the actual process is not executed for the storage system even if the command is executed.

This option is available to specify the setting of \$HORCC_NO_EXEC environment variable and \$HORCC_CTX_CHK environment variable.

[-store <file name>]

Specifies the file name of the configuration file to be created for implementation check.

[-nomsg]

Prevents displaying messages. This option is required to be defined on the beginning of the command parameter.

[-zt <file name>]

Specifies the script file.

[-load <file name>]

Specifies the file name of the file (Configuration file) that is created for implementation check.

[-resource <resource_grp_id>...]

Limits the operations coverage to the specified resource groups. When this option is specified in reference commands, objects that are in the specified resource groups are displayed. When the resource group specified by this option does not contain the resource specified for executing the command, the command is not executed because of the EX_EGPERM error.

Returned values

Unless otherwise stated, the **raidcom** commands set the following returned values:

- **0**: Normal termination.
- **other than 0**: Abnormal termination

For details, see the section describing command error messages in *Command Control Interface User and Reference Guide*.

Examples

Performing user authentication (login) by the user ID: USER01 and the password: PASS01.

```
# raidcom -login USER01 PASS01
```

Performing log-out.

```
# raidcom -logout
```

Performing user authentication (login) to instance 99 by the user ID: USER01 and the password: PASS01.

```
# raidcom -login USER01 PASS01 -I99
```

Performing syntax check and the context check of the script file (the actual processing is not executed).

```
# raidcom -zt <script file> -checkmode precheck
```

Performing syntax check, the context check, and the implementation check of the script file (the actual processing is not executed).

```
# raidcom -zt <script file> -load <work file> -checkmode precheck
```



Note: If the user who does not have the write permission for the HORCM directory and the directories under the directory performs a raidcom command, the following WARNING might be output. If the WARNING is output, confirm whether you have the write permission for the directory, and then retry the command. If only the WARNING occurs, the returned value will be 0 (normal termination).

```
raidcom: [WARNING] Couldn't create User file for Serial#(<XXXXXX>).
```

Serial#: Displays the product serial number.

Method for specifying LDEV number

The methods for specifying LDEV numbers are shown in the following table.

Specification method	Example
Decimal number	-ldev_id 300
Hexadecimal number	-ldev_id 0x12C -ldev_id 01:2C

Methods for specifying multiple LDEVs

Methods for specifying multiple LDEVs by LDEV IDs

The methods for specifying multiple LDEV IDs in a single command are shown below. Some methods do not apply to some commands.

```
-ldev_id 300-305
-ldev_id 0x12C-0x131
-ldev_id 01:2C-01:31
-ldev_id 300 -cnt 6
-ldev_id 300 301 302 303 304 305
```

Method for specifying LDEVs by device group

When you configure multiple LDEVs for pools or journals (for example, changing the resource group), you must issue the command to each LDEV. However, if you define LDEVs that configure pools or journals, you can issue the command to all LDEVs defined as a device group by specifying each device group at a time.

If you issue the command to LDEVs by specifying a device group, the command is executed in each LDEV registered in the device group. If an error occurs in an LDEV while the command is being executed, the execution of command stops at the LDEV where the error occurred. The command is not executed in the remaining LDEVs. If an error occurs, solve the error. Then restore the LDEV and issue the command to all remaining LDEVs.

- Creating a pool by specifying a device group:

```
# raidcom add device_grp -device_grp_name dg_pool1 data1
-ldev_id 512 513 514 515
```

```
# raidcom add dp_pool -pool_id 2 -grp_opt ldev -device_grp_name dg_pool1
```

- Creating a journal by specifying a device group:

```
# raidcom add device_grp -device_grp_name dg_jnl1 data1
-ldev_id 512 513 514 515
```

```
# raidcom add journal -journal_id 2 -grp_opt ldev
-device_grp_name dg_jnl1
```

Operations where multiple LDEVs can be specified

You can use "-ldev_id <ldev#>" to specify multiple LDEVs at the same time in the following operations only:

(a) Displaying LDEV information

```
# raidcom get ldev -ldev_id 100-103
# raidcom get ldev -ldev_id 100 -cnt 4
```

**Note:**

Specifying multiple LDEV as follows cannot be performed: #raidcom get ldev -ldev_id 100 101 103

(b) Creating a journal

```
# raidcom add journal -journal_id 1 -ldev_id 265 266
# raidcom add journal -journal_id 1 -ldev_id 265-266
# raidcom add journal -journal_id 1 -ldev_id 265 -cnt 2
```

(c) Creating a pool

Creating a pool for Copy-on-Write Snapshot:

```
# raidcom add snap_pool -pool 1 -ldev_id 365 366 367
# raidcom add snap_pool -pool 1 -ldev_id 365-367
# raidcom add snap_pool -pool 1 -ldev_id 365 -cnt 3
```

Creating a pool for Dynamic Provisioning or Dynamic Provisioning for Mainframe:

```
# raidcom add dp_pool -pool 1 -ldev_id 465 466 467
```

```
# raidcom add dp_pool -pool 1 -ldev_id 465-470
```

```
# raidcom add dp_pool -pool 1 -ldev_id 465 -cnt 5
```

(d) Creating a device group

```
# raidcom add device_grp -device_grp_name DevG2 dev101 -ldev_id 101 105 201
```

```
# raidcom add device_grp -device_grp_name DevG2 dev101 -ldev_id 101-105
```

```
# raidcom add device_grp -device_grp_name DevG2 dev101 -ldev_id 101 -cnt 5
```

(e) Deleting a device group

```
# raidcom delete device_grp -device_grp_name DevG3 -ldev_id 200 201 204
```

```
# raidcom delete device_grp -device_grp_name DevG3 -ldev_id 200-204
```

```
# raidcom delete device_grp -device_grp_name DevG3 -ldev_id 200 -cnt 5
```

Specifying and displaying storage system serial numbers

Note the following when you specify the serial number of a storage system for running commands, or when you check the serial number in output result.

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.
- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).



Note: When you specify the serial number for the storage system in the GUI (for example, HDvM - SN, Hitachi Storage Advisor, or Hitachi Command Suite), use only the five-digit serial number. Do not add "3" or "5" in front of the serial number.

Maximum number of acceptable asynchronous commands

The following list shows the maximum numbers of asynchronous commands that the storage system can accept. When you continuously run a large number of asynchronous commands, make sure not to exceed these maximum numbers. For details, see the topic regarding asynchronous commands in the *Command Control Interface User and Reference Guide*. If the following number is exceeded, split the commands so that the asynchronous commands exceeding the maximum number of acceptable asynchronous commands are not run on the storage system.

- VSP 5000 series: 2,560
- VSP G1x00 and VSP F1500: 1,280
- VSP One B20, VSP E series: 2,560
- VSP G130, G/F350, G/F370, G/F700, G/F900: 2,560
- VSP G200, G400, G600, G800 and VSP F400, F600, F800: 2,560

Command execution with MP blade or controller failure

If a failure occurs in a processor on an MP blade or controller, some commands cannot be executed on any MP blades or controllers, even if no failure occurs on those MP blades or controllers. The following table lists the commands.

Command group	Commands
raidcom <action> clpr	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> user_system_opt	<ul style="list-style-type: none"> ▪ modify
raidcom <action> copy_grp	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> device_grp	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> drive	<ul style="list-style-type: none"> ▪ modify
raidcom <action> external_grp	<ul style="list-style-type: none"> ▪ add ▪ check_ext_storage ▪ delete ▪ disconnect ▪ modify
raidcom <action> external_storage	<ul style="list-style-type: none"> ▪ discover
raidcom <action> external_iscsi_name	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> external_chap_user	<ul style="list-style-type: none"> ▪ modify
raidcom <action> initiator_chap_user	<ul style="list-style-type: none"> ▪ modify
raidcom <action> journal	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> ldev	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ extend¹

Command group	Commands
	<ul style="list-style-type: none"> ▪ initialize ▪ modify
raidcom <action> license	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> quorum	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify ▪ replace
raidcom <action> local_replica_opt	<ul style="list-style-type: none"> ▪ modify
raidcom <action> remote_replica_opt	<ul style="list-style-type: none"> ▪ modify
raidcom <action> lun	<ul style="list-style-type: none"> ▪ add²
raidcom <action> path	<ul style="list-style-type: none"> ▪ add ▪ check_ext_storage ▪ delete ▪ disconnect ▪ modify
raidcom <action> pool	<ul style="list-style-type: none"> ▪ delete ▪ initialize ▪ modify ▪ monitor ▪ reallocate ▪ rename
raidcom <action> system	<ul style="list-style-type: none"> ▪ modify
raidcom <action> port	<ul style="list-style-type: none"> ▪ modify³
raidcom <action> parity_grp	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ initialize ▪ modify

Command group	Commands
raidcom <action> rcu	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> rcu_iscsi_port	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> rcu_path	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> ssid	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> resource	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ map ▪ modify ▪ unmap
raidcom <action> snap_pool	<ul style="list-style-type: none"> ▪ add
raidcom <action> dp_pool	<ul style="list-style-type: none"> ▪ add
raidcom <action> server	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> qos_grp	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> nvm_subsystem	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
raidcom <action> nvm_subsystem_port	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> host_nqn	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify

Command group	Commands
raidcom <action> namespace_path	<ul style="list-style-type: none"> ▪ add ▪ delete
raidcom <action> namespace	<ul style="list-style-type: none"> ▪ add ▪ delete ▪ modify
<ol style="list-style-type: none"> 1. For the raidcom extend ldev command, this restriction applies only when you run the command with the following options specified: <code>raidcom extend ldev - request_id auto</code> 2. For the raidcom add lun command, this restriction applies only when you run the command with the following options specified: <code>raidcom add lun -request_id auto</code> 3. For the raidcom modify port command, this restriction applies only when you run the command with the following options specified: <code>raidcom modify port {-t10pi {enable disable} -iscsi_virtual_port_mode {enable disable} -port_mode {scsi nvme} -request_id auto -delete_login_host_nqn}</code> 	

Resource group operation

When you have multiple resource group authorities, use the `-resource` option to see the resource group information.

Examples

```
# raidcom get resource
```

```
RS_GROUP      RGID   stat    Lock_owner  Lock_host   Serial#
meta_resource    0  Unlocked  -           -           64568
RSG_CLI1         1  Unlocked  -           -           64568
RSG_CLI2         2  Unlocked  -           -           64568
```

```
# raidcom get port
```

```
PORT  TYPE  ATTR SPD LPID FAB CONN SSW SL Serial# WWN      PHY_PORT
CL1-A FIBRE ELUN AUT  EF N  FCAL N  0  64568 50060e8006fc3800 -
CL1-B FIBRE TAR  AUT  EF N  FCAL N  0  64568 50060e8006fc3801 -
```

```
CL1-C FICON TAR - - - - - 0 64568 50060e8006fc3802 -
CL1-D FICON TAR - - - - - 0 64568 50060e8006fc3803 -
```

```
# raidcom get port -resource 1
```

```
PORT  TYPE  ATTR SPD LPID FAB  CONN SSW SL Serial# WWN      PHY_PORT
CL1-B FIBRE TAR  AUT   EF N   FCAL N   0 64568 50060e8006fc3801 -
CL1-C FICON TAR  -   -   -   -   -   0 64568 50060e8006fc3802 -
```

```
# raidcom get port -resource 2
```

```
PORT  TYPE  ATTR SPD LPID FAB  CONN SSW SL Serial# WWN      PHY_PORT
CL1-A FIBRE ELUN AUT   EF N   FCAL N   0 64568 50060e8006fc3800 -
CL1-D FICON TAR  -   -   -   -   -   0 64568 50060e8006fc3803 -
```

Resource lock operation

When multiple users perform operations for a single resource, lock the resource group to which the target resource is allocated before performing operations. This applies to both the setting and viewing operations.

The following examples show locking resource groups, executing commands, and then unlocking resource groups.

Examples

- Performing a user authentication by User ID: USER01, Password: PASS01.

```
# raidcom -login USER01 PASS01
```

- Locking a resource group: rsg001.

```
# raidcom lock resource -resource_name rsg001
```

- Creating LDEV#100 and #101.

```
# raidcom add ldev -parity_grp_id 5-2 -ldev_id 100 -capacity 10g
```

```
# raidcom add ldev -parity_grp_id 5-2 -ldev_id 101 -capacity 10g
```

- Unlocking resource group: rsg001

```
# raidcom unlock resource -resource_name rsg001
```

Resource locking and CCI commands

If you execute a CCI command when the specified resource is locked, the specified resource groups cannot be used by other users. Commands can be executed when the specified resources are not locked. However, if another user locks the resource, CCI commands will result in error. The following tables show the relationship between CCI commands and resources that need to be locked, except for the required options for the commands. In the following tables:

- Res. group: resource group
- Lock/auth: Resource locking and resource authority are required.
- Auth: Only resource authority check is specified.
- (VSP only) Lock by user: If another user locks the applicable resource, the command execution fails. To execute the command, the applicable resource authority is not required.

Relationship between commands and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/VDEV	LDEV	Command	Option
-	-	-	-	-	-	Lock/auth	raidcom add copy_grp	-
-	-	-	-	-	-	Lock/auth	raidcom delete copy_grp	-
-	-	-	-	-	-	Auth	raidcom get copy_grp	-
-	-	-	-	-	-	Lock/auth	raidcom add device_grp	-
-	-	-	-	-	-	Lock/auth	raidcom delete device_grp	-
-	-	-	-	-	-	Auth	raidcom get device_grp	-
-	Lock/auth	-	-	-	-	Lock/auth	raidcom add external_grp	-
-	-	-	-	-	Lock/auth	-	raidcom check_ext_storage external_grp	-
-	-	-	-	-	Lock/auth	Auth	raidcom check_ext_storage external_grp	-ldev_id <ldev#>

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	Lock/auth	-	raidcom delete external_grp	-
-	-	-	-	-	Lock/auth	-	raidcom disconnect external_grp	-
-	-	-	-	-	Lock/auth	Auth	raidcom disconnect external_grp	-ldev_id <ldev#>
-	-	-	-	-	Auth	-	raidcom get external_grp	-
-	-	-	-	-	Lock/auth	-	raidcom modify external_grp	-
-	Auth	-	-	-	-	-	raidcom discover external_storage	-
-	-	Lock/ auth	-	-	-	-	raidcom add host_grp	-
-	-	Lock/ auth	-	-	-	-	raidcom delete host_grp	-
-	Auth	Auth	-	-	-	-	raidcom get host_grp	-
-	-	Lock/ auth	-	-	-	-	raidcom modify host_grp	-
-	-	Lock/ auth	-	-	-	-	raidcom add hba_wwn	-
-	-	Lock/ auth	-	-	-	-	raidcom delete hba_wwn	-
-	Auth	Auth	-	-	-	-	raidcom get hba_wwn	-
-	-	-	-	-	-	Lock/ auth	raidcom add journal	-
-	-	-	-	-	-	Lock/ auth	raidcom delete journal	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	Auth	raidcom get journal	-
-	-	-	-	-	-	Lock/ auth	raidcom modify journal	-
-	-	-	-	Lock/ auth	-	Lock/ auth	raidcom add ldev	-parity_ grp_id <gno- sgno>
-	-	-	-	-	Lock/auth	Lock/ auth	raidcom add ldev	-external_ grp_id <gno- sgno>
-	-	-	-	-	-	Lock/ auth	raidcom add ldev	-
-	-	-	-	-	-	Lock/ auth	raidcom delete ldev	-
-	-	-	-	-	-	Lock/ auth	raidcom extend ldev	-
-	-	-	-	-	-	Auth	raidcom get ldev	-
-	-	-	-	-	-	Lock/ auth	raidcom initialize ldev	-
-	-	-	-	-	-	Lock/ auth	raidcom modify ldev	-
-	-	Lock/ auth	-	-	-	Lock/ auth	raidcom add lun	-
-	-	Lock/ auth	-	-	-	Lock/ auth	raidcom delete lun	-
-	Auth	-	-	-	-	-	raidcom discover lun	-
-	-	Auth	-	-	-	-	raidcom get lun	-
-	-	Lock/ auth	-	-	-	Lock/ auth	raidcom modify lun	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	Lock/auth	-	-	-	-	-	raidcom add path	-
-	Lock/auth	-	-	-	-	-	raidcom check_ext_storage path	-
-	Lock/auth	-	-	-	-	-	raidcom delete path	-
-	Lock/auth	-	-	-	-	-	raidcom disconnect path	-
-	-	-	-	-	-	Auth	raidcom get path	-
-	-	-	-	-	-	Lock/auth	raidcom delete pool	-
-	-	-	-	-	-	Auth	raidcom get pool	-
-	-	-	-	-	-	Lock/auth	raidcom modify pool	-
-	-	-	-	-	-	Lock/auth	raidcom rename pool	-
-	Auth	-	-	-	-	-	raidcom get port	-
-	Lock/auth	-	-	-	-	-	raidcom modify port	-
-	-	-	-	Auth	-	-	raidcom get parity_grp	-
-	Lock/auth	-	-	-	-	-	raidcom add rcu	-
-	Lock/auth	-	-	-	-	-	raidcom delete rcu	-
-	Auth	-	-	-	-	-	raidcom get rcu	-
-	Lock/auth	-	-	-	-	-	raidcom modify rcu	-
-	Lock/auth	-	-	-	-	-	raidcom add rcu_path	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	Lock/ auth	-	-	-	-	-	raidcom delete rcu_path	-
Auth	-	-	-	-	-	-	raidcom get resource	-
Auth	-	-	-	-	-	-	raidcom lock resource	-
Auth	-	-	-	-	-	-	raidcom unlock resource	-
-	-	-	-	-	-	Lock/ auth	raidcom add dp_pool	-
-	-	-	-	-	-	Auth	raidcom get dp_pool	-
-	-	Lock/ auth	-	-	-	-	raidcom set hba_wwn	-
-	-	Lock/ auth	-	-	-	-	raidcom reset hba_wwn	-
-	-	-	-	-	-	-	raidcom get command_status	-
-	-	-	-	-	-	-	raidcom reset command_status	-
-	-	-	-	-	-	-	raidcom add resource	-
Lock/ auth	-	-	-	-	-	Lock/ auth	raidcom add resource	-ldev_id <ldev#>
Lock/ auth	Lock/ auth	-	-	-	-	-	raidcom add resource	-port <port#>
Lock/ auth	-	Lock/ auth	-	-	-	-	raidcom add resource	-port <port#> <host group name>

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/VDEV	LDEV	Command	Option
Lock/auth	-	-	-	Lock/auth	-	-	raidcom add resource	- parity_grp_id <gno-sgno>
Lock/auth	-	-	-	-	Lock/auth	-	raidcom add resource	-external_grp_id <gno-sgno>
Auth	-	-	-	-	-	-	raidcom delete resource	-
Lock/auth	-	-	-	-	-	Lock/auth	raidcom delete resource	-ldev_id <ldev#>
Lock/auth	Lock/auth	-	-	-	-	-	raidcom delete resource	-port <port#>
Lock/auth	-	Lock/auth	-	-	-	-	raidcom delete resource	-port <port#> <host group name>
Lock/auth	-	-	-	Lock/auth	-	-	raidcom delete resource	- parity_grp_id <gnosgno>
Lock/auth	-	-	-	-	Lock/auth	-	raidcom delete resource	-external_grp_id <gno-sgno>
Lock/auth	-	-	-	-	-	-	raidcom modify resource	-
-	Lock/auth	-	-	-	-	Lock/auth	raidcom map resource	-
-	Lock/auth	-	-	-	-	-	raidcom unmap resource	-
-	-	-	-	-	-	-	raidcom get error_message	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	Lock/ auth	-	-	-	-	raidcom add hba_iscsi	-port <port#> [<host group name>]
-	-	Lock/ auth	-	-	-	-	raidcom delete hba_iscsi	-port <port#> [<host group name>]
-	-	Lock/ auth	-	-	-	-	raidcom set hba_iscsi	-port <port#> [<host group name>]
-	-	Lock/ auth	-	-	-	-	raidcom reset hba_iscsi	-port <port#> [<host group name>]
-	-	Auth	-	-	-	-	raidcom get hba_iscsi	-port <port#> [<host group name>]
-	-	Lock/ auth	-	-	-	-	raidcom add chap_user	-port <port#> [<host group name>]
-	-	Lock/ auth	-	-	-	-	raidcom delete chap_user	-port <port#> [<host group name>]

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/VDEV	LDEV	Command	Option
-	-	Lock/auth	-	-	-	-	raidcom set chap_user	-port <port#> [host group name>]
-	-	Lock/auth	-	-	-	-	raidcom reset chap_user	-port <port#> [host group name>]
-	-	Auth	-	-	-	-	raidcom get chap_user	-port <port#> [host group name>]
-	Auth	-	-	-	-	-	raidcom send ping	-port <port#>
-	Lock/auth	-	-	-	-	-	raidcom add external_iscsi_name	-
-	Lock/auth	-	-	-	-	-	raidcom delete external_iscsi_name	-
-	Lock/auth	-	-	-	-	-	raidcom modify external_chap_user	-
-	Lock/auth	-	-	-	-	-	raidcom modify initiator_chap_user	-
-	Auth	-	-	-	-	-	raidcom get external_iscsi_name	-
-	Auth	-	-	-	-	-	raidcom get initiator_iscsi_name	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	Auth	-	-	-	-	-	raidcom discover external_iscsi_name	-
-	Auth	-	-	-	-	-	raidcom check external_iscsi_name	-
-	Lock/auth	-	-	-	-	-	raidcom add rcu_iscsi_port	-
-	Lock/auth	-	-	-	-	-	raidcom delete rcu_iscsi_port	-
-	Auth	-	-	-	-	-	raidcom get rcu_iscsi_port	-
-	-	-	-	Lock/auth	-	-	raidcom modify parity_grp	-
-	-	-	-	-	-	-	raidcom modify local_replica_opt	-
-	-	-	-	-	-	-	raidcom get local_replica_opt	-
-	-	-	-	-	-	-	raidcom get license	-
-	-	-	-	-	-	Lock/auth	raidcom modify quorum	-quorum_id <quorum id>
-	-	-	-	-	-	Auth	raidcom get quorum	-quorum_id <quorum id>
-	-	-	-	-	-	Lock/auth	raidcom replace quorum	-quorum_id <quorum id> -ldev_id <ldev#>
-	-	-	-	Auth	Auth	Auth	raidcom get clpr	-key pg, -key vvol
-	-	-	-	-	-	Auth	raidcom add snapshot	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	Auth	raidcom delete snapshot	Except: - range tree
-	-	-	-	-	-	Lock/ auth	raidcom delete snapshot	-range tree
-	-	-	-	-	-	Auth	raidcom modify snapshot	-
-	-	-	-	-	-	Auth	raidcom map snapshot	-
-	-	-	-	-	-	Auth	raidcom unmap snapshot	-
-	-	-	-	-	-	Auth	raidcom get snapshot	-
-	-	-	-	-	-	Auth	raidcom replace snapshot	-
-	-	-	-	-	-	Lock/ auth	raidcom initialize pool	-ppid {<pool ID#> <pool naming>}
-	-	-	-	-	-	-	raidcom get system	-
-	-	-	-	-	-	Lock/ auth	raidcom add quorum	-ldev_id <ldev#>
-	-	-	-	-	-	Lock/ auth	raidcom delete quorum	-quorum_id <quorum id>
-	-	-	-	-	-	-	raidcom get apn	-
-	-	-	-	-	-	-	raidcom get ctg	-

* If another user locks the applicable resource, the command execution fails. To execute the command, the applicable resource authority is not required.

Relationship between commands supported only by VSP 5000 series and VSP One B20, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	Lock/ auth	raidcom add qos_grp	-ldev_id <ldev#>
-	-	-	-	-	-	Lock/ auth	raidcom delete qos_grp	-ldev_id <ldev#>
-	-	-	-	-	-	Lock/ auth	raidcom modify qos_grp	-
-	-	-	-	-	-	Auth	raidcom get qos_grp	-key resource
-	-	-	-	-	-	-	raidcom monitor resource	-

Relationship between commands supported only by VSP 5000 series, VSP One B20, VSP E1090, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume/ VDEV	LDEV	Command	Option
-	-	-	Lock/auth	-	-	-	raidcom add nvm_subsystem	-
-	-	-	Lock/auth	-	-	-	raidcom modify nvm_subsystem	-
-	-	-	Lock/auth	-	-	-	raidcom delete nvm_subsystem	-
-	Auth	-	Auth	-	-	-	raidcom get nvm_subsystem	-key namespace -key port
-	Lock/ auth	-	Lock/auth	-	-	-	raidcom add nvm_subsystem_port	-
-	Lock/ auth	-	Lock/auth	-	-	-	raidcom delete nvm_subsystem_port	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume/ VDEV	LDEV	Command	Option
-	-	-	Auth	-	-	-	raidcom get nvm_subsystem_ port	-
-	-	-	Lock/auth	-	-	-	raidcom add host_nqn	- nvm_subsy stem_id <nvm subsystem id>
-	-	-	-	-	-	-	raidcom modify host_nqn	-
-	-	-	Lock/auth	-	-	-	raidcom delete host_nqn	-
-	-	-	Auth	-	-	-	raidcom get host_nqn	-
-	-	-	Lock/auth	-	-	Lock/ auth	raidcom add namespace_path	-
-	-	-	Lock/auth	-	-	Lock/ auth	raidcom delete namespace_path	-
-	-	-	Auth	-	-	Auth	raidcom get namespace_path	-ldev_id <ldev#>
-	-	-	Lock/auth	-	-	Lock/ auth	raidcom add namespace	-
-	-	-	Lock/auth	-	-	Lock/ auth	raidcom delete namespace	-
-	-	-	Lock/auth	-	-	Lock/ auth	raidcom modify namespace	-
-	-	-	Auth	-	-	Auth	raidcom get namespace	-

Relationship between commands supported only by VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800 and resources that need to be locked

Res. group	Port	Host group	Parity group	External volume/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	raidcom add parity_grp	-
-	-	-	Lock/ auth	-	Lock/ auth	raidcom delete parity_grp	-

Relationship between commands supported by VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	Lock/ auth	raidcom add snap_pool	-
-	-	-	-	-	-	Auth	raidcom get snap_pool	-
-	-	-	-	-	-	Lock/ auth	raidcom monitor_pool	-
-	-	-	-	-	-	Lock/ auth	raidcom reallocate pool	-
-	Auth	-	-	-	-	-	raidcom add spm_wwn	-port <port#>
-	Auth	-	-	-	-	-	raidcom add spm_group	-port <port#>
-	_*	Auth	-	-	-	-	raidcom add spm_group	-port <port#> [-host group name>] - spm_host_grp

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	Auth	-	-	-	-	-	raidcom delete spm_wwn	-port <port#>
-	Auth	-	-	-	-	-	raidcom delete spm_group	-port <port#>
-	_*	Auth	-	-	-	-	raidcom delete spm_group	-port <port#> [host group name>] - spm_host_grp
-	Auth	-	-	-	-	-	raidcom modify spm_wwn	-port <port#>
-	Auth	-	-	-	-	-	raidcom modify spm_group	-port <port#>
-	_*	Auth	-	-	-	-	raidcom modify spm_group	-port <port#> [host group name>] - spm_host_grp
-	Auth	-	-	-	-	-	raidcom get spm_wwn	-port <port#>
-	Auth	-	-	-	-	-	raidcom get spm_group	-port <port#>
-	-	Auth	-	-	-	-	raidcom get spm_group	-port <port#> [host group name>] - spm_host_grp
-	Auth	-	-	-	-	-	raidcom monitor spm_wwn	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	Auth	-	-	-	-	-	raidcom monitor spm_group	-
-	-	-	-	-	-	Auth	raidcom modify spm_ldev	-ldev_id <ldev#>
-	-	-	-	-	-	Auth	raidcom delete spm_ldev	-ldev_id <ldev#>
-	-	-	-	-	-	Auth	raidcom monitor spm_ldev	-ldev_id <ldev#>
-	-	-	-	-	-	Auth	raidcom get spm_ldev	-ldev_id <ldev#>
-	-	-	-	Lock/ auth	Lock/auth	Lock/ auth	raidcom modify clpr	-ldev_id <ldev#> - parity_grp_id <gnosgno> - external_grp_id <gnosgno>

* If another user locks the applicable resource, the command execution fails. To execute the command, the applicable resource authority is not required.

Relationship between commands supported by VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	-	raidcom modify remote_replica_opt	-
-	-	-	-	-	-	-	raidcom get remote_replica_opt	-

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	-	raidcom modify path	-

Relationship between commands supported by VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	-	raidcom add server	-
-	-	-	-	-	-	-	raidcom delete server	-
-	-*	Lock/ auth	-	-	-	-*	raidcom modify server	-
-	-	-	-	-	-	-	raidcom get server	-
* If another user locks the applicable resource, the command execution fails. To execute the command, the applicable resource authority is not required.								

Relationship between commands supported by VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	Auth	raidcom modify system_opt	-
-	-	-	-	-	-	Auth	raidcom get system_opt	-key destage - cu <CU#>

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	Auth	-	-	raidcom get drive	-
-	-	-	-	Lock/ auth	-	-	raidcom initialize parity_grp	-
-	-	-	-	-	-	-	raidcom add license	-
-	-	-	-	-	-	-	raidcom delete license	-
-	-	-	-	-	-	-	raidcom modify license	-

Relationship between commands supported by VSP 5000 series, VSP G1x00, VSP F1500, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	-	raidcom modify system	-

Relationship between commands supported by VSP G800 and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	-	raidcom modify drivebox	-
-	-	-	-	-	-	-	raidcom get drivebox	-
-	-	-	-	-	-	-	raidcom check drivebox	-

Relationship between commands supported by VSP 5000 series, VSP G1x00, VSP F1500, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	-	raidcom get ssid	-

Relationship between commands supported by VSP 5000 series and VSP G800, and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume group/ VDEV	LDEV	Command	Option
-	-*	-	-	-	-	-	raidcom add ssid	-
-	-*	-	-	-	-	-	raidcom delete ssid	-

* If another user locks the applicable resource, the command execution fails. To execute the command, the applicable resource authority is not required.

Relationship between commands supported only by VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, and resources that need to be locked

Res. group	Port	Host group	Parity group	External volume/ VDEV	LDEV	Command	Option
-	-	-	-	-	-	raidcom modify drive*	-

* If another user locks a resource, the command fails.

Relationship between commands supported only by VSP One B20 and resources that need to be locked

Res. group	Port	Host group	NVM subsystem	Parity group	External volume/VDEV	LDEV	Command	Option
-	-	-	-	Lock/auth	-	-	raidcom extend parity_grp	-
-	-	-	-	Lock/auth	-	Lock/auth	raidcom extend parity_grp	- ldev_range <range>

Request ID function

Request ID is an identification number that manages the execution results of asynchronous commands that can specify the `-request_id auto` option. If you specify the `-request_id auto` option for an asynchronous command, a request ID is automatically assigned. For details about the asynchronous commands that can specify the `-request_id auto` option, see each asynchronous command.

Even if the execution of the asynchronous command fails, a request ID might be assigned and the request ID might be output. Each time the request ID is output, confirm the execution result of the command using the `raidcom get command_status` command.

Up to 65,280 request IDs can be assigned per storage system. If no request ID is available, the command fails with EX_IDEXHA.

If EX_IDEXHA fails, release unnecessary request IDs in the following procedure, and then retry the command. Or, release the request ID by another user using the request ID, and then retry the command.

1. Specify the request ID and confirm the execution result using the `raidcom get command_status` command.
2. Release the request ID whose execution result is confirmed by using the `raidcom reset command_status -request_id <request #>` command or the `raidcom reset command_status -request_id all` command.



Note:

- Use the `raidcom reset command_status -request_id <request #>` command to release the specific request ID.
- Use the `raidcom reset command_status -request_id all` command to release all the request IDs.

If you specify `-request_id all` for the `raidcom reset command_status` command, the request ID of the command whose execution result is not confirmed is also deleted.

raidcom add clpr

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900

Adds CLPR. If no option is specified, no operation is performed.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add clpr -clpr <clpr#> [-clpr_name <clpr name>] -cache_size <Cache Size(MB)>
```

Options and parameters

-clpr <clpr#>

Specifies the CLPR number to add. You can specify a number that is not assigned. An error occurs for the following:

- 0 or a number greater than the maximum CLPR number is specified as a CLPR number.
- The specified CLPR number is already used.

-clpr_name <clpr name>

Sets a CLPR name. An error occurs if you specify any of the following:

- A character other than alphanumeric characters
- A name consisting of more than 16 characters
- A reserved CLPR name
- A CRPR name that is already used

-cache_size <Cache Size(MB)>

Specifies the cache size (megabytes). An error occurs for the following:

- 0 to 4095 MB is specified.
- The size that is larger than 4095 MB is specified, but it is not a multiple of 2048 MB.
- Capacity exceeding the maximum capacity is specified.

Example

Creating a CLPR of CLPR number: 2, CLPR name: TEST02, and cache size: 8192 MB.

```
#raidcom add clpr -clpr 2 -clpr_name TEST02 -cache_size 8192
```

raidcom delete clpr

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900

Deletes CLPR. If no option is specified, no operation is performed.

This command is executed asynchronously with the command input. Check the completion of this process on the raidcom get command_status command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete clpr -clpr <clpr#>
```

Options and parameters**-clpr <clpr#>**

Specifies the CLPR number to delete. An error occurs if you specify either of the following:

- 0 or a number greater than the maximum CLPR number is specified as a CLPR number.
- A CLPR for which the parity group, LDEV, or external group is registered is specified.

Example

Deleting CLPR of CLPR number: 2.

```
#raidcom delete clpr -clpr 2
```

raidcom get clpr

Displays the CLPR information for the storage system.

Syntax

```
raidcom get clpr
```

Options and parameters

None

Examples

```
# raidcom get clpr
```

CLPR	CLPR_NAME	TC_CAP(MB)	TU_CAP(MB)	WP_CAP(MB)	SF_CAP(MB)	U(%)	W(%)	S(%)
001	Oracle_DB_PROD	20000	10000	2000	0	50	10	0
003	Oracle_DB_BACK	10000	5000	500	0	50	5	0

Description of the `raidcom get clpr` output:

CLPR

CLPR ID (decimal)

CLPR_NAME

Nickname of the CLPR

TC_CAP(MB)

Capacity (megabytes) of cache memory of the CLPR

TU_CAP(MB)

Used capacity (megabytes) of cache memory of the CLPR

WP_CAP(MB)

Capacity (megabytes) of write pending data of the CLPR

SF_CAP(MB)

Capacity (megabytes) of sidefiles of the CLPR

U(%)

Usage rate of cache memory of the CLPR

W(%)

Rate of write pending data of the CLPR

S(%)

Usage rate of sidefiles of the CLPR

raidcom modify clpr

Modifies a CLPR.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify clpr -clpr <clpr#> { -ldev_id <ldev#> | -parity_grp_id <gno-sgno> | -external_grp_id <gno-sgno>} | {[-clpr_name <clpr name>] [-cache_size <Cache Size (MB)>]}
```

Options and parameters

-clpr <clpr#>

Specifies a CLPR number.

For example:

- -clpr 2

-ldev_id <ldev#>

Specifies an LDEV number (0-65279).

For example:

- -ldev_id 200

You cannot use this option to specify an LDEV in a LUSE volume. For the other considerations, see the *Performance Guide*.

If all of the following conditions are met, the CLPR of the journal volume can be moved. However, do not move the CLPR of the journal volume by using this command unless instructed otherwise:

- The CCI version is 01-76-03/03 or later.
- A reserve journal volume does not exist in the journal to which the target journal volume belongs.

If the reserve journal volume exists, run this command after deleting the reserve journal volume. Add the deleted reserve journal volume after changing the CLPR by using this command.

If the reserve journal volume exists, when the CLPR of the journal volume is moved, the command fails with EX_CMDRJE (SSB1=2E30, SSB2=000F).

-parity_grp_id <gno-sgno>

Specifies a parity group number (gno: 1-52, sgno: 1-32).

For example:

- 3-1

-external_grp_id <gno-sgno>

Specifies an external volume group number (gno: 1-16384, sgno: 1-4096).

For example:

- 52-11

-clpr_name <clpr name>

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specify this option to change the CLPR name. This option cannot be specified in the same command line as the `-ldev_id`, `-parity_grp_id`, and `-external_grp_id` options. An error occurs if you specify any of the following:

- A character other than alphanumeric characters
- A name consisting of more than 16 characters
- A reserved CLPR name. However, you can restore the default value of each CLPR.
- A CLPR name that is already used

-cache_size <Cache Size(MB)>

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specify this option to change the cache size (megabytes). This option cannot be specified in the same command line as the `-ldev_id`, `-parity_grp_id`, and `-external_grp_id` options. An error occurs if you specify any of the following:

- CLPR0
- 0 to 4095 MB
- Larger than 4095 MB, but not a multiple of 2048 MB
- Capacity exceeding the maximum capacity

Examples

Moving the LDEV 02:00 to the CLPR ID 2.

```
# raidcom modify clpr -clpr 2 -ldev_id 0x0200
```

Moving the parity group 5-2 to the CLPR ID 2

```
# raidcom modify clpr -clpr 2 -parity_grp_id 5-2
```

Moving the external volume group 1-1 to the CLPR ID 2.

```
# raidcom modify clpr -clpr 2 -external_grp_id 01-01
```

Changing the CLPR name of CLPR ID 2 to TEST02. (Only for VSP 5000 series, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900)

```
# raidcom modify clpr -clpr 2 -clpr_name TEST02
```

Changing the cache size of CLPR ID 2 to 4 GB. (Only for VSP 5000 series, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900)

```
# raidcom modify clpr -clpr 2 -cache_size 4096
```

raidcom get system_opt

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Views system options.

Syntax

```
raidcom get system_opt [-key {destage -cu <CU#> | mode -lpr <system|clpr#>}]
```

Options and parameters

-key <value>

Displays the setting status of a system option.

- destage -cu <CU#>: Displays the setting status of the write through operation. Nothing is displayed if the CU which is not installed is specified. For VSP G200, G400, G600, G800 and VSP F400, F600, F800, this option cannot be specified.
- mode -lpr <system|clpr#>: Displays the setting status of a system option mode. Nothing is displayed if a CLPR number which is not assigned is specified. For VSP G200, G400, G600, G800 and VSP F400, F600, F800, make sure to specify this option. If this option is not specified, a syntax error occurs.

Examples

Displaying the basic information of the system option (without the -key option).

```
#raidcom get system_opt
Serial# : 563528
SDR : interleave
DCP : medium
CC : Enable
DS : Enable
LFT : 255
```

Output of the raidcom get system_opt command:

Serial#: Product serial number.

SDR: Spare disk recover. Displays the speed of copying data to a spare drive.

- interleave: I/Os from the host are prioritized. Copies data for slots, and then stops the copy processing for a certain period of time.
- fullspeed: Copy processing is prioritized.

DCP: Disk copy pace. When interleave is specified for SDR, displays the speed of copying data to a spare drive.

- faster: Copying data to the drive takes precedence over jobs from the host.
- medium: Optimal mode. The copying time is determined by the load of host I/Os
- slower: Jobs from the host takes precedence over the drive copy.

CC: Correction copy. Specifies the operation to be performed when the drive is blocked.

- Enable: When the drive is blocked, correction copy to the spare drive is performed.
- Disable: When the drive is blocked, correction copy to the spare drive is not performed.

DS: Dynamic sparing. Displays the behavior when a drive failure occurred because the threshold was exceeded.

- Enable: If a drive failure occurs because the threshold was exceeded, data is copied to the spare drive automatically.
- Disable: If a drive failure occurs because the threshold was exceeded, data is not copied to the spare drive automatically.

LFT: Link failure threshold. Displays the threshold (0 to 255) for notifying a link failure.

Displaying the setting status of the write through operation for each LDEV of CU#: 0.

```
#raidcom get system_opt -key destage -cu 0
Serial# LDEV# DESTAGE
563528 100 E
563528 101 E
563528 102 E
563528 103 E
```

Output of the raidcom get system_opt command:

DESTAGE: Displays the setting status of the write through operation.

- E: The write through operation setting is disabled (the write through suppression setting is enabled).
- D: The write through operation setting is enabled (the write through suppression setting is disabled).

Specifying " system " and displaying the setting status of system option mode.

```
#raidcom get system_opt -key mode -lpr system
Serial# : 563528
LPR : system
CACHE_T : level2
COM_CTL : 0xffffffff
MODE_ID : 117 256 2047
```

Output of the raidcom get system_opt command:

LPR: Displays the specified " system " or the CLPR number.

CACHE_T: CACHE tuning. Displays the setting status (level1 to level5) of the cache tuning level. A hyphen (-) is displayed if LPR is not " system " .

COM_CTL: Command control. Displays the setting status (0x00000000 to 0xffffffff) of information for switching read ahead conditions. A hyphen (-) is displayed if LPR is not " system " .

MODE_ID: Lists the numbers of the system option modes set to ON.

raidcom modify user_system_opt

Supported storage systems:

- VSP One B20
- VSP E590, VSP E790, and VSP E1090

Changes the advanced system settings.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify user_system_opt -opt_id <Option ID> -opt <enable|disable>
```

Options and parameters

-opt_id <Option ID>

Specifies the option ID of the advanced system settings in decimal.

If you want to change the enable or disable setting of option IDs not listed below, change the setting in the **Edit Advanced System Settings** window.

For VSP One B20, the option IDs 15, 16, 17, and 18 can be specified by using the **raidcom modify user_system_opt** command. For details about the functions of the option IDs 15, 16, 17, and 18, see the *Hitachi Universal Replicator User Guide*.

- 5: If you enable this option, the differential data is maintained in a pool with which a DP-VOL that has the user capacity (up to 4,194,304 MB) is linked when a new TC, UR, or GAD pair is created using the DP-VOL, or resynchronized (differential data management). Note that differential data is maintained, regardless of this setting, in a pool linked with the DP-VOL that has the user capacity (greater than 4,194,304 MB) when a new TC, UR, or GAD pair using the DP-VOL is created.
- 6: If you enable this option, the differential data is maintained in a pool with which a DP-VOL that has the user capacity (up to 4,194,304 MB) is linked when a new TC, UR, or GAD pair is created using the DP-VOL (differential data management). Note that differential data is maintained, regardless of this setting, in a pool linked with the DP-VOL that has the user capacity (greater than 4,194,304 MB) when a new TC, UR, or GAD pair using the DP-VOL is created.
- 9: If you enable this option, when tier relocation is suspended by the system, an alert is issued to users. For details about an alert (SIM) to be issued, see the *Provisioning Guide*.



Note: If both the option IDs 5 and 6 of the advanced system settings are enabled, the setting of the option ID 5 is ignored.

--opt <enable|disable>

Specifies the settings to enable or disable use of the advanced system settings option specified with -opt_id.

- enable: Enables the specified option of the advanced system settings.
- disable: Disables the specified option of the advanced system settings.

Examples

Enables the option ID 5 of the advanced system settings.

```
# raidcom modify user_system_opt -opt_id 5 -opt enable
```

raidcom get user_system_opt

Supported storage systems:

- VSP One B20
- VSP E590, VSP E790, and VSP E1090

Views the enabled option of the advanced system settings.

Syntax

```
raidcom get user_system_opt
```

Options and parameters

None

Examples

Viewing the enabled option of the advanced system settings.

```
# raidcom get user_system_opt
OPT_ID : 5 6
```

Description of each column in output example:**OPT_ID**

Displays the enabled option ID of the advanced system settings.

raidcom get command_status

It displays error information of the configuration setting command (asynchronous command) to be executed asynchronously.

When an error occurs with the execution of an asynchronous command, the total number of errors or error information such as error codes (SSB1 and SSB2) are stored in the storage system at the first occurrence. After executing asynchronous command, check the error information by executing this command.

The error information is deleted if you execute the **raidcom reset command_status** command, or log out from the storage system. For details about **raidcom reset command_status**, see [raidcom reset command_status \(on page 307\)](#). When you log out from the system, all error information for the session you used is deleted.

The error codes of the asynchronous commands are stored in different areas depending on request ID is specified or not. The options at the time of command execution and the status for each storage area are as follows.

Asynchronous command type	Storage area per login session ¹	Storage area per request ID ²
-request_id auto is specified	Stored	Stored
-request_id auto is not specified	Stored	Not stored
Notes:		

Asynchronous command type	Storage area per login session ¹	Storage area per request ID ²
<ol style="list-style-type: none"> 1. This area stores the error information prepared in the storage system for each login session. This area stores the error of the first occurrence during the login session. However, for an error of the second occurrence, error codes of SSB1 and SSB2 are not stored. To be able to refer the error code when an error occurs, reset the error information that is stored in the storage area for each login session by executing raidcom reset command_status before executing asynchronous commands. 2. This area stores the error information prepared for each request ID. The stored error code can be checked individually by specifying the request ID. 		

Syntax

```
raidcom get command_status [-time <time(sec)>] [-request_id <request#>]
```

Options and parameters

[-time <time(sec)>]

Specifies waiting time (0 to 21474836) to complete the process of asynchronous command.

If this option is omitted, the default waiting time (CMD_DEF_TMOUT: 7200 sec.) is set.

[-request_id <request#>]

Specifies the request ID. Specifying this option displays error information of the command specified by the request ID. If you specify the request ID of the command issued by another user, this command ends without displaying the information.

This command interprets <request#> as a hexadecimal number. If the specified <request#> satisfies either of the following conditions, EX_INVARG is returned:

- <request#> contains characters other than numbers and letters (a to z, A to Z).
- <request#> contains nine or more characters.

Returned Values

Either of the following returned values is returned to exit (), which allows users to check the execution results using a user program or a script.

- **0**: Normal Termination
- **1**: One or more errors occurred. Abnormal termination.

Examples

Displaying error information of the asynchronous command:

```
# raidcom get command_status
```

```
HANDLE SSB1 SSB2 ERR_CNT Serial# Description
7E30    2E20 6000      4   64034 The pool ID is not installed
```

Specifying request ID: 1, and displaying error information of the asynchronous command:

```
# raidcom get command_status -request_id 1
REQID R SSB1 SSB2 Serial# ID Description
00000001 - 2E20 6000 64034 3 The pool ID is not installed
```

Description of the `raidcom get command_status` output:

HANDLE

Handle number that uniquely identifies the user

SSB1

SSB1 error code. For details about error codes, see the *Command Control Interface User and Reference Guide*.

SSB2

SSB2 error code. For details about error codes, see the *Command Control Interface User and Reference Guide*.

ERR_CNT

Total number of errors in this user handle

Serial#

Serial number.

REQID

Request ID of the command.

R

Displays if the error is caused by a failure of another command.

- 0: The error occurred due to an error caused by a failure of another command. SSB1, SSB2, and the description show error codes and error information of the failure of another command.
- T: The command failed due to the command specified by the `-request_id` option. SSB1, SSB2, and the description show error codes and error information of the command specified by the option.
- - (hyphen): This information is not available for this error.

ID

Displays ID of the object related to the request ID. A hyphen (-) is displayed if the ID information is not available.

- An LDEV number is displayed if the request ID that is output in the following cases is specified for the `-request_id` option.
 - When the `-ldev_id auto -request_id auto` option is specified in the **raidcom add ldev** command.
 - When any of the following options are specified in the **raidcom modify ldev** command.
 - `-upper_throughput_io <upper throughput io> -request_id auto`
 - `-upper_data_trans_mb <upper data trans mb> -request_id auto`
 - `-upper_alert_time <upper alert time> -request_id auto`
 - `-lower_throughput_io <lower throughput io> -request_id auto`
 - `-lower_data_trans_mb <lower data trans mb> -request_id auto`
 - `-lower_alert_time <lower alert time> -request_id auto`
 - `-response_priority <#priority> -request_id auto`
 - `-response_alert_time <response alert time> -request_id auto`
- An LDEV number is displayed if the request ID that is output when the `-ldev_id auto -request_id auto` option is specified for the **raidcom add ldev** command or the `-request_id auto` option is specified for the **raidcom extend ldev** command is specified for the `-request_id` option.
- A hyphen (-) is displayed if the request ID that is output in the following cases is specified for the `-request_id` option:
 - When the `-ldev_id auto -request_id auto -ddp_ldev` option is specified for the **raidcom add ldev** command
 - When the `-ddp_ldev -request_id auto` option is specified for the **raidcom delete ldev** command
 - When the `-request_id auto` option is specified for the **raidcom map resource** command, the **raidcom unmap resource** command, the **raidcom add parity_grp** command, the **raidcom modify parity_grp** command, or the **raidcom extend parity_grp** command
- A pool number is displayed if the request ID that is output when the `-parity_grp_id <gno-sgno>-ddp_ldev -request_id auto` option is specified for the **raidcom add dp_pool** command or the **raidcom delete pool** command is specified for the `-request_id` option.

- An LU number is displayed if the request ID that is output when the `lun_id auto -request_id auto` option is specified for the **raidcom add lun** command is specified for the `-request_id` option.
- A server ID is displayed if the request ID that is output when executing the **raidcom add server** command, **raidcom delete server** command, or **raidcom modify server** command is specified for the `-request_id` option.
- A QoS group ID is displayed if the request ID that is output when executing the **raidcom add qos_grp** command, **raidcom delete qos_grp** command, or **raidcom modify qos_grp** command is specified for the `-request_id` option.
- An NVM subsystem ID is displayed if the request ID that is output when running the **raidcom add nvm_subsystem** command, **raidcom modify nvm_subsystem** command, **raidcom delete nvm_subsystem** command, **raidcom add nvm_subsystem_port** command, **raidcom delete nvm_subsystem_port** command, **raidcom add host_nqn** command, **raidcom modify host_nqn** command, **raidcom delete host_nqn** command, **raidcom add namespace_path** command, or **raidcom delete namespace_path** command is specified for the `-request_id` option.
- A namespace ID is displayed if the request ID that is output when running the **raidcom add namespace** command, **raidcom modify namespace** command, or **raidcom delete namespace** command is specified for the `-request_id` option.

Description

Error information. If no error occurred, a hyphen (-) is displayed.

EX_EWSTOT

Timeout while waiting the result of the command execution

For details, see the section describing **Command error messages** in the *Command Control Interface User and Reference Guide*.

raidcom reset command_status

Resets the error information of the configuration setting command that is stored in the storage system and executed asynchronously (Asynchronous command).

The command options and the information to be reset are as follows:

Option type	Storage area per login session	Storage area per request ID
No option	Deleted	-
<code>-request_id <request#></code>	-	Deleted ¹
<code>-request_id all</code>	-	Deleted ²
Notes:		

Option type	Storage area per login session	Storage area per request ID
<ol style="list-style-type: none"> 1. Deletes only the error information of the specified <request #> used in the login session in which the command is executed. If you specify <request #> used in other login session, the deletion is not be performed. 2. Deletes all error information for each request ID used in the login session in which the command is executed. 		

Syntax

```
raidcom reset command_status [-request_id <request#> | -request_id all]
```

Options and parameters

[-request_id <request#> | -request_id all]

Specifies the request ID of a command whose error information is cleared.

- -request_id <request#>: Specifies the request ID of a command whose error information is cleared.
- -request_id all: Clears all error information of the raidcom add ldev command which is specified with the -ldev_id auto option and executed by the user who executes the raidcom reset command_status command.

For <request#>, type the request ID output by the **raidcom add ldev** command. This command interprets <request#> as a hexadecimal number. If the specified <request#> satisfies any of the following conditions, EX_INVARG or EX_CMDRJE is returned:

- <request#> contains characters other than numbers and letters (a to z, A to Z) (EX_INVARG is returned).
- <request#> contains nine or more characters (EX_INVARG is returned).
- <request#> is a hexadecimal number, and it is interpreted as 0x00000000 (EX_CMDRJE is returned).
- <request#> is a hexadecimal number, and it is interpreted as a value between 0x0000ff01 and 0xffffffff (EX_CMDRJE is returned).

Examples

Resetting the error information of the asynchronous command.

```
# raidcom reset command_status
```

Resetting the error information of the command for request ID: 1.

```
# raidcom reset command_status -request_id 1
```

Resetting the error information of the raidcom add ldev command with the -ldev_id auto option executed by the user.

```
# raidcom reset command_status -request_id all
```

raidcom add copy_grp

Creates a copy group.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add copy_grp -copy_grp_name <copy group name>
    <device group name> [<device group name>] [-mirror_id <mu#>]
    [-journal_id <journal ID#>]
```

Options and parameters

-copy_grp_name <copy group name><device group name>[<device group name>]

Specifies the device group (maximum 32 characters) configuring a copy group (maximum 32 characters).

You can specify up to two device group names. If you specify more than two, the option is ignored.

For a copy group for an ShadowImage pair, specify two device groups.

For a copy group for a TrueCopy pair, specify only one device group for the relevant storage system side (primary/main or secondary/remote).

[-mirror_id <mu#>]

Specifies the mirror ID.

If this option is omitted (by **raidcom get copy_grp**), "-" is displayed.

The mirror ID to be registered is as follows depending on whether the environment variable is set or not and whether the option is specified or not.

Option	HORCC_MRCF environment variable is not set	HORCC_MRCF environment variable is set
-IM and -IH options are not specified	hx	x
-IM option is specified	x	x
-IH option is specified	hx	hx
x: The mirror ID specified in <mu#>		

[-journal_id <journal ID#>]

Specifies the journal number (0-255).

The journal number can be specified in decimal or hexadecimal. If you specify a hexadecimal number, add 0x at the beginning of the journal number.

If this option is omitted (by raidcom get copy_grp), "-" is displayed.

Example

Creating a copy group (ora) by device groups (grp1, grp2).

```
# raidcom add copy_grp -copy_grp_name ora grp1 grp2
```

raidcom delete copy_grp

Deletes the specified copy group.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete copy_grp -copy_grp_name <copy group name>
```

Options and parameters

-copy_grp_name <copy group name>

Specifies the name of the copy group (maximum 32 characters).

Example

Deleting the copy group: ora.

```
# raidcom delete copy_grp -copy_grp_name ora
```

raidcom get copy_grp

Displays the information of the specified copy group.

Syntax

```
raidcom get copy_grp
```

Options and parameters

None.

Example

Displaying copy group information.

```
# raidcom get copy_grp
```

```
COPY_GROUP LDEV_GROUP MU# JID# Serial#
ora grp1      0 -      64034
ora grp2      0 -      64034
```

Description of the raidcom get copy_grp output:**COPY_GROUP**

Copy group name

LDEV_GROUP

Device group name that composes copy group

MU#

Mirror ID to which the device group belongs. If -mirror_id is not specified at the creation, "-" is displayed.

JID#

Journal number to which device group belongs in decimal. If -journal_id is not specified at the creation, "-" is displayed.

Serial#

Product serial number. The display of the product serial number varies depending on storage systems.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

raidcom add device_grp

Assigns a device name to the specified LDEV, and creates a device group. If the group already exists, the LDEV is added to the group. If the LDEV also already exists in the specified device group, the specified LDEV name is set.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add device_grp -device_grp_name <device group name>
<device name> -ldev_id <ldev#>... [-cnt <count>]
```

Options and parameters**device_grp_name <device group name><device name>**

Specifies the device group name (maximum 32 characters) and the device name in the device group (maximum 32 characters). If multiple LDEVs are specified, the same device name is set for all of them.

-ldev_id <ldev#> ...

Specifies the LDEV number (0-65279).

When you specify an LDEV that is part of a LUSE volume, all LDEVs in the LUSE volume have the same name. For example:

- -ldev_id 200
- -ldev_id 100-110
- -ldev_id 100 -cnt 10

Up to 64 of LDEVs can be specified.

When you specify an LDEV in a LUSE volume, the number of LDEVs in the LUSE volume is included.

[-cnt <count>]

Specifies the count (2-64).

If this option is omitted, the count is set to one.

Example

Assigning a device name: data1 to an LDEV: 400 and adding it to the device group: grp1.

```
# raidcom add device_grp -device_grp_name grp1 data1 -ldev_id 400
```

raidcom delete device_grp

Deletes the specified LDEV from the specified group. When the last LDEV is deleted, the device group is also deleted.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete device_grp -device_grp_name <device group name>
                        -ldev_id <ldev#>... [-cnt <count>]
```

Options and parameters**-device_grp_name <device group name>**

Specifies the device group name (maximum 32 characters)

-ldev_id <ldev#> ...

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200
- -ldev_id 100-110
- -ldev_id 100 -cnt 10

Up to 64 of LDEVs can be specified. When you specify an LDEV in a LUSE volume, the number of LDEVs in the LUSE volume is included.

[-cnt <count>]

Specifies the count (2-64).

If this option is omitted, the count is set to one.

Example

Deleting the LDEV400 from the device group: grp1.

```
# raidcom delete device_grp -device_grp_name grp1 -ldev_id 400
```

raidcom get device_grp

Displays the LDEV information for the specified device group, or lists all device groups.

Syntax

```
raidcom get device_grp [-device_grp_name <device group name>]
```

Options and parameters

[-device_grp_name <device group name>]

Displays the device (LDEV) information for the specified device group (maximum 32 characters).

If this option is omitted, the list of the registered device groups is displayed.

Example for displaying device group information

```
# raidcom get device_grp
```

```
LDEV_GROUP Serial#
grp1 64034
grp2 64034
grp3 64034
```

Example for displaying device group information: grp1

```
# raidcom get device_grp -device_grp_name grp1
```

```
LDEV_GROUP LDEV_NAME LDEV# Serial#
grp1      data1      400    64034
grp1      data2      401    64034
```

Description of the `raidcom get device_grp` output:

LDEV_GROUP

Device group name

LDEV_NAME

Device name in the device group

LDEV#

LDEV number

Serial#

Product serial number.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

raidcom get drive

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Displays drive information.

Syntax

```
raidcom get drive [-parity_grp_id <gno-sgno> | -usage <usage>]
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno: 1 to 52, sgno: 1 to 32)

Example:

3-1

-usage <usage>

Specifies the drive usage.

Specify the following character strings for <usage>:

- data: Data drive
- spare: Spare drive
- free: Unused drive

Example

Displaying drive information.

```
#raidcom get drive
```

LOCATION	TYPE	RPM	TOTAL_CAP (GB)	CODE	USAGE	STS	GROUP
1-1	SAS	15000	300	DKS5C-K300SS	DATA	NML	1-1

Description of the **raidcom get drive** output:

LOCATION

Displays the location of the drive in the xx-yy format. In Device Manager - Storage Navigator, the drive location is displayed in the HDDxx-yy format.

TYPE

Displays the drive type.

(VSP 5000 series) SSD or SCM is displayed as a drive type of a drive whose drive type code (CODE) is SPx5x-YxxxNC, where x is any character string or number. The displayed character string depends on the DKCMAIN microcode version.

RPM

Displays the number of revolutions of a drive in rpm. For SSD, a hyphen (-) is displayed.

TOTAL_CAP(GB)

Displays the capacity of a drive in gigabytes.

CODE

Displays the drive type code.

USAGE

Displays the drive usage.

- DATA: Data drive
- SPARE: Spare drive
- FREE: Unused drive

STS

Displays the drive status.

- NML: Normal.
- WAR: A blocked part exists.
- CPY: Drive copy is in process.
- CPI: Copy is incomplete.
- RSV: Spare disk is unusable.
- FAI: Blocked due to a failure.
- BLK: Blocked due to maintenance.
- UNK: The status is unknown.

GROUP

Displays the parity group number of the drive if it is contained in a parity group. If the drive is not contained in a parity group, a hyphen (-) is displayed.

raidcom modify drive

Supported storage systems:

- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Configures spare drives or cancels the settings.

This command is executed asynchronously with the command input. Use the **raidcom get command_status** command to check if the command is completed.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify drive -drive_location <drive location> -spare {enable|disable}
```

Options and parameters

-drive_location <drive location>

Specifies the drive location.

Example:

When the drive location is HDD00-01 (Drive00-01): 0-1

-spare {enable|disable}

Sets a spare drive or cancels the setting.

- enable: Sets a drive as a spare drive.
- disable: Cancels the spare drive setting.

Example

Setting a drive whose drive location is HDD00-01 as a spare drive.

```
# raidcom modify drive -drive_location 0-1 -spare enable
```

raidcom get error_message

Displays the error message for the specified error code.

Syntax

```
raidcom get error_message -ssb <ssb1> <ssb2>
```

Options and parameters

-ssb <ssb1> <ssb2>

Specifies the error code as a hexadecimal number (add the 0x prefix).

- <ssb1>: Specifies SSB1 of the error code.
- <ssb2>: Specifies SSB2 of the error code.

Example

Displaying the error message for the error code whose SSB1 is 0x2E00 and SSB2 is 0x0023.

```
# raidcom get error_message -ssb 0x2E00 0x0023
CAUSE : Volume capacity is too small.
```

raidcom add external_grp

Maps the volume of the external storage system to the external volume group on the local storage system. Only one volume can be mapped in a single operation.

Maps the specified volume on the external storage system to the specified external volume group by using the path that is established between the specified local storage system port and the external storage system port specified by WWN. If the specified path group already exists, the external volume group is added to the path group. If the specified path group does not exist, a new path group is created.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.

This command is rejected by EX_ENOOBJ in the following cases:

- The specified iSCSI port cannot be found.
- The iSCSI virtual port mode is enabled, but the specified iSCSI virtual port ID is not correct.

To fix this error, specify the correct port and iSCSI virtual port ID.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Note: If the local storage system port is a Fiber Channel Bidirectional port, it is recommended that you always execute the command by specifying the `-safety_check enable` option unless instructed by the Command Control Interface User and Reference Guide.

If you do not specify the `-safety_check enable` option, the following problems might occur. For the details, see the Hitachi Universal Volume Manager User Guide.

- If you use a non-Hitachi external storage system:
The I/O path from the external storage system to the local storage system that uses the route between the specified local storage system port and the external storage system port might be disconnected.
- If you use a Hitachi external storage system:
The external or remote path connection that uses the route between the specified local storage system port and the external storage system port might be temporarily disconnected (the external or remote path will be reconnected immediately and no blockage occurs).

Syntax

```
raidcom add external_grp -path_grp <path group#>
    -external_grp_id <gno-sgno> -port <port#> [-external_wwn
    <wwn strings> | -external_iscsi_name <external iscsi name>
    -external_address <IP address> [-iscsi_virtual_port_id
    <iSCSI virtual port ID>]] -lun_id <lun#>
    [-emulation <emulation type>] [-clpr <clpr#>]
    [-external_attribute migration] [-data_direct_mapping enable]
    [-command_device y -ldev_id <ldev#>] [-safety_check enable]
```

Options and parameters

-path_grp <path group#>

Specifies the external VOL path group number (0-63231).

-external_grp_id <gno-sgno>

Specifies the external volume group number (gno:1-16384, sgno:1-4096). For example:

- 52-11

-port <port#>

Specifies the port number on the local storage system. Specifies the port number whose attribute is External. For example:

- CL1-A

Displays an external port.

-external_wwn <wwn strings>

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The 17th digit and later are ignored. The value can be split in units of 4 bytes by ",", (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-external_iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: `iqn.` and the subsequent maximum 219 characters.
- eui format: `eui.` and the subsequent 16 characters in hexadecimal notation.

**Note:**

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

-external_address <IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).



Note:

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the **raidcom get external_iscsi_name** command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

-lun_id <lun#>

Displays LUN (0-4095) of the external storage system port.

[-emulation <emulation type>]

Specifies the emulation type. If this option is omitted, OPEN-V is used.

The valid values for <emulation type> are:

- OPEN-3, OPEN-8, OPEN-9, OPEN-E, OPEN-K, OPEN-L, OPEN-V
- 3390-1, 3390-2, 3390-3, 3390-A, 3390-3A, 3390-3B, 3390-3C, 3390-3R, 3390-9, 3390-9A, 3390-9B, 3390-9C, 3390-L, 3390-LA, 3390-LB, 3390-LC, 3390-M, 3390-MA, 3390-MB, 3390-MC, 3390-V
- 3380-3, 3380-3A, 3380-3B, 3380-3C

Some emulation types cannot be specified according to the type of device.



Caution:

You can specify 3390-3 or 3390-3R as the emulation type, but these are unable to be mixed. You can specify the 3380 series or the 3390 series as the emulation type for each parity group, but they are unable to be mixed for each 32 address boundary because of the OS restriction.

[-clpr <clpr#>]

Specifies the CLPR number.

[-external_attribute migration]

Specifies if an attribute of NDM functions is set.

[-data_direct_mapping enable]

Specifies to set the data direct mapping attribute. If you do not specify this option, the data direct mapping attribute is not set. You can specify either of the following parameters in this option:

- enable or yes: Sets the data direct mapping attribute.
- disable or no: Does not set the data direct mapping attribute.

Note that the data direct mapping attribute is automatically set to an LDEV which is created in the external volume group having the data direct mapping attribute. For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

[-command_device y -ldev_id <ldev#>]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies to map a volume on the external storage system as a remote command device. The LDEV number specified by -ldev_id <ldev#> is set for the remote command device.

[-safety_check enable]

Specify the Fiber Channel bidirectional port for the -port option. If you specify any other port, this option is ignored. When this option is specified, processing that might cause the disconnection between the Bidirectional port specified by the -port option and the Fiber Channel port of the external storage system specified by the -external_wwn option is suppressed. When processing is suppressed, a volume on the external storage system is not mapped.

Examples

Mapping an LU: 0 defined to the external storage system port: 50060e80,05fa0f36 connected to the External (ELUN) port: CL1-A by the External Volume Group#1-1 and the path group#1.

```
# raidcom add external_grp -path_grp 1 -external_grp_id 1-1
-port CL1-A -external_wwn 50060e80,05fa0f36 -lun_id 0 -safety_check enable
```

Mapping an LU:0 defined to the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100) which is connected to the port: CL1-A (iSCSI port) of the local storage system by the External Volume Group#1-1 and the path group#1.

```
# raidcom add external_grp -path_grp 1 -external_grp_id 1-1 -port CL1-A -
external_iscsi_name iqn.z2 -external_address 158.214.135.100 -lun_id 0
```

Mapping an LU:0 defined to the external storage system port: 50060e80,05fa0f36 connected to the External (ELUN) port: CL1-A by the External Volume Group#1-1 and the path group#1, and setting the attribute for the NDM function and the data direct mapping attribute.

```
# raidcom add external_grp -path_grp 1 -external_grp_id 1-1 -port CL1-A -external_wnn
50060e80,05fa0f36 -lun_id 0 -external_attribute migration -data_direct_mapping enable
-safety_check enable
```

Mapping an LU:0 defined to the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100) connected to the iSCSI port: CL1-A, iSCSI virtual port ID: 2 of the local storage system by the External Volume Group#1-1 and the path group#1.

```
# raidcom add external_grp -path_grp 1 -external_grp_id 1-1 -port CL1-A
-iscsi_virtual_port_id 2 -external_iscsi_name iqn.z2 -external_address
158.214.135.100 -lun_id 0
```

Mapping an LU: 0 defined to the external storage system port: 50060e80, 05fa0f36 connected to the External (ELUN) port: CL1-A by the External Volume Group#1-1 and the path group#1, and setting the LDEV number: 1.

```
# raidcom add external_grp -path_grp 1 -external_grp_id 1-1 -port CL1-A
-external_wnn 50060e80,05fa0f36 -lun_id 0 -command_device y -ldev_id 1 -safety_check
enable
```

raidcom check_ext_storage external_grp

Specifies the external volume group, check the connection to a volume on the external storage system, and then restart using. Only one external volume group is operated in each operation.

An LDEV or device group can be specified instead of an external volume group:

- If an LDEV is specified, CCI finds the external volume groups to which the specified LDEV belongs and displays the result.
- If a device group is specified, CCI finds the external volume groups to which the specified device group belongs and displays the result.

If no LDEV exists in the external volume, the command is rejected with EX_ENOOBJ.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom check_ext_storage external_grp {-external_grp_id
<gno-sgno> | -ldev_id <ldev#>} | -grp_opt <group option>
-device_grp_name <device group name> [<device name>]}
```

Options and parameters**-external_grp_id <gno-sgno>**

Specifies the external volume group number (gno:1-16384, sgno:1-4096). For example:

- 52-11

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information about LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all of the LDEVs in the device group are operated.

Examples

Executing the confirmation of existence and the LDEV recovery for the external volume group #1-1.

```
# raidcom check_ext_storage external_grp -external_grp_id 1-1
```

Executing the confirmation of connection and the LDEV recovery for the external volume group including the external volume (LDEV:200).

```
# raidcom check_ext_storage external_grp -ldev_id 200
```

Executing the confirmation of connection and the LDEV recovery for the external volume group including the LDEV belonging to the device group: grp1.

```
# raidcom check_ext_storage external_grp -grp_opt ldev -device_grp_name grp1
```

raidcom delete external_grp

Deletes the mapped volume on the external storage system to delete the registered external volume group from the configuration. Only one external volume group is deleted in each operation. When the last external volume group is deleted, the path group is also deleted.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete external_grp -external_grp_id <gno-sgno> [-forcible]
```

Options and parameters

-external_grp_id <gno-sgno>

Specifies the external volume group number (gno:1 to 16384, sgno:1 to 4096). For example:

- 52-11

[-forcible]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specify this option to delete an external volume group whose connection is not disconnected. Specifying this option deletes an external volume group without writing data on cache to the volume on the external storage system. If your microcode version does not support this option, this option is ignored even if it is specified.

Examples

Deleting the external volume group #1-1.

```
# raidcom delete external_grp -external_grp_id 1-1
```

Deleting external volume group 1-1 whose connection is not disconnected without writing data on cache to the volume on the external storage system.

```
# raidcom delete external_grp -external_grp_id 1-1 -forcible
```

raidcom disconnect external_grp

Disconnects the connection to the volume on the external storage system. Only one external volume group is operated in each operation.

You can specify an LDEV defined for the external volume group or a device group to which the LDEV in the external volume group belongs instead of the external volume group.

- If an LDEV is specified, CCI finds the external volume groups to which the specified LDEV belongs and displays the result.
- If a device group is specified, CCI finds the external volume groups to which the LDEV in the specified device group belongs and displays the result.

If no LDEV exists in the external volume group, the command is rejected with EX_ENOOBJ.

Before finishing the write processing from the cache to the volume on the external storage system, the processing of **raidcom disconnect external_grp** command ends. Check the status (STS) using the **raidcom get path** command, and confirm the finishing of the write processing (destaging). (destaging). The following are the statuses (STS) after executing the **raidcom disconnect external_grp** command.

- NML: It means the previous status of receiving the request by the **raidcom disconnect external_grp** command.
- SYN: Write processing (destaging) is in process.
- DSC: Write processing (destaging) has finished.
- BLK: Write processing (destaging) has failed.

For details, see **raidcom get path**.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom disconnect external_grp {-external_grp_id <gno-sgno>
| -ldev_id <ldev#>} | -grp_opt <group option>
-device_grp_name <device group name> [<device name>]}
```

Options and parameters

-external_grp_id <gno-sgno>

Specifies the external volume group number (gno:1-16384, sgno:1-4096). For example:

- 52-11

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group.

Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all of the LDEVs belonging in the device group are operated.

Examples

Disconnect the connection to the external volume group #1-1.

```
# raidcom disconnect external_grp -external_grp_id 1-1
```

Disconnect the connection to the external volume group including the external volume (LDEV:200) to "blocked".

```
# raidcom disconnect external_grp -ldev_id 200
```

Disconnect the connection to the external volume group including the LDEV belonging to the device group: grp1.

```
# raidcom disconnect external_grp -grp_opt ldev -device_grp_name grp1
```

raidcom get external_grp

Displays the information of registered external volume groups.

Syntax

```
raidcom get external_grp [-external_grp_id <gno-sgno>]
```

Options and parameters**[-external_grp_id <gno-sgno>]**

Specifies the external volume group number (gno:1-16384, sgno:1-4096).

If this option is omitted, the list of the registered external volumes is displayed.

If this option is specified, the LDEV information defined for the specified external volume group is displayed. For example:

52-11

Example

Displaying external volume group information (The display might not be in ascending order).

```
#raidcom get external_grp
```

T	GROUP	Num_LDEV	U(%)	AV_CAP(GB)	R_LVL	E_TYPE	SL	CL	DRIVE_TYPE	M
E	1-1	0	0	100	-	OPEN-V	0	0	OPEN-V	N
E	1-2	0	0	30	-	OPEN-V	0	0	OPEN-V	N

Description of the each column in output example:**T**

Type of the volume group

R

Parity group, E: External volume group

GROUP

External volume group number

Num_LDEV

Number of LDEV assigned to the external volume group

U(%)

Usage rate of the external volume group

AV_CAP(GB)

Available capacity (free space) for the external volume group

R_LVL

RAID level of the parity group. As the external volume group is not relevant, "-" (bar) is displayed.

E_TYPE

Base emulation type of the external volume group

SL

SLPR to which the external volume group belongs (always displays 0)

CL

CLPR to which the external volume group belongs

DRIVE_TYPE

Product ID included in the SCSI Inquiry command response of the volume on the external storage system

M

Displays whether the allocation of external volume group to CLPR is changed.

- Y: The allocation of external volume group is being changed.
- N: The allocation of external volume group is not changed.
- - (hyphen): The information is not available (not supported).

Displaying the external volume group information by specifying the external volume group:

```
# raidcom get external_grp -external_grp_id 01-01
```

```
T GROUP P_NO LDEV# STS LOC_LBA SIZE_LBA Serial# SP
E 1-1 0 - NML 0x000000000000 0x000000003f00 64034 -
E 1-1 1 200 NML 0x000000003f00 0x000000010000 64034 R
E 1-1 2 201 REG 0x000000013f00 0x000000010000 64034 V
E 1-1 3 - DEL 0x000000023f00 0x0000f0000000 64034 -
```

Description of each column in the above output example:

T

Type of the volume group

R

Parity group, E: External volume group

GROUP

External volume group number.

P_NO

Partition number in this external volume group.

LDEV#

LDEV number assigned to this external volume group.

STS

Displays the following status.

- NML: an LDEV is installed.
- REG: an LDEV is being created.
- DEL: an LDEV is being deleted.

LOC_LBA

Starting point of LBA for this partition on this external volume group, in blocks (512 bytes).

SIZE_LBA:

Partition size of this external volume group, in blocks (512 bytes).

Serial#

Product serial number.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

SP

Displays whether the LDEV uses the expanded space of the parity group.

- V: LDEV uses the expanded space.
- R: LDEV does not use the expanded space.
- - (hyphen): LDEV is not mounted.

raidcom modify external_grp

Changes the attribute of external volume group (cache mode, cache inflow control mode, and MP blade ID setting).



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify external_grp -external_grp_id <gno-sgno>
    {-cache_mode {y|n|through|sync} | -cache_inflow {y|n}
    | -mp_blade_id <mp#> | -load_balance <mode>
    | -alua_switch {y|n}}
```

Options and parameters**-external_grp_id <gno-sgno>**

Specifies the external volume group number (gno:1-16384, sgno:1-4096). For example:

- 52-11

-cache_mode {y|n|through|sync}

Specifies the settings for the cache mode for an external volume group. If the relevant external volume group is split into the multiple LDEVs and these LDEVs are allocated to the pools, the cache mode settings might not be changed. For details, see the user guide for each software.

- y or yes: Write cache enabled (E)
- n or no: Write cache disabled (D)

- through: Cache through (T). Available only when the attribute of the relevant external volume group is NDM.
- sync: Write Sync mode (S). Available only when the attribute of the relevant external volume group is NDM.

-cache_inflow {y|n}

Specifies whether to use the Cache Inflow Control mode for an external volume group. If the external volume group consists of multiple LDEVs and these LDEVs are allocated to the pool, you cannot change this parameter.

- y or yes: Cache Inflow Control mode enabled (E)
- n or no: Cache Inflow Control mode disabled (D)

-mp_blade_id <mp#>

Specifies the MP blade ID (0-15). For example:

```
-mp_blade_id 2
```

-load_balance <mode>

Specifies load distribution mode of the alternate paths.

- normal: normal round robin
- extended: extended round robin
- disable: disables the alternate paths

-alua_switch {y|n}

Specifies whether the ALUA mode is used.

- y or yes: The ALUA mode is used.
- n or no: The ALUA mode is not used.

Examples

Turning the cache mode of the external volume group #01-01 ON.

```
# raidcom modify external_grp -external_grp_id 01-01 -cache_mode y
```

Enabling the Cache Inflow Control mode of the external volume group #01-01 ON.

```
# raidcom modify external_grp -external_grp_id 01-01 -cache_inflow y
```

Changing the MP blade ID of the external volume group #01-01 to "3".

```
# raidcom modify external_grp -external_grp_id 01-01 -mp_blade_id 3
```

raidcom discover external_storage

Searches the port information on the external storage system connected to the External (ELUN) port.

If you execute the `raidcom discover external_storage` command while the `raidcom discover lun` command or another `raidcom discover external_storage` command is being executed, the external storage system might not be displayed. If this happens, confirm that the storage system in which the command is being executed and the external storage system are connected correctly, and the LU of the external storage system is configured correctly.

If both the storage system and the external storage system have no problem, wait a while, and then execute one command at a time for a storage system.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Note:

If the local storage system port is a Fiber Channel Bidirectional port, it is recommended that you always execute the command by specifying the `-safety_check enable` option unless instructed by the user guide.

If you do not specify the `-safety_check enable` option, the following problems might occur. For the details, see the Hitachi Universal Volume Manager User Guide.

- If you use a non-Hitachi external storage system:
The I/O path from the external storage system to the local storage system that uses the route between the specified local storage system port and the searched external storage system port might be disconnected.
- If you use the Hitachi external storage system:
The external or remote path connection that uses the route between the specified local storage system port and the searched external storage system port might be temporarily disconnected (the external or remote path will be reconnected immediately and no blockage occurs).
- If you run multiple instances of the `raidcom discover external_storage` command, `raidcom discover lun` command, `raidcom check external_iscsi_name` command, or `raidcom discover external_iscsi_name` command at the same time, an EX_CMDIOE error might occur. If an EX_CMDIOE error occurs, wait for a while, and then run the commands again one by one. For details, see the *Hitachi Universal Volume Manager User Guide*.

Syntax

```
raidcom discover external_storage -port <port#> [-safety_check enable] [-
discovery_external_wnn <wnn strings>]
```

Options and parameters

-port <port#>

Specifies the port number. Specifies the port number where the attribute is "ELUN(External)". For example:

- CL1-A

Displays an external port.

[-safety_check enable]

Specify the Fiber Channel bidirectional port for the `-port` option. If you specify any other port, this option is ignored.

When this option is specified, processing that might cause a disconnection between the bidirectional port specified by the `-port` option and the Fiber Channel port of the external storage system is suppressed. When processing is suppressed, the following information is displayed for Serial#, VENDOR_ID, and PRODUCT_ID.

- Serial#: "-" (hyphen)
- VENDOR_ID: "Unknown"
- PRODUCT_ID: "Unknown"

[-discovery_external_wwn <wwn strings>]

If you specify this option, information only for the specified external storage system port is displayed. Specify the external storage system WWN (Hexadecimal) in 8 bytes. You can separate 4 bytes using "," (comma).

(Example)

- 21000e08b0256f8
- 21000e0,8b0256f8

Use this option in one of the following methods:

1. Specify the `-discovery_external_wwn <wwn strings>` option without the `-safety_check enable` option.

Use this method to check the Serial #, VENDOR_ID, and PRODUCT_ID information of the port specified by the `-port` option as a bidirectional port and the external storage system port.

When using this method, make sure that the route is not applicable to the precautions described in the Hitachi Universal Volume Manager User Guide before executing the command. If applicable, this method might temporarily disconnect the external or remote path that uses the route between the external storage system port specified by `-discovery_external_wwn` and the bidirectional port specified by the `-port` option. If the external storage system is made by a third party, the I/O path from the external storage system to the local storage system might be disconnected.

- Specify the `-discovery_external_wnn <wnn strings>` option with the `-safety_check enable` option.

Use this method to check whether the external storage port specified by the `-discovery_external_wnn` option is connected when the route is applicable to the route mentioned above.

Example

Displaying the external storage system ports from the port: CL1-A.

```
# raidcom discover external_storage -port CL1-A
```

```
PORT WWN PM USED Serial# VENDOR_ID PRODUCT_ID
CL1-A 50060e8005fa0f36 M YES 60010 HITACHI VSP
CL1-A 50060e8005fa0f38 M YES 60010 HITACHI VSP
```

Displaying the external storage system ports from the port: CL1-A by specifying the `-safety_check enable` option.

```
# raidcom discover external_storage -port CL1-A -safety_check enable
PORT WWN PM USED Serial# VENDOR_ID PRODUCT_ID
CL1-A 50060e8005fa0f36 M No - Unknown Unknown
CL1-A 50060e8005fa0f38 M YES 60010 HITACHI VSP
```

Checking the Serial#, VENDOR_ID, and PRODUCT_ID of the external storage system port 50060e8005fa0f36 from the port: CL1-A.

```
# raidcom discover external_storage -port CL1-A -discovery_external_wnn
50060e8005fa0f36
PORT WWN PM USED Serial# VENDOR_ID PRODUCT_ID
CL1-A 50060e8005fa0f36 M YES 60010 HITACHI VSP
```

Checking the connection with the external storage system port 50060e8005fa0f36 from the port: CL1-A.

```
# raidcom discover external_storage -port CL1-A -safety_check enable -
discovery_external_wnn 50060e8005fa0f36
PORT WWN PM USED Serial# VENDOR_ID PRODUCT_ID
CL1-A 50060e8005fa0f36 M No - Unknown Unknown
```

Description of each column in output example:

PORT

Displays the external port number of the storage system.

WWN

Displays the WWN which can be referred to from the port.

PM:

Displays the path mode for external path.

- M: Multi
- S: Single
- A: APLB

USED

Displays whether this target WWN is used.

- YES: Used
- NO: Not used

Serial#

Product serial number of the external storage system.

If the `-safety_check enable` option is specified, "-" (hyphen) is displayed when processing that might cause disconnection is suppressed.

VENDOR_ID

Displays the vendor name of the external storage system. "OTHER" is displayed if an unsupported external storage system is connected.

If the `-safety_check enable` option is specified, "Unknown" is displayed when processing that might cause disconnection is suppressed.

PRODUCT_ID

Displays the system name of the external storage system. "OTHER" is displayed if an unsupported external storage system is connected.

If the `-safety_check enable` option is specified, "Unknown" is displayed when processing that might cause disconnection is suppressed.

raidcom add host_grp

Creates a host group or an iSCSI target on the specified port.

- If the port type is FIBRE or FCoE (Fibre Channel over Ethernet), this command creates a host group.
- If the port type is iSCSI, this command creates the iSCSI target (equivalent of a host group) and the iSCSI name.

If the specified port does not exist, the command is rejected with EX_ENOOBJ.

If the specified port/host group already exists, the existing host group name is changed to the specified host group name.

If the specified port, iSCSI target, and iSCSI name already exist, the existing information is changed to the specified information.

The specified host group name must be unique in a port.

Syntax

```
raidcom add host_grp -port <port#> -host_grp_name <host group name> [-iscsi_name
<target iscsi name>]
```

Options and parameters

-port <port#>

Specifies the port number and the host group ID. The host group ID can be omitted. For example:

- CL1-A-g
A is the port number, and g is the host group ID (0-254).
- CL1-A

If the host group ID is not specified, an unused host group ID is automatically assigned. However, if multiple commands without host group ID are executed simultaneously for the same port, the same host group ID might be assigned to multiple newly created host groups.

To avoid this behavior, use the **raidcom lock resource** command in advance to lock the resource group to which the unused host group ID belongs. If you lock the resource group, commands executed by other users cannot take out unused host group IDs from the resource group. If you set multiple host groups for the same port, execute the next **raidcom add host_grp** command after the execution of the current **raidcom add host_grp** command is complete.

After the execution of all the **raidcom add host_grp** commands are complete, use the **raidcom unlock resource** command to unlock the resource group.

-host_grp_name <host group name>

If the port type is FIBRE or FCoE, specifies the HOST group name. Up to 64 characters can be set by CCI. If more than 64 characters are set, commands that specify host group name by CCI cannot be executed.

If the port type is iSCSI, specifies the iSCSI target name. Up to 32 English one-byte characters can be set by CCI.

-iscsi_name <target iscsi name>

Specifies the iSCSI name by using either one of two formats, iqn or eui.

- iqn format: "iqn." followed by up to 219 English one-byte characters. The permitted characters are:
 - Alphabet
 - Number (0-9)
 - Period (.)
 - Hyphen (-)
 - Colon (:)
- eui format: "eui" (case sensitive) followed by a 16-digit hexadecimal value.

If this option is omitted, the default settings are specified. The default value depends on the serial number, the port number, or the target ID.



Note:

- You cannot register multiple iSCSI names to the same port if the only difference among them is the case (lowercase or uppercase). For example, if `iqn.win2k8.example.of.iqn.form` is already registered as an iSCSI name, you cannot register `iqn.win2k8.example.of.iqn.FORM` as a different iSCSI name.
- If you input an iSCSI name in iqn format, use lowercase letters only.
- If you input an iSCSI name in eui format, use lowercase letters for "eui.", and use uppercase letters for the following hexadecimal value.

Examples

Creating a host group ID: 3, the host group name: a host group of Win_export, to the port: CL4-E.

```
# raidcom add host_grp -port CL4-E-3 -host_grp_name Win_export
```

Creating a host group ID: allocated automatically, the host group name: a host group of Win_export, to the port: CL4-E.

```
# raidcom add host_grp -port CL4-E -host_grp_name Win_export
```

Creating an iSCSI name: `iqn.2014-04.jp.co.hitachi:xxx.h70.i.62510.1A.FF`, iSCSI target name: Target00, to the port: CL4-E.

```
# raidcom add host_grp -port CL4-E -host_grp_name Target00
-iscsi_name iqn.2014-04.jp.co.hitachi:xxx.h70.i.62510.1A.FF
```

raidcom delete host_grp

Deletes the specified host group or iSCSI target. This command also deletes the WWN/iQN or LUNs settings of the host registered for the host group or iSCSI target.

If the port type of the port with the specified host group is FIBRE or FCoE (Fibre Channel over Ethernet), this command deletes the host group, and WWN and LUNs settings of the host registered to the host group.

If the port type of the port with the specified iSCSI target is iSCSI, this command deletes the iSCSI target, and WWN and LUNs settings of the host (initiator) registered to the iSCSI target. However, if the host group ID of the host group or the target ID of the iSCSI target is 0, this command changes the settings back to the default. If the LUSE configuration is defined, this command releases the LUSE.

When the specified host group is associated with an SPM group, the following information is deleted after the host group is deleted:

- Association between the SPM group and the host group
- SPM settings of the WWN registered in the host group
- Registration of the WWN in the SPM group

If no WWN is registered in the SPM group, the SPM group itself is also deleted.

If the specified port does not exist, the command is rejected with EX_ENOOBJ.

If a specific host mode option is set for the specified host group, the command might be rejected with EX_CMDRJE. For details about specific host mode options, see the *Provisioning Guide* for your storage system.

Syntax

```
raidcom delete host_grp -port <port#> [<host group name>]
```

Options and parameters

-port <port#> [<host group name>]

Specifies the port number, host group ID, and host group name (iSCSI target name for iSCSI). If the host group name or the iSCSI target name is more than 64 characters, use the host group ID or the iSCSI target ID. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86
- CL1-A Target00

Examples

Deleting the host group of port CL4-E, host group #7.

```
# raidcom delete host_grp -port CL4-E-7
```

Deleting the host group of port CL4-E, host group name: host group of Win_export.

```
# raidcom delete host_grp -port CL4-E Win_export
```

Deleting the target of port CL4-E, iSCSI target name: Target01.

```
# raidcom delete host_grp -port CL4-E Target01
```

raidcom get host_grp

Displays the information about all host groups or an iSCSI target that are defined on the specified port.

Syntax

```
raidcom get host_grp {-port <port#> [<host group name>]| -allports} [-key <keyword>]
```

Options and parameters

-port <port#> [<host group name>]

Specifies the port number, host group ID, and host group name (iSCSI target name for iSCSI). If the host group name or the iSCSI target name is more than 64 characters, use the host group ID or the iSCSI target ID.

For example:

- CLI-A
- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86
- CL1-A Target00

If you specified the above settings, the information about all host groups that are defined on the specified port is displayed regardless of the examples.

-allports

Displays the information about the host groups or the iSCSI target that is set to all ports.

[-key <keyword>]

Specifies the information to be displayed. Shows display keywords can be specified in <keyword> as follows:

- host_grp: Specifies the host group IDs that are not installed.
- server: Specifies the information about the server managed by Storage Advisor Embedded or VSP One Block Administrator on which the host group or the iSCSI target is registered.

This parameter can be specified only for VSP One B20, VSP E series and VSP G130, G/F350, G/F370, G/F700, G/F900.
- detail: Specifies the option to display the resource group ID of the host groups or iSCSI targets. -key detail and -allports cannot be specified at the same time. When the -key detail is specified, host group information about both implemented and unimplemented host groups is displayed.

Example

Displaying the host group information set for the port: CL4-E (in case of FIBRE, FCoE (Fibre Channel over Ethernet), HNASS, or HNASU).

```
# raidcom get host_grp -port CL4-E
```

PORT	GID	GROUP_NAME	Serial#	HMD	HMO_BITS
CL4-E	0	Linux_x86	63528	LINUX/IRIX	2 13

CL4-E	1	Solaris	63528	SOLARIS	2 22
CL4-E	2	HP-UX	63528	HP-UX	40
CL4-E	3	Win_export	63528	WIN_EX	7
CL4-E	5	DEC	63528	TRU64	14
CL4-E	6	OpenVMS	63528	OVMS	
CL4-E	254	RMTEST	63528	LINUX	7

```
# raidcom get host_grp -port CL4-E -key host_grp
```

PORT	GID	GROUP_NAME	Serial#	HMD	HMO_BITS
CL4-E	0	Linux_x86	63528	LINUX/IRIX	2 13
CL4-E	1	Solaris	63528	SOLARIS	2 22
CL4-E	2	HP-UX	63528	HP-UX	40
CL4-E	3	Win_export	63528	WIN_EX	7
CL4-E	5	DEC	63528	TRU64	14
CL4-E	6	OpenVMS	63528	OVMS	
CL4-E	10	-	63528	-	
CL4-E	11	-	63528	-	
CL4-E	12	-	63528	-	
CL4-E	13	-	63528	-	
CL4-E	254	RMTEST	63528	LINUX	7

Displaying the iSCSI target information set for the port: CL4-E (in case of iSCSI)

```
# raidcom get host_grp -port CL4-E
```

PORT	GID	GROUP_NAME	IQN	AMD	D	Serial#	HMD	HMO_BITS
CL4-E	0	Linux_x86	iqn.z1...	CHAP	S	63528	LINUX/IRIX	2 13
CL4-E	1	Solaris	iqn.z2...	CHAP	S	63528	SOLARIS	2 22
CL4-E	2	HP-UX	iqn.z3...	CHAP	S	63528	HP-UX	40

Displaying the information about the host groups or the iSCSI targets that are set for the port: CL4-E (in case of iSCSI and Fibre Channel)

```
# raidcom get host_grp -port CL4-E -key server
```

PORT	GID	GROUP_NAME	Serial#	SRVID	SRV_NAME
CL4-E	0	Linux_x86	63528	N	-
CL4-E	1	Solaris	63528	1	"server01"
CL4-E	2	HP-UX	63528	N	-
CL4-E	3	Win_export	63528	N	-
CL4-E	5	DEC	63528	N	-

CL4-E	6	OpenVMS	63528	N	-
CL4-E	254	RMTEST	63528	N	-

Displaying the resource group information about the host groups or the iSCSI targets that are set for the port: CL4-E (in case of iSCSI or Fibre Channel)

```
# raidcom get host_grp -port CL4-E -key detail
PORT GID RGID GROUP_NAME Serial# HMD HMO_BITS
CL4-E 0 0 "Linux_x86" 63528 LINUX/IRIX 2:13
CL4-E 1 0 "Solaris" 63528 SOLARIS 2:22
CL4-E 2 1 "HP-UX" 63528 HP-UX 40
CL4-E 3 0 "Win_export" 63528 WIN_EX 7
CL4-E 5 0 "DEC" 63528 TRU64 14
CL4-E 6 0 "OpenVMS" 63528 OVMS -
CL4-E 10 0 - 63528 - -
CL4-E 11 0 - 63528 - -
CL4-E 12 0 - 63528 - -
CL4-E 13 0 - 63528 - -
CL4-E 254 0 "RMTEST" 63528 LINUX -
```

Displaying the information about the host groups or the iSCSI targets that are set to all ports

```
# raidcom get host_grp -allports
```

PORT	GID	GROUP_NAME	Serial#	HMD
CL1-A	0	Linux_x86	63528	LINUX/IRIX
CL1-A	1	Solaris	63528	SOLARIS
CL1-A	2	HP-UX	63528	HP-UX
CL1-A	3	Win_export	63528	WIN_EX
CL1-A	5	DEC	63528	TRU64
CL1-A	6	OpenVMS	63528	OVMS
CL1-A	254	RMTEST	63528	LINUX
CL1-B	0	Linux_x86	63528	LINUX/IRIX
CL1-B	1	Solaris	63528	SOLARIS
CL1-B	2	HP-UX	63528	HP-UX
CL1-B	3	Win_export	63528	WIN_EX
CL1-B	5	DEC	63528	TRU64
CL1-B	6	OpenVMS	63528	OVMS
CL1-B	254	RMTEST	63528	LINUX

```
# raidcom get host_grp -allports -key host_grp
```

PORT	GID	GROUP_NAME	Serial#	HMD
CL1-A	0	Linux_x86	63528	LINUX/IRIX
CL1-A	1	Solaris	63528	SOLARIS

```

CL1-A 2      HP-UX      63528  HP-UX
CL1-A 3      Win_export 63528  WIN_EX
CL1-A 5      DEC       63528  TRU64
CL1-A 6      OpenVMS   63528  OVMS
CL1-A 10     -         63528  -
CL1-A 11     -         63528  -
CL1-A 12     -         63528  -
CL1-A 13     -         63528  -
CL1-A 254    RMTEST    63528  LINUX
CL1-B 0      Linux_x86  63528  LINUX/IRIX
CL1-B 1      Solaris   63528  SOLARIS
CL1-B 2      HP-UX     63528  HP-UX
CL1-B 3      Win_export 63528  WIN_EX
CL1-B 5      DEC       63528  TRU64
CL1-B 6      OpenVMS   63528  OVMS
CL1-B 10     -         63528  -
CL1-B 11     -         63528  -
CL1-B 12     -         63528  -
CL1-B 13     -         63528  -
CL1-B 254    RMTEST    63528  LINUX

```

```
# raidcom get host_grp -allports -key server
```

PORT	GID	GROUP_NAME	Serial#	SRVID	SRV_NAME
CL1-A	0	Linux_x86	63528	N	-
CL1-A	1	Solaris	63528	1	"server01"
CL1-A	2	HP-UX	63528	N	-
CL1-A	3	Win_export	63528	N	-
CL1-A	5	DEC	63528	N	-
CL1-A	6	OpenVMS	63528	N	-
CL1-A	254	RMTEST	63528	N	-
CL1-B	0	Linux_x86	63528	N	-
CL1-B	1	Solaris	63528	N	-
CL1-B	2	HP-UX	63528	N	-
CL1-B	3	Win_export	63528	N	-
CL1-B	5	DEC	63528	N	-
-					

Displaying only the host group IDs that are allocated to available resource groups for users.

Description of each column in output example:

PORT

Displays the port number.

GID

Displays the host group ID of a port.

GROUP_NAME

Displays the host group name of a port or target alias.

For unimplemented host groups or target aliases, a hyphen (-) is displayed. If the `-key` option is specified, a double quotation (") is added to the beginning and end of the name of the implemented host groups or target aliases.

IQN

Displays the iSCSI Qualified Name of the port.

AMD

Displays the authentication mode of the iSCSI target.

- CHAP: CHAP authentication is enabled.
- NONE: Authentication is disabled.
- BOTH: Both CHAP authentication and connection by no-authentication are enabled.

D

Displays the direction of the authentication mode of the iSCSI target.

- S: Unidirectional (An initiator is recognized from the target side)
- D: Bidirectional (An initiator is recognized from the target side, and a target is recognized from the initiator)

Serial#

Product serial number.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

HMD

Displays the following HOST MODE for the host adapter setting on host group.

- HP-UX, SOLARIS, AIX, WIN, LINUX/IRIX, TRU64, DYNIX, OVMS, NETWARE, HI-UX
- VMWARE, HP-XP, VMWARE_EX, WIN_EX, UVM

HMO_BITS

Displays the host mode options of the host group. For details, see the *Provisioning Guide* for the storage system.

If the host mode options are two or more, when case the `-key detail` option is specified, the host mode options which are set are displayed separately by a colon (:). If the `-key detail` option is specified, when there is no host mode option that is set, a hyphen (-) is displayed.

SRVID

Displays the server ID for the Storage Advisor Embedded or VSP One Block Administrator on which the host groups or the iSCSI targets are registered in decimal.

Displays "N" if they are not registered on the server. Displays a hyphen (-) if the microcode version does not support the server ID indication.

SRV_NAME

Displays the nickname of the server identified by SRVID. A double quotation (") is added at the beginning and end of the server nickname. A hyphen (-) is displayed if SRVID is "N" or hyphen (-).

RSGID

Displays the resource group ID of the host groups or the iSCSI targets.

raidcom modify host_grp

Sets a host mode to the host group or an iSCSI target on the specified port.

If the specified host group does not exist, the command is ignored.

In case of iSCSI, set the CHAP authentication (enable/disable, or unidirectional/bidirectional).

Syntax

```
raidcom modify host_grp -port <port#> [<host group name>]
    -host_mode <host mode> [-host_mode_opt <host mode option>... |
    -set_host_mode_opt <host mode option>... | -reset_host_mode_opt]
    [-authmethod {CHAP|NONE|BOTH}] [-mutual {enable|disable}]
```

Options and parameters

-port <port#>[<host group name>]

Specifies the port number, host group ID, and host group name (iSCSI target name for iSCSI). If the host group name or the iSCSI target name is more than 64 characters, use the host group ID or the iSCSI target name.

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86
- CL1-A Target00

-host_mode <host mode>

Specifies the host mode using the following strings and values. For the supported host modes, see the *Provisioning Guide* for the storage system.

- LINUX or IRIX (0x00)
- VMWARE (0x01)
- HP-UX (0x03)
- OVMS (0x05)
- TRU64 (0x07)

- SOLARIS (0x09)
- NETWARE (0x0a)
- WIN (0x0c)
- AIX (0x0f)
- VMWARE_EX (0x21)
- WIN_EX (0x2c)
- UVM (0x4c)

[-host_mode_opt <host mode option>...]

This parameter remains for the compatibility with the old version. Use the `-set_host_mode_opt` option and `-reset_host_mode_opt` option.

[-set_host_mode_opt <host mode option> ...]

Specifies the host mode option. The other host mode options which you do not specify is cleared. For details about the host mode option, see the *Provisioning Guide* for the storage system.

[-reset_host_mode_opt]

Resets all host mode options. For details about the host mode option, see the *Provisioning Guide* for the storage system.

[-authmethod {CHAP|NONE|BOTH}]

Specifies the CHAP authentication mode. Even if the CHAP user name has not been specified yet, the CHAP authentication mode can be specified. This option must be specified in parallel with specifying the host mode.

- CHAP: CHAP authentication is enabled.
- NONE: Authentication is disabled.
- BOTH: Both CHAP authentication and connection by no-authentication are enabled.

[-mutual {enable|disable}]

Specifies the CHAP authentication mode: unidirectional authentication, or bidirectional authentication. Even if the CHAP authentication mode is specified to NONE, CHAP authentication can be specified (when the authentication mode will be changed to CHAP, or BOTH, the specified mode will be enabled). This option must be specified in parallel with specifying the host mode.

- enable: specifies bidirectional CHAP authentication (an initiator is recognized from the target side, and a target is recognized from the initiator).
- disable: specifies unidirectional CHAP authentication (an initiator is recognized from the target).

Examples

Setting the host mode: HP-UX for the port: CL4-E, the host group: #2.

```
# raidcom modify host_grp -port CL4-E-2 -host_mode HP-UX
```

Setting the host mode: HP-UX and the host mode option: 12, 60 for the port: CL4-E, the host group: #2.

```
# raidcom modify host_grp -port CL4-E-2 -host_mode HP-UX -set_host_mode_opt 12 60
```

Clearing the host mode options of the host mode: HP-UX of the port: CL4-E, the host group: #2.

```
# raidcom modify host_grp -port CL4-E-2 -host_mode HP-UX -reset_host_mode_opt
```

Setting the host mode: HP-UX and the bidirectional CHAP authentication enabled for the port: CL4-E, the host group: #2.

```
# raidcom modify host_grp -port CL4-E-2 -host_mode HP-UX -authmethod CHAP -mutual
enable
```

raidcom add host_nqn

Supported storage systems:

- VSP 5000 series
- VSP E1090

Registers a host NQN on the NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom add host_nqn {-nvm_subsystem_id <nvm subsystem id>} -host_nqn
<host_nqn> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>
Specifies the NVM subsystem ID.

-host_nqn <host_nqn>
Specifies the host NQN.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Adding a host NQN: `nqn.xxx` to the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom add host_nqn -nvm_subsystem_id 1 -host_nqn nqn.xxx -request_id
auto
```

raidcom get host_nqn

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays the host NQN information.

Syntax

```
raidcom get host_nqn {-nvm_subsystem_id <nvm subsystem id>} [-key
<keyword>]
```

Options and parameters**-nvm_subsystem_id <nvm subsystem id>**

Specifies the NVM subsystem ID.

[-key <keyword>]

Specifies a display keyword:

- `opt`: Displays the detailed information about the host NQN.

Examples

Displaying the information about a host NQN registered on the NVM subsystem ID 1.

```
# raidcom get host_nqn -nvm_subsystem_id 1
NVMSS_ID NVMSS_NAME HOST_NQN
1 my_nvm_subsystem nqn.xxx
1 my_nvm_subsystem nqn.yyy
```

Displaying the detailed information about a host NQN registered on the NVM subsystem ID 1.

```
# raidcom get host_nqn -nvm_subsystem_id 1 -key opt
NVMSS_ID NVMSS_NAME NICK_NAME
HOST_NQN
1 my_nvm_subsystem my_host
```

```
nqn.xxx
1 my_nvm_subsystem host2 nqn.yyy
```

Description of each column in output example:

NVMSS_ID

Displays the NVM subsystem ID.

NVMSS_NAME

Displays the NVM subsystem name.

HOST_NQN

Displays the host NQN.

NICK_NAME

Displays the nickname of the host NQN. If no nicknames are specified, a hyphen (-) is displayed.

raidcom modify host_nqn

Supported storage systems:

- VSP 5000 series
- VSP E1090

Specifies a host NQN nickname.

If the specified host NQN is registered on multiple NVM subsystems, the setting is applied to all NVM subsystems.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom modify host_nqn -host_nqn <host_nqn> -nvm_subsystem_id <nvm
subsystem id> {-set_host_name <host nick name> | -reset_host_name} -
request_id auto
```

Options and parameters**-nvm_subsystem_id <nvm subsystem id>**

Specifies the NVM subsystem ID.

-host_nqn <host_nqn>

Specifies the host NQN.

-set_host_name <host nick name>

Specifies the host NQN nickname.

-reset_host_name

Deletes the host NQN nickname.

-request_id auto<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).**Example**

Changing the host NQN nickname from nqn.xxx to my_host.

```
# raidcom modify host_nqn -host_nqn nqn.xxx -nvm_subsystem_id 1 -
set_host_name my_host -request_id auto
```

raidcom delete host_nqn

Supported storage systems:

- VSP 5000 series
- VSP E1090

Deletes a host NQN from the NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom delete host_nqn {-nvm_subsystem_id <nvm subsystem id>} -host_nqn
<host_nqn> -request_id auto
```

Options and parameters**-nvm_subsystem_id <nvm subsystem id>**

Specifies the NVM subsystem ID.

-host_nqn <host_nqn>

Specifies the host NQN.

-request_id auto<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).**Example**

Deleting the host NQN: nqn.xxx from the NVM subsystem ID 1.

```
# raidcom delete host_nqn -nvm_subsystem_id 1 -host_nqn nqn.xxx -
request_id auto
```

raidcom add chap_user

This command sets the CHAP user name for the specified iSCSI target. Also this command registers the CHAP user name of the host on the initiator set in the specified iSCSI target. If the specified CHAP user name of the host on the initiator already exists, the registration is ignored.

Syntax

```
raidcom add chap_user -port <port#> [<host group name>]
    {-target_chap_user <user name>|-initiator_chap_user
    <user name>}
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

-target_chap_user <user name>

Specifies the CHAP user name for the iSCSI target. You can specify up to 223 characters. The maximum number of CHAP user names of the iSCSI target is 1 for each iSCSI target.

For example: storage01

-initiator_chap_user <user name>

Specifies the CHAP user name on the initiator that is set as the iSCSI target. You can specify up to 223 characters.

For example: Linux-abc

Examples

To set the CHAP user name "storage01" to the iSCSI target whose port is CL4-E and the host group ID is 0:

```
# raidcom add chap_user -port CL4-E-0 -target_chap_user storage01
```

To set the CHAP user name "storage02" to the host whose iSCSI target name is target00 and port is CL4-E:

```
# raidcom add chap_user -port CL4-E Target00 -target_chap_user storage02
```

To register the CHAP user name "Linux-abc" to the host on the initiator whose iSCSI target name is target00 and port is CL4-E:

```
# raidcom add chap_user -port CL4-E Target00 -initiator_chap_user Linux-abc
```

raidcom delete chap_user

This command deletes the CHAP user name from the specified iSCSI target. Also this command deletes the CHAP user name from the host on the initiator set in the specified iSCSI target. If the specified CHAP user name does not exist, the command is ignored.

Syntax

```
raidcom delete chap_user -port <port#> [<host group name>]
    {-target_chap_user <user name>|-initiator_chap_user
    <user name>}
```

Options and parameters

-port <port#>[<host group name>]

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

-target_chap_user <user name>

Specifies the CHAP user name of the iSCSI target.

For example: storage01

-initiator_chap_user <user name>

Specifies the CHAP user name of the host on the initiator.

For example: Linux-abc

Examples

To delete the CHAP user name "storage01" from the iSCSI target whose port is CL4-E and the host group ID is 0:

```
# raidcom delete chap_user -port CL4-E-0 -target_chap_user storage01
```

To delete the CHAP user name "storage02" from the host whose iSCSI target name is target00 and port is CL4-E:

```
# raidcom delete chap_user -port CL4-E Target00 -target_chap_user storage02
```

To delete the CHAP user name "Linux-abc" from the host on the initiator whose iSCSI target name is target00 and port is CL4-E:

```
# raidcom delete chap_user -port CL4-E Target00 -initiator_chap_user Linux-abc
```

raidcom set chap_user

This command sets the password, called "secret", for the specified CHAP user. Both the host on the initiator and the host on the target are set by this command. To avoid a secret is given as an argument directly, the proper prompt is displayed to enter the secret. If the secret already exists for the specified CHAP user, the secret is overwritten.

Syntax

```
raidcom set chap_user -port <port#> [<host group name>] {-target_chap_user <user name> -secret | -initiator_chap_user <user name> -secret}
```

Options and parameters

-port <port#>[<host group name>]

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

-target_chap_user <user name>

Specifies the CHAP user name of the iSCSI target.

For example: storage01

-initiator_chap_user <user name>

Specifies the CHAP user name of the host on the initiator.

For example: Linux-abc

-secret

Displays the prompt for entering a secret.

Input characters as "secret", within the range of 12 characters to 32 characters, or an error occurs.

Examples

To set the "iSCSI-secret" for the "secret" to the user whose CHAP user name is storage01, port is CL4-E, and the target ID of the iSCSI target is 0:

```
# raidcom set chap_user -port CL4-E-0 -target_chap_user storage01 -secret
Enter Secret :
```

(Enter "iSCSI-secret" after the "Enter secret: " above. The entered characters are not displayed on the screen.)

To register the secret: Linux-secret for the CHAP user name "Linux-abc" from the host on the initiator whose iSCSI target name is target00 and port is CL4-E:

```
# raidcom set chap_user -port CL4-E Target00 -initiator_chap_user Linux-abc -secret
Enter Secret :
```

(Enter "Linux-secret" after the "Enter secret: " above. The entered characters are not displayed on the screen.)

raidcom reset chap_user

This command removes the secret from the specified CHAP user. Both the host on the initiator and the host on the target are set by this command.

Syntax

```
raidcom reset chap_user -port <port#> [<host group name>]
    {-target_chap_user <user name>|-initiator_chap_user
    <user name>}
```

Options and parameters

-port <port#>[<host group name>]

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

-target_chap_user <user name>

Specifies the CHAP user name of the iSCSI target.

For example: storage01

-initiator_chap_user <user name>

Specifies the CHAP user name of the host on the initiator.

For example: Linux-abc

Examples

To delete the secret for the CHAP user name "storage01" from the iSCSI target whose port is CL4-E and the target ID is 0:

```
# raidcom reset chap_user -port CL4-E-0 -target_chap_user storage01
```

To delete the secret for the CHAP user name "Linux-abc" from the initiator host whose port is CL4-E:

```
# raidcom reset chap_user -port CL4-E Target00 -initiator_chap_user Linux-abc
```

raidcom get chap_user

This command indicates the CHAP user name of the iSCSI target on the specified port and the CHAP user name of the host bus adapter on the initiator that is registered in the iSCSI target.

Syntax

```
raidcom get chap_user -port <port#> [<host group name>]
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

Example

To display the CHAP user name whose port is CL4-E and whose host group ID is 0:

```
# raidcom get chap_user -port CL4-E-0
```

PORT	GID	GROUP_NAME	CHAP_user	Serial#	WAY	Sec
CL4-E	0	Linux_x86	raidmanager	63528	INI	*
CL4-E	0	Linux_x86	raidmanager1	63528	INI	*
CL4-E	0	Linux_x86	raidmanager2	63528	INI	*
CL4-E	0	Linux_x86	oracle	63528	TAR	*

Description of each column in the output example:**PORT**

Displays the port.

GID

Displays the host group ID of the port.

GROUP_NAME

Displays the iSCSI target name of the port.

CHAP_user

It indicates the CHAP user name of the iSCSI target and the CHAP user name of the host bath adapter which is registered in the iSCSI target.

Serial#

Displays the Seq#.

WAY

Indicates whether the CHAP user name on the iSCSI target or the CHAP user name on the host bus adapter (initiator).

- TAR: iSCSI target side
- INI: Host bus adapter (initiator) side

Sec

An asterisk (*) is always displayed.

raidcom add hba_wwn

Registers the WWN of the host adapter to the host group on the specified port to add a host.

If the specified WWN already exists, this command is ignored. (VSP and VSP G1x00 and VSP F1500 only) If the specified host group is associated with a n SPM group, the WWN of the host adapter is registered. After that, the WWN is registered in the SPM group, and then the SPM information is set for the WWN.

Syntax

```
raidcom add hba_wwn -port <port#> [<host group name>]
                        -hba_wwn <WWN strings>
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, host group ID, and host group name. If the host group name is more than 64 characters, use the host group ID. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86

-hba_wwn <WWN strings>

Specifies the WWN (16-digit hexadecimal value) of the host adapter. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

Example

Setting the WWN of host adapter: 210000e0,8b0256f8 to the port: CL4-E, the host group #0.

```
# raidcom add hba_wwn -port CL4-E-0 -hba_wwn 210000e0,8b0256f8
```

raidcom delete hba_wwn

Deletes the specified host (WWN) from the host group.

If the specified WWN does not exist, this command is ignored. (VSP and VSP G1x00 and VSP F1500 only) If the specified host group is associated with a n SPM group, the WWN of the host adapter is deleted, and then the SPM settings of the WWN and the registration of the WWN in the SPM group are deleted.

Syntax

```
raidcom delete hba_wwn -port <port#> [<host group name>]
                        -hba_wwn <WWN strings>
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, host group ID, and host group name. If the host group name is more than 64 characters, use the host group ID. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86

-hba_wwn <WWN strings>

Specifies the WWN (16-digit hexadecimal value) of the host adapter. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

Example

Deleting the WWN of host adapter: 210000e0,8b039800 set for the port: CL4-E, the host group #0.

```
# raidcom delete hba_wwn -port CL4-E-0 -hba_wwn 210000e0,8b039800
```

raidcom get hba_wwn

Displays the WWN of the HBA registered to the host group.

Syntax

```
raidcom get hba_wwn -port <port#> [<host group name>]
```

Options and parameters**-port <port#>[<host group name>]**

Specifies a port number, a host group ID, and a host group name. It cannot be specified when more than 64 characters is set for the host group name. Use the host group ID. If setting the host group ID or the host group name is omitted, the information about host group ID 0 is displayed. For example,

- CLI-A
- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86

Example

Displaying the WWN of host adapter set for the port: CL4-E, the host group 0.

```
# raidcom get hba_wwn -port CL4-E-0
```

PORT	GID	GROUP_NAME	HWWN	Serial#	NICK_NAME
CL4-E	0	Linux_x86	210000e08b0256f8	63528	ORA_NODE0_CTL_0
CL4-E	0	Linux_x86	210000e08b039c15	63528	ORA_NODE1_CTL_0

Description of each column in output example:**PORT**

Displays the port number.

GID

Displays the host group ID of a port.

GROUP_NAME

Displays the host group name of a port.

HWWN

Displays the WWN of registered host adapter.

Serial#

Displays the storage system serial number:

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

NICK_NAME

Displays the Nick Name of the WWN of host adapter.

raidcom reset hba_wwn

Deletes the nickname from the specified WWN on the specified port.

If there is no specified port, the command is rejected with EX_ENOOBJ.

Syntax

```
raidcom reset hba_wwn -port <port#>[<host group name>]
                    -hba_wwn <WWN strings>
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number. For example:

- CL1-A-g (g: 0-254)
- CL1-A Linux_X86

-hba_wwn <WWN strings>

Specifies the WWN (16-digit hexadecimal value) of the host adapter. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

Example

Deleting the nickname that is given to the connection host "WWN:210000e0,8b0256f8" to which the port: CL4-E and host group #0 are set.

```
# raidcom reset hba_wwn -port CL4-E-0 -hba_wwn 210000e0,8b0256f8
```

raidcom set hba_wwn

Sets a nickname (maximum 64 characters) to the specified WWN on the specified port.

If the specified port does not exist, the command is rejected with EX_ENOOBJ.

If a nickname exists for the specified WWN, it is changed as NEW.

Syntax

```
raidcom set hba_wwn -port <port#>[<host group name>]
                    -hba_wwn <WWN strings> -wnn_nickname <WWN Nickname>
```

Options and parameters

-port <port#>[<host group name>]

Specifies the port number. For example:

- CL1-A-g (g: 0-254)
- CL1-A Linux_X86

-hba_wwn <WWN strings>

Specifies the WWN (16-digit hexadecimal value) of the host adapter. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-wnn_nickname <WWN Nickname>

Specifies the nickname (maximum 64 characters) to be assigned to the WWN of a specific port.

You cannot specify the same nickname to another WWN in the same port.

Example

Giving the nickname of "ORA_NODE0_CTL_0" to the connection host "WWN: 210000e0,8b0256f8" to which the port: CL4-E and host group #0 are set.

```
# raidcom set hba_wwn -port CL4-E-0 -hba_wwn 210000e0,8b0256f8
-wnn_nickname ORA_NODE0_CTL_0
```

raidcom add hba_iscsi

This command registers the iSCSI name (on the initiator) of the host bus adapter on the iSCSI target of the specified port in order to add hosts. If the specified iSCSI name already exists, the command is ignored.

Syntax

```
raidcom add hba_iscsi -port <port#> [<host group name>]
                        -hba_iscsi_name <initiator iscsi name>
```

Options and parameters

-port <port#>[<host group name>]

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator). You can specify within 223 characters. The maximum number of the host bus adapter is 255 for each port. For example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

Notes:

- You cannot register multiple iSCSI names to the same port if the only difference among them is the case (lowercase or uppercase). For example, if iqn.win2k8.example.of.iqn.form is already registered as an iSCSI name, you cannot register iqn.win2k8.example.of.iqn.FORM as a different iSCSI name.
- If you input an iSCSI name in iqn format, use lowercase letters only.
- If you input an iSCSI name in eui format, use lowercase letter s for "eui.", and use uppercase letters for the following hexadecimal value.

Examples

To set the iSCSI name "iqn.win2k8.example.of.iqn.form" of the host bus adapter whose port is CL4-E and whose target ID is 0:

```
# raidcom add hba_iscsi -port CL4-E-0 -hba_iscsi_name iqn.win2k8.example.of.iqn.form
```

To set the iSCSI name "eui.0123456789ABCDEF" of the host bus adapter whose port is CL4-E and whose iSCSI target name is Target00:

```
# raidcom add hba_iscsi -port CL4-E Target00 -hba_iscsi_name eui.0123456789ABCDEF
```

raidcom delete hba_iscsi

This command deletes the host (initiator iSCSI name) from the host group. If the specified initiator iSCSI name does not exist, the command is ignored.

Syntax

```
raidcom delete hba_iscsi -port <port#> [<host group name>]
                        -hba_iscsi_name <initiator iscsi name>
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Target00

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator). For example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

Examples

To delete the iSCSI name "iqn.win2k8.example.of.iqn.form" of the host bus adapter from the host group whose port is CL4-E and the target ID is 0:

```
# raidcom delete hba_iscsi -port CL4-E-0 -hba_iscsi_name
iqn.win2k8.example.of.iqn.form
```

To delete the iSCSI name "eui.0123456789ABCDEF" of the host bus adapter from the host group whose port is CL4-E and the iSCSI target name is Target00:

```
# raidcom delete hba_iscsi -port CL4-E Target00 -hba_iscsi_name eui.0123456789ABCDEF
```

raidcom set hba_iscsi

This command sets a nickname for the iSCSI name of the initiator on the specified port. If the specified port does not exist, the command is rejected with EX_ENOOBJ. If the nickname already exists for the specified initiator iSCSI name, the existing nickname is overwritten.

Syntax

```
raidcom set hba_iscsi -port <port#>[<host group name>]
                        -hba_iscsi_name <initiator iscsi name>
                        -iscsi_nickname <initiator iscsi Nickname>
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g g is from 0 to 254)
- CL1-A Target00

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator). For example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

-iscsi_nickname <initiator iscsi Nickname>

Specifies the nickname given to the iSCSI name of the initiator. You can specify up to 32 characters.

Example

This command sets the nickname: ORA_NODE0_CTL_0 for the connection host iSCSI name: iqn.win2k8.example.of.iqn.form whose port is CL4-E and the target ID is 0.

```
# raidcom set hba_iscsi -port CL4-E-0 -hba_iscsi_name iqn.win2k8.example.of.iqn.form -iscsi_nickname ORA_NODE0_CTL_0
```

raidcom reset hba_iscsi

This command removes the nickname from the iSCSI name of the initiator on the specified port. When the specified port does not exist, the command will be rejected with EX_ENOOBJ.

Syntax

```
raidcom reset hba_iscsi -port <port#> [<host group name>] -hba_iscsi_name <initiator iscsi name>
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, target ID, or iSCSI target name. For example:

- CL1-A-g g is from 0 to 254)
- CL1-A Linux_X86

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator). For example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

Examples

This command removes the nickname from the connection host iSCSI name:
iqn.win2k8.example.of.iqn.form whose port ID is CL4-E and the host group is 0.

```
# raidcom reset hba_iscsi -port CL4-E-0 -hba_iscsi_name iqn.win2k8.example.of.iqn.form
```

raidcom get hba_iscsi

This command displays the iSCSI name of the host bus adapter on the initiator for each iSCSI target, which is registered as an iSCSI target.

Syntax

```
raidcom get hba_iscsi -port <port#> [<host group name>]
```

Options and parameters

- port <port#>[<host group name>]**
Specifies the port number, target ID, or iSCSI target name. For example:
- CL1-A-g (g is from 0 to 254)
 - CL1-A Target00

Examples

To display the iSCSI name of the host bus adapter whose port is CL4-E and the iSCSI target ID is 0:

```
# raidcom get hba_iscsi -port CL4-E-0
```

PORT	GID	GROUP_NAME	IQN	Serial#	NICK_NAME
CL4-E	0	Linux_x86	iqn.z1...	63528	ORA_NODE0_CTL_0
CL4-E	0	Linux_x86	iqn.z2...	63528	ORA_NODE1_CTL_0

Description of each column in the output example:

- PORT**
Displays the port.
- GID**
Displays the target ID of the port.
- GROUP_NAME**
Displays the iSCSI target name of the port.
- IQN**
Displays the iSCSI name of the registered host bus adapter.

Serial#

Displays the product serial number.

NICK_NAME

Displays the nickname of the iSCSI name for the host bus adapter.

raidcom add external_iscsi_name

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command registers the iSCSI name of the iSCSI target on the external storage system to the iSCSI port of the local storage system.

When the iSCSI name has been registered in the iSCSI port of the specified local storage system, this command sets the CHAP authentication mode, and enables or disables the mutual CHAP authentication mode. When the iSCSI name has been registered in the port other than the iSCSI port of the specified local storage system, the command registers the iSCSI name to the iSCSI port of the specified local storage system. In this case, the information about the CHAP authentication mode and the mutual CHAP authentication which have been set to the iSCSI target is used by the local storage system.

If the specified iSCSI port does not exist, the command is rejected with EX_ENOOBJ.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add external_iscsi_name -port <port#> -iscsi_name
    <external iscsi name> -address <external IP address>
    [-authmethod {CHAP|NONE}] [-mutual {enable|disable}]
    [-tcp_port <value>]
    [-iscsi_virtual_port_id <iSCSI virtual port ID>] [-header_digest {enable|disable}]
    [-data_digest {enable|disable}]
```

Options and parameters**-port <port#>**

Specifies the port number of the local storage system. For example:

- CL1-A

-iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: `iqn.` and the subsequent maximum 219 characters.
- eui format: `eui.` and the subsequent 16 characters in hexadecimal notation.

-address <external IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-authmethod {CHAP|NONE}]

Can configure the CHAP authentication mode. Even if the CHAP user name has not been specified yet, the CHAP authentication mode can be configured.

- CHAP: CHAP authentication mode is enabled.
- NONE: The authentication mode is not configured.

If this option is omitted, the current setting value is maintained. The initial value is "NONE".

[-mutual {enable|disable}]

Enables or disables the mutual CHAP authentication mode. Even if the `-authmethod` option is configured to NONE, the mutual CHAP authentication mode can be enabled or disabled. When the `-authmethod` option changes to CHAP from NONE, this option setting becomes to be enabled.

- enable or yes: The mutual CHAP authentication is enabled. An iSCSI target recognizes the iSCSI initiator and vice versa.
- disable or no: The mutual CHAP authentication is disabled. An iSCSI target recognizes the iSCSI target.

If this option is omitted, the current setting value is maintained. The initial value is "disable".

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

[-tcp_port <value>]

Specifies the TCP port number of the iSCSI target on the external storage system. If this option is omitted, the TCP port number of the iSCSI target port which is specified with the `-port` option is configured.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

[-header_digest {enable|disable}]

Specifies whether a header digest is enabled or disabled.

This option is only available if a port on the local storage system is iSCSI 25G CHB. If this option is omitted, the header digest is enabled.

- enable: Enables the header digest.
- disable: Disables the header digest.

[-data_digest {enable|disable}]

Specifies whether a data digest is enabled or disabled.

This option is only available if a port on the local storage system is iSCSI 25G CHB. If this option is omitted, the data digest is enabled.

- enable: Enables the data digest.
- disable: Disables the data digest.

Examples

When you register the iSCSI name of the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system to the iSCSI port: CL4-E of the local storage system:

```
# raidcom add external_iscsi_name -port CL4-E -iscsi_name iqn.z1 -address
158.214.135.100
```

When you register the iSCSI name of the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system to the iSCSI port: CL4-E, virtual port ID: 1 of the local storage system:

```
#raidcom add external_iscsi_name -port CL4-E -iscsi_name iqn.z1 -address
158.214.135.100 -iscsi_virtual_port_id 1
```

When you change the CHAP authentication mode of the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system which is registered in the iSCSI port: CL4-E of the local storage system to "CHAP" and the mutual CHAP authentication mode to "enable":

```
# raidcom add external_iscsi_name -port CL4-E -iscsi_name iqn.z1 -address
158.214.135.100 -authmethod CHAP -mutual enable
```

When you register all iSCSI names of the iSCSI target on the external storage system which is registered in the port: CL2-E of the local storage system to the port: CL4-E:

```
# raidcom get external_iscsi_name | rmawk @1-eq:CL2-E exe="raidcom add
external_iscsi_name -port CL4-E -address @3 -iscsi_name @4"
```

When you register all iSCSI names of the iSCSI target on the external storage system which is registered in the port: CL2-E of the local storage system to the port: CL4-E. At this time, configure whether to enable the CHAP authentication mode and the mutual CHAP authentication mode:

```
# raidcom get external_iscsi_name | rmawk @1-eq:CL2-E | @7-eq:D exe="raidcom add
external_iscsi_name -port CL4-E -address @3 -iscsi_name @4 -authmethod @6!u -mutual
enable" -n exe="raidcom add external_iscsi_name -port @1 -address @3 -iscsi_name @4 -
authmethod @6!u -mutual disable"
```

When you search the iSCSI name of the iSCSI target which exists in the iSCSI port (IP address: 10.213.60.111) on the external storage system, and register the detected iSCSI name to the iSCSI port: CL4-E of the local storage system.

```
# raidcom discover external_iscsi_name -port CL4-E -address 10.213.60.111 | rmawk @5-
eq:N exe="raidcom add external_iscsi_name -port @1 -address @3 -iscsi_name @6"
```

raidcom delete external_iscsi_name

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command deletes the iSCSI name of the iSCSI target on the external storage system which is registered in the specified iSCSI port of the local storage system.

If the specified iSCSI port of the local storage system does not exist, the command is rejected with EX_ENOOBJ. If the iSCSI name of the iSCSI target on the specified external storage system does not registered in the iSCSI port of the specified local storage system, this command is ignored.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete external_iscsi_name -port <port#> -iscsi_name
    <external iscsi name> -address <external IP address>
    [-iscsi_virtual_port_id <iSCSI virtual port ID>]
```

Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

-iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: `iqn.` and the subsequent maximum 219 characters.
- eui format: `eui.` and the subsequent 16 characters in hexadecimal notation.

-address <external IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Examples

When you delete the iSCSI name: iqn.z1 of the iSCSI target (IP address: 158.214.135.100) on the external storage system from the iSCSI port: CL4-E of the local storage system:

```
# raidcom delete external_iscsi_name -port CL4-E -iscsi_name iqn.z1 -address
158.214.135.100
```

When you delete the iSCSI name: iqn.z1 of the iSCSI target (IP address: 158.214.135.100) on the external storage system from the iSCSI port: CL4-E, virtual port ID: 1 of the local storage system:

```
#raidcom delete external_iscsi_name -port CL4-E -iscsi_name iqn.z1 -address
158.214.135.100 -iscsi_virtual_port_id 1
```

When you delete all iSCSI names of the iSCSI target on the external storage system which are registered in the iSCSI port: CL4-E of the local storage system:

```
# raidcom get external_iscsi_name | rmawk @1-eq:CL4-E exe="raidcom delete
external_iscsi_name -port @1 -address @3 -iscsi_name @4"
```

When you delete all iSCSI names of the iSCSI target which exists on the iSCSI port (IP address: 158.214.135.100) of the external storage system from the iSCSI port of the local storage system:

```
# raidcom get external_iscsi_name | rmawk @3-eq:158.214.135.100 exe="raidcom delete
external_iscsi_name -port @1 -address @3 -iscsi_name @4"
```

raidcom get external_iscsi_name

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command displays the iSCSI name of the iSCSI target on the external storage system which is registered in the iSCSI port of the specified local storage system.

If the iSCSI port does not exist in the specified local storage system, the command is rejected with EX_ENOOBJ.

Only the iSCSI names registered in the port to which the user who executes the command can refer are output. For details about the port to which the user can refer, see the descriptions of the relationship of the resource group and the command operation in the *Command Control Interface User and Reference Guide*.

Syntax

```
raidcom get external_iscsi_name [-port <port#> [-iscsi_virtual_port_id <iSCSI virtual
port ID>]]
```

Options and parameters

[-port <port#>]

Specifies the port number of the local storage system. For example:

- CL1-A

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the iSCSI virtual port mode of the specified port is enabled. Only the iSCSI initiator of the port that can be referenced by the user who executes the command is output. If the microcode or firmware version does not support this option, nothing is output.

Examples

Displays all iSCSI targets of the external storage system which are registered in all iSCSI ports of the local storage system.

```
# raidcom get external_iscsi_name
```

PORT	Serial#	IP_ADDR	IQN	WWN(pseudo)	AMD	D	CHAP_user	Sec
ISCSI_VP_ID	IP_PORT#	HEADER_DIGEST	DATA_DIGEST	HEADER_DIGEST_SESSION				
DATA_DIGEST_SESSION								
CL4-E	63528	158.214.135.100	iqn.z1	50060e80070a3640	CHAP	D	Win_SQL_EX	
*	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL2-E	63528	158.214.135.100	iqn.z2	50060e80070a3641	CHAP	S	-	
-	1	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL2-E	63528	158.214.135.102	iqn.z3	50060e80070a3642	CHAP	S	-	
-	1	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.100	iqn.z2	50060e80070a3643	CHAP	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.102	iqn.z3	50060e80070a3644	CHAP	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.102	iqn.z4	50060e80070a3645	NONE	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.102	iqn.z5	50060e80070a3646	NONE	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED

Displays all iSCSI names of the iSCSI target on the external storage system which are registered in iSCSI port: CL4-E of the local storage system.

```
# raidcom get external_iscsi_name -port CL4-E
```

PORT	Serial#	IP_ADDR	IQN	WWN(pseudo)	AMD	D	CHAP_user	Sec
ISCSI_VP_ID	IP_PORT#	HEADER_DIGEST	DATA_DIGEST	HEADER_DIGEST_SESSION				
DATA_DIGEST_SESSION								
CL4-E	63528	158.214.135.100	iqn.z1	50060e80070a3640	CHAP	D	Win_SQL_EX	
*	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.100	iqn.z2	50060e80070a3643	CHAP	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.102	iqn.z3	50060e80070a3644	CHAP	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.102	iqn.z4	50060e80070a3645	NONE	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED
CL4-E	63528	158.214.135.102	iqn.z5	50060e80070a3646	NONE	S	-	
-	-	3260	ENABLED	ENABLED	ENABLED			ENABLED

Displays all iSCSI names of the iSCSI target of the external storage system registered in the iSCSI port: CL2-E, virtual port ID: 1 of the local storage system.

```
#raidcom get external_iscsi_name -port CL2-E -iscsi_virtual_port_id 1
```

```
PORT Serial# IP_ADDR IQN WWN(pseudo) AMD D CHAP_user Sec ISCSI_VP_ID IP_PORT#
HEADER_DIGEST DATA_DIGEST HEADER_DIGEST_SESSION DATA_DIGEST_SESSION
CL2-E 63528 158.214.135.100 iqn.z2 50060e80070a3641 CHAP S - - 1 3260
ENABLED ENABLED ENABLED ENABLED
CL2-E 63528 158.214.135.102 iqn.z3 50060e80070a3642 CHAP S - - 1 3260
ENABLED ENABLED ENABLED ENABLED
```

Description of each column in output example:

PORT

Displays the port number.

Serial#

Displays the product serial number.

IP_ADDR

Displays the IP address of the iSCSI target on the external storage system.

IQN

Displays the iSCSI name of the iSCSI target on the external storage system.

WWN(pseudo)

Displays the pseudo WWN of the iSCSI target on the external storage system. The pseudo WWN matches the iSCSI name of the iSCSI target and the IP address on the external storage system. The pseudo WWN is managed by each storage system. Therefore, if the iSCSI target on an external storage is shared within multiple storage systems, the pseudo WWN which is corresponded with an iSCSI target depends on the storage system.

AMD

Displays the authentication mode of the iSCSI target on the external storage system.

- CHAP: CHAP authentication is enabled.
- NONE: The authentication mode is not configured.

D

Displays the direction of the authentication mode of the iSCSI target.

- S: single directional (An iSCSI target recognizes the iSCSI target.)
- D: dual-directional (An iSCSI target recognizes the iSCSI initiator and vice versa.)

CHAP_user

Displays the CHAP user name of the iSCSI target on the external storage system. If the CHAP user name is not set, a hyphen (-) is displayed.

Sec

Displays an asterisk (*) when the secret is set to the iSCSI target on the external storage system. Otherwise, a hyphen (-) is displayed.

ISCSI_VP_ID

Displays the virtual port number when the virtual port mode is enabled. A hyphen (-) is displayed when the virtual port mode is disabled.

IP_PORT#

Displays the TCP port number of the iSCSI target on the external storage system.

HEADER_DIGEST

Specifies whether a header digest is enabled or disabled.

- ENABLED: Enables the header digest setting.
- DISABLED: Disables the header digest setting.
- - (hyphen): This information is not valid.

DATA_DIGEST

Specifies whether a data digest is enabled or disabled.

- ENABLED: Enables the data digest setting.
- DISABLED: Disables the data digest setting.
- - (hyphen): This information is not valid.

HEADER_DIGEST_SESSION

Specifies whether a header digest is enabled or disabled during session establishment.

- ENABLED: Enables the header digest during session establishment.
- DISABLED: Disables the header digest during session establishment.
- NOT_ESTABLISHED: There is no session established.
- - (hyphen): This information is not valid.

DATA_DIGEST_SESSION

Specifies whether a data digest is enabled or disabled during session establishment.

- ENABLED: Enables the data digest during session establishment.
- DISABLED: Disables the data digest during session establishment.
- NOT_ESTABLISHED: There is no session established.
- - (hyphen): This information is not valid.

raidcom get initiator_iscsi_name

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command displays the iSCSI initiator of the iSCSI port on the specified local storage system.

If the iSCSI port does not exist in the specified local storage system, the command is rejected with EX_ENOOBJ.

Only iSCSI initiators of the port to which the user who executes the command can refer are output. For details about the port to which the user can refer, see the descriptions of the relationship of the resource group and the command operation in the *Command Control Interface User and Reference Guide*.

Syntax

```
raidcom get initiator_iscsi_name -port <port#> [-iscsi_virtual_port_id <iSCSI virtual port ID>]
```

Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the iSCSI virtual port mode of the specified port is enabled. Only the iSCSI initiator of the port that can be referenced by the user who executes the command is output. If the microcode version does not support this option, nothing is output.

Examples

Displays the iSCSI initiator of the iSCSI port: CL4-E on the local storage system.

```
# raidcom get initiator_iscsi_name -port CL4-E
```

PORT	Serial#	IP_ADDR	IQN	IP_PORT#	CHAP_user	Sec	ISCSI_VP_ID
CL4-E	63528	158.214.197.100	iqn.z1	3260	Elun_INI_4E	*	-

Displays the iSCSI initiator of the iSCSI port: CL2-E on the local storage system.

```
# raidcom get initiator_iscsi_name -port CL2-E
```

```
PORT      Serial# IP_ADDR      IQN          IP_PORT#  CHAP_user  Sec  ISCSI_VP_ID
CL2-E     63528 158.214.197.101 iqn.zx      3260  Elun_INI_2E *      -
```

Displays the iSCSI initiator of the iSCSI port: CL2-E, virtual port: 1 on the local storage system.

```
# raidcom get initiator_iscsi_name -port CL2-E -iscsi_virtual_port_id 1
```

```
PORT      Serial# IP_ADDR      IQN          IP_PORT#  CHAP_user  Sec  ISCSI_VP_ID
CL2-E     63528 158.214.197.102 iqn.zy      3260  Elun_INI_2E *      1
```

Description of each column in output example:

PORT

Displays the port number.

Serial#

Displays the product serial number.

IP_ADDR

Displays the IP address which is set to the iSCSI initiator of the iSCSI port on the local storage system. See the table below for more information. To display the details of the IP address, use the `-key opt` option of the **raidcom get port** command.

IPv6 mode status	Address acquisition mode	Address status	IP address to display
Invalid	-	-	IP address of IPv4
Valid	Manual acquisition (MM)	-	Link local address
	Automatic acquisition (AM)	Valid (VAL)	-(hyphen)
		Invalid (INV)	
		Acquiring (ACQ)	
		Duplication (DUP)	

IQN

Displays the iSCSI name which is set to the iSCSI initiator of the iSCSI port on the local storage system.

IP_PORT#

Displays the TCP port number of the iSCSI target which is registered in the iSCSI port on the local storage system.

CHAP_user

Displays the CHAP user name which is set in the iSCSI initiator of the iSCSI port on the local storage system. If the CHAP user name is not set, a hyphen (-) is displayed.

Sec

Displays an asterisk (*) when the secret is set to the iSCSI initiator of the iSCSI port on the local storage system. Otherwise, a hyphen (-) is displayed.

ISCSI_VP_ID

Displays the virtual port number when the virtual port mode is enabled. A hyphen (-) is displayed when the virtual port mode is disabled.

raidcom discover external_iscsi_name

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

From the iSCSI port of the local storage system, this command searches the iSCSI targets which is registered in the port of the external storage system, and then displays the iSCSI name of the iSCSI target.

To execute this command, the host which executes the command has to support IPv6. If the host does not support an IPv6, the command is rejected with EX_ENOSUP.



Note: If you run multiple instances of the **raidcom discover external_storage** command, **raidcom discover lun** command, **raidcom check external_iscsi_name** command, or **raidcom discover external_iscsi_name** command at the same time, an EX_CMDIOE error might occur. If an EX_CMDIOE error occurs, wait for a while, and then run the commands again one by one. For details, see the *Hitachi Universal Volume Manager User Guide*.

Syntax

```
raidcom discover external_iscsi_name -port <port#>
  -address <external IP address> [-tcp_port <value>]
  [-iscsi_virtual_port_id <iSCSI virtual port ID>]
```


Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

-address <external IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-tcp_port <value>]

Specifies the TCP port number of the iSCSI target on the external storage system. If this option is omitted, the TCP port number of the iSCSI target port which is specified with -port option is configured.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Examples

Searches the iSCSI name of the iSCSI target which is registered in the iSCSI port (IP address: 10.213.60.111) on the external storage system from the iSCSI port: CL4-E of the local storage system, and then displays the iSCSI name:

```
# raidcom discover external_iscsi_name -port CL4-E -address 10.213.60.111
```

```
PORT    Serial# IP_ADDR          IP_PORT# R  IQN    ISCSI_VP_ID
```

CL4-E	63528	10.213.60.111	3260	N	iqn.z1	-
CL4-E	63528	10.213.60.111	3260	N	iqn.z2	-

Searches the iSCSI name of the iSCSI target which is registered in the iSCSI port (IP address: 10.213.60.112) on the external storage system from the iSCSI port: CL4-E, iSCSI virtual port ID: 1 of the local storage system, and then displays the iSCSI name:

```
#raidcom discover external_iscsi_name -port CL4-E -address 10.213.60.112 -
iscsi_virtual_port_id 1
```

PORT	Serial#	IP_ADDR	IP_PORT#	R	IQN	ISCSI_VP_ID
CL4-E	63528	10.213.60.112	3260	N	iqn.z1	1
CL4-E	63528	10.213.60.112	3260	N	iqn.z2	1

Description of each column in output example:

PORT

Displays the port number.

Serial#

Displays the product serial number.

IP_ADDR

Displays the IP address of the iSCSI target on the external storage system.

IP_PORT#

Displays the TCP port number of the iSCSI target on the external storage system.

R

Displays whether or not the searched iSCSI target of the external storage system has been registered in the iSCSI port of the local storage system.

- Y: The iSCSI target has been registered in the iSCSI port.
- N: The iSCSI target has not been registered in the iSCSI port.

IQN

Displays the searched iSCSI name of the iSCSI target on the external storage system.

ISCSI_VP_ID

Displays the virtual port number when the virtual port mode is enabled. A hyphen (-) is displayed when the virtual port mode is disabled.

raidcom check external_iscsi_name

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command attempts to login to the iSCSI target on the external storage system which has been registered in the local storage system, and then displays the result of the login.

To execute this command, the host which executes the command has to support IPv6. If the host does not support an IPv6, the command is rejected with EX_ENOSUP.



Note: If you run multiple instances of the `raidcom discover external_storage` command, `raidcom discover lun` command, `raidcom check external_iscsi_name` command, or `raidcom discover external_iscsi_name` command at the same time, an EX_CMDIOE error might occur. If an EX_CMDIOE error occurs, wait for a while, and then run the commands again one by one. For details, see the *Hitachi Universal Volume Manager User Guide*.

Syntax

```
raidcom check external_iscsi_name [-port <port#> [-iscsi_name <external iscsi name> -
address <external IP address>]
[-iscsi_virtual_port_id <iSCSI virtual port ID>]]
```

Options and parameters

[-port <port#>]

Specifies the port number of the local storage system. For example:

- CL1-A

[-iscsi_name <external iscsi name> -address <external IP address>]

Specify this option to display the login result when attempting to log in to the iSCSI target of the specified external storage system.

For `-iscsi_name <external iscsi name>`, specify the iSCSI name of the iSCSI target on the external storage system in the iqn or eui format:

- iqn format: iqn. followed by up to 219 characters.
- eui format: eui. followed by 16 characters in hexadecimal notation.

For -address <external IP address>, specify the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified. If the specified iSCSI target is not registered in the specified port or iSCSI virtual port, the command is rejected with EX_ENOOBJ.

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the iSCSI virtual port mode is enabled. When the iSCSI virtual port mode is enabled, if both this option and the -iscsi_name <external iscsi name> -address <external IP address> option are not specified, information including all iSCSI virtual port IDs is displayed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Examples

Attempts to log in to the iSCSI target of all external storage systems which are registered in the iSCSI port of the local storage system, and then displays the result of login:

```
# raidcom check external_iscsi_name
```

PORT	Serial#	IP_ADDR	IQN	IP_PORT#	AMD	D	CHAP_user	Sec	LOGIN
ISCSI_VP_ID									
CL4-E	63528	158.214.135.100	iqn.z1	3260	CHAP	D	Win_SQL_EX	*	OK
-									
CL2-E	63528	158.214.135.100	iqn.z2	3260	CHAP	S	-	-	OK
-									
CL2-E	63528	158.214.135.102	iqn.z3	3260	CHAP	S	-	-	OK
-									
CL4-E	63528	158.214.135.100	iqn.z2	3260	CHAP	S	-	-	OK
-									
CL4-E	63528	158.214.135.102	iqn.z3	3260	CHAP	S	-	-	OK
-									
CL4-E	63528	158.214.135.102	iqn.z4	3260	NONE	S	-	-	NG
-									
CL4-E	63528	158.214.135.102	iqn.z5	3260	NONE	S	-	-	NG
-									

Attempts to log in to the iSCSI target of the external storage systems which is registered in the iSCSI port: CL4-E of the local storage system, and then displays the result of login:

```
# raidcom check external_iscsi_name -port CL4-E
```

PORT	Serial#	IP_ADDR	IQN	IP_PORT#	AMD	D	CHAP_user	Sec	LOGIN	ISCSI_VP_ID
CL4-E	63528	158.214.135.100	iqn.z1	3260	CHAP	D	Win_SQL_EX	*	OK	
-										
CL4-E	63528	158.214.135.100	iqn.z2	3260	CHAP	S	-	-	OK	
-										
CL4-E	63528	158.214.135.102	iqn.z3	3260	CHAP	S	-	-	OK	
-										
CL4-E	63528	158.214.135.102	iqn.z4	3260	NONE	S	-	-	NG	
-										
CL4-E	63528	158.214.135.102	iqn.z5	3260	NONE	S	-	-	NG	
-										

Attempts to log in to the iSCSI target (iSCSI Name: iqn.z1, IP Address: 158.214.135.100) on the external storage system which is registered in the iSCSI port: CL4-E on the local storage system, and then displays the result of login:

```
# raidcom check external_iscsi_name -port CL4-E -iscsi_name iqn.z1 -address 158.214.135.100
```

PORT	Serial#	IP_ADDR	IQN	IP_PORT#	AMD	D	CHAP_user	Sec	LOGIN	ISCSI_VP_ID
CL4-E	63528	158.214.135.100	iqn.z1	-	-	-	-	-	OK	

Attempts to log in to the iSCSI target of the external storage system which is registered in the iSCSI port: CL2-E, iSCSI virtual port ID: 0 of the local storage system, and then displays the result of login:

```
# raidcom check external_iscsi_name -port CL2-E -iscsi_virtual_port_id 0
```

PORT	Serial#	IP_ADDR	IQN	IP_PORT#	AMD	D	CHAP_user	Sec	LOGIN	ISCSI_VP_ID
CL2-E	63528	158.214.135.100	iqn.z2	3260	CHAP	S	-	-	OK	0
CL2-E	63528	158.214.135.102	iqn.z3	3260	CHAP	S	-	-	OK	0

Description of each column in output example:

PORT

Displays the port number.

Serial#

Displays the product serial number.

IP_ADDR

Displays the IP address of the iSCSI target on the external storage system.

IQN

Displays the iSCSI name of the iSCSI target on the external storage system.

IP_PORT#

Displays the TCP port number of the iSCSI target on the external storage system.

AMD

Displays the authentication mode of the iSCSI target on the external storage system.

- CHAP: CHAP authentication is enabled.
- NONE: The authentication mode is not configured.

D

Displays the direction of the authentication mode of the iSCSI target.

- S: single directional (An iSCSI target recognizes the iSCSI target.)
- D: dual-directional (An iSCSI target recognizes the iSCSI initiator and vice versa.)

CHAP_user

Displays the CHAP user name of the iSCSI target on the external storage system. If the CHAP user name is not set, a hyphen (-) is displayed.

Sec

Displays an asterisk (*) when the secret is set to the iSCSI target of the external storage system. Otherwise, a hyphen (-) is displayed.

LOGIN

Displays the result of the login.

ISCSI_VP_ID

Displays the iSCSI virtual port ID. A hyphen (-) is displayed when the iSCSI virtual port is disabled.

A hyphen (-) is also displayed if the -iscsi_name <external iscsi name> -address <external IP address> option is specified and the -iscsi_virtual_port_id <value> option is omitted for the enabled iSCSI virtual port.

raidcom modify external_chap_user

This command sets the CHAP user name and the secret (password) to the iSCSI target of the specified external storage system.

When you omit both the CHAP user name and secret, the CHAP user name and the secret which are set in the iSCSI target are deleted. If the iSCSI target of the specified external storage system has been registered in multiple iSCSI ports of the local storage system, the settings are applied to all iSCSI ports of the local storage system.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify external_chap_user -port <port#> -iscsi_name
    <external iscsi name> -address <external IP address>
    [-iscsi_virtual_port_id <iSCSI virtual port ID>]
    [-chap_user <user name> ] [-secret]
```

Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

-iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: *iqn.* and the subsequent maximum 219 characters.
- eui format: *eui.* and the subsequent 16 characters in hexadecimal notation.

-address <external IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

[-chap_user <user name>]

Specifies the CHAP user name of the iSCSI target on the external storage system. You can input up to 223 characters for the name.

For example: storage01

[-secret]

Displays the prompt for entering a secret. You can specify the secret within the range of 12 to 32 characters.

If the characters you input for the secret are less than 12 characters, an error occurs. If the characters you input for the secret are more than 32 characters, the first 32 characters are valid.

For the details of available characters for the secret when you use CCI, see the supported characters topic. Note that the space and back slash cannot be used.

Examples

When you set the CHAP user name: Elun_TAR_4E and the secret to the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system which is registered in the port: CL1-A of the local storage system:

```
# raidcom modify external_chap_user -port CL1-A -iscsi_name iqn.z1 -address
158.214.135.100 -chap_user Elun_TAR_4E -secret
Enter Secret :
```

When you set the secret to the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system which is registered in the port: CL1-A of the local storage system:

```
# raidcom modify external_chap_user -port CL1-A -iscsi_name iqn.z1 -address
158.214.135.100 -secret
Enter Secret :
```

When you delete the registered CHAP user name and the secret from the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system which is registered in the port: CL1-A of the local storage system:

```
# raidcom modify external_chap_user -port CL1-A -iscsi_name iqn.z1 -address
158.214.135.100
```


When you set the CHAP user name and the secret for the iSCSI target (iSCSI name: iqn.z1, IP address: 158.214.135.100) on the external storage system which is registered in the port: CL1-A, iSCSI virtual port ID: 2 of the local storage system:

```
# raidcom modify external_chap_user -port CL1-A -iscsi_virtual_port_id 2 -iscsi_name
iqn.z1 -address 158.214.135.100 -chap_user Elun_TAR_4E -secret
Enter Secret :
```

raidcom modify initiator_chap_user

This command sets the CHAP user name and the secret to the iSCSI initiator of the specified local storage system. When you omit both the CHAP user name and secret, the CHAP user name and the secret which are set in the iSCSI initiator are deleted.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify initiator_chap_user -port <port#>
[-chap_user <user name> ] [-secret]
```

Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

[-chap_user <user name>]

Specifies the CHAP user name of the iSCSI initiator on the external storage system. You can input up to 223 characters for the name. For example:

storage01

[-secret]

Displays the prompt for entering a secret. You can specify the secret within the range of 12 to 32 characters.

If the characters you input for the secret are less than 12 characters, an error occurs. If the characters you input for the secret are more than 32 characters, the first 32 characters are valid.

For the details of available characters for the secret when you use CCI, see the topic Supported characters. Note that the space and back slash cannot be used.

Examples

When you set the CHAP user name: Elun_INI_4E of the iSCSI initiator and the secret to the iSCSI port: CL4-E of the local storage system:

```
# raidcom modify initiator_chap_user -port CL4-E -chap_user Elun_INI_4E -secret
Enter Secret :
```

When you set the secret of the iSCSI initiator to the iSCSI port: CL4-E of the local storage system:

```
# raidcom modify initiator_chap_user -port CL4-E -secret
Enter Secret :
```

When you delete the CHAP user name and the secret of the iSCSI initiator from the iSCSI port: CL4-E of the local storage system:

```
# raidcom modify initiator_chap_user -port CL4-E
```

raidcom add journal

Registers (adds) a volume to a journal.

This makes a journal for open systems (Open System) or mainframe (M/F System) according to the specified LDEVs.

If the journal already exists, the specified LDEV is added to the journal.

If the journal does not exist, you must create it first, and then add an LDEV.

If the `-timer_type` option is specified, this command makes the journal as "M/F System". If not, this command makes the journal as "Open System".

A device group can also be specified instead of an LDEV.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If cache memory is blocked, the operation cannot be performed.

Syntax

```
raidcom add journal -journal_id <journal ID#> {-ldev_id <ldev#> ...
    [-cnt <count>] | -grp_opt <group option>
    -device_grp_name <device group name> [<device name>]}
    [-mp_blade_id <mp#> | -timer_type <timer type> ]
```

Options and parameters**-journal_id <journal ID#>**

Specifies the journal number (0-255).

The journal number can be specified in decimal or hexadecimal. If you specify a hexadecimal number, add 0x at the beginning of the journal number.

-ldev_id <ldev#> ...

Specifies the LDEV number (0-65279).

Up to 64 LDEVs can be specified at a time. For example:

- -ldev_id 200
- -ldev_id 100-110
- -ldev_id 100 -cnt 10

[-cnt <count>]

Specifies the count (2-64).

If this option is omitted, "1" is used as the count.

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all of the LDEVs belonging in the device group are operated.

[-mp_blade_id <mp#>]

Specifies the MP blade ID number (0-15).

If this option is omitted, the device automatically allocates an MP blade ID number. For example:

- -mp_blade_id 2

[-timer_type <timer type>]

Specifies the timer type: system, local, or None. For example:

- -timer_type system
- -timer_type local

For details, see the Universal Replicator manual for the storage system.

Examples

Examples for open systems:

Creating a journal #1 of LDEVs: 265, 266.

```
# raidcom add journal -journal_id 1 -ldev_id 265 266
```

Creating a journal #1 with an LDEV belonging to the device group: grp1.

```
# raidcom add journal -journal_id 1 -grp_opt ldev -device_grp_name grp1
```

Example for mainframe:

Creating a journal #1 with LDEVs: 265, 266. As a timer type, the system clock of the mainframe host is used. (When LDEV is added to the journal of M/F, it is required to specify the system clock of the mainframe host to the timer type.)

```
# raidcom add journal -journal_id 1 -ldev_id 265 266
-timer_type system
```

raidcom delete journal

Deletes a journal from the specified journal.

A device group can also be specified instead of an LDEV. If the LDEV and device group are not specified, a journal is deleted.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete journal -journal_id <journal ID#> [-ldev_id
<ldev#> | -grp_opt <group option> -device_grp_name
<device group name> [<device name>]]
```

Options and parameters

-journal <journal ID#>

Specifies the journal number (0-255).

The journal number can be specified in decimal or hexadecimal. If you specify a hexadecimal number, add 0x at the beginning of the journal number.

[-ldev_id <ldev#>]

Specifies the LDEV number (0-65279). If the both LDEV and device group are not specified, journal is deleted. For example:

- -ldev_id 200

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all of the LDEVs belonging in the device group are operated. If the both LDEV and device group are not specified, journal is deleted.

Examples

Deleting the specified journal.

```
# raidcom delete journal -journal_id 6
```

Deleting the specified LDEV from the journal.

```
# raidcom delete journal -journal_id 6 -ldev_id 265
```

Deleting the LDEV belonging to the device group: grp1 from the journal.

```
# raidcom delete journal -journal_id 6 -grp_opt ldev  
-device_grp_name grp1
```

raidcom get journal

Displays the information of registered journal.

Syntax

Displaying journal information

```
raidcom get journal [-key <keyword>]
```

Displaying timer-related information

```
raidcom get journalt
```

Options and parameters**[-key <keyword>]**

Specifies the display keyword. Specify opt as <keyword>.

Examples

Displaying journal information.

```
# raidcom get journal
```

JID	MU	CTG	JNLS	AP	U(%)	Q-Marker	Q-CNT	D-SZ(BLK)	Seq#	Num	LDEV#
001	0	1	PJNN	4	21	43216fde	30	512345	62500	4	265
002	1	2	PJNF	4	95	3459fd43	52000	512345	62500	3	270
002	2	2	SJNS	4	95	3459fd43	52000	512345	62500	3	270
003	0	3	PJSN	4	0	-	-	512345	62500	1	275
004	0	4	PJSF	4	45	1234f432	78	512345	62500	1	276
005	0	5	PJSE	0	0	-	-	512345	62500	1	277

Displaying Timer related information of the journal.

```
# raidcom get journalt
```

JID	MU	CTG	JNLS	AP	U(%)	Q-Marker	Q-CNT	D-SZ(BLK)	Seq#	DOW	PBW	APW
001	0	1	PJNN	4	21	43216fde	30	512345	63528	20	300	40
002	1	2	PJNF	4	95	3459fd43	52000	512345	63528	20	300	40
003	0	3	PJSN	4	0	-	-	512345	63528	20	300	40

Displaying option information of the journal.

```
# raidcom get journal -key opt
```

JID	MU	CTG	JNLS	TYPE	TTYE	MODE	IF	DOW(S)	PBW(M)	CR	CS(bps)	DM
MP#	Seq#	T	RCMD#									
000	0	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
0	302614	N	NA									
000	1	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
0	302614	N	NA									
000	2	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
0	302614	N	NA									
000	3	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
0	302614	N	NA									
001	0	1	PJSN	OPEN	-	CACHE	E	60	5	L	256	Y
4	302614	Y	100									
001	1	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
4	302614	N	NA									
001	2	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
4	302614	N	NA									
001	3	0	SMPL	OPEN	-	CACHE	E	60	5	L	256	Y
4	302614	N	NA									

```

002 0 1 SJSN M/F system CACHE E 60 5 L 256 Y
0 302614 N NA
002 1 0 SMPL M/F system CACHE E 60 5 L 256 Y
0 302614 N NA
002 2 0 SMPL M/F system CACHE E 60 5 L 256 Y
0 302614 N NA
002 3 0 SMPL M/F system CACHE E 60 5 L 256 Y
0 302614 N NA

```

Description of each column in output example:

JID

Journal number in decimal.

MU

Mirror ID on Universal Replicator.

CTG

Consistency group ID.

JNLS

Status in the journal:

- SMPL: a journal volume that does not have a pair, or is being deleted.
- P(S)JNN: "P(S)vol Journal Normal Normal".
- P(S)JNS: "P(S)vol Journal Normal Suspend" created with `-nocsus` option.
- P(S)JSN: "P(S)vol Journal Suspend Normal".
- P(S)JNF: "P(S)vol Journal Normal Full".
- P(S)JSF: "P(S)vol Journal Suspend Full".
- P(S)JSE: "P(S)vol Journal Suspend Error" including link failure.
- P(S)JES: "P(S)vol Journal Error Suspend" created with `-nocsus` option.

AP

Number of active link paths of Universal Replicator.

U(%)

Usage rate of journal volumes assuming the entire relevant volume is 100%.

Q-Marker

The P-VOL journal volume shows the newest sequential number (Q-marker) at the time of receiving WRITE data. The S-VOL journal volume shows the newest sequential number (Q-marker) written to the cache.

Q-CNT

Number of Q-markers remaining in the P-VOL journal volume.

D-SZ(BLK)

Capacity of the data block size of the journal volume in units of 512 bytes.

For details about the displayed capacity, see *Hitachi Universal Replicator User Guide*.

Seq#

Serial number (Seq#).

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Num

Number of LDEVs composing the journal volume.

LDEV#

LDEV number of a journal volume for storage systems that support reserve journal volumes:

- LDEV number of the master journal volume for the master journal
- LDEV number of the restore journal volume for the restore journal. Whether to support reserve journal volumes depends on the model of your storage system. For the support status, see the *Hitachi Universal Replicator User Guide* for your model.

DOW

Data Overflow Watch timer setting (in seconds) per the Journal.

PBW

Path Blockade Watch timer setting (in seconds) per the Journal. If the setting is more than 3600 seconds, it displays 6000 seconds.

APW

Active Path Watch timer setting (in seconds) to detect any link failures.

TYPE

Open or M/F system.

TTYE

Timer type for mainframe: System or Local.

MODE

Status of the journal:

- HDD: Store the journal data to the journal volume (cache mode disabled).
- CACHE: Store the journal data to the cache (cache mode enabled).

IF: Inflow control for journal:

- E: Enable
- D: Disable

DOW(S)

Data Overflow Watch timer setting (in seconds) per the Journal.

PBW(M)

Path Blockade Watch timer setting (in minutes) per the Journal. If the setting is more than 60 minutes, it displays 100 minutes.

CR

Copy rate: L (Low), M (Medium), or H (High).

CS(bps)

Copy speed in Mbps: 3 (3 Mbps), 10 (10 Mbps), 100 (100 Mbps), or 256 (256 Mbps).

DM

Copy mode "Y" or "N" under failure of the delta resync:

- Y: copying ALL data.
- N: No copying.

MP#

MP blade ID.

T

Displays whether to transfer the path watch time for the master journal to the secondary side (RCU side) of the mirror.

- Y: Transfer.
- N: Do not transfer.
- - (hyphen): This information is not available for this journal.

RCMD#

LDEV number of a remote command device. Displays "NA" if a remote command device is not allocated, or "NU" if a remote command device is allocated with no LDEV number specified. Displays a hyphen (-) if display of remote command device information is not supported. If the `-fx` option is specified, an LDEV number is displayed in hexadecimal.

raidcom modify journal

Changes the option of Universal Replicator to be used at journal, and sets the specified control parameter to the journal.

If you specify the `-mp_blade_id` option, you cannot specify the other options (`-data_overflow_watch`, `-path_blocked_watch`, `-cache_mode`, or `-timer_type`).

The `-timer_type` option must be specified to the journal on mainframe systems. This means that this option cannot be used for changing from "Open System" to "M/F System".



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

When changing the data overflow watching time of journal data area, cache mode or the timer type.

```
# raidcom modify journal -journal_id <journal ID#>
    {[-data_overflow_watch <time>][-cache_mode {y | n}]
    [-timer_type <timer type>]}
```

When changing the path blocked watch.

```
# raidcom modify journal -journal_id <journal ID#>
    -path_blocked_watch <time> [-mirror_id <mu#>]
```



Note: This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

When setting the path blocked watch, path watch time transfer, copy pace, transfer speed, and operation when delta resync fails by specifying the mirror ID.

```
raidcom modify journal -journal_id <journal ID#> [-mirror_id <mu#>] {[-
path_blocked_watch <time>] [-path_blocked_watch_transfer {y|n}] [-
copy_size <size>] [-transfer_speed <speed>] [-entire_copy {y|n}]}
```

When setting or releasing the remote command device.

```
raidcom modify journal -journal_id <journal ID#> [-mirror_id <mu#>]
{-command_device y [-ldev_id <ldev#>] | -command_device n}
```

When changing the MP blade ID.

```
# raidcom modify journal -journal_id <journal ID#>
    -mp_blade_id <mp#>
```

Options and parameters

-journal_id <journal ID#>

Specifies the journal number (0-255).

The journal number can be specified in decimal or hexadecimal. If you specify a hexadecimal number, add 0x at the beginning of the journal number.

-data_overflow_watch <time(sec)>

The data overflow watch timer (0-600) (second).

Specifies the watch time for the journal data area being full.

If 0 is specified, the setting of the data overflow watch timer is changed and the inflow control for journals is set to disabled.

-path_blocked_watch <time(min)>

The watch for the path blockage (1-60) (minute).

Disable the option ID 16 of the advanced system settings by using the **raidcom modify user_system_opt** command to set a value of more than 60 minutes. This setting stops the path blockage watching. For the storage systems other than VSP One B20, Device Manager - Storage Navigator can also be used to disable the advanced system settings.

If 0 is specified, the time of the watch for the path blockage does not change, and setting for transferring the path blockage time from the primary side of the mirror (MCU side) to the secondary side of the mirror (RCU side) is invalid.

If the path blockage time (1 to 60) (minutes) is specified, setting for transferring the path watch time from the primary side of the mirror (MCU side) to the secondary side of the mirror (RCU side) is enabled.

[-cache_mode {y | n}]

Specifies whether to use the cache mode.

- y or yes: Cache mode enabled (E)
- n or no: Cache mode disabled (D)

[-timer_type <timer type>]

Specifies the timer type: system, local, or None. For example:

- -timer_type system
- -timer_type local

For details, see the Universal Replicator manual for the storage system.

-mp_blade_id <mp#>

Specifies the MP blade ID (0-15).

Do not change the MP blade ID during initial copying or I/O processing of Universal Replicator.

If you change MP blade IDs for multiple journals, wait at least 10 minutes before changing another MP blade ID.

To change the MP blade ID again for the same journal, wait for more than 30 minutes after changing the MP blade ID. For example:

- -mp_blade_id 2

[-mirror_id <mu#>]

Specifies the Mirror ID.

If the setting is omitted, "0" is used.

[-copy_size <size>]

Specify the <size> from 1 to 15 as follows. If you specify the large value as this option, the copy time can be reduced, but the I/O performance might deteriorate.

- 1 or 2: low-speed
- 3: medium-speed
- 4 or more: high-speed

-transfer_speed <speed>

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specify the data transfer speed by Mbps. You can specify 256, 100, 10 or 3.

-path_blocked_watch_transfer {y|n}

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specify whether to transfer the path watch time from the primary side of the mirror (MCU side) to the secondary side of the mirror (RCU side).

- y or yes: Transfer
- n or no: Not transfer

-entire_copy {y|n}

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specify the operation when delta resync cannot be performed.

- y or yes: All the data of the P-VOL is copied to the S-VOL.
- n or no: No operation. The S-VOL is not updated.

-command_device y [-ldev_id <ldev#>]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Sets a remote command device. A remote command device specified by -ldev_id <ldev#> is set for the specified mirror ID.

-command_device n

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Releases a remote command device. The remote command device of the specified mirror ID is released.

Examples

Changing the data overflow watch time for journal "6" to 15 seconds.

```
# raidcom modify journal -journal_id 6 -data_overflow_watch 15
```

Changing the settings for journal 6: data overflow watch time to 15 seconds, and the timer type to the system clock of the mainframe host.

```
# raidcom modify journal -journal_id 6 -data_overflow_watch 15
-timer_type system
```

Changing the the MP blade ID for journal "6" to 2.

```
# raidcom modify journal -journal_id 6 -mp_blade_id 2
```

Changing the setting for journal "6" to store journal data in the secondary journal to the cache.

```
# raidcom modify journal -journal_id 6 -cache_mode y
```

Changing the path block monitoring time of mirror ID 1 for journal "6" to 59 minutes.

```
# raidcom modify journal -journal_id 6 -path_blocked_watch 59
-mirror_id 1
```

Changing the setting for journal "6" to set the remote command device of LDEV number 1 for mirror ID 1.

```
# raidcom modify journal -journal_id 6 -mirror_id 1 -command_device y -ldev_id 1
```

Changing the setting for journal "6" to release the remote command device of mirror ID 1.

```
# raidcom modify journal -journal_id 6 -mirror_id 1 -command_device n
```

raidcom add ldev

Adds an LDEV to the specified parity group or the external volume group. Or this adds V-VOL for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, active flash for mainframe, Thin Image, or Copy-on-Write Snapshot to the specified pool.

Alternatively, this command creates V-VOL for Dynamic Provisioning associated with a pool volume having the data direct mapping attribute.

When an LDEV is created in the specified DDP parity group, all available space in the parity group is assigned to the LDEV. The number and capacity of LDEVs are set automatically. The created LDEV is only available for a pool volume.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#)



Note:

For VSP 5000 series, VSP One B28, VSP E1090, VSP E990, or VSP E790, if the total cache initialization capacity for open system DP-VOLs created at once is large, the processing might time out.

If the number of DP-VOLs to be created by using this command exceeds 256, 256 DP-VOLs are created per creation on the storage system.

To prevent this processing from timing out, the total cache initialization capacity for up to 256 DP-VOLs to be created should be at most 8 PB. If 8 PB is exceeded, limit the total cache initialization capacity for DP-VOLs to be created at once to less than 8 PB, verify that the processing is complete by using the **raidcom get command_status** command, and then run the command to create the next DP-VOLs.

For more information about how to calculate the initialization capacity of the cache, see the *Provisioning Guide* for your storage system.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If cache memory is blocked, a DP-VOL cannot be created. If CHB, DKB, or cache memory is blocked, a V-VOL in Thin Image cannot be created.

Syntax

When creating an LDEV or V-VOL in the specified parity group, external volume group, or pool.

```
raidcom add ldev {-parity_grp_id <gno-sgno>|
  -external_grp_id <gno-sgno> | -pool {<pool ID#> |
    <pool naming> | snap}} {-ldev_id <ldev#> |
  -tse_ldev_id <ldev#> | -ese_ldev_id <ldev#> | -ldev_id auto -request_id auto
  [-resource_id <resource group id>] [-ldev_range <range>]}
  {-capacity <size> | -offset_capacity <size> | -cylinder <size>}
  [-emulation <emulation type>][-location <lba>]
  [-mp_blade_id <mp#>][-clpr <clpr#>]
  [-tl0pi_enable]
  [-status {enable_fullallocation | disable_fullallocation} | -capacity_saving
  {compression | deduplication_compression}
  [-capacity_saving_mode <saving_mode>] [-compression_acceleration {enable |
  disable} -request_id auto] | [-drs -request_id auto]
  -capacity_saving disable]
```

When creating a V-VOL for Dynamic Provisioning associated with a pool volume having the data direct mapping attribute.

```
raidcom add ldev -ldev_id <ldev#> -mapping_ldev_id <ldev id>
```

When creating an LDEV in a DDP parity group and automatically assigning the LDEV ID.

```
raidcom add ldev -parity_grp_id <gno-sgno> -ldev_id auto [-resource_id <resource  
group id>] [-ldev_range <range>] [-emulation <emulation type>] [-mp_blade_id <mp#>] -  
ddp_ldev -request_id auto
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno:1-52, sgno:1-32) (for example, 3-1).

-external_grp_id <gno-sgno>

Specifies the external volume group number (gno:1-16384, sgno:1-4096) (for example, 52-11 ("E" is not required)).

-pool {<pool ID#> | <pool naming> | snap}

Specifies a Pool ID or Pool name for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe when V-VOL is created to Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe pool.

When only a number is specified, it is recognized as a pool ID. Therefore, when you specify a pool that the pool name is only a number, specify the pool ID instead of the pool name.

When creating V-VOL for a Thin Image or Copy-on-Write Snapshot pool, specify "snap".

-ldev_id <ldev#> | -tse_ldev_id <ldev#> | -ese_ldev_id <ldev#>

Specifies the LDEV number (0-65279).

For VSP E590, the range is 0-32,767. For VSP E790, the range is 0-49,151.

If you specify the `-tse_ldev_id` option, the TSE-VOL that is used by FlashCopy SE is created. When you specify the `-ese_ldev_id` option, ESE-VOL is created. You can specify the `-ese_ldev_id` option to VSP 5000 series, and VSP G1x00 and VSP F1500 only. TSE-VOL and ESE-VOL can be created in the pool of Dynamic Provisioning for Mainframe only. If you want to create TSE-VOL and ESE-VOL, specify the pool of Dynamic Provisioning for Mainframe by using the `-pool` option. For example:

- `-ldev_id 200`
- `-tse_ldev_id 400`
- `-ese_ldev_id 600`

If you use the `-tse_ldev_id` or `-ese_ldev_id` option in a command other than the `raidcom add ldev` command, that option functions as the `-ldev_id` option.

-ldev_id auto -request_id auto [-resource_id <resource group id>] [-ldev_range <range>]

This option can be specified only when a V-VOL of which emulation type is OPEN-V is created in a Dynamic Provisioning or Dynamic Tiering pool.

Specify this option to automatically assign an LDEV number to the LDEV you created. From the unused numbers of the LDEVs for which the user has operation permission, the smallest LDEV number is assigned. To limit the range of the LDEV numbers to be assigned, use any of the following options:

- `-request_id auto`: Effective only when "auto" is specified as the `-request_id` option. If other option than "auto" is specified, EX_INVARG or EX_REQARG is replied.
- `-resource_id <resource grp id>`: Assigns the smallest LDEV number from the numbers of the unused LDEVs with the specified resource group ID.
- `-ldev_range <range>`: Assigns the smallest LDEV number from the numbers of the unused LDEVs with the LDEV number specified by <range>. If the `-resource_id <resource grp id>` is specified with this option, the smallest LDEV number from the numbers of the unused LDEVs that have the LDEV number in <range> and the specified resource group number is assigned.

Specify <range> in the format of *<starting-LDEV-number-ending-LDEV-number>*. LDEV numbers must be specified in decimal or hexadecimal. To specify in hexadecimal, add 0x at the beginning of the LDEV number.

For example, to specify LDEV numbers from 300 to 305:

- `-ldev_range 300-305`
- `-ldev_range 0x12c-0x131`
- `-ldev_range 0x12c-305`

(You can combine decimal and hexadecimal numbers.)

You can check the selected LDEV numbers using the `raidcom get command_status` command.

The request ID is output in the following format after completion of the **raidcom add ldev** command:

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed.
For details, see [Request ID function \(on page 293\)](#).

Unused LDEV numbers in the resource group where you create an LDEV and in the LDEV number range must satisfy the following requirements:

- The unused LDEV numbers are not in the same range (grouped by every 32 LDEV numbers) as the installed LDEV in an emulation group other than OPEN-V. If this requirement is not satisfied, the command might fail with EX_CMDRJE (SSB1=2E30, SSB2=0026).
- The unused LDEV numbers are not used as an alias of PAV. If this requirement is not satisfied, the command might fail with EX_CMDRJE (SSB1=2E30, SSB2=0025).
- The unused LDEV numbers can be used in the shared memory installed in the storage system. If this requirement is not satisfied, the command might fail with EX_CMDRJE (SSB1=2E23, SSB2=0001).
- When using a mainframe volume in the same storage system, SSID that is not used by the mainframe volume is assigned to the unused LDEV numbers.

The storage system might execute multiple **raidcom add ldev** commands at the same time. In such a batch operation, if one command fails, the others also fail. You can display command error information by executing **raidcom get command_status - request_id <request#>**, where you must specify the Request ID output when the command is executed. For errors caused due to batch operation, 'O' is displayed in column R, while for errors caused for other reason, 'T' is displayed in the same column. For example, if a user issues two **raidcom add ldev -ldev_id auto** commands at the same time to a storage system having only one free LDEV number, the system attempts to execute the commands in a batch, causing both of them to fail. Check the command status by executing **raidcom get command_status - request_id <request#>** with the Request ID. If 'O' is displayed under column R, retry the command.

The command cannot be run if either of the following conditions is met:

1. There are not enough unimplemented LDEVs to create an LDEV in the resource group with the smallest resource group ID for which the user has operation permission or in the resource group specified in the **-resource_id <resource group id>** option.
2. When the **-ldev_range <range>** option is specified, in addition to condition 1, there are not enough unimplemented LDEVs to create an LDEV among the LDEV numbers specified in **<range>**.

Before running the command, use the `raidcom get ldev -ldev_list undefined -resource <resource_grp_id>` command to verify whether there are enough unimplemented LDEVs to create the LDEV. If there are not enough unimplemented LDEVs to create the LDEV, use the `raidcom add resource -ldev_id <ldev#>` command or the `raidcom delete resource -ldev_id <ldev#>` command to move the unimplemented LDEV numbers to the target resource group, or use the `raidcom delete ldev` command to delete unnecessary LDEVs.

-capacity <size>

Specifies the capacity. The size can be specified in bytes or blocks. When specifying in bytes, the unit is t/T (terabyte), g/G (gigabyte), m/M (megabyte), or k/K (kilobyte). If this unit is omitted, block (512 bytes) is used.

When specifying in bytes, note the following:

- There is no size correction.
- If the capacity of LDEVs that are created by GUI and by CLI are the same, a copy pair might not be created. To create a pair with an LDEV that was created by GUI, create the LDEV by specifying blocks.

Example of specification:

1GB (gigabyte) is:

-capacity 1G, -capacity 1g, -capacity 1024M, -capacity 1024m, -capacity 1048576K, -capacity 1048576k, -capacity 2097152

In the case of volumes other than OPEN-V, more free space than the capacity that was actually specified is required. For details, see the *Provisioning Guide* for the storage system.

When you specify "all" instead of digits to create an OPEN-V LDEV and no LDEV has been created in the specified parity group* or the external volume group, the system allocates all empty space for the LDEV. If the empty space is larger than the maximum capacity of an LDEV, the system creates the LDEV with the maximum capacity and leaves the rest of the space as an empty space.

*: "all" cannot be specified depending on the combination of the drive type and drive level because of the LDEV control area allocation, and so on. In this case specify bytes or blocks.

-offset_capacity <size>

Specifies capacity. This option is only used when using HDvM - SN.

The same correction is performed as when Capacity Compatibility Mode (Offset boundary) is specified on HDvM - SN. The capacity can be specified in bytes or blocks. When specifying byte, the unit is t/T (terabyte), g/G (gigabyte), m/M (megabyte), or k/K (kilobyte).

-cylinder <size>

Specifies the size in cylinder unit. When specifying, the unit is t/T (teracylinder), g/G (gigacylinder), m/M (megacylinder), or k/K (kilocylinder). Cylinder is applied if you do not specify the unit.

When an open-systems emulation type is specified, this option cannot be specified.

[-location <lba>]

Specifies the Location (the starting point of the LDEV to be created in the parity group/ external volume group). If this specification is omitted, create a LDEV and close up in the free space.

[-emulation <emulation type>]

Specifies the emulation type (for example, OPEN-V).

If this specification is omitted, OPEN-V is specified.

If this specification is omitted when the virtual volume is created, the following emulation type is specified.

- OPEN-V: Dynamic Provisioning, Dynamic Tiering, or active flash
- 3390-A: Dynamic Provisioning for Mainframe, Dynamic Tiering for Mainframe, or active flash for mainframe

The values that can be specified for <emulation type> are:

- OPEN-3, OPEN-8, OPEN-9, OPEN-E, OPEN-L, OPEN-V
- 3390-1, 3390-2, 3390-3, 3390-A, 3390-3A, 3390-3B, 3390-3C, 3390-3R, 3390-9, 3390-9A, 3390-9B, 3390-9C, 3390-L, 3390-LA, 3390-LB, 3390-LC, 3390-M, 3390-MA, 3390-MB, 3390-MC, 3390-V
- 3380-3, 3380-3A, 3380-3B, 3380-3C

Some emulation types cannot be specified depending on the emulation type.

**Caution:**

You can specify 3390-3 or 3390-3R as the emulation type, but these are unable to be mixed. You can specify the 3380 series or the 3390 series as the emulation type for each parity group, but they are unable to be mixed for each 32 address boundary because of the OS restriction.

[-mp_blade_id <mp#>]

Specifies the MP blade ID (0-15). If this specification is omitted, it is allocated automatically.

[-clpr <clpr#>]

When you create virtual volumes by specifying the `-pool <pool#>` option, specify the CLPR ID. If you omit this option, the ID number of CLPR to which the pool is allocated is used.

[-status {enable_fullallocation | disable_fullallocation}]

Specifies an availability of Full Allocation when the virtual volume is used for Dynamic Provisioning, Dynamic Tiering, or active flash. If this specification is omitted, Full Allocation is disabled.

- `enable_fullallocation`: Enables Full Allocation. If all areas equivalent to the sum of the pool capacities the specified volume requires can be reserved, it is guaranteed that all areas of DP-VOL are writable.
- `disable_fullallocation`: Disables Full Allocation.

When you create a volume whose capacity saving setting is compression, or deduplication and compression, this option cannot be specified to enable.

[-t10pi_enable]

Enables the T10 PI attribute.

-mapping_ldev_id <ldev id>

Creates the V-VOL for Dynamic Provisioning associated with the pool volume having the data direct mapping attribute. When you specify this option, the data direct mapping attribute is automatically set to the V-VOL to be created.

[-compression_acceleration {enable | disable} -request_id auto]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E1090

Specifies whether compression accelerator is enabled or disabled. If you omit this option, it will be enabled when COMPRESSION_ACCELERATION of raidcom get system -key dedupe_compression is AVAILABLE. It will be disabled when COMPRESSION_ACCELERATION of raidcom get system -key dedupe_compression is UNAVAILABLE.

- enable: Compression accelerator is enabled.
- disable: Compression accelerator is disabled.

<request #> is the Request ID assigned for each command execution. For details, see [Request ID function \(on page 293\)](#).

[-capacity_saving <capacity saving>]

Specifies the capacity saving setting. If you omit this option, a volume is created with the capacity saving function disabled.

- compression: Enables compression.
- deduplication_compression: Enables deduplication and compression.
- disable: Disables capacity saving.

For VSP One B20, this option must be specified for a volume that is not a journal volume in Universal Replicator. Specify compression or deduplication_compression. Omit this option or specify disable only when creating a volume for a journal volume.

For VSP 5000 series, VSP One B20, and VSP E series, Compressed volumes for accelerated compression, or volumes for deduplication and compression can be created only for pools whose subscription limit is set to 65535 (unlimited). An error occurs if compression or deduplication_compression is specified for a pool whose subscription limit is other than 65535 (unlimited).

[-capacity_saving_mode <saving mode>]

When you enable capacity saving, specifies the post-process mode or inline mode as the capacity saving processing mode for DP-VOLs. For details about the capacity saving processing modes, see the *Provisioning Guide* for the storage system.

- post_process: Sets the post-process mode for capacity saving processing.
- inline: Sets the inline mode for capacity saving processing.

For VSP G1x00, VSP F1500, VSP G200, G400, G600, G800, and VSP F400, F600, F800, post_process is the default. For VSP 5000 series, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, inline is the default.

[-drs -request_id auto]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series

Specify this option to create a data reduction shared volume (DRS-VOL).

<request #> is the Request ID assigned for each command execution. For details, see [Request ID function \(on page 293\)](#).

-ddp_ldev -request_id auto

Specify this option to create an LDEV in a DDP parity group.

-request_id is a request ID assigned for each time the command is run. For details, see [Request ID function \(on page 293\)](#).

Examples

Creating an LDEV: 100 of size 10 GB (gigabytes) in a parity group: 5-2.

Location in the parity group: automatic allocation, LDEV Emulation type: OPEN-V

```
# raidcom add ldev -parity_grp_id 5-2 -ldev_id 100 -capacity 10G
```

Creating an LDEV with the following conditions: Parity group: 5-3, location of the parity group: automatic allocation, Emulation type: 3390-3, Size: 10 M cylinder (Cylinder specification), LDEV: 120.

```
# raidcom add ldev -parity_grp_id 5-3 -ldev_id 120 -cylinder 10m -emulation 3390-3
```

Creating an LDEV in a DDP parity group: 1-2, and then automatically assigning an unused LDEV number.

```
# raidcom add ldev -parity_grp_id 1-2 -ldev_id auto -ddp_ldev -request_id auto
```

Creating an LDEV of position in the external volume group: allocated automatically, emulation type: OPEN-V, External volume group: 01-02, Capacity: 200 MB (megabytes), and LDEV number: 200.

```
# raidcom add ldev -external_grp_id 01-02 -ldev_id 200 -capacity 200m
```

Creating an LDEV of external volume group: 01-03, position in the external volume group: allocated automatically, emulation type: OPEN-V, External volume size: takeover, and LDEV number: 220.

```
# raidcom add ldev -external_grp_id 01-03 -ldev_id 220 -capacity all
```

Creating a V-VOL of Capacity: 300 MB and number: 300, to a pool for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.

```
# raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m
```

Creating a V-VOL of Capacity: 300 MB and LDEV number: 400, to a pool for Thin Image or Copy-on-Write Snapshot.

```
# raidcom add ldev -pool snap -ldev_id 400 -capacity 300m
```

Creating a V-VOL of Capacity: 300 MB, LDEV number: 300, and Full Allocation: enable, to the pool ID: 4 for Dynamic Provisioning.

```
# raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m -status enable_fullallocation
```

Creating a V-VOL (LDEV ID: 44:44) for Dynamic Provisioning associated with the pool volume (LDEV ID: 22:22) that has the data direct mapping attribute in the Dynamic Provisioning pool.

```
#raidcom add ldev -ldev_id 44:44 -mapping_ldev_id 22:22
```

Creating a V-VOL of Capacity: 300 MB, LDEV number: 300, and the T10 PI attribute: valid, in a pool (pool ID: 4) for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.

```
# raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m -t10pi_enable
```

Creating a V-VOL of Capacity: 300 MB, LDEV number: 300, and capacity saving setting: compression, in the pool (ID: 4) for Dynamic Provisioning.

```
# raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m -capacity_saving compression
```

Creating a V-VOL of Capacity: 300 MB, LDEV number: 300, capacity saving setting: compression, and compression accelerator: enable, in the pool (ID: 4) for Dynamic Provisioning.

```
# raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m -capacity_saving compression -
compression_acceleration enable -request_id auto
```

Creating a V-VOL of Capacity: 300 MB, LDEV number: 300, capacity saving setting: deduplication and compression, and capacity saving mode: inline, in the pool (ID: 4) for Dynamic Provisioning.

```
#raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m -capacity_saving
deduplication_compression -capacity_saving_mode inline
```

Creating a V-VOL of capacity: 10 GB for a Dynamic Provisioning or Dynamic Tiering pool (pool ID: 4), and assigning an unused LDEV number automatically.

```
# raidcom add ldev -pool 4 -ldev_id auto -request_id auto -capacity 10G
REQID : 1
```

Creating a DRS-VOL with capacity: 300 MB, LDEV number: 300, and capacity saving setting: deduplication and compression in a Dynamic Provisioning pool (pool ID: 4).

```
#raidcom add ldev -pool 4 -ldev_id 300 -capacity 300m -capacity_saving
deduplication_compression -drs -request_id auto
```



Caution: The following message is shown whenever an LDEV larger than 4 TB is specified. It does not necessarily indicate that there is a problem:

```
raidcom : Please check if LDEV over 4 TB will be used for replication
products, and then check the supported capacity.
```

raidcom delete ldev

Deletes the specified LDEVs or V-VOLs. A device group can also be specified instead of an LDEV.

This command is executed asynchronously with the command input. Check the completion of this process on the raidcom get command_status command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax for deleting an LDEV that does not belong to a DDP parity group

```
raidcom delete ldev {-ldev_id <ldev#> | -grp_opt <group option>
    -device_grp_name <device group name> [<device name>] }
    [-operation initialize_capacity_saving]
```

Syntax for deleting an LDEV belonging to a DDP parity group

```
raidcom delete ldev -parity_grp_id <gno-sgno> -ddp_ldev -request_id auto
```

Options and parameters**-ldev_id <ldev#> (0-65279)**

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

-parity_grp_id <gno-sgno> -ddp_ldev -request_id auto

Specifies the parity group number (gno: 1-52, sgno: 1-32). For example:

3-1

Specify this option to delete an LDEV belonging to a DDP parity group.

-request_id is a request ID assigned each time the command is run. For details, see [Request ID function \(on page 293\)](#).

[-operation initialize_capacity_saving]

Deletes an LDEV for which the capacity saving setting is enabled. Specify this option if either of the following conditions is met:

- The LDEV is a DRS-VOL.
- The LDEV is not a DRS-VOL and CSV_Status displayed by using the **raidcom get ldev** command is ENABLED, ENABLING, REHYDRATING, DELETING, or CONVERTING.

When you specify this option, use the **raidcom get command_status** command to check if the LDEV deletion started, and then use the **raidcom get ldev** command to check if VOL_TYPE of the LDEV is changed from REMOVING to NOT DEFINED. If you perform another operation during or immediately after deletion of the LDEV, the operation might fail. In this case, wait a while and then retry the operation.

(VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) When there are no pool-associated LDEVs for which the deduplicated data is enabled, the deduplication system data volumes are automatically deleted asynchronously with the deletion of the LDEVs. To verify that a deduplication system data volume is deleted, use the **raidcom get pool -key saving** command to check the LDEV number, and use the **raidcom get ldev** command to verify that VOL_TYPE of the LDEV is changed from REMOVING to NOT DEFINED.



Note:

- Deleting a deduplication volume using this command might take a while. Also, more pool/physical capacity might be used. You cannot stop the deletion operation while it is in process.
- If you want to delete all of the pool-associated volumes for which capacity saving is enabled, first execute the **raidcom modify ldev** command to block both the pool-associated volumes and the deduplication system data volumes, and then execute the **raidcom initialize pool** command. By executing the **raidcom initialize pool** command before deleting the volumes, you can shorten the time for the volume deletion process and prevent the increase of pool capacity to be used.

Examples

Deleting an LDEV: 200.

```
raidcom delete ldev -ldev_id 200
```

Deleting an LDEV belonging to the device group: grp1.

```
raidcom delete ldev -grp_opt ldev -device_grp_name grp1
```

Deleting an LDEV belonging to a DDP parity group: 1-2.

```
# raidcom delete ldev -parity_grp_id 1-2 -ddp_ldev -request_id auto
```

Deleting an LDEV: 200, for which capacity saving setting is enabled.

```
# raidcom delete ldev -ldev_id 200 -operation initialize_capacity_saving
```

raidcom extend ldev

Extends the capacity of a V-VOL for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.

If the specified LDEV is not a V-VOL of Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe, the command is rejected with EX_ENOOBJ.

A device group can also be specified instead of an LDEV.

When the `-request_id auto` option is specified, if a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If you specify the `-request_id` option for the asynchronous command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the `raidcom get command_status` command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom extend ldev {-ldev_id <ldev#> | -grp_opt <group option>
    -device_grp_name <device group name> [<device name>]}
    -capacity <size> | -offset_capacity <size> | -cylinder <size> [-request_id auto]
```

Options and parameters

-ldev_id <ldev#> (0-65279)

Specifies the LDEV number (0-65279). Specify LDEVs for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe. For example:

- `-ldev_id 200`

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

-capacity <size>

The increment size of capacity can be specified in bytes or blocks. When specifying in bytes, the unit is t/T (terabyte), g/G (gigabyte), m/M (megabyte), or k/K (kilobyte). If this unit is omitted, block (512 bytes) is used.

Example of specification:

1GB (gigabyte) is:

-capacity 1G, -capacity 1g, -capacity 1024M, -capacity 1024m, -capacity 1048576K, -capacity 1048576k, -capacity 2097152

-offset_capacity <size>

Specifies capacity. This option is only used when using HDvM - SN.

The same correction is performed as when Capacity Compatibility Mode (Offset boundary) is specified on HDvM - SN. The capacity can be specified in bytes or blocks.

-cylinder <size>

Specifies the size in cylinder unit. When specifying, the unit is g/G (gigabyte), m/M (megabyte), or k/K (kilobyte). Cylinder is applied if you do not specify the unit.

When the emulation type for the Open System is specified, this option cannot be specified.

-request_id auto

(VSP 5000 series, VSP G1x00, VSP F1500, VSP One B20, VSP E series, and VSP G130, G/F350, G/F370, G/F700, G/F900 only) Specifying this option is recommended when running the command. Specify this option to expand the capacity of a DRS-VOL. <request#> is a Request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

If you refresh the Device Manager - Storage Navigator window during the capacity expansion of a DRS-VOL, refreshing the window might not be performed. For more information, see the Troubleshooting section in the System Administrator Guide.

Examples

Extending the capacity of a Dynamic Provisioning or Dynamic Provisioning for Mainframe V-VOL: 200 by 10 GB (gigabytes).

```
# raidcom extend ldev -ldev_id 200 -capacity 10G
```

Extending the capacity of a Dynamic Provisioning or Dynamic Provisioning for Mainframe V-VOL: 201 by 200 MB (megabytes).

```
# raidcom extend ldev -ldev_id 201 -capacity 200M
```

Extending the capacity of Dynamic Provisioning or Dynamic Provisioning for Mainframe V-VOL belonging to the device group: grp1 by 200 MB.

```
# raidcom extend ldev -grp_opt ldev -device_grp_name grp1 -capacity 200M
```

Extending the capacity of a Dynamic Provisioning or Dynamic Provisioning for Mainframe V-VOL: 201 by 200 MB.

```
# raidcom extend ldev -ldev_id 201 -capacity 200M -request_id auto
REQID : 1
```

raidcom get ldev

Displays the information of the specified LDEV or the device file. A device group can also be specified instead of an LDEV.

Syntax

```
raidcom get ldev {-ldev_id <ldev#> ... [-cnt <count>] |
  -grp_opt <group option> -device_grp_name <device group
  name> [<device name>] | -ldev_list <ldev list option>}
  [-key <keyword>][{-check_status | -check_status_not}
  <string>... [-time <time>]] [-time_zone <time zone>]
```

Options and parameters

-ldev_id <ldev#> ...

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200
- -ldev_id 100-110
- -ldev_id 100 -cnt 10

[-cnt <count>]

Specifies the count (2-65280).

If this option is omitted, the count is set to one.

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

-ldev_list <ldev list option>

Specifies the type of LDEVs to display. Specify one of the following LDEV list options.
If the LDEV that meets the specified condition does not exist, [EX_ENODEV] No such device is displayed.

- **defined:** Displays all implemented LDEVs.
- **dp_volume:** Displays LDEVs that have Dynamic Provisioning attributes.
Combining this with -pool_id <pool id> specification displays LDEVs that have Dynamic Provisioning attributes related to the specified POOL.
- **external_volume:** Displays external volumes.
- **undefined:** Displays all LDEV numbers that are not implemented.
- **mapped:** Displays all LDEVs to which LU paths are defined.
Combining this with -pool_id <pool id> specification displays LDEVs (defined the LU path) relate to the specified POOL.
- **mapped_nvme:** Displays all LDEVs that are assigned to namespaces.
- **unmapped:** Displays all LDEVs to which LU paths are defined. However, LDEVs, mainframe volumes, and namespaces that are not implemented are not displayed because they cannot specify the LU path.
Combining this with -pool_id <pool id> specification displays LDEVs (not defined the LU path) relate to the specified POOL.
- **journal -journal_id <journal id>:** Displays LDEVs that belong to the specified journal.
- **pool -pool_id <pool id>:** Displays LDEVs that belong to the specified pool.
If a -pool_id option is non-numeric, the specified value is recognized as a pool name to identify the pool ID.
- **parity_grp -parity_grp_id <parity group id>:** Displays LDEVs that belong to the specified parity group.
- **mp_blade -mp_blade_id <mp#>:** Displays LDEVs that are set to the specified MP blades.
- **quorum:** Displays LDEVs set for the quorum disks.
- **clpr -clpr_id <clpr#> :** Displays LDEVs belonging to the specified CLPR.

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G/F350, G/F370, G/F700, G/F900

[-key <keyword>]

Specifies a display keyword.

If this option is omitted, basic LDEV information is displayed. If this option is specified, the following information is displayed. The following display keywords can be specified:

- `front_end`: Front-end information
- `front_end_nvme`: Front-end NVMe (NVMe-oF or NVMe/TCP) information
- `parity_grp`: Parity group information
- `external`: External volume information
- `tier`: Tier information for the Dynamic Tiering or active flash V-VOL. This display keyword can only be specified for VSP 5000 series, VSP G1x00, VSP F1500, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800.
- `software_saving`: Information on the capacity saved by the capacity saving feature for volumes with the capacity saving feature enabled.
- `naa`: The NAA identifier of the LU (LUN WWN)
- `nguid`: Namespace identifier (NGUID: Namespace Global Unique Identifier)
- `qos`: Specifies the qos when displaying the QoS setting information. This display keyword can be specified only for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900.
- `qos_monitor`: Specifies the qos_monitor when displaying the performance monitor information related to the QoS of the target LDEV. This display keyword can be specified only for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900.
- `reservation_holder`: Specify this keyword to display the information about the host that sets a SCSI reservation.
- `cmd`: Specify this keyword to display the command device information.
- `ldev_status`: Specify this keyword to display the LDEV status information.

`[-check_status <string>... [-time <time>]]`

Check if the LDEV is in the same state as the specified in `<string>`. If the option contains multiple states, the OR condition check is performed and verifies that the LDEV is in one of the states contained in the option.

The following strings are specified in the <string>.

- STS
 - NML: Normal
 - BLK: Blocked
 - BSY: Status is changing
- OPE_TYPE
 - FMT: Formatting
 - QFMT: Quick formatting
 - CCOPY: Data copying (Correction copy / Copy back / Drive copy / Dynamic sparing)
 - CACCS: Accessing to collections
 - NONE: Not in operation
 - SHRD: Shredding
 - ZPD: Page discarding
 - SHRPL: Deleting from the pool
 - RLC: Pool relocating
 - RBL: Pool rebalancing
- VOL_TYPE
 - NOT DEFINED: An LDEV is not installed
 - DEFINING: An LDEV is being created
 - REMOVING: An LDEV is being deleted

If "-time" is specified, the status of the LDEV is checked every three seconds until the end of the specified <time> (seconds).

When this option is specified, the returned values are as follows.

- The LDEV is in one of the specified states: 0
- The LDEV is in none of the specified states (without -time option): 1
- The LDEV is in none of the specified states (when the specified <time> passed): EX_EWSTOT

[-check_status_not <string>... [-time <time>]]

Check that the LDEV is not in the same state as the specified in <string>. If the option contains multiple states, the NOR condition check is performed and verifies that the LDEV is not in any of the states contained in the option.

The following strings are specified in the <string>.

- STS
 - NML: Normal
 - BLK: Blocked
 - BSY: Status is changing
- OPE_TYPE
 - FMT: Formatting
 - QFMT: Quick formatting
 - CCOPY: Data copying (Correction copy / Copy back / Drive copy / Dynamic sparing)
 - CACCS: Accessing to collections
 - NONE: Not in operation
 - SHRD: Shredding
 - ZPD: Page discarding
 - SHRPL: Deleting from the pool
 - RLC: Pool relocating
 - RBL: Pool rebalancing
- VOL_TYPE
 - NOT DEFINED: An LDEV is not installed
 - DEFINING: An LDEV is being created
 - REMOVING: An LDEV is being deleted

If "-time" is specified, the status of the LDEV is checked every three seconds until the end of the specified <time> (seconds).

When this option is specified, the returned values are as follows.

- The LDEV is not in any of the specified states: 0
- The LDEV is in one of the specified states (without -time option): 1
- The LDEV is in one of the specified states (when the specified <time> passed): EX_EWSTOT

[-time_zone <time zone>]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the time zone displayed in storage systems. If this option is omitted, the time in the time zone set for the storage system is displayed. This option is valid only when it is specified with the `-key qos` option or `-key qos_monitor` option.

- `utc`: Displays the time in Coordinated Universal Time (UTC).

A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

Some keywords might not be displayed depending on the LDEV attribute as shown below.

LDEV attribute		Front end	Parity group	External	Tier
Volume that is neither a Pool volume nor a Journal volume	Internal volume	Y	Y	N	Y
	External volume	Y	N	Y	Y
Pool volume for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe	Internal volume	N	Y	N	N
	External volume	N	N	Y	N
Journal volume	Internal volume	N	Y	N	N
	External volume	N	N	Y	N
Legend Y: Displayed, N: Not displayed					

Examples

- Internal volume examples
- External volume examples
- Dynamic Provisioning V-VOL examples
- Deduplication system data volume example
- Volume deletion examples
- Dynamic Tiering V-VOL examples
- Thin Image primary volume example
- Pool volume example
- `rmawk` command examples
- Capacity saving example

Internal volume examples

Example 1

Displaying the information of the LDEV number 577 (internal VOL).

```
# raidcom get ldev -ldev_id 577
```

```
Serial# : 63502 PHY_Serial# : 302594
LDEV : 577 PHY_LDEV : 600
SL : 0
CL : 0
VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 2181120
VOL_Capacity (cyl): 22720 (Displayed only for the mainframe volume)
NUM_LDEV : 1
LDEVs : 577
NUM_PORT : 2
PORTs : CL2-E-0 1 Linux_X86 : CL2-E-1 1 Solaris
F_POOLID : NONE
VOL_ATTR : CVS
RAID_LEVEL : RAID1
RAID_TYPE : 2D+2D
NUM_GROUP : 1
RAID_GROUPS : 02-01
CMP : Enable
EXP_SPACE : V

DRIVE_TYPE : DKS2C-K072FC
DRIVE_Capa : 141822798
LDEV_NAMING : Oracle_data_1
STS : NML
OPE_TYPE : QFMT
OPE_RATE : 100
FMT_STS : QFMT
MP# : 2
ASSIGNED_MP# : 2
SSID:001F
ALUA : Enable
RSGID : 0
PWSV_S : -
CL_MIG : Y
```

Example 2

Displaying the NAA identification information of the LDEV number 577 (internal VOL).

```
# raidcom get ldev -ldev_id 577 -key naa
```

```
Serial# : 63502 PHY_Serial# : 302594
LDEV : 577 PHY_LDEV : 600
SL : 0
CL : 0
VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 2181120
VOL_Capacity(cyl) : 22720 (Displayed for the mainframe volume only)
NUM_LDEV : 1
LDEVs : 577
NUM_PORT : 2
PORTs : CL2-E-0 1 Linux_X86 : CL2-E-1 1 Solaris
LUN_WWN : 60060e80160164000001016400000893
F_POOLID : NONE
VOL_ATTR : CVS
RAID_LEVEL : RAID1
RAID_TYPE : 2D+2D
NUM_GROUP : 1
RAID_GROUPS : 02-01
CMP : Enable
EXP_SPACE : V
DRIVE_TYPE : DKS2C-K072FC
DRIVE_Capa : 141822798
LDEV_NAMING : Oracle_data_1
STS : NML
OPE_TYPE : QFMT
OPE_RATE : 100
FMT_STS : QFMT
MP# : 2
ASSIGNED_MP# : 2
SSID : 001F
ALUA : Enable
RSGID : 0
PWSV_S : -
CL_MIG : Y
```

Description of each column in output example 1:**Serial#**

Product serial number.

The serial number of the virtual storage machine is displayed when you specify a volume that is virtualized by the global storage virtualization function using the virtual LDEV number.

PHY_Serial#

If you specify a volume that is virtualized by the global storage virtualization function using the virtual LDEV number, the actual serial number of the storage system is displayed. In other cases, this item is not displayed.

LDEV

LDEV number.

PHY_LDEV

Actual LDEV number of the storage system.

This item is displayed only when you specify a volume that was virtualized by the global storage virtualization function using the virtual LDEV number.

VIR_LDEV

Virtual LDEV number. This item is displayed only when you specify a volume that was virtualized by the global storage virtualization function using the actual LDEV number and it is different from the virtual LDEV number. When the virtual LDEV number is not given to the volume, "FF:FE(65534)" is displayed. When you set the reserve attribute of global-active device to the volume, "FF:FF(65535)" is displayed. This item is displayed only when the volume is virtualized by global storage virtualization, and this item is displayed on the place where PHY_LDEV is displayed in the example (Displaying the information of the LDEV number 577 (internal VOL)). This item and PHY_LDEV are not displayed at the same time.

SL

SLPR information ("0" is displayed for all storage systems except USP V/VM).

CL

CLPR information.

VOL_TYPE

Emulation type of the relevant LDEV. Displays the same name as the product ID of the Inquiry command. For open-systems LDEVs, a string is appended to the emulation type to indicate the LDEV attribute:

- *n: LUN Expansion (LUSE)
- -CVS: Virtual LVI/LUN
- -A: ALU
- -S: SLU
- -CM: command device



Note: You can check the emulation type of a volume that was virtualized by the global storage virtualization function by running the command specifying the following instances:

- The instance for which the command device belonging to the same virtual storage machine as the virtualized volume is specified for HORCM_CMD in the configuration definition file
- The instance for which the serial number of the virtual storage machine that the virtualized volume belongs to is defined for HORCM_VCMD in the configuration definition file

If the LDEV is not installed or is in the process of being created or deleted, the status is one of the following:

- NOT DEFINED: An LDEV is not installed.
- DEFINING: An LDEV is being created.
- REMOVING: An LDEV is being deleted.

VOL_Capacity (BLK)

Capacity of LDEV in block size.

VOL_Capacity(cyl)

Capacity of LDEV in cylinder size. Displayed only when the attribute is MF-VOL.

NUM_LDEV

Number of LDEVs that compose the LU where the specified LDEV belongs.

LDEVs

Number of LDEVs in the LU.

NUM_PORT

Number of ports defined to the paths for relevant LDEV. When the relevant LDEV is a mainframe volume or a multiplatform volume without the LU path definition, the number of ports for the dummy LU is displayed. For details about the dummy LU, see the *Command Control Interface User and Reference Guide*.

PORTs

Ports defined to the paths for the relevant LDEV. It lists up the ports defined to the paths for the relevant LDEV. When the relevant LDEV is a mainframe volume or a multiplatform volume without the LU path definition, the port number of the dummy LU is displayed.

If the number of characters for the host group name is 17 or more, a maximum of 16 characters can be displayed. Execute the `raidcom get host_grp` command to display the host group name with 17 characters or more.

LUN_WWN

The NAA identifier (LUN WWN) of the first LU defined for the LDEV. A hyphen (-) is displayed if no value can be displayed for LUN_WWN. For details about the cases where LUN WWNs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#). This information is displayed only when the `-key naa` option is specified.

F_POOLID

Pool ID if the LDEV is a component of the pool. In other cases, "NONE" is displayed.

VOL_ATTR

Attributes of the LDEV.

- CMD: Command device (open-systems only)
- CLUN: Cache LUN (DCR)
- CVS: CVS volume
- LUSE: LUSE volume
- ALUN: Volume Migration volume
- ELUN: External volume
- OLG: Open LDEV Guard volume
- VVOL: V-VOL
- HORC: The remote copy pair volumes:
 - TrueCopy/TrueCopy for Mainframe, Universal Replicator/Universal Replicator for Mainframe, global-active device: P-VOL or S-VOL
- MRCF: ShadowImage volume (P-VOL or S-VOL)
- QS: Thin Image or Copy-on-Write Snapshot volume (P-VOL or S-VOL)
- JNL: JNL volume
- HDP: volume for Dynamic Provisioning or Dynamic Provisioning for Mainframe
- HDT: volume for Dynamic Tiering (DT), Dynamic Tiering for Mainframe, active flash, or active flash for mainframe
- POOL: POOL volume
- QRD: Quorum disk
- ENCD: Encryption disk
- SYSD: System disk
- TSE: Dynamic Provisioning for Mainframe volumes that are used in FlashCopy SE
- ESE: a virtual volume capable of page release by the User Directed Space Release function
- GAD: Global-active device volume
- HNASS: a volume used as the system LU of Hitachi NAS
- HNASU: a volume used as a user LU of Hitachi NAS
- MG: a volume for the data migration
- T10PI: a volume of which T10 PI is enabled
- DSD: Deduplication system data volume
- DS: Deduplication system data volume that stores duplicated data (data store)

- RCMD: Remote command device
- DRS: Data reduction shared volume
- VCP: a volume with the vClone Parent attribute
- VC: a volume with the vClone attribute

RAID_LEVEL

RAID level (RAID1, RAID5, or RAID6).

RAID_TYPE

Configuration of the drives.

NUM_GROUP

Number of parity groups to which the relevant LDEV belongs.

RAID_GROUPS

Parity groups to which the relevant LDEVs belongs.

CMP

Accelerated compression setting for the parity group to which the LDEV belongs.

This item is displayed only for the VSP Fx00 models and VSP Gx00 models.

- Enable: Accelerated compression is enabled.
- Disable: Accelerated compression is disabled.
- NA: Accelerated compression is not supported.
- - (hyphen): This information is invalid for this LDEV.

For VSP G1x00, VSP F1500, a hyphen (-) is always displayed.

EXP_SPACE

Displays whether the LDEV uses the accelerated compression area of the parity group.

This item is displayed only for the VSP Fx00 models and VSP Gx00 models.

- V: The LDEV uses the accelerated compression.
- R: The LDEV does not use accelerated compression.
- - (hyphen): This information is invalid for this LDEV.

For VSP G1x00, VSP F1500, a hyphen (-) is always displayed.

DRIVE_TYPE

Drive type code which is set when the parity group to which the LDEV belongs is set.

This drive type code is a typical code of a drive that is compatible with drives contained in the parity group, and therefore might not match the actual code of the drive contained in the parity group.

To view the drive type code of the drive in the parity group, execute the **raidcom get drive** command.

DRIVE_Capa

Capacity of relevant HDD in the number of block (512 bytes) (decimal number).

LDEV_NAMING

Nickname of the LDEV.

STS

Status of the LDEV:

- NML: Normal
- BLK: Blocked
- BSY: Status is changing
- NONE: unknown state (not supported)

OPE_TYPE

Current operation.

- FMT: Formatting
- QFMT: Quick formatting
- CCOPY: Data copying (Correction copy / Copy back / Drive copy / Dynamic sparing)
- CACCS: Accessing to collections
- NONE: Not in operation
- SHRD: Shredding
- ZPD: Page Discarding
- SHRPL: Deleting from the pool
- RLC: Pool relocating
- RBL: Pool rebalancing

OPE_RATE

Progress of the format or shred operation. When the status is other than formatting or shredding, 100 is displayed. If the process ended abnormally, "BSY" is displayed on the "STS".

FMT_STS

Status of the format, quick format, or shredding operation for the LDEV.

This item is not displayed if the storage system does not support the display of this information.

The OPE_TYPE option cannot be used to determine the actual LDEV status because the OPE_TYPE option includes the information of the standby task to be run as an asynchronous command. This information is used to determine the actual LDEV status.

- FMT: Formatting
- QFMT: Quick formatting
- QFMTP: Preparing quick format
- SHRD: Shredding
- - (hyphen): Other than above

MP#

MP blade ID.

ASSIGNED_MP#

Assigned MP blade ID

SSID

Storage subsystem ID number (hexadecimal). VSP One B20, VSP E series, VSP Gx00 models, VSP Fx00 models, and Unified Storage VM do not support SSID but display the specified value.



Note: You can check the SSID of a volume that was virtualized by the global storage virtualization function by running the command specifying the following instances:

- The instance for which the command device belonging to the same virtual storage machine as the virtualized volume is specified for HORCM_CMD in the configuration definition file
- The instance for which the serial number of the virtual storage machine that the virtualized volume belongs to is defined for HORCM_VCMD in the configuration definition file

ALUA

ALUA mode

- Enable: The ALUA mode is enabled.
- Disable: The ALUA mode is disabled.

RSGID

Displays the resource group ID of the resource group to which the LDEV belongs.

CL_MIG

Displays whether the allocation of LDEV to CLPR is changed.

- Y: The allocation of LDEV is being changed.
- N: The allocation of LDEV is not changed.

Example 3

Displaying front-end information of the LDEV number 577

```
# raidcom get ldev -ldev_id 577 -key front_end
```

```
Serial# LDEV# SL CL VOL_TYPE VOL_Cap(BLK) PID ATTRIBUTE Ports
PORT_No:LU#:GRPNAME ...
63502      577  0  0 OPEN-V-CVS 2181120 - CVS 2
CL2-E-0:1:Linux_X86 CL2-E-1:1:Solaris
```

Description of each column in output example 3**PID**

Pool ID for a virtual volume. If the volume is not a virtual volume, a hyphen (-) is displayed.

ATTRIBUTE

LDEV attribute. Same as VOL_ATTR.

Example 4

Displaying front-end nvme (NVMe-oF) information of the LDEV number 577

```
# raidcom get ldev -ldev_id 577 -key front_end_nvme
Serial# LDEV# SL CL VOL_TYPE VOL_Cap(BLK) PID ATTRIBUTE NVMS_ID
NS_ID Ports PORT_No
63502      577  0  0 OPEN-V-CVS 2181120 - CVS 1
1 2 CL2-E CL2-F
```

Description of each column in output example 4**NVMS_ID**

NVM subsystem ID.

NSID

Namespace ID.

PORT_No

NVM subsystem port number.

If the NVM subsystem port number is not assigned, a hyphen (-) is displayed.

Example 5

Displaying back-end (parity group) information of the LDEV number 577.

```
# raidcom get ldev -ldev_id 577 -key parity_grp
```

```
Serial# LDEV# SL CL PID ATTRIBUTE R_LVL RAID_TYPE DRV_TYPE DRV_Cap
```

```
GRPs RAID_GRP ...
63502 577 0 0 - CVS RAID1 2D+2D DKS2C-K072FC 141822798 1 02-01
```

Example 6

Check if the LDEV number 577 is in normal status.

```
# raidcom get ldev -ldev_id 577 -check_status NML
```

Example 7

Check if the LDEV number 577 is in blocked status.

```
# raidcom get ldev -ldev_id 577 -check_status BLK
```

Example 8

Set in wait status for 30 seconds until the formatting process of LDEV number 577 is complete.

```
# raidcom get ldev -ldev_id 577 -check_status_not FMT -time 30
raidcom:[EX_EWSTOT]Timeout waiting for specified status
```

Refer to the command log(/HORCM/log0/horcc_rmhost.log) for details.

(The command fails if the FMT does not change (is not complete) in 30 seconds.)

External volume examples**Example 1**

Displaying the information of the LDEV number 160 (external VOL).

```
# raidcom get ldev -ldev_id 160
Serial# : 63502
LDEV : 160
SL : 0
CL : 0
VOL_TYPE : OPEN-V
VOL_Capacity(BLK) : 4385280
NUM_PORT : 0
PORTs :
F_POOLID : NONE
VOL_ATTR : ELUN : RCMD

E_VendorID : HITACHI
E_ProductID : OPEN-V
E_VOLID :
4849544143484920523530304638304530303641000000000000000000000000000000
E_VOLID_C : HITACHI R500F80E006A.....
NUM_E_PORT : 1
E_PORTs : CL2-G-0 0 50060e8004f80e34
```

```

LDEV_NAMING : Oracle_data_1
STS : NML
OPE_TYPE : QFMT
OPE_RATE : 70
FMT_STS : QFMT
MP# : 2
ASSIGNED_MP# : 2
SSID:001F
ALUA : Enable
RSGID : 0
CL_MIG : Y

```

Description of columns in output example 1:

E_VendorID

Vendor name that the external storage system's volume notifies to the host.

E_ProductID

System name that the external storage system's volume notifies to the host.

E_VOLID

Number to identify the volume on the external storage system (hexadecimal).

E_VOLID_C

Number to identify the volume on the external storage system (ASCII display).

NUM_E_PORT

Number of alternate paths.

E_PORTS:

List of defined alternate paths.

Example 2

If the LDEV number 160 is an external volume, its back-end (RAID group) information is displayed.

```
# raidcom get ldev -ldev_id 160 -key external
```

```

Serial#  LDEV#  SL CL PID ATTRIBUTE  E_VendorID  E_ProductID
E_VOLID
"E_VOLID_C"                                E_PORTS  PORT_NO:LU#:WWN ...
  63502    160   0  0   - ELUN             HITACHI    OPEN-V
48495441434849205235303046383045303036410000000000000000000000000000 "HITACHI
R500F80E006A....."                      1 CL2-G-0:0:50060e8004f80e34

```

Example 3

Displaying the information of the LDEV number 39320 (quorum disk).

```
# raidcom get ldev -ldev_id 39320
```

```
Serial#   : 302656
LDEV      : 39320
SL        : 0
CL        : 0
VOL_TYPE  : OPEN-V-CVS
VOL_Capacity(BLK) : 31457280
NUM_PORT  : 0
PORTs     :
F_POOLID  : NONE
VOL_ATTR  : CVS : ELUN : QRD
E_VendorID : HITACHI
E_ProductID : OPEN-V
E_VOLID   : 48495441434849203530333030413536313230300000000000000000000000000000
E_VOLID_C : HITACHI 50300A561200.....
NUM_E_PORT : 1
E_PORTS   : CL1-B-0 0 50060e80070a5630
LDEV_NAMING :
STS       : BLK
OPE_TYPE  : NONE
OPE_RATE  : 100
FMT_STS   : -
MP#       : 0
ASSIGNED_MP# : 2
SSID      : 004D
QRDID     : 31
QRP_Serial# : 302646
QRP_ID    : R8
ALUA      : Disable
RSGID     : 0
CL_MIG    : Y
```

Description of a column in output example 2 and 3**QRDID**

ID of the quorum disk for HAM or GAD.

QRP_Serial#

Shows the serial number of the storage system when the external volume is the quorum disk for HAM or GAD.

QRP_ID

Shows the ID for identifying storage systems when the external volume is the quorum disk for HAM or GAD.

- R9: VSP 5000 series
- R8: VSP G1x00 and VSP F1500
- VSP One B20
- M8: VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, and VSP F400, F600, F800

Dynamic Provisioning V-VOL examples

Example

Displaying the information of the LDEV number 4368 (V-VOL of Dynamic Provisioning).

```
# raidcom get ldev -ldev_id 4368
```

```
Serial#   : 302614
LDEV      : 4368
SL        : 0
CL        : 0
VOL_TYPE  : OPEN-V-CVS
VOL_Capacity(BLK) : 16777216
NUM_PORT  : 0
PORTs     :
F_POOLID  : NONE
VOL_ATTR  : CVS : HDP
B_POOLID  : 0
S_POOLID  : 6
LDEV_NAMING :
STS       : BLK
D_PM      : E
D_PM(%)   : 50
OPE_TYPE  : NONE
OPE_RATE  : 100
FMT_STS   : -
MP#       : 0
ASSIGNED_MP# : 2
SSID      : 0006
Used_Block(BLK) : 0

FLA(MB)   : Disable
RSV(MB)   : 0
CSV_Status : ENABLING
CSV_PROGRESS(%) : 60
CSV_Mode   : DEDUP+COMPRESS
COMPRESSION_ACCELERATION : ENABLED
```

```

COMPRESSION_ACCELERATION_STATUS : ENABLED
CSV_PROCESS_MODE : POST_PROCESS
DEDUPLICATION_DATA : ENABLED
ALUA : Disable
RSGID : 0
DM_LDEV : 4096
PARENT LDEV#: 5
PWSV_S:PS
CL_MIG : Y

```

Description of columns in output example:

B_POOL ID

Pool ID to which the LDEV is associated.

S_POOLID

Snapshot pool ID to which the LDEV is associated. This item is displayed only for the S-VOL whose LDEV attribute is HDP and QS. When S_POOLID is displayed, B_POOL ID indicates the ID of the pool for Dynamic Provisioning, Dynamic Tiering, or active flash.

D_PM

Processing status of the pool for each virtual volume

If multiple statuses coexist, the statuses are displayed separated by commas (,). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

- N: Processing of the pool is not running.
- WS: The pool is waiting to shrink.
- S: The pool is shrinking.
- WRB: The pool is waiting to rebalance.
- RB: The pool is rebalancing.
- E: The pool is being expanded due to DDP expansion.
- PE: The pool expansion due to DDP expansion has been suspended.
- - (hyphen): The information is not available for this pool.

D_PM(%)

Percentage of progress of the processing status of the pool for each virtual volume

If multiple statuses coexist, the percentage of progress for each status is displayed separated by a comma (,). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported or if the percentage of progress is not valid.

Used_Block(BLK)

Number of blocks used in the pool. This number of blocks includes the reserved blocks by Full Allocation. When Full Allocation is changed to enabled from disabled, the value of the Used_Block(BLK) will be increased by an amount equal to the number of blocks which are reserved by Full Allocation.

FLA(MB)

Shows the capacity (megabytes) which is reserved by Full Allocation or Proprietary Anchor. When Full Allocation is disabled, "Disable" is displayed.

RSV(MB)

Shows the capacity (megabytes) that is reserved by Full Allocation or Proprietary Anchor.

CSV_Status

Shows the status of capacity saving:

- DISABLED: Capacity saving is disabled.
- ENABLED: Capacity saving is enabled.
- ENABLING: Capacity saving is being enabled.
- REHYDRATING: Capacity saving is being disabled.
- DELETING: A volume which has capacity saving enabled is being deleted.
- FAILED: Consistency of the data in the deduplication system data volume cannot be guaranteed.
- CONVERTING: Compression method is being changed.

CSV_PROGRESS(%)

Shows the progress rate when the capacity saving status is ENABLING, REHYDRATING, or DELETING. For other statuses, a hyphen (-) is displayed.



Note: When the status is REHYDRATING, post-processing takes time, and a hyphen (-) might be displayed for CSV_PROGRESS(%) for a while.

CSV_Mode

Shows the capacity saving setting.

- DISABLED: Capacity saving is disabled.
- COMPRESS: Compression is enabled.
- DEDUP+COMPRESS: Deduplication and compression are enabled.

CSV_PROCESS_MODE

Shows the capacity saving processing mode.

- POST_PROCESS: Post-process mode.
- INLINE: Inline mode.
- - (hyphen): Capacity saving is disabled.

COMPRESSION_ACCELERATION

Shows whether the compression accelerator setting is enabled or disabled. DKCMAIN microcode versions that do not support compression accelerator display - (hyphen).

- ENABLED: Compression accelerator is enabled.
- DISABLED: Compression accelerator is disabled.
- - (hyphen): This information is invalid for this LDEV.

COMPRESSION_ACCELERATION_STATUS

Shows whether compression accelerator is applied or not to the data in the volume. DKCMAIN microcode versions that do not support compression accelerator display - (hyphen).

- ENABLED: Compression accelerator is applied to all data on the volume.
- DISABLED: Compression accelerator is not applied to any data on the volume.
- HYBRID: The volume contains data that is compressed with compression accelerator and data that is compressed without using compression accelerator.
- - (hyphen): This information is invalid for this LDEV.

DEDUPLICATION_DATA

Shows whether the deduplicated data exists on a virtual volume.

- DISABLED: The deduplicated data does not exist on a virtual volume.
- ENABLED: The deduplicated data exists on a virtual volume.

If DEDUP+COMPRESS is selected as the capacity saving setting for a virtual volume, ENABLED is shown even when the amount of used capacity is 0. Also, when the capacity saving setting is being disabled and the capacity saving status for a virtual volume is DISABLED, ENABLED is shown when deduplicated data exists.

If the capacity saving setting for a virtual volume is changed from deduplication_compression to compression with the following DKCMAIN versions, ENABLED is shown:

- VSP One B20: A3-04-01-xx/xx or later
- VSP E series: 93-06-88-xx/xx or later and earlier than 93-07-01-xx/xx, or 93-07-26-xx/xx or later
- VSP 5000 series: 90-08-89-xx/xx or later and earlier than 90-09-01-xx/xx, or 90-09-27-xx/xx or later
- - (hyphen): The volume is not a virtual volume, or the DKCMAIN version does not support the deduplication function.

DM_LDEV

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Shows the LDEV number of the pool volume associated with a V-VOL for Dynamic Provisioning that has the data direct mapping attribute. All LBAs of the V-VOL for Dynamic Provisioning are mapped to LBAs of the pool volume one-on-one. DM_LDEV is displayed only when the LDEV has the data direct mapping attribute.

PARENT LDEV#

If the LDEV is a volume with the vClone attribute, a volume with the vClone Parent attribute, or a volume of a Thin Image Advanced pair, displays the LDEV number of a parent (a volume with the vClone Parent attribute, a root volume with the vClone Parent attribute enabled, or a root volume) in a family to which the LDEV belongs.

This item is not displayed if the LDEV is not a volume with the vClone attribute, a volume with the vClone Parent attribute, or a volume of a Thin Image Advanced pair, or if the DKCMAIN microcode version is not supported.

A hyphen (-) is displayed if all of the following conditions are met:

- HORCM_VCMD is defined in the configuration definition file. Alternatively, a command device that belongs to a virtual storage machine is defined in HORCM_CMD.
- The LDEV displayed in LDEV and the LDEV displayed in this item belong to different virtual storage machines.

Example

Displaying the command device information about the LDEV number 16 (V-VOL of Dynamic Provisioning).

```
#raidcom get ldev -ldev_id 16 -I1 -key cmd
Serial# : 800089
LDEV : 16
SL : 0
CL : 0
VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 2097152
NUM_PORT : 2
PORTs : CL1-A-1 16 allhosts : CL2-A-1 16 allhosts
F_POOLID : NONE
VOL_ATTR : CVS : DP
CMD_SECURITY : -
CMD_USER_AUTHENTICATION : -
CMD_DEVICE_GROUP_DEFINITION : -
CMP : -
EXP_SPACE : -
B_POOLID : 1
LDEV_NAMING :
STS : NML
OPE_TYPE : NONE
OPE_RATE : 100
FMT_STS : -
MP# : 0
ASSIGNED_MP# : 0
SSID : 0004
Used_Block(BLK) : 0
FLA(MB) : Disable
RSV(MB) : 0
CSV_Status : DISABLED
```

```

CSV_PROGRESS(%) : -
CSV_Mode : DISABLED
COMPRESSION_ACCELERATION : -
COMPRESSION_ACCELERATION_STATUS : -
CSV_PROCESS_MODE : -
DEDUPLICATION_DATA : DISABLED
ALUA : Disable
RSGID : 0
PWSV_S : -
CL_MIG : N

```

Description of columns in output example:

CMD_SECURITY

Displays the command device security. This item is not displayed if displaying this information is not supported on the storage system:

- ENABLED: The command device security is enabled.
- DISABLED: The command device security is disabled.
- - (hyphen): This volume is not a command device. (A hyphen (-) is also displayed for the remote command device attribute.)

CMD_USER_AUTHENTICATION

Displays the user authentication settings for the command device. This item is not displayed if displaying this information is not supported on the storage system:

- ENABLED: The user authentication for the command device is enabled.
- DISABLED: The user authentication for the command device is disabled.
- - (hyphen): This volume is not a command device. (A hyphen (-) is also displayed for the remote command device attribute.)

CMD_DEVICE_GROUP_DEFINITION

Displays the device group definition settings for the command device. This item is not displayed if displaying this information is not supported on the storage system:

- ENABLED: The definition setting for the command device is enabled.
- DISABLED: The definition setting for the command device is disabled.
- - (hyphen): This volume is not a command device. (A hyphen (-) is also displayed for the remote command device attribute.)

Dynamic Provisioning V-VOL namespace examples

Examples

Displaying the information about the LDEV number 4368 (namespace of Dynamic Provisioning V-VOL) (VSP 5000 series and VSP E1090).

```
# raidcom get ldev -ldev_id 4368
```

```
Serial# : 302614
LDEV : 4368
SL : 0
CL : 0
VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 16777216
NUM_PORT : 0
PORTs : -
NVM_NUM_PORT : 2
NVM_PORTs : CL1-A CL2-B
NSID : 2
NVMSS_ID : 1
F_POOLID : NONE
VOL_ATTR : CVS : HDP
B_POOLID : 0
S_POOLID : 6
LDEV_NAMING :
STS : BLK
OPE_TYPE : NONE
OPE_RATE : 100
FMT_STS : -
MP# : 0
ASSIGNED_MP# : 2
SSID : 0006
Used_Block(BLK) : 0
FLA(MB) : Disable
RSV(MB) : 0
CSV_Status : ENABLING
CSV_PROGRESS(%): 60
CSV_Mode : DEDUP+COMPRESS
COMPRESSION_ACCELERATION : ENABLED
COMPRESSION_ACCELERATION_STATUS : ENABLED
CSV_PROCESS_MODE : POST_PROCESS
DEDUPLICATION_DATA : ENABLED
ALUA : Disable
RSGID : 0
DM_LDEV:4096
PWSV_S : PS
CL_MIG : Y
```

Displaying the NGUID of the LDEV number 4368 (namespace of Dynamic Provisioning V-VOL) (VSP 5000 series and VSP E1090).

```
# raidcom get ldev -ldev_id 4368 -key nguid
```

```
Serial# : 302614
LDEV : 4368
SL : 0
CL : 0
VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 16777216
NUM_PORT : 0
PORTs : -
NVM_NUM_PORT : 2
NVM_PORTs : CL1-A CL2-B
NSID : 2
NVMSS_ID : 1
LDEV_NGUID : 12345...
F_POOLID : NONE
VOL_ATTR : CVS : HDP
B_POOLID : 0
S_POOLID : 6
LDEV_NAMING :
STS : BLK
OPE_TYPE : NONE
OPE_RATE : 100
FMT_STS : -
MP# : 0
ASSIGNED_MP# : 0
SSID : 0006
Used_Block(BLK) : 0
FLA(MB) : Disable
RSV(MB) : 0
CSV_Status : ENABLING
CSV_PROGRESS(%): 60
CSV_Mode : DEDUP+COMPRESS
COMPRESSION_ACCELERATION : ENABLED
COMPRESSION_ACCELERATION_STATUS : ENABLED
CSV_PROCESS_MODE : POST_PROCESS
DEDUPLICATION_DATA : ENABLED
ALUA : Disable
RSGID : 0
DM_LDEV:4096
PWSV_S : PS
CL_MIG : Y
```

Description of columns in output example:**NVM_NUM_PORT**

Number of NVM subsystem ports to which the LDEV is defined.

NVM_PORTs

NVM subsystem port number to which the LDEV is defined.

LDEV_NGUID

Namespace identifier (NGU ID). A hyphen (-) is displayed if the specified LDEV is used with a protocol other than NVMe-oF or NVMe/TCP. "Unknown" is displayed if no value can be displayed for NGUID for other reasons. For details about the cases where NGU IDs cannot be displayed, see [Displaying NAA identifiers \(LUN WWNs\) and namespace identifiers \(NGU IDs\) \(on page 21\)](#). This item is not displayed for DKCMAIN microcode versions that do not support NVMe-oF or NVMe/TCP.

NSID

Namespace ID.

NVMSS_ID

NVM subsystem ID.

Example

Displaying the information about the host that sets a SCSI reservation on the LU of the LDEV number 4368

```
#raidcom get ldev -ldev_id 4368 -key reservation_holder
Serial#  LDEV#  PORT  LUN  WWN                      RSV_KEY
400007   4368   CL1-A-1  1    1234567890123456  1234567890123456
```

Description of columns in output example:**PORT**

Displays the host group for the LU set with the SCSI reservation in the port-host-group-ID format.

LUN

Displays the ID of the LU set with the SCSI reservation.

WWN

Displays the WWN of the host that sets the SCSI reservation. A hyphen (-) is displayed if the host uses iSCSI to set the SCSI reservation.

RSV_KEY

Displays the key information about the host that sets the SCSI reservation if the type of the SCSI reservation is persistent reservation. A hyphen (-) is displayed for the following cases:

- The type of the SCSI reservation is not persistent reservation.
- Persistent reservation is in progress.

For details about the types of SCSI reservations, see the “Description of each column in output example” in [raidcom get lun \(on page 500\)](#).

In the following cases, the information about the hosts that set SCSI reservations is not displayed.

- No SCSI reservation is set on any LU of any LDEV number.
- The DKCMAIN microcode version does not support display of the information about hosts that set SCSI reservations.



Caution:

- Use the `raidcom get lun` command to check whether SCSI reservations are set. If the information about the host that sets a SCSI reservation is not displayed even if the SCSI reservation is set on the LU, check the DKCMAIN microcode version. The following DKCMAIN microcode versions support display of the information about hosts that set SCSI reservations (`-key reservation holder` option):
 - For VSP 5000 series:
 - 90-09-21-xx/xx or later
 - 90-09-02-xx/xx or later and earlier than 90-09-21-xx/xx
 - 90-08-84-xx/xx or later and earlier than 90-09-01-xx/xx
 - 90-08-64-xx/xx or later and earlier than 90-08-81-xx/xx
 - For VSP G1x00, VSP F1500:
 - 80-06-94-xx/xx or later
 - For VSP One B20:
 - All versions
 - For VSP E series:
 - 93-07-21-xx/xx or later
 - 93-07-02-xx/xx or later and earlier than 93-07-21-xx/xx
 - 93-06-83-xx/xx or later and earlier than 93-07-01-xx/xx
 - 93-06-64-xx/xx or later and earlier than 93-06-81-xx/xx
 - For VSP G130, G/F350, G/F370, G/F700, G/F900:
 - 88-08-11-xx/xx or later
 - For VSP G100, G200, G400, G600, G800, and VSP F400, F600, F800:
 - 83-06-21-xx/xx or later, and 83-05-49-xx/xx or later and earlier than 83-06-01-xx/xx
- When a host sets a SCSI reservation on a volume contained in a GAD pair, the volume targeted by the host's SCSI reservation command and the volume paired with the target volume are reserved.
- For GAD pair volumes, the `raidcom get ldev` command displays the information about the host that sets a SCSI reservation only when you specify the volume for which the host issues the SCSI reservation command.

- To view the information about GAD pair volumes, specify the volume for which the host issues the SCSI reservation command when you execute the **raidcom get ldev** command.
- If you cannot identify the volume for which the host issued the SCSI reservation command, execute the **raidcom get ldev** command for both the P-VOL and the S-VOL of the target GAD pair.
- If a host sets a SCSI reservation on an external volume whose cache mode is set to cache through mode, the **raidcom get ldev** command does not display the information about the host that sets the SCSI reservation.

Examples

Displaying the LDEV status information about the LDEV number 4368.

```
#raidcom get ldev -ldev_id 4368 -key ldev_status
Serial# LDEV# FMT_STS
400007 4368 QFMT
```

Description of columns in output example:

FMT_STS

Displays the LDEV status of formatting, quick formatting, or shredding:

- FMT: Formatting the LDEV is in process.
- QFMT: Quick formatting the LDEV is in process.
- QFMTP: Quick formatting the LDEV is being prepared.
- SHRD: Shredding the LDEV is in process.
- - (hyphen): The LDEV is in a different status or displaying this information is not supported on the storage system.

Deduplication system data volume example

Example

Displaying the information of the LDEV number 4368 (deduplication system data volume).

```
# raidcom get ldev -ldev_id 4368
```

```
Serial#   : 302614
LDEV     : 4368
SL       : 0
CL       : 0
VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 16777216
NUM_PORT : 0
PORTs    :
F_POOLID : NONE
```

```

VOL_ATTR : CVS : HDP : DSD
B_POOLID : 0
LDEV_NAMING :
STS : BLK
OPE_TYPE : NONE
OPE_RATE : 100
FMT_STS : -
MP# : 0
ASSIGNED_MP# : 2
SSID : 0006
Used_Block(BLK) : 0
FLA(MB) : Disable
RSV(MB) : 0
CSV_Status : DISABLED
CSV_PROGRESS(%): -
CSV_Mode : DISABLED
COMPRESSION_ACCELERATION : -
COMPRESSION_ACCELERATION_STATUS : -
CSV_PROCESS_MODE : -
DEDUPLICATION_DATA : DISABLED
ALUA : Disable
RSGID : 0
CL_MIG : Y

```

Volume deletion examples

Example 1

Displaying the information of the LDEV being deleted (VOL_ATTR is other than HDP).

```
# raidcom get ldev -ldev_id 4096
```

```

Serial# : 64568
LDEV : 4096
SL : -
CL : -
VOL_TYPE : REMOVING
SSID : 0005
RSGID : 0

```

Example 2

Displaying the information of the LDEV being deleted by the raidcom delete ldev command without the `-operation initialize_capacity_saving` option (VOL_ATTR is HDP).

```
# raidcom get ldev -ldev_id 4096
```

```
Serial# : 64568
LDEV : 4096
SL : -
CL : -
VOL_TYPE : REMOVING
CSV_Status : DISABLED
CSV_PROGRESS(%) : -
SSID : 0005
RSGID : 0
```

Example 3

Displaying the information of the LDEV being deleted by the raidcom delete ldev command with the `-operation initialize_capacity_saving` option (VOL_ATTR is HDP).

```
# raidcom get ldev -ldev_id 4096
```

```
Serial# : 64568
LDEV : 4096
SL : -
CL : -
VOL_TYPE : REMOVING
CSV_Status : DELETING
CSV_PROGRESS(%) : 30
SSID : 0005
RSGID : 0
```

Dynamic Tiering V-VOL examples**Example 1**

Displaying the information of the LDEV number 640 (V-VOL of Dynamic Tiering).

```
# raidcom get ldev -ldev_id 640
```

```
Serial# : 63502
LDEV : 640
SL : 0
CL : 0
```

```

VOL_TYPE : OPEN-V-CVS
VOL_Capacity(BLK) : 2181120
NUM_LDEV : 1
LDEVs : 640
NUM_PORT : 1
PORTs : CL2-E-0 14 Linux_X86
F_POOLID : NONE
VOL_ATTR : CVS : HDP : HDT
B_POOLID : 5
LDEV_NAMING : Oracle_data_1
STS : NML
OPE_TYPE : NONE
OPE_RATE : 100
MP# : 2
FMT_STS : -
SSID : 001F
ASSIGNED_MP# : 2
Used_Block(BLK) : 218112
TIER_Relocation : Enable
TIER_LEVEL: 6
TIER#1(MB) : 1120
TIER#2(MB) : 3000
:
:
TIER_Alloc_level : H
TIER#1_Alloc_rate : MAX : 50 : MIN : 30
TIER#3_Alloc_rate : MAX : 50 : MIN : 30
FLA(MB) : 980
RSV(MB) : 980
CSV_Status : DISABLED
CSV_PROGRESS(%): -
CSV_Mode : DISABLED
COMPRESSION_ACCELERATION : -
COMPRESSION_ACCELERATION_STATUS : -
CSV_PROCESS_MODE : -
DEDUPLICATION_DATA : DISABLED
ALUA : Enable
RSGID : 0
CL_MIG : Y

```

Description of columns in output example 1:

TIER_Relocation

Status of the relocation setting:

- Enable: The status where the relocation is enabled
- Disable: The status where the relocation is disabled

TIER_LEVEL

Shows the level of tiering policy that is used for reallocation.

- all: Shows that all tiers of the pool to which the relevant LDEV is allocated are being used.
- 1-5: Shows the tiering policy level that is configured to the relevant LDEV.
- 6-31: Shows the tiering policy (customized policy (1-26)) that is configured to the relevant LDEV. For details, see the *Provisioning Guide* or *LUN Manager User Guide* for the storage system.

TIER#n(MB)

Shows the capacity (megabytes) allocated to each tier.

TIER_Alloc_level

Shows the tier level of the new mapped page.

- H: High
- M: Middle
- L: Low

TIER#1_Alloc_rate

Shows the Tier1 Max or Min value that is set in the tiering policy.

TIER#3_Alloc_rate

Shows the Tier3 Max or Min value that is set in the tiering policy.

FLA(MB)

Shows the capacity (megabytes) that is reserved by Full Allocation. When Full Allocation is disabled, "Disable" is displayed.

Example 2

Displays the tier information about the LDEV number 640 (V-VOL of Dynamic Tiering).

```
# raidcom get ldev -ldev_id 640 -key tier
```

```
Serial# LDEV#  SL  CL VOL_TYPE VOL_Cap(BLK)  PID ATTRIBUTE VOL_Used(BLK)  TR TL T#1(MB)
T#2(MB)  ...
63502  640  0  0 OPEN-V-CVS 204800 13 CVS|2181120 - CVS|VVOL|HDT 218112 E 6 1120 3000
```

Description of each column in output example 2:**TR**

Displays the enabled or disabled of the tier relocation.

- E: Enabled
- D: Disabled

TL

Displays the tier level.

- 0: ALL
- 1-5: Level (1-5)
- 6-31: Customized policy (1-26)

T#x(MB)

Displays the allocated volume of the relevant LDEV for Tier x. Displays up to 5 tiers. If there is no tier, "0" is displayed.

Thin Image primary volume example

Example

Displays the information about the LDEV number 1000 (primary volume of Thin Image).

```
# raidcom get ldev -ldev_id 1000
```

```
Serial#   : 64568
LDEV      : 1000
SL        : 0
CL        : 0
VOL_TYPE  : OPEN-V-CVS
VOL_Capacity(BLK) : 204800
NUM_PORT  : 2
PORTs     : CL1-A-0 0 1A-G00 : CL5-B-0 1 5B-G00
F_POOLID  : NONE
VOL_ATTR  : CVS : QS
RAID_LEVEL : RAID1
RAID_TYPE  : 2D+2D
NUM_GROUP : 1
RAID_GROUPS : 01-02
DRIVE_TYPE : DKR2G-K146SS
DRIVE_Capa : 285177528
LDEV_NAMING :
STS       : NML
OPE_TYPE  : NONE
OPE_RATE  : 100
FMT_STS   : -
MP#       : 2
ASSIGNED_MP# : 2
SSID      : 0009
ALUA      : Enable
RSGID     : 0
Snap_Used_Pool(MB) : 100
CL_MIG    : Y
SNAP_USED(MB) : 126
```

```

SNAP_GARBAGE(MB) : 15
DELETING_SNAP_GARBAGE : NONE
DELETING_SNAP_GARBAGE(%) : -

```

Description of column in output example:

Snap_Used_Pool(MB)

Displays the capacity (megabytes) used for the snapshot data out of the snapshot capacity assigned from the pool. If the used capacity is less than 1 MB, the displayed value is rounded up.

- For a volume with the vClone Parent attribute, displays the snapshot used capacity of the family to which the volume with the vClone Parent attribute belongs.
- For a Thin Image root volume with the vClone Parent attribute or vClone attribute disabled, displays the snapshot used capacity of the snapshot tree to which the volume belongs.
- In other cases, this information is not displayed.

SNAP_USED(MB)

For the Thin Image (TI only) root volume, displays the capacity (megabytes) in the root volume for the snapshot data area. If the used capacity is less than 1 MB, the displayed value is rounded up. For a Thin Image Advanced pair, a hyphen (-) or zero is always displayed.

SNAP_GARBAGE(MB)

Displays the garbage data volume (megabytes) in the snapshot data area of the specified Thin Image root volume. If the garbage data volume is less than 1 MB, the displayed value is rounded up. While the snapshot garbage data is being deleted (defragmentation is in progress), a hyphen (-) is displayed. This item is displayed only for VSP 5000 series. For other models, this item is not displayed, or a hyphen (-) is always displayed. For a Thin Image Advanced pair, a hyphen (-) is always displayed.

DELETING_SNAP_GARBAGE

Displays the execution status of the snapshot garbage data deletion process (defragmentation process).

This item is displayed only for VSP 5000 series. For other models, this item is not displayed, or a hyphen (-) is always displayed.

- NONE: Not in operation
- PROCESSING: Processing is in progress.

The execution status is PROCESSING even when processing is suspended.

See the Hitachi Thin Image User Guide for the conditions when the processing is suspended.

- STOPPING: Defragmentation is being stopped (transition from PROCESSING to NONE).

For a Thin Image Advanced pair, a hyphen (-) is always displayed.

DELETING_SNAP_GARBAGE(%)

Displays the progress rate of the snapshot garbage data deletion processing (defragmentation processing). This item is displayed only for VSP 5000 series. For other models, this item is not displayed, or a hyphen (-) is always displayed. For a Thin Image Advanced pair, a hyphen (-) is always displayed.

Pool volume example

Example

Displays the information about the LDEV number 4096 (pool volume).

```
# raidcom get ldev -ldev_id 4096
```

```
Serial#   : 64568
LDEV      : 4096
SL        : 0
CL        : 0
VOL_TYPE  : OPEN-V-CVS
VOL_Capacity(BLK) : 16777216
NUM_LDEV  : 1
LDEVs     : 4096
NUM_PORT  : 0
PORTs     :
F_POOLID  : 127
VOL_ATTR  : CVS : POOL
RAID_LEVEL : RAID5
RAID_TYPE  : 3D+1P
NUM_GROUP : 1
RAID_GROUPS : 01-01
DRIVE_TYPE : DKR2G-K146SS
DRIVE_Capa : 285177528
LDEV_NAMING :
STS       : BLK
OPE_TYPE  : NONE
OPE_RATE  : 100
FMT_STS   : -
MP#       : 0
ASSIGNED_MP# : 2
SSID      : 0005
ALUA      : Disable
RSGID     : 0
DM_LDEV   : 640
CL_MIG    : Y
MAPPED_PAGE_CAP(PAGE) : 470
VOL_EFFECTIVE(PAGE) : 500
```

Description of column in output example:**DM_LDEV**

Shows the LDEV number of the V-VOL for Dynamic Provisioning associated with a pool volume that has the data direct mapping attribute. When any V-VOL for Dynamic Provisioning is not associated with the pool volume, "NONE" is displayed. DM_LDEV is displayed only when the LDEV has the data direct mapping attribute.

MAPPED_PAGE_CAP(PAGE)

Displays the value corrected for each page of the allocated capacity of a pool volume in pages. This item is displayed only for the pool volume. This item is not displayed if displaying this information is not supported on the storage system.

VOL_EFFECTIVE(PAGE)

Displays the effective pool volume capacity. This item is displayed only when the storage system supports displaying this information and for the pool volume.

rmawk command examples**Example 1**

Displays the information of the used Tier 1 capacity for the pool ID 73 of the DT volume (using the **rmawk** command).

```
# raidcom get ldev -ldev_list dp_volume -pool_id 73 -key tier |rmawk -EC @L-ne:0 @l2?
=ad:@l2 -n exe="print Total = @l2?"

Total = 8064
```

Example 2

Displays the information of each used Tier capacity for the pool ID 73 of the DT volume (using the **rmawk** command).

```
# raidcom get ldev -ldev_list dp_volume -pool_id 73 -key tier |rmawk -EC @L-ne:0
@l2=ad:@l2? @l3=ad:@l3? @l4=ad:@l4? exe="print @0" -n exe="print " exe="print Total
= T#1(MB): @l2? T#2(MB): @l3? T#3(MB): @l4?"
```

Serial#	LDEV#	SL	CL	VOL_TYPE	VOL_Cap(BLK)	PID	ATTRIBUTE	VOL_Used(BLK)	TR	TL
T#1(MB)	T#2(MB)	...								
64558	29440	0	0	OPEN-V-CVS	4042752	73	CVS HDP HDT	404 2752	E	5
0	0		1974	0	0					
64558	29441	0	0	OPEN-V-CVS	10924032	73	CVS HDP HDT	1092 4032	E	6
5334	0		0	0	0					
64558	29442	0	0	OPEN-V-CVS	10924032	73	CVS HDP HDT	1092 4032	E	13
1596	1638		2100	0	0					
64558	29443	0	0	OPEN-V-CVS	10924032	73	CVS HDP HDT	1092 4032	E	18
1134	3654		546	0	0					

```
Total = T#1(MB):8064 T#2(MB):5292 T#3(MB):4620
```

Example 3

Displays information of the LDEV (when VOL_ATTR is other than HDP) that is being deleted.

```
# raidcom get ldev -ldev_id 4096
Serial# : 64568
LDEV : 4096
SL : -
CL : -
VOL_TYPE : REMOVING
SSID : 0005
RSGID : 0
```

Example 4

Displays information of the LDEV that is being deleted by the **raidcom delete ldev** command without **-operation initialize_capacity_saving** (when VOL_ATTR is other than HDP).

```
# raidcom get ldev -ldev_id 4096
Serial# : 64568
LDEV : 4096
SL : -
CL : -
VOL_TYPE : REMOVING
CSV_Status : DISABLED
CSV_PROGRESS(%) : -
SSID : 0005
RSGID : 0
```

Example 5

Displays information of the LDEV that is being deleted by the **raidcom delete ldev** command without **-operation initialize_capacity_saving** (when VOL_ATTR is other than HDP).

```
# raidcom get ldev -ldev_id 4096
Serial# : 64568
LDEV : 4096
SL : -
CL : -
VOL_TYPE : REMOVING
CSV_Status : DELETING
CSV_PROGRESS(%) : 30
SSID : 0005
RSGID : 0
```

Capacity saving examples

If the specified LDEV is a volume with the vClone attribute or an S-VOL of a Thin Image Advanced pair, no information is displayed. The capacity saving effect information on the volume with the vClone attribute or the S-VOL of the Thin Image Advanced pair is aggregated and displayed under a LDEV of a parent (a volume with the vClone Parent attribute, a root volume with the vClone Parent attribute enabled, or a root volume) in a family to which the LDEV belongs.

Example

Displays the information on the capacity saved by the capacity saving feature for the LDEV number 18 (volume with the capacity saving feature enabled).

For details about the capacity saved by the capacity saving feature, see the *Provisioning Guide* for your storage system.

```
# raidcom get ldev -ldev_id 18 -key software_saving
```

LDEV#	TLS_R	TOTAL_SAVING (BLK)	CMP (BLK)	DDP (BLK)	
RECLAIM (BLK)		SYSTEM (BLK)	PRE_USED (BLK)	POOL_USED (BLK)	UNCMP (BLK)
18	4.00	15728640	6291456	8388608	
2097152		1048576	20971520	5242880	111

Description of columns in output example:**TLS_R**

Displays the ratio of the amount of data reduced by the capacity saving feature against the capacity before reduction.

For example, if the capacity before and after reduction is 4.00 to 1, 4.00 is displayed.

TOTAL_SAVING(BLK)

Displays the capacity reduced by the capacity saving feature in blocks. The capacity includes the amount of zero data reduction, metadata, and garbage data.

CMP(BLK)

Displays the capacity reduced by compression of the capacity saving feature in blocks. The reduced capacity does not include the amount of metadata and garbage data.

DDP(BLK)

Displays the capacity reduced by deduplication of the capacity saving feature in blocks. The reduced capacity does not include the amount of metadata and garbage data.

RECLAIM(BLK)

Displays the capacity reduced by reclaiming the fixed data pattern using the capacity saving feature in blocks. The reduced capacity does not include the amount of metadata and garbage data.

SYSTEM(BLK)

Displays the amount of consumed system data (the amount of metadata and garbage data) of the capacity saving feature in blocks. The amount of metadata capacity and garbage data for deduplication system data volumes is not included .

PRE_USED(BLK)

Displays the capacity before capacity saving of data to be deleted by the capacity saving feature in blocks.

POOL_USED(BLK)

Displays the used capacity of the pool volumes used by volumes.

For VSP G1x00, VSP F1500, a hyphen (-) is always displayed.

UNCMP(BLK)

Displays the amount of data that cannot be compressed. If the storage system does not support displaying this information, a hyphen (-) is displayed.

QoS examples

Displays the QoS setting information of LDEV number 200 (VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900)

Nothing is displayed if QoS is not set. If the acquisition of QoS setting information fails, "NA" is displayed for items other than LDEV number.

```
# raidcom get ldev -ldev_id 200 -key qos
LDEV# UPPER_THROUGHPUT(Iops) UPPER_DATA_TRANS(MBps) UPPER_ALERT(s)
UPPER_ALERT_NOTICE LOWER_THROUGHPUT(Iops) LOWER_DATA_TRANS(MBps)
LOWER_ALERT(s) LOWER_ALERT_NOTICE PRIORITY RESPONSE_TARGET(ms)
RESPONSE_ALERT(s) RESPONSE_ALERT_NOTICE
200 80 100 30
2019-07-31T10:15:20 20
- 0 2019-07-31T10:15:20 1
80 100 2019-07-31T10:15:20
```

Description of each column in output example:**UPPER_THROUGHPUT(IOps)**

Displays the upper limit of the throughput.

For more information on the upper limit of the throughput per second, see the *Performance Guide*.

UPPER_DATA_TRANS(MBps)

Displays the upper limit of the data transfer volume.

For more information on the upper limit of the data transfer volume, see the *Performance Guide*.

UPPER_ALERT(s)

Displays the following alert notification time (second):

- The alert notification time (second) when the number of received commands continuously exceeds the upper limit of the throughput per second.
- The alert notification time (second) when the transfer volume of received commands continuously exceeds the upper limit of the data transfer volume per second.

UPPER_ALERT_NOTICE

Displays the last alert time of the following alert.

- The alert when the number of received commands continuously exceeds the upper limit of the throughput per second.
- The alert when the transfer volume of received commands continuously exceeds the upper limit of the data transfer volume per second.

If the last alert time does not exist, a hyphen (-) is displayed.

LOWER_THROUGHPUT(IOps)

Displays the lower limit of the throughput.

For more information on the lower limit of the throughput per second, see the *Performance Guide*.

LOWER_DATA_TRANS(MBps)

Displays the lower limit of the data transfer volume.

For more information on the lower limit of the data transfer volume per second, see the *Performance Guide*.

LOWER_ALERT(s)

Displays the following alert notification times (second):

- The alert notification time (second) when the number of received commands has not continuously reached the lower limit of the throughput per second.
- The alert notification time (second) when the transfer volume of received commands has not continuously reached the lower limit of the data transfer volume per second.

LOWER_ALERT_NOTICE

Displays the last alert time of the following alerts:

- The alert when the number of received commands has not continuously reached the lower limit of the throughput per second.
- The alert when the transfer volume of received commands has not continuously reached the lower limit of the data transfer volume per second.

A hyphen (-) is displayed if the last alert time does not exist.

PRIORITY

Displays the I/O processing priority. 1 represents the lowest priority and a greater number represents a higher priority.

RESPONSE_TARGET(ms)

Displays the target response time.

RESPONSE_ALERT(s)

Displays the alert notification time (second) when the target response time is not reached.

If the target response time has not been continuously reached for longer than the alert notification time, an alert is issued.

RESPONSE_ALERT_NOTICE

Displays the last alert time when the target response time has not been continuously reached.

A hyphen (-) is displayed if the last alert time does not exist.

Displays I/Os monitoring information of the target LDEV (VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900)

Nothing is displayed if all the monitoring information is invalid. If the acquisition of QoS monitoring information fails, "NA" is displayed for items other than LDEV number.

```
# raidcom get ldev -ldev_id 200 -key qos_monitor
LDEV# RCV_CMD (IOps) RCV_CMD_TRANS (KBps) THROUGHPUT (IOps)
DATA_TRANS (KBps) RESPONSE (us) MONITOR_TIME
200          3840          15          2650
10           5    2019-07-31T10:15:20
```


Description of each column in output example:**RCV_CMD(IOps)**

Displays the number of received commands.

RCV_CMD_TRANS(KBps)

Displays the data transfer volume of received commands.

THROUGHPUT(IOps)

Displays the throughput.

DATA_TRANS(KBps)

Displays the data transfer volume.

RESPONSE(us)

Displays the response time.

MONITOR_TIME

Displays the monitoring collection time.

raidcom initialize ldev

Formats LDEVs. You can specify Quick Format, Normal Format, or Shredding.

A device group can also be specified instead of an LDEV.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

**Caution:**

- Formatting a deduplication volume might take a while. Also, more pool/physical capacity might be used. You cannot stop the formatting operation while it is in process.
- If you want to format all of the pool-associated volumes for which capacity saving is enabled, first execute the **raidcom modify ldev** command to block the deduplication system data volumes, and then execute the **raidcom initialize pool** command. By executing the **raidcom initialize pool** command before formatting the volumes, you can shorten the time for the formatting operation and prevent the increase of pool capacity to be used.

**Caution:**

- Make sure to run the `raidcom initialize ldev` command after the completion of the asynchronous commands including the `raidcom initialize ldev` command. The status of the asynchronous commands can be verified by using the `raidcom get command_status` command. For more information about whether the command is run either synchronously or asynchronously with the command input, see each command section.

If the `raidcom initialize ldev` command is run before the completion of the asynchronous commands, the formatting or shredding operation might not be complete.

- If the Normal Format operation is performed by using CCI, the following operations might be in the standby status until the formatting is complete:
 - Run the different asynchronous commands.
 - Perform the operations by using VSP One Block Administrator.
 - Perform the operations by using Storage Advisor Embedded.
 - Perform the operations by using REST API.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Note: A volume is formatted asynchronously. This allows you to format multiple V-VOLs in parallel. On the other hand, if you format a V-VOL whose defined capacity is 15 TiB or more and whose capacity saving is enabled, the processing might time out if the total defined capacity of the V-VOLs to be formatted in parallel exceeds 1000 TiB. For this reason, make sure to avoid overlapping processing within the same time, such as by staggering formatting times so that the total defined capacity is 1000 TiB or less. You can verify that the formatting is in process or complete by using the `raidcom get command_status` command in CCI.

If the V-VOL is used as an S-VOL in one of the following pairs, formatting the V-VOL whose capacity saving is enabled might cause the processing to time out.

- TrueCopy
- Universal Replicator
- global-active device
- Thin Image (clone attribute)
- Thin Image Advanced

Therefore, make sure to format the V-VOL after changing the pair status to PAIR.

If a timeout error occurs during the volume formatting, a SIM is reported. If only SIM=410002 is reported, reformat the V-VOL that remains blocked. If any other SIMs are reported, contact customer support to take actions.

Syntax

```
raidcom initialize ldev {-ldev_id <ldev#>
  | -grp_opt <group option> -device_grp_name <device group
  name> [<device name>]} -operation <type>
```

Syntax for stopping the formatting in process

```
raidcom initialize ldev -operation stop_fmt
```

Options and parameters

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

-operation <type>

Instructs the operation.

The following operations can be specified. For more information about the supported formatting types, see the *System Administrator Guide* and the *Provisioning Guide* for the storage system.

- fmt: Normal Format
- qfmt: Quick Format

- **shrd [<pattern>]: Shredding**

Formats the LDEV three times according to the shredding pattern which is specified to "pattern".

If <pattern> is specified, format the LDEV three times according to the following order. Specify "pattern" in the range of 0x00000000 to 0xFFFFFFFF.

- 0x00000000
- The specified shredding pattern
- 0x00000000

If <pattern> is omitted, format the LDEV three times according to the following order.

- 0x00000000
- 0xFFFFFFFF
- 0x00000000

- **stop:** Stops shredding. The processing for all LDEVs stops. However, normal and quick format processing cannot be stopped.
- **stop_fmt:** Stops all normal formatting and parity-group formatting operations in process. This option is specified only for VSP One B20. The LDEV number and the device group name cannot be specified at the same time as specifying this option.

Quick Format cannot be stopped. When running this command, lock all storage resource groups. Use the following procedure to run this command:

1. Lock the resource group by using the **raidcom lock resource** command.
2. Run this command.
3. Wait for five minutes for the formatting to be stopped.
4. If the formatting is not stopped after waiting for five minutes, go back to step 2.
5. Release the resource group lock by using the **raidcom unlock resource** command.

Examples

Performing Quick Format for an LDEV: 200.

```
# raidcom initialize ldev -operation qfmt -ldev_id 200
```

Performing Quick Format for an LDEV belonging to the device group: grp1.

```
# raidcom initialize ldev -operation qfmt -grp_opt ldev
-device_grp_name grp1
```

Performing Normal Format for an LDEV: 200.

```
# raidcom initialize ldev -operation fmt -ldev_id 200
```

Performing Shredding (Pattern: 0x55aa55aa) for an LDEV: 200.

```
# raidcom initialize ldev -operation shrd 0x55aa55aa -ldev_id 200
```

Stopping Shredding.

```
# raidcom initialize ldev -operation stop -ldev_id 200
```

Stopping all LDEV formatting operations in process.

```
# raidcom lock resource
# raidcom initialize ldev -operation stop_fmt
# rmawk exe="raidcom get ldev -ldev_list defined -key ldev_status " @3-ne:FMT -AT
timeout=300
# raidcom unlock resource
```

raidcom modify ldev

Changes the following LDEV attributes:

- Blocking LDEV and restoring LDEV.
- Setting LDEV nickname.
- Setting MP blade ID of LDEV.
- Setting the Tiering policy, the new page assignment tier, or enabling or disabling of the tier relocation for the Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe V-VOL. (Only for VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800)
- Instructing the page discarding of V-VOL for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.
- Setting the SSID.
- Setting the command device attribute.
- Setting the quorum disk and releasing setting of the quorum disk.
- Setting capacity saving.
- Setting the capacity saving processing mode (post process or inline).
- Setting compression accelerator (enable or disable) (only for VSP 5000 series, VSP One B20, and VSP E1090).
- Setting QoS (only for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900).
- ESE-VOL setting

When blocking LDEV, restoring LDEV, setting the quorum disk, releasing setting of the quorum disk, setting capacity saving, setting compression accelerator, or setting QoS, this command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax for specifying LDEV ID

```
raidcom modify ldev -ldev_id <ldev#> {-status <status> [<level>]
  [-forcible -password <One Time Password>]
  | -ldev_name <ldev naming> | -mp_blade_id <mp#> | -ssid <value> |
  -command_device {y | n} [Security value] | -quorum_enable <serial#>
  <id> -quorum_id <quorum id> | -quorum_disable |
  -alua {enable|disable} | -capacity_saving <capacity saving>
  | -capacity_saving_mode <saving mode>| -compression_acceleration {enable |
  disable} -request_id auto | -upper_throughput_io<upper throughput io> -request_id
  auto | -upper_data_trans_mb <upperdata trans mb> -request_id auto | -upper_alert_time
  <upper alert time> -request_id auto | -lower_throughput_io <lower throughput io> -
  request_id
  auto | -lower_data_trans_mb <lower data trans mb> -request_id auto | -
  lower_alert_time <lower alert time> -request_id auto | -
  response_priority <priority> -request_id auto | -response_alert_time
  <response alert time> -request_id auto | -ese {enable | disable} -request_id auto}
```

Syntax for specifying pool ID

```
raidcom modify ldev -pool {<pool ID#> | <pool naming>} -compression_acceleration
{enable | disable} -request_id auto
```

Options and parameters

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-status <status> [<level>]

Specifies the LDEV status. The following LDEV statuses can be specified.

- nml: Changes the LDEV status to Normal
- blk: Changes the LDEV status to Blockade

- {enable_reallocation [<level>]|enable_relocation [<level>}: Relocation of LDEV (V-VOL for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe) enabled. Defines the tier that is used for reallocation depending on the value of the level.
 - all: Uses all tiers in the pool.
 - 1-5: Specifies the level of the tier for use. For details, see the *Provisioning Guide* for your storage system.

This parameter can only be specified for VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800.

- disable_reallocation | disable_relocation: Relocation of LDEV (V-VOL for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe) disabled

This parameter can only be specified for VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800.

- discard_zero_page: Discarding 0 page of LDEV(V-VOL) for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.
- stop_discard_zero_page (only for VSP 5000 series, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900): Stops the discard of LDEV (virtual volume) pages for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.
- enable_relocation_policy <policy_id>: Enables the LDEV relocation for Dynamic Tiering or active flash, and sets the level or customized policy of the Tier to be used for the relocation by Policy ID.

<policy_id>:

- all: Uses all tiers in the pool.
- 1-5: Specifies the level of the tier for use.
- 6-31: Specifies the customized policy of the tier for use.

This parameter can only be specified for VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800.

- new_page_allocation: Sets the Tier when new page is allocated to the LDEV for Dynamic Tiering or active flash (V-VOL). The value is one of high/middle/low.

This parameter can only be specified for VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800.

- `enable_fullallocation`: Enables Full Allocation. If all areas equivalent to the sum of the pool capacities the specified volume requires can be reserved, it is guaranteed that all areas of DP-VOL are writable. Full Allocation cannot be enabled for a volume whose capacity saving setting is compression, or deduplication and compression.
- `disable_fullallocation`: Disables Full Allocation.

For example:

- `-status nml`

[-forcible -password <One Time Password>]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specify this option with a one-time password to restore an LDEV status forcibly without considering data consistency. For details about the password, contact customer support. This option is valid only if the `-status nml` option is specified.

-ldev_name <ldev naming>

Specifies the LDEV nickname (maximum 32 characters).

-mp_blade_id <mp#>

Specifies the MP blade ID (0-15).

Do not change the MP blade ID during initial copying or I/O processing of TrueCopy, ShadowImage, Universal Replicator, or global-active device.

To change the MP blade ID again for the same LDEV, wait for more than 30 minutes after changing the MP blade ID. For example:

```
-mp_blade_id 2
```

When you change the MP blade ID where the LDEV is allocated, both before and after the changing should be executed during as off-peak hours for the rate of write pending data of the MP blade as possible. It is recommended to execute when the rate of write pending data of all CLPRs is less than 50%.

Do not change the MP blade ID in regard to a lot of LDEVs at the same time. The number of LDEVs and I/O workload for which you can change the MP blade ID at the same time is lower than 10% of the total number of LDEVs where the same MP blade ID is allocated as a guideline.

-ssid <value>

Specifies SSID (hexadecimal number).

When specifying SSID, specify not only unallocated SSID but also LDEV ID. In this case, LDEV ID must be ID for the undefined LDEV in the area where SSID is not allocated. For example:

```
-ssid 0x1234 -ldev_id 200
```




Note: This option is for enterprise storage systems. You do not need to specify SSID for VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, or VSP G200, G400, G600, G800 and VSP F400, F600, F800, but this option is enabled.

-command_device {y | n} [Security value]

Configures command device attribute.

y or yes: Command device attribute enabled.

n or no: Command device attribute disabled.

Specifies the value of command device security (0-7).

You can specify 0-7 to Security value as the command device security setting:

- 0: Security: OFF, User authentication: OFF, Group information acquisition: OFF
- 1: Security: OFF, User authentication: OFF, Group information acquisition: ON
- 2: Security: OFF, User authentication: ON, Group information acquisition: OFF
- 3: Security: OFF, User authentication: ON, Group information acquisition: ON
- 4: Security: ON, User authentication: OFF, Group information acquisition: OFF
- 5: Security: ON, User authentication: OFF, Group information acquisition: ON
- 6: Security: ON, User authentication: ON, Group information acquisition: OFF
- 7: Security: ON, User authentication: ON, Group information acquisition: ON

-quorum_enable <serial#> <id>

Sets quorum disk for global-active device configuration. You must also set the -quorum_id parameter.

serial#: Serial number (of the migration source storage system)

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

id: The identifier of the source storage system

- R900: VSP 5000 series
- R800: VSP G1x00, VSP F1500
- M800: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, VSP F400, F600, F800

-quorum_id <quorum id>

Specifies the quorum ID for setting the quorum disk.

-quorum_disable

Releases setting of the quorum disk.

-alua {enable|disable}

Specifies the ALUA mode. You must enable ALUA mode only when you use ALUA by global-active device.

- enable or yes: The ALUA mode is enabled.
- disable or no: The ALUA mode is disabled.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

[-capacity_saving <capacity saving>]

Specifies the capacity saving setting:

- disable: Capacity saving is disabled.
- compression: Compression is enabled.
- deduplication_compression: Deduplication and compression are enabled.

disable cannot be specified for a DRS-VOL. compression cannot be specified for a V-VOL whose capacity saving setting is deduplication_compression. However, you can specify compression for a V-VOL whose capacity saving setting is deduplication_compression with the following DKCMAIN versions:

- VSP One B20: A3-04-01-xx/xx or later
- VSP E series: 93-06-88-xx/xx or later and earlier than 93-07-01-xx/xx, or 93-07-26-xx/xx or later
- VSP 5000 series: 90-08-89-xx/xx or later and earlier than 90-09-01-xx/xx, or 90-09-27-xx/xx or later

(VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) When there are no pool-associated LDEVs for which the deduplicated data is enabled, the deduplication system data volumes are automatically deleted asynchronously with the deletion of the LDEVs. To verify that a deduplication system data volume is deleted, use the **raidcom get pool -key saving** command to check the LDEV number, and use the **raidcom get ldev** command to verify that VOL_TYPE of the LDEV is changed from REMOVING to NOT DEFINED.



Note: Disabling the capacity saving setting for a deduplication volume might take a while to complete due to its data extension operation. Also, more pool/physical capacity might be used. You cannot stop the operation of disabling the capacity saving setting while it is in process.

For VSP 5000 series, VSP One B20, and VSP E series, compressed volumes for accelerated compression, or volumes for deduplication and compression can be created only for pools whose subscription limit is set to 65535 (unlimited). An error occurs if compression or deduplication_compression is specified for a pool whose subscription limit is other than 65535 (unlimited).

[-capacity_saving_mode <saving mode>]

Specifies post-process mode or inline mode as the capacity saving processing mode for DP-VOLs. For details about the capacity saving processing modes, see the *Provisioning Guide* for the storage system.

- post_process: Sets the post-process mode for capacity saving processing.

- inline: Sets the inline mode for capacity saving processing.

For VSP G1x00, VSP F1500, VSP G200, G400, G600, G800, and VSP F400, F600, F800 post_process is the default. For VSP 5000 series, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, inline is the default.

-upper_throughput_io <upperthroughput io> -request_id auto

Specifies the upper limit of the throughput per second. To disable it, specify 0. This option is available only for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900. For more information on the upper limit of the throughput per second, see the *Performance Guide*.

The request ID is output in the following format when the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

-upper_data_trans_mb <upper datatrans mb> -request_id auto

Specifies the upper limit of the data transfer volume in megabytes (MB) per second. To disable it, specify 0. This option is available only for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900. For more information on the upper limit of the data transfer volume (MB) per second, see the *Performance Guide*.

The request ID is output in the following format when the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

-upper_alert_time <upper alerttime> -request_id auto

Specifies the following alert notification time (second). To disable it, specify 0. This option is available only for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900. If either of the following conditions has continued for the alert notification time or longer, an alert is issued.

- The alert notification time (second) when the number of received commands continuously exceeds the upper limit of the throughput per second.
- The alert notification time (second) when the transfer volume of received commands continuously exceeds the upper limit of the data transfer volume per second.

The request ID is output in the following format when the execution of the **raidcom modify ldev** command ends:

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

-lower_throughput_io <lowerthroughput io> -request_id auto

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Specifies the lower limit of the throughput per second. To disable it, specify 0. For more information on the lower limit of the throughput per second, see the *Performance Guide*.

The request ID is output in the following format when the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed.

For details, see [Request ID function \(on page 293\)](#).

- lower_data_trans_mb <lower data trans mb> -request_id auto

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Specifies the lower limit of the data transfer volume in megabytes (MB) per second. To disable it, specify 0. For more information on the lower limit of the data transfer volume (MB) per second, see the *Performance Guide*.

The request ID is output in the following format when the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed.

For details, see [Request ID function \(on page 293\)](#).

- lower_alert_time <lower alert time> -request_id auto

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Specifies the following alert notification time (second). To disable it, specify 0.

- The alert notification time (second) when the number of received commands has not continuously reached the lower limit of the throughput per second.
- The alert notification time (second) when the transfer volume of received commands has not continuously reached the lower limit of the data transfer volume per second.

For more information on the lower limit of the alert notification time per second, see the *Performance Guide*.

The request ID is output in the following format when the execution of the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed.

For details, see [Request ID function \(on page 293\)](#).

-response_priority <priority> -request_id auto

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Specifies the I/O processing priority. The valid ranges are 1-3. To

disable it, specify 0. For more information on the I/O processing priority, see the *Performance Guide*.

The request ID is output in the following format when the execution of the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed.

For details, see [Request ID function \(on page 293\)](#).

-response_alert_time <responsealert time> -request_id auto

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Specifies the alert notification time (second) when the target response time is not reached . To disable it, specify 0. If the target response time has not been continuously reached for more than the alert notification time, an alert is issued. For more information on the alert notification time (second) when the target response time is not reached , see the *Performance Guide*.

The request ID is output in the following format when the execution of the **raidcom modify ldev** command ends.

- REQID : <request#>

Where <request#> is a request ID assigned each time the command is executed.

For details, see [Request ID function \(on page 293\)](#).

-compression_acceleration {enable | disable} -request_id auto

Supported storage systems:

- VSP 5000 series
- VSP One B26 and VSP One B28
- VSP E1090

Specifies whether compression accelerator is enabled or disabled.

If you change -capacity_saving from disable to compression or deduplication_compression, and omit this this option, it will be enabled when COMPRESSION_ACCELERATION of raidcom get system -key dedupe_compression is AVAILABLE. It will be disabled when COMPRESSION_ACCELERATION of raidcom get system -key dedupe_compression is UNAVAILABLE.

- enable: Compression accelerator is enabled.
- disable: Compression accelerator is disabled.

<request #> is the Request ID assigned for each command execution. For details, see [Request ID function \(on page 293\)](#).

-ese {enable | disable} -request_id auto

Supported storage systems:

- VSP 5000 series
- VSP E1090

Specifies whether to set or cancel the ESE attribute for the LDEV.

- enable: Set the ESE attribute
- disable: Cancel the ESE attribute

<request#> is the request ID assigned each time the command is run. For details, see [Request ID function \(on page 293\)](#)

Examples

Restoring the LDEV: 200.

```
# raidcom modify ldev -status nml -ldev_id 200
```

Blocking the LDEV: 200.

```
# raidcom modify ldev -status blk -ldev_id 200
```

Assigning an LDEV nickname: my_volume to LDEV: 200

```
# raidcom modify ldev -ldev_id 200 -ldev_name my_volume
```

Setting the LDEV owner MP blade ID of LDEV: 200 to 2

```
# raidcom modify ldev -ldev_id 200 -mp_blade_id 2
```

Enabling relocation of LDEV (Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe V-VOL): 200

```
# raidcom modify ldev -ldev_id 200 -status enable_reallocation
```

Disabling relocation of LDEV (Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe V-VOL): 200

```
# raidcom modify ldev -ldev_id 200 -status disable_reallocation
```

Discarding the zero page of LDEV (Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe V-VOL): 200

```
# raidcom modify ldev -ldev_id 200 -status discard_zero_page
```

Specifying 0x1234 to LDEV SSID of LDEV: 200.

**Note:**

Specify undefined LDEV number: 200 in the area that is not allocated SSID, and assign new SSID: 0x1234. If you specify an LDEV number that is already defined or an SSID that is already registered, an error occurs.

```
# raidcom modify ldev -ssid 0x1234 -ldev_id 200
```

Enabling LDEV command device attribute of LDEV: 200. Specifies 2 to the value of command device security.

```
# raidcom modify ldev -command_device y 2 -ldev_id 200
```

Disabling LDEV command device attribute of LDEV: 200.

```
# raidcom modify ldev -command_device n -ldev_id 200
```

Setting the relocation of LDEV for LDEV:200 by the customized policy 6.

```
# raidcom modify ldev -ldev_id 200 -status enable_relocation_policy 6
```

Setting the tier from which the new mapped page of LDEV: 200 is allocated to High.

```
# raidcom modify ldev -ldev_id 200 -status new_page_allocation high
```

Setting LDEV: 200 as the quorum disk whose quorum ID is 10.

```
# raidcom modify ldev -ldev_id 200 -quorum_enable 65384 R800 -quorum_id 10
```

Enabling the ALUA mode for LDEV: 200.

```
# raidcom modify ldev -ldev_id 200 -alua enable
```

Setting the capacity saving setting to compression for LDEV: 200.

```
# raidcom modify ldev -ldev_id 200 -capacity_saving compression
```

Enabling the compression accelerator setting for LDEV: 200.

```
# raidcom modify ldev -ldev_id 200 -compression_acceleration enable -request_id auto
```

Enabling all compression accelerator settings associated with pool ID: 1.

```
# raidcom modify ldev -pool 1 -compression_acceleration enable -request_id auto
```

Enabling the compression accelerator settings for all LDEVs associated with from pool ID: 0 to maximum 128 pools.

```
# raidcom get dp_pool | rmawk @L-gt:1 exe="raidcom modify ldev -pool @1
-compression_acceleration enable -request_id auto"
```

Changing the capacity saving processing mode for LDEV: 0 to the inline mode.

```
# raidcom modify ldev -ldev_id 0 -capacity_saving_mode inline
```

Changing the capacity saving processing mode for LDEV: 0 to the post-process mode.

```
# raidcom modify ldev -ldev_id 0 -capacity_saving_mode post_process
```

Setting the upper limit of the data transfer volume to 100 MB/s for LDEV ID: 200.

```
# raidcom modify ldev -ldev_id 200 -upper_data_trans_mb 100 -request_id
auto
```

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Setting the lower limit of the throughput per second to 1500 IOPS for LDEV ID: 200.

```
# raidcom modify ldev -ldev_id 200 -lower_throughput_io 1500 -request_id auto
```

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Setting the I/O processing priority to 1 for LDEV ID: 200.

```
# raidcom modify ldev -ldev_id 200 -response_priority 1 -request_id auto
```

(For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Setting the alert notification time when the lower limit is not reached to 10 seconds for LDEV ID: 200.

```
# raidcom modify ldev -ldev_id 200 -lower_alert_time 10 -request_id auto
```

(For VSP 5000 series) Enabling ESE for LDEV: 200.

```
# raidcom modify ldev -ldev_id 200 -ese enable -request_id auto
```


raidcom add license

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command installs the license.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add license -keycode <key code>
```

Options and parameters

-keycode <key code>

Specify the key code of the license.

Examples

Specify PXPQRS275WMYZ as the key code of the license, and then install the license key.

```
#raidcom add license -keycode PXPQRS275WMYZ
```

raidcom delete license

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command removes the license.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete license -product_id <product ID>
```

Options and parameters

-product_id <product ID>
Specify the software ID.

Examples

Remove the license of ID: 4102.

```
#raidcom delete license -product_id 4102
```

raidcom modify license

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command changes the term license status.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify license -product_id <product ID> -license_status <status>
```

Options and parameters

-product_id <product ID>
Specify the software ID.

-license_status <status>
Specify the term license status.

- enable or yes: Enables the term license.
- disable or no: Disables the term license.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

Examples

Enable the term license of ID: 4102.

```
#raidcom modify license -product_id 4102 -license_status enable
```

raidcom get license

This command acquires the license information.

Syntax

```
raidcom get license [-key opt]
```

Options and parameters

[-key opt]

Displays the capacity of the internal and external volumes created in the storage system in GB, and the serial number.

Examples

Display the license information.

```
#raidcom get license
```

PRO_ID	STS	Type	L	Cap_Perm(TB)	Cap_Used(GB)	-	Term	Name
34049	INS	PER	L	50	10	-	-	"Cache Residency Manager"
34055	INS	TEM	U	-	-	-	-	"Dynamic Provisioning"

```
#raidcom get license -key opt
Serial# : 302656
Cap_Mounted(GB) : 1229000
```

Description of each column in output example:

PRO_ID

Displays the software ID.

STS

Displays the installation status of the software. For details about installation statuses, see the *System Administrator Guide*.

- INS: Installed
- DIS: Installed but the license is disabled
- NIN: Not installed

- NEL: Installed but the license capacity is insufficient
- GRP: The license capacity is insufficient because an LDEV was added, a copy pair was created, or a pool volume was added. The license will expire in 30 days.
- EXT: The temporary key is expired

Type

Displays the license key type. For details about license key types, see the *System Administrator Guide*.

- PER: Permanent
- TER: Term
- TEM: Temporary
- EME: Emergency
- - (hyphen): License is not installed. But if the license period is still available and you removed temporary key, TEM is displayed.

L

Displays if there is an upper limit for the installed permitted capacity.

- U: There is no upper limit
- L: There is an upper limit
- T: There is no upper limit. Displayed only when the license key type is temporary or emergency and the license key does not overwrite the installed license key.

Cap_Perm(TB)

Displays the installed permitted capacity in TB. If the software does not have the capacity limit, or if the license is not installed, a hyphen is displayed.

Cap_Used(GB):

Displays the size of the volume used by the software in GB. If the licensed capacity type is other than Used capacity, or if the license is not installed, a hyphen is displayed. For details about the license capacity, see the *System Administrator Guide*.

Term

Displays the remaining days until the term key, temporary key, or emergency key expires. If the license period is still available and you removed the temporary key, Term displays the remaining days of the license period. After the temporary key expires, Term displays the remaining days until the license can be installed again. If the term is unlimited and the license is not installed, a hyphen (-) is displayed.

Name

Displays the product name. Double quotation marks (") are added at the beginning and end of the product name.

Serial#

Displays the serial number of the unit.

Cap_Mounted(GB)

Displays the capacity of the internal and external volumes created in the storage system in GB.

raidcom add quorum

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Sets a quorum disk. If an LDEV number is not allocated to a quorum disk, no volume is set for the quorum disk.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add quorum -quorum_id <quorum id> -request_id auto - remote_storage <serial#>
<id> [-ldev_id <ldev#>]
```

Options and parameters

-quorum_id <quorum id> - request_id auto

Specifies the quorum ID.

The request ID is output in the following format when the execution of the **raidcom add quorum** command ends:

- REQID : <request#>

<request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

-remote_storage <serial#> <id>

Specifies the remote storage system for GAD pairs monitored by the quorum disk.

- serial#: Specifies the product serial number of the remote storage system.
 - When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
 - When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.
- id: Specifies the type identifier of the remote storage system.
 - R800: VSP G1x00, VSP F1500
 - R900: VSP 5000 series
 - M800: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, VSP F400, F600, F800

[-ldev_id <ldev#>]

Specifies an LDEV number of a quorum disk. If an LDEV number is not allocated to a quorum disk, no volume is set for the quorum disk.

Examples

Setting a quorum disk whose quorum ID is 0, remote storage serial number is 310001, type identifier is R800, and LDEV number is 0x0010.

```
#raidcom add quorum -quorum_id 0 -request_id auto -remote_storage 310001
R800 -ldev_id 0x0010
REQID : 0
```

Setting a quorum disk whose quorum ID is 0, remote storage serial number is 310001, type identifier is R800 without setting an LDEV.

```
#raidcom add quorum -quorum_id 0 -request_id auto -remote_storage 310001
R800
REQID : 0
```

raidcom delete quorum

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Deletes a quorum disk.

This command is executed asynchronously with the command input. Check the completion of this process using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete quorum -quorum_id <quorum id> -request_id auto
```

Options and parameters

-quorum_id <quorum id> -request_id auto

Specifies the quorum ID.

The request ID is output in the following format when the execution of the **raidcom delete quorum** command ends:

- REQID : <request#>

<request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

Example

Deleting a quorum disk whose quorum ID is 0.

```
# raidcom delete quorum -quorum_id 0 -request_id auto
```

raidcom modify quorum

For the quorum disk, this command sets the period of Read Response Guaranteed Time When Quorum Monitoring Stopped.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify quorum -quorum_id <quorum id> -timeout <timeout>
```

Options and parameters

-quorum_id <quorum id>

Specifies a quorum ID (0 to 31) in decimal or hexadecimal. For hexadecimal IDs, prefix 0x.

Example:

- -quorum_id 10
- -quorum_id 0x0a

-timeout <timeout>

Specifies the period of time in seconds (5 to 100) for Read Response Guaranteed Time When Quorum Monitoring Stopped.

Examples

For the quorum disk (quorum ID: 1), specify 40 seconds as the period of Read Response Guaranteed Time When Quorum Monitoring Stopped.

```
#raidcom modify quorum -quorum_id 1 -timeout 40
```

raidcom get quorum

This command displays information about the quorum disk.

Syntax

```
raidcom get quorum [-quorum_id <quorum id>] [-fx]
```

Options and parameters**[-quorum_id <quorum id>]**

Specifies a quorum ID (0 to 31) in decimal or hexadecimal. For hexadecimal IDs, prefix 0x. If this option is omitted, information of all quorum disks is displayed.

Example:

- -quorum_id 10
- -quorum_id 0x0a

[-fx]

Specifies this option to display an LDEV number in hexadecimal.

Examples

Displays the information of the quorum disk (quorum ID: 1).

```
#raidcom get quorum -quorum_id 1
```

```
QRDID : 1
LDEV : 2045
QRP_Serial# : 302646
QRP_ID : R8
```



```
Timeout(s) : 30
STS: REPLACING
```

Description of each column in output example:

QRDID

Displays the quorum disk ID for GAD.

LDEV

Displays the LDEV number of a quorum disk for GAD. If no volume is set for the quorum disk, no LDEV number is allocated to the quorum disk. Therefore, 65535 (0xffff) is displayed.

QRP_Serial#

Displays the serial number of a remote storage system assigned to serve as a quorum disk of GAD devices.

When the storage system is VSP 5000 series, 5 is added to the beginning of the serial number.

When the storage system is VSP G1x00 or VSP F1500, 3 is added to the beginning of the serial number.

QRP_ID

Displays the identifier of a remote storage system assigned to serve as a quorum disk of GAD devices.

- R8: VSP G1x00 or VSP F1500
- R9: VSP 5000 series
- M8: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, VSP F400, F600, F800

Timeout(s)

Displays the time in seconds for Read Response Guaranteed Time When Quorum Monitoring Stopped set for the quorum disk of GAD.

STS

Displays the status of a quorum disk for GAD. A hyphen (-) is displayed if no volume is set for the quorum disk.

- NORMAL: The quorum disk is in normal status.
- TRANSITIONING: The status of the quorum disk is being changed.
- BLOCKED: The quorum disk is blocked.
- REPLACING: The quorum disk is being replaced.
- FAILED: The quorum disk is in abnormal state.
- - (hyphen): This information is not available for this quorum disk.

raidcom replace quorum

This command replaces a quorum disk.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom replace quorum -quorum_id <quorum id> -ldev_id <ldev#>
```

Options and parameters

-quorum_id <quorum id>

Specifies the quorum ID.

-ldev_id <ldev#>

Specifies the LDEV number (0 to 65279). For example:

- -ldev_id 200

Examples

Change the LDEV number of the quorum disk (LDEV) associated with the quorum ID 1 to the LDEV number 200.

```
#raidcom replace quorum -quorum_id 1 -ldev_id 200
```

raidcom modify local_replica_opt

This command sets the local replica option. For details about the local replica option, see the user manuals for ShadowImage, ShadowImage for Mainframe, Thin Image, Compatible FlashCopy®, Volume Migration, and nondisruptive migration.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify local_replica_opt -opt_type <option type>
{-set_system_opt <system option> | -reset_system_opt <system option>}
```

Options and parameters**-opt_type <option type>**

Specify the local replication type to be set.

- open: Local replication option for ShadowImage, Thin Image, Volume Migration, nondisruptive migration.
- mainframe: Local replication option for ShadowImage for Mainframe, Compatible FlashCopy®, Volume Migration.

-set_system_opt <system option>

Specifies the local replica option ID that you are enabling. For details about the local replica option ID, see the applicable user manual.

-reset_system_opt <system option>

Specifies the local replica option ID that you are disabling. For details about the local replica option ID, see the applicable user manual.

Examples

Enable local replica option: 1

```
raidcom modify local_replica_opt -opt_type open -set_system_opt 1
```

Disable local replica option: 1

```
raidcom modify local_replica_opt -opt_type open -reset_system_opt 1
```

raidcom get local_replica_opt

This command lets you view the local replica option.

Syntax

```
raidcom get local_replica_opt -opt_type <option type>
```

Options and parameters**-opt_type <option type>**

Specifies the local replica option type that you are going to view:

- open: Local replication option for ShadowImage, Thin Image, Volume Migration, and nondisruptive migration.
- mainframe: Local replication option for ShadowImage for Mainframe, Compatible FlashCopy®, and Volume Migration.

Examples

View the local replica option of ShadowImage, Thin Image, Volume Migration, and nondisruptive migration.

```
#raidcom get local_replica_opt -opt_type open
```

```
Serial# : 563528
Type : open
Option : 1 14
```

View the local replica option of ShadowImage for Mainframe, Compatible FlashCopy[®], and Volume Migration.

```
#raidcom get local_replica_opt -opt_type mainframe
```

```
Serial# : 563528
Type : mainframe
Option : 1 14
```

Description of each column in output example:

Serial#

Displays the serial number.

Type

Displays the local replica option type.

Option

Displays the local replica option that is set to enable.

raidcom get remote_replica_opt

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Displays remote replication options.

Syntax

```
raidcom get remote_replica_opt [-opt_type <option type>]
[-key cu [-cu <cu#>]]]
```

Options and parameters

[-opt_type <option type>]

Specify the remote replication type to be displayed.

- tc: TrueCopy
- ur: Universal Replicator
- gad: global-active device

[-key cu]

Specify this option to view the remote replication option of each CU. This option is valid when the -opt_type tc option is specified.

[-cu <cu#>]

Specify this option to view the remote replication option of a specific CU number.

Examples

Displaying information about path blocked watch and path blocked watch SIM.

```
#raidcom get remote_replica_opt
Serial# : 563528
PBW(s) : 40
PBW_SIM(s) : 70
```

Viewing the remote replication options of TrueCopy.

```
# raidcom get remote_replica_opt -opt_type tc
Serial# : 563528
Type : TC
CAS : CU
MAX_CA : 2
```

Viewing the remote replication options of Universal Replicator.

```
# raidcom get remote_replica_opt -opt_type ur
Serial# : 563528
Type : UR
MAX_CA : 2
```

Viewing the remote replication options of global-active device.

```
# raidcom get remote_replica_opt -opt_type gad
Serial# : 563528
Type : GAD
MAX_CA : 2
```

Description of each column in the output example:

Serial#

Displays the product serial number.

Type

Displays the remote replication option type.

CAS

Displays the unit of setting the maximum number of initial copy activities.

- System: Set per system.
- CU: Set per CU.

MAX_CA

Displays the maximum number of initial copy activities.

PBW(s)

Displays the path blocked watch time (shown in seconds).

PBW_SIM(s)

Displays the path blocked watch sim time (shown in seconds).

Examples

Viewing the remote replication options of TrueCopy for each CU.

```
# raidcom get remote_replica_opt -opt_type tc -key cu
CU# MAX_CA
00 1
01 1
...
fe 1
```

Viewing the remote replication options of TrueCopy for CU number: 1.

```
# raidcom get remote_replica_opt -opt_type tc -key cu -cu 1
CU# MAX_CA
01 1
```

Description of each column in the output example:

CU#

Displays the CU number. The number is displayed in hexadecimal notation, and not displayed if it cannot be used.

MAX_CA

Displays the maximum number of initial copy activities. Displayed only when the unit for the maximum number of initial copy activities is CU.

raidcom modify remote_replica_opt

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Sets remote replication options.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

Setting the `path_blocked_watch` or `path_blocked_watch_sim` option.

```
raidcom modify remote_replica_opt { -path_blocked_watch <time(sec)>
| -path_blocked_watch_sim <time(sec)>}
```

Setting other remote replication options.

```
raidcom modify remote_replica_opt -opt_type <option type> {-copy_activity_setting
{system|cu}
| -copy_activity <number of activities> [-cu <cu#>]}
```

Options and parameters

-opt_type <option type>

Specifies the remote replication type to be set.

- tc: TrueCopy
- ur: Universal Replicator
- gad: global-active device

-copy_activity_setting {system|cu}

Specifies the unit (per system or CU) of managing the maximum number of initial copy activities.

- system: Managed per system.
- cu: Managed per CU.

-copy_activity <number of activities>

Specifies the maximum number of initial copy activities.

[-cu <cu#>]

Specifies the CU number for which the `-copy_activity <number of activities>` option is set, in decimal or hexadecimal. For hexadecimal numbers, prefix 0x. For example:

- `-cu 128`
- `-cu 0x80`

-path_blocked_watch <time(sec)>

Specifies the path blocked watch time in seconds.

-path_blocked_watch_sim <time(sec)>

Specifies the path blocked watch SIM time in seconds.

Examples

Specifying 30 seconds as the path blocked watch time .

```
# raidcom modify remote_replica_opt -path_blocked_watch 30
```

Managing the maximum number of initial copy activities of TrueCopy per system.

```
# raidcom modify remote_replica_opt -opt_type tc -copy_activity_setting system
```

Specifying 2 as the maximum number of initial copy activities of TrueCopy for CU number: 1.

```
# raidcom modify remote_replica_opt -opt_type tc -copy_activity 2 -cu 1
```

Specifying 2 as the maximum number of initial copy activities of Universal Replicator.

```
# raidcom modify remote_replica_opt -opt_type ur -copy_activity 2
```

raidcom add lun

To set the LU path, this maps the specified LDEV to a LUN on a host group on the specified port and creates an LU path or alternate path. A device group can also be specified instead of an LDEV.

If the specified port or host group does not exist, this command is rejected with EX_ENOOBJ(EX_CMDRJE).

If the specified LUN or LDEV already exists, this command is ignored.

If a LUN is not specified, an empty LUN is assigned automatically.

Not allowed:

- Mapping the same LDEV to another LUN in the same host group.
- Overwriting the same LUN to another LDEV.

When the `-request_id auto` option is specified, if a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If you specify the `-request_id` option for the asynchronous command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).



Caution:

- LDEVs in LUSE volumes cannot be configured with this command.
- If you execute this command on an LDEV with the command device attribute already set and the LDEV already has a defined/configured path, the command device attribute will be released.

Syntax

To set LU path by specifying the LDEV:

```
raidcom add lun -port <port#> [<host group name>]
    -ldev_id <ldev#> [-lun_id <lun#> | -lun_id auto
    -request_id auto]
```

To set LU path by specifying the device group:

```
raidcom add lun -port <port#> [<host group name>]
    -grp_opt ldev -device_grp_name <device group name> [<device name>]
```

To set LU path by specifying multiple port numbers simultaneously:

```
raidcom add lun -port <port#> [<host group name>]
    {-ldev_id <ldev#> [-lun_id <lun#> | -lun_id auto
    -request_id auto ] -additional_port <additional port>...}
```

Options and parameters

-port <port#> [<host group name>]

Specifies the Port number and the host group. You can specify the host group ID or the host group name for the host group. If you specify neither the host group ID nor the host group name, the host group 0 is used. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-grp_opt ldev

The information of LDEV belonging to the device group is used. Specify 'ldev' whenever.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

[-lun_id <lun#> | -lun_id auto -request_id auto]

Specifies the LU number. If this option is omitted, a free LU number is assigned automatically.

- -lun_id <lun#>: Specifies the LU number (0 to 2047). For VSP 5000 series, the LUN that can be specified is (0 to 4095).
- -lun_id auto: The storage system assigns an available LU number.
- -request_id auto: Effective only when "auto" is specified as the `-request_id` option. If an option other than "auto" is specified, EX_INVARG or EX_REQARG is replied.

The request ID is output in the following format when the execution of the **raidcom add lun** command ends:

- REQID : <request#>

<request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

If multiple commands without LU numbers are executed simultaneously for the same host group, the same LU number might be assigned to multiple newly created LU paths. To avoid this behavior, perform one of the following:

- Specify the `-lun_id auto` option.
- Use the **raidcom lock resource** command to lock the resource group to which the host group belongs. The host group sets the LU number.

If you lock the resource group, commands executed by other users cannot take out unused LU numbers from the resource group. If you set multiple LU paths for the same host group, execute the next **raidcom add lun** command after the execution of the current **raidcom add lun** command is complete.

After the execution of all the **raidcom add lun** commands are complete, use the **raidcom unlock resource** command to unlock the resource group.

**Note:**

- You cannot specify this option if you want to specify a device group as the operation target.
- You cannot omit this option if you specify multiple port numbers at the same time.

additional_port <additional port>...

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Up to 5 port numbers can be specified to set the LU path additionally.

The LU path is set to the host group that you specified in the `-port <port#>` option. When you add ports, you cannot specify the host group name by using the `[<host group name>]` option.

If the `-lun_id auto` option is specified, a common LU number that is not used by the host groups to which LU paths are added is selected. If there are multiple unused LU numbers, the smallest LU number is selected.

Examples

Map the LDEV: 200 to the LU numbers: 1 for the port: CL1-A, the host group #0.

```
# raidcom add lun -port CL1-A-0 -lun_id 1 -ldev_id 200
```

Map the LDEV: 200 for the port: CL1-A, the host group #0. The LU numbers are automatically assigned.

```
# raidcom add lun -port CL1-A-0 -ldev_id 200
```

Map the LDEV for the port: CL1-A, the host group #0, and the one belonging to the device group: grp1. The LU numbers are automatically assigned.

```
# raidcom add lun -port CL1-A-0 -grp_opt ldev -device_grp_name grp1
```

Map the LDEV: 200 for the port: CL1-A, CL2-A, CL3-A, the hosts group #0, and the LU number 1.

```
# raidcom add lun -port CL1-A-0 -ldev_id 200 -lun_id 1 -additional_port CL2-A CL3-A
```

raidcom delete lun

Deletes the LU path on the host group on the specified port. An LDEV or a device group can also be specified instead of a LUN. If an LDEV does not exist on the specified port /host group/LUN, this command is rejected with EX_ENLDEV or EX_ENOOBJ.

When deleting the LU path, stop the I/O for the LU path to be deleted. In the last one path, LDEV must be specified as SMPL volume.



Note:

LDEVs in LUSE volumes cannot be configured with this command. Do not execute this command to an LDEV whose command device attribute is set. If the command is executed, the command device attribute is released.

Syntax

To delete LU path with specifying the LUN:

```
raidcom delete lun -port <port#> [<host group name>]
                    -lun_id <lun#>
```

To delete LU path with specifying the LDEV:

```
raidcom delete lun -port <port#> [<host group name>]
                    -ldev_id <ldev#>
```

To delete LU path with specifying the device group:

```
raidcom delete lun -port <port#> [<host group name>]
                    -grp_opt <group option> -device_grp_name <device group
                    name> [<device name>]
```

To delete LU path with specifying the multiple port numbers simultaneously :

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

```
raidcom delete lun -port <port#> [<host group name>]
                    {-lun_id <lun#> | -ldev_id <ldev#>}
                    -additional_port <additional port>...
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the Port number and the host group. You can specify the host group ID or the host group name for the host group. If you specify neither the host group ID nor the host group name, the host group 0 is used. For example:

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86

-lun_id <lun#>

Specifies the LUN number (0-2047). For VSP 5000 series, the LUN that can be specified is (0 to 4095).

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-grp_opt ldev

The information of LDEV belonging to the device group is used. Specify 'ldev' whenever.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all of the LDEVs belonging in the device group are operated.

-additional_port <additional port>...

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Up to 5 port numbers can be specified to delete the LU path additionally.

The LU path is deleted from the host group that you specified into -port <port#> [<host group name>] option. If you use this option, you cannot specify the host group name by using the -port <port#> [<host group name>] option.

If you use this option with the -ldev_id <ldev#> option, all LUNs of the LU path which you delete must be the same.

If you use this option with the `-lun_id <lun#>` option, all LU path which you delete must be set to the same LDEV.

Examples

Deleting LUN: 1(LDEV number 200) on the port: CL1-A-0

```
# raidcom delete lun -port CL1-A-0 -lun_id 1
# raidcom delete lun -port CL1-A-0 -ldev_id 200
```

Deleting LDEV belonging to the device group: grp1 on the port: CL1-A-0

```
# raidcom delete lun -port CL1-A-0 -grp_opt ldev -device_grp_name grp1
```

Deleting LUN: 1(LDEV number 200) on the port CL1-A, CL2-A, CL3-A of the host group number #0

```
# raidcom delete lun -port CL1-A-0 -lun_id 1 -additional_port CL2-A CL3-A
# raidcom delete lun -port CL1-A-0 -ldev_id 200 -additional_port CL2-A CL3-A
```

raidcom discover lun

Searches volume on the external storage system. Displays a list of LUs that can be referred to from the External port of a specific external storage system.

This command is rejected by EX_ENOOBJ in the following cases:

- The specified iSCSI port cannot be found.
- The iSCSI virtual port mode is enabled, but the specified iSCSI virtual port ID is not correct.

To fix this error, specify the correct port and iSCSI virtual port ID.

If you execute the **raidcom discover lun** command while the **raidcom discover external_storage** command or another **raidcom discover lun** command is being executed, the LU of the external storage system might not be displayed. If this happens, confirm that the storage system in which the command is being executed and the external storage system are connected correctly, and the LU of the external storage system is configured correctly.

If both the storage system and the external storage system have no problem, wait a while, and then execute one command at a time for a storage system.



Note:

If the local storage system port is a Fiber Channel Bidirectional port, it is recommended that you always execute the command by specifying the `-safety_check enable` option unless instructed by the user guide.

If you do not specify the `-safety_check enable` option, the following problems might occur. For the details, see the Hitachi Universal Volume Manager User Guide.

- If you use a non-Hitachi external storage system:
The I/O path from the external storage system to the local storage system that uses the route between the specified local storage system port and the searched external storage system port might be disconnected.
- If you use the Hitachi external storage system:
The external or remote path connection that uses the route between the specified local storage system port and the searched external storage system port might be temporarily disconnected (the external or remote path will be reconnected immediately and no blockage occurs).
- If you run multiple instances of the `raidcom discover external_storage` command, `raidcom discover lun` command, `raidcom check external_iscsi_name` command, or `raidcom discover external_iscsi_name` command at the same time, an EX_CMDIOE error might occur. If an EX_CMDIOE error occurs, wait for a while, and then run the commands again one by one. For details, see the *Hitachi Universal Volume Manager User Guide*.

Syntax

```
raidcom discover lun -port <port#> {-external_wwn <wwn strings>
| -external_iscsi_name <external iscsi name>
-external_address <IP address>
[-iscsi_virtual_port_id <iSCSI virtual port ID>]} [-safety_check enable]
```

Options and parameters

-port <port#>

Specifies the Port number. It specifies a port of which attribute is External. For example:

- CL1-A



Note: When you specify a non-bidirectional port with an attribute other than External for the `-port` option, the `raidcom discover lun` command works as the `raidcom get lun` command.

-external_wwn <wwn strings>

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The 17th digit and later are ignored. The value can be split in units of 4 bytes by ",", (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-external_iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn or eui format.

- iqn format: `iqn.` and the subsequent maximum 219 characters.
- eui format: `eui.` and the subsequent 16 characters in hexadecimal notation.



Note: You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option or the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the **raidcom get external_iscsi_name** command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

-external_address <IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

The following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

The following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).



Note: You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option or the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the **raidcom get external_iscsi_name** command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

[-safety_check enable]

Specify the Fiber Channel bidirectional port for the `-port` option. If you specify any other port, this option is ignored. When this option is specified, processing that might cause the disconnection between the Bidirectional port specified by the `-port` option and the Fiber Channel port of the external storage system specified by the `-external_wwn` option is suppressed. When processing is suppressed, the LU information is not displayed.

Example 1

Displaying LUs defined to the external storage system port: 50060e80, 06fc3850 connected to External (ELUN) port: CL5-A when the external storage system is connected to a FIBRE port.

```
# raidcom discover lun -port CL5-A -external_wwn 50060e8006fc3850 -safety_check enable
```

PORT	WWN	LUN	VOL_Cap (BLK)	PRODUCT_ID	E_VOL_ID_C
CL5-A	50060e8006fc3850	0	102400	OPEN-V	HITACHI R500FC381000
CL5-A	50060e8006fc3850	1	102400	OPEN-V	HITACHI R500FC381001
CL5-A	50060e8006fc3850	2	102400	OPEN-V	HITACHI R500FC381002
CL5-A	50060e8006fc3850	3	102400	OPEN-V	HITACHI R500FC381003
CL5-A	50060e8006fc3850	4	102400	OPEN-V	HITACHI R500FC381004
CL5-A	50060e8006fc3850	5	102400	OPEN-V	HITACHI R500FC381005
CL5-A	50060e8006fc3850	6	102400	OPEN-V	HITACHI R500FC381006
CL5-A	50060e8006fc3850	7	102400	OPEN-V	HITACHI R500FC381007
CL5-A	50060e8006fc3850	8	102400	OPEN-V	HITACHI R500FC381008
CL5-A	50060e8006fc3850	9	102400	OPEN-V	HITACHI R500FC381009
CL5-A	50060e8006fc3850	10	102400	OPEN-V	HITACHI R500FC38100A

Description of each column in output example 1:**PORT**

External (ELUN) port number.

WWN

WWN on the external storage system. When the external storage system is connected to an iSCSI port, this item displays the pseudo WWN of the external storage system.

LUN

LUN of the port on external storage system.

VOL_Cap (BLK)

Capacity of the volume on the external storage system in units of block (1 block = 512 bytes).

PRODUCT_ID

Product_ID included in the SCSI Inquiry command response of the volume on the external storage system. The displayed contents and format depend on the connected volume.

For LUs that cannot be used as an external volume, "OTHER" is displayed.

E_VOL_ID_C

Volume identifier included in the SCSI inquiry command of the volume on the external storage system. The displayed contents and format depend on the connected volume.

Example 2

Displaying LUs defined to the iSCSI target of the external storage system (iSCSI name: iqn.z2, IP address: 158.214.135.100) connected to the port: CL1-B, iSCSI virtual port ID: 2.

```
# raidcom discover lun -port CL1-B -iscsi_virtual_port_id 2 -external_iscsi_name
iqn.z2 -external_address 158.214.135.100
```

PORT	WWN	LUN	VOL_Cap (BLK)	PRODUCT_ID	E_VOL_ID_C
CL1-B	50060e8006fc3d60	16	545280	OPEN-V HITACHI	R500FC3D0210
CL1-B	50060e8006fc3d60	17	545280	OPEN-V HITACHI	R500FC3D0211
CL1-B	50060e8006fc3d60	18	545280	OPEN-V HITACHI	R500FC3D0212
CL1-B	50060e8006fc3d60	19	545280	OPEN-V HITACHI	R500FC3D0213

Example 3

Displaying LUs defined to the iSCSI target of the external storage system (iSCSI name: iqn.z2, IP address: 158.214.135.100) connected to port: CL5-A when the external storage system is connected to an iSCSI port.

```
# raidcom discover lun -port CL5-A -external_iscsi_name iqn.z2
```

```
-external_address 158.214.135.100
```

PORT	WWN	LUN	VOL_Cap (BLK)	PRODUCT_ID	E_VOL_ID_C
CL1-B	50060e8006fc3d60	16	545280	OPEN-V HITACHI	R500FC3D0210
CL1-B	50060e8006fc3d60	17	545280	OPEN-V HITACHI	R500FC3D0211
CL1-B	50060e8006fc3d60	18	545280	OPEN-V HITACHI	R500FC3D0212
CL1-B	50060e8006fc3d60	19	545280	OPEN-V HITACHI	R500FC3D0213


Getting the external storage system's iSCSI target information corresponding to the pseudo WWN

To get the iSCSI name and IP address of the iSCSI target of the external storage system which corresponds to the pseudo WWN, execute the **raidcom get external_iscsi_name** command.

```
# raidcom get external_iscsi_name
```

PORT	Serial#	IP_ADDR	IQN	WWN(pseudo)	AMD	D	CHAP_user	Sec
CL4-E	63528	158.214.135.100	iqn.z1	50060e80070a3640	CHAP	D	Win_SQL_EX	*
CL2-E	63528	158.214.135.102	iqn.z3	50060e80070a3642	CHAP	S	-	-
CL1-B	63528	158.214.135.100	iqn.z2	50060e8006fc3d60	CHAP	S	-	-

The pseudo WWN is managed by each storage system. Therefore, when two storage systems share the iSCSI target of an external storage system as shown in the following figure, the pseudo WWN corresponding to iSCSI target 1 of the storage system A is different from the pseudo WWN corresponding to iSCSI target 1 of the storage system B.



Examples

The following example shows, in the configuration shown above, how to get the iSCSI name and the IP address of the external storage system, and how to display the list of LUNs, by specifying the pseudo WWN which is managed by the storage system A.

```
# raidcom get external_iscsi_name -s 64562 | rmawk @5-eq:50060e8006fc3d60
exe="raidcom discover lun -s 34562 -port CL1-b -external_address@3 -iscsi_name @4"
```

The following examples show how to display the list of LUNs by the storage system B by getting the pseudo WWN of the storage system B which corresponds to the pseudo WWN managed by the storage system A.

Windows example

```
C:\horcm\etc>raidcom get external_iscsi_name -s 64562 | rmawk @5-eq: 50060e8006fc3d60
exe="raidcom get external_iscsi_name -s 34562 | rmawk @@3-eq:@3 -a @@4-eq:@4 exe=
\"raidcom discovery lun -s 34562 -port CL1-b -external_wnn @@5\""
```

UNIX example

```
# raidcom get external_iscsi_name -s 64562 | rmawk @5-eq: 50060e8006fc3d60
exe="raidcom get external_iscsi_name -s 34562 | rmawk @@3-eq:@3 -a @@4-eq:@4
exe="raidcom discovery lun -s 34562 -port CL1-b -external_wnn @@5"
```

raidcom get lun

Displays the LU path information defined in the specified port and host group.

If the specified port does not exist, this command is rejected with EX_ENOOBJ. If an external port is specified, it is rejected with EX_REQARG.

Syntax

```
raidcom get lun -port <port#> <host group name>
[-key <keyword>]
```

Options and parameters**-port <port#>[<host group name>]**

Specifies the port number, host group ID, and host group name. If the host group name is more than 64 characters, use the host group ID. If the host group ID and host group name are omitted, the LU path information for host group ID 0 is displayed.

For example,

- CLI-A
- CL1-A-g (g: 0-254)
- CL1-A Linux_X86

-key <keyword>]

Specifies the information about the LU to be displayed.

- opt: Displays the status of the host reservation.
- opt_page1: Displays the information about the specified ALUA mode.

Examples

Displaying the LU information defined to the port CL4-E, the host group #0.

```
# raidcom get lun -port CL4-E-0
```

PORT	GID	HMD	LUN	NUM	LDEV	CM	Serial#	HMO_BITS
CL4-E	0	LINUX/IRIX	0	1	0	CM	63528	2 13
CL4-E	0	LINUX/IRIX	2	1	2	-	63528	2 13
CL4-E	0	LINUX/IRIX	3	1	3	-	63528	2 13
CL4-E	0	LINUX/IRIX	4	1	4	-	63528	2 13
CL4-E	0	LINUX/IRIX	5	1	992	-	63528	2 13
CL4-E	0	LINUX/IRIX	6	1	993	-	63528	2 13

```
#raidcom get lun -port CL4-E-0 -key opt
```

PORT	GID	HMD	LUN	NUM	LDEV	CM	Serial#	OPKMA	HMO_BITS
CL4-E	0	LINUX/IRIX	0	1	0	CM	63528	-Y---	2 13
CL4-E	0	LINUX/IRIX	2	1	2	-	63528	-Y---	2 13
CL4-E	0	LINUX/IRIX	3	1	3	-	63528	-Y---	2 13
CL4-E	0	LINUX/IRIX	6	1	993	-	63528	-Y---	2 13

```
# raidcom get lun -port CL4-E-0 -key opt_page1
```

PORT	GID	HMD	LUN	NUM	LDEV	CM	Serial#	AL	AAS
CL4-E	0	LINUX/IRIX	0	1	0	CM	63528	E	AO
CL4-E	0	LINUX/IRIX	2	1	2	-	63528	D	AO
CL4-E	0	LINUX/IRIX	3	1	3	-	63528	E	AO

Description of each column in output example:

PORT

Displays the port number.

GID

Displays the host group ID on the port.

HMD

Displays the following HOST MODE for the host adapter setting on host group.

- HP-UX, SOLARIS, AIX, WIN, LINUX/IRIX, TRU64, DYNIX, OVMS, NETWARE, HI-UX
- VMWARE, HP-XP, VMWARE_EX, WIN_EX, UVM

LUN

Displays LUN number on host group mapping LDEV.

NUM

Displays the number of LDEVs composing an LUSE.

LDEV

Displays the LDEV number.

CM

Displays the command device.

Serial#

Product serial number (Serial#).

OPKMA

Displays the host-reserved LU.

- O: The LU is reserved by an open system¹.
 - P: The LU is reserved by a persistent group¹.
 - K: The LU is registered by a PGR key².
 - M: The LU is reserved by a mainframe¹.
 - A: The LU is reserved by ACA¹.
1. If Y is displayed under each character, the LU is reserved. If a hyphen (-) is displayed, the LU is not reserved.
 2. If Y is displayed under each character, the PGR key is registered. If a hyphen (-) is displayed, the PGR key is not registered.

HMO_BITS

Displays the host mode options of the host groups.

For details, see the *Provisioning Guide* for the storage system.

AL

Displays the information about the ALUA mode.

- E: The ALUA mode is enabled.
- D: The ALUA mode is disabled.

AAS

Displays the setting value of the asymmetric access state for ALUA. The displayed setting value shows whether or not the LU is accessed from the host preferentially. If the ALUA mode is enabled, the setting value displayed under AAS is reported to the host as the value of the asymmetric access state.

- AO: Active and optimized LU. The host accesses the LU preferentially.
- AN: Active and non-optimized LU. When an LU whose setting value of the asymmetric access state is AO cannot be used, the host accesses the LU.
- - (hyphen) : The setting for the asymmetric access state is not supported.

raidcom modify lun

Modifies the LU attribute.

When releasing LU host reserves, this command is executed asynchronously with command input. Check the completion of the process using the `raidcom get command_status` command.



Caution: When releasing the LU host reserve, do not operate from Device Manager - Storage Navigator and CCI simultaneously. If you do so, the LU host reserve might not be released. If you fail to release the host reserve, specify the LU again and retry the operation from CCI. When you want to release the host reserve on an LU that is used in a global-active device pair, make sure to release the host reserves for both the P-VOL and the S-VOL. If you release only one host reserve, a Reserve Conflict might occur when the host is replaced.



Caution:

(VSP 5000 series only) If you set the asymmetric access state for LU with the `raidcom modify lun` command, an EX_CMDIOE error might occur. If the EX_CMDIOE error occurs, wait for approximately 40 seconds, and then execute the `raidcom get lun -key opt_page1` command to verify that the specified value is set in the asymmetric access state. If the specified value is in the asymmetric access state, the `raidcom modify lun` command is ended correctly.

If the specified value is not in the asymmetric access state, set the timeout value in the configuration definition file to 80 seconds or more, and then retry the `raidcom modify lun` command. If the error occurs again, see the command error messages in the Command Control Interface User and Reference Guide to take appropriate actions. After the asymmetric access state migration is completed correctly, return the timeout value in the configuration definition file to the original value.

Syntax

```
raidcom modify lun {-port <port#> [<host group name>]
                  -lun_id {all | <lun#>} | -ldev_id <ldev#>} {-asymmetric_access_state {optimized |
non_optimized}
                  | -reservation release}
```

Options and parameters

-port <port#> [<host group name>]

Specifies the port number, and host group ID or host group name (iSCSI target alias if iSCSI is used). If the number of characters for the host group name is more than 64, specify the host group ID or the iSCSI target alias.

For example,

- CL1-A-g (g is from 0 to 254)
- CL1-A Linux_X86
- CL1-A Target00

-lun_id {all | <lun#>}

Specifies the LU.

- all: Specifies all LUs in the specified host group. Users who execute the command must have authority to the specified host group, and to all the LDEVs mapped to the LUs of the host group.
- <lun#>:
Specifies the LUN number (0 to 2047). For VSP 5000 series, the LUN that can be specified is (0 to 4095). Users who execute the command must have authority to the specified host group, and to all the LDEVs mapped to the specified LUs. <lun#> cannot be specified together with the `-asymmetric_access_state` option.

-ldev_id <ldev#>

Specifies the LDEV ID. When releasing the host reserve by specifying this option, specify an LDEV with an emulation type of OPEN-X (X is optional).

For example:

- -ldev_id 200
- -ldev_id 0xc8
- -ldev_id 00:c8

Users who execute the command must have the following authority:

- The authority of LDEVs.
- The authority of the host group connected to the LU path.
- The authority of the iSCSI target.

If the LU path of the target LDEV is changed during the command execution, releasing the LU path host reserve of the non-target LDEV might be executed. To avoid this behavior, use the **raidcom lock resource** command in advance to lock the host group or iSCSI target connected to the LU path. As a result, this prevents LU path change operations of other users. After the execution of all the **raidcom modify lun** commands are complete, use the **raidcom unlock resource** command to unlock the host group or iSCSI target.

When releasing the host reserve is executed by specifying the `-ldev_id` option, the **host reserve release** command is executed by the `-port -lun_id` option for all the LUs of the LDEV specified internally.

This internally executed command is processed asynchronously with the completion of the **raidcom modify lun** command specifying the `-ldev_id` option. When releasing the host reserve of the LDEV where the LU exceeding the following number is defined, release the host reserve by specifying the `-lun_id` option. The number of asynchronous commands that can be accepted by the storage system is as follows.

- VSP 5000 series: 2,560
- VSP G1x00 and VSP F1500: 1,280
- VSP One B20, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP E series: 2,560
- VSP G200, G400, G600, G800 and VSP F400, F600, F800: 2,560

When releasing the host reserve by specifying the `-ldev_id` option, perform the following procedure:

1. Every time the **raidcom modify lun** command is executed, execute the **raidcom get command_status** command to confirm the completion of the host reserve release.
2. If asynchronous commands that exceed the number of acceptable asynchronous commands are executed for the storage system, an EX_CMDRJE (SSB1 = 2E11, SSB2 = 9400) error is displayed. If this error is displayed, wait a while, and then retry the command.

-asymmetric_access_state {optimized | non_optimized}

Specifies the asymmetric access state for the LU. For the LU mapped to an LDEV whose ALUA mode is enabled, the value specified by the `-asymmetric_access_state` option is reported to the host as the value of the asymmetric access state.

Specify the `-asymmetric_access_state` option together with the `-lun_id all` option.

- `optimized` or `yes`: Active and optimized LU. The host accesses the LU preferentially.
- `non_optimized` or `no`: Active and non-optimized LU. When an LU whose asymmetric access state is "optimized" cannot be used, the host accesses the LU.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

-reservation release

Releases the host reserve of the LU specified by the `-lun_id` option or `-ldev_id <ldev#>`.

Examples

Set the asymmetric access state for all LU on the host group (host group ID: 2) on the port (port ID: CL4-E) to optimized.

```
# raidcom modify lun -port CL4-E-2 -lun_id all -asymmetric_access_state optimized
```

Release the host reserve of all LUs on the host group (host group ID: 2) on the port (port ID: CL4-E).

```
# raidcom modify lun -port CL4-E-2 -lun_id all -reservation release
```

Release the host reserve of all LUs on the host group (host group ID: 0x01)

```
# raidcom modify lun -ldev_id 1 -reservation release
```

raidcom add path

Adds and changes an external path connected to a volume on the external storage system to the existing path group. Only one path is operated in one operation. The order of priority for the path is allocated in accordance with the order of adding paths.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

This command is rejected by EX_ENOOBJ in the following cases:

- The specified iSCSI port cannot be found.
- The iSCSI virtual port mode is enabled, but the specified iSCSI virtual port ID is not correct.

To fix this error, specify the correct port and iSCSI virtual port ID.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Note: If the local storage system port is a Fiber Channel Bidirectional port, it is recommended that you always execute the command by specifying the `-safety_check enable` option unless instructed by the Command Control Interface User and Reference Guide .

If you do not specify the `-safety_check enable` option, the following problems might occur. For the details, see the Hitachi Universal Volume Manager User Guide.

- If you use a non-Hitachi external storage system:
The I/O path from the external storage system to the local storage system that uses the route between the specified local storage system port and the external storage system port might be disconnected.
- If you use a Hitachi external storage system:
The external or remote path connection that uses the route between the specified local storage system port and the external storage system port might be temporarily disconnected (the external or remote path will be reconnected immediately and no blockage occurs).

Syntax

```
raidcom add path -path_grp <path group#> -port <port#>
    {-external_wnn <wnn strings> | -external_iscsi_name
    <external iscsi name> -external_address <IP address>
    [-iscsi_virtual_port_id <iSCSI virtual port ID>]} [-safety_check enable]
```

Options and parameters

-path_grp <path group#>

Specifies the external VOL path group number (0-63231).

-port <port#>

Specifies the Port number. Specifies the number of the port whose attribute is External. For example:

- CL1-A

-external_wnn <wnn strings>

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-external_iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: `iqn.` and the subsequent maximum 219 characters.
- eui format: `eui.` and the subsequent 16 characters in hexadecimal notation.



Note:

You can specify the pseudo WWN of the iSCSI target to the `-external_wnn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

-external_address <IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).



Note:

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the **raidcom get external_iscsi_name** command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

[-safety_check enable]

Specify the Fiber Channel bidirectional port for the `-port` option. If you specify any other port, this option is ignored. When this option is specified, processing that might cause the disconnection between the Bidirectional port specified by the `-port` option and the Fiber Channel port of the external storage system specified by the `-external_wwn` option is suppressed. When processing is suppressed, the external path is not added.

Examples

Adding a path of External (ELUN) port CL1-A, external storage system port 50060e80,05fa0f36 to an external volume path group number: 1.

```
# raidcom add path -path_grp 1 -port CL1-A -external_wnn 50060e80,05fa0f36 -
safety_check enable
```

Adding the path between the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100) and the iSCSI port: CL1-A of the local storage system to the external volume path group number: 1.

```
# raidcom add path -path_grp 1 -port CL1-A -external_iscsi_name iqn.z2 -
external_address 158.214.135.100
```

Adding the path between the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100) and the iSCSI port: CL1-A, iSCSI virtual port ID: 2 of the local storage system to the external volume path group number: 1.

```
# raidcom add path -path_grp 1 -port CL1-A -iscsi_virtual_port_id 2 -
external_iscsi_name iqn.z2 -external_address 158.214.135.100
```

raidcom get path

Displays the information about the external path connected to a volume on the external storage system.

Syntax

```
raidcom get path [-path_grp <path group#> |
  -external_grp_id <gno-sgno> | -ldev_id <ldev#>]
[{-check_status | -check_status_not} <string>...
[-time <time>][-key <keyword>]]
```

Options and parameters

[-path_grp <path group#>]

Specifies the external VOL path group number (0-63231).

When you specify this option, only the external path information in the specified path group is displayed.

[-external_grp_id <gno-sgno>]

Specifies the external volume group number (gno:1-16384, sgno: 1-4096). For example:

- 52-11

When you specify this option, only the external path information of the specified external volume group is displayed.

[-ldev_id<ldev#>]

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

When you specify this option, only the external path information to the external volume of the specified LDEV is displayed.

[-check_status <string>... [-time <time>]]

Check if the external volume group is in the same state as specified in <string>. If the option contains multiple states, the OR condition check is performed and verifies that the external volume is in one of the states contained in the option. You can check the state of the external volume group displayed at the beginning when executing the **raidcom get path** command without specifying the **-check_status** option.

The following strings are specified in the <string>:

- NML: The external volume group is in the normal status.
- CHK: The external volume group is being checked.
- SYN: The cached data is being written to the volume on the external storage system.
- DSC: The connection to the volume on the external storage system is disconnected.
- BLK: The external path is blocked.
- UNK: The external volume group is in the unknown status.
- WAR: The status of the external volume group is at the warning level.

If you specify the **-time** option, this command checks the status of the external volume group every three seconds until the end of the specified <time> (seconds).

When this option is specified, the returned values are as follows:

- The external volume group is in one of the specified states: 0
- The external volume group is not in any of the specified states (without **-time** option): 1
- The external volume group is not in any of the specified states (when the specified <time> passed): EX_EWSTOT

[-check_status_not <string> [-time <time>]]

Check if the external volume group is not in the same state as specified in <string>. If the option contains multiple states, the NOR condition check is performed and verifies that the external volume is not in any of the states contained in the option. You can check the state of the external volume group displayed at the beginning when executing the **raidcom get path** command without specifying the **-check_status_not** option.

The following strings are specified in the <string>.

- NML: The external volume group is in the normal status.
- CHK: The external volume group is being checked.

- SYN: The cached data is being written to the volume on the external storage system.
- DSC: The connection to the volume on the external storage system is disconnected.
- BLK: The external path is blocked.
- UNK: The external volume group is in the unknown status.
- WAR: The state of the external volume group is at the warning level.

If you specify the `-time` option, this command checks the status of the external volume every three seconds until the end of the specified `<time>` (seconds).

When this option is specified, the returned values are as follows:

- The external volume group is not in any of the specified states: 0
- The external volume group is in one of the specified states (without `-time` option): 1
- The external volume group is in one of the specified states (when the specified `<time>` passed): EX_EWSTOT

[-key <keyword>]

Displays the detailed information about a path. Specify `opt` as `<keyword>`.

Example 1

Displaying the external path (group) information

```
# raidcom get path
```

PHG	GROUP	STS	CM	IF	MP#	PORT	WWN		PR	LUN	PHS	Serial#	PRODUCT_ID	LB
PM	DM	QD	TO(s)	PBW(s)	PDBM	PDBT(%)								
1	1-1	NML	E	D	2	CL1-A	50060e8005fa0f36		1	3	NML	60010 VSP N	M E	8
15	10			D 90										
1	1-1	NML	E	D	2	CL2-A	50060e8005fa0f38		2	3	NML	60010 VSP N	M D	8
15	10			D 90										
5	2-1	NML	E	D	0	CL3-B	50060e8006fc3222		1	0	NML	64562 VSP N	M E	8
15	10			D 90										
5	2-1	NML	E	D	0	CL5-A	50060e8006fc4150		-	-	UNK	64562 VSP N	M E	8
15	10			D 90										
5	2-2	NML	E	D	2	CL3-B	50060e8006fc3222		-	-	UNK	64562 VSP N	A D	8
15	10			D 90										
5	2-2	NML	E	D	2	CL5-A	50060e8006fc4150		1	0	NML	64562 VSP N	A D	8
15	10			D 90										

Displaying the information of the external volume path group number: 1

```
# raidcom get path -path_grp 1
```

PHG	GROUP	STS	CM	IF	MP#	PORT	WWN	PR	LUN	PHS	Serial#	PRODUCT_ID	LB
PM	DM	QD	TO(s)	PBW(s)	PDBM	PDBT(%)							
1	1-1	NML	E	D	2	CL1-A	50060e8005fa0f36	1	3	NML	60010 VSP N	M E	8
15		10		D 90									
1	1-1	NML	E	D	2	CL2-A	50060e8005fa0f38	2	3	NML	60010 VSP N	M D	8
15		10		D 90									

Displaying the information of the external volume path group number: 5

```
# raidcom get path -path_grp 5
```

PHG	GROUP	STS	CM	IF	MP#	PORT	WWN	PR	LUN	PHS	Serial#	PRODUCT_ID	LB
PM	DM	QD	TO(s)	PBW(s)	PDBM	PDBT(%)							
5	1-1	NML	E	D	0	CL3-B	50060e8006fc3222	1	0	NML	64562 VSP N	M E	8
15		10		D 90									
5	1-1	NML	E	D	0	CL5-A	50060e8006fc4150	-	-	UNK	64562 VSP N	M E	8
15		10		D 90									
5	1-2	NML	E	D	2	CL3-B	50060e8006fc3222	-	-	UNK	64562 VSP N	A D	8
15		10		D 90									
5	1-2	NML	E	D	2	CL5-A	50060e8006fc4150	1	0	NML	64562 VSP N	A D	8
15		10		D 90									

Displays the detailed information about a path

```
# raidcom get path -key opt -I15
```

PHG	GROUP	STS	CM	IF	MP#	PORT	WWN	PR	LUN	PHS	Serial#	PRODUCT_ID	LB
PM	DM	QD	TO(s)	PBW(s)	PDBM	PDBT(%)							
0	1-1	NML	E	D	0	CL3-A	50060e8028005920	1	160	40	800089 VSP One	B20 N	
M	D	8	15	10		D 90							
0	1-1	NML	E	D	0	CL4-A	50060e8028005930	2	160	1	800089 VSP One	B20 N	
M	D	8	15	10		D 90							
0	1-1	NML	E	D	0	CL5-A	50060e8028005940	3	160	1	800089 VSP One	B20 N	
M	D	8	15	10		D 90							
0	1-1	NML	E	D	0	CL6-A	50060e8028005950	4	160	1	800089 VSP One	B20 N	
M	D	8	15	10		D 90							
1	1-2	NML	E	D	1	CL3-A	50060e8028005920	1	161	40	800089 VSP One	B20 N	
M	D	8	15	10		D 90							

Displays the detailed information about a path whose detailed status is 40

```
# raidcom get path -key opt | rmawk @L=eq:1 -o @11=eq:40
```

PHG	GROUP	STS	CM	IF	MP#	PORT	WWN	PR	LUN	PHS	Serial#	PRODUCT_ID	LB
-----	-------	-----	----	----	-----	------	-----	----	-----	-----	---------	------------	----


```

PM DM QD TO(s) PBW(s) PDBM PDBT(%)
0 1-1 NML E D 0 CL3-A 50060e8028005920 1 160 40 800089 VSP One B20 N M D 8 15 10 D 90
1 1-2 NML E D 1 CL3-A 50060e8028005920 1 161 40 800089 VSP One B20 N M D 8 15 10 D 90

```

Description of each column in output example 1:

PHG

Displays the path group number for the external volume.

GROUP

Displays the external volume group number.

STS

Displays the following status of the external volume group.

- NML: The external volume group is in the normal status.
- CHK: The external volume group is being checked.
- SYN: The cached data is being written to the volume on the external storage system.
- DSC: The connection to the volume on the external storage system is disconnected.
- BLK: The external path is blocked.
- UNK: The external volume group is in the unknown status.
- WAR: The state of the external volume group is at the warning level.

CM

Displays the cache mode for external volume group.

- E: Write cache enabled.
- D: Write cache disabled.
- EM: NDM attribute and Write cache enabled.
- DM: NDM attribute and Write cache disabled.
- TM: NDM attribute and cache through mode.
- SM: NDM attribute and Write sync mode.

IF

Displays the cache inflow control for external volume group.

E: Enable, D: Disable

MP#

Displays the MP blade ID for the external volume group.

PORT

Displays the port number on the local storage system.

WWN

Displays the target WWN of the port on the external storage system. For VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, when the external storage system is connected by iSCSI, this item displays the pseudo WWN of the external storage system.

PR

Displays the priority number in the external volume path group.

LUN

Displays the LUN in the port on the external storage system side.

PHS

Displays the following status of the external path if the `-key opt` option is not specified.

- NML: "Normal" status in external path.
- CHK: "temporary blockading " status in external path.
- BLK: "blockading" status in external path.
- DSC: "disconnected" status in external path.
- UNK: "Unknown" status in external path.

Displays the detailed status of the path in hexadecimal if the `-key opt` option is specified. For details about the meaning of the value, see the *Hitachi Universal Volume Manager User Guide*. A hyphen (-) is displayed if displaying the detailed status is not supported on the storage system.

Serial#

Displays the serial number of external storage system.

PRODUCT_ID

Displays the PRODUCT_ID of the external storage system.

LB

Displays the following I/O load balance mode for the external storage system.

- N: "normal round robin" mode.
- E: "extended round robin" mode.
- D: Executes with one path without load balance mode.

If Single is used for the path mode or the load balance mode is not supported, a hyphen (-) is displayed.

PM

Displays the path mode for the external storage system.

- M: Multiple path mode
- S: Single path mode
- A: APLB mode

- AL: ALUA mode
- MA: Multiple path mode (changeable to ALUA mode)
- SA: Single path mode (changeable to ALUA mode)

DM

Displays whether the data direct mapping attribute is set to the external volume group.

- E: The data direct mapping attribute is set.
- D: The data direct mapping attribute is not set.

QD

Displays the number of **read-write** commands that can be issued (queued) at a time to the volume on the external storage system. (For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900)

- - (hyphen): This information is invalid for this external path.

For VSP G1x00 and VSP F1500, a hyphen (-) is always displayed.

TO (s)

Displays the setting value of the I/O time over to the volume on the external storage system in seconds. (For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900)

- - (hyphen): This information is invalid for this external path.

For VSP G1x00 and VSP F1500, a hyphen (-) is always displayed.

PBW (s)

Displays the time in seconds after all external path connections in the external volume path group are disconnected until the external volume is blocked. For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900)

- - (hyphen): This information is invalid for this external path.

For VSP G1x00 and VSP F1500, a hyphen (-) is always displayed.

PDBM

Displays whether the path delay block mode function is enabled or disabled if a response delay occurs.

- E: The path delay block mode function is enabled.
- D: The path delay block mode function is disabled.
- - (hyphen): This information is invalid for this external path.

PDBT(%)

When the path delay block mode function is available, displays a percentage (25 through 100%) of an I/O timeout value that is used to determine a response delay. A hyphen (-) is displayed if this information is invalid.

Example 2

Waiting until the status of the external volume group #1-1 changes to DSC.

When the status has changed to DSC, the command ends with the return value 0. If the status does not change to DSC within 30 minutes, the command times out, and ends with the return value EX_EWSTOT.

```
# raidcom get path -external_grp_id 1-1 -check_status DSC -time 1800
```

Example 3

Waiting until the status of the external volume group with LDEV #0x10 changes to DSC.

When the status has changed to DSC, the command ends with the return value 0. If the status does not change to DSC within 30 minutes, the command times out, and ends with the return value EX_EWSTOT.

```
# raidcom get path -ldev_id 0x10 -check_status DSC -time 1800
```

Example 4

Checking if the external volume group #1-1 is in DSC status.

When the status is in DSC, the command ends with the return value 0. If the status is not in DSC, the command ends with the return value 1.

```
# raidcom get path -external_grp_id 1-1 -check_status DSC
```

Example 5

Executing the **raidcom disconnect external_grp** command to the external volume group #1-1, and then waiting until the status of the external volume group #1-1 changes to DSC.

```
# raidcom disconnect external_grp -external_grp_id 1-1
# raidcom get path -external_grp_id 1-1 -check_status DSC -time 1800
```

raidcom check_ext_storage path

Restores an external path to a volume on the external storage system. Only one path is operated in one operation.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

This command is rejected by EX_ENOOBJ in the following cases:

- The specified iSCSI port cannot be found.
- The iSCSI virtual port mode is enabled, but the specified iSCSI virtual port ID is not correct.

To fix this error, specify the correct port and iSCSI virtual port ID.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom check_ext_storage path -path_grp <path group#>
    -port <port#> {-external_wnn <wnn strings> |
    -external_iscsi_name <external iscsi name>
    -external_address <IP address>
    [-iscsi_virtual_port_id <iSCSI virtual port ID>]}
```

Options and parameters

-path_grp <path group#>

Specifies the external VOL path group number (0-63231).

-port <port#>

Specifies the port number. Specifies the number of the port whose attribute is ELUN (External). For example:

- CL1-A

-external_wnn <wnn strings>

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The 17th digit and later are ignored. The value can be split in units of 4 bytes by ",", (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-external_iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: *iqn.* and the subsequent maximum 219 characters.
- eui format: *eui.* and the subsequent 16 characters in hexadecimal notation.



Note:

You can specify the pseudo WWN of the iSCSI target to the `-external_wnn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

-external_address <IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).



Note:

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the **raidcom get external_iscsi_name** command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Examples

Restoring a path of external volume path group number: 1, External (ELUN) port CL1-A, and external storage system port 50060e80,05fa0f36.

```
# raidcom check_ext_storage path -path_grp 1 -port CL1-A
  -external_wwn 50060e80,05fa0f36
```

Restoring the path of the external volume path group number: 1, iSCSI port: CL1-A of the local storage system, and the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100).

```
# raidcom check_ext_storage path -path_grp 1 -port CL1-A -external_iscsi_name iqn.z2 -
external_address 158.214.135.100
```

Restoring a path of the external volume path group number: 1, the iSCSI port: CL1-A, iSCSI virtual port ID: 2 of the local storage system, and the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100).

```
# raidcom check_ext_storage path -path_grp 1 -port CL1-A -iscsi_virtual_port_id 2 -
external_iscsi_name iqn.z2 -external_address 158.214.135.100
```

raidcom delete path

Deletes the external path or alternative path to a volume on the external storage system. Only one path is operated in one operation.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

This command is rejected by EX_ENOOBJ in the following cases:

- The specified iSCSI port cannot be found.
- The iSCSI virtual port mode is enabled, but the specified iSCSI virtual port ID is not correct.

To fix this error, specify the correct port and iSCSI virtual port ID.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete path -path_grp <path group#> -port <port#>
{-external_wnn <wnn strings> | -external_iscsi_name
<external iscsi name> -external_address <IP address>
[-iscsi_virtual_port_id <iSCSI virtual port ID>]}
```

Options and parameters

-path_grp <path group#>

Specifies the external VOL path group number (0-63231).

-port <port#>

Specifies the Port number. Specifies the number of the port whose attribute is ELUN (External). For example:

- CL1-A

-external_wwn <wwn strings>

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The 17th digit and later are ignored. The value can be split in units of 4 bytes by ",", (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-external_iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: `iqn.` and the subsequent maximum 219 characters.
- eui format: `eui.` and the subsequent 16 characters in hexadecimal notation.

**Note:**

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

-external_address <IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

**Note:**

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the **raidcom get external_iscsi_name** command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Examples

Deleting a path of External (ELUN) port CL1-A, and external storage system port 50060e80,05fa0f36 from the external volume path group number: 1.

```
# raidcom delete path -path_grp 1 -port CL1-A -external_wwn 50060e80,05fa0f36
```

Deleting the path of the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100) and the iSCSI port: CL1-A of the local storage system from the external volume path group number: 1.

```
# raidcom delete path -path_grp 1 -port CL1-A -external_iscsi_name iqn.z2 -
external_address 158.214.135.100
```

Deleting the path of the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100), the iSCSI port: CL1-A, iSCSI virtual port ID: 2 of the local storage system from the external volume path group number: 1.

```
# raidcom delete path -path_grp 1 -port CL1-A -iscsi_virtual_port_id 2 -
external_iscsi_name iqn.z2 -external_address 158.214.135.100
```

raidcom disconnect path

Blocks the usage of external paths to the volume on the external storage system. Only one path is operated in one operation.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

This command is rejected by EX_ENOOBJ in the following cases:

- The specified iSCSI port cannot be found.
- The iSCSI virtual port mode is enabled, but the specified iSCSI virtual port ID is not correct.

To fix this error, specify the correct port and iSCSI virtual port ID.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom disconnect path -path_grp <path group#>
    -port <port#> {-external_wwn <wwn strings> |
    -external_iscsi_name <external iscsi name>
    -external_address <IP address>
    [-iscsi_virtual_port_id <iSCSI virtual port ID>]}
```

Options and parameters

-path_grp <path group#>

Specifies the external VOL path group number (0-63231).

-port <port#>

Specifies the port number. Specifies the number of the port whose attribute is ELUN (External). For example:

- CL1-A

-external_wwn <wwn strings>

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The 17th digit and later are ignored. The value can be split in units of 4 bytes by ",", (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-external_iscsi_name <external iscsi name>

Specifies the iSCSI name of the iSCSI target on the external storage system in an iqn format or an eui format.

- iqn format: *iqn.* and the subsequent maximum 219 characters.
- eui format: *eui.* and the subsequent 16 characters in hexadecimal notation.

**Note:**

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

-external_address <IP address>

Specifies the IP address of the iSCSI target on the external storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

**Note:**

You can specify the pseudo WWN of the iSCSI target to the `-external_wwn` option instead of specifying the iSCSI target by using the `-external_iscsi_name` option and the `-external_address` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Examples

Blocks the usage of a path of the external volume path group number: 1, the External (ELUN) port CL1-A, and the external storage port 50060e80,05fa0f36.

```
# raidcom disconnect path -path_grp 1 -port CL1-A -external_wnn 50060e80,05fa0f36
```

Blocks the usage of the path between the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100), and the iSCSI port: CL1-A of the local storage system, and the external volume path group number: 1.

```
# raidcom disconnect path -path_grp 1 -port CL1-A -external_iscsi_name iqn.z2 -
external_address 158.214.135.100
```

Blocks the usage of the path for the external storage system iSCSI target (iSCSI name: iqn.z2, IP address: 158.214.135.100), the iSCSI port: CL1-A, iSCSI virtual port ID: 2 of the local storage system, and the external volume path group number: 1.

```
# raidcom disconnect path -path_grp 1 -port CL1-A -iscsi_virtual_port_id 2 -
external_iscsi_name iqn.z2 -external_address 158.214.135.100
```

raidcom modify path

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Changes the values set for the external path to the volume on the external storage system.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If HIE or a logical path inside the storage is blocked, the operation cannot be performed.

Syntax

```
raidcom modify path -external_wnn <WWN strings> {-qdepth <Q depth>
| -timeout <timeout> | -path_blocked_watch <time(sec)>| -path_delay_block_mode
{enable | disable} | -path_delay_block_threshold <threshold %>}
```

Options and parameters**-external_wwn <WWN strings>¹**

Specifies the WWN (16-digit hexadecimal value) of the external storage system. The value can be split in units of 4 bytes by a "," (comma).

For example:

- 210000e08b0256f8
- 210000e0, 8b0256f8

-qdepth <Q depth>

Specifies the number of Read/Write commands that can be issued (queued) at a time for a volume on the external storage system. The same setting value is applied to all the external paths that use the WWN or iSCSI name of the specified external storage system.

-timeout <timeout>

Specifies a timeout value (5-240) in seconds for I/O to a volume on the external storage system.

-path_blocked_watch <time(sec)>

Specifies a time (5-180) (in seconds) from when all the external path connections in the external volume path group are disconnected until when the external volume group is blocked. The same setting value is applied to all the external paths that use the WWN or iSCSI name of the specified external storage system.

-path_delay_block_mode {enable | disable}²

Specifies whether the path delay block mode function is enabled or disabled. The default is disabled.

- enable: Enables the path delay block mode function.
- disable: Disables the path delay block mode function.

If enable is specified, a threshold value (time) is set based on an I/O timeout value specified in the `-timeout <timeout>` option and a percentage specified in the `-path_delay_block_threshold <threshold %>` option. When the response time from the external storage system exceeds the threshold value (time), if the response delay occurs 100 times per hour, the external path that uses the WWN specified in the `-external_wwn <wwn strings>` option becomes blocked.

For example:

If the I/O timeout value of 15 seconds and the percentage of 90% are specified, the threshold value is set to 13.5 seconds (15 seconds × 0.9 = 13.5 seconds).

-path_delay_block_threshold <threshold %>²

Specifies a percentage (25 through 100%) of the I/O timeout value specified in the `-timeout <timeout>` option. The default is 100%.

This option works when the `-path_delay_block_mode` option is enabled.

**Note:**

1. When using iSCSI connections, specify the pseudo WWN of the iSCSI target to the `-external_wwn` option. The pseudo WWN corresponds to the iSCSI name and IP address of the iSCSI target on the external storage system. By executing the `raidcom get external_iscsi_name` command, you can confirm the suitability between the pseudo WWN and the iSCSI name and IP address of the iSCSI target on the external storage system.
2. This function is supported by the following DKCMAIN firmware versions:
 - VSP One B20: A3-04-21-x0 or later
 - VSP E series: 93-07-27 or later
 - VSP 5000 series: 90-09-28 or later

For VSP 5000 series, this function is available only for open systems. It does not work on the external path for which the emulation type for the mainframe system is set.

Example

Specifying 15 as the timeout value for I/O to a volume on the external storage system, for all the external paths that use the port 50060e80 or 05fa0f36 of the external storage subsystem.

```
# raidcom modify path -external_wwn 50060e80,05fa0f36
-timout 15
```

raidcom delete pool

Deletes the specified Pool for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.

When LDEVs or device groups are specified, LDEVs are deleted from the specified pools.

If you want to delete an LDEV in a pool containing a DDP parity group, specify the parity group. The LDEV is deleted for each DDP parity group.

This command is executed asynchronously with the command input. Check the completion of this process on the `raidcom get command_status` command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If HIE or a logical path inside the storage is blocked, the operation cannot be performed.

Syntax

```
raidcom delete pool -pool {<pool ID#> | <pool naming>}
[-ldev_id <ldev#> |-grp_opt <group option> -device_grp_name <device group name>
[<device name>]| -parity_grp_id <gno-sgno> -ddp_ldev -request_id auto | -
delete_volume {yes|no}]
```

Options and parameters

-pool {<pool ID#> | <pool naming>}

Specifies the Pool ID (0-127) or pool name for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.

When you specify only a number, it is recognized as a pool ID. Therefore, to specify a pool whose name is a number, use the pool ID instead of the pool name.

-ldev_id <ldev#>

Specifies the LDEV number (0 to 65279), for example:

- -ldev_id 200

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of the LDEV (maximum 32 characters).

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

-parity_grp_id <gno-sgno> -ddp_ldev -request_id auto

Specifies the parity group number.

Specify this option to delete an LDEV belonging to a DDP parity group from a pool. For example: 3-1

-request_id is a request ID assigned each time the command is run. For details, see [Request ID function \(on page 293\)](#).

-delete_volume {yes|no}

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies whether to delete the pool volume when deleting a pool.

- yes (recommended): The pool volume is deleted when a pool is deleted.
- no (default): The pool is deleted, but the pool volume is not deleted.

The external volume is not deleted even if "yes" is specified. To delete it, first delete the pool using the **raidcom delete pool** command, and then execute the **raidcom delete ldev** command.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

Examples

Deleting a pool of pool ID 5:

```
# raidcom delete pool -pool 5
```

Deleting a pool of pool name "my_aou_pool":

```
# raidcom delete pool -pool my_aou_pool
```

Deleting a pool volume when deleting a pool of pool name "my_aou_pool":

```
# raidcom delete pool -pool my_aou_pool -delete_volume yes
```

Deleting (Shrinking) a pool volume belonging to a DDP parity group: 3-1 from a pool of pool name "my_aou_pool":

```
# raidcom delete pool -pool my_aou_pool -parity_grp_id 3-1 -ddp_ldev -request_id auto
```



Note: "Aou" (allocation on use) refers to Dynamic Provisioning.

raidcom get pool

Displays pool information for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, and active flash for mainframe.

Syntax

```
raidcom get pool [[-pool {<pool ID#> | <pool naming>}] -key <keyword>] [-fx] [-pcap] [-time_zone <time zone>]
```

Options and parameters

-pool {<pool ID#> | <pool naming>}

Specifies a pool ID (0-127) or a pool name for Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, or Dynamic Tiering.

If you specify only numeric characters, they can be recognized as the pool ID. Therefore, if you want to specify a pool whose name contains only numeric characters, specify the pool ID instead of the pool name.

This option can be specified only when the `-key resource` option is specified.

[-key <keyword>]

Specify this option to display the following information that can be specified for <keyword>:

- `opt`: Pool name.
- `basic`: Basic information about the pool.
- `fmc`: Information about parity groups that support accelerated compression.
FMC stands for flash module compression (a drive in a parity group that supports accelerated compression).
- `software_saving`: Information about the amount of capacity saved by deduplication and compression.

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

If the microcode version does not support this option, nothing is displayed when this option is specified.

If you use the capacity saving function, the saving ratio is calculated against the amount of metadata and garbage data generated internally in addition to user data.

- `saving`: Information about the amount of saved capacity.

- **total_saving:** Information about the amount of capacity saved by deduplication and compression, and accelerated compression of the parity group.

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- If the microcode version does not support this option, nothing is displayed when this option is specified.
 - If you use the capacity saving function, the saving ratio is calculated against the amount of metadata and garbage data generated internally in addition to user data.
 - This option is ignored when the microcode version does not support the option.
- **resource:** Information about the parity group configuring the pool or external volume group. When specifying this display keyword, specify the `-pool` option. Otherwise, the `-pool` option cannot be specified.

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

If the microcode version does not support this keyword, nothing is displayed when this keyword is specified.

- **efficiency:** Information about the data saving effect, snapshot effect, provisioning effect, and their total for the pool.

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

[-fx]

Displays the LDEV number in hexadecimal.

[-pcap]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Displays the information about the actual capacity. This option is valid only when the `-key basic` option is specified.

For pools for which accelerated compression is enabled, and for pools to which pool volumes can be added automatically, specify this option to check the capacity information.

The following GUI displays the actual capacity information as the capacity of the accelerated compression enabled pool. Specify this option to confirm the capacity equivalent to the capacity displayed in the GUI.

- Device Manager - Storage Navigator (VSP 5000 series, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900)
- Storage Advisor Embedded or VSP One Block Administrator

This option is ignored when the microcode version does not support the option.

[-time_zone <time zone>]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the time zone displayed in storage systems. If this option is omitted, the time in the time zone set for the storage system is displayed.

This option is valid only when it is specified with the `-key efficiency` option.

`utc`: Displays the time in Coordinated Universal Time (UTC).

Example1

Displaying pool information.

```
# raidcom get pool
```

PID	POLS	U(%)	SSCNT	Available(MB)	Capacity(MB)	Seq#	Num	LDEV#	H(%)	FMT_CAP(MB)
001	POLN	10	330	10000000	1000000000	62500	2	365	80	100
002	POLF	95	9900	100000	1000000000	62500	3	370	70	100
003	POLS	100	10000	100	1000000000	62500	1	375	70	100

```
004 POLE 0 0 0 0 62500 0 0 80 100
005 POLN 10 330 10000000 1000000000 62500 2 365 80 100
```

Example 2

Displaying pool name.

```
# raidcom get pool -key opt
```

```
PID POLS U(%) POOL_NAME Seq# Num LDEV# H(%) VCAP(%) TYPE PM
PT AUTO_ADD_PLV
001 POLN 10 my_aou_pool 62500 2 265 80 65500 OPEN S
HDP E
002 POLF 95 New_Pool_2 62500 3 270 70 65534 OPEN S
HDP -
003 POLS 100 my_ss_pool 62500 1 275 70 - OPEN N
TI D
004 POLN 0 New_Pool_4 62500 2 280 80 0 M/F N
CW -
005 POLE 0 New_Pool_5 62500 4 0 80 100 M/F S
DM -
```

Description of each column in output examples:

PID

Displays the pool ID.

POLS

Displays the status of the pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is blocked in the failure status.) In this status, the pool information cannot be displayed.

U(%)

Displays the usage rate of the pool.

SSCNT

Displays the number of volumes in the pool.

POOL_NAME

Displays the pool name.

Available (MB)

Displays the capacity (megabytes) available to the volume data in the pool.

Capacity (MB)

Displays the total capacity (megabytes) of the pool.

Seq#

Displays the product serial number (Seq#).

Num

Displays the number of LDEVs composing the pool.

LDEV#

Displays the first number of LDEV composing the pool. 65535(ffff) is displayed if the pool is being created.

H(%)

Displays threshold for the pool.

VCAP(%)

Displays the percentage of the subscription limit of V-VOL and Thin Image pair to the pool capacity. For VSP G130, G/F350, G/F370, G/F700, G/F900, a hyphen (-) indicating invalid is displayed.

TYPE

Displays the platform type of pools.

- Open: Shows that it is a Dynamic Provisioning pool.
- M/F: Shows that it is a Dynamic Provisioning for Mainframe pool.

PM

Displays the pool status.

- N: Normal status.
- S: Shrinking or rebalancing.
- NT: The pool for Thin Image is in the normal status.
- ST: (VSP G1x00 and VSP F1500, VSP Gx00 models, and VSP Fx00 models only)
The pool for Thin Image is shrinking or rebalancing.

PT

Displays the pool type. Any one of the following types is displayed.

- HDP: Pool for Dynamic Provisioning
- HDT: Pool for Dynamic Tiering
- RT: Pool for active flash
- TI: Pool for Thin Image
- CW: Pool for Copy-on-Write Snapshot
- DM: Pool for Dynamic Provisioning that has the data direct mapping attribute

FMT_CAP(MB)

Displays the formatted pool capacity (megabytes). If there is no valid information for the pool, a hyphen (-) is displayed.

AUTO_ADD_PLV

Displays whether to automatically add pool volumes according to the compression ratio of a parity group for which accelerated compression is enabled.

- E: Pool volumes are added automatically according to the compression ratio of the accelerated compression enabled parity group.
- D: Pool volumes in the parity group for which accelerated compression is enabled are not added automatically.
- - (hyphen): This information is invalid for this pool.

Example 3

Displaying basic information about a pool

```
# raidcom get pool -key basic
```

```
PID POLS U(%) LCNT SSCNT Available(MB) Capacity(MB) Snap_Used(MB)
TL_CAP(MB) BM TR_CAP(MB) RCNT Seq# Num LDEV# W(%) H(%) STIP
VCAP(%) TYPE PM PT POOL_NAME D_PM D_PM(%) EFFECTIVE(MB)
000 POLN 0 11001 11001 46998 46998 0
2432398 NB 0 0 300050 1 0 70 80 YES
UNLIMITED OPEN N HDP dp_ti_pool N - 46998
001 POLN 0 - 11001 46998 46998 -
- - - - 300050 1 1 - 80 -
- OPEN N TI tipool N - 46998
```

Displaying basic information about a pool. The actual capacity is displayed.

```
# raidcom get pool -key basic -pcap
```

```
PID POLS U(%) LCNT SSCNT Available(MB) Capacity(MB) Snap_Used(MB)
TL_CAP(MB) BM TR_CAP(MB) RCNT Seq# Num LDEV# W(%) H(%) STIP VCAP(%)
TYPE PM PT POOL_NAME D_PM D_PM(%) EFFECTIVE(MB)
000 POLN 0 11001 11001 46998 46998 -
2432398 NB 0 0 300050 1 0 70 80 YES -
OPEN N HDP dp_ti_pool N - 46998
001 POLN 0 - 11001 46998 46998
- - - - 300050 1 1 - 80
- - OPEN N TI tipool N - 46998
```

Description of each column in output example 3:**PID**

Displays the pool ID. Three digits padded with leading zeros.

POLS

Displays the status of the pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is blocked in the failure status.) In this status, the pool information cannot be displayed.

U(%)

Displays the usage rate of the pool.

Displays the actual usage rate of the pool when the `-pcap` option is specified.

LCNT

Displays the total number of Dynamic Provisioning virtual volumes mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

SSCNT

Displays the total number of snapshot data items mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

Available(MB)

Displays the available capacity (megabytes) for the volume data in the pool.

Displays the actual capacity that can be used for volume data in the pool when the `-pcap` option is specified.

Capacity(MB)

Displays the total capacity (megabytes) of the pool. When the `-pcap` option is specified, displays the total actual capacity of the pool.

Snap_Used(MB)

Displays the capacity (megabytes) used for Thin Image data. If the used capacity is less than 1 MB, the value is rounded up. A hyphen (-) is displayed if the information is not available for this pool.

Displays a hyphen (-) when the `-pcap` option is specified.

TL_CAP(MB)

Displays the total capacity (megabytes) of all Dynamic Provisioning virtual volumes and Thin Image pairs mapped to the pool. A hyphen (-) is displayed if the information is not available for this pool.

BM

Displays the I/O Blocking Mode of the pool.

- PF (Pool Full): If the pool is full, you cannot read from or write to the target DP-VOL. If the pool VOL is blocked, you can read from or write to the target DP-VOL.
- PB (Pool vol Blockade): If the pool VOL is blocked, you cannot read from or write to the target DP-VOL. If the pool is full, you can read from or write to the target DP-VOL.
- FB (Full or Blockade): If the pool is full or pool VOL is blocked, you cannot read from or write to the target DP-VOL.
- NB (No Blocking): If the pool is full or pool VOL is blocked, you can read from or write to the target DP-VOL.
- - (Not supported): The configuration does not support the I/O Blocking Mode.

TR_CAP(MB)

Displays the sum of the pool capacities (megabytes) reserved for the volumes for which Full Allocation or Proprietary Anchor is enabled. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

RCNT

Displays the number of volumes for which Full Allocation is enabled that are mapped to a pool. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

Seq#

Displays the serial number (Seq#).

Num

Displays the number of LDEVs belonging to the pool.

LDEV#

Displays the number of the first LDEV in the pool. "65535 (ffff)" is shown while the pool is being created.

W(%)

Displays the threshold value for WARNING set for the pool. A hyphen (-) is displayed if the information is not available for this pool.

H(%)

Displays the threshold value set for the pool as the high water mark.

STIP

Displays the setting for Thin Image pairs when the high water mark threshold is exceeded.

- YES: Thin Image pairs are suspended.
- NO: Thin Image pairs are not suspended.
- - (hyphen): The information is not available for this pool.

VCAP(%)

Displays the percentage of the subscription limit of V-VOL and Thin Image to the pool capacity. For VSP G130, G/F350, G/F370, G/F700, G/F900, a hyphen (-) indicating invalid is displayed.

- UNLIMITED: Unlimited.
- - (hyphen): The information is not available for this pool.

TYPE

Displays the platform type of the pool.

- OPEN: Pool for open systems
- M/F: Pool for mainframe systems

PM

Displays the pool status.

- N: Normal status.
- S: Shrinking or rebalancing.
- NT: The pool for Thin Image is in the normal status.
- ST: The pool for Thin Image is shrinking or rebalancing.

PT

Displays the pool type. One of the following types is displayed.

- HDP: Pool for Dynamic Provisioning
- HDT: Pool for Dynamic Tiering
- RT: Pool for active flash
- TI: Pool for Thin Image
- CW: Pool for Copy-on-Write Snapshot
- DM: Pool for Dynamic Provisioning that has the data direct mapping attribute

POOL_NAME

Displays the pool name.

D_PM

Processing status of the pool

If multiple statuses coexist, the statuses are displayed separated by commas (.). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

- N: Processing of the pool is not running.
- WS: The pool is waiting to shrink.
- S: The pool is shrinking.
- WRB: The pool is waiting to rebalance.
- RB: The pool is rebalancing.
- E: The pool is being expanded due to DDP expansion.
- PE: The pool expansion due to DDP expansion has been suspended.
- - (hyphen): The information is not available for this pool.

D_PM(%)

Percentage of progress of the processing status of the pool

If multiple statuses coexist, the percentage of progress for each status is displayed separated by a comma (.). A hyphen (-) is displayed if the DKCMAIN microcode version is not supported or if the percentage of progress is not valid.

EFFECTIVE(MB)

Effective capacity of the pool

A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

Example 4

Displaying information about an accelerated compression-enabled parity group that configures a pool. This information is not displayed if the functionality to show this information is not supported.

For details about the capacity saved by the capacity saving feature, see the *Provisioning Guide* for your storage system.

```
# raidcom get pool -key fmc
```

```
PID U(%) ACT_AV(MB) ACT_TP(MB) FR(%) FMC_LOG_USED(BLK)
FMC_ACT_USED(BLK) FMC_ACT_TP(BLK) FMC_PLV_USED(BLK) FMC_LOG_TP(BLK)
001 90 10000000 100000000 0 0
0 0 0 0
002 99 10000 100000000 200 100000000
100000000 51200000000 409608601600 409600000000
004 90 10000000 100000000 150 100000000
100000000 38400000000 307201720320 307200000000
```

Description of each column in output example 4:

PID

Displays the pool ID.

U(%)

Displays the actual usage rate of the pool.

ACT_AV(MB)

Displays the available actual capacity (megabytes) of the volume mapped to this pool for Dynamic Provisioning, Dynamic Provisioning for Mainframe, and Dynamic Tiering.

ACT_TP(MB)

Displays the total actual capacity (megabytes) of the pool.

FR(%)

Displays the percentage of the logical capacity against the actual capacity of an accelerated compression-enabled parity group in the pool. 0 is displayed if the pool does not contain an accelerated compression-enabled parity group.

FMC_LOG_USED(BLK)

Displays the logically used capacity of an accelerated compression-enabled parity group in the pool, in the unit of 512 bytes. 0 is displayed if the pool does not contain an accelerated compression-enabled parity group.

FMC_ACT_USED(BLK)

Displays the actually used capacity of an accelerated compression-enabled parity group, in the unit of 512 bytes. 0 is displayed if the pool does not contain an accelerated compression-enabled parity group.

FMC_ACT_TP(BLK)

Displays the total physical capacity of the pool volume that supports accelerated compression used in the pool. 0 is displayed if the pool does not contain a parity group that supports accelerated compression.

FMC_PLV_USED(BLK)

Displays the logically used page capacity of the pool volume that supports accelerated compression in the pool in the unit of 512 bytes. 0 is displayed if the pool does not contain a parity group that supports accelerated compression.

FMC_LOG_TP(BLK)

Displays the total logical capacity of the pool volume that supports accelerated compression used in the pool. 0 is displayed if the pool does not contain a parity group that supports accelerated compression.

Example 5

Displaying information about the reduced capacity of data in a pool.

```
# raidcom get pool -key saving
```

PID	SE_SAVING (BLK)	PL_SAVING (BLK)	PL_PRE_USED (BLK)	SES (%)	DDP (%)
CMP (%)	PLS (%)	Num	LDEV#		
000	12582912	9437184	28311552	20	15
001	0	0	- -	-	0

Description of each column in output example 5:**PID**

Displays the pool ID.

SE_SAVING(BLK)

Displays the capacity reduced by both capacity saving and accelerated compression of the parity group in number of blocks of the parity group.

The saved capacity does not include the amount of zero data reduction, metadata, garbage data, metadata reduction, and garbage data reduction.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility. See the description of `SE_SAVING(BLK)` displayed by `# raidcom get pool -key total_saving`.

PL_SAVING(BLK)

Displays the capacity reduced by capacity saving in number of blocks.

The saved capacity does not include the amount of zero data reduction, metadata, garbage data, metadata reduction, and garbage data reduction.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility. See the description of `PL_SAVING(BLK)` displayed by `# raidcom get pool -key software_saving`.

PL_PRE_USED(BLK)

Displays the capacity of data (prior to reduction) to be reduced by capacity saving in number of blocks.

The capacity prior to reduction does not contain the amount of zero data.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility. See the description of `PL_PRE_USED(BLK)` displayed by `# raidcom get pool -key software_saving`.

SES(%)

Displays the percentage of capacity reduced by both capacity saving and accelerated compression against the capacity used by the virtual volume mapped to this pool before the reduction.

The saved capacity does not include the amount of zero data reduction, metadata, garbage data, metadata reduction, and garbage data reduction.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility. See the description of `SES(%)` displayed by `# raidcom get pool -key total_saving`.

DDP(%)

Displays the percentage of capacity reduced by deduplication of capacity saving against the capacity used by the virtual volume mapped to this pool before the reduction.

The saved capacity does not include the amount of metadata and garbage data.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility.

CMP(%)

Displays the percentage of capacity reduced by compression of capacity saving or the percentage of capacity of the parity group reduced by accelerated compression against the capacity used by the virtual volume mapped to this pool before the reduction.

The saved capacity does not include the amount of metadata, garbage data, metadata reduction, and garbage data reduction.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility.



Note: The capacity used by the virtual volume before reduction, which is the denominator for calculating SES(%), DDP(%), CMP(%), and PLS(%), does not include the capacity used by or reserved for the deduplication system data volume.

PLS(%)

Displays the percentage of capacity reduced by capacity saving against the amount before the reduction, which is calculated as follows:

$$PL_SAVING(BLK) / PL_PRE_USED(BLK) \times 100 [\%]$$

The saved capacity does not include the amount of zero data reduction, metadata, and garbage data.

For VSP 5000 series, VSP F1500, VSP G1x00, VSP One B20, VSP E series, and VSP G/F350, G/F370, G/F700, G/F900, this information is provided for compatibility.

See the description of PLS(%) displayed by # **raidcom get pool -key software_saving**.

Num

Displays the number of deduplication system data volumes in the pool.

LDEV#

Displays the LDEV number of the deduplication system data volume in the pool. A hyphen (-) is displayed if the number of deduplication system data volumes is zero.

If there is more than one deduplication system data volume, LDEV numbers are separated by a space.

Example 6

Displaying information about the capacity saving effect achieved by the capacity saving function and accelerated compression of a parity group.

For details about the capacity saving effect achieved by the capacity saving feature and accelerated compression of a parity group, see the *Provisioning Guide* for your storage system.

```
# raidcom get pool -key total_saving
PID SE_SAVING(BLK) SES(%) DATAVOL_USED(BLK)
000 12582912 20 56623104
001 0 0 42467328
```

Description of each column in output example:

- **SE_SAVING(BLK)**: Displays the capacity saved by both capacity saving and accelerated compression of the parity group in blocks. The reduced amount is calculated as the difference between the amount of data before and after the capacity saving function is applied. The amount of data after capacity saving includes the amount of data (metadata and garbage data) that is generated internally by using the capacity saving function. The reduced amount also includes the amount of data reduced by reclaiming zero data pages.

A hyphen (-) indicating an invalid value might be displayed if the amount of used data volume before the reduction is smaller than the used pool capacity.

- **SES(%)**: Displays the percentage of capacity saved by both capacity saving and accelerated compression of the parity group against the capacity used by the virtual volume mapped to this pool before the reduction, which is calculated as follows:

$$\text{SE_SAVING(BLK)} / \text{DATAVOL_USED(BLK)} \times 100[\%]$$

When the capacity saving function is used, the saving ratio is calculated against the amount of metadata and garbage data generated internally in addition to user data.

A hyphen (-) indicating an invalid value might be displayed if the amount of used data volume before the reduction is smaller than the used pool capacity.

- **DATAVOL_USED(BLK)**: Displays the capacity used by a virtual volume in blocks. The capacity used by the deduplication system data volume is not included.

Example 7

Displaying information about the capacity of data stored in a pool saved by the capacity saving function.

For details about the capacity saved by the capacity saving feature, see the *Provisioning Guide* for your storage system.

```
# raidcom get pool -key software_saving
PID PLS(%) PL_SAVING(BLK) CMP(BLK) DDP(BLK) RECLAIM(BLK)
SYSTEM(BLK) PL_PRE_USED(BLK) PRE_CMP_USED(BLK) PRE_DDP_USED(BLK)
000 52 100931731456 0 64424509440 42949672960
6442450944 193273528320 0 193273528320
```

Description of each column in output example:

- PLS(%): Displays the percentage of capacity compressed by capacity saving against the amount of data before the compression, which is calculated as follows:

$$PL_SAVING(BLK) / PL_PRE_USED(BLK) \times 100[\%]$$

When the capacity saving function is used, the saving ratio is calculated against the amount of metadata and garbage data generated internally in addition to user data.

A hyphen (-) indicating an invalid value might be displayed if the amount of used data volume before the reduction is smaller than the used pool capacity.

- PL_SAVING(BLK): Displays the capacity saved by capacity saving in blocks. The reduced amount is calculated as the difference between the amount of data before and after the capacity saving function is applied. The amount of data after capacity saving includes the amount of data (metadata and garbage data) that is generated internally by using the capacity saving function. The reduced amount also includes the amount of data reduced by reclaiming zero data pages.

A hyphen (-) indicating an invalid value might be displayed if the amount of used data volume before the reduction is smaller than the used pool capacity.

- CMP(BLK): Displays the capacity compressed by capacity saving in blocks. The reduced amount is calculated as the difference between the amount of data before and after compression of the capacity saving function is applied. The amount of data after compression does not include the amount of data (metadata and garbage data) that is generated internally by using the capacity saving function.
- DDP(BLK): Displays the capacity saved by deduplication of capacity saving in blocks. The reduced amount is calculated as the difference between the amount of data before and after deduplication of the capacity saving function is applied. The amount of data after deduplication does not include the amount of data (metadata and garbage data) that is generated internally by using the capacity saving function.
- RECLAIM(BLK): Displays the capacity saved by reclaiming the fixed data pattern using capacity saving in blocks. The reduced amount is calculated as the difference between the amount of data before and after fixed data pattern reclamation of the capacity saving function is applied. The amount of data after reclaiming the fixed data pattern does not include the amount of data (metadata and garbage data) that is generated internally by the capacity saving function.
- SYSTEM(BLK): Displays the amount of consumed system data (metadata and garbage data) for capacity saving in blocks.
- PL_PRE_USED(BLK): Displays the capacity of data (before reduction) to be reduced by capacity saving in blocks.
- PRE_CMP_USED(BLK): Displays the capacity of data (before compression) to be compressed by capacity saving in blocks.
- PRE_DDP_USED(BLK): Displays the capacity of data (before deduplication) to be deduplicated by capacity saving in blocks.

Displaying information about the data saving effect , snapshot effect, provisioning effect, and their total for the pool.

For details about each effect, see the *Provisioning Guide* for your storage system.

```
#raidcom get pool -key efficiency
PID TOTAL_EFF_R TLS_R PLS_R PLS_CMP_R
PLS_DDP_R PLS_RECLAIM_R FMD_SAVING_R FMD_CMP_R
FMD_RECLAIM_R SNAPSHOT_EFF_R PROVISIONING_EFF(%)
CALCULATION_START CALCULATION_END
000 15.18 3.79 2.25 1.34
1.64 1.09 2.22 2.18
1.06 7.25 80
2015-05-15T10:05 2015-05-15T10:38
001 9.78 1.67 1.44 1.21
1.20 1.01 1.26 1.26
1.00 5.85 40
2015-05-15T10:05 2015-05-15T10:38
```




Note: The following values are the maximum of each saving effect ratio available to display in HDvM - SN or CCI:

- HDvM - SN: 99999999999999.99
- CCI: 92233720368547758.07 (a hundredth value of hexadecimal 0x7FFFFFFFFFFFFFFF.)

If no data is stored in the pool as shown in the following example, the saving effect ratio displays the maximum available value. If data is stored, the saving effect ratio calculated from the amount of actual data reduction is displayed. The following examples show the conditions where the maximum values are displayed:

Example 1. Total Efficiency Ratio after the pool and DP-VOL creations, but before data writing.

- HDvM - SN: 99999999999999.99
- CCI: 92233720368547758.07

Example 2. Snapshot Efficiency Ratio applied when the pool includes no Thin Image pairs, but only Thin Image Advanced pairs in the PAIR status or ones with no differential data.

- HDvM - SN: 99999999999999.99
- CCI: 92233720368547758.07

▪ TOTAL_EFF_R

Displays the ratio for the total (excluding system data) of the data saving effect, snapshot effect, and provisioning effect.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

▪ TLS_R

Displays the ratio of the data saving effect (excluding system data) achieved by the capacity saving feature and accelerated compression of the parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

▪ PLS_R

Displays the ratio of the data saving effect (excluding system data) of the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- PLS_CMP_R

Displays the ratio of the data saving effect (excluding system data) achieved by compression of the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- PLS_DDP_R

Displays the ratio of the data saving effect (excluding system data) achieved by deduplication of the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- PLS_RECLAIM_R

Displays the ratio of the saving effect (excluding system data) achieved by reclaiming the fixed data pattern using the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- FMD_SAVING_R

Displays the ratio of the data saving effect (excluding system data) achieved by accelerated compression of a parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- FMD_CMP_R

Displays the ratio of the data saving effect (excluding system data) achieved by compression of accelerated compression for a parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- FMD_RECLAIM_R

Displays the ratio of the data saving effect (excluding system data) achieved by reclaiming the fixed data pattern using accelerated compression of a parity group.

- NC: Calculate information is not available.
- - (hyphen): The information is not available for this pool.

- SNAPSHOT_EFF_R

Displays the ratio of the snapshot effect (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- PROVISIONING_EFF (%)

Displays the percentage of the provisioning effect (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- CALCULATION_START

Displays the start date and time for calculating the data saving effect, snapshot effect, provisioning effect, and their total (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

- CALCULATION_END

Displays the end date and time for calculating the data saving effect, snapshot effect, provisioning effect, and their total (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

Displaying information about the parity group configuring the pool and the external volume group.

```
# raidcom get pool -pool 0 -key resource
PID T GROUP AV_CAP(GB) TOTAL_CAP(GB) R_LVL R_TYPE TNO V E AUTO_ADD_PLV
000 R 5-2 300 300 RAID1 2D+2D 1 D D -
000 E 1-1 300 300 - - 2 D - -
```

Description of each column in output example:

- T: Displays the type of the volume group.
 - R: Parity group
 - E: External volume group
 - P: DDP
- Group: Displays the parity group number or the external volume group number.
- AV_CAP(GB): Displays the remaining available capacity (physically available area) (gigabytes) of the parity group or the external volume group. A value less than 1 GB is rounded down and 0 is displayed.
- TOTAL_CAP(GB): Displays the total capacity (gigabytes) that can be used by the parity group or external volume group.
- R_LVL: Displays the RAID level of the parity group or external volume group. A hyphen is displayed for external volume groups.
- R_TYPE: Displays the RAID type of the parity group or external volume group. A hyphen is displayed for external volume groups.
- TNO: Displays the number of the tier to which the parity group or external volume group belongs. The tier number is displayed only for Dynamic Tiering pools. A hyphen (-) is displayed for other pools.
- V: Displays whether capacity virtualization of the parity group is enabled or disabled.
 - E: Capacity virtualization is enabled.
 - D: Capacity virtualization is disabled.
- E: Displays whether encryption is enabled or disabled.
 - E: Encryption is enabled.
 - D: Encryption is disabled.
 - - (hyphen): Displaying encryption status is not supported.
- AUTO_ADD_PLV: Displays whether to automatically add pool volumes for which accelerated compression is enabled according to the compression ratio of a parity group for which accelerated compression is enabled.
 - E: Automatically adds pool volumes for which accelerated compression is enabled according to the compression ratio of a parity group for which accelerated compression is enabled.
 - D: Pool volumes for which accelerated compression is enabled are not added automatically.
 - - (hyphen): This information is invalid for this parity group. If accelerated compression is disabled for the parity group, a - (hyphen) is displayed.

raidcom modify pool

Sets the options of a pool for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe. This option also changes the setting of the data direct mapping attribute of the pool for Dynamic Provisioning.

When the `-status` or `-deduplication` option is specified, the operation is executed asynchronously with the command input. Check the completion of this process by using the `raidcom get command_status` command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify pool -pool {<pool ID#> | <pool naming>}
    [-status {nml|stop_shrinking} | -user_threshold <threshold_1>
    [<threshold_2>] | -tier <Tier number> [<ratio>]
    [-tier_buffer_rate <%>] | -subscription <%> |
    -pool_attribute <pool_attribute> | -monitor_mode <Monitor mode> |
    -blocking_mode <IO blocking mode> | -data_direct_mapping {enable|disable}
    | -deduplication yes -ldev_id <ldev#>... [-ssids <ssid> ...] |
    -deduplication no | -suspend_tipair {yes | no} |
    -delete dsd_volumes} | -auto_add_poolvol enable
    [-password <One Time Password>] | -auto_add_poolvol disable [-password
    <One Time Password>]
```

Options and parameters

-pool {<pool ID#> | <pool naming>}

Specifies the Pool ID (0-127) or pool name for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.

When only a number is specified, it is recognized as a pool ID. Therefore, to specify a pool whose name is a number, use the pool ID instead of the pool name.

-status <status>

Supported storage systems:

- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specifies the status of the pool. To restore the pool status, specify "nml".

-status {nml|stop_shrinking}

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the status of the pool:

- nml: Pool status is restored.
- stop_shrinking: Stops shrinking the pool.

-user_threshold <threshold_1> [<threshold_2>]

Sets a user-defined threshold.

- For Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, and active flash for mainframe, you may specify two user-defined thresholds, and their valid ranges are 1-100%, respectively. If you specify <threshold_1> and <threshold_2>:

the value of <threshold_1> is set as the threshold for WARNING specified to a pool.

the value of <threshold_2> is set as the threshold for High water mark specified to a pool.

If you specify only <threshold_1>, your specified value and the system default value (80%) are applied.

The setting in which you specify only <threshold_1> is supported to maintain backward compatibility with microcode before 70-02-0x-xx/xx. Once you specify two user-defined thresholds, you must continue to specify two user-defined thresholds thereafter.

- The valid range for Thin Image or Copy-on-Write Snapshot is 20-95%. You may specify only <threshold_1>. Even if you specify the value for <threshold_2>, the value is ignored.

-tier <Tier number> [<ratio>]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specifies the pool information for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe. This is specified when changing the free space percentage for a new allocation by each tier. When this option is specified, pool attribute changes into manual relocation.

<Tier number>: Tier number (1-3)

<ratio>: Free space percentage for new allocation(0-50) [%]

[-tier_buffer_rate <%>]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specifies the pool information for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe. Specifies the amount of reallocation buffer (2-40) for each tier in percent (%).

-subscription <%>

For VSP G130, G/F350, G/F370, G/F700, G/F900, an error occurs if you specify this option.

You can specify this option for VSP 5000 series, VSP G1x00, and VSP F1500, VSP One B20, VSP E series. However, for VSP 5000 series, VSP One B20, or VSP E series, an error occurs if a value other than 65535 (unlimited) is specified for the following pools:

- The pool that contains virtual volumes for which capacity saving feature, deduplication and compression are enabled.
- The pool that contains pool volumes that belongs to a parity group for which accelerated compression is enabled.

Sets the percentage (0-65535) of the subscription limit of V-VOL and Thin Image pair to the pool capacity.

- 0 to 65534: Specified percentage
- 65535: Unlimited

-pool_attribute <pool_attribute>

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specifies when changing the pool from Dynamic Provisioning/Dynamic Provisioning for Mainframe to Dynamic Tiering/Dynamic Tiering for Mainframe, or from Dynamic Tiering/Dynamic Tiering for Mainframe to Dynamic Provisioning/Dynamic Provisioning for Mainframe.

Specifies the following value depend on the type of pool desired to be changed.

- `dt_manual`: Changes the pool from Dynamic Provisioning/Dynamic Provisioning for Mainframe to Dynamic Tiering/Dynamic Tiering for Mainframe. It is changed to manual relocation.
- `dp`: Changes the pool from Dynamic Tiering/Dynamic Tiering for Mainframe to Dynamic Provisioning/Dynamic Provisioning for Mainframe.

-monitor_mode <monitor mode>

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specifies the monitoring mode for a Dynamic Tiering/Dynamic Tiering for Mainframe pool. Or specifies the availability of active flash.

- `period`: Performs monitoring periodically.
- `continuous`: Performs monitoring continuously.
- `realtime_tiering`: Enables active flash.
- `non_realtime_tiering`: Disables active flash.

-blocking_mode <IO blocking mode>

Sets the I/O activity (availability for read/write access) when the pool for Dynamic Provisioning, Dynamic Tiering, or active flash is full and when the pool for Dynamic Provisioning, Dynamic Tiering, or active flash is blocked. This option is ignored when the microcode version does not support the option.



Note: If the microcode version does not support this option, "-" is displayed in the BM column when you execute the `raidcom get dp_pool` command.

- `pool_full`: If the pool is full, read/write access for the target DP-VOL is disabled. If the pool is blocked, read/write access for the target DP-VOL is enabled.
- `pool_vol_blockade`: If the pool-VOL is blocked, read/write access for the target DP-VOL is disabled. If the pool-VOL is full, read/write access for the target DP-VOL is enabled.
- `full_or_blockade`: If the pool is full or blocked, read/write access for the target DP-VOL is disabled.
- `no_blocking`: If the pool is full and/or blocked, read/write access for the target DP-VOL is enabled.

-data_direct_mapping {enable| disable}

For VSP 5000 series, VSP G1x00, VSP F1500, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800: Changes the setting of the data direct mapping attribute of the pool for Dynamic Provisioning.

- enable or yes: Changes a pool for Dynamic Provisioning to a pool for Dynamic Provisioning that has the data direct mapping attribute.
- disable or no: Changes a pool for Dynamic Provisioning that has the data direct mapping attribute to a pool for Dynamic Provisioning.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

-deduplication yes -ldev_id <ldev#> ... [-ssids <ssid> ...]

Enables deduplication of the pool. When this option is specified, a deduplication system data volume that has an LDEV number (0 to 65279) specified for "-ldev_id <ldev#> ..." is created.

For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, an error occurs if you specify this option.

If you want to assign an SSID to the deduplication system data volume to be created, specify "-ssids <ssid> ..." in hexadecimal. SSIDs specified by "-ssids <ssid> ..." are assigned to deduplication system data volumes specified for "-ldev_id <ldev#> ..." in the order they were specified. If you omit this option, or specify auto for <ssid>, SSIDs are assigned automatically to the corresponding deduplication system data volumes.

The -ssids option can be specified only for VSP G1x00, VSP F1500.

If you specify multiple LDEVs, an error occurs because the current microcode does not support this function. If this happens, deduplication of the pool is not enabled.

-deduplication no

Disables deduplication of the pool. When this option is specified, a deduplication system data volume in this pool is deleted.

For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, an error occurs if you specify this option.

[-suspend_tipair {yes | no}]

Specifies or cancels the setting for suspending Thin Image pairs when the high water mark threshold is exceeded.

- yes: Suspends Thin Image pairs.
- no: Cancels the setting of suspending Thin Image pairs.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

-delete dsd_volumes

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Deletes the deduplication system data volume.

-auto_add_poolvol enable [-password <One Time Password>]

Contact customer support if you want to specify this option in the following storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

For VSP G1x00 and VSP F1500, "-password" does not need to be specified.

Automatically adds pool volumes for which accelerated compression of the parity group is enabled according to the compression ratio of the parity group for which accelerated compression is enabled.

-auto_add_poolvol disable [-password <One Time Password>]

Contact customer support if you want to specify this option in the following storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

For VSP G1x00 and VSP F1500, "-password" does not need to be specified.

Pool volumes for which accelerated compression of the parity group is enabled are not added automatically.

Examples

Restoring the status of a pool ID: 6.

```
# raidcom modify pool -pool 6 -status nml
```

Restoring the status of a pool name: my_ss_pool.

```
# raidcom modify pool -pool my_ss_pool -status nml
```

Changing the user-defined thresholds of the pool ID:6 of the pool for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe, WARNING to 70% and High water mark to 80%.

```
# raidcom modify pool -pool 6 -user_threshold 70 80
```

Changing the user-defined threshold of the pool ID: 6 of the pool for Thin Image or Copy-on-Write Snapshot to 80%.

```
# raidcom modify pool -pool 6 -user_threshold 80
```

Changing the free space percentage for a new allocation to the tier number 1 of the pool ID:6 for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe to 30%.

```
# raidcom modify pool -pool 6 -tier 1 30
```

Changing the amount of reallocation buffer to the tier number 1 of the pool ID:6 for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe to 20%.

```
# raidcom modify pool -pool 6 -tier 1 -tier_buffer_rate 20
```

Changing a Dynamic Provisioning/Dynamic Provisioning for Mainframe pool of the pool ID:6 to a Dynamic Tiering/Dynamic Tiering for Mainframe pool.

```
# raidcom modify pool -pool 6 -pool_attribute dt_manual
```

Changing a Dynamic Tiering/Dynamic Tiering for Mainframe pool of the pool ID:6 to a Dynamic Provisioning/Dynamic Provisioning for Mainframe pool.

```
# raidcom modify pool -pool 6 -pool_attribute dp
```

Changing the automatic relocation of the pool (Pool name: my_pool) for Dynamic Provisioning to manual relocation.

```
# raidcom modify pool -pool my_pool -pool_attribute dt_manual
```

Changing the monitoring mode of Dynamic Tiering/Dynamic Tiering for Mainframe pool of the pool ID: 6 to continuous.

```
# raidcom modify pool -pool 6 -monitor_mode continuous
```

When the pool is full, changing read/write of Dynamic Provisioning pool ID: 6 to rejected.

```
# raidcom modify pool -pool 6 -blocking_mode pool_full
```

Enabling active flash of Dynamic Tiering pool ID: 6.

```
# raidcom modify pool -pool 6 -monitor_mode realtime_tiering
```

Changing a pool (pool ID: 6) for Dynamic Provisioning to a pool for Dynamic Provisioning that has the data direct mapping attribute.

```
# raidcom modify pool -pool 6 -data_direct_mapping enable
```

Enabling deduplication of the pool (pool ID: 6) for Dynamic Provisioning to create a deduplication system data volume of LDEV: 400. For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, an error occurs.

```
# raidcom modify pool -pool 6 -deduplication yes -ldev_id 400
```

Enabling deduplication of the pool (pool ID: 6) for Dynamic Provisioning to create a deduplication system data volume of LDEV: 400, and assigning SSID: 0x0004 to the deduplication system data volume. For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, an error occurs.

```
# raidcom modify pool -pool 6 -deduplication yes -ldev_id 400 -ssids 0x0004
```

Disabling deduplication of the pool (pool ID: 6) for Dynamic Provisioning to delete a deduplication system data volume. For VSP 5000 series, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, an error occurs.

```
# raidcom modify pool -pool 6 -deduplication no
```

Deleting the deduplication system data volume of the pool (pool ID: 6) for Dynamic Provisioning.

```
# raidcom modify pool -pool 6 -delete dsd_volumes
```

Suspending Thin Image pairs when the high water mark threshold of the pool (pool ID: 6) for Dynamic Provisioning is exceeded.

```
# raidcom modify pool -pool 6 -suspend_tipair yes
```

Canceling the setting for suspending Thin Image pairs when the high water mark threshold of the pool (pool ID: 6) for Dynamic Provisioning is exceeded.

```
# raidcom modify pool -pool 6 -suspend_tipair no
```

Disabling the setting for automatically adding pool volumes for which accelerated compression is enabled in the pool (pool ID: 6).

```
# raidcom modify pool -pool 6 -auto_add_poolvol disable -password
<One Time Password>
```

Disabling the setting for automatically adding pool volumes for which accelerated compression is enabled in the pool (pool ID: 6).

```
# raidcom modify pool -pool 6 -auto_add_poolvol disable
```

Enabling the setting for automatically adding pool volumes for which accelerated compression is enabled in the pool (pool ID: 6).

```
# raidcom modify pool -pool 6 -auto_add_poolvol enable -password <One Time Password>
```

Enabling the setting for automatically adding pool volumes for which accelerated compression is enabled in the pool (pool ID: 6).

```
# raidcom modify pool -pool 6 -auto_add_poolvol enable
```

raidcom monitor pool

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Sets the start or stop of performance monitoring of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom monitor pool -pool {<pool ID#> | <pool naming>}
                        -operation <type>
```

Options and parameters

-pool {<pool ID#> | <pool naming>}

Specifies the pool ID (0-127) or pool name of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.

When specifying just a number, it is recognized as a pool ID. Therefore, when you specify a pool that the pool name is only a number, operate by specifying pool ID instead of specifying pool name.

-operation <type>

Instructs the operation of performance monitoring.

The operational types that can be specified are shown below.

- start: Start the performance monitoring.
- stop: Stop the performance monitoring.

Examples

Starting the performance monitoring of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool ID: 20.

```
# raidcom monitor pool -pool 20 -operation start
```

Starting the performance monitoring of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool Name: my_tier_pool.

```
# raidcom monitor pool -pool my_tier_pool -operation start
```

Stopping the performance monitoring of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool ID: 20.

```
# raidcom monitor pool -pool 20 -operation stop
```

Stopping the performance monitoring of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool Name: my_tier_pool.

```
# raidcom monitor pool -pool my_tier_pool -operation stop
```

raidcom reallocate pool

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Sets the start or stop of tier relocation of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom reallocate pool -pool {<pool ID#> | <pool naming>}
                        -operation <type>
```

Options and parameters

-pool {<pool ID#> | <pool naming>}

Specifies the pool ID (0-127) or pool name of a Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe pool.

When only a number is specified, it is recognized as a pool ID. Therefore, when you specify a pool whose name is only a number, use the pool ID instead of the pool name.

-operation <type>

Relocation operation instruction:

- start: Start the tier relocation.
- stop: Stop the tier relocation.

Examples

Starting the tier relocation of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool: 20.

```
# raidcom reallocate pool -pool 20 -operation start
```

Starting the tier relocation of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool Name: my_tier_pool.

```
# raidcom reallocate pool -pool my_tier_pool -operation start
```

Stopping the relocation of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool: 20.

```
# raidcom reallocate pool -pool 20 -operation stop
```

Stopping the relocation of a pool for Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe with Pool Name: my_tier_pool.

```
# raidcom reallocate pool -pool my_tier_pool -operation stop
```

raidcom rename pool

Changes the pool name of a pool for Thin Image, Copy-on-Write Snapshot, Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom rename pool -pool_id <pool ID#> -pool_name <pool naming>
```

Options and parameters

-pool_id <pool ID#>

Specifies a pool ID (0-127).

If a -pool_id option is non-numeric, the specified value is recognized as a pool name to identify the pool ID.

-pool_name <pool naming>

Specifies a new pool name. You can specify up to 32 characters.

The -pool_name option cannot be specified by configuring only in numeric because numeric values are preferentially identified as a pool ID. Specifies a pool ID with the '-pool_id<pool ID#>' option.

Examples

Changing the pool name of the pool (ID: 1) to my_pool.

```
# raidcom rename pool -pool_id 1 -pool_name my_pool
```

raidcom initialize pool

This command initializes deduplication system data volumes and volumes that contain deduplicated data. This command is processed asynchronously with the command input. Use the **raidcom get command_status** command to check if the **raidcom initialize pool** command completes.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom initialize pool -pool {<pool ID#> | <pool naming>} -operation <type>
```

Options and parameters

-pool {<pool ID#> | <pool naming>}

Specifies the pool ID (0 to 127) or pool name.

-operation <type>

Supported storage systems:

- VSP G1x00 and VSP F1500
- VSP G/F350, G/F370, G/F700, G/F900

Specifies the operation to perform. Specify the following:

`initialize_deduplication`: Initializes the deduplication system data volume specified by the `-pool` option and the volumes that contain deduplicated data.

-operation initialize_deduplication

Supported storage systems:

- VSP G200, G400, G600, G800, and VSP F400, F600, F800

Initializes the deduplication system data volume specified by the `-pool` option and the volumes that contain deduplicated data.

Examples

Initializing the deduplication system data volume and the volumes that contain deduplicated data in the pool with pool ID: 1.

```
#raidcom initialize pool -pool 1 -operation initialize_deduplication
```

raidcom modify system

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Configures the following settings:

- The CPU power saving settings of Dynamic Carbon Reduction
- The description setting on the storage system
- The setting to assign the MP unit to resources
- The expiration lock setting for Data Retention Utility

When configuring the setting to assign the MP unit to resources or the expiration lock setting for Data Retention Utility, this command is run asynchronously with the command input. Use the **raidcom get command_status** command to check if the command is complete.

To configure the setting to assign the MP unit to resources or the expiration lock setting for Data Retention Utility, lock all resource groups on the storage system.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify system -system_operation {set_description -description <description> |
reset_description | cpu_power_saving -cpu_power_saving_mode {enable | disable} -
request_id auto | cpu_power_saving -cpu_power_saving_monitor_time <time(min)> -
request_id auto | set_mpu_assignment -mp_blade_id <mp#> -mpu_auto_assignment {enable
| disable} | set_druexp_lock -druexp_lock {enable|disable}}
```

Options and parameters

-system_operation set_description -description <description>

Sets the description of the storage system using 1 to 255 characters. You can specify all characters that can be used for raidcom commands. For details about the characters, see [Supported characters \(on page 12\)](#).

-system_operation reset_description

Deletes the description of the storage system.

-system_operation cpu_power_saving -cpu_power_saving_mode {enable | disable} -request_id auto

Configures the setting to enable or disable the CPU power saving mode of Dynamic Carbon Reduction:

- enable: Enable the CPU power saving mode of Dynamic Carbon Reduction.
- disable: Disable the CPU power saving mode of Dynamic Carbon Reduction.

This option is only available for VSP One B20.

-system_operation cpu_power_saving -cpu_power_saving_monitor_time<time(min)> -request_id auto

Specifies the load monitoring period (10 through 180 minutes) for switching the CPU power saving status of Dynamic Carbon Reduction to POWER_SAVING.

This option is only available for VSP One B20.

-system_operation set_mpu_assignment -mp_blade_id <mp#> -mpu_auto_assignment {enable | disable}

Configures the setting to automatically assign the MP unit to resources.

-mp_blade_id <mp#> specifies the target MP unit ID.

-mpu_auto_assignment {enable | disable} specifies how the MP unit is assigned to resources. This option is specified only for VSP One B20.

If all MP units displayed are set to disable, all MP units work the same way as when they are set to enable.

- enable: The specified MP unit can be automatically assigned to resources.
- disable: The specified MP unit cannot be automatically assigned to resources.

-system_operation set_druexp_lock -druexp_lock {enable|disable}

Configures the expiration lock setting. This option is specified only for VSP One B20. When the expiration lock setting is enabled, the access attribute cannot be changed to read/write attribute even if the time periods during which changing to the read/write attribute is prohibited reach zero days of Data Retention Utility.

`-druexp_lock {enable|disable}` specifies either of the following settings:

- `enable` : The expiration lock is enabled.
- `disable` : The expiration lock is disabled.

Example 1

Setting the character string, "This is a storage system." as the description of the storage system:

```
raidcom modify system -system_operation set_description -description "This is a
storage system."
```

Example 2

Deleting the description of the storage system.

```
raidcom modify system -system_operation reset_description
```

Example 3

Enabling the CPU power saving mode of Dynamic Carbon Reduction.

```
# raidcom modify system -system_operation cpu_power_saving -cpu_power_saving_mode
enable -request_id auto
```

Example 4

Setting the automatic assignment of the MP unit ID 0 to resources.

```
# raidcom modify system -system_operation set_mpu_assignment -mp_blade_id 0 -
mpu_auto_assignment enable
```

Example 5

Setting the expiration lock to enable.

```
# raidcom modify system -system_operation set_druexp_lock -druexp_lock enable
```

raidcom get system

Displays the system status, data saving effect, snapshot effect, provisioning effect, availability of compression accelerator, and CPU power saving information of Dynamic Carbon Reduction.

Syntax

```
# raidcom get system [-key <keyword>] [-time_zone <time zone>]
```

Options and parameters

[-key <keyword>]

Specify the following keyword for <keyword>.

- efficiency (VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900)

Displays the data saving, snapshot, and provisioning effects of the system.

- information (VSP 5000 series, VSP One B20, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900)

Displays the information set for the storage system.

- dedupe_compression (VSP 5000 series, VSP One B20, VSP E1090)

Displays the information on the availability of compression accelerator.

- cpu_power_saving: Displays the information on the CPU power saving settings of Dynamic Carbon Reduction.

(This is only available for VSP One B20.)

- mpu: Displays the MP unit information.
- dru: Displays the Data Retention Utility information.
- replication: Displays the pair information.

[-time_zone <time zone>]

Specifies the time zone displayed in storage systems. If this option is omitted, the time in the time zone set for the storage system is displayed.

(VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) This option is valid only when it is specified with the -key efficiency option.

utc: Displays the time in Coordinated Universal Time (UTC).

Example 1

Displays the status of the system (VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, VSP F400, F600, F800).

```
#raidcom get system
Serial# : 12345
SYSTEM_TIME : 2016-09-21T08:30:20
AVE(W) : 500
MODEL : RH10HG
NUMOF_CACHE_MANAGEMENT_DEVICE : 640
MAX_CACHE_MANAGEMENT_DEVICE : 65280
```

Description of each column in output examples:**Serial#**

Displays the storage system serial number:

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00, VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

SYSTEM_TIME

Displays the controller system time when the acquisition instruction is issued. It is represented in yyyy-mm-ddThh:mm:ss format:

- yyyy (year): 4-digit decimal number in A.D.
- mm (month): 2-digit decimal number (01-12)
- dd (day): 2-digit decimal number (01-31)
- hh (hour): Time in 24 hours, 2-digit decimal number (00-23)
- mm (minute): 2-digit decimal number (00-59)
- ss (second): 2-digit decimal number (00-59)

AVE(W)

Displays the average power consumption per minute collected at the time of command execution.

If an invalid value is returned, a hyphen (-) is displayed.

If a failure occurs on the storage system, a value smaller than the actual power consumption or a hyphen (-) might be displayed.

MODEL

Displays the model information of storage system.

If the storage system model information display is not supported, - (hyphen) is displayed.

The following model information is displayed:

- R9F: VSP 5100, VSP 5500
- R9G: VSP 5100H, VSP 5500H
- RH10HF: VSP 5200, VSP 5600
- RH10HG: VSP 5200H, VSP 5600H

NUMOF_CACHE_MANAGEMENT_DEVICE

Displays the number of cache management devices used in the storage system. This item is not displayed if the storage system does not support display of this information.

MAX_CACHE_MANAGEMENT_DEVICE

Displays the maximum number of cache management devices in the storage system. This item is not displayed if the storage system does not support display of this information.

Example 2

Displays information about the data saving effect, snapshot effect, provisioning effect, and their total.

For details about each effect, see the *Provisioning Guide* for your storage system.

```
#raidcom get system -key efficiency
```

```
Serial# TOTAL_EFF_R TLS_R PLS_R PLS_CMP_R
PLS_DDP_R PLS_RECLAIM_R FMD_SAVING_R FMD_CMP_R
FMD_RECLAIM_R SNAPSHOT_EFF_R PROVISIONING_EFF(%)
CALCULATION_START CALCULATION_END
12345 11.64 1.84 1.52 1.15
1.34 1.11 2.21 2.14
1.08 10.37 70
2015-05-15T10:05 2015-05-15T10:38
```

Description of each column in output examples:



Note: The following values are the maximum of each saving effect ratio available to display in HDvM - SN or CCI:

- HDvM - SN: 99999999999999.99
- CCI: 92233720368547758.07 (a hundredth value of hexadecimal 0x7FFFFFFFFFFFFFFF.)

If no data is stored in the pool as shown in the following example, the saving effect ratio displays the maximum available value. If data is stored, the saving effect ratio calculated from the amount of actual data reduction is displayed. The following examples show the conditions where the maximum values are displayed:

Example 1. Total Efficiency Ratio after the pool and DP-VOL creations, but before data writing.

- HDvM - SN: 99999999999999.99
- CCI: 92233720368547758.07

Example 2. Snapshot Efficiency Ratio applied when the pool includes no Thin Image pairs, but only Thin Image Advanced pairs in the PAIR status or ones with no differential data.

- HDvM - SN: 99999999999999.99
- CCI: 92233720368547758.07

TOTAL_EFF_R

Displays the ratio for the total (excluding system data) of data saving effect, snapshot effect, and provisioning effect.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

TLS_R

Displays the ratio of the data saving effect (excluding system data) achieved by the capacity saving feature and accelerated compression of a parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

PLS_R

Displays the ratio of the data saving effect (excluding system data) achieved by the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

PLS_CMP_R

Displays the ratio of the data saving effect (excluding system data) achieved by compression of the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

PLS_DDP_R

Displays the ratio of the data saving effect (excluding system data) achieved by deduplication of the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

PLS_RECLAIM_R

Displays the ratio of the saving effect (excluding system data) achieved by reclaiming the fixed data patterns using the capacity saving feature.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

FMD_SAVING_R

Displays the ratio of the data saving effect (excluding system data) achieved by accelerated compression of a parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

FMD_CMP_R

Displays the ratio of the data saving effect (excluding system data) achieved by compression of accelerated compression for a parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

FMD_RECLAIM_R

Displays the ratio of the data saving effect (excluding system data) achieved by reclaiming the fixed data patterns using accelerated compression of a parity group.

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

SNAPSHOT_EFF_R

Displays the ratio of the snapshot effect (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

PROVISIONING_EFF(%)

Displays the percentage of the provisioning effect (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

CALCULATION_START

Displays the start date and time for calculating the data saving effect, snapshot effect, provisioning effect, and their total (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

CALCULATION_END

Displays the end date and time for calculating the data saving effect, snapshot effect, provisioning effect, and their total (excluding system data).

- NC: Calculated information is not available.
- - (hyphen): The information is not available for this pool.

Example 3

Displays the information set for the storage system.

```
#raidcom get system -key information
```

```
DESCRIPTION: This is a storage system.
```


Description of each column in output examples:**DESCRIPTION**

Displays the description of the storage system.

Example 4

Displays availability of compression accelerator.

```
#raidcom get system -key dedupe_compression
COMPRESSION_ACCELERATION: AVAILABLE
```

Description of each column in output examples:**COMPRESSION_ACCELERATION**

Displays availability of compression accelerator.

- AVAILABLE: Compression accelerator is available.
- UNAVAILABLE: Compression accelerator is unavailable.

Example 5

Displays the information on the CPU power saving settings of Dynamic Carbon Reduction.

```
#raidcom get system -key cpu_power_saving
AVE(W) : 500
CPU_POWER_SAVING_MODE : ENABLED
CPU_POWER_SAVING_MONITORING_TIME(M) : 10
CPU_STATUS(CTL01) : POWER_SAVING
CPU_STATUS(CTL02) : NON_POWER_SAVING
```

Description of each column in output examples:**AVE (W)**

Displays the average power consumption per minute collected when running the command. If an invalid value is returned, a hyphen (-) is displayed.

If an error occurs on the storage system, a value smaller than the actual power consumption or a hyphen (-) might be displayed.

CPU_POWER_SAVING_MODE

Displays the CPU power saving mode settings of Dynamic Carbon Reduction:

- ENABLED: The CPU power saving mode of Dynamic Carbon Reduction is enabled.
- DISABLED: The CPU power saving mode of Dynamic Carbon Reduction is disabled.
- - (hyphen): This information is not available.

CPU_POWER_SAVING_MONITORING_TIME(M)

Displays the load monitoring period (minute) for switching the CPU power saving status of Dynamic Carbon Reduction to POWER_SAVING. This output item is available only if the CPU power saving mode of Dynamic Carbon Reduction is enabled. If the CPU power saving mode is disabled, a hyphen is displayed.

CPU_STATUS(CTLxx)

Displays the CPU power saving status of Dynamic Carbon Reduction on CTLxx:

- POWER_SAVING: The CPU power saving status of Dynamic Carbon Reduction is enabled.
- NON_POWER_SAVING: The CPU power saving status of Dynamic Carbon Reduction is disabled.
- - (hyphen): This information is not available.

Example 6

Displays the DRU information configured on the storage system.

```
#raidcom get system -key dru
EXP_LOCK : ENABLED
```

Description of columns in output example:**EXP_LOCK**

Displays the status of the expiration lock setting:

- ENABLED: The expiration lock setting is enabled. The access attribute cannot be changed to read/write attribute even if the time periods during which changing to the read/write attribute is prohibited reach zero days.
- DISABLED: The expiration lock setting is disabled. The access attribute can be changed to read/write attribute even if the time periods during which changing to the read/write attribute is prohibited reach zero days.
- - (hyphen): This information is not valid.

Example 7

Displays the MP unit information.

```
# raidcom get system -key mpu
MP# AUTO_ASSIGNMENT
0 ENABLED
1 DISABLED
```

Description of columns in output example:**MP#**

Displays the MP unit ID.

AUTO_ASSIGNMENT

Displays the status of the setting to automatically assign the MP unit to resources. This item is not displayed if displaying this information is not supported on the storage system:

- **ENABLED:** The MP unit can be automatically assigned to resources.
- **DISABLED:** The MP unit cannot be automatically assigned to resources.

If all MP units displayed are set to **DISABLED**, all MP units work the same way as when they are set to **ENABLED**.

Example 8

Displays the information about differential tables or pair tables.

```
#raidcom get system -key replication
NUMOF_LOCAL_REPLICATION_DIFFERENTIAL_TABLE : 14
MAX_LOCAL_REPLICATION_DIFFERENTIAL_TABLE : 419200
NUMOF_LOCAL_REPLICATION_PAIR_TABLE : 14
MAX_LOCAL_REPLICATION_PAIR_TABLE : 32768
NUMOF_SNAPSHOT_PAIR_TABLE : 32
MAX_SNAPSHOT_PAIR_TABLE : 1048575
```

Description of columns in output example:**NUMOF_LOCAL_REPLICATION_DIFFERENTIAL_TABLE**

Displays the number of differential tables in ShadowImage and Volume Migration on the storage system. This item is not displayed if displaying this information is not supported on the storage system.

MAX_LOCAL_REPLICATION_DIFFERENTIAL_TABLE

Displays the maximum number of differential tables in ShadowImage and Volume Migration on the storage system. This item is not displayed if displaying this information is not supported on the storage system.

NUMOF_LOCAL_REPLICATION_PAIR_TABLE

Displays the number of pair tables in ShadowImage and Volume Migration on the storage system. This item is not displayed if displaying this information is not supported on the storage system.

MAX_LOCAL_REPLICATION_PAIR_TABLE

Displays the maximum number of pair tables in ShadowImage and Volume Migration on the storage system. This item is not displayed if displaying this information is not supported on the storage system.

NUMOF_SNAPSHOT_PAIR_TABLE

Displays the number of pair tables in Thin Image on the storage system. This item is not displayed if displaying this information is not supported on the storage system.

MAX_SNAPSHOT_PAIR_TABLE

Displays the maximum number of pair tables in Thin Image on the storage system.

This item is not displayed if displaying this information is not supported on the storage system.

raidcom get port

Displays port information.

This queries the setting information on all ports.

Syntax

```
raidcom get port [-key detail]
```

```
raidcom get port [-port <port> [-key <keyword> [-iscsi_virtual_port_id  
<iSCSI Virtual Port ID>]]]
```

```
raidcom get port -port <port> [-key <keyword> [-virtual_port_id <Virtual Port ID>]]
```

Options and parameters

[-port <port>]

Specifies the port number (for example, CL1-A). The port type you specify must be FIBRE, FCoE, or iSCSI. If you specify a port for which LUN security is enabled, the following items are displayed:

- If the target port type is FIBRE or FCoE: LOGIN_WWN
- If the target port type is iSCSI: LOGIN_ISCSI_NAME

[-key <keyword>]

Specifies one of the following values:

- **opt**: Displays the detailed information of FIBRE, FCoE, or iSCSI. To specify this option, you must specify the `-port <port>` option at the same time.
- **login_host_nqn**: Specifying this option for a port whose operation mode is NVMe mode displays the login status of the host NQN that is logged in to an NVMe subsystem port set on the target port. To specify this option, you must specify the `-port <port>` option at the same time.
- **login_mf_host**: Specifying this option displays the login status of the host when the type of the specified port is FICON. To specify this option, you must specify the `-port <port>` option at the same time. This option displays the information only about 32-Gbps FICON® ports (4Mx32(Mfibre) and 4Mx32E(MfibreE)). The information about the other ports is not displayed. To check the type of the CHB, see the **Port Condition** window in HDvM - SN.
- **detail**: Displays a column showing the additional information. To specify this option, you must not specify the `-port <port>` option at the same time.

[-iscsi_virtual_port_id <iSCSI Virtual Port ID>]

Displays detailed information when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, information about the virtual port ID: 0 is displayed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

[-virtual_port_id <Virtual Port ID>]

Displays the detailed information by specifying this option when the virtual port mode of the specified port is enabled. If the specified virtual port mode is enabled and this option is omitted, the information about the virtual port number: 0 is displayed.

An error occurs if you specify this option when any of the following conditions are met:

- The specified virtual port does not exist.
- The virtual port mode is disabled.
- The storage system does not support the virtual port mode.

Example 1

Displaying port information.

For a mainframe port (ESCON, FICON®), "-" is displayed for the values from SPD to SSW. If the transfer speed is displayed, SPD displays the setting value of the transfer speed for the FICON® port.

For a port used for Hitachi NAS (HNASS or HNASU), "-" is displayed for the values from ATTR to SSW and for WWN.

```
# raidcom get port
```

PORT	TYPE	ATTR	SPD	LPID	FAB	CONN	SSW	SL	Serial#	WWN	PHY_PORT
CL1-A	FIBRE	TAR	AUT	EF	N	FCAL	N	0	64568	50060e8006fc3800	-
CL1-B	FIBRE	TAR	AUT	EF	N	FCAL	N	0	64568	50060e8006fc3801	-
CL1-C	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3802	-
CL1-D	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3803	-
CL2-A	FIBRE	TAR	AUT	D9	N	FCAL	N	0	64568	50060e8006fc3810	-
CL2-B	FIBRE	TAR	AUT	D3	N	FCAL	Y	0	64568	50060e8006fc3811	-
CL2-C	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3812	-
CL2-D	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3813	-
CL3-A	FIBRE	MCU	AUT	E8	N	FCAL	N	0	64568	50060e8006fc3820	-
CL3-B	FIBRE	TAR	AUT	E0	N	FCAL	Y	0	64568	50060e8006fc3821	-
CL3-J	ISCSI	TAR	AUT	00	N	UNKN	N	0	64015	-	-
CL4-A	HNASS	-	-	-	-	-	-	0	64568	-	-
CL4-B	HNASU	-	-	-	-	-	-	0	64568	-	-

Description of each column in output example:

PORT

Displays the port numbers.

TYPE

Displays the following port type: FIBRE, SCSI, ISCSI, NVME_TCP, ENAS, ESCON, FICON[®], FCoE, HNASS, HNASU.

- HNASS means the port is for the system LU of Hitachi NAS.
- HNASU means the port is for a user LU of Hitachi NAS.

ATTR

Displays the following attribute setting on a port. One of the following items is displayed. If the port is a bidirectional port, the following four attributes are all displayed for each port. If you specify raidcom modify port-port_attribute TAR for the bidirectional port in VSP 5000 series, only the attribute TAR is displayed. When the raidcom modify port-port attribute ALL is specified, the following attributes are displayed.

- TAR: Fibre target port (target port)
- MCU: MCU initiator port (initiator port)
- RCU: RCU target port (RCU target port)
- ELUN: External initiator port (External port)



Note: If the port type is mainframe (TYPE is ESCON or FICON), the displayed port attribute settings indicate the following:

- TAR: HTP port (Hyper Transfer Processor port)
- MCU: FNP port (FICON Native Channel Processor port)

SPD

Displays the transfer rate setting on a port. One of the following values is displayed: AUT(AUTO), 1G, 2G, 4G, 8G, 10G...

LPID

Displays the ALPA (Arbitrated Loop Physical Address) setting value on a port.

FAB

Displays the fabric mode setting on a port as Y(YES) or N(NO).

CONN

Displays the following topology setting on a port: FCAL/PtoP/UNKN/-. If the port does not support the topology setting, UNKN or a hyphen is displayed.

SSW

Displays the LUN security setting on a port as Y (enabled) or N (disabled).

SL

Displays the SLPR number to which the port belongs.

Serial#

Product serial number (Serial#).

WWN

Displays the external WWN.

PHY_PORT

Displays the port number of the physical port in a resource ID takeover configuration.

Example 2

Displaying the additional port information

```
raidcom get port -key detail
```

PORT	TYPE	ATTR	SPD	LPID	FAB	CONN	SSW	SL	Serial#	WWN
PHY_PORT	PORT_MODE	RSGID	SECURITY							
CL1-A	FIBRE	TAR	AUT	EF	N	FCAL	N	0	64568	50060e8006fc3800
-	SCSI	0	-							
CL1-B	FIBRE	TAR	AUT	EF	N	FCAL	N	0	64568	50060e8006fc3801
-	NVME	0	-							
CL1-C	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3802
-	-	0	-							
CL1-D	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3803
-	-	0	ENFORCED							
CL2-A	FIBRE	TAR	AUT	D9	N	FCAL	N	0	64568	50060e8006fc3810
-	SCSI	0	-							
CL2-B	FIBRE	TAR	AUT	D3	N	FCAL	Y	0	64568	50060e8006fc3811
-	SCSI	0	-							
CL2-C	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3812
-	-	0	-							
CL2-D	FICON	TAR	-	-	-	-	-	0	64568	50060e8006fc3813
-	-	0	-							
CL3-A	FIBRE	MCU	AUT	E8	N	FCAL	N	0	64568	50060e8006fc3820
-	SCSI	1	-							
CL3-B	FIBRE	TAR	AUT	E0	N	FCAL	Y	0	64568	50060e8006fc3821
-	SCSI	1	-							
CL3-J	ISCSI	TAR	10G	00	N	UNKN	N	0	64015	-
-	-	1	-							
CL4-A	HNASS	-	-	-	-	-	-	0	64568	-
-	-	1	-							
CL4-B	HNASU	-	-	-	-	-	-	0	64568	-
-	-	1	-							

Description of each column in output example:**PORT_MODE**

Operation mode of the port. Displays one of the following:

- SCSI: SCSI mode
- ISCSI: iSCSI mode
- NVME: NVMe mode
- NVME_TCP: NVMe/TCP mode
- - (hyphen): The information is not available for this port.

RSGID

Displays the resource group ID.

SECURITY

Displays the endpoint security encryption setting for mainframe ports.

- ENFORCED: Endpoint security encryption is enabled. The connection cannot be established if endpoint security encryption is not supported on the host.
- ENABLED: Endpoint security encryption is enabled. The connection can be established even if endpoint security encryption is not supported on the host.
- DISABLED: Endpoint security encryption is disabled.
- – (hyphen): The information is not available for this port.

Example 3

Displaying the information of the port CL4-E (in case the TYPE is FIBRE or FCoE)

```
# raidcom get port -port CL4-E
```

```
PORT  LOGIN_WWN      Serial#  -
CL4-E 210000e08b0256f8 63528   OLA_NODE0_CTL_0
CL4-E 210000e08b039c15 63528   OLA_NODE1_CTL_0
```

Description of each column in output example:

PORT

Displays the port numbers.

LOGIN_WWN

Displays the WWN of the host adapter login to this port.

The WWN nickname specified in Performance Manager, or the SPM name specified in Server Priority Manager is displayed.

The WWN nickname specified with the **raidcom set hba_wwn** command is not displayed.



Note: Only the currently connected WWN is displayed unlike the display on Device Manager - Storage Navigator.

Serial#

Product serial number (Serial#).

Example 4

Displaying the information of the port CL4-E (in case of TYPE is ISCSI)

```
# raidcom get port -port CL4-E
```

```
PORT  LOGIN_IQN  Serial#  -
CL4-E iqn.zl...   63528   OLA_NODE0_CTL_0
```

```
CL4-E  iqn.z2...  63528  OLA_NODE1_CTL_0
```

Description of each column in output example:

PORT

Displays the port numbers.

LOGIN_IQN

Displays the iSCSI name for the host adapter log-in to this port that is currently being connected.

CCI displays only the iSCSI name of the port that is currently being connected, differently from Storage Navigator.

Serial#

Product serial number.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

Example 5

Example for getting information on FCoE option:

```
# raidcom get port -port CL4-E -key opt
```

```
PORT  ENMA          VLAN_ID  FPMA          VPS  VP_I  FCF_I
CL4-E  e3:00:00:e0:8b:02  0x03fe  e2:00:00:e0:8b:02  DWN  0x00  0x0000
```

Description of each column in output example:

PORT

Displays the port number.

ENMA

Displays Enode MAC address setting to this port.

VLAN_ID

Displays the VLAN identifier. This value is displayed even if Tag VLAN is disabled.

FPMA

Displays FP MAC address setting to this port.

VPS

Displays the virtual port status.

- DWN: the Link status is in Link Down.
- LOT: the Link status is in Link Up and Log-Out.
- LIN: the Link status is in Link Up and Log-In.

VP_I

Displays the virtual port index (zero is currently displayed).

FCF_I

Displays the FCoE index (zero is currently displayed).

Example 6

Displaying the example of iSCSI.

```
# raidcom get port -port CL4-E -key opt
```

```
PORT : CL4-E
TCP_OPT : IPV6_E : SACK_E : DACK_E : INS_E : VTAG_E
TCP_MTU : 1500
WSZ : 64KB
KA_TIMER : 30
TCP_PORT : 3260
IPV4_ADDR : 158.214.135.100
IPV4_SMSK : 255.255.255.255
IPV4_GWAD : 158.214.135.101
IPV6_ADDR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GADR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GWAD_INF : STS : fe80::209:6bff:febe:3c17 : fe80::209:6bff:febe:3c17
ISNS_PORT : 3260
ISNS_ADDR : 158.214.135.101
VLAN_ID : 0001
ISCSI_VP_MODE : E
ISCSI_VP_ID : 0 1 15
IPV6_GADR2_INF : STS : AM : fe80::209:6bff:febe:3c18
MAC_ADDR : e3:00:00:e0:8b:02
LINK_STATUS : UP
LINK_SPEED(Gbps) : 10
```

Description of each column in output example:**PORT**

Displays the port numbers.

TCP_OPT

Displays whether each option for iSCSI communication (IPv6 mode, Selective ACK mode, delayed ACK mode, iSNS mode, and Tag VLAN) is enabled or disabled. The meanings of the displayed values are as follows:

- IPV6_E: IPv6 mode is enabled.
- IPV6_D: IPv6 mode is disabled.
- SACK_E: Selective ACK mode is enabled.
- SACK_D: Selective ACK mode is disabled.
- DACK_E: Delayed ACK mode is enabled.
- DACK_D: Delayed ACK mode is disabled.
- INS_E: iSNS service is enabled.
- INS_D: iSNS service is disabled.
- VTAG_E: Tag VLAN is enabled.
- VTAG_D: Tag VLAN is disabled.

TCP_MTU

Displays the MTU value for iSCSI communication.

WSZ

Displays the window size for iSCSI communication.

KA_TIMER

Displays the Keep Alive Timer value for iSCSI communication.

TCP_PORT

Displays the TCP port number for iSCSI communication.

IPV4_ADDR

Displays IPv4 address.

IPV4_SMSK

Displays IPv4 subnet mask.

IPV4_GWAD

Displays IPv4 address of the gateway to use for iSCSI communication.

IPV6_ADDR_INF

Displays the status of IPv6 link local address, and the value of address acquiring mode and address. If the address information is available to display, the address is displayed even if IPv6 mode is disabled. The details of STS and AM in the displayed example are described in the table below.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.	INV	Invalid
		VAL	Valid
		ACQ	acquiring
		DUP	duplicated
AM	address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

IPV6_GADR_INF

Displays the status of IPv6 Global address, and the value of address acquiring mode and address. If the address information is available to display, the address is displayed even if IPv6 mode is disabled. The details of STS and AM in the displayed example are described in the table below.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.	INV	Invalid
		VAL	Valid
		ACQ	acquiring
		DUP	duplicated
AM	address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

IPV6_GWAD_INF

Displays the IPv6 Global address of the gateway to use for the iSCSI communication. If the address information is available to display, the address is displayed even if IPv6 mode is disabled. The values of address are displayed in order of address and current address. The details of STS in the displayed example are described in the table below.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.	INV	Invalid
		VAL	Valid
		ACQ	acquiring
		DUP	duplicated

ISNS_PORT

Displays the TCP port number of iSNS server. If iSNS is invalid, this item is not displayed.

ISNS_ADDR

Displays the address of iSNS server. If iSNS is invalid, this item is not displayed.

VLAN_ID

Displays VLAN ID in decimal. If VLAN ID is not set, a hyphen (-) is displayed. This value is displayed even if Tag VLAN is disabled.

ISCSI_VP_MODE

Displays the iSCSI virtual port mode status.

- E: Enabled
- D: Disabled

When the iSCSI virtual port mode is not supported, a hyphen (-) is displayed.

ISCSI_VP_ID

Lists the iSCSI virtual port IDs defined for the specified ports in decimal. When the iSCSI virtual port mode is disabled, this item is not displayed.

IPv6_GADR2_INF

Displays the IPv6 Global address2 status, and the value of address acquiring mode and address. Enabled for the virtual port of which physical port or iSCSI virtual port ID is 0. When the iSCSI virtual port ID is from 1 to 15, "- : - : -" is displayed. The details of STS and AM in the displayed example are described in the following table.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
	If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.	DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

MAC_ADDR

Displays the MAC address of the port. When the MAC address output function is not supported, a hyphen (-) is displayed.

LINK_STATUS

Displays the link status of the port. One the following status is displayed:

- UP: The port is link up status.
- DOWN: The port is not link up status.
- - (hyphen): The display of the link status is not supported.

LINK_SPEED(Gbps)

Displays the current transfer speed of the port in Gbps. If the display of the current transfer speed is not supported or the port link is not up, a hyphen (-) is displayed.

Example 7

Example of detailed information when the port is iSCSI and the iSCSI virtual port mode is enabled.

```
# raidcom get port -port CL4-E -key opt -iscsi_virtual_port_id 15
```

```
PORT : CL4-E
ISCSI_VP_ID : 15
TCP_OPT : IPV6_E : SACK_E : DACK_E : INS_E : VTAG_E
TCP_MTU : 1500
WSZ : 64KB
KA_TIMER : 30
TCP_PORT : 3260
IPV4_ADDR : 158.214.135.100
IPV4_SMSK : 255.255.255.255
IPV4_GWAD : 158.214.135.101
IPV6_ADDR_INF : STS : AM : fe80::209:6bff:febe:3c17
```

```

IPV6_GADR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GWAD_INF : STS : fe80::209:6bff:febe:3c17 : fe80::209:6bff:febe:3c17
ISNS_PORT : 3260
ISNS_ADDR : 158.214.135.101
VLAN_ID : 0001
IPV6_GADR2_INF : - : - : -

```

Description of each column in output example:

PORT

Displays the port numbers.

ISCSI_VP_ID

Displays the iSCSI virtual port ID.

TCP_OPT

Displays whether each option for iSCSI communication (IPv6 mode, Selective ACK mode, delayed ACK mode, iSNS mode, and Tag VLAN) is enabled or disabled. The meanings of the displayed values are as follows:

- IPV6_E: IPv6 mode is enabled.
- IPV6_D: IPv6 mode is disabled.
- SACK_E: Selective ACK mode is enabled.
- SACK_D: Selective ACK mode is disabled.
- DACK_E: Delayed ACK mode is enabled.
- DACK_D: Delayed ACK mode is disabled.
- INS_E: iSNS service is enabled.
- INS_D: iSNS service is disabled.
- VTAG_E: Tag VLAN is enabled.
- VTAG_D: Tag VLAN is disabled.

TCP_MTU

Displays the MTU value for iSCSI communication.

WSZ

Displays the window size for iSCSI communication.

KA_TIMER

Displays the Keep Alive Timer value for iSCSI communication.

TCP_PORT

Displays the TCP port number for iSCSI communication.

IPV4_ADDR

Displays IPv4 address.

IPV4_SMSK

Displays IPv4 subnet mask.

IPV4_GWAD:

Displays IPv4 address of the gateway to use for iSCSI communication.

IPV6_ADDR_INF

Displays the status of IPv6 link local address, and the value of address acquiring mode and address. If the address information is available to display, the address is displayed even if IPv6 mode is disabled. The details of STS and AM in the displayed example are described in the table below.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status*	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

* If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.

IPV6_GADR_INF

Displays the status of IPv6 Global address, and the value of address acquiring mode and address. If the address information is available to display, the address is displayed even if IPv6 mode is disabled. The details of STS and AM in the displayed example are described in the following table.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status*	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
* If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.			

IPv6_GWAD_INF

Displays the IPv6 Global address of the gateway to use for the iSCSI communication. If the address information is available to display, the address is displayed even if IPv6 mode is disabled. The values of address are displayed in order of address and current address. The details of STS in the displayed example are described in the following table.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status*	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
* If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.			

ISNS_PORT

Displays the TCP port number of the iSNS server. If iSNS is disabled, this item is not displayed.

ISNS_ADDR

Displays the address of the iSNS server. If iSNS is invalid, this item is not displayed.

VLAN_ID

Displays the VLAN ID in decimal. If the VLAN ID is not set, a hyphen (-) is displayed. This value is displayed even if Tag VLAN is disabled.

IPv6_GADR2_INF

Displays the IPv6 Global address2 status, and the value of address acquiring mode and address. Enabled for the virtual port of which physical port or iSCSI virtual port ID is 0. When the iSCSI virtual port ID is from 1 to 15, "- : - : -" is displayed. The details of STS and AM in the displayed example are described in the following table.

Item in the displayed example	Meanings of the item	Displayed value	Meanings of the value
STS	Status*	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

* If IPv6 mode is disabled, the displayed status depends on the status that IPv6 mode is enabled. This note can be ignored if IPv6 mode is disabled.

Example 8

Displaying the detailed NVMe/TCP information.

```
# raidcom get port -port CL4-E -key opt
PORT : CL4-E
TCP_OPT : IPV6_E : SACK_E : DACK_E : INS_E : VTAG_E
TCP_MTU : 1500
WSZ : 64KB
TCP_PORT_IO_CONTROLLER : 4420
TCP_PORT_DISCOVERY_CONTROLLER : 8009
IPV4_ADDR : 158.214.135.100
IPV4_SMSK : 255.255.255.255
IPV4_GWAD : 158.214.135.101
IPV6_ADDR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GADR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GWAD_INF : STS : fe80::209:6bff:febe:3c17 : fe80::209:6bff:febe:3c17
VLAN_ID : 0001
VP_MODE : E
VP_ID : 0 1 15
IPV6_GADR2_INF : STS : AM : fe80::209:6bff:febe:3c18
MAC_ADDR : e3:00:00:e0:8b:02
LINK_STATUS : UP
LINK_SPEED(Gbps) : 10
```

Description of each column in output example:

PORT

Displays the port numbers.

TCP_OPT

Displays whether each option (IPv6 mode, Selective ACK mode, Delayed ACK mode, iSNS mode, or Tag VLAN) is enabled or disabled during NVMe/TCP communications. The following list shows what the displayed values stand for:

- IPV6_E: IPv6 mode is enabled.
- IPV6_D: IPv6 mode is disabled.
- SACK_E: Selective ACK mode is enabled.
- SACK_D: Selective ACK mode is disabled.
- DACK_E: Delayed ACK mode is enabled.
- DACK_D: Delayed ACK mode is disabled.
- INS_E: iSNS service is enabled.
- INS_D: iSNS service is disabled.
- VTAG_E: Tag VLAN is enabled.
- VTAG_D: Tag VLAN is disabled.

TCP_MTU

Displays the MTU value during NVMe/TCP communications.

WSZ

Displays the window size value during NVMe/TCP communications.

TCP_PORT_IO_CONTROLLER

Displays the TCP port number for the storage access service during NVMe/TCP communications.

TCP_PORT_DISCOVERY_CONTROLLER

Displays the TCP port number for the discovery service during NVMe/TCP communications.

IPV4_ADDR

Displays the IPv4 address.

IPV4_SMSK

Displays the IPv4 subnet mask.

IPV4_GWAD

Displays the IPv4 address of the gateway used for NVMe/TCP communications.

IPV6_ADDR_INF

Displays the status of IPv6 link local address, the address acquiring mode, and the address values. The following table describes the details of STS and AM in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

IPV6_GADR_INF

Displays the status of IPv6 Global address, the address acquiring mode, and the address values. The following table describes the details of STS and AM in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

IPV6_GWAD_INF

Displays the IPv6 Global address values of the gateway used for NVMe/TCP communications. The address values are displayed in order of address and current address. The following table describes the details of STS in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated

VLAN_ID

Displays the VLAN ID in decimal. If the VLAN ID is not set, a hyphen (-) is displayed.

VP_MODE

Displays the NVMe/TCP virtual port mode status:

- E: Enabled
- D: Disabled

If the NVMe/TCP virtual port mode is not supported, a hyphen (-) is displayed.

VP_ID

Displays the NVMe/TCP virtual port ID.

IPV6_GADR2_INF

Displays the IPv6 Global address2 status, the address acquiring mode, and the address values. This item is available for the virtual port whose physical port ID or the NVMe/TCP virtual port ID is 0. If the NVMe/TCP virtual port ID is 1 through 15, "- : - : -" is displayed. The following table describes the details of STS and AM in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

MAC_ADDR

Displays the MAC address of the port. If the MAC address output function is not supported, a hyphen (-) is displayed.

LINK_STATUS

Displays the link status of the port. One of the following statuses is displayed:

- UP: The port is link up status.
- DOWN: The port is not link up status.
- - (hyphen): The display of the link status is not supported.

LINK_SPEED(Gbps)

Displays the current transfer speed of the port in Gbps. If the display of the current transfer speed is not supported or the port link is not up, a hyphen (-) is displayed.

Example 9

Displaying the detailed information when NVMe/TCP and the NVMe/TCP virtual port mode are enabled.

```
# raidcom get port -port CL4-E -key opt -virtual_port_id 15
PORT : CL4-E
ISCSI_VP_ID : 15
TCP_OPT : IPV6_E : SACK_E : DACK_E : INS_E : VTAG_E
TCP_MTU : 1500
WSZ : 64KB
TCP_PORT_IO_CONTROLLER : 4420
TCP_PORT_DISCOVERY_CONTROLLER : 8009
```

```

IPV4_ADDR : 158.214.135.100
IPV4_SMSK : 255.255.255.255
IPV4_GWAD : 158.214.135.101
IPV6_ADDR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GADR_INF : STS : AM : fe80::209:6bff:febe:3c17
IPV6_GWAD_INF : STS : fe80::209:6bff:febe:3c17 : fe80::209:6bff:febe:3c17
VLAN_ID : 0001
IPV6_GADR2_INF : - : - : -

```

Description of each column in output example:

PORT

Displays the port number.

ISCSI_VP_ID

Displays the iSCSI virtual port ID.

TCP_OPT

Displays whether each option (IPv6 mode, Selective ACK mode, Delayed ACK mode, iSNS mode, or Tag VLAN) is enabled or disabled during NVMe/TCP communications. The following list shows what the displayed values stand for.

- IPV6_E: IPv6 mode is enabled.
- IPV6_D: IPv6 mode is disabled.
- SACK_E: Selective ACK mode is enabled.
- SACK_D: Selective ACK mode is disabled.
- DACK_E: Delayed ACK mode is enabled.
- DACK_D: Delayed ACK mode is disabled.
- INS_E: iSNS service is enabled.
- INS_D: iSNS service is disabled.
- VTAG_E: Tag VLAN is enabled.
- VTAG_D: Tag VLAN is disabled.

TCP_MTU

Displays the MTU value during NVMe/TCP communications.

WSZ

Displays the window size value during NVMe/TCP communications.

TCP_PORT_IO_CONTROLLER

Displays the TCP port number for the storage access service during NVMe/TCP communications.

TCP_PORT_DISCOVERY_CONTROLLER

Displays the TCP port number for the discovery service during NVMe/TCP communications.

IPV4_ADDR

Displays the IPv4 address.

IPV4_SMSK

Displays the IPv4 subnet mask.

IPV4_GWAD

Displays the IPv4 address of the gateway used for NVMe/TCP communications.

IPV6_ADDR_INF

Displays the status of IPv6 link local address, the address acquiring mode, and the address values. The following table describes the details of STS and AM in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

IPV6_GADR_INF

Displays the status of IPv6 Global address, the address acquiring mode, and the address values. The following table describes the details of STS and AM in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

IPV6_GWAD_INF

Displays the IPv6 Global address values of the gateway used for NVMe/TCP communications. The address values are displayed in order of address and current address. The following table describes the details of STS in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated

VLAN_ID

Displays the VLAN ID in decimal. If the VLAN ID is not set, a hyphen (-) is displayed.

IPV6_GADR2_INF

Displays the IPv6 Global address2 status, the address acquiring mode, and the address values. This item is available for the virtual port whose physical port ID or the NVMe/TCP virtual port ID is 0. If the NVMe/TCP virtual port ID is 1 through 15, "- : - : -" is displayed. The following table describes the details of STS and AM in the displayed example.

Item in the displayed example	Meaning of the item	Displayed value	Meaning of the value
STS	Status	INV	Invalid
		VAL	Valid
		ACQ	Acquiring
		DUP	Duplicated
AM	Address acquiring mode	AM	Automatic acquiring
		MM	Manual acquiring

Example 10

Displaying the example of FIBRE.

```
# raidcom get port -port CL4-E -key opt
```

```
PORT  S  LKNSPD  CURADR  T
CL4-E  U      16  821A00  D
```

Description of each column in output example:**PORT**

Displays the port numbers.

S

Displays the link status of the port.

- U: The port is link up status.
- D: The port is not link up status.
- - (hyphen): The display of the link status is not supported. The display of **LKNSPD** and **CURADR** are also not supported.

The item is displayed for VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800.

LKNSPD

Displays the current transfer speed of the port in Gbps. If the display of the current transfer speed is not supported or the port link is not up, a hyphen (-) is displayed.

CURADR

Displays the current port address of the port in hexadecimal number. If the display of the current port address is not supported or the port link is not up, a hyphen (-) is displayed.

T

Displays the setting of the T10 PI mode of the port.

- E: T10 PI mode is enabled.
- D: T10 PI mode is disabled.
- - (hyphen): T10 PI is not supported.

Example 11

Displaying the login status of the host NQN that is logged in to the port CL4-E.

```
# raidcom get port -port CL4-E -key login_host_nqn
PORT LOGIN_STATUS HOST_NQN
CL4-E LOGIN
nqn.2014-08.org.nvmexpress:uuid:24ffc04c-74ed-413d-9261-0a40bfb4585f
```

```
CL4-E LOGOUT
nqn.2014-08.org.nvmexpress:uuid:24ffc04c-74ed-413d-9261-0a40bfb4584f
```

Description of each column in output example:

HOST_NQN

Displays Host NQN.

Example 12

Displaying the login status of the host that is logged in to the port CL4-E.

```
# raidcom get port -port CL4-E -key login_mf_host
PORT    LOGIN_TYPE          HOST_TYPE  MODEL    SERIAL_NUMBER    TAG
CL4-E   DISABLED_UNAVAILABLE  002964    N63      8400000009EDF7   807E
```

Description of each column in output example:

LOGIN_TYPE

Displays the login type of the connected host.

- ENCRYPTED: Indicates that the host is logged in with authentication and encryption, and endpoint security encryption is enabled on the mainframe port.
- AUTHENTICATED: Indicates that the host is logged in with authentication and without encryption, and endpoint security encryption is enabled on the mainframe port.
- DISABLED_AVAILABLE: Indicates that the host is logged in without authentication and without encryption but endpoint security encryption is enabled on the mainframe port.
- DISABLED_UNAVAILABLE: Indicates that the host is logged in with endpoint security encryption disabled on the mainframe port or on the host.

HOST_TYPE

Displays the type of the connected host.

MODEL

Displays the model name of the connected host.

SERIAL_NUMBER

Displays the serial number of the connected host.

TAG

Displays the tag of the connected host.

raidcom modify port

Sets the attribute of the specified port.

When you set a port attribute by options other than a `-port_attribute` option, the port type of the specified port must be one of the following:

- FIBRE
- FCoE
- iSCSI
- NVMe/TCP
- FICON[®]

If these conditions are not satisfied, this command is rejected with EX_ENOOBJ.

When using an FCoE package, specify Port speed: 10G, Port topology: f_port. In this case, the port attribute cannot be changed.

When you set the T10 PI mode or specify the `-request_id` option, this command is executed asynchronously with the command input. Check the completion of this process using the `raidcom get command_status` command.

There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

If the port type is FICON[®], you can change the following settings for 32-Gbps FICON[®] ports:

- 4Mx32E(MfibreE) port: The port speed and the endpoint security encryption setting can be changed.
- 4Mx32(Mfibre) port: Only the port speed can be changed.

To check the CHB type, see the **Port Condition** window of Device Manager - Storage Navigator. For more information about how to display the Port Condition window and the contents displayed in the window, see the *System Administrator Guide*.

For more information about how to display the **Port Condition** window and the contents displayed in the window, see the *System Administrator Guide*.

The syntax is separated as follows. If you execute a command that includes options in syntax 3 with options in syntax 2, and syntaxes 4 through 9, an error occurs.

Syntax

Syntax 1

```
raidcom modify port -port <port#>{[-port_speed <value>]
  [-loop_id <value>][-topology <topology>][-security_switch
  {y | n}] | -port_attribute <port attribute> | -t10pi {enable|disable}
  | -iscsi_virtual_port_mode {enable|disable}| -port_mode {scsi | nvme} -request_id
  auto |
  -delete_login_host_nqn}
```

Syntax 2 (Available for iSCSI packages or NVMe/TCP packages)

```
raidcom modify port -port <port#> [-mtu <value>]
  [-vlan_tagging_mode {enable|disable}] [-add_vlan_id <value>]
  [-delete_vlan_id <value>] [-modify_vlan_id <new value>]
```

```
[-ipv4_address <address>] [-ipv4_subnetmask <subnet mask>]
[-ipv4_gateway_address <address>] [-ipv6_mode {enable|disable}]
[-ipv6_local_address {auto|<address>}]
[-ipv6_global_address {auto|<address>}]
[-ipv6_gateway_address <address>] [-tcp_port <value>]
[-tcp_port_io_controller <value>] [-tcp_port_discovery_controller <value>]
[-selective_ack_mode {enable|disable}]
[-delayed_ack_mode {enable|disable}]
[-window_size <size>] [-keep_alive_timer <value>]
[-ipv6_global_address2 <address>]
```

Syntax 3 (Available only for iSCSI packages)

```
raidcom modify port -port <port#> [-isns_mode {enable|disable}]
[-isns_server_address <IPv4 address|IPv6 address>]
[-isns_port <iSNS TCP Port number>]
```

Syntax 4 (Available only for iSCSI packages)

```
raidcom modify port -port <port> -add_iscsi_virtual_port <value>
-ipv6_mode {enable|disable} [-mtu <value>]
[-vlan_tagging_mode {enable|disable}] [-add_vlan_id <value>]
[-ipv4_address <address>] [-ipv4_subnetmask <subnet mask>]
[-ipv4_gateway_address <address>]
[-ipv6_local_address {auto|<address>}]
[-ipv6_global_address {auto|<address>}]
[-ipv6_gateway_address <address>] [-tcp_port <value>]
[-selective_ack_mode {enable|disable}]
[-delayed_ack_mode {enable|disable}]
[-window_size <size>] [-keep_alive_timer <value>]
```

Syntax 5 (Available only for iSCSI packages)

```
raidcom modify port -port <port> -modify_iscsi_virtual_port <value>
[-mtu <value>] [-vlan_tagging_mode {enable|disable}]
[-add_vlan_id <value>] [-delete_vlan_id <value>]
[-modify_vlan_id <new value>]
[-ipv4_address <address>] [-ipv4_subnetmask <subnet mask>]
[-ipv4_gateway_address <address>] [-ipv6_mode {enable|disable}]
[-ipv6_local_address {auto|<address>}] [-ipv6_global_address {auto|<address>}]
[-ipv6_gateway_address <address>] [-tcp_port <value>]
[-selective_ack_mode {enable|disable}] [-delayed_ack_mode {enable|disable}]
[-window_size <size>] [-keep_alive_timer <value>]
[-ipv6_global_address2 <address>]
```

Syntax 6 (Available only for iSCSI packages)

```
raidcom modify port -port <port> -delete_iscsi_virtual_port <value>
```

Syntax 7 (Available for iSCSI packages or NVMe/TCP packages)

```
raidcom modify port -port <port> -add_virtual_port <value>
-ipv6_mode {enable|disable} [-mtu <value>] [-vlan_tagging_mode {enable|disable}]
[-add_vlan_id <value>]
[-ipv4_address <address>] [-ipv4_subnetmask <subnet mask>]
[-ipv4_gateway_address <address>]
[-ipv6_local_address {auto|<address>}]
[-ipv6_global_address {auto|<address>}]
[-ipv6_gateway_address <address>] [-tcp_port <value>]
[-tcp_port_io_controller <value>] [-tcp_port_discovery_controller <value>]
[-selective_ack_mode {enable|disable}]
[-delayed_ack_mode {enable|disable}]
[-window_size <size>] [-keep_alive_timer <value>]
```

Syntax 8 (Available for iSCSI packages or NVMe/TCP packages)

```
raidcom modify port -port <port> -modify_virtual_port <value>
[-mtu <value>] [-vlan_tagging_mode {enable|disable}]
[-add_vlan_id <value>] [-delete_vlan_id <value>]
[-modify_vlan_id <new value>]
[-ipv4_address <address>] [-ipv4_subnetmask <subnet mask>]
[-ipv4_gateway_address <address>] [-ipv6_mode {enable|disable}]
[-ipv6_local_address {auto|<address>}]
[-ipv6_global_address {auto|<address>}]
[-ipv6_gateway_address <address>] [-tcp_port <value>]
[-tcp_port_io_controller <value>] [-tcp_port_discovery_controller <value>]
[-selective_ack_mode {enable|disable}]
[-delayed_ack_mode {enable|disable}]
[-window_size <size>] [-keep_alive_timer <value>]
[-ipv6_global_address2 <address>]
```

Syntax 9 (Available for iSCSI packages or NVMe/TCP packages)

```
raidcom modify port -port <port> -delete_virtual_port <value>
```

Syntax 10

```
raidcom modify port -port <port> {-port_speed <value> | -endpoint_security {enforce|
enable|disable} [-forcible] -request_id auto}
```

Options and parameters

-port <port#>

Specifies the port number. For example:

- CL1-A

[-port_speed <value>]

Specifies the port speed (0, 1, 2, 4, 8, 10, 16, 25, 32, 64, 100) as follows.

- 0: AUTO
- 1: 1G
- 2: 2G
- 4: 4G
- 8: 8G
- 10: 10G
- 16: 16G
- 25: 25G
- 32: 32G
- 64: 64G
- 100: 100G

The specifiable values differ depending on the type of the channel board or channel adapter.

If the target port is either of the following, the value you specify using this option is ignored even if you specify a valid value.

- For VSP 5000 series, the port type is iSCSI.
- For VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, the type of the channel board or channel adapter for the target port is 10iSCSI2o(CHB) (10G iSCSI(Optic)) respectively.

If the port type is iSCSI, AUTO can only be set for ports on channel boards or channel adapters whose type is 10iSCSI2c(CHB) (10G iSCSI(Copper)). For more information about the available port speeds, see the *Provisioning Guide* for your storage system.

To check the type of your channel board or channel adapter, see the **Port Condition** window of Device Manager - Storage Navigator or the **Controller Chassis** window of the maintenance utility.

[-loop_id <value>]

Specifies the Loop ID (0x01-0xEF) of the Port.

If the port type is iSCSI, the value you specify using this option is ignored even if you specify a valid value.

[-topology <topology>]

Specifies the topology of the Port as follows:

- fl_port: fabric on and fcal
- f_port: fabric on and PtoP
- nl_port: fabric off and fcal
- n_port: fabric off and PtoP

When you change the topology setting of a port, delete the remote path and external path defined for that port beforehand. If an LU path is defined for the port, make sure in advance that the host is not performing I/O processing and the state is not the host reserve (mounted).

If the port type is iSCSI, the value you specify using this option is ignored even if you specify a valid value.

[-security_switch {y | n}]

Specifies whether to enable LUN security as follows:

- y or yes: LUN security is enabled.
- n or no: LUN security is disabled.

-port_attribute <port attribute>

Specifies the Port attribute as follows:

- TAR: Fibre target port (target port)
- MCU: MCU initiator port (initiator port)
- RCU: RCU target port (RCU target port)
- ELUN: External initiator port (External port)
- ALL: Resets (removes the TAR attribute from) a bidirectional port.

If the port is a bidirectional port of VSP 5000 series, the storage system detects the appropriate attribute of the port, and then operates the port with the detected attribute. If TAR is specified as the attribute value for the bidirectional port, the port attribute is changed to TAR (Target), and the storage system treats the port as a target port. By specifying ALL as the attribute value for the port specified in the TAR attribute, the user can restore the operation with the appropriate port attribute.

If the port is a bidirectional port of VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, you cannot change the attribute value. The storage system detects the appropriate attribute of the port, and then operates the port with the detected attribute.

If the package type of VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800 is iSCSI, the user cannot set the port attribute.

If this port attribute is changed from Target or RCU Target to Initiator or External, the host group belonging to this port belongs to meta_resource.

-t10pi {enable|disable}

Specifies the T10 PI mode as follows:

- enable: T10 PI mode is enable
- disable: T10 PI mode is disable

When you change the T10 PI mode, the T10 PI mode of the ports which share the settings with the port are also changed. The user who executes the command must have authority of the specified port and the ports which share the settings with the port. Details about the ports which share the settings with the specified port, see the *Provisioning Guide* for the storage system.

-iscsi_virtual_port_mode {enable|disable}

Specifies the iSCSI virtual port mode as follows:

- enable: iSCSI virtual port mode is enabled
- disable: iSCSI virtual port mode is disabled

The setting of the iSCSI virtual port mode is shared among all ports on the channel adapter (CHA) or channel board (CHB) to which ports specified for -port <port> belong.

When you change the iSCSI virtual port mode, the iSCSI virtual port mode of the ports that share the settings with the port is also changed. The user who executes the command must have authority of the specified port and the ports that share the settings with the port. For details about the ports that share the settings with the specified port, see the *Provisioning Guide* for the storage system.

If you specify this option when the device does not support the iSCSI virtual port, an error occurs.

For this command, you can specify "e" as an alias of "enable" or "d" as an alias of "disable".

-port_mode {scsi | nvme} -request_id auto

Changes the operation mode of the port.

- scsi: SCSI mode
- nvme: NVMe mode

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

-delete_login_host_nqn

Deletes the host NQN login information.

However, the login information about the host NQN of a host logged in to an NVM subsystem port or a discovery controller (which is an NVMe controller for discovering an NVM subsystem and whose NQN is fixed as nqn.2014-08.org.nvmexpress.discovery) on the specified port is not deleted.

To check the login status, verify if the connection between the host and the NVM subsystem or the NVMe controller is established by using the command that the host provides.

For details about how to verify the connection between the host and the NVM subsystem or the NVMe controller, see the procedure for verifying the subsystem NQN for the NVMe controller to which the host recognizes the connection in the *Provisioning Guide for Open Systems* or the *Provisioning Guide*.

[-mtu <value>]

Specifies the MTU value (1500/4500/9000) used during iSCSI communication. When you omit the specification, the MTU value is not changed.

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-vlan_tagging_mode {enable|disable}]

Specifies the availability of Tag VLAN. When you omit the specification, the settings for the availability of Tag VLAN is not changed.

- enable: enable Tag VLAN
- disable: disable Tag VLAN

For this command, you can specify "e" as an alias of "enable" or "d" as an alias of "disable".

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-add_vlan_id <value>]

Specifies the adding VLAN ID (1-4094). When you omit the specification, the VLAN ID is not added.

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

This value can be set even if Tag VLAN is disabled. The setting is used when Tag VLAN is enabled.

If you specify a virtual port, Tag VLAN is enabled.

[-delete_vlan_id <value>]

Specifies the deleting VLAN ID (1-4094). Also, Tag VLAN is disabled. When you omit the specification, the VLAN ID is not deleted.



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

The VLAN ID can be deleted even if Tag VLAN is disabled. If the VLAN ID is deleted, to enable Tag VLAN, make sure to add VLAN ID when Tag VLAN is enabled.

[-modify_vlan_id <new value>]

Specifies the VLAN ID (1 to 4094). The VLAN ID is changed to the specified ID. When you omit this specification, the VLAN ID is not changed.



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

This value can be set even if Tag VLAN is disabled. The setting is used when Tag VLAN is enabled.

If you specify this option for a port or virtual port for which VLAN ID is not set, an error occurs.

[-ipv4_address <address>]

Specifies the IPv4 address. When you omit the specification, the IPv4 address is not changed.

This option is required if you specify both -add_iscsi_virtual_port and -ipv6_mode disable.

You cannot specify following IPv4 addresses:

- Network address (For example 192.168.10.0 or 0.120.10.1)
- Broadcast address (For example 255.255.255.255 or 10.1.255.255)
- Loopback address (For example 127.0.0.1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-ipv4_subnetmask <subnet mask>]

Specifies the IPv4 subnet mask. When you omit the specification, the IPv4 subnet mask is not changed.

This option is required if you specify both -add_iscsi_virtual_port and -ipv6_mode disable.



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-ipv4_gateway_address <address>]

Specifies the IPv4 default gateway address. When you omit the specification, the IPv4 default gateway address is not changed.

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-ipv6_mode {enable|disable}]

Specifies the availability of IPv6 mode. When you omit the specification, the settings for the availability of IPv6 mode is not changed.

- enable: enable IPv6 mode
- disable: disable IPv6 mode

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

This option is required if you specify `-add_iscsi_virtual_port`. The operation modes for iSCSI virtual ports are as follows:

- enable: IPv6 address is enabled.
- disable: IPv4 address is enabled. In this case, specify both `-ipv4_address <address>` and `-ipv4_subnetmask <subnet mask>`.

For this command, you can specify "e" as an alias of "enable" or "d" as an alias of "disable".

[-ipv6_local_address {auto|<address>}]

Specifies the IPv6 link local address. When you omit the specification, the IPv6 link local address is not changed.

When you specify "auto", the address is set automatically. When you set the address manually, enter the IPv6 address for `<address>`.

You cannot specify following IPv6 addresses:

- Not set (For example ::)
- Multicast address (For example ff00:1024:1215::01)
- Loopback address (For example ::1)



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

This setting can be configured even if IPv6 mode is disabled. The configured value is used when IPv6 mode is enabled.

[-ipv6_global_address {auto|<address>}]

Specifies the IPv6 global address. When you omit the specification, the IPv6 global address is not changed.

When you specify "auto", the address is set automatically. When you set the address manually, enter the IPv6 address for <address>.

You cannot specify following IPv6 addresses:

- Multicast address (For example ff00:1024:1215::01)
- Loopback address (For example ::1)

If you use Router Advertisement (RA) to automatically set the default gateway of the 25-Gbps iSCSI channel board, omit specifying this option.



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

This setting can be configured even if IPv6 mode is disabled. The configured value is used when IPv6 mode is enabled.

[-ipv6_gateway_address <address>]

Specifies the IPv6 gateway address. When you omit the specification, the IPv6 gateway address is not changed.

You cannot specify following IPv6 addresses:

- Multicast address (For example ff00:1024:1215::01)
- Loopback address (For example ::1)

If you use Router Advertisement (RA) to automatically set the default gateway of the 25-Gbps iSCSI channel board, omit specifying this option.

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

This setting can be configured even if IPv6 mode is disabled. The configured value is used when IPv6 mode is enabled.

[-tcp_port <value>]

Specifies the TCP port number (1 - 65535) during iSCSI communication. When you omit the specification, the TCP port number is not changed.

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-tcp_port_io_controller <value>]

Specifies the TCP port number (1 through 65535) for the storage access service during NVMe/TCP communications. When you omit the specification, the TCP port number is not changed. For more information about a TCP port number that cannot be specified, see the *Provisioning Guide* for your storage system.

[-tcp_port_discovery_controller <value>]

Specifies the TCP port number (1 through 65535) for the discovery service during NVMe/TCP communications. When you omit the specification, the TCP port number is

not changed. For more information about a TCP port number that cannot be specified, see the *Provisioning Guide* for your storage system.

[-selective_ack_mode {enable|disable}]

Specifies the availability of selective ACK. When you omit the specification, the settings of the selective ACK is not changed.

- enable: enable selective ACK
- disable: disable selective ACK



Note:

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

If the type of the channel board or channel adapter is 25iSCSI2o(CHB) (25G iSCSI(Optic)), Ethernet 25Gbps Channel Board, or Ethernet 100Gbps Channel Board, you cannot disable this option.

To check the type of your channel board or channel adapter, see the **Port Condition** window of Device Manager - Storage Navigator or the **Controller Chassis** window of the maintenance utility.

For this command, you can specify "e" as an alias of "enable" or "d" as an alias of "disable".

[-delayed_ack_mode {enable|disable}]

Specifies the availability of delayed ACK. When you omit the specification, the settings of the delayed ACK is not changed.

- enable: enable delayed ACK
- disable: disable delayed ACK

For this command, you can specify "e" as an alias of "enable" or "d" as an alias of "disable".

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-window_size <size>]

Specifies the size of window. You can specify 64KB, 128KB, 256KB, 512KB, 1024KB, or 2048KB. The specifiable unit is m or M for mega byte, and k or K for kilo byte. When you omit the unit, Block (512 byte) is used. For example:

- When you specify 1,024KB:

```
-window_size 1M, -window_size 1m, -window_size 1024K, -
window_size 1024k, or -window_size 2048
```

- When you specify 256KB:

```
-window_size 256K, -window_size 256k, -window_size 512
```

**Note:**

If this option is specified, the communications through the port are interrupted.

When you set the attribute value related to iSCSI, the communications through the specified port are interrupted, and an I/O error is reported to the SCSI command which is being executed. Therefore, if the attribute value related to iSCSI is set more than once continuously, troubles on I/O, for example the decreasing I/O throughput, and the failure in reading and writing, might occur. We recommend that you set the attribute value for iSCSI only when the I/O from the hosts does not exist.

[-keep_alive_timer <value>]

Specifies the value of the Keep Alive Timer (30 - 64800 seconds) during iSCSI communication.

When you omit the specification, the value of the Keep Alive Timer is not changed.

[-ipv6_global_address2 <address>]

Specifies the IPv6 global address2. You can specify this option for virtual ports for which the physical port ID or iSCSI virtual port ID is 0. If a value from 1 to 15 is specified as the iSCSI virtual port ID, an error occurs. When you omit this option, the IPv6 global address2 is not changed.

If you specify the IPv6 address automatically, specify auto for -ipv6_global_address. An EX_INVARG error occurs if you specify auto again when it is already specified.

If you want to specify IPv6 global address2 manually, you must also specify the IPv6 global address manually. When the IPv6 global address is already set manually, you do not need to change it. Type the IPv6 address for <address>.

The following IPv6 addresses cannot be set:

- Multicast Address (Example: ff00:1024:1215::01)
- Loopback Address (Example: ::1)

If you set this option when IPv6 mode is disabled, the setting is ignored.

If IPv6 has changed from enabled to disabled, the setting of this option is initialized.

[-isns_mode {enable|disable}]

Specifies the availability of iSNS service.

- enable: enable iSNS service
- disable: disable iSNS service

When you omit the specification, the settings of the iSNS service is not changed.

For this command, you can specify "e" as an alias of "enable" or "d" as an alias of "disable".

[-isns_server_address <IPv4 address|IPv6 address>]

Specifies the IP address of the iSNS server. You can specify both IPv4 and IPv6 address.

You cannot specify following IPv4 addresses:

- Broadcast address (For example 255.255.255.255 or 10.1.255.255)
- Loopback address (For example 127.0.0.1)

You cannot specify following IPv6 addresses:

- Not set (For example ::)
- Multicast address (For example ff00:1024:1215::01)
- Loopback address (For example ::1)

When you omit the specification, the IP address of the iSNS server is not changed.

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-isns_port <iSNS TCP Port number>]

Specifies the value of the TCP port in the iSNS server (1 - 65535).

When you omit the specification, the value of the TCP port in the iSNS server is not changed.

-add_iscsi_virtual_port <value>

Adds the virtual port to the port for which iSCSI virtual port mode is enabled. Specify the iSCSI virtual port ID (1 to 15) for <value>. Either IPv4 or IPv6 must be set. When you specify IPv4, select disable for -ipv6_mode, and specify both -ipv4_address <address> and -ipv4_subnetmask <subnet mask>. For IPv6, select enable for -ipv6_mode.

An error occurs if you specify this option in any of the following cases:

- The specified iSCSI virtual port is already registered.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

-modify_iscsi_virtual_port <value>

Edits virtual port information of the port for which iSCSI virtual port mode is enabled.

Specify the iSCSI virtual port ID (0 to 15) for <value>.

When you change settings from IPv6 to IPv4 for iSCSI virtual ports whose IDs are from 1 to 15, select disable for -ipv6_mode, then set both -ipv4_address <address> and -ipv4_subnetmask <subnet mask>.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

-delete_iscsi_virtual_port <value>

Deletes the virtual port from the port for which iSCSI virtual port mode is enabled.

Specify the iSCSI virtual port ID (1 to 15) for <value>.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

-add_virtual_port <value>

Adds the virtual port to the port whose virtual port mode is enabled. Specify the virtual port ID (1 through 15) for <value>. Set either of IPv4 or IPv6.

To set IPv4, set disable for -ipv6_mode, and then set both of -ipv4_address <address> and -ipv4_subnetmask <subnet mask>. To set IPv6, set enable for -ipv6_mode. An error occurs if you specify this option when any of the following conditions are met:

- The specified virtual port is already registered.
- The virtual port mode is disabled.
- The storage system does not support the virtual port mode.

-modify_virtual_port <value>

Modifies the virtual port information about the port whose virtual port mode is enabled. Specify the virtual port ID (0 through 15) for <value>. An error occurs if you specify this option when any of the following conditions are met:

- The virtual port with the specified virtual port ID does not exist.
- The virtual port mode is disabled.
- The storage system does not support the virtual port mode.

To change the settings from IPv6 to IPv4 for the virtual ports whose IDs are 0 through 15, set `disable` for `-ipv6_mode`, and then set both of `-ipv4_address <address>` and `-ipv4_subnetmask <subnet mask>`.

-delete_iscsi_virtual_port <value>

Deletes the virtual port from the ports whose virtual port mode is enabled. Specify the virtual port ID (1 through 15) for `<value>`. An error occurs if you specify this option when any of the following conditions are met:

- The specified virtual port does not exist.
- The virtual port mode is disabled.
- The storage system does not support the virtual port mode.

-endpoint_security{enforce|enable|disable}[-forcible] -request_id auto

Specifies the endpoint security encryption setting for mainframe ports.

- `enforce`: Enables endpoint security encryption. The connection cannot be established if endpoint security encryption is not supported on the host.
- `enable`: Enables endpoint security encryption. The connection can be established even if endpoint security encryption is not supported on the host.
- `disable`: Disables endpoint security encryption.

Specify `"-forcible"` (which enables forced configuration mode) if you want to permit disconnection of the host when the endpoint security encryption setting is changed.

`<request#>` (request ID) is assigned each time the command is run with the `-request_id auto` option. For more information, see [Request ID function \(on page 293\)](#).

Returned values

One of the values shown below is returned to `exit()` so that you can check the execution results using a user program or a script.

- **0**: Normal termination
- **EX_ENOSUP**: The host does not support IPv6. For details, see the *Command Control Interface Installation and Configuration Guide*.
- **Other than 0 and EX_ENOSUP**: Abnormal termination. For details, see the *Command Control Interface User and Reference Guide*.

Examples

Changing the port attributes (the Loop ID and the topology of the port) of a port CL3-E.

```
# raidcom modify port -port CL3-E -loop_id 0xAB -topology fl_port
```

Changing the port attributes of the port CL3-E to the External initiator port (ELUN).

```
# raidcom modify port -port CL3-E -port_attribute ELUN
```

When you enable the IPv6 of the port CL3-E, and automate the settings of the Global address.

```
# raidcom modify port -port CL3-E -ipv6_mode enable
-ipv6_global_address auto
```

When you enable the iSNS service of the port CL3-E, and set the IP address to the iSNS server:

```
# raidcom modify port -port CL3-E -isns_mode enable
-isns_server_address fe80::209:6bff:febe:3c17
```

When you enable the T10 PI mode of the port CL3-E:

```
# raidcom modify port -port CL3-E -t10pi enable
```

When you enable the virtual port mode of the port CL3-E:

```
# raidcom modify port -port CL3-E -iscsi_virtual_port_mode enable
```

When you add virtual port 1 to the port CL3-E using IPv4 to set IP address:10.213.46.63, subnet mask: 255.255.254.0:

```
#raidcom modify port -port CL3-E -add_iscsi_virtual_port 1 -ipv6_mode disable -
ipv4_address 10.213.46.63 -ipv4_subnetmask 255.255.254.0
```

When you change the operation mode of the port CL3-E to NVMe mode:

```
# raidcom modify port -port CL3-E -port_mode nvme -request_id auto
```

When you delete the login information about the host NQN that is logged in to the port CL3-E:

```
# raidcom modify port -port CL3-E -delete_login_host_nqn
```

When you change the endpoint security encryption setting for the port CL3-E to “enable” (the connection can be established even if endpoint security encryption is not supported on the host):

```
# raidcom modify port -port CL3-E -endpoint_security enable -request_id auto
```

raidcom add parity_grp

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Creates parity groups. For VSP One B20, traditional parity groups cannot be created . Only DDP groups can be created.

This command is executed asynchronously with the command input. Use the **raidcom get command_status** command to check if the command is completed.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If a drive, cache memory, HIE, or a logical path inside the storage is blocked, the operation cannot be performed.

Syntax for a parity group whose attribute is not DDP

```
raidcom add parity_grp {-parity_grp_id <gno-sgno> |
-concatenated_parity_grp_id <gno-sgno>... } -drive_location <drive location>...
-raid_type <raid type> [-encryption {enable|disable}]
[-copy_back {enable|disable}] [-accelerated_compression {enable|disable}]
[-clpr <clpr#>]
```

Syntax for a DDP parity group

```
raidcom add parity_grp -parity_grp_id <gno-sgno> -drive_location <drive location>... -
raid_type <raid type> -ddp [-encryption {enable|disable}] [-clpr <clpr#>] -request_id
auto
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno: 1-52, sgno: 1-32).

Example:

3-1

-concatenated_parity_grp_id <gno-sgno>...

Specifies the number of the parity group to be concatenated if a distributed parity group is created.

Drives specified by **-drive_location <drive location>...** are used in the order they were specified to create parity groups to be concatenated.

Parity group numbers specified by `-concatenated_parity_grp_id <gno-sgno>...` are assigned in the order they were specified to parity groups to be created.

-drive_location <drive location>...

Specifies the drive location.

Example:

When specifying drive location HDD00-01 (Drive00-01): 0-1

-raid_type <raid type>

Specifies the RAID type. The following types can be specified:

- For RAID type 2D+2D: 2D2D
- For RAID type 3D+1P: 3D1P
- For RAID type 4D+1P: 4D1P
- For RAID type 6D+1P: 6D1P
- For RAID type 7D+1P: 7D1P
- For RAID type 6D+2P: 6D2P
- For RAID type 12D+2P: 12D2P
- For RAID type 14D+2P: 14D2P

[-encryption {enable|disable}]

Enables or disables encryption. If you omit this option, disable is set.

- enable: Encryption is enabled.
- disable: Encryption is disabled.

[-copy_back {enable|disable}]

Enables or disables the copy back mode. If you omit this option, enable is set.

- enable: Copy back is enabled.
- disable: Copy back is disabled.

For VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, specifying `-copy_back disable` is recommended.

The copy back mode is always enabled for DDP groups. Specifying this option is not required.

[-accelerated_compression {enable|disable}]

Enables or disables accelerated compression of parity groups. If you omit this option, disable is set.

- enable or yes: accelerated compression is enabled.
- disable or no: accelerated compression is disabled.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

[-clpr <clpr#>]

Specifies the CLPR number. If you omit this option, 0 is set.

Example:

-clpr 2

-ddp -request_id auto

This option is used to create a DDP parity group.

-request_id is a request ID assigned each time the command is run. For details, see [Request ID function \(on page 293\)](#).

Examples

Creating parity group 1-1 with RAID type 3D+1P using drives of which drive locations are HDD00-00 through HDD00-03.

```
# raidcom add parity_grp -parity_grp_id 1-1 -drive_location 0-0 0-1 0-2 0-3 -
raid_type 3D1P
```

Creating parity groups 1-1, 1-2, 1-3, 1-4 with RAID type 7D+1P using drives of which drive locations are from HDD00-00 to HDD00-31, and configuring distributed parity groups. This command uses HDD00-00 to HDD00-07 to create parity group 1-1, HDD00-08 to HDD00-15 to create parity group 1-2, HDD00-16 to HDD00-23 to create parity group 1-3, and HDD00-24 to HDD00-31 to create parity group 1-4.

```
# raidcom add parity_grp -concatenated_parity_grp_id 1-1 1-2 1-3 1-4
-drive_location 0-0 0-1 0-2 0-3 0-4 0-5 0-6 0-7 0-8 0-9 0-10 0-11 0-12 0-13 0-14
0-15 0-16 0-17 0-18 0-19 0-20 0-21 0-22 0-23 0-24 0-25 0-26 0-27 0-28 0-29 0-30
0-31 -raid_type 7D1P
```

Creating a DDP parity group 1-1 with RAID type 6D+2P using drives of which drive locations are Drive 00-00 through Drive 00-08.

```
# raidcom add parity_grp -parity_grp_id 1-1 -drive_location 0-0 0-1 0-2 0-3 0-4 0-5 0-
6 0-7 0-8 -raid_type 6D2P -ddp -request_id auto
```

raidcom delete parity_grp

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Deletes parity groups.

This command is executed asynchronously with the command input. Use the **raidcom get command_status** command to check if the command is completed.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If a drive, cache memory, HIE, or a logical path inside the storage is blocked, the operation cannot be performed.



Note: Do not delete a parity group that is being moved to a CLPR. If you delete the parity group being moved to a CLPR, the deletion of the source and destination CLPRs might fail with the error of SSB=2E11, 7004. If the error, SSB=2E11, 7004, occurs, perform the following steps.

1. Use the **raidcom get parity_grp** command to make sure there are no parity groups being moved to CLPRs.

The M column of the command output shows information about whether or not the parity group is being moved to CLPR.

2. Wait for a while (about 20 minutes) and then execute the **raidcom delete clpr** command.

If you get the same error again, repeat step 2.

Syntax

```
raidcom delete parity_grp -parity_grp_id <gno-sgno>
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno: 1 to 52, sgno: 1 to 32). If the specified parity group is concatenated with other parity groups, all the concatenated parity groups are deleted.

Example:

3-1

Examples

Deleting parity group 1-1.

```
# raidcom delete parity_grp -parity_grp_id 1-1
```

raidcom get parity_grp

Displays parity group information.

Syntax

```
raidcom get parity_grp [-parity_grp_id <gno-sgno> | -key <keyword> ] [-pcap]
```

Options and parameters

[-parity_grp_id <gno-sgno>]

Specifies the parity group number (gno: 1 to 52, sgno: 1 to 32).

If this option is specified, the LDEV and free space information defined in the specified parity group is displayed. For example:

- 3-1



Note: If this option is omitted, the list of parity groups defined in the storage system is displayed.

[-key <keyword>]

Specifies a display keyword. Specify opt for <keyword>. When this option is specified, the following information is displayed:

- Total capacity of the parity group
- Information whether accelerated compression setting for the parity group is enabled or disabled
- Information whether copy back mode is enabled or disabled
- Information whether encryption is enabled or disabled
- When a parity group is contained in a distributed parity group, the number of the parity group that has the smallest number and the operation status of the capacity saving feature.



Note: If this option is omitted, the list of parity groups defined in the storage system is displayed.

[-pcap]

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Displays the actual capacity. This option is valid only when the `-key <keyword>` option is not specified, or when the `-key opt` option is specified. For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, specify this option to check the capacity information.

If the microcode version does not support this option, nothing is displayed when this option is specified.

Example 1

Displaying parity group information.

```
# raidcom get parity_grp
```

```
T GROUP  Num_LDEV  U(%)  AV_CAP(GB)  R_LVL  R_TYPE  SL  CL  DRIVE_TYPE
M E_TYPE
R 5-2      4    45    140000  RAID1  2D+2D  0  0  DKS2C-K072FC
Y OPEN-V
R 5-3      4    45    140000  RAID1  2D+2D  0  0  DKS2C-K072FC
N OPEN-V
```

```
# raidcom get parity_grp -key opt
T GROUP TOTAL_CAP(GB) V C E C_GROUP PWSV_S RGID NUM_SPARE NUM_REMAIN_SPARE
PROCESS_STATUS PROGRESS(%)

R 1-1    300 D E E 1-1 - 0 - - - -
R 1-2    300 D E E 1-1 - 0 - - - -
R 1-3    300 D E E 1-1 - 1 - - - -
R 1-4    300 D E E 1-1 - 2 - - - -
R 2-16   500 E D D - - 3 - - - -
R 3-1    300 D E E - PS 1 - - - -
```

Description of each column in output example:**T**

Displays the type of the volume group.

- R: parity group
- P: DDP

GROUP

Displays the parity group number.

Num_LDEV

Displays the number of LDEVs assigned to this parity group.

U(%)

Displays the usage rate of this parity group. Displays the actual usage rate of the parity group when the `-pcap` option is specified.

AV_CAP(GB)

Displays the available capacity (free space) for this parity group. A value less than 1 GB is rounded down and 0 is displayed. When the `-pcap` option is specified, displays the actual capacity that can be used for the parity group (free area).

R_LVL

Displays the RAID level of the parity group.

R_TYPE

Displays the RAID type of the parity group.

SL

Displays the SLPR number to which the parity group belongs.

CL

Displays the CLPR number to which the parity group belongs.

DRIVE_TYPE

Drive type code which is set when the parity group is set. This drive type code is a typical code of a drive that is compatible with drives contained in the parity group, and therefore might not match the actual code of the drive contained in the parity group.

To view the drive type code of the drive in the parity group, execute the **raidcom get drive** command.

M

Displays whether the allocation of parity group to CLPR is changed.

- Y: The allocation of parity group is being changed.
- N: The allocation of parity group is not changed.
- - (hyphen): a hyphen (-) is always displayed for VSP G200, G400, G600, G800, and VSP F400, F600, F800.

E_TYPE

Displays the base emulation type of the parity group.

TOTAL_CAP(GB)

Displays the total capacity of the parity group.

For VSP 5000 series, VSP G1x00, and VSP F1500, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, displays the actual usage rate of the parity group when the `-pcap` option is specified.

V

Displays whether accelerated compression of the parity group is enabled or disabled.

- E: Accelerated compression is enabled.
- D: Accelerated compression is disabled.

C

Displays whether the copy back mode is enabled or disabled.

- E: Copy back mode is enabled.
- D: Copy back mode is disabled.
- - (hyphen): Displaying copy back mode is not supported.

E

Displays whether encryption is enabled or disabled.

- E: Encryption is enabled.
- D: Encryption is disabled.
- - (hyphen): Displaying the encryption status is not supported.

C_GROUP

When a parity group is contained in a distributed parity group, displays the smallest parity group number in the distributed parity group. If parity groups are not contained in distributed parity groups, or the function for displaying the parity group numbers is not supported, a hyphen (-) is displayed.

RGID

Displays the resource group ID. A hyphen (-) is displayed if the DKCMAIN microcode version is not supported.

NUM_SPARE

Displays the number of times that a backup is available in the spare area.

A hyphen (-) is displayed if this information is not valid in this parity group.

NUM_REMAIN_SPARE

Displays the number of times that a backup is available in the remaining spare area.

A hyphen (-) is displayed if this information is not valid in this parity group.

PROCESS_STATUS

Processing status of the parity group

- NONE: Processing of expanding the parity group is not running.
- EXTENDING: The parity group is being expanded.
- SUSPEND_EXTEND: Expansion of the parity group is stopped.
- - (hyphen): The information is not available for this parity group.

PROGRESS(%)

Percentage of the progress of the processing status of the parity group

- (hyphen): The information is not available for this parity group. A hyphen (-) is displayed if PROCESS_STATUS is NONE.

Example 2

Displaying information on LDEVs and the free space defined in the parity group.

```
# raidcom get parity_grp -parity_grp_id 5-2
```

T	GROUP	P_NO	LDEV#	STS	LOC_LBA	SIZE_LBA	Serial#	SP
R	5-2	0	-	NML	0x00000000000000	0x0000000003f00	64034	-
R	5-2	1	100	NML	0x0000000003f00	0x0000000010000	64034	R

```
R 5-2      2    101 REG  0x0000000013f00  0x0000000010000  64034  V
R 5-2      3      - DEL  0x0000000023f00  0x0000f00000000  64034  -
```

Example 3

Displaying information on the parity group 33-16 whose attribute is not DDP.

```
# raidcom get parity_grp -parity_grp_id 33-16
T GROUP P_NO LDEV# STS LOC_LBA SIZE_LBA Serial# SP
R 33-16 0 - NML 0x0000000000000 0x0000000003f00 - 64034 -
R 33-16 1 100 NML 0x0000000003f00 0x0000000010000 64034 R
R 33-16 2 101 REG 0x0000000013f00 0x0000000010000 64034 V
R 33-16 3 - DEL 0x0000000023f00 0x0000f00000000 - 64034 -
```

Example 4

Displaying information on the DDP parity group 1-2 that contains no LDEVs.

```
# raidcom get parity_grp -parity_grp_id 1-2
T GROUP P_NO LDEV# STS LOC_LBA SIZE_LBA Serial# SP
P 1-2 0 - NML - 0x0000000000000 800089 -
```

Example 5

Displaying information on the DDP parity group 1-2 that contains LDEVs.

```
# raidcom get parity_grp -parity_grp_id 1-2
c:\HORCM\etc>raidcom get parity_grp -parity_grp_id 1-2 -I1
T GROUP P_NO LDEV# STS LOC_LBA SIZE_LBA Serial# SP
P 1-2 0 226 NML - 0x0000a4e20000 800089 R
P 1-2 1 227 NML - 0x0000a4e20000 800089 R
P 1-2 2 228 NML - 0x0000a4e20000 800089 R
P 1-2 3 229 NML - 0x0000a4e20000 800089 R
P 1-2 4 230 NML - 0x0000a4e20000 800089 R
P 1-2 5 231 NML - 0x0000a4e20000 800089 R
P 1-2 6 232 NML - 0x0000a4e20000 800089 R
P 1-2 7 233 NML - 0x0000a4e20000 800089 R
P 1-2 8 234 NML - 0x0000a4e20000 800089 R
P 1-2 9 235 NML - 0x0000a4e20000 800089 R
P 1-2 10 - NML - 0x0000000000000 800089 -
```

Description of each column in output example:

T

Displays the type of the volume group.

- R: parity group
- P: DDP

GROUP

Displays the parity group number.

P_NO

Displays the partition number partitioning this parity group.

LDEV#

Displays LDEV number.

A hyphen (-) is displayed if no LDEVs are defined for the partition.

For a DDP parity group, a partition whose SIZE_LBA is 0x000000000000 with no LDEVs defined is displayed at the end. For details, see the display example above.

STS

Displays the following status.

- NML: LDEV is installed or free space is settled.
- REG: LDEV is being created.
- DEL: LDEV is being deleted.

LOC_LBA

Displays the Start of LBA for this partition on this parity group, in blocks (512 bytes).

A hyphen (-) is displayed if the parity group is a DDP parity group.

SIZE_LBA

Displays the size for this partition on this parity group, in blocks (512 bytes).

Serial#

Product serial number.

SP

Displays whether the LDEV uses the expanded space of the parity group.

- V: LDEV uses the expanded space.
- R: LDEV does not use the expanded space.
- -: LDEVs are not mounted.

raidcom initialize parity_grp

Supported storage systems:

- VSP 5000 series
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

This command formats all areas in the drives of the specified parity group (the parity-group formatting operation).

This command is executed asynchronously with the command input. Check the completion of this process using the `raidcom get command_status` command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Caution: This is a destructive operation. Verify the correct parity group ID before issuing this command. The user is responsible for backing up the data in the parity group, if necessary, before issuing this command.

Syntax

```
raidcom initialize parity_grp -parity_grp_id <gno-sgno> -operation <type>
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group ID of the drives to be formatted (gno:1-52, sgno:1-32).

Example:

- 3-1

You can see the progress of the format by OPE_RATE of the `raidcom get ldev` command.

-operation <type>

Specifies "fmt" in <type> to format all areas in the drives of the specified parity group.

Examples

Formats all areas in the drives of parity group: 1-1:

```
# raidcom initialize parity_grp -parity_grp_id 1-1 -operation fmt
```

raidcom modify parity_grp

Enables or disables accelerated compression of a parity group. This command also stops expansion of a parity group (stops DDP expansion) or resumes expansion of a parity group (resumes DDP expansion).

This command is executed asynchronously with the command input. Check the completion of this process using the `raidcom get command_status` command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify parity_grp -parity_grp_id <gno-sgno>
    { -accelerated_compression {enable|disable} | -copy_back {enable|disable} | -
operation <type> -request_id auto}
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno:1-52, sgno:1-32) (for example, 3-1).

-accelerated_compression <enable|disable>

Enables or disables accelerated compression of the specified parity group.

For VSP 5000 series, an error occurs if enable is specified for the parity group that belongs to the pool whose subscription limit is other than 65535 (unlimited).

- enable or yes: Enables accelerated compression.
- disable or no: Disables accelerated compression.

For this command, you can specify "y" as an alias of "yes" or "n" as an alias of "no".

-copy_back {enable|disable}

Enables or disables the copy back mode. This option is specified only for VSP One B20.

- enable: Enables copy back.
- disable: Disables copy back.

The copy back mode is always enabled and cannot be changed for a DDP parity group.

-operation <type> -request_id auto

Supported storage systems:

- VSP One B20

Specifies the operation. You can specify the following operations:

- suspend_extend: Stops expanding the parity group.
- resume_extend: Resumes expanding the parity group.

-request_id is a request ID assigned for each time the command is run. For details, see [Request ID function \(on page 293\)](#).

Examples

Enables accelerated compression of parity group: 1-1:

```
# raidcom modify parity_grp -parity_grp_id 1-1 -accelerated_compression enable
```

Enables copy back of parity group: 1-1:

```
# raidcom modify parity_grp -parity_grp_id 1-1 -copy_back enable
```

Stops expansion of parity group 1-1:

```
# raidcom modify parity_grp -parity_grp_id 1-1 -operation suspend_extend
-request_id auto
```

Resumes expansion of parity group 1-1:

```
# raidcom modify parity_grp -parity_grp_id 1-1 -operation resume_extend -
request_id auto
```

raidcom extend parity_grp

Supported storage systems:

- VSP One B20

Expands the data capacity of a parity group (pool capacity expansion per drive). When expanding the pool capacity per drive, an LDEV is automatically created in the specified parity group, and the LDEV is automatically added to the pool to which the parity group belongs.

This command is run asynchronously with the command input. After running the **raidcom get command_status** command to confirm that the processing has started, confirm that **PROCESS_STATUS** in the result of the **raidcom get parity_grp** command has changed from **EXTENDING** to **NONE**.

The number of asynchronous commands that can be accepted at the same time is limited. For details about the maximum number of acceptable asynchronous commands, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom extend parity_grp -parity_grp_id <gno-sgno> -drive_location
<drive location>... [-ldev_range <range>] -request_id auto
```

Options and parameters

-parity_grp_id <gno-sgno>

Specifies the parity group number (for example, 3-1).

-drive_location <drive location>...

Specifies the location of the drive to be added to the parity group (for example, when specifying drive location Drive00-01: 0-1).

[-ldev_range <range>]

Specifies the range of the LDEV numbers of LDEVs that are automatically created.

When the `-ldev_range <range>` option is specified, among the LDEV numbers of unimplemented LDEVs that are specified in `<range>`, this option assigns the largest LDEV number in the resource group that contains the pool volumes in the pool to which the specified parity group belongs. Specify the LDEV number in decimal or hexadecimal. When specifying the LDEV number in hexadecimal, add "0x" at the beginning of the LDEV number.

For example, when specifying the range of LDEV number from 300 to 305:

```
-ldev_range 300-305
```

```
-ldev_range 0x12c-0x131
```

You can also specify the mixture of decimal and hexadecimal numbers. For example, `-ldev_range 0x12c-305`

When the `-ldev_range <range>` option is not specified, the largest LDEV number in the resource group that contains the pool volumes in the pool to which the specified parity group belongs is assigned.

The command cannot be run if either of the following conditions is met:

1. There are not enough unimplemented LDEVs to create an LDEV in the resource group for which the user has operation permission.
2. When the `-ldev_range <range>` option is specified, in addition to condition 1, there are not enough unimplemented LDEVs to create an LDEV among the LDEV numbers specified in `<range>`.

Before running the command, use the `raidcom get ldev -ldev_list undefined -resource <resource_grp_id>` command to verify whether there are enough unimplemented LDEVs to create the LDEV. If there are not enough unimplemented LDEVs to create the LDEV, use the `raidcom add resource -ldev_id <ldev#>` command or the `raidcom delete resource -ldev_id <ldev#>` command to move the unimplemented LDEV numbers to the target resource group, or use the `raidcom delete ldev` command to delete unnecessary LDEVs.

-request_id auto

`-request_id` is a request ID assigned for each time the command is run. For details, see [Request ID function \(on page 293\)](#).

Examples

Adds the drive whose drive location is Drive02-01 to parity group: 1-2 to expand the data capacity.

```
# raidcom extend parity_grp -parity_grp_id 1-2 -drive_location 2-1 -request_id auto
```

raidcom add rcu

Registers RCUs.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add rcu {-rcu <serial#> <mcu#> <rcu#> <id> -ssid <ssid>
| -cu_free <serial#> <id> <pid> -mcu_port <port#>
-rcu_port <port#>}
```

Options and parameters

-rcu <serial#> <mcu#> <rcu#> <id>

Specifies the CU specified by serial number, <mcu#>, <rcu#>, and <id>. You can use hexadecimal numbers (add "0x" prefix) or decimal numbers for <mcu#> and <rcu#>.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system type as follows:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

-ssid <ssid>

Specifies storage subsystem IDs. Up to 4 SSIDs can be specified.

You can use hexadecimal numbers (add "0x" prefix) or decimal numbers for <ssid>.

Note: This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

-cu_free <serial#><id><pid>

This parameter is used to specify CU free specified by serial number, ID, and PID.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

<pid> specifies the path group ID (1-255).



Note:

The RCU is displayed by default on the GUI of TrueCopy, Universal Replicator, or Universal Replicator for Mainframe when "0" is specified for <pid>.

-mcu_port <port#>

Specifies the port number of the MCU.

-rcu_port <port#>

Specifies the port number on the RCU (storage system port on the remote side).

Specify the port that the attribute is MCU Initiator port (MCU) or RCU Target port (RCU).

Examples

Register VSP of serial number: 64034 with CU free. Sets the path group ID: 0, the port on MCU: CL1-A, and the port on RCU: CL1-B.

```
# raidcom add rcu -cu_free 64034 R800 0 -mcu_port CL1-A -rcu_port CL1-B
```

raidcom delete rcu

Deletes the RCU.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete rcu {-rcu <serial#> <mcu#> <rcu#> -ssid <ssid> |  
-cu_free <serial#> <id> <pid>
```

Options and parameters

-rcu <serial#> <mcu#> <rcu#>

Specifies the CU specified by serial number, <mcu#>, <rcu#>. You can use hexadecimal numbers (add "0x" prefix) or decimal numbers for <mcu#> and <rcu#>.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-ssid <ssid>

Specifies the storage subsystem ID.

You can use hexadecimal numbers (add "0x" prefix) or decimal numbers for <ssid>.

Note: This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

-cu_free <serial#><id><pid>

This parameter is used to specify CU free specified by serial number, ID, and PID.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

<pid> specifies the path group ID (1-255).



Note:

The RCU is displayed by default on the GUI of TrueCopy, Universal Replicator, or Universal Replicator for Mainframe when "0" is specified for <pid>.

Examples

Deleting the VSP for which the product number: 64034, the RAID type: R800 and the path group ID: 1 are set.

```
# raidcom delete rcu -cu_free 64034 R800 1
```

raidcom get rcu

Displays MCU/RCU information.

Syntax

```
raidcom get rcu [-rcu <serial#> <mcu#> <rcu#> [-ssid <ssid>] |
-cu_free <serial#> <id> <pid>]
```

Options and parameters

-rcu <serial#> <mcu#> <rcu#>

Specifies the CU specified by serial number, <mcu#>, <rcu#>. You can use hexadecimal numbers (add "0x" prefix) or decimal numbers for <mcu#> and <rcu#>.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

[-ssid <ssid>]

Specifies the storage subsystem ID.

You can use hexadecimal numbers (add "0x" prefix) or decimal numbers for <ssid>.

Note: This option is for enterprise storage systems. You need not specify the -ssid option when you use Unified Storage VM, but this option is enabled.

-cu_free <serial#><id><pid>

This parameter is used to specify CU free specified by serial number, ID, and PID.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

<pid> specifies the path group ID (1-255).

**Note:**

The RCU is displayed by default on the GUI of TrueCopy, Universal Replicator, or Universal Replicator for Mainframe when "0" is specified for <pid>.

Example 1

Displaying RCU information.

```
# raidcom get rcu
```

Serial#	ID	PID	MCU	RCU	M/R	T	STS	MP	NP	IM	FZ	RTO(s)	RTT(ms)	PBWT	PBW_SIM(s)
64034	R7	-	1C	23	RCU	F	NML	4	8	MR	D	15	20	R	70
64034	R7	-	1C	23	RCU	F	NML	4	8	RO	E	15	20	R	70
64034	R7	1	-	-	MCU	E	UNK	4	8	-	-	15	20	S	70

Description of each column in output example:**Serial#**

Displays the product serial number.

ID

Displays the ID for identifying RAID type:

- R9: VSP 5000 series
- R8: VSP G1x00 and VSP F1500
- M8: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, VSP F400, F600, F800

PID

Displays a path group ID. If CU is specified for RCU, "-" is displayed.

MCU

Displays the CU number (hexadecimal) as MCU. If CU free is specified for RCU, "-" is displayed.

RCU

Displays the CU number (hexadecimal) as RCU. If CU free is specified for RCU, "-" is displayed.

M/R

Displays the CU type as MCU/RCU.

MCU is displayed in the storage system that has TrueCopy or TrueCopy for Mainframe S-VOL after the initial copy was performed.

T

Displays the type of physical path:

- F: Fibre
- E: ESCON
- I: iSCSI
- M: path types are mixed
- - (hyphen): cannot identify the path type

For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, a hyphen (-) is always displayed if M/R is MCU. For other storage systems, a hyphen (-) is displayed if M/R is MCU and MCU might be connected via a path other than the Fibre Channel. To determine if MCU is connected via a path other than Fibre Channel, check the existence of an RCU target port other than FIBRE in the storage system.

STS

Displays the following status of the CU.

- NML: Normal
- WAR: Warning
- ERR: Failing
- UNK: Unknown, displayed when the target of the pair is MCU.

MP

Displays the number of path as minimum.

NP

Displays the number of path setting between MCU and RCU.

IM

Displays the incident mode setting to RCU.

- MR: it sends incident to MCU host and RCU host
- RO: it sends incident only to RCU host
- If CU free is specified for RCU, "-" is displayed.

FZ

Displays the freeze option.

- D: the freeze option is disabled.
- E: the freeze option is enabled.
- -: "-" is displayed when CU free is specified for RCU.

RTO (s)

Displays the timeout value for RIO (Remote IO) setting between MCU and RCU.

RTT (ms)

Displays the round trip time value between MCU and RCU.

PBWT

Displays the type of a value used as the setting time of path blocked watch SIM.

- R: Uses the path blocked watch SIM time that is set on an RCU basis.
- S: Uses the path blocked watch SIM time that is set in the remote replication option on a system basis.
- - (hyphen): This information is not valid.

PBW_SIM(s)

Displays the setting time of path blocked watch SIM in seconds.

- If PWBT is R, displays the path blocked watch SIM time that is set on an RCU basis.
- If PWBT is S, displays the path blocked watch SIM time that is set in the remote replication option.
- - (hyphen): This information is not valid.

Example 2

```
# raidcom get rcu -cu_free 64034 R800 1
```

```
Serial# ID PID MCU RCU M/R T PNO MPORT RPORT STS_CD SSIDs ...
64034 R7 1 - - RCU F 1 CL1-A CL1-B NML_01 -
64034 R7 1 - - RCU F 2 CL1-A CL1-B NML_01 -
64034 R7 1 - - RCU F 3 CL1-A CL1-B NML_01 -
```

Description of each column in output example:**Serial#**

Product serial number.

ID

Displays the ID for identifying RAID type:

- R9: VSP 5000 series
- R8: VSP G1x00 and VSP F1500
- M8: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, VSP F400, F600, F800

PID

Displays the path group ID. If CU is specified for RCU, "-" is displayed.

MCU

Displays the CU number (hexadecimal) as MCU. If CU free is specified for RCU, "-" is displayed.

RCU

Displays the CU number (hexadecimal) as RCU. If CU free is specified for RCU, "-" is displayed.

M/R

Displays the CU type as MCU/RCU.

T

Displays the type of physical path:

- F: Fibre
- E: ESCON
- I: iSCSI
- M: path types are mixed
- - (hyphen): cannot identify the path type

For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, a hyphen (-) is always displayed if M/R is MCU. For other storage systems, a hyphen (-) is displayed if M/R is MCU and MCU might be connected via a path other than the Fibre Channel. To determine if MCU is connected via a path other than Fibre Channel, check the existence of an RCU target port other than FIBRE in the storage system.

PNO

Displays the path number.

MPORT

Displays the MCU port number.

RPORT

Displays the RCU port number.

STS_CD

Displays the following path status:

- NML_01: Normal
- ERR_02: Initialization error
- ERR_03: Communication timeout
- ERR_04: Logical blockade
- ERR_05: Resource shortage
- ERR_06: Serial number mismatch
- ERR_10: Invalid port
- ERR_80: Invalid port number on the remote storage system
- ERR_81: Invalid port attribute on the remote storage system
- ERR_82: Communication error
- If path creation or path deletion is in progress, "-" is displayed.

SSIDs

Displays the SSIDs (hexadecimal) setting to RCU. If CU free is specified for RCU, "-" is displayed.

raidcom modify rcu

This sets the control parameters to specified CU that is specified using two way.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify rcu {-rcu <serial#> <mcu#> <rcu#> -ssid <ssid> |
-cu_free <serial#> <id> <pid>} [-rcu_option <mpth> <rto>
<rtt> [fzd | fze]] [-path_blocked_watch_sim <time(sec)>]
```

Options and parameters**-rcu <serial#> <mcu#> <rcu#>**

Specifies the CU specified by serial number, <mcu#>, and <rcu#>. You can use hexadecimal numbers (add 0x prefix) or decimal numbers for the <mcu#> and <rcu#>.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-ssid <ssid>

Specifies the storage subsystem ID.

You can use hexadecimal numbers (add 0x prefix) or decimal numbers for the <ssid>.

Note: This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

-cu_free <serial#><id><pid>

This parameter is used to specify CU free specified by serial number, ID, and PID.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

<pid> specifies the path group ID (1-255).



Note:

The RCU is displayed by default on the GUI of TrueCopy, Universal Replicator, or Universal Replicator for Mainframe when "0" is specified for <pid>.

[-rcu_option <mpth> <rto> <rtt> [fzd | fze]]

Specifies CU control parameters:

- <mpth>: minimum number of paths (1-8)
- <rto>: RIO timeout value (10-100) (second) for RIO (Remote IO) setting between MCU and RCU.
- <rtt> is used to set the round trip time value (1-500) (millisecond) between MCU and RCU.
- [fzd | fze]: Specify *fze* to enable the freeze option, or *fzd* to disable it.

If the freeze option is not specified to the RCU with CU units, the freeze option is disabled.

[-path_blocked_watch_sim <time(sec)>]

This option can be specified only for VSP One B20.

Specifies the path blocked watch SIM time in seconds.

If 0 is specified, the setting in the remote replication option on a system basis is used.

Examples

For the RCU for which the product number: 64034, the RAID type: R800 and the path ID: 1 are set, setting the options: the minimum number of paths <mpth>4, RIO MIH time<rto>15 seconds, and round trip time 20 milliseconds are set.

```
# raidcom modify rcu -cu_free 64034 R800 1 -rcu_option 4 15 20
```

raidcom add rcu_iscsi_port

This command registers the RCU-side iSCSI port of the remote storage system to the MCU-side iSCSI port of the local storage system. If the iSCSI port of the specified local storage system does not exist, the command is rejected with EX_ENOOBJ.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Note:

- To create an iSCSI path between the local storage system and a remote storage system, you must register the iSCSI target port on the MCU-side iSCSI port of the local storage system first.
- If you register an RCU port connected via iSCSI to the MCU-side iSCSI port, you cannot perform TCz or URz remote copy using iSCSI path connection.

Syntax

```
raidcom add rcu_iscsi_port -port <port#> -rcu_port <port#>
    -rcu_id <serial#> <id> -rcu_address <IP address>
    [-tcp_port <value>][-header_digest {enable|disable}] [-data_digest {enable|
    disable}]
```

Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

-rcu_port <port#>

Specifies the iSCSI port number of the remote storage system.

-rcu_id <serial#> <id>

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system model:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, and VSP F400, F600, F800

-rcu_address <IP address>

Specifies the IP address of the iSCSI target on the remote storage system. An IPv4 address or an IPv6 address can be specified.

Note that the following addresses cannot be specified when you specify the IPv4 address:

- Network address (for example: 192.168.10.0, 0.120.10.1)
- Broadcast address (for example: 255.255.255.255, 10.1.255.255)
- Loop back address (for example: 127.0.0.1)

Note that the following addresses cannot be specified when you specify the IPv6 address:

- Unspecified address (for example: ::)
- Multicast address (for example: ff:1024:1215::01)
- Loop back address (for example: ::1)

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-tcp_port <value>]

Specifies the TCP port number of the iSCSI target on the remote storage system. If this option is omitted, the TCP port number of the iSCSI target port which is specified with -port option is configured.

[-header_digest {enable|disable}]

Specifies whether a header digest is enabled or disabled.

This option is only available if a port on the local storage system is iSCSI 25G CHB. If this option is omitted, the header digest is enabled.

- enable: Enables the header digest.
- disable: Disables the header digest.

[-data_digest {enable|disable}]

Specifies whether a data digest is enabled or disabled.

This option is only available if a port on the local storage system is iSCSI 25G CHB. If this option is omitted, the data digest is enabled.

- enable: Enables the data digest.
- disable: Disables the data digest.

Examples

Registers the iSCSI port: CL1-A (IP address: 158.214.135.100) of the remote storage system (product serial number: 400031, model: VSP G200, G400, G600, G800 and VSP F400, F600, F800) in the iSCSI port: CL4-E of the local storage system:

```
# raidcom add rcu_iscsi_port -port CL4-E -rcu_port CL1-A -rcu_id 400031 M800 -
rcu_address 158.214.135.100
```


Registers iSCSI ports of the remote storage system which is registered in the iSCSI port: CL1-E of the local storage system to the iSCSI port: CL4-E of the local storage system:

```
# raidcom get rcu_iscsi_port | rmawk @1-eq:CL1-E exe="raidcom add rcu_iscsi_port -
port CL4-E -rcu_port @4 -rcu_id @2 @3 -rcu_address @5"
```

raidcom delete rcu_iscsi_port

This command deletes the iSCSI port of the remote storage system that is registered as the RCU side port from MCU-side iSCSI port of the local storage system.

If the specified iSCSI port does not exist in the local storage system, the command is rejected with EX_ENOOBJ. If the specified port of the remote storage system does not exist, the command is ignored.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete rcu_iscsi_port -port <port#> -rcu_port <port#>
-rcu_id <serial#> <id>
```

Options and parameters

-port <port#>

Specifies the port number of the local storage system. For example:

- CL1-A

-rcu_port <port#>

Specifies the iSCSI port number of the remote storage system.

-rcu_id <serial#> <id>

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a "5" at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a "3" at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system model:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, and VSP F400, F600, F800

Examples

Deletes the iSCSI port: CL1-A (IP address: 158.214.135.100) of the remote storage system (product serial number: 400031, model: VSP G200, G400, G600, G800 and VSP F400, F600, F800) from the iSCSI port: CL4-E of the local storage system:

```
# raidcom delete rcu_iscsi_port -port CL4-E -rcu_port CL1-A -rcu_id 400031 M800
```

Deletes the iSCSI port of the remote storage system which is registered in the iSCSI port: CL1-E of the local storage system from the iSCSI port: CL4-E of the local storage system.

```
# raidcom get rcu_iscsi_port | rmawk @1-eq:CL1-E exe="raidcom delete rcu_iscsi_port -port CL4-E -rcu_port @4 -rcu_id @2 @3"
```

raidcom get rcu_iscsi_port

This command displays the RCU-side iSCSI port of the remote storage system which is registered in the MCU-side iSCSI port of the local storage system.

If the specified iSCSI port does not exist in the local storage system, the command is rejected with EX_ENOOBJ.

Only the remote storage port registered in the port to which the user who executes the command can refer are output. For details about the port to which the user can refer, see the descriptions of the relationship of the resource group and the command operation in the *Command Control Interface User and Reference Guide*.

Syntax

```
raidcom get rcu_iscsi_port
```

Options and parameters

None.

Examples

Displays the port of the remote storage system which are registered in the iSCSI port of the local storage system:

```
#raidcom get rcu_iscsi_port
```

PORT	Serial#	ID	RPORT	IP_ADDR	IP_PORT#	HEADER_DIGEST	DATA_DIGEST
CL4-E	400031	M8	CL1-A	158.214.135.100	3260	ENABLED	ENABLED
ENABLED			ENABLED				
CL2-E	400031	M8	CL1-A	158.214.135.100	3260	ENABLED	ENABLED
ENABLED			ENABLED				

CL1-E	400031	M8	CL1-A	158.214.135.100	3260	ENABLED	DISABLED
DISABLED			DISABLED				

Description of each column in output example:**PORT**

Displays the port number.

Serial#

Displays the product serial number of the remote storage system.

ID

Displays the model of the remote storage system.

- R8: VSP G1x00 and VSP F1500
- R9: VSP 5000 series
- M8: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

RPORT

Displays the port number of the remote storage system.

IP_ADDR

Displays the IP address of the remote storage system.

IP_PORT#

Displays the TCP port number of the port of the remote storage system.

HEADER_DIGEST

Specifies whether a header digest is enabled or disabled.

- ENABLED: Enables the header digest setting.
- DISABLED: Disables the header digest setting.
- - (hyphen): This information is not valid.

DATA_DIGEST

Specifies whether a data digest is enabled or disabled.

- ENABLED: Enables the data digest setting.
- DISABLED: Disables the data digest setting.
- - (hyphen): This information is not valid.

HEADER_DIGEST_SESSION

Specifies whether a header digest is enabled or disabled during session establishment.

- ENABLED: Enables the header digest during session establishment.
- DISABLED: Disables the header digest during session establishment.
- NOT_ESTABLISHED: There is no session established.
- - (hyphen): This information is not valid.

DATA_DIGEST_SESSION

Specifies whether a data digest is enabled or disabled during session establishment.

- ENABLED: Enables the data digest during session establishment.
- DISABLED: Disables the data digest during session establishment.
- NOT_ESTABLISHED: There is no session established.
- - (hyphen): This information is not valid.

raidcom add rcu_path

Adds logical paths to RCUs.

This command is executed asynchronously with the command input. Check the completion of this process on the raidcom get command_status command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add rcu_path {-rcu <serial#> <mcu#> <rcu#> -ssid <ssid>
| -cu_free <serial#> <id> <pid>} -mcu_port <port#>
-rcu_port <port#>
```

Options and parameters**-rcu <serial#> <mcu#> <rcu#>**

Specifies the CU specified by serial number, <mcu#>, and <rcu#>. You can use hexadecimal numbers (add 0x prefix) or decimal numbers for the <mcu#> and <rcu#>.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-ssid <ssid>

Specifies the storage subsystem ID.

You can use hexadecimal numbers (add 0x prefix) or decimal numbers for the <ssid>.

Note: This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

-cu_free <serial#><id><pid>

This parameter is used to specify CU free specified by serial number, ID, and PID.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

<pid> specifies the path group ID (1-255).



Note:

The RCU is displayed by default on the GUI of TrueCopy, Universal Replicator, or Universal Replicator for Mainframe when "0" is specified for <pid>.

-mcu_port <port#>

Specifies the port number on the MCU.

-rcu_port <port#>

Specifies the port number on the RCU (storage system port on the remote side).

Specify the port that the attribute is MCU initiator port (MCU) or RCU target port (RCU).

Examples

To add a path to the RCU for which the serial number is 64034, the storage ID is R800 and the path ID is 1 (from MCU port: CL1-A to RCU port: CL1-B).

```
# raidcom add rcu_path -cu_free 64034 R800 1 -mcu_port CL1-A
-rcu_port CL1-B
```

raidcom delete rcu_path

Deletes logical paths from a specified RCU.

This command is executed asynchronously with the command input. Check the completion of this process on the raidcom get command_status command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete rcu_path {-rcu <serial#> <mcu#> <rcu#>
                        -ssid <ssid> | -cu_free <serial#> <id> <pid>} -mcu_port
                        <port#> -rcu_port <port#>}
```

Options and parameters

-rcu <serial#> <mcu#> <rcu#>

Specifies the CU specified by serial number, <mcu#>, and <rcu#>. You can use hexadecimal numbers (add 0x prefix) or decimal numbers for the <mcu#> and <rcu#>.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-ssid <ssid>

Specifies the storage subsystem ID.

You can use hexadecimal numbers (add 0x prefix) or decimal numbers for the <ssid>.

Note: This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

-cu_free <serial#><id><pid>

This parameter is used to specify CU free specified by serial number, ID, and PID.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

<id> specifies the storage system:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One SDS Block, VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800

<pid> specifies the path group ID (1-255).



Note:

The RCU is displayed by default on the GUI of TrueCopy, Universal Replicator, or Universal Replicator for Mainframe when “0” is specified for <pid>.

-mcu_port <port#>

Specifies the port number on the MCU.

-rcu_port <port#>

Specifies the port number on the RCU (storage system port on the remote side).

Specifies the port that the attribute is MCU initiator port (MCU) or RCU target port (RCU).

Examples

From the RCU for which the product number: 64034, the RAID type: R800 and the path ID: 1 are set, deleting RCU path (MCU port: CL1-A and RCU port: CL1-B).

```
# raidcom delete rcu_path -cu_free 64034 R800 1 -mcu_port CL1-A
-rcu_port CL1-B
```

raidcom add ssid

Adds the specified SSID to the RCU that is specified by serial number, <mcu#>, and <rcu#>.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add ssid -rcu <serial#> <mcu#> <rcu#> <id> -ssid <ssid>
```

Options and parameters**-rcu <serial#> <mcu#> <rcu#> <id>**

Specifies the CU specified by serial number, <mcu#>, <rcu#>, and <id>. You can use hexadecimal (add the prefix 0x) or decimal for the <mcu#> and <rcu#> numbers.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

Use the following strings for <id> to specify the RAID storage system type:

- R900: VSP 5000 series
- R800: VSP G1x00 and VSP F1500
- M800: VSP One B20, VSP E series, VSP G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800, and VSP F400, F600, F800

-ssid <ssid>

Specifies the storage subsystem ID (SSID) to add to the RCU.

You can use hexadecimal (add the 0x prefix) or decimal for the <ssid> number.

**Note:**

This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

Examples

Adding SSID:345 to the RCU with serial number: 64034, RAID type: R800, MCU#:0, and RCU#:1 are set.

```
raidcom add ssid -rcu 64034 0 1 R800 -ssid 345
```

raidcom delete ssid

Deletes the specified SSID from the RCU that is specified by a serial number, <mcu#>, and <rcu#>.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete ssid -rcu <serial#> <mcu#> <rcu#> -ssid <ssid>
```

Options and parameters**-rcu <serial#> <mcu#> <rcu#>**

This parameter is used to specify a CU that is specified by a serial number, <mcu#>, and <rcu#>. This option specifies <mcu#> and <rcu#> with hexadecimal numbers (adding 0x) or decimal numbers.

<serial#> specifies the storage system serial number:

- When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
- When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

-ssid <ssid>

Specifies the storage subsystem ID to be deleted from the RCU.

You can specify the ssid as a hexadecimal number (add the 0x prefix) or a decimal number.

Note: This option is for enterprise storage systems. You need not specify the -ssid option for the Unified Storage VM, but this option is enabled.

Examples

Deleting SSID:345 from the RCU where the serial number: 64034, MCU#:0, and RCU#:1 are set.

```
raidcom delete ssid -rcu 64034 0 1 -ssid 345
```

raidcom get ssid

Supported storage systems:

- VSP 5000 series
- VSP G1x00 or VSP F1500

Displays the SSID allocation information for the boundary area of the CU number and LDEV number.

Syntax

```
raidcom get ssid
```

Examples

Displays the SSID allocation information for the boundary area of the CU number and LDEV number.

```
# raidcom get ssid
CU LDEV_B SSID
00 00-FF 0004
01 00-FF 0005
02 00-FF      -
03 00-FF 0012
04 00-FF      -
...
FE 00-FF FEFF
```

Description of each column in output example:

CU

Displays the CU number in hexadecimal notation.

LDEV_B

Displays the boundary area of the LDEV number.

SSID

Displays the SSID in hexadecimal notation.

A hyphen (-) is displayed if the SSID is not specified

raidcom add resource

Creating resource groups. If you specify only the resource group name, an empty resource group is created. If you specify the resource group name and the information on the virtual storage machine, an empty resource group that corresponds to the virtual storage machine is created.

When you input a resource group name and a resource group ID, the current name of the resource group whose ID you specify is changed to the new resource group name.

When you specify resource group name, LDEV number, port number, host group number, parity group ID or an external group ID, the specified resource is registered to the specified resource group. If the specified resource group does not exist, an error occurs. When the resource group is already created, the specified resource is added to the resource group. You can specify a device group name instead of an LDEV number.

When the relevant LDEVs configure the pool, journal, and LUSE, all LDEVs must be added to the same resource group.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

When creating a blank resource group

```
raidcom add resource -resource_name <resource group name>
```

When creating a virtual storage machine and a blank resource group which correspond to the virtual storage machine

```
raidcom add resource -resource_name <resource group name> -virtual_type <serial#> <id>
```

When changing the resource group name of the specified resource group ID

```
raidcom add resource -resource_name <resource group name> -resource_id <resource group_id>
```

When registering a resource (LDEV, port, host group, parity group, external volume group, or NVM subsystem) to the resource group (An error occurs if the specified resource group is not available)

```
raidcom add resource -resource_name <resource group name> [-ldev_id <ldev#> | -port <port#> [  
<host group name>] | -parity_grp_id <gno-sgno>| -external_grp_id <gno-sgno>|  
-nvm_subsystem_id <nvm subsystem id>]
```

When registering an LDEV which belongs to the device group, to the resource group (An error occurs if the specified resource group is not available)

```
raidcom add resource -resource_name <resource group name> -grp_opt ldev -  
device_grp_name <device group name> [<device name>]
```



Note: If the corresponding LDEV is a volume that includes pool, journal, or LUSE, all LDEVs need to be assigned to the same resource group.

Options and parameters

-resource_name <resource group name>

Specifies the resource group name. Up to 32 characters can be specified.

[-virtual_type <serial#> <id>]

Specifies the serial number and the system name (type identifier) of the virtual storage machine.

- serial#: Serial number of the virtual storage machine.
 - When specifying the serial number for VSP 5000 series, add a “5” at the beginning of the serial number. For example, for serial number 12345, enter 512345.
 - When specifying the serial number for VSP G1x00 and VSP F1500, add a “3” at the beginning of the serial number. For example, for serial number 12345, enter 312345.

- id: Type identifier of the virtual storage machine.
 - RH10HF: VSP 5200, VSP 5600
 - RH10HG: VSP 5200H, VSP 5600H
 - R900F: VSP 5500, VSP 5100
 - R900G: VSP 5100H, VSP 5500H
 - R800: VSP G1x00, VSP F1500
 - M900S: VSP E590
 - M900SG: VSP E590H
 - M900M: VSP E790
 - M900MG: VSP E790H
 - M900H: VSP E990
 - RH10MHF: VSP E1090
 - RH10MHG: VSP E1090H
 - RH20ML2: VSP One B24
 - RH20M2: VSP One B26
 - RH20MH2: VSP One B28
 - M850XS: VSP G130
 - M850S1: VSP G350
 - M850S1F: VSP F350
 - M850S2: VSP G370
 - M850S2F: VSP F370
 - M850M3: VSP G700
 - M850M3F: VSP F700
 - M850H: VSP G900
 - M850HF: VSP F900
 - M800S: VSP G200
 - M800M: VSP G400, VSP G600, VSP F400, VSP F600
 - M800H: VSP G800, VSP F800

Storage System	Virtual storage machine that can be created in the storage system: Identifier of the virtual storage machine									
	R500 RK500	R600 RK600	R700 M700	R800	M800 M800H	M850XS M850S1 M850S1F M850S2 M850S2F M850M3 M850M3F M850H M850HF	RH10HG RH10MH G RH10HF RH10MHF	R900G R900F	M900H M900M M900S M900MG M900SG	RH20MH2 RH20M2 RH20ML2
VSP G1000, VSP G1500, VSP F1500	Y	Y	Y	Y	Y	N	N	N	N	N
VSP G100, VSP G200, G400, G600, G800 and VSP F400, F600, F800	Y	Y	Y	Y	Y	N	N	N	N	N
VSP G130, G/ F350, G/F370, G/F700, G/ F900	Y	Y	Y	Y	Y	Y	N	N	N	N
VSP 5100, VSP 5500, VSP 5100H, and VSP 5500H	Y	Y	Y	Y	Y	Y	Y	Y	N	N

Storage System	Virtual storage machine that can be created in the storage system: Identifier of the virtual storage machine									
	R500 RK500	R600 RK600	R700 M700	R800	M800 M800H	M850XS M850S1 M850S1F M850S2 M850S2F M850M3 M850M3F M850H M850HF	RH10HG RH10MH G RH10HF RH10MHF	R900G R900F	M900H M900M M900S M900MG M900SG	RH20MH2 RH20M2 RH20ML2
VSP 5200, VSP 5600, VSP 5200H, and VSP 5600H	Y	Y	Y	Y	Y	Y	Y	Y	Y	N
VSP E590, VSP E790, VSP E990, VSP E1090, VSP E590H, and VSP E790H, VSP E1090H	Y	Y	Y	Y	Y	Y	Y	Y	Y	N
VSP One B24, VSP One B26, VSP One B28	N	N	Y	Y	Y	Y	Y	Y	Y	Y
Legend Y: Supported, N: Not supported										

-resource_id <resource_group_id>

Specifies the resource group ID (1-1023) (for example, 5).

-ldev_id <ldev#>

Specifies the LDEV number (0-65279) (for example, -ldev_id 200).

-port <port#> [<host group name>]

Specifies the port number, host group ID, and host group name. If the host group name is more than 64 characters, use the host group ID instead of the host group name (for example, CL1-A or CL1-A-g, where g is from 0 to 255).

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno:1-52, sgno:1-32) (for example, 3-1).

-external_grp_id <gno-sgno>

Specifies the external volume group number (gno:1-16384, sgno:1-4096) (for example, 52-11, "E" is not required).

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

If you want to register an NVM subsystem in a virtual storage machine created by a user, specify an unimplemented NVM subsystem.

-grp_opt ldev -device_grp_name <device group name> [<device name>]

Specifies the name of the device group (maximum 32 characters).

To specify an LDEV in the device group, use the device name of the LDEV (maximum 32 characters).

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, this command is applied to all LDEVs in the specified device group.

Examples

Creating a resource group of resource group name: sql_srv.

```
# raidcom add resource -resource_name sql_srv
```

Creating a virtual storage machine: rsg_vir, and the serial number of the virtual storage machine: 1000.

```
# raidcom add resource -resource_name rsg_vir -virtual_type 1000 R800
```

Changing the resource group name of the resource group ID:5 to sql_srv.

```
# raidcom add resource -resource_name sql_srv -resource_id 5
```


Add LDEV: 400 to the resource group: sql_srv.

```
# raidcom add resource -resource_name sql_srv -ldev_id 400
```

Adding a port of CL1-A to a resource group: sql_srv.

```
# raidcom add resource -resource_name sql_srv -port CL1-A
```

Adding a host group of CL1-A-0 to a resource group: sql_srv.

```
# raidcom add resource -resource_name sql_srv -port CL1-A-0
```

Adding a parity group:5-2 to a resource group: sql_srv.

```
# raidcom add resource -resource_name sql_srv -parity_grp_id 5-2
```

Adding an external volume group:01-02 to a resource group: sql_srv.

```
# raidcom add resource -resource_name sql_srv -external_grp_id 01-02
```

Adding an NVM subsystem ID: 1 to a resource group: sql_srv.

```
# raidcom add resource -resource_name sql_srv -nvm_subsystem_id 1
```

raidcom modify resource

Sets the virtual storage mode of a resource group. This is an asynchronous command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify resource -resource_name <resource group name>
                        -virtual_switch <y/n>
```

Options and parameters

-resource_name <resource group name>

Specifies the resource group name. Up to 32 characters can be specified.

-virtual_switch {y|n}

Sets the virtual storage mode of the resource group to ON/OFF.

- y or yes: Enable the virtual storage mode
- n or no: Disable the virtual storage mode

Examples

Resource group: Enable the virtual storage mode of sql_srv.

```
# raidcom modify resource -resource_name sql_srv -virtual_switch y
```

raidcom delete resource

Deletes resource groups. You can delete a resource group only after all resources that are registered to that resource group have been deleted.

LDEV number, port number, host group number, parity group, and external group are deleted from the specified resource groups. The deleted resources are moved to resource group 0. Even if the specified resources do not exist in the specified resource group, the specified resources are moved to resource group 0, and then the command is complete. You can specify a device group name instead of a LDEV number.

If an LDEV is a volume that configures a pool, journal, or LUSE, all LDEVs of the pool, journal, or LUSE must be assigned to the same resource group.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Note:

- For VSP One B20, release the resource group by using the following REST API before or after running the **raidcom delete resource** command:

```
[POST<base URL>/v1/objects/user-groups/<object ID>/actions/remove-resource-group/invoke
```

For details, see the *VSP One Block Administrator System Administrator Guide*.

- For the storage systems other than VSP One B20, use Device Manager - Storage Navigator to delete resource groups that are assigned to user groups where All Resource Groups Assigned is set to No. Do not use the **raidcom delete resource** command by CCI.

Procedure for deleting resource groups:

1. Release the assigned information to the user group where the All Resource Groups Assigned is set to No in the **Edit Resource Groups Assignment** window.
2. Delete the resource group.

If the resource group is deleted before the assigned information is released, the deleted resource group ID is displayed in the **Edit Resource Groups Assignment** window. The deleted resource group name is not displayed.

Syntax

```
raidcom delete resource -resource_name <resource group name>
    [-ldev_id <ldev#> | -port <port#> [<host group name>] | -parity_grp_id <gno-
sgno> | -external_grp_id
    <gno-sgno> | -grp_opt <group option> -device_grp_name <device group name>
[<device name>]] -nvm_subsystem_id <nvm subsystem id>]
```

Options and parameters

-resource_name <resource group name>

Specifies the resource group name. Up to 32 characters can be specified.

-ldev_id <ldev#>

Specifies the LDEV number (0-65279). For example:

- -ldev_id 200

-port <port#> [<host group name>]

Specifies the port number, host group ID, and host group name. If the host group name is more than 64 characters, use the host group ID instead of the host group name. For example:

- CL1-A or CL1-A-g (g is from 0 to 255)

-parity_grp_id <gno-sgno>

Specifies the parity group number (gno:1-52, sgno:1-32). For example:

- 3-1

-external_grp_id <gno-sgno>

Specifies the external volume group number (gno:1-16384, sgno:1-4096). For example:

- 52-11 ("E" is not required)

-grp_opt <group option>

Specifies the device information about the LDEV in the device group. Specify "ldev" (fixed). The information about the LDEV in the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be deleted.

To specify the specific LDEV in the device group, specify the device name of the LDEV (maximum 32 characters) in the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, the command is applied to all LDEVs in the device group.

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

If you want to delete an NVM subsystem from a virtual storage machine created by a user, specify an unimplemented NVM subsystem.

Examples

Deleting the LDEV: 400 from the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv -ldev_id 400
```

Deleting a port of CL1-A from the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv -port CL1-A
```

Deleting a host group of CL1-A-0 from the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv -port CL1-A-0
```

Deleting a parity group: 5-2 from the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv -parity_grp_id 5-2
```

Deleting an external volume group: 01-02 from the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv -external_grp_id 01-02
```

Deleting the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv
```

Deleting the NVM subsystem ID: 0 from the resource group: sql_srv.

```
# raidcom delete resource -resource_name sql_srv -nvm_subsystem_id 0
```

raidcom get resource

Displays resource group information.

Syntax

```
raidcom get resource -key <option>
```

Options and parameters

-key <option>

Specify "opt" as the option to display resource group information on the virtual mode.

Examples

Displaying a resource group and resource group lock information.

```
# raidcom get resource
```

RS_GROUP	RGID	stat	Lock_owner	Lock_host	Serial#
meta_resource	0	Unlocked	-	-	64556

Displaying a resource group status on the virtual mode.

```
#raidcom get resource -key opt
```

RS_GROUP	RGID	V_Serial#	V_ID	V_IF	Serial#
meta_resource	0	302624	R8	Y	302624
USP_002	1	64035	R5	Y	302624

Description of each column in output example:

RS_GROUP

Displays the resource group name.

RGID

Displays the resource group ID. RGID=0 is used for meta resource group.

stat

Displays the locking status of the resource group name.

Lock_owner

Displays the (authorized) user who locks the resource group. A hyphen (-) is displayed if the user who locks the resource group has already logged out, or the resource group is not locked.

Lock_host

Displays the host name of a user who locks the resource group. A hyphen (-) is displayed if the user who locks the resource group has already logged out, or the resource group is not locked.

If REST API locks the resource group, the IP address that is used internally in the storage system might be displayed.

Serial#

Displays the product serial number.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

V_Serial#

Displays the product serial number on the virtual storage mode.

- The serial number for VSP 5000 series is displayed with a "5" at the beginning ("512345" = serial number 12345).
- The serial number for VSP G1x00 and VSP F1500 is displayed with a "3" at the beginning ("312345" = serial number 12345).

V_ID

Displays the storage system type:*

- RH10HF: VSP 5200, VSP 5600
- RH10HG: VSP 5200H, VSP 5600H
- R9F: VSP 5100, VSP 5500
- R9G: VSP 5100H, VSP 5500H
- R8: VSP F1500, VSP G1x00
- RH20ML2: VSP One B24
- RH20M2: VSP One B26
- RH20MH2: VSP One B28
- RH10MHF: VSP E1090
- RH10MHG: VSP E1090H
- M9H: VSP E990
- M9MG: VSP E790H
- M9M: VSP E790
- M9S: VSP E590
- M9SG: VSP E590H
- M850HF: VSP F900
- M850H: VSP G900
- M850M3F: VSP F700
- M850M3: VSP G700
- M850S2F: VSP F370
- M850S2: VSP G370
- M850S1F: VSP F350
- M850S1: VSP G350
- M850XS: VSP G130
- M8H: VSP G800, VSP F800
- M8M: VSP G400, G600, VSP F400, F600
- M8S: VSP G200

* The storage system uses the storage system type information to determine the model information to respond to the host. The storage system type information for a resource group that does not belong to a virtual storage machine matches the actual model information of the storage system. In addition, when the storage system controller has been upgraded and changed, the storage system type information before the controller upgrade is displayed and does not match the actual storage model. To see the actual model information of a storage system whose controller has been upgraded, refer to the "MODEL" information of the **raidcom get system** command.

V_IF

Displays the enabled status of the virtual storage mode:

- Y: The virtual storage mode is enabled.
- N: The virtual storage mode is disabled.

raidcom lock resource

This locks the specified resource group name.

When you perform these commands, lock the resource group to which resource is allocated before executing the command.

- **add**
- **delete**
- **modify**
- **initialize**
- **check_ext_storage**
- **disconnect**
- **set**
- **reset**
- **reallocate**
- **monitor**

If the specified resource group name does not exist, this command is rejected with EX_ENOOBJ. If the resource group to be locked is deleted while executing the **raidcom lock resource** command, the command might fail with the error EX_CMDRJE (SSB1=2E30, SSB2=0026). If the command fails with this error, retry the command.

In authentication mode, the user executing this command must have a permission for the resource group name.

Syntax

```
raidcom lock resource [-resource_name <resource group name>]
                    [-time <time(sec)>] [-automatic_unlock <time>]
```

Options and parameters**[-resource_name <resource group name>]**

Specifies the name of resource group (maximum 32 characters).

Specify defined resource group names.

If this parameter is not specified, all resource groups that are assigned to the user are locked.

[-time <time(sec)>]

This parameter is used for specifying the latency (0 to 21474836) until the specified resource is locked.

The TOV time of the lock instruction is specified.

When <time> is specified as "0", it is executed as "nowait (no waiting time)" mode.

If this parameter is not specified, the default waiting time (7200 seconds) is used.

[-automatic_unlock <time>]

Specifies the latency until resource lock is released automatically. If this option is specified, resource lock is obtained, and then released automatically if another raidcom command for configuration setting and reference (except raidcom -login or raidcom -h) is not executed within the time period specified by <time>.

<time> must be specified in seconds. When 0 is specified, resource lock is not released automatically. If the specified value is 10 or smaller, resource lock might be released automatically during command execution.

Examples

Resource group: Locking the resource of the meta_resource.

```
# raidcom lock resource -resource_name meta_resource
```

raidcom unlock resource

This unlocks the specified resource group name.

If the specified resource group name does not exist, this command is rejected with EX_ENOOBJ.

In authentication mode, a user executing this command must have a permission for the resource group name.

Syntax

```
raidcom unlock resource [-resource_name <resource group name>]
```


Options and parameters

[-resource_name <resource group name>]

Specifies the name of resource group (maximum 32 characters).

Specify defined resource group names.

If this parameter is not specified, all resource groups that are assigned to the user are unlocked.

Examples

Resource group: Unlocking the resource of the meta_resource.

```
# raidcom unlock resource -resource_name meta_resource
```

raidcom map resource

Arrange a resource to the virtual storage machine. This is a synchronous command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If you specify the `-request_id` option for the asynchronous command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the `raidcom get command_status` command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom map resource {-ldev_id <ldev#> -virtual_ldev_id
    {<ldev#>|reserve} [-ssid<ssid>] [-emulation <emulation type>][-request_id auto]
    | -port <port#> -virtual_port <port#>}
```

Options and parameters

-ldev_id <ldev#>

Specify an LDEV ID (0-65279).

Example:

- -ldev_id 400

-virtual_ldev_id {<ldev#>| reserve}

Specify an LDEV ID (0-65279) to be used in the virtual storage machine.

If you specify "reserve" instead of the LDEV ID, the reserve attribute of global-active device is set.

Example:

- -virtual_ldev_id 100
- -virtual_ldev_id reserve

-ssid <ssid>

Specify an SSID related to an LDEV in the virtual storage machine.

[-emulation <emulation type>]

Specify the emulation type of a relevant LDEV on the virtual storage machine. This setting is reflected in the inquiry response.

Specify the emulation type by adding "*"n" in the LUSE configuration or by adding "-CVS" in the CVS configuration. ("n" shows the number of LUSE components.)

Apply in order from "*"n" (LUSE configuration) to "-CVS" (CVS configuration) when it is LUSE configuration and CVS configuration.

Example:

- -emulation OPEN-3-CVS
- -emulation OPEN-3*6
- -emulation OPEN-3*6-CVS

[-request_id auto]

Supported storage systems:

- VSP One B20 (DKCMAIN firmware: All versions)
- VSP 5000 series (DKCMAIN firmware version 90-08-41-XX/XX or later)
- VSP E series (DKCMAIN firmware version 93-06-41-XX/XX or later)

It is recommended to specify this option when running the command. The command is run as an asynchronous command only when this option is specified.

<request #> is the Request ID assigned for each command run. For details, see [Request ID function \(on page 293\)](#).

-port <port#>

Supported storage systems:

- VSP
- HUS VM

Specify a port number. Specify the port number whose attribute is Target.

Example:

- CL1-A

-virtual_port <port#>

Supported storage systems:

- VSP
- HUS VM

Specify a port number to be used in the virtual storage machine.

Example:

- CL3-B

Examples

Create the virtual LDEV100 in the LDEV400 (run as a synchronous command).

```
# raidcom map resource -ldev_id 400 -virtual_ldev_id 100
```

Create the virtual LDEV100 in the LDEV400 (run as an asynchronous command).

```
raidcom map resource -ldev_id 400 -virtual_ldev_id 100 -request_id auto REQID : 1
```

Set the global-active device reserve attribute to the LDEV400.

```
# raidcom map resource -ldev_id 400 -virtual_ldev_id reserve
```

raidcom unmap resource

Cancel the resource arrangement in the virtual storage machine. This is a synchronous command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If you specify the `-request_id` option for the asynchronous command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom unmap resource {-ldev_id <ldev#> -virtual_ldev_id  
{<ldev#>|reserve}[-request_id auto] | -port <port#> -virtual_port <port#>}
```

Options and parameters

-ldev_id <ldev#>

Specify an LDEV ID (0-65279).

Example:

- -ldev_id 400

-virtual_ldev_id {<ldev#>| reserve}

Specify an LDEV ID (0-65279) to be used in the virtual storage machine.

If you specify "reserve" instead of the LDEV ID, the reserve attribute of global-active device is released.

Example:

- -virtual_ldev_id 100
- -virtual_ldev_id reserve

[-request_id auto]

Supported storage systems:

- VSP One B20 (DKCMAIN firmware: All versions)
- VSP 5000 series (DKCMAIN firmware version 90-08-41-XX/XX or later)
- VSP E series (DKCMAIN firmware version 93-06-41-XX/XX or later)

It is recommended to specify this option when running the command. The command is run as an asynchronous command only when this option is specified.

<request #> is the Request ID assigned for each command run. For details, see [Request ID function \(on page 293\)](#).

-port <port#>

Specify a port number. Specify the port number whose attribute is Target.

Example:

- CL1-A

-virtual_port <port#>

Specify a port number to be used in the virtual storage machine.

Example:

- CL3-B

Examples

Cancel a virtual LDEV100 in an LDEV400 (run as a synchronous command).

```
# raidcom unmap resource -ldev_id 400 -virtual_ldev_id 100
```

Cancel a virtual LDEV100 in an LDEV400 (run as an asynchronous command).

```
# raidcom unmap resource -ldev_id 400 -virtual_ldev_id 100 -request_id auto
REQID : 1
```

Release the global-active device reserve attribute to the LDEV400.

```
# raidcom unmap resource -ldev_id 400 -virtual_ldev_id reserve
```

Cancel a virtual port CL2-B in a port CL1-A.

```
# raidcom unmap resource -port CL1-A -virtual_port CL2-B
```

raidcom add snap_pool

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Creates pools and adds pool VOLs for Thin Image and Copy-on-Write Snapshot by the specified resource.

When specifying a pool that is already created for Thin Image or Copy-on-Write Snapshot, the specified resource is added as a pool volume. A parity group, an LDEV, and a device group can be specified as a resource.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.



Tip: For VSP 5000 series, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, specifying a parity group is recommended. When a parity group is specified, an LDEV is created and a pool volume is added. It is not necessary to create an LDEV in advance. In addition, you cannot add an existing LDEV as a pool volume.

Specify either one of Pool ID or Pool Name certainly. If both the Pool ID and Pool Name options are omitted, this command is rejected with EX_REQARG.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

Syntax

```
raidcom add snap_pool {-pool_id <pool ID#> | -pool_name <pool
naming> | -pool_id <pool ID#>
-pool_name <pool naming> }
{-parity_grp_id <gno-sgno> [-resource_id <resource group_id >] | -ldev_id <ldev#> ...[-
cnt<count>] | -grp_opt <group option>
-device_grp_name <device group name> [<device name>]}
[-user_threshold <%> ] [-thinsnap] [-auto_add_poolvol {enable|disable}]
```

Options and parameters

-pool_id <pool ID#>

Specifies the Pool ID (0-127) of a Thin Image or Copy-on-Write Snapshot pool.

If a `-pool_id` option is non-numeric, the specified value is recognized as a pool name to identify the pool ID.

When you omit specifying the `-pool_id` option, you need to specify `-pool_name` option.

When specifying the `-pool_name` option without a `-pool_id` option, the specified pool is searched first. If the pool exists, the pool volume is added to the corresponding pool. If the pool does not exist, a new pool is created and is assigned a pool ID. However, if the pool does not exist and multiple commands without the `-pool_id` option are executed simultaneously, the same pool ID might be assigned to multiple newly created pools.

To avoid this behavior, use the `raidcom lock resource` command in advance to lock all resource groups. If all the resource groups are locked, commands executed by other users cannot take out unused pool IDs from all the resource groups. If you create multiple pools, execute the next `raidcom add snap_pool` command after the execution of the current `raidcom add snap_pool` command is complete.

After the execution of all the `raidcom add snap_pool` commands are complete, use the `raidcom unlock resource` command to unlock all the resource groups.

-pool_name <pool naming>

Specifies the pool name of a pool for Thin Image or Copy-on-Write Snapshot. Up to 32 characters can be specified.

When the specification of Pool Name is omitted, a Pool ID must be specified. When the Pool ID is specified and the Pool Name is omitted, a pool name is allocated automatically in the form of "New Pool<number>".

The `-pool_name` option cannot be specified by configuring only in numeric because numeric values are preferentially identified as a pool ID. Specifies a pool ID with the `-pool_id<pool ID#>` option.

The existing pool name is changed if all of the following conditions are met:

- The pool ID is specified.
- The pool name is specified.
- A pool with the specified pool ID already exists.
- The `-parity_grp_id` option is not specified.

-parity_grp_id <gno>-<sgno> [-resource_id <resource group_id >]

Supported storage systems:

- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the parity group number (gno: 1 to 52, sgno: 1 to 32).

If this option is specified, an LDEV is automatically created in the specified parity group, and then a pool is created or expanded on the created LDEV.

For example: 3-1

When `-resource_id <resource grp id>` is specified, the LDEV that is not installed in the specified resource group and has the largest LDEV number should be selected to create a pool volume. This option is valid only when a pool is created.

When `-resource_id <resource grp id>` is not specified, the largest LDEV number in the resource group for which the user has operation permission should be selected to create a pool volume.

When expanding a pool, the largest LDEV number in the resource group to which an existing pool volume belongs should be selected to create the pool volume.

The command cannot be run if there are not enough unimplemented LDEVs to create the pool volume in the target resource group. In that case, use the `raidcom add resource -ldev_id <ldev#>` command or the `raidcom delete resource -ldev_id <ldev#>` command to move the unimplemented LDEV numbers to the target resource group, or use the `raidcom delete ldev` command to delete unnecessary LDEVs.

-ldev_id <ldev#>

Specifies the LDEV number (0 to 65279). Up to 64 of LDEVs can be specified at a time. For example:

- `-ldev_id 100`
- `-ldev_id 100-110`
- `-ldev_id 100 -cnt 10`

[-cnt <count>]

Specifies the count (2 to 64).

The count becomes singular if not specified.

Up to 64 of LDEVs can be specified at a time.

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging to the device group are operated.

When the `-pool_name` option is omitted, the device group name changes into the pool name.

[-user_threshold <%>]

Specifies the user defined threshold value (20 to 95) (%).

If this option is omitted, "80" is used.

When you add a pool volume, this option is ignored even if it is specified. If you want to change the user defined threshold value of the additional pool volume, execute the **raidcom modify pool** command.

[-auto_add_poolvol {enable|disable}]

Specifies whether to automatically add pool volumes for which accelerated compression is enabled. This option is valid only for creating pools. This option can be specified for VSP G1x00 or VSP F1500.

- **enable:** Automatically adds pool volumes for which accelerated compression is enabled according to the compression ratio of a parity group for which accelerated compression is enabled.
- **disable:** The pool volumes for which accelerated compression is enabled are not added automatically.

When these options are not specified, the default values will be applied as follows:

- **enable:** When the storage system is VSP 5000 series, VSP E series, VSP G/F350, G/F370, G/F700, G/F900.
- **disable:** When the storage system is VSP G1x00 or VSP F1500.

Examples

Using LDEVs:400, 401, and 402, creating a pool of Pool ID:1, Pool Name: my_ss_pool for Thin Image (TI only).

```
# raidcom add snap_pool -pool_id 1 -pool_name my_ss_pool -ldev_id 400 401 402
```

Using LDEVs:410, 411, and 412, creating a pool of Pool ID:3, Pool Name: my_ss_pool for Thin Image (TI only).

```
# raidcom add snap_pool -pool_id 3 -pool_name my_ss_pool -ldev_id 410 411 412 -
thinsnap
```

Using LDEVs:500, 501, and 502, creating a pool of Pool ID: Allocated automatically, Pool Name: my_ss_pool for Thin Image (TI only).

```
# raidcom add snap_pool -pool_name my_ss_pool -ldev_id 500 501 502
```

Using LDEVs:600, 601, and 602, creating a pool of Pool ID: 2, Pool Name: Allocated automatically for Thin Image (TI only).

```
# raidcom add snap_pool -pool_id 2 -ldev_id 600 601 602
```


Using LDEV belonging to the device group: grp1, creating a pool of Pool ID: 1, Pool Name: Allocated automatically for Thin Image (TI only).

```
# raidcom add snap_pool -pool_id 1 -grp_opt ldev -device_grp_name grp1
```

raidcom get snap_pool

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Displays pool information for Thin Image or Copy-on-Write Snapshot.

Syntax

```
raidcom get snap_pool
```

Options and parameters

None.

Examples

Displaying pool information for Thin Image or Copy-on-Write Snapshot.

```
# raidcom get snap_pool
```

```
PID POLS U(%) SSCNT Available(MB) Capacity(MB) Seq# Num LDEV# H(%) FMT_CAP(MB)
003 POLS 100 10000          100 10000000000 62500 1 375 70 100
```

Description of each column in output example:

PID

Pool ID

POLS

Displays status of the pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)

- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is suspended in the failure status.) In this status, the pool information cannot be displayed.

U(%)

Usage rate of the pool, including the mapped capacity and the capacity for Full Allocation

SSCNT

Number of volumes in the pool

Available (MB)

Available capacity (megabytes) for the data volumes in the pool

Capacity (MB)

Total capacity (megabytes) of the pool.

Seq#

Serial number.

Num

Number of LDEVs in the pool

LDEV#

Number of the first LDEV in the pool

H(%):

Threshold value for the pool

FMT_CAP(MB)

Formatted capacity (megabytes) of the pool

- (hyphen): This information is not available for this pool.

raidcom add snapshot

Add a combination of the specified LDEV number and Pool ID to a snapshot group. If there is no name of specified snapshot group, create a new snapshot group. Do not run the **pairxx** command for the pair created by using this command.

Syntax

```
raidcom add snapshot -ldev_id <ldev#(P)> <ldev#(S)> -pool
{<pool ID#> | <pool naming>}
-snapshotgroup <name> [-snap_mode <mode>] [-mirror_id <mu#>]
```

Options and parameters**-ldev_id <ldev#(P)> <ldev#(S)>**

Specifies the LDEV number with the snapshot data created. LDEV numbers for P-VOL must be included. If VOL_ATTR of the specified LDEV is a DRS-VOL and cascade is specified for the `-snap_mode` option, a Thin Image Advanced pair is created.

-pool {<pool ID#> | <pool naming>}

Specifies the pool ID or the pool name created for Snapshot.

-snapshotgroup <name>

Specifies a name to be given for snapshot group.

[-mirror_id <mu#>]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the mirror ID.

If this option is omitted, an unused mirror ID is allocated automatically.

[-snap_mode <mode>]

Specifies the mode to create a snapshot group and the mode to add a P-VOL or S-VOL to the snapshot group. If you specify multiple modes, use a space as a separator.

The following modes can be specified:

- `<mode>= CTG`: For creating in CTG mode. The consistency group number is allocated internally by itself.

If this option is omitted, the snapshot group is created in normal mode.

This option is effective only when a new snapshot group is to be created.

This option is ignored if this is specified for the existing snapshot group.

- `<mode> = clone`: For adding a P-VOL or S-VOL as a volume for which you can create a pair with the clone attribute (clone mode). clone cannot be specified in Thin Image Advanced.
- `<mode>= cascade`: For adding a P-VOL or S-VOL as a volume in which you can create a cascade pair (cascade mode).

Example:

- `-snap_mode CTG clone`
- `-snap_mode CTG`

Examples

Adding a combination of the P-VOL (LDEV number 10:10), the S-VOL (LDEV number 20:20), and the Pool (SnapPool00) to the snapshot group (db1).

```
# raidcom add snapshot -ldev_id 0x1010 0x2020 -pool SnapPool00 -snapshotgroup db1
```

raidcom map snapshot

Maps the specified snapshot data to the S-VOL. The S-VOL to be mapped snapshot data must be created before it is specified.

For a Thin Image Advanced pair, the pair status changes to CPYP while the specified operation is in process, and then changes to PSUS when the process is complete. To verify that the operation is complete, run the `raidcom get snapshot` command to verify that STAT has changed from CPYP to other than CPYP.

Syntax

```
raidcom map snapshot -ldev_id <ldev#(P)> <ldev#(S)> {-snapshotgroup <name> | -mirror_id <mu#>}
```

Options and parameters

-ldev_id <ldev#(P)> <ldev#(S)>

Specifies the LDEV number of P-VOL and S-VOL.

-snapshotgroup <name>

Specifies the name of the snapshot group in which the snapshot data as the operation target is included.

One MU that corresponds to the specified P-VOL is mapped from the specified snapshot group.



Note:

Because the MU is selected automatically, an unexpected MU could be mapped. So if you intend to specify the specific snapshot data, specify the snapshot data by an MU number instead of a snapshot group.

-mirror_id <mu#>

Specifies the mirror ID of a snapshot data to be a target.

Examples

Mapping the snapshot data of the LDEV number 10:10 that is included in the snapshot group (db1) to the LDEV number 20:00.

```
# raidcom map snapshot -ldev_id 0x1010 0x2000 -snapshotgroup db1
```

Mapping the snapshot data of the LDEV number 10:10 and the Mirror ID 10 to the LDEV number 20:00.

```
# raidcom map snapshot -ldev_id 0x1010 0x2000 -mirror_id 10
```

raidcom unmap snapshot

Unmaps the S-VOL which is mapping the snapshot data.

For a Thin Image Advanced pair, the pair status changes to CPYP while the specified operation is in process, and then changes to PSUS when the process is complete. To verify that the operation is complete, run the **raidcom get snapshot** command to verify that STAT has changed from CPYP to other than CPYP.



Caution: Do not run this command in a Thin Image pair managed in the CCI configuration definition file.

If this command is run, an LDEV number of a volume paired with the P-VOL is displayed as 65535 (0xffff), the pair status of the S-VOL is displayed as SMPL, and the Thin Image pair operations cannot be performed by using the **pairXX** command.

If you have run this command, re-map the S-VOL by using the **raidcom map snapshot** command. Alternatively, delete the Thin Image pair from the configuration definition file, and then restart the CCI instance.

Syntax

```
raidcom unmap snapshot -ldev_id <ldev#> [-snapshotgroup <name> | -mirror_id <mu#>]
```

Options and parameters

-ldev_id <ldev#>

Specifies the LDEV number to be unmapped.

Snapshot group name or MU number must be specified when you specify the LDEV number of P-VOL to identify the snapshot data.

Do not specify the snapshot group name and MU number when you specify the LDEV number of S-VOL.

[-snapshotgroup <name>]

Specifies the name of the snapshot group in which the snapshot data as the operation target is included.

One MU which corresponds to the specified P-VOL is unmapped from the specified snapshot group.

**Note:**

Because the MU is selected automatically, an unexpected MU could be unmapped. So if you intend to specify the specific snapshot data, specify the snapshot data by an MU number instead of a snapshot group.

[-mirror_id <mu#>]

Specifies the mirror ID of the snapshot data to be a target of unmapping when you specify the LDEV number of P-VOL.

Examples

Unmapping the S-VOL (LDEV number 20:00).

```
# raidcom unmap snapshot -ldev_id 0x2000
```

Unmapping the snapshot data of the LDEV number 10:10 that is included in the snapshot group (db1).

```
# raidcom unmap snapshot -ldev_id 0x1010 -snapshotgroup db1
```

Unmapping the P-VOL (LDEV number 10:10) and the mirror ID 10.

```
# raidcom unmap snapshot -ldev_id 0x1010 -mirror_id 10
```

raidcom delete snapshot

Deletes the snapshot data and the snapshot group. The relevant snapshot data of the LDEV is deleted from the snapshot group by specifying LDEV number. When no snapshot data left in the snapshot group, the snapshot group is deleted.

Syntax

```
raidcom delete snapshot {-snapshotgroup <name> | -ldev_id <ldev#>
{-mirror_id <mu#> |-snapshotgroup <name> | -range tree}}
```

Options and parameters**-snapshotgroup <name>**

Specifies the snapshot group in which the target data to be deleted is included.

If the snapshot group is specified as the target, all the snapshot data and the snapshot group are deleted.

-ldev_id <ldev#>

Specifies the LDEV number of P-VOL or S-VOL for the snapshot data to be deleted.

When P-VOL is specified, specify the snapshot data by specifying the MU number or the snapshot group (Specifying the MU number or the snapshot group is mandatory).

When you specify the S-VOL, do not specify a MU number or a Snapshot group. If you specify the MU number or the Snapshot group, the P-VOL of specified LDEV number becomes the subject of deletion.

-mirror_id <mu#>

Specifies the Mirror ID of the snapshot data to be deleted.

-snapshotgroup <name>

Specifies the snapshot group to be deleted.

The smallest number of MU in the snapshot group becomes the subject to be deleted.

-range tree

Specify this option to delete snapshot data in all layers of a snapshot tree for which the root volume is specified by -ldev_id <ldev#>. This option is not supported by Thin Image Advanced. To specify this option, the following conditions must be met:

- Root volume is an LDEV specified by -ldev_id <ldev#>.
- The pair to be deleted was created in the cascade or clone mode.
- The user has the operation authority for volumes of all pairs to be deleted.



Note:

When multiple snapshot trees are cloned, specify the root volume of the first snapshot tree for -ldev_id <ldev#>, and execute the command. Only the first snapshot tree is deleted, and the status of the pairs under the first snapshot tree changes to PSUE. Specify the root volume of a snapshot tree under the first snapshot tree for -ldev_id <ldev#>, and execute the command again. When this option is specified, the SIM code (4b3xxx) that indicates abnormal termination of the Thin Image option might be output. However, the snapshot tree has been deleted normally and there is no problem.

Examples

Deleting the snapshot data of the snapshot group (db1).

```
# raidcom delete snapshot -snapshotgroup db1
```

Deleting the snapshot data of the P-VOL (LDEV number 10:10) and the Mirror ID (10).

```
# raidcom delete snapshot -ldev_id 0x1010 -mirror_id 10
```

Deleting the snapshot data of the LDEV number 10:10 that is included in the snapshot group (db1).

```
# raidcom delete snapshot -ldev_id 0x1010 -snapshotgroup db1
```

```
# raidcom delete snapshot -snapshotgroup db1 -ldev_id 0x1010
```

Deleting the snapshot data of the S-VOL (LDEV number 20:10).

```
# raidcom delete snapshot -ldev_id 0x2010
```

Deleting the snapshot data in all layers of a snapshot tree whose root volume has LDEV number 20:10.

```
# raidcom delete snapshot -ldev_id 0x2010 -range tree
```

raidcom modify snapshot

Operate the specified snapshot group.

Syntax

```
raidcom modify snapshot -ldev_id <ldev#> {-snapshotgroup <name>
| -mirror_id <mu#>} -snapshot_data <op> [-copy_pace <copy pace>][-retention
<retention period (hour)>]
```

```
raidcom modify snapshot -snapshotgroup <name> -snapshot_data <op>
[-copy_pace <copy pace>][-retention <retention period (hour)>]
```

```
raidcom modify snapshot -ldev_id <ldev#> -snapshot_data <op>
[-copy_pace <copy pace>][-retention <retention period (hour)>]
```

```
raidcom modify snapshot -snapshotgroup <current name>
<new name> -snapshot_data rename
```

Options and parameters

-ldev_id <ldev#>

Specifies the LDEV number of P-VOL (or S-VOL) to be performed.

-snapshotgroup <name>

Specifies the snapshot group name in which the snapshot data is included.

-mirror_id <mu#>

Specifies the mirror ID of a snapshot data.

-snapshot_data <op>

Specifies the operation to be performed for the specified snapshot group. The parameter of the operation to be specified is the following. (Note that parameters clone, delete_garbage, and stop_deleting_garbage cannot be used in Thin Image Advanced.):

- create or split: Takes snapshot data by splitting a pair.^{1, 4, 6}

If the `-retention <retention period (hour)>` option is specified and the `-snapshotgroup <name>` option is used to specify a snapshot group that is not a consistency group, the snapshot data retention period is set in order for all pairs belonging to the specified snapshot group. As a result, the time that the snapshot data retention period expires might differ for each pair. If you want to set the same time for the expiration time, use the consistency group.

- resync: Deletes the created snapshot data.¹
- restore: Restores the Thin Image pair.¹

For a Thin Image Advanced pair, perform the following steps:

1. Verify that the status of the pair to be restored is PSUS or SSUS.
2. Run the `raidcom modify snapshot -snapshot_data restore`, and then verify that the command is complete normally.
3. Wait until the status of the pair to be restored migrates to a status that is not RCPY.
4. If the pair status changes to PSUS or SSUS in step 3, the pair restore operation is successfully complete.

- clone: Clones pairs.^{2, 5}
- delete_garbage (VSP 5000 series): Starts the snapshot garbage data deletion processing (defragmentation)³ of the specified Thin Image root volume.^{3, 5}
- stop_deleting_garbage (VSP 5000 series): Stops the snapshot garbage data deletion (defragmentation)³ of the specified Thin Image root volume.^{3, 5}
- renew_retention: (VSP 5000 series, VSP One B20, VSP E series) Updates the snapshot data retention period to a value specified in the `-retention <retention period (hour)>` option.

Specify the `-retention <retention period (hour)>` option so that the value is the same as or more than the existing snapshot data retention period. If a value less than the existing snapshot data retention period is specified, the command is rejected. If the `-retention <retention period (hour)>` option is not specified, the command is also rejected.

If a snapshot group is specified by using the `-snapshotgroup <name>` option, the snapshot data retention period is set in order for all pairs belonging to the specified snapshot group. As a result, the time that the snapshot data retention period expires might differ for each pair. This behavior is the same even when using consistency groups.

If this parameter and the `-ldev_id <ldev#>` option are specified at the same time, also specify the `-snapshotgroup <name>` or `-mirror_id <mu#>` option.

- **vcclone:** (VSP One Block 20) Takes snapshot data by creating a virtual clone. This operation can only be performed in a Thin Image Advanced pair in the PAIR status that uses the S-VOL with the vClone Parent attribute disabled. If this operation is performed on a pair in the PSUP status, the command ends without taking any actions.

If this operation is specified, `-copy_pace <copy pace>`, `-read_only enable`, or `-retention <retention period (hour)>` option cannot be specified.

- **ss2vc:** (VSP One Block 20) Changes a Thin Image Advanced pair in the pair split (PSUS) status to a volume with the vClone attribute (conversion to vClone volumes). The data on the S-VOL is maintained. This operation can only be performed in a Thin Image Advanced pair in the PSUS status that uses the S-VOL with the vClone attribute disabled. If this operation is performed on a pair in the PSUP status, the command ends without taking any actions.

If this operation is specified, `-copy_pace <copy pace>`, `-read_only enable`, or `-retention <retention period (hour)>` option cannot be specified.

- **vc2ss:** (VSP One Block 20) Changes a volume with the vClone attribute to a Thin Image Advanced pair in the pair split (PSUS) status (conversion to snapshots). The data on the S-VOL is maintained. This operation can only be performed in a Thin Image Advanced pair in the PAIR status that uses the S-VOL with the vClone attribute enabled. If this operation is performed on a pair in the PSUP or PSUS status, the command ends without taking any actions.

If this operation is specified, the `-copy_pace <copy pace>` or `-retention <retention period (hour)>` options cannot be specified.



Note:

1. Clone mode pairs do not support this operation.
2. Cascade mode pairs do not support this operation.
3. For the details about the defragmentation, see the Hitachi Thin Image User Guide.
4. For a Thin Image Advanced pair, the pair status changes to PSUP while the specified operation is in process, and then changes to PSUS when the processing is complete. To verify that the operation is complete, run the `raidcom get snapshot` command to verify that STAT has changed from PSUP to PSUS, SSUS, or PSUE.
5. This operation is not available for a Thin Image Advanced pair.
6. For a Thin Image Advanced pair, the group operation with the `-snapshotgroup` option specified might not be performed if all of the following conditions are met. If the following conditions are met, perform the operation on one pair at a time by specifying the `-ldev_id` and `-mirror_id` options.
 - The snapshot group is not in the CTG mode.
 - The snapshot group contains two or more pairs created using the same P-VOL.

[-copy_pace <copy pace>]

Specifies the copy speed. Enabled only when `clone` is specified for `-snapshot_data`. The following speeds can be specified:

- slower
- medium
- faster

When `clone` is specified for `-snapshot_data`, if you did not specify the copy speed, `medium` is set. When you specify a value other than `clone`, if you specify the copy speed, this option is not enabled.

-snapshotgroup <current name> <new name> -snapshot_data rename

Specifies when you change the snapshot group name. The snapshot group name specified in `<current name>` is changed to the name specified in `<new name>`.

[-retention <retention period (hour)>]

Specifies the snapshot data retention period (hour).

This option can be specified only for VSP 5000 series, VSP One B20, VSP E series. This option is enabled if all of the following conditions are met:

- The operation must be either of the following:
 - Create a new snapshot data by specifying `create` or `split` for the `-snapshot_data` option.
 - Update the set snapshot data retention period (hour) by specifying `renew_retention` for the `-snapshot_data` option.
- The operation target pair must be a Thin Image Advanced pair.

If `create` or `split` is specified for the `-snapshot_data` option, zero cannot be specified for this option.

The following shows the operation of the snapshot data with the combination of options and parameters:

When taking the snapshot data by splitting a pair (when specifying create/split)

#	Ways to specify the parameter	CTG mode	Normal mode
1	LDEV number and snapshot group	P-VOL in the snapshot group	P-VOL in the snapshot group and the smallest number of MU
2	LDEV number and MU number	Specified P-VOL and MU number	Specified P-VOL and MU number

#	Ways to specify the parameter	CTG mode	Normal mode
3	LDEV number only (Specifying S-VOL)	P-VOL and MU number that are mapped by the specified S-VOL	P-VOL and MU number that are mapped by the specified S-VOL
4	Snapshot group	All the P-VOLs related to the snapshot group The consistency is guaranteed.	All the P-VOLs related to the snapshot group* The consistency is not guaranteed.
<p>* CCI performs the processing sequentially for each pair in the snapshot group. If an error occurs on a certain pair during the processing and the processing cannot be performed (for example, if a snapshot data retention period is set for some pairs in a snapshot group, the processing for those pairs cannot be performed and is suspended), the command stops the processing at that point and returns an error. In this case, the processing is not performed for any pairs after the pair with the error. After resolving the cause of the error, retry the command, or run the command for each pair after the pair with the error individually.</p>			

When discarding or restoring the snapshot data (when specifying resync/restore)

#	Ways to specify the parameter	CTG mode	Normal mode
1	LDEV number and snapshot group	P-VOL in the snapshot group	P-VOL in the snapshot group and the smallest number of MU
2	LDEV number and MU number	Specified P-VOL and MU number	Specified P-VOL and MU number
3	LDEV number only (Specifying S-VOL)	P-VOL and MU number that are mapped by the specified S-VOL	P-VOL and MU number that are mapped by the specified S-VOL
4	Snapshot group	All the P-VOLs related to the snapshot group	All the P-VOLs related to the snapshot group*
<p>* CCI performs the processing sequentially for each pair in the snapshot group. If an error occurs on a certain pair during the processing and the processing cannot be performed (for example, if a snapshot data retention period is set for some pairs in a snapshot group, the processing for those pairs cannot be performed and is suspended), the command stops the processing at that point and returns an error. In this case, the processing is not performed for any pairs after the pair with the error. After resolving the cause of the error, retry the command, or run the command for each pair after the pair with the error individually.</p>			

When deleting (defragmenting) the snapshot garbage data (when specifying delete_garbage/stop_deleting_garbage) (VSP 5000 series)

#	Ways to specify the parameter	CTG mode	Normal mode
1	LDEV number and snapshot group	Cannot be specified	Cannot be specified
2	LDEV number and MU number	Cannot be specified	Cannot be specified
3	LDEV number only (Specifying root volume)	Snapshot data area of the specified root volume	Snapshot data area of the specified root volume
4	Snapshot group	Cannot be specified	Cannot be specified

Update for the snapshot data retention period (when specifying renew_retention) (VSP 5000 series, VSP One B20, VSP E series)

#	Ways to specify the parameter	CTG mode	Normal mode
1	LDEV number and snapshot group	P-VOL in the snapshot group	P-VOL in the snapshot group and the smallest MU number
2	LDEV number and MU number	The specified P-VOL and MU number	The specified P-VOL and MU number
3	Only LDEV number (specifying S-VOL)	Cannot be specified	Cannot be specified
4	Snapshot group	All P-VOLs associated with the snapshot group	All P-VOLs associated with the snapshot group*

* CCI performs the processing sequentially for each pair in the snapshot group. If an error occurs on a certain pair during the processing and the processing cannot be performed (for example, if a snapshot data retention period is set for some pairs in a snapshot group, the processing for those pairs cannot be performed and is suspended), the command stops the processing at that point and returns an error. In this case, the processing is not performed for any pairs after the pair with the error. After resolving the cause of the error, retry the command, or run the command for each pair after the pair with the error individually.

**When taking the snapshot data by creating a virtual clone (when specifying vclone)
(VSP One B20)**

#	Ways to specify the parameter	CTG mode	Normal mode
1	LDEV number and snapshot group	P-VOL in the snapshot group	P-VOL in the snapshot group and the smallest MU number
2	LDEV number and MU number	The specified P-VOL and MU number	The specified P-VOL and MU number
3	Only LDEV number (specifying S-VOL)	P-VOL and MU number to be mapped to the specified S-VOL	P-VOL and MU number to be mapped to the specified S-VOL
4	Snapshot group	All P-VOLs associated with the snapshot group The data consistency is guaranteed.	All P-VOLs associated with the snapshot group The data consistency is not guaranteed.

When performing conversion to vClone volumes or snapshots (when specifying vc2ss/ss2vc) (VSP One B20)

#	Ways to specify the parameter	CTG mode	Normal mode
1	LDEV number and snapshot group	P-VOL in the snapshot group	P-VOL in the snapshot group and the smallest MU number
2	LDEV number and MU number	The specified P-VOL and MU number	The specified P-VOL and MU number
3	Only LDEV number (specifying S-VOL)	P-VOL and MU number to be mapped to the specified S-VOL	P-VOL and MU number to be mapped to the specified S-VOL
4	Snapshot group	All P-VOLs associated with the snapshot group	All P-VOLs associated with the snapshot group

Examples

Creating a snapshot data for the P-VOL (LDEV number 10:10) that is included in the snapshot group (db1).

```
# raidcom modify snapshot -ldev_id 0x1010 -snapshotgroup db1 -snapshot_data create
```

Creating a snapshot data for the P-VOL (LDEV number 10:10) and the mirror ID 10.

```
# raidcom modify snapshot -ldev_id 0x1010 -mirror_id 10 -snapshot_data create
```

Creating a snapshot data for the S-VOL (LDEV number 20:10).

```
# raidcom modify snapshot -ldev_id 0x2010 -snapshot_data create
```

Creating a snapshot data for all the P-VOLs that are included in the snapshot group (db1).

```
# raidcom modify snapshot -snapshotgroup db1 -snapshot_data create
```

Discarding the snapshot data for the P-VOL (LDEV number 10:10) and the mirror ID 10.

```
# raidcom modify snapshot -ldev_id 0x1010 -mirror_id 10 -snapshot_data resync
```

Restoring the snapshot data for the S-VOL (LDEV number 20:10).

```
# raidcom modify snapshot -ldev_id 0x2010 -snapshot_data restore
```

Changing the snapshot group name (db1) to db2.

```
# raidcom modify snapshot -snapshotgroup db1 db2 -snapshot_data rename
```

Starting the snapshot garbage data deletion (defragmentation) of the root volume (LDEV number 10:10) (only for VSP 5000 series).

```
# raidcom modify snapshot -ldev_id 0x1010 -snapshot_data delete_garbage
```

Stopping the snapshot garbage data deletion (defragmentation) of the root volume (LDEV number 10:10) (only for VSP 5000 series).

```
# raidcom modify snapshot -ldev_id 0x1010 -snapshot_data stop_
deleting_garbage
```

Splitting the snapshot data of all P-VOLs in the snapshot group (db1) and setting the snapshot data retention period to 72 hours (VSP 5000 series, VSP One B20, VSP E series).

```
# raidcom modify snapshot -snapshotgroup db1 -snapshot_data split -
retention 72
```

Updating the snapshot data retention period of all P-VOLs in the snapshot group (db1) to 100 hours (VSP 5000 series, VSP One B20, VSP E series).

```
# raidcom modify snapshot -snapshotgroup db1 -snapshot_data renew_retention -
retention 100
```

raidcom get snapshot

Displays the information of snapshot group and snapshot data that are defined in the device. If this option is omitted, the list of snapshot group is displayed.

Syntax

```
raidcom get snapshot [-ldev_id <ldev#> [-key opt] | -snapshotgroup <name>
[-key {opt|detail|details}]] [-vclone] [-format_time] [{-check_status | -
check_status_not} <string>...
[-time <time>]]
```



Note: You can verify if a Thin Image pair exists in either of the following ways:

- When running the **raidcom get snapshot -snapshotgroup -key detail** command, if the returned value from the command is 205 (EX_ENOOBJ), the Thin Image pair does not exist.
- When running the **raidcom get snapshot -ldev_id <LDEV#> -mirror_id <mu#> -key detail** command, if the returned value from the command is 0 and no information is displayed in the standard output, the Thin Image pair does not exist.

Options and parameters

[-ldev_id <ldev#>]

Specifies the LDEV number to be displayed the snapshot data information. Specifies either one of P-VOL or S-VOL for the LDEV number.

[-snapshotgroup <name>]

Specifies the snapshot group in which you want to display the snapshot data information.

[-key opt]

Specify the **-key opt** option together with the **-ldev_id** or the **-snapshotgroup** option to display the following:

- SMPP as the snapshot data status
- (VSP 5000 series, VSP G130, G/F350, G/F370, G/F700, G/F900) SLU attribute and information of the Diff Clone pair that is being used in the VASA environment
- (VSP 5000 series, VSP G1x00 and VSP F1500, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Status of expansion processing of DP-VOL
- (VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Status of the Read Only attribute of the Snapshot data.

[-key detail]

Specify the `-key detail` option together with the `-ldev_id` or the `-snapshotgroup` option to display the following:

- SMPP as the snapshot data status
- (VSP 5000 series, VSP G130, G/F350, G/F370, G/F700, G/F900) SLU attribute and information of the Diff Clone pair that is being used in the VASA environment
- Status of expansion processing of DP-VOL
- (VSP 5000 series, VSP G1x00 and VSP F1500, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900) Status of the Read Only attribute of the Snapshot data
- (VSP 5000 series, VSP One B20, VSP E series) MODE (R) indicating snapshot data of a Thin Image Advanced pair in the snapshot group mode

Specifying this option expands the display area of MODE.

The following lists the differences in the display format depending on whether to specify this option:

- When the `-key detail` option is not specified:

MODE

-N--

(Displayed in a format available to display up to 4 modes.)

- When the `-key detail` option is specified:

MODE

---AR-----

(Displayed in a format available to display up to 12 modes.)

[-key details]

Specify this option to display the information about the snapshot data retention period in addition to the information displayed when the `-key detail` option is specified.

This option can be specified only for VSP 5000 series, VSP One B20, VSP E series.

To specify this option, also specify the `-ldev_id` or `-snapshotgroup` option.

[-vclone]

(VSP One B20) Specify this option to display the vclone information. If the `-key details` and `-ldev_id <ldev#>` options are specified at the same time as the `-vclone` option, displays the list of volumes with the vClone attribute and Thin Image Advanced pairs in the family to which the specified LDEV belongs. If only the `-vclone` option is specified, displays the list of LDEV numbers of volumes with the vClone Parent attribute on the storage system.

If the DKCMAIN microcode version is not supported the virtual cloning feature, the command ends without displaying anything.

[--format_time]

Specify this option to display the time when a snapshot data (SPLT-TIME) is taken, with the format¹ below. The time zone used in the storage system² is displayed. Zero "0" is added to the beginning if the value of each element is shorter than the regulated length.

Format: YYYY-MM-DDThh:mm:ss

Where:

- YYYY-MM-DD = year-month-date
- hh = time in 24 hours, range 0-23
- mm:ss = minute:second

For example, December 2, 2016, one o'clock p.m. is displayed as 2016-12-02T13:00:00

Notes:

1. "T" in the format separates the date and the time.
2. For VSP 5000 series, VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, the time zone set in the device is used.

For other models, UTC+0 is used.

[-check_status <string> [-time <time>]]

Check if the snapshot group or the snapshot data is in the same state as specified in <string>. If the option contains multiple states, the OR condition check is performed and verifies that the snapshot group or snapshot data is in one of the states contained in the option.

The following strings are specified in the <string>:

- COPY: Copy status.
- PAIR: Pair status.
- PSUS: Suspend status.
- PSUE: Suspend failure status.
- PFUL: Indicates that the Thin Image or Copy-on-Write Snapshot pool exceeds the threshold in the PAIR status.
- PFUS: Indicates that the Thin Image or Copy-on-Write Snapshot pool exceeds the threshold in the PSUS status.
- RCPY: Shows that the copying is in progress by resynchronization.
- SMPP: Indicates that the pair is being deleted. When specifying SMPP, specify -key opt as well.
- PSUP: Indicates that the pair is being split.

If "-time" is specified, the status of the snapshot group or the snapshot data is checked every three seconds until the end of the specified <time> (seconds).

When this option is specified, the returned values are as follows:

- The snapshot group or the snapshot data is in one of the specified states: 0
- The snapshot group or the snapshot data is in none of the specified states (without -time option): 1
- The snapshot group or the snapshot data is in none of the specified states (when the specified <time> passed): EX_EWSTOT

Note:

When you specify -snapshotgroup <name>, and the pair is deleted, EX_ENOOBJ is returned.

[-check_status_not <string> [-time <time>]]

Check if the snapshot group or the snapshot data is not in the same state as specified in <string>. If the option contains multiple states, the OR condition check is performed and verifies that the snapshot group or snapshot data is not in any of the states contained in the option. The following strings are specified in the <string>.

- COPY: Copy status.
- PAIR: Pair status.
- PSUS: Suspend status.
- PSUE: Suspend failure status.
- PFUL: Indicates that the Thin Image or Copy-on-Write Snapshot pool exceeds the threshold in the PAIR status.
- PFUS: Indicates that the Thin Image or Copy-on-Write Snapshot pool exceeds the threshold in the PSUS status.
- RCPY: Shows that the copying is in progress by resynchronization.
- SMPP: Indicates that the pair is being deleted. When specifying SMPP, specify -key opt as well.
- PSUP: Indicates that the pair is being split.

If "-time" is specified, the status of the snapshot group or the snapshot data is checked every three seconds until the end of the specified <time> (seconds).

When this option is specified, the returned values are as follows:

- The snapshot group or the snapshot data is not in any of the specified states: 0
- The snapshot group or the snapshot data is in one of the specified states (without -time option): 1
- The snapshot group or the snapshot data is in one of the specified states (when the specified <time> passed): EX_EWSTOT



Note:

When you specify -snapshotgroup <name>, and the pair is deleted, EX_ENOOBJ is returned.

Examples

Displaying the list of snapshot groups.

```
# raidcom get snapshot
```

SnapShot_name	P/S	STAT	Serial#	LDEV#	MU#	P-LDEV#	PID	%	MODE	SPLT-TIME
snap1	-	-	85000098	-	-	-	-	-	----	-
snap2	-	-	85000098	-	-	-	-	-	----	-
snap3	-	-	85000098	-	-	-	-	-	----	-

Displaying the snapshot data related to the specific P-VOL (LDEV number: 14536).

```
# raidcom get snapshot -ldev_id 14536
```

SnapShot_name	P/S	STAT	Serial#	LDEV#	MU#	P-LDEV#	PID	%	MODE	SPLT-TIME
snap1	P-VOL	PAIR	85000098	14536	1010	-	2	100	----	-
snap2	P-VOL	PAIR	85000098	14536	1011	13000	2	100	G---	-
snap3	P-VOL	PAIR	85000098	14536	1012	-	2	100	----	-

Displaying the snapshot data related to the specific P-VOL (LDEV number: 10).

```
# raidcom get snapshot -ldev_id 10 -key detail
SnapShot_name P/S STAT Serial# LDEV# MU# P-LDEV#
PID % MODE SPLT-TIME SLU C_LDEV# P R D_STAT
snap1 P-VOL PAIR 85000098 10 0
20 0 - ---AR----- - Y - - D PAIR
```

Displaying the snapshot data related to the specific P-VOL (LDEV number: 10) (VSP 5000 series, VSP One B20, VSP E series).

```
# raidcom get snapshot -ldev_id 10 -key details
SnapShot_name P/S STAT Serial# LDEV# MU# P-LDEV# PID %
MODE SPLT-TIME SLU C_LDEV# P R D_STAT RETENTION(H)
snap1 P-VOL PAIR 85000098 10 0 20 0 -
--A-R----- - Y - - D PAIR 72
```

Displaying the snapshot data related to the specific S-VOL (LDEV number: 13000).

```
# raidcom get snapshot -ldev_id 13000
```

SnapShot_name	P/S	STAT	Serial#	LDEV#	MU#	P-LDEV#	PID	%	MODE	SPLT-TIME
snap2	S-VOL	PAIR	85000098	13000	1011	14536	2	100	G---	-

Displaying the snapshot data included in the specific snapshot group.

```
# raidcom get snapshot -snapshotgroup snap2
```

```
SnapShot_name P/S STAT Serial# LDEV# MU# P-LDEV# PID % MODE SPLT-TIME
snap2         P-VOL PAIR 85000098 14536 1011 13000 2 100 G--- -
snap2         P-VOL PAIR 85000098 14537 1011 13001 2 100 G--- -
snap2         P-VOL PAIR 85000098 14538 1011 13002 2 100 G--- -
```

Displaying SPLT-TIME in YYYY-MM-DDThh:mm:ss format.

```
# raidcom get snapshot -ldev_id 14356 -format_time
```

```
SnapShot_name  P/S STAT   Serial# LDEV#   MU# P-LDEV# PID    % MODE SPLTTIME
      snap1 P-VOL PSUS 85000098 14536 1010      -      2 100 ---- 2016-07-
22T10:18:20
      snap2 P-VOL PSUS 85000098 14536 1011    13000 2 100 G--- 2016-07-22T10:18:20
      snap3 P-VOL PSUS 85000098 14536 1012      -      2 100 ---- 2016-07-22T10:18:20
```

Displays the snapshot data with the SLU attribute related to the specific S-VOL (LDEV # 40960) by specifying the `-key opt` option.

```
# raidcom get snapshot -ldev_id 40960 -key opt
SnapShot_name P/S STAT Serial# LDEV# MU# P-LDEV# PID %
MODE SPLT-TIME SLU C_LDEV# P R
Snap2         S-VOL PAIR 85000098 40960 1012 258 2 100
A---          -      Y      - - D
```

Displays lists of volumes with the vClone attribute and Thin Image Advanced pairs in the family to which LDEV# 3080 belongs.

```
raidcom get snapshot -ldev_id 0x3080 -key details -vclone
```

Snapshot_name	P/S	STAT	Serial#	LDEV#	MU#	P-LDEV#	PID	%	
MODE	SPLT-TIME	SLU	C_LDEV#	P	R	D_STAT	RETENTION(H)	PARENT_LDEV#	
ROW		P-VOL	PSUS	800001	3080	3	3081	6	- --A-
RV-----	-	Y	-	N	D	PSUS	-	3080	
ROW		P-VOL	PSUS	800001	3081	3	3082	6	- --A-
RV-----	-	Y	-	N	D	PSUS	-	3080	
-		-	-	800001	3083	-	-	6	- --A-
RV-----	67ad9b5e	Y	-	N	D	-	-	3080	
ROW		P-VOL	PAIR	800001	3084	3	3085	6	- --A-
RV-----	-	Y	-	N	D	PAIR	-	3080	
ROW		P-VOL	PSUS	800001	3084	4	3087	6	- --A-
RV-----	-	Y	-	N	D	PSUS	-	3080	
-		-	-	800001	3084	-	-	6	- --A-
RV-----	67ad9b5e	Y	-	N	D	-	-	3080	
ROW		P-VOL	PAIR	800001	3085	3	3086	6	- --A-
RV-----	-	Y	-	N	D	PAIR	-	3080	
-		-	-	800001	3085	-	-	6	- --A-
RV-----	67ad9b5e	Y	-	N	D	-	-	3080	
-		-	-	800001	3086	-	-	6	- --A-
RV-----	67ad9b5e	Y	-	N	D	-	-	3080	

Description of each column in output example:

Snapshot_name

Displays the name of snapshot group defined in the storage system. For a pair that does not belong to a snapshot group and the information about volumes with the vClone attribute, a hyphen (-) is displayed.

P/S

Displays the attribute of the target LDEV. It displays P-VOL for the P-VOL and S-VOL for the S-VOL. For a snapshot list and information about volumes with the vClone attribute, a hyphen (-) is displayed.

STAT

Displays the following status of each snapshot data:

- SMPP: The pair is being deleted.
- PSUP: The pair is being suspended.
- A hyphen (-) is displayed for the information about volumes with the vClone attribute.
- Other statuses: See **pairdisplay** command.

Serial#

Displays the product serial number.

LDEV#

Displays the LDEV number related to the snapshot.

MU#

Displays the Mirror ID of the P-VOL for the snapshot. A hyphen (-) is displayed for the information about volumes with the vClone attribute.

P-LDEV#

Displays the LDEV number of the volume (P-VOL or S-VOL) of the pair associated with the snapshot data. If the LDEV that is paired belongs to a different virtual storage machine, "----" is displayed. A hyphen (-) is displayed for the information about volumes with the vClone attribute.

PID

Displays the pool ID.

%

When MODE is clone or cascade mode, displays one of the following according to the status of STAT:

- When STAT is COPY, CPYP, RCPY, SMPP, or PSUP: Displays the progress rate of each processing.
- For other statuses: Displays the concordance rate between the P-VOL and the S-VOL of a pair.

When MODE is other than clone or cascade mode, the concordance rate between the P-VOL and the S-VOL is always displayed regardless of the STAT status.

When MODE is R that indicates snapshot data of a Thin Image Advanced pair, a hyphen (-) is always displayed in the concordance rate. A hyphen (-) is displayed for the information about volumes with the vClone attribute.

The following shows the concordance rate for each pair status.

- PSUS: 0% to 100%
- RCPY: 0% to 100%¹
- PAIR: 100%
- COPY: 0% to 100%²

**Note:**

1. Displays the concordance rate before changing to RCPY status. The displayed value might decrease during RCPY status because the copy works.
2. Displays the concordance rate before changing to COPY status. The displayed value does not change during COPY status.

MODE

Displays the status of snapshot data:

- G: The snapshot data created in CTG mode.
- W: The status when the data are written in the secondary volume from the host in the PSUS/PFUS status.
- C: The snapshot data created in clone mode.
- A: The snapshot data created in cascade mode.
- R: The snapshot data of a Thin Image Advanced pair.
- S: The snapshot data retention period is enabled (VSP 5000 series, VSP One B20, VSP E series).
- V: The volume displayed in LDEV# column is the one with the vClone or vClone Parent attribute enabled.
- P: The volume displayed in LDEV# column is the one with the vClone Parent attribute enabled.

SPLT-TIME

Displays the time when a snapshot data is taken by splitting a pair, or a snapshot data is taken by creating a virtual clone. For a Thin Image pair, displays the time when the snapshot data is taken by splitting the pair, or for a volume with the vClone attribute, displays the time when the snapshot data is taken by creating a virtual clone. When `-format_time` option is specified, displays in YYYY-MM-DDThh:mm:ss format. When `-format_time` option is not specified, indicates accumulated time in seconds from January 1, 1970 (GMT).

SLU

Displays whether the SLU attribute is set for the snapshot:

- Y: The SLU attribute is set.
- N: Non-SLU attribute is set.

A hyphen (-) is displayed for unsupported DKCMAIN microcode version.

C_LDEV#

Displays the LDEV number of the diff compare volume (base volume), which is used for comparing the difference between the Diff Clone pair volumes in the VASA environment.

A hyphen (-) is displayed for unsupported microcode version. A hyphen (-) is also displayed for information about volumes with the vClone attribute.

P

Displays the status of expansion processing of DP-VOL (VSP 5000 series, VSP G1x00 and VSP F1500, VSP One B20, VSP E series, and VSP G130, G/F350, G/F370, G/F700, G/F900 only).

- E: Expansion is in progress.
- N: Not processed.
- - (hyphen): This information is not available.

A hyphen (-) is displayed if the microcode version is not supported. A hyphen (-) is also displayed for information about volumes with the vClone attribute.

R

Displays the status of the Read Only attribute of the snapshot data (VSP 5000 series, VSP One B20, VSP E series, and VSP G130, G/F350, G/F370, G/F700, G/F900 only).

- E: Read Only attribute is enabled.
- D: Read Only attribute is disabled.
- - (hyphen): This information is not available.

A hyphen (-) is displayed if the microcode version is not supported. A hyphen (-) is also displayed for information about volumes with the vClone attribute.

D_STAT

Displays the detailed status of each snapshot data. This status information might be changed or added in the future to make it more detailed.

The following status are displayed in this item. For the other status, the same information as Status is displayed. A hyphen (-) is displayed for information about volumes with the vClone attribute.

- SMPP: The pair is being deleted.
- CPPD: The pair is being created.
COPY is displayed if displaying CPPD is not supported on the storage system.
- CPRS: The pair is being resynchronized.
- PSUP: The pair is being suspended.
- CPYD: The differential data is in a copy process.
COPY is displayed if displaying CPYD is not supported on the storage system.
- CPYP: The snapshot data is being relocated.

RETENTION(H)

Displays the remaining time of the snapshot data retention period (VSP 5000 series, VSP One B20, VSP E series).

The snapshot data retention period (hour) is displayed by rounding up to the nearest whole number. A hyphen (-) is displayed for information about volumes with the vClone attribute.

PARENT LDEV#

Displays the LDEV number of a parent (a volume with the vClone Parent attribute, a root volume with the vClone Parent attribute, or a root volume) in a family to which a volume with the vClone attribute or a Thin Image Advanced pair belongs.

A hyphen (-) is displayed if all of the following conditions are met:

- HORCM_VCMD is defined in the configuration definition file. Alternatively, a command device that belongs to a virtual storage machine is defined in HORCM_CMD.
- The LDEV displayed in LDEV# and the LDEV displayed in this item belong to different virtual storage machines.

Displaying a list of LDEV numbers of volumes with the vClone Parent attribute.

```
# raidcom get snapshot -vclone
Serial#  LDEV#
800013    1
800013   10
```

Description of columns in output example:**Serial#**

Displays the serial numbers of the storage systems.

- If the target storage systems are VSP G1x00 and VSP F1500, the serial number added with 300,000 is displayed.
- If the target storage system is VSP 5000 series, the serial number added with 500,000 is displayed.

LDEV#

Displays a list of LDEV numbers of volumes with the vClone Parent attribute.

raidcom replace snapshot

Replaces the snapshot data that is mapped to the S-VOL.

For a Thin Image Advanced pair, the pair status changes to CPYP while the specified operation is in process, and then changes to PSUS when the processing is complete. To verify that the operation is complete, run the raidcom get snapshot command to verify that STAT has changed from CPYP to the status other than CPYP.

Syntax

```
raidcom replace snapshot -ldev_id <ldev#> {-snapshotgroup <name> | -mirror_id <mu#> }
```

Options and parameters**-ldev_id <ldev#>**

Specifies the LDEV number of the S-VOL to be replaced.

-snapshotgroup <name>

Specifies the name of the snapshot group in which the snapshot data as the operation target is included.

One MU which corresponds to the snapshot group which is specified by the P-VOL, corresponding to the specified S-VOL, is mapped.

**Note:**

Because the MU is selected automatically, an unexpected MU could be mapped. So if you intend to specify the specific snapshot data, specify the snapshot data by an MU number instead of a snapshot group.

-mirror_id <mu#>

Specifies the mirror ID of the specified snapshot data.

The specified snapshot data is mapped to the S-VOL.

Examples

Replacing the snapshot data of S-VOL (LDEV number 20:00) to the snapshot group snap3.

```
# raidcom replace snapshot -ldev_id 0x2000 -snapshotgroup snap3
```

raidcom add spm_wwn

Specifies the Server Priority Manager name for preferred/non-preferred WWNs.

Syntax

```
raidcom add spm_wwn -port <port#> -spm_name <nick_name> -hba_wwn <wwn_strings>
```

Options and parameters**-port <port#>**

Specifies the port number with the Target attribute, for example:

- CL1-A

-spm_name <nick_name>

Specifies the SPM name.

Up to 64 characters can be specified by CLI.

SPM names are managed uniquely in the entire system.

-hba_wwn <wwn_strings>

Specifies the monitored WWN.

Before setting the SPM name, the WWN must be registered as preferred or non-preferred.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

Examples

Specifies the SPM name (WWN_NICK_LINUX) to WWN (50060e8005fa0f36).

```
# raidcom add spm_wwn -port CL4-E -spm_name WWN_NICK_LINUX -hba_wwn 50060e80,05fa0f36
```

raidcom delete spm_wwn

Deletes WWN from the Server Priority Manager targets.

Syntax

```
raidcom delete spm_wwn -port <port#> [-hba_wwn <wwn_string> | -spm_name <nick_name>]
```

Options and parameters

-port <port#>

Specifies the port number whose attribute is Target. For example:

- CL1-A

-hba_wwn <wwn_string>

Specifies the WWN to be deleted.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-spm_name <nick_name>

Specifies the SPM name to be deleted.

Up to 64 characters can be specified by CLI.

Examples

Deletes the SPM name (WWN_NICK_LINUX) from the SPM targets.

```
# raidcom delete spm_wwn -port CL4-E -spm_name WWN_NICK_LINUX
```

Deletes the WWN (50060e8005fa0f36) from the SPM targets.

```
# raidcom delete spm_wwn -port CL4-E -hba_wwn 50060e80,05fa0f36
```

raidcom modify spm_wwn

Specifies the Server Priority Manager information to the Server Priority Manager target WWN.



Note: For VSP 5000 series, VSP G1x00, VSP F1500, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, this command cannot be executed if SPM information is already set for the storage system by the **raidcom modify spm_ldev** command.

Syntax

```
raidcom modify spm_wwn -port <port#> [-spm_priority <y/n>]
{-limit_io | -limit_kb | -limit_mb } <value>
{-hba_wwn <wwn_strings> | -spm_name <nick_name>}
```

Options and parameters

-port <port#>

Specifies the port number whose attribute is Target. For example:

- CL1-A

-spm_priority {y|n}

Specifies preferred/non-preferred WWN.

- y or yes: preferred WWN
- n or no: non-preferred WWN

{-limit_io | -limit_kb | -limit_mb } <value>

Specifies maximum value/threshold value by the I/O rate or the transmission rate.

- -limit_io: I/O rate, maximum value: 2,147,483,647 [IOPS]
- -limit_kb: transmission rate (units in kilobytes (KB)), maximum value: 2,147,483,647 [KB]
- -limit_mb: transmission rate (unit in megabytes (MB)), maximum value: 2,097,151 [MB]

(If the value is specified by MB, 1 MB is calculated as 1024 KB.)

The threshold value for the prioritized WWN is the same as the threshold value for the entire system.

-hba_wwn <wwn_strings>

Specifies WWN that you set the SPM information.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-spm_name <nick_name>

Specifies the SPM name.

Up to 64 characters can be specified by CLI.

Examples

Specifies WWN (50060e8005fa0f36) as non-preferred WWN, and specifies 5000/[I/O] (I/O rate) as the maximum value.

```
# raidcom modify spm_wwn -port CL4-E -spm_priority n -limit_io 5000 -hba_wwn 50060e80,05fa0f36
```

Specifies the SPM name (WWN_NICK_LINUX) as non-preferred WWN, and specifies 500[MB/s] (transmission rate in megabytes per second) as the maximum value.

```
# raidcom modify spm_wwn -port CL4-E -spm_priority n -limit_mb 500 -spm_name WWN_NICK_LINUX
```

Specifies WWN (50060e8005fa0f36) as preferred WWN, and specifies 5000/[I/O] (I/O rate) as the threshold value.

```
# raidcom modify spm_wwn -port CL4-E -spm_priority y -limit_io 5000 -hba_wwn 50060e80,05fa0f36
```

Specifies the SPM name (WWN_NICK_LINUX) as preferred WWN, and specifies 500000[KB/s] (transmission rate) as the maximum value.

```
# raidcom modify spm_wwn -port CL4-E -spm_priority y -limit_kb 500000 -spm_name WWN_NICK_LINUX
```

Deletes the SPM information of WWN (50060e8005fa0f36).

```
# raidcom modify spm_wwn -port CL4-E -hba_wwn 50060e80,05fa0f36
```

Deletes the SPM information of the SPM name (WWN_NICK_LINUX).

```
# raidcom modify spm_wwn -port CL4-E -spm_name WWN_NICK_LINUX
```

raidcom get spm_wwn

Gets the Server Priority Manager information of the Server Priority Manager target WWN.

Syntax

```
raidcom get spm_wwn -port <port#> [-hba_wwn <wwn_strings>
| -spm_name <nick_name>
```

Options and parameters

-port <port#>

Specifies the port number whose attribute is Target. For example:

- CL1-A

-hba_wwn <wwn_strings>

Specifies WWN that you get the SPM information.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-spm_name <nick_name>

Specifies the SPM name that you get the SPM information.

Up to 64 characters can be specified by CLI.

Examples

Gets the SPM information of WWN that belongs to the specified port (CL4-E).

```
#raidcom get spm_wwn -port CL4-E
PORT  SPM_MD          SPM_WWN NICK_NAME  GRP_NAME  Serial#
CL4-E  WWN 210000e08b0256f8 WWN_NICK_LINUX_0 OLA_NODE0_CTL 63528
CL4-E  WWN 210000e08b0256f7 WWN_NICK_LINUX_1 OLA_NODE0_CTL 63528
```

Specifies WWN (50060e8005fa0f36) to get the SPM information.

```
# raidcom get spm_wwn -port CL4-E -hba_wwn 50060e80,05fa0f36
```

```
PORT  SPM_MD PRI IOps KBps Serial#
CL4-E WWN      Y 5000   - 63528
```

Specifies the SPM name (WWN_NICK_LINUX) to get the SPM information.

```
# raidcom get spm_wwn -port CL4-E -spm_name WWN_NICK_LINUX
```

```
PORT  SPM_MD PRI IOps KBps Serial#
CL4-E WWN      Y    - 5000 63528
```

Description of each column in output example:**PORT**

Displays the port to which the WWN is set.

SPM_MD

Displays the Control mode of the SPM.

- WWN: WWN control
- PORT: Port control

SPM_WWN

Displays the set WWN.

NICK_NAME

Displays the SPM name set to the WWN. If the SPM name is not set, a hyphen (-) is displayed.

GRP_NAME

Displays the SPM group name to which the WWN belongs. If the WWN does not belong to the group, a hyphen (-) is displayed.

Serial#

Displays the Seq#.

PRI

Displays the settings of the priority.

- Y: Preferred
- N: Non-preferred

IOps

If the PRI setting is Y, it displays the threshold. If the PRI setting is N, it displays the maximum value. If the value is set by the transferred rate in kilobytes (KBps), it displays a hyphen (-).

KBps

If the PRI setting is Y, it displays the threshold. If the PRI setting is N, it displays the maximum value. If the value is set by the IO rate (IOPS), it displays a hyphen (-). If the value is set by megabytes (MB), 1 MB is calculated as 1024 kilobytes (KB), and displays the value by KB.

raidcom monitor spm_wwn

Gets the monitoring information of Server Priority Manager target WWN.

Syntax

```
raidcom monitor spm_wwn {-hba_wwn <wwn_strings> | -spm_name <nick_name>}
```


Options and parameters

-hba_wwn <wwn_strings>

Specifies WWN that you get the monitoring information.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-spm_name <nick_name>

Specifies the SPM name that you get the monitoring information.

Up to 64 characters can be specified by CLI.

Examples

Specifies WWN (50060e8005fa0f36) to get the monitoring information.

```
# raidcom monitor spm_wwn -hba_wwn 50060e80,05fa0f36
```

```
PORT  SPM_MD IOps      KBps Serial#
CL4-E PORT  5000 5000000 63528
```

Specifies the SPM name (WWN_NICK_LINUX) to get the monitoring information.

```
# raidcom monitor spm_wwn -spm_name WWN_NICK_LINUX
```

```
PORT  SPM_MD IOps      KBps Serial#
CL4-E PORT  5000 5000000 63528
```

Description of each column in output example:

PORT

Displays the port to which the WWN is set.

SPM_MD

Displays the Control mode of the SPM.

- WWN: WWN control
- PORT: Port control

IOps

Displays the current IO rate (IOPS) of the specified WWN of or the specified SPM name.

KBps

Displays the current transferred rate in kilobytes per second (KBps) of the specified WWN or the specified SPM name.

Serial#

Displays the Seq#.

raidcom add spm_group

Registers Server Priority Manager target WWN to Server Priority Manager group. For VSP 5000 series, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, VSP G200, G400, G600, G800 and VSP F400, F600, F800, this command cannot be executed if SPM information is already set for the storage system by the **raidcom modify spm_ldev** command.

Syntax

For registering the WWN to the SPM group:

```
raidcom add spm_group -port <port#> -spm_group <group_name>
    {<nick_name> | -hba_wnn <wnn_strings>}
```

(VSP and VSP G1x00 and VSP F1500 only) For associating the host group with the SPM group:

```
raidcom add spm_group -spm_group <group name> -port <port#>
    [<host group name>] -spm_host_grp
```

Options and parameters**-port <port#>**

Specifies the port number whose attribute is Target. For example:

- CL1-A

-spm_group <group_name> [<nick_name>]

Supported storage systems:

- VSP 5000 series
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900
- VSP G200, G400, G600, G800 and VSP F400, F600, F800

Specifies the SPM group name and the SPM name.

Up to 64 characters can be specified by CLI. If <nick_name> is omitted, -hba_wnn is required. The SPM group name must be unique across the system.

-spm_group <group_name>

Specifies the SPM group name.

Up to 64 characters can be specified by CLI.

<nick_name>

Specifies the SPM target WWN using the SPM name.

Up to 64 characters can be specified by CLI.

-hba_wwn <wwn_strings>

Specifies SPM target WWN.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-port <port#> [<host group name>] -spm_host_grp

VSP and VSP G1x00 and VSP F1500 only. When registering the WWN registered in the host group to the SPM group, specifies the port number and the host group ID, or the host group name. If you register the WWN by this method, the host group is associated with the SPM group. When adding or deleting the WWN to or from the host group, the WWN is also added or deleted to or from the SPM group. Also, if you delete a host group, all WWNs registered in the host group are deleted from the SPM group.

Examples:

- CL1-A-g, where g is a value from 0 to 254.
- CL1-A Linux_X86

Examples

Registers WWN (50060e8005fa0f36) as the SPM group name (WWN_GRP_LINUX).

```
# raidcom add spm_group -port CL4-E -spm_group WWN_GRP_LINUX
-hba_wwn 50060e80,05fa0f36
```

Registers the SPM name (WWN_NICK_LINUX) as the SPM group name (WWN_GRP_LINUX).

```
# raidcom add spm_group -port CL4-E -spm_group WWN_GRP_LINUX WWN_NICK_LINUX
```

(VSP and VSP G1x00 and VSP F1500 only) Registers the host group (CL4-E-1) as the SPM group name (WWN_GRP_LINUX).

```
# raidcom add spm_group -port CL4-E-1 -spm_host_grp
-spm_group WWN_GRP_LINUX
```

raidcom delete spm_group

Deletes the WWN of the specified port or the host group from the Server Priority Manager group to delete the SPM information of the WWN. If no WWN is registered in the SPM group, the SPM group itself is also deleted.

Syntax

```
raidcom delete spm_group -port <port#> -spm_group <group_name>
```

VSP and VSP G1x00 and VSP F1500 only:

```
raidcom delete spm_group {-port <port#> -spm_group <group_name> |  
-port <port#> [<host group name>] -spm_host_grp}
```

Options and parameters

-port <port#>

Specifies the port number whose attribute is Target. For example:

- CL1-A

-spm_group <group_name>

Specifies the SPM group name.

Up to 64 characters can be specified by CLI.

-port <port#> [<host group name>] -spm_host_grp

VSP and VSP G1x00 and VSP F1500 only. When setting SPM information for the SPM group associated with the host group, specifies the port number and the host group ID, or the host group name.

Examples:

- CL1-A-g, where g is a value from 0 to 254.
- CL1-A Linux_X86

Examples

Deletes the SPM group (WWN_GRP_LINUX).

```
# raidcom delete spm_group -port CL4-E -spm_group WWN_GRP_LINUX
```

(VSP and VSP G1x00 and VSP F1500 only) Deletes SPM associated with the host group (CL4-E-1).

```
# raidcom delete spm_group -port CL4-E-1 -spm_host_grp
```

raidcom modify spm_group

Specifies the Server Priority Manager information to the Server Priority Manager target group.

Syntax

```
raidcom modify spm_group -port <port#> -spm_group <group_name> [-spm_priority <y/n>]
    {-limit_io | -limit_kb | -limit_mb } <value>
```

VSP and VSP G1x00 and VSP F1500 only:

```
raidcom modify spm_group {-port <port#> -spm_group <group_name> |
    -port <port#> [<host group name>] -spm_host_grp} [-spm_priority <y/n>]
    {-limit_io | -limit_kb | -limit_mb } <value>
```

Options and parameters

-port <port#>

Specifies the port number whose attribute is Target. For example:

- CL1-A

-port <port#> [<host group name>] -spm_host_grp

VSP and VSP G1x00 and VSP F1500 only. When you specify SPM information for the SPM group associated with the host group, specify the port number, host group ID, or host group name.

Examples:

- CL1-A-g, where g is a value from 0 to 254.
- CL1-A Linux_X86

-spm_priority {y|n}

Specifies preferred/non-preferred WWN.

- y or yes: preferred WWN
- n or no: non-preferred WWN

{-limit_io | -limit_kb | -limit_mb } <value>

Specifies maximum value/threshold value by the I/O rate or the transmission rate.

- -limit_io: I/O rate, maximum value: 2,147,483,647 [IOPS]
- -limit_kb: transmission rate (units in kilobytes (KB)), maximum value: 2,147,483,647 [KB]
- -limit_mb: transmission rate (units in megabytes (MB)), maximum value: 2,097,151 [MB]

If you specify the value per MB, 1 MB is calculated as 1024 KB.

The threshold value for the prioritized WWN is the same as the threshold value for the entire system.

-spm_group <group_name>

Specifies the SPM group name that you set the SPM information.

Up to 64 characters can be specified by CLI.

Examples

Specifies the SPM group name (WWN_GRP_LINUX) as non-preferred WWN, and specifies 5000[I/O] (I/O rate) as the maximum value.

```
# raidcom modify spm_group -port CL4-E -spm_priority n -limit_io 5000 -spm_group
WWN_GRP_LINUX
```

Specifies the SPM group name (WWN_GRP_LINUX) as non-preferred WWN, and specifies 500[MB/s] (transmission rate in megabytes per second) as the maximum value.

```
# raidcom modify spm_group -port CL4-E -spm_priority n -limit_mb 500 -spm_group
WWN_GRP_LINUX
```

Specifies the SPM group name (WWN_GRP_LINUX) as preferred WWN, and specifies 5000[I/O] (I/O rate) as the threshold value.

```
# raidcom modify spm_group -port CL4-E -spm_priority y -limit_io 5000 -spm_group
WWN_GRP_LINUX
```

Specifies the SPM group name (WWN_GRP_LINUX) as preferred WWN, and specifies 500000[KB/s] (transmission rate) as the threshold value.

```
# raidcom modify spm_group -port CL4-E -spm_priority y -limit_kb 500000 -spm_group
WWN_GRP_LINUX
```

Deletes the SPM information of the SPM group name (WWN_GRP_LINUX).

```
# raidcom modify spm_group -port CL4-E -spm_group WWN_GRP_LINUX
```

(VSP and VSP G1x00 and VSP F1500 only) Specifies the SPM group associated with the host group (CL4-E-1) as non-preferred WWN, and specifies 5000[I/O] (I/O rate) as the maximum value.

```
# raidcom modify spm_group -port CL4-E-1 -spm_host_grp -spm_priority n
-limit_io 5000
```

raidcom get spm_group

Gets the Server Priority Manager information of the Server Priority Manager target WWN in the specified port by the Server Priority Manager group unit.

Syntax

```
raidcom get spm_group -port <port#> -spm_group <group_name>
```

VSP and VSP G1x00 and VSP F1500 only:

```
raidcom get spm_group {-port <port#> -spm_group <group_name> |  
-port <port#> [<host group name>] -spm_host_grp}
```

Options and parameters

-port <port#>

Specifies the port number whose attribute is Target. For example:

- CL1-A

-spm_group <group_name>

Specifies the SPM group name.

Up to 64 characters can be specified by CLI.

-port <port#> [<host group name>] -spm_host_grp

VSP and VSP G1x00 and VSP F1500 only. When obtaining SPM information of the SPM group associated with the host group, specifies the port number and the host group ID, or the host group name.

Examples:

- CL1-A-g, where g is a value from 0 to 254.
- CL1-A Linux_X86

Examples

Specifies the port (CL4-E) and the SPM group name (WWN_GRP_LINUX) to get the SPM information.

```
# raidcom get spm_group -port CL4-E -spm_group WWN_GRP_LINUX
```

```
PORT  SPM_MD PRI  IOps KBps Serial#  
CL4-E PORT      Y 5000   - 63528
```

(VSP and VSP G1x00 and VSP F1500 only) Specifies the host group (CL4-E-1) to obtain SPM information.

```
# raidcom get spm_group -port CL4-E-1 -spm_host_grp
```

```
PORT  SPM_MD PRI      IOps KBps  Serial#  
CL4-E  WWN    Y 1000000000    -    63528
```

Description of each column in output example:**PORT**

Displays the port to which the WWN is set.

SPM_MD

Displays the Control mode of the SPM.

- WWN: WWN control
- PORT: Port control

PRI

Displays the settings of the priority.

- Y: Preferred
- N: Non-preferred

IOps

If the PRI setting is Y, it displays the threshold. If the PRI setting is N, it displays the maximum value. If the value is set by the transferred rate in kilobytes per second (KBps), it displays a hyphen (-).

KBps

If the PRI setting is Y, it displays the threshold. If the PRI setting is N, it displays the maximum value. If the value is set by the IO rate (IOPS), it displays a hyphen (-). If the value is set by megabytes (MB), 1 MB is calculated as 1024 kilobytes (KB), and displays the value by KB.

Serial#

Displays the Seq#.

raidcom monitor spm_group

Gets the monitoring information of Server Priority Manager target WWN by the Server Priority Manager group unit.

Syntax

```
raidcom monitor spm_group -spm_group <group_name>
```

Options and parameters**-spm_group <group_name>**

Specifies the SPM group name.

Up to 64 characters can be specified by CLI.

Examples

Specifies the SPM group name (WWN_GRP_LINUX) to get the monitoring information.

```
# raidcom monitor spm_group -spm_group WWN_GRP_LINUX
```

```
PORT  SPM_MD IOps KBps      Serial#
CL4-E PORT    5000 5000000 63528
```

Description of each column in output example:

PORT

Displays the port to which the WWN is set.

SPM_MD

Displays the Control mode of the SPM.

- WWN: WWN control
- PORT: Port control

IOps

Displays the current IO rate (IOPS) of the specified WWN or the specified SPM name.

KBps

Displays the current transferred rate (KBps) of the specified WWN or the specified SPM name.

Serial#

Displays the Seq#.

raidcom modify spm_ldev

Sets the Server Priority Manager information for a combination of an LDEV and WWN, or an LDEV and iSCSI name.



Note:

- This command cannot be executed if SPM information is already set for the storage system by the `raidcom modify spm_wwn` or `raidcom add spm_group` command.
- When the number of LDEVs with SPM configured exceeds 4,096, I/O response for the 4,097th and subsequent LDEVs is slower than for the other 4,096 LDEVs. To improve host I/O response, delete the existing SPM information so that the total number of LDEVs with SPM configured will be 4,096 or fewer.

Syntax

```
raidcom modify spm_ldev -ldev_id
    <ldev#> {-hba_wwn <wwn strings>
    | -hba_iscsi_name <initiator iscsi name>}
    [-spm_priority {y/n}] [{-limit_io | -limit_mb} <value>]
```

Options and parameters

-ldev_id <ldev#>

Specifies the LDEV number (0 to 65279) whose SPM information you want to set.

Example: -ldev_id 200

-hba_wwn <wwn_strings>

Specifies the WWN of the host bus adapter (initiator) for which you set the SPM information.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator) for which you set the SPM information. A maximum of 223 characters can be specified.

Example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

-spm_priority {y|n}

Specifies the priority (prioritized or non-prioritized) for the combination of an LDEV and WWN, or an LDEV and iSCSI name.

- y or yes: prioritized. Obtains monitoring information, and the upper limit is not set. This combination is prioritized.
- n or no: non-prioritized. Obtains monitoring information, but the upper limit is set. This combination is not prioritized.

Non-prioritized is set if omitted. If y is set, only monitoring is performed.

For details about operations for setting SPM with LDEVs and WWNs, or LDEVs and iSCSI names, see the *Command Control Interface User and Reference Guide*.

{-limit_io | -limit_mb } <value>

Specifies the upper limit for the I/O rate or the transmission rate. If omitted, the I/O rate or the transmission rate is not suppressed by the upper limit.

- -limit_io: I/O rate
Maximum value: 65,535 [IOPS]
- -limit_mb: Transmission rate in megabytes (MB)
Maximum value: 31 [MB]

Examples

The following shows an example of setting n (non-prioritized) for the combination of LDEV: 1024 and WWN: 50060e8005fa0f36, and I/O rate: 5000 [IOPS] as the upper limit:

```
# raidcom modify spm_ldev -ldev_id 1024 -hba_wnn 50060e80,05fa0f36 -spm_priority n -
limit_io 5000
```

The following shows an example of setting n (non-prioritized) for the combination of LDEV: 1024 and iSCSI name: iqn.z1, and I/O rate: 5000 [IOPS] as the upper limit:

```
# raidcom modify spm_ldev -ldev_id 1024 -hba_iscsi_name iqn.z1 - spm_priority n -
limit_io 5000
```

raidcom delete spm_ldev

Deletes a combination of the SPM target LDEV and WWN, or the SPM target LDEV and iSCSI name from the SPM target.

Syntax

```
raidcom delete spm_ldev -ldev_id <ldev#> {-hba_wnn <wnn strings>
| - hba_iscsi_name <initiator iscsi name>}
```

Options and parameters**-ldev_id <ldev#>**

Specifies the LDEV number (0 to 65279) whose SPM information you want to delete.

Example: -ldev_id 200

-hba_wnn <wnn_strings>

Specifies the WWN of the host bus adapter (initiator) whose SPM information you want to delete.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator) whose SPM information you want to delete. A maximum of 223 characters can be specified.

Example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

Examples

The following shows an example of deleting the combination of LDEV: 1024 and WWN: 50060e8005fa0f36 from the SPM target:

```
# raidcom delete spm_ldev -ldev_id 1024 -hba_wwn 50060e80,05fa0f36
```

The following shows an example of deleting the combination of LDEV: 1024 and iSCSI name: iqn.z1 from the SPM target:

```
# raidcom delete spm_ldev -ldev_id 1024 -hba_iscsi_name iqn.z1
```

raidcom monitor spm_ldev

Obtains the monitoring information for the combination of an LDEV and WWN, or an LDEV and iSCSI name. This command displays information when both the following conditions are met:

- SPM information is set for the specified LDEV.
- The current I/O rate or transfer rate for the LDEV with the specified WWN or iSCSI name is not 0.

Syntax

```
raidcom monitor spm_ldev -ldev_id <ldev#> {-hba_wwn <wwn strings>
| - hba_iscsi_name <initiator iscsi name>}
```

Options and parameters**-ldev_id <ldev#>**

Specifies the LDEV number (0 to 65279) whose monitoring information you want to obtain.

Example: -ldev_id 200

-hba_wwn <wwn_strings>

Specifies the WWN of the host bus adapter (initiator) whose monitoring information you want to obtain.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator) whose monitoring information you want to obtain. A maximum of 223 characters can be specified.

Example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

Examples

The following shows an example of obtaining monitoring information for the combination of LDEV: 1024 and WWN: 50060e8005fa0f36:

```
# raidcom monitor spm_ldev -ldev_id 1024 -hba_wwn 50060e80,05fa0f36
```

Serial#	LDEV	IOps	KBps	WWN
63528	1024	5000	5000000	50060e8005fa0f36

The following shows an example of obtaining monitoring information for the combination of LDEV: 1024 and iSCSI name: iqn.z1:

```
# raidcom monitor spm_ldev -ldev_id 1024 -hba_iscsi_name iqn.z1
```

Serial#	LDEV	IOps	KBps	IQN
63528	1024	5000	5000000	iqn.z1

Description of each column in output example:

LDEV

Displays the LDEV number.

IOps

Displays the current I/O rate (IOPS) of the LDEV identified by the specified WWN or iSCSI name.

KBps

Displays the current transfer rate (KBps) for the LDEV identified by the specified WWN or iSCSI name.

WWN

Displays the WWN of the host bus adapter (initiator).

IQN

Displays the iSCSI name of the host bus adapter (initiator).

raidcom get spm_ldev

Displays the Server Priority Manager information for a combination of an LDEV and WWN or an LDEV and iSCSI name.

Syntax

```
raidcom get spm_ldev -ldev_id <ldev#> {-hba_wwn <wwn strings>
| -hba_iscsi_name <initiator iscsi name>}
```

Options and parameters

-ldev_id <ldev#>

Specifies the LDEV number (0 to 65279) whose SPM information you want to obtain.

Example: -ldev_id 200

This command terminates without displaying information if SPM information is not set for the combination of an WWN and the specified LDEV, or an iSCSI name and the specified LDEV.

-hba_wwn <wwn_strings> |

Specifies the WWN of the host bus adapter (initiator) whose SPM information you want to obtain.

The WWN is specified as a 16-digit hexadecimal value. The 17th digit and later are ignored. The value can be split in units of 4 bytes by "," (comma). For example:

- 210000e08b0256f8
- 210000e0,8b0256f8

This command terminates without displaying information if SPM information is not set for the combination of an LDEV and the specified WWN.

-hba_iscsi_name <initiator iscsi name>

Specifies the iSCSI name of the host bus adapter (initiator) whose SPM information you want to obtain. A maximum of 223 characters can be specified.

Example:

- iqn.win2k8.example.of.iqn.form
- eui.0123456789ABCDEF

This command terminates without displaying information if SPM information is not set for the combination of an LDEV and the specified iSCSI name.

Examples

The following shows an example of displaying SPM information for the combination of LDEV: 1024 and WWN, and LDEV: 1024 and iSCSI name:

```
# raidcom get spm_ldev -ldev_id 1024
```

Serial#	LDEV	PRI	IOps	KBps	T	Name
63528	1024	Y	-	-	W	50060e8005fa0f3
63528	1024	N	-	31744	I	iqn.z1

The following shows an example of displaying SPM information for the combination of an LDEV and iSCSI name: iqn.z1:

```
# raidcom get spm_ldev -hba_iscsi_name iqn.z1
```

Serial#	LDEV	PRI	IOps	KBps	T	Name
---------	------	-----	------	------	---	------

63528	1024	Y	-	-	I	iqn.z1
63528	1025	N	-	31744	I	iqn.z1

The following shows an example of displaying SPM information for the combination of an LDEV and WWN: 50060e8005fa0f36:

```
# raidcom get spm_ldev -hba_wwn 50060e8005fa0f3
```

Serial#	LDEV	PRI	IOps	KBps	T	Name
63528	1024	Y	-	-	W	50060e8005fa0f3
63528	1025	N	-	31744	W	50060e8005fa0f3

The following shows an example of displaying SPM information for the combination of an LDEV and WWN, and an LDEV and iSCSI name:

```
# raidcom get spm_ldev
```

Serial#	LDEV	PRI	IOps	KBps	T	Name
63528	1024	Y	-	-	W	50060e8005fa0f3
63528	1024	N	-	31744	I	iqn.z1
63528	1025	N	5000	-	I	iqn.z1

Description of each column in output example:

Serial#

Displays the product serial number.

LDEV

Displays the LDEV number.

PRI

Displays the priority.

- Y: Prioritized
- N: Non-prioritized

IOps

For the non-prioritized setting, displays the upper limit. For the prioritized setting, displays a hyphen (-) if the transfer rate is set.

KBps

Displays the upper limit for the non-prioritized setting. For the prioritized setting, displays a hyphen (-) if the I/O rate is set. In addition, if the value is set in megabytes, it is converted to kilobytes by using the following formula: 1 MB = 1024 KB

T

Displays the type of name to be displayed in the Name column:

- W: WWN
- I: iSCSI name

Name

Displays the WWN or iSCSI name.

raidcom add dp_pool

Create a pool for Dynamic Provisioning or Dynamic Provisioning for Mainframe by the specified resource. Or, add pool volumes to the pools for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe. A parity group, an LDEV, and a device group can be specified as a resource.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

If a cache memory, HIE, or a logical path inside the storage is blocked, the operation cannot be performed.



Tip: For VSP One B20, VSP E series, VSP G130, G/F350, G/F370, G/F700, G/F900, specifying a parity group is recommended. When a parity group is specified, an LDEV is created and a pool volume is added. It is not necessary to create an LDEV in advance. In addition, you cannot add an existing LDEV as a pool volume.

When specifying a pool that is already created for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe, the specified resource is added to the pool.

Rebalancing is performed automatically if the pool capacity is expanded. However, some metadata of the capacity saving function cannot be moved by rebalancing. Therefore, the loads of the I/O processing, the asynchronous deduplication processing, and garbage collection on V-VOLs with capacity saving enabled that have been defined before the pool capacity expansion might not be evenly distributed among the parity groups in the pool.

Specify either one of Pool ID or Pool Name certainly. If both the Pool ID and Pool Name options are omitted, this command is rejected with EX_REQARG.

This command is executed asynchronously with the command input. Check the completion of this process on the **raidcom get command_status** command.

Syntax for adding an LDEV belonging to a parity group whose attribute is not DDP, or creating and adding the LDEV

```
raidcom add dp_pool {-pool_id <pool ID#> | -pool_name <pool
    naming> | -pool_id <pool ID#>
    -pool_name <pool naming>}}
{-parity_grp_id <gnosgno> [-resource_id <resource group id >]
| -ldev_id <ldev#> ...[-cnt <count>] | -grp_opt <group option>
-device_grp_name <device group name> [<device name>]}
[ -user_threshold <threshold_1>[<threshold_2>] ]
[-suspend_tipair {yes | no}]
[-auto_add_poolvol {enable|disable}]
```

Syntax for adding an LDEV belonging to a DDP parity group, or creating and adding the LDEV

```
raidcom add dp_pool {-pool_id <pool ID#> | -pool_name <pool naming> | -pool_id <pool
ID#> -pool_name <pool naming>} {-parity_grp_id <gno-sgno> [-resource_id <resource
group_id>] | -parity_grp_id <gno-sgno> -ddp_ldev -request_id auto}{ -user_threshold
<threshold_1> [<threshold_2>] ] [-suspend_tipair {yes | no}]
```

Options and parameters

-pool_id <pool ID#>

Specifies the Pool ID (0-127) for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe. When the specification of Pool ID is omitted, a Pool Name must be specified.

If a **-pool_id** option is non-numeric, the specified value is recognized as a pool name to identify the pool ID.

When you omit specifying **-pool_id** option, you need to specify **-pool_name** option.

When specifying the `-pool_name` option without a `-pool_id` option, the specified pool is searched first. If the pool exists, the pool volume is added to the corresponding pool. If the pool does not exist, a new pool is created and is assigned a pool ID. However, if the pool does not exist and multiple commands without the `-pool_id` option are executed simultaneously, the same pool ID might be assigned to multiple newly created pools.

To avoid this behavior, use the `raidcom lock resource` command in advance to lock all resource groups. If all the resource groups are locked, commands executed by other users cannot take out unused pool IDs from all the resource groups. If you create multiple pools, execute the next `raidcom add snap_pool` command after the execution of the current `raidcom add snap_pool` command is complete.

After the execution of all the `raidcom add snap_pool` commands are complete, use the `raidcom unlock resource` command to unlock all the resource groups.

-pool_name <pool naming>

Specifies a pool name of a pool for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe. Up to 32 characters can be specified.

When the specification of Pool Name is omitted, a Pool ID must be specified. When the Pool ID is specified and the Pool Name and `-device_grp_name` option are omitted, a pool name is allocated automatically in the form of "New Pool<number>".

The `-pool_name` option cannot be specified by configuring only in numeric because numeric values are preferentially identified as a pool ID. Specifies a pool ID with the `-pool_id<pool ID#>` option.

The existing pool name is changed if all of the following conditions are met:

- The pool ID is specified.
- The pool name is specified.
- A pool with the specified pool ID already exists.
- The `-parity_grp_id` option is not specified.

-parity_grp_id <gno-sgno> [-resource_id <resource group id >]

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the parity group number (gno: 1 to 52, sgno: 1 to 32).

If this option is specified, an LDEV is automatically created in the specified parity group, and then a pool is created or expanded on the created LDEV.

For example: 3-1

When `-resource_id <resource grp id>` is specified, the LDEV which is not installed in the specified resource group and has the largest LDEV number should be selected to create a pool volume. This option is valid only when a pool is created.

When `-resource_id <resource grp id>` is not specified, the largest LDEV number in the resource group for which the user has operation permission should be selected to create a pool volume.

When expanding a pool, the largest LDEV number in the resource group to which an existing pool volume belongs should be selected to create the pool volume.

The command cannot be run if there are not enough unimplemented LDEVs to create the pool volume in the target resource group. In that case, use the **raidcom add resource -ldev_id <ldev#>** command or the **raidcom delete resource -ldev_id <ldev#>** command to move the unimplemented LDEV numbers to the target resource group, or use the **raidcom delete ldev** command to delete unnecessary LDEVs.

-parity_grp_id <gno-sgno> -ddp_ldev -request_id auto

Specifies the parity group number.

Specify this option to create and expand a pool containing a DDP parity group with an LDEV created. To create and expand a pool containing the DDP parity group with no LDEVs created, specify the `-parity_grp_id <gno-sgno> [-resource_id <resource group id >]` option listed above. For example:

3-1

`-request_id` is a request ID assigned each time the command is run. For details, see [Request ID function \(on page 293\)](#).

-ldev_id <ldev#> ...

Specifies the LDEV number (0 to 65279). Up to 64 of LDEVs can be specified at a time. For example:

- `-ldev_id 100`
- `-ldev_id 100 - 110`
- `-ldev_id 100 -cnt 10`

[-cnt <count>]

Specifies the count (2 to 64).

If this specification is omitted, "1" is used.

Up to 64 of LDEVs can be specified at a time.

-grp_opt <group option>

Specifies the device information extracted from the LDEV belonging to the device group. Specify "ldev" (fixed). The information of LDEV belonging to the device group is used.

-device_grp_name <device group name> [<device name>]

Specifies the name of device group (maximum 32 characters) to be operated.

To operate an LDEV in the device group, specify the device name of LDEV (maximum 32 characters) within the device group.

If you specify the target LDEV by the device name, all devices whose names begin with the specified string are subject to the operation.

If the device name is omitted, all the LDEVs belonging in the device group are operated.

When the -pool_name option is omitted, the device group name changes into the pool name.

[-user_threshold <threshold_1> [<threshold_2>]]

You may specify two user-defined thresholds, and their valid ranges are 1-100%, respectively. If you specify <threshold_1> and <threshold_2>:

the value of <threshold_1> is set as the threshold for WARNING specified to a pool.

the value of <threshold_2> is set as the threshold for High water mark specified to a pool.

If you specify only <threshold_1>, your specified value and the system default value (80%) are applied. If you omit to specify the value, 70% and 80% are applied, automatically.

When you add a pool volume, this option is ignored even if it is specified. If you want to change the user defined threshold value of the additional pool volume, execute the raidcom modify pool command.

[-suspend_tipair {yes | no}]

Specifies whether to suspend Thin Image pairs when the threshold for high water mark is exceeded. This option is valid only for creating pairs.

- yes: Thin Image pairs are suspended.
- no: Thin Image pairs are not suspended.

If this option is omitted , "yes" is set.

[-auto_add_poolvol {enable|disable}]

Specifies whether to automatically add pool volumes for which accelerated compression is enabled. This option is valid only for creating pools. This option can be specified for VSP G1x00 or VSP F1500.

- enable: Automatically adds pool volumes for which accelerated compression is enabled according to the compression ratio of a parity group for which accelerated compression is enabled.
- disable: The pool volumes for which accelerated compression is enabled are not added automatically.

When these options are not specified, the default values will be applied as follows:

- enable: When the storage system is VSP 5000 series, VSP One B20, VSP E series or VSP G/F350, G/F370, G/F700, G/F900.
- disable: When the storage system is VSP G1x00 or VSP F1500.

Examples

By using LDEVs:400, 401, and 402, creating a pool of Pool ID:1, Pool Name: my_pool for Dynamic Provisioning/Dynamic Provisioning for Mainframe.

```
# raidcom add dp_pool -pool_id 1 -pool_name my_pool -ldev_id 400 401 402
```

By using LDEVs:500, 501, and 502, creating a pool of Pool ID: Allocated automatically, Pool Name: my_pool for Dynamic Provisioning/Dynamic Provisioning for Mainframe.

```
# raidcom add dp_pool -pool_name my_pool -ldev_id 500 501 502
```

By using LDEVs:600, 601, and 602, creating a pool of Pool ID: 2, Pool Name: Allocated automatically for Dynamic Provisioning/Dynamic Provisioning for Mainframe.

```
# raidcom add dp_pool -pool_id 2 -ldev_id 600 601 602
```

By using LDEVs: 700, 701, and 702, creating a pool of Pool ID: 3, Pool Name: my_pool for Dynamic Provisioning, Dynamic Provisioning for Mainframe, with the user-defined thresholds WARNING specified to 70% and High water mark to 80%.

```
# raidcom add dp_pool -pool_id 3 -pool_name my_pool
-ldev_id 700 701 702 -user_threshold 70 80
```

Add LDEV: 368 to the pool ID: 10 for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe.

```
# raidcom add dp_pool -pool_id 10 -ldev_id 368
```

By using LDEV belonging to the device group: grp1, creating a pool of Pool ID:1, pool name: Allocated automatically for Dynamic Provisioning/Dynamic Provisioning for Mainframe.

```
# raidcom add dp_pool -pool_id 1 -grp_opt ldev -device_grp_name grp1
```

By using an LDEV belonging to a DDP parity group: 1-2, creating a pool for Dynamic Provisioning with pool ID: 0 and with pool name automatically assigned.

```
# raidcom add dp_pool -pool_id 0 -parity_grp_id 1-2 -ddp_ldev -
request_id auto
```

By using LDEVs: 800, 801, and 802, creating a pool of Pool ID: 4, Pool Name: my_pool for Dynamic Provisioning, with the user-defined thresholds WARNING specified to 70% and High water mark to 80%, and Thin Image pairs are suspended when the high water mark threshold is exceeded.

```
# raidcom add dp_pool -pool_id 4 -pool_name my_pool -ldev_id 800 801 802 -
user_threshold 70 80 -suspend_tipair yes
```

Using an LDEV ID of resource group ID: 0, creating an LDEV for parity group: 1-1, and then creating a pool of Pool ID: 1, Pool Name: Allocated automatically for Dynamic Provisioning or Dynamic Provisioning for Mainframe.

```
# raidcom add dp_pool -pool_id 1 -parity_grp_id 1-1 -resource_id 0
```

Using LDEVs: 500, 501, and 502, creating a pool of Pool ID: Allocated automatically and Pool Name: my_pool for Dynamic Provisioning for which automatic addition settings of pool volumes whose accelerated compression is enabled: enabled.

```
# raidcom add dp_pool -pool_name my_pool -ldev_id 500 501 502 -auto_add_poolvol enable
```

raidcom get dp_pool

Displays pool information for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, and active flash for mainframe.

Syntax

```
raidcom get dp_pool [-key <keyword> [-fx]]
                    [-pool {<pool_id> | <pool naming>} ]] [-pcap]
```

Options and parameters

[-key <keyword> [-fx]]

Specifies the display keyword. Specify opt as <keyword>. If you specify opt, information about Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe pool is displayed.

If you specify the -fx option, TL_RANGE and TD_RANGE are displayed in hexadecimal notation.

[-pool {<pool_id> | <pool naming>}]

Specifies the pool ID or the pool name of Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe when you want to display the range for each tiering level of Dynamic Tiering, Dynamic Tiering for Mainframe, active flash, or active flash for mainframe pools. When you specify only digits, it is recognized as a pool ID. When the pool name consists of digits, specify the pool ID instead of the pool name. This option is effective only when the -key option is specified.

[-pcap]

Supported storage systems:

- VSP 5000 series
- VSP G1x00 and VSP F1500
- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Specifies the actual capacity. This option is valid only when the -key opt option is specified. For the pool for which automatic addition settings of pool volumes whose accelerated compression is enabled are enabled, check the pool capacity by specifying this option.

If the microcode version does not support this option, nothing is displayed when this option is specified.

Example 1

Displaying the pool information of the pool for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, and active flash for mainframe.

```
# raidcom get dp_pool
```

PID	POLS	U(%)	AV_CAP(MB)	TP_CAP(MB)	W(%)	H(%)	Num	LDEV#	LCNT	TL_CAP(MB)	BM	TR_CAP(MB)
001	POLN	10	45000000	50000000	50	80	2	265	33	65000000	PF	
4000000		1										
002	POLF	95	10000	100000000	50	80	3	270	900	100000000	PF	
0	0											
004	POLN	0	10000000	100000000	80	90	2	280	0	200000000	PF	
0	0											

Description of each column in output example:

PID

Displays the pool ID of Dynamic Provisioning/Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.

POLS

Displays the status of the pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is suspended in the failure status.) In this status, the pool information cannot be displayed.

U(%)

Displays the usage rate of the pool (including the mapped capacity and the capacity for Full Allocation).

AV_CAP(MB)

Displays the available capacity (megabytes) for the volumes of Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe mapped to this pool.

TP_CAP(MB)

Displays the total capacity (megabytes) of the pool.

W(%)

Displays the threshold value for WARNING set for the pool.

H(%)

Displays the threshold value set for the pool as the high water mark.

Num

Displays the number of LDEVs composing the pool.

LDEV#

Displays the LDEV number of a pool-VOL that includes the pool management area. 65535(ffff) is displayed if the pool is being created.

LCNT

Displays the total number of Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe volumes mapped to the pool.

TL_CAP(MB)

Displays the total capacity (megabytes) of all Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe virtual volumes and Thin Image pairs mapped to the pool.

BM

Displays the I/O Blocking Mode of the pool.

- PF (Pool Full): If the pool is full, you cannot read from the target DP-VOL or write to the target DP-VOL. If the pool VOL is blocked, you can read from the target DP-VOL or write to the target DP-VOL.
- PB (Pool vol Blockade): If the pool VOL is blocked, you cannot read from the target DP-VOL or write to the target DP-VOL. If the pool is full, you can read from the target DP-VOL or write to the target DP-VOL.
- FB (Full or Blockade): You cannot read from the target DP-VOL or write to the target DP-VOL if the pool is full and/or pool VOL is blocked.
- NB (No Blocking): You can read from the target DP-VOL or write to the target DP-VOL even if the pool is full or pool VOL is blocked.
- - (Not supported): The configuration that does not support the I/O Blocking Mode.

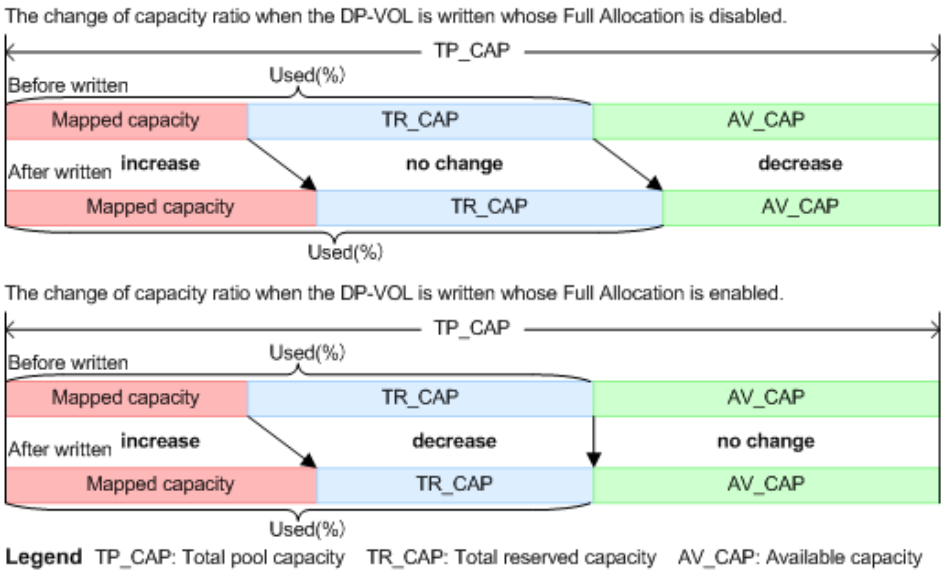
TR_CAP(MB)

Shows the sum of the pool capacities (megabytes) which are reserved for the volumes on which Full Allocation or Proprietary Anchor is enabled. For a configuration that does not support Full Allocation, a hyphen (-) is displayed.

RCNT

Shows the number of volumes for which Full Allocation is enabled that are mapped to a pool. In the configuration that does not support Full Allocation, a hyphen (-) is displayed.

The following figures show the difference between the capacity ratio changing according to whether Full Allocation is enabled or disabled when DP-VOL is written. The "Mapped capacity" in the figure is the total capacity of the user data in each virtual volumes and the page capacity which is storing the control information.



Example 2

Displays the tier information of the pool for Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe. For details about the parameters, see the *Provisioning Guide* for the storage system.

```
# raidcom get dp_pool -key opt
```

PID	POLS	MODE	STS	DAT	TNO	TL_RANGE	TD_RANGE	TU_CAP (MB)	TT_CAP (MB)	T (%)	P (%)	R (%)	B (%)
MM	MC (h)												
001	POLN	DEF	STP	VAL 1	1	00005000	00003000	200000	1000000	80	54	98	40
PM	24												
001	POLN	DEF	STP	VAL 2	2	00003000	00002000	400000	1000000	80	54	98	30
PM	24												
001	POLN	DEF	STP	VAL 3	3	00002000	00002000	600000	1000000	80	54	98	40
PM	24												
002	POLF	AUT	MON	PND 1	1	-	-	500000	1000000	80	54	100	2
PM	8												

Example 3

When you specify -pool, displays the range for each tiering level of corresponding pools.

```
# raidcom get dp_pool -pool 1 -key opt
```

PID	POLS	MODE	STS	DAT	TNO	TL_RANGE	TD_RANGE	TU_CAP (MB)	TT_CAP (MB)	T (%)	P (%)	R (%)
-----	------	------	-----	-----	-----	----------	----------	-------------	-------------	-------	-------	-------

```

B (%) MM MC (h)
001 POLN DEF STP VAL 0 4294967294 00000000 8064 8190 10 28 100 2 PM 24
001 POLN DEF STP VAL 0 4294967294 00000000 3948 3990 10 100 100 2 PM 24
001 POLN DEF STP VAL 0 00000000 00000000 504 8190 10 33 100 2 PM 24
001 POLN DEF STP VAL 1 00000000 00000000 8064 8190 10 28 100 2 PM 24
001 POLN DEF STP VAL 1 00000000 00000000 3948 3990 10 100 100 2 PM 24
001 POLN DEF STP VAL 1 00000000 00000000 504 8190 10 33 100 2 PM 24
001 POLN DEF STP VAL 2 00000039 00000000 8064 8190 10 28 100 2 PM 24
001 POLN DEF STP VAL 2 00000000 00000000 3948 3990 10 100 100 2 PM 24
001 POLN DEF STP VAL 2 00000000 00000000 504 8190 10 33 100 2 PM 24
001 POLN DEF STP VAL 3 00000100 00000000 8064 8190 10 28 100 2 PM 24
001 POLN DEF STP VAL 3 00000041 00000000 3948 3990 10 100 100 2 PM 24
001 POLN DEF STP VAL 3 00000000 00000000 504 8190 10 33 100 2 PM 24
001 POLN DEF STP VAL 4 00000001 00000000 8064 8190 10 28 100 2 PM 24
001 POLN DEF STP VAL 4 00000000 00000000 3948 3990 10 100 100 2 PM 24
001 POLN DEF STP VAL 4 00000000 00000000 504 8190 10 33 100 2 PM 24
001 POLN DEF STP VAL 5 00000001 00000000 8064 8190 10 28 100 2 PM 24
001 POLN DEF STP VAL 5 00000001 00000000 3948 3990 10 100 100 2 PM 24
001 POLN DEF STP VAL 5 00000000 00000000 504 8190 10 33 100 2 PM 24

```

Description of each column in output example:

PID

Displays the pool ID for Dynamic Provisioning, Dynamic Provisioning for Mainframe, Dynamic Tiering, active flash, Dynamic Tiering for Mainframe, or active flash for mainframe.

POLS

Displays the status of the pool.

- POLN: "Pool Normal" (The pool is in the normal status.)
- POLF: "Pool Full" (The pool is in the overflow status exceeding the threshold.)
- POLS: "Pool Suspend" (The pool is in the overflow status exceeding the threshold and is blocked.)
- POLE: "Pool failure" (The pool is suspended in the failure status.) In this status, the pool information cannot be displayed.

MODE

Displays the execution mode.

- DEF: The start/termination of the monitor is performed by the instruction of the CCI, and the Tier range setting is performed by the automatic calculation of the storage system.
- AUT: The start/termination of the monitor is performed by time specification, and the Tier range setting is performed by the automatic calculation of the storage system.

**Note:**

AUT cannot be instructed from the CCI. AUT can only be performed if displayed when set from the SN2.

STS

Displays the operational status of the performance monitor and the tier relocation.

- STP: The performance monitor and the tier relocation are stopped.
- RLC: The performance monitor is stopped. The tier relocation is waiting or operating.
- MON: The performance monitor is operating. The tier relocation is stopped.
- RLM: The performance monitor is operating. The tier relocation is waiting or operating.

DAT

Displays the status of the monitor information.

- VAL: Valid.
- INV: Invalid.
- PND: Being calculated.

TNO

Tiering number. When you specify the pool, displays the tiering level (0 means "all" is set to the level of tiering policy).

TL_RANGE

Lowest limit value of the Tier in IOPH.

0 (0x00000000) to 4294967294 (0xFFFFFFFF)

When "all" is set to the level of tiering policy, the lowest limit value of the tier is displayed. When there is no V-VOL where "all" is set to the level of tiering policy, or when there is no V-VOL related to the pool, the invalid value "-" (0xFFFFFFFF) is displayed. If the -fix option is specified, this displays the value in parentheses given in hexadecimal.

TD_RANGE

Delta value of the Tier in IOPH. 0 (0x00000000) to 4294967294 (0xFFFFFFFFE): When "all" is set to the level of tiering policy, the lowest limit value of the tier is displayed. When there is no V-VOL where "all" is set to the level of tiering policy, or when there is no V-VOL related to the pool, the invalid value "-" (0xFFFFFFFF) is displayed. If the -fix option is specified, this displays the value in parentheses given in hexadecimal.

TU_CAP(MB)

Tier capacity (Usage amount). Displays the actual capacity (megabytes) of the tier (usage amount) when the `-pcap` option is specified.

TT_CAP(MB)

Tier capacity (Total capacity). Displays the actual capacity (megabytes) of the tier (total capacity) when the `-pcap` option is specified.

T(%)

The free space percentage for the new allocation

P(%)

Performance working ratio

R(%)

Progress percentage of relocation.

- 0 to 99: Shows one of the following statuses.
 - When the value of STS is RLC or RLM: relocation is waiting or in progress.
 - When the value of STS is STP or MON: relocation is suspended.
- 100: Shows if the relocation operation is not in progress, or the relocation is complete.

B(%)

Displays the amount of buffer for the tier reallocation.

MM

Displays the mode of performance monitoring and the availability of active flash:

- PM: Periodical mode.
- CM: Continuous mode.
- RPM: Periodical mode and active flash is enabled.
- RCM: Continuous mode and active flash is enabled.

MC(h)

Displays the execution cycle of performance monitoring and tier reallocation. A hyphen (-) is displayed if MODE is not AUT.

For VSP G1x00, VSP F1500, VSP, VSP G200, G400, G600, G800, and VSP F400, F600, F800, a hyphen (-) is always displayed.

raidcom send ping

Sends a ping from the specified port to the specified host, and then displays the result.

Syntax

```
raidcom send ping -port <port#> -address <IP address>
[-iscsi_virtual_port_id <iSCSI virtual port ID>]
```

Options and parameters

-port <port#>

Specifies the port number. For example:

- CL1-A

-address <IP address>

Specifies the IP address of the host that is the destination of the ping. You can specify IPv4 address or IPv6 address.

For details about how to specify an IP address, see [Methods for specifying IP addresses \(on page 20\)](#).

[-iscsi_virtual_port_id <iSCSI virtual port ID>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

An error occurs if you specify this option in any of the following cases:

- The virtual port of the specified iSCSI virtual port ID is not found.
- The iSCSI virtual port mode is disabled.
- The device does not support the iSCSI virtual port mode.

Returned values

The `raidcom send ping` command sets one of the following returned values in `exit()`, which allows users to check the execution results using a user program or script.

▪ **Normal termination:**

- **0:** One or more packets were received.
- **1:** A ping could not be sent, or a ping was sent, but no response packets have been received within the timeout period.

▪ **Abnormal termination:**

- **EX_XXX:** The operation terminated abnormally. For details, see the descriptions of the command error messages in the *Command Control Interface User and Reference Guide*.

Examples

To send a ping from the port: CL4-E to the host whose IP address is 10.213.60.111, and then display the result (Normal).

```
# raidcom send ping -port CL4-E -address 10.213.60.111
raidcom : 5 packets transmitted, 5 packets received.
```

To send a ping from the port: CL4-E, iSCSI virtual port ID: 0 to the host whose IP address is 10.213.60.111, and then display the result (Normal).

```
# raidcom send ping -port CL4-E -address 10.213.60.111 -iscsi_virtual_port_id 0
raidcom : 5 packets transmitted, 5 packets received.
```

To send a ping from the port: CL4-E to the host whose IP address is 10.213.60.111, and then display the result (No response from the host).

```
# raidcom send ping -port CL4-E -address 10.213.60.111
raidcom : 5 packets transmitted, 0 packets received.
```

To send a ping from the port: CL4-E to the host whose IP address is 10.213.60.111, and then display the result (Fail to send a ping by an internal error etc).

```
# raidcom send ping -port CL4-E -address 10.213.60.111
raidcom : 0 packets transmitted.
```

Description of output example:

***n* packets transmitted**

Number of packets sent to the host.

***n* packets received**

Number of packets responded from the host normally.

raidcom add server

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Creates a server managed by Storage Advisor Embedded or VSP One Block Administrator. To execute this command, you need to log in to the storage system with the user account belonging to the user group to which "all resource groups" are set. In addition, you must be able to perform operations with all implemented resource groups. Therefore, this command cannot be executed if any of the following conditions is met:

- The command device belonging to a resource group other than meta_resource is used in the in-band management.
- The virtual storage machine that contains only some of the implemented resource groups is specified as HORCM_VCMD in the configuration file.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom add server -server_name <name> -request_id auto
```

Options and parameters

-server_name <name> -request_id auto

Specifies the server nickname with one character through 229 characters.

The following are the characters that can be used:

0 through 9, A through Z, a through z, -, ., /, @, \, _ single-byte space

This command creates a server with no OS type set. 255 servers are created in the storage system.

The server ID is automatically selected. You can check the selected server ID using the **raidcom get command_status** command.

The request ID is output in the following format when the execution of the **raidcom add server** command ends:

- REQID : <request#>

<request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

Example

Creating a server with the server name: server0.

```
#raidcom add server -server_name server0 -request_id auto
REQID:0
```

raidcom delete server

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Deletes a server managed by Storage Advisor Embedded or VSP One Block Administrator. To execute this command, you need to log in to the storage system with the user account belonging to the user group to which "all resource groups" are set. In addition, you must be able to perform operations with all implemented resource groups. Therefore, this command cannot be executed if any of the following conditions is met:

- The command device belonging to a resource group other than meta_resource is used in the in-band management.
- The virtual storage machine that contains only some of the implemented resource groups is specified as HORCM_VCMD in the configuration file.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom delete server {-server_id <id>|-server_name <name>} -request_id auto
```

Options and parameters

{-server_id <id>|-server_name <name>} -request_id auto

Specifies the ID or nickname of the server to be deleted.

You need to specify the server ID (0 through 4095) in decimal or hexadecimal in -server_id <id>. To specify in hexadecimal, add "0x" at the beginning of the server ID.

For example:

- -server_id 128
- -server_id 0x80

To specify the server nickname, specify the server nickname in -server_name <name> with one character through 229 characters.

The following are the characters that can be used:

0 through 9, A through Z, a through z, -, ., @, \, _ single-byte space

The specified server is deleted by executing this command. The registered host groups or iSCSI targets are not deleted.

The request ID is output in the following format when the execution of the **raidcom delete server** command ends:

- REQID : <request#>

<request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

Example

Deleting a server with the server ID: 0.

```
#raidcom delete server -server_id 0 -request_id auto
REQID:0
```

raidcom modify server

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Modifies a server managed by Storage Advisor Embedded or VSP One Block Administrator. To execute this command, you need to log in to the storage system with the user account belonging to the user group to which "all resource groups" are set. In addition, you must be able to perform operations with all implemented resource groups. Therefore, this command cannot be executed if any of the following conditions is met:

- The command device belonging to a resource group other than meta_resource is used in the in-band management.
- The virtual storage machine that contains only some of the implemented resource groups is specified as HORCM_VCMD in the configuration file.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom modify server {-server_id <id>|-server_name <name>} -request_id auto {-
server_operation {add_host_grp|delete_host_grp} -port <port> [<host group name>]}
```

Options and parameters

{-server_id <id>|-server_name <name>} -request_id auto

Specifies the ID or nickname of the server to be modified.

You need to specify the server ID (0 through 4095) in decimal or hexadecimal in -server_id <id>. To specify in hexadecimal, add "0x" at the beginning of the server ID. For example:

- -server_id 128
- -server_id 0x80

To specify the server nickname, specify the server nickname in -server_name <name> with one character through 229 characters.

The following are the characters that can be used:

0 through 9, A through Z, a through z, -, /, @, \, _ single-byte space

The specified server is deleted by executing this command. The registered host groups or iSCSI targets are not deleted.

The request ID is output in the following format when the execution of the **raidcom modify server** command ends:

- REQID : <request#>

<request#> is a request ID assigned each time the command is executed. For details, see [Request ID function \(on page 293\)](#).

-server_operation add_host_grp -port <port> [<host group name>]

Registers the host groups or the iSCSI targets specified with the -port <port> [<host group name>] on the server specified for the -server_id <id> or the -server_name <name>.

If the WWN or iSCSI name of the host is registered in the specified host groups or iSCSI targets, the WWN or iSCSI name of the host is registered on the server. In addition, if an LU path is defined for a specified host groups or iSCSI targets, the LDEV associated with the LU path is registered on the server.

For -port <port> [<host group name>], specify the port number and host groups (iSCSI target ID for iSCSI), or host group name (target alias for iSCSI).

If the host group name or the target alias is more than 64 characters, the host group name or the target alias cannot be specified. Use the host group ID or the target ID. For example:

- CL1-A -g (g is from 1 to 254)
- CL1-A Linux_X86
- CL1-A Target00

To check if host groups or iSCSI targets are registered on the server, execute the **raidcom get host_grp** command with the `-key server` option specified.

If any of the following conditions is met, the command execution fails:

- The specified server and host group/iSCSI target are not installed.
- The IDs of the specified host groups or iSCSI targets are 0.
- The host group/iSCSI target, LU path, or LDEV information on the specified server is being updated.
- The iSCSI target is specified even if the specified server protocol is Fibre. Alternatively, a host group is specified even if the specified server protocol is iSCSI.
- The host group/iSCSI target which has already been registered on other servers is specified.
- Another host group/iSCSI target on the same port as the one used for the specified host group/iSCSI target has already been registered on the specified server.
- If 32 WWNs or iSCSI names of the host have been registered on the specified server, the host group/iSCSI target including WWNs or iSCSI names that are not registered on the server is specified.
- The specified host group/iSCSI target, the port to which the host group/iSCSI target belongs, or the resource group to which the LDEV associated with the LU path of the host group/iSCSI target belongs is locked by another user.
- The LU path to an LDEV other than that for Dynamic Provisioning, Dynamic Tiering, or active flash exists in the specified host group/iSCSI target.
- The LDEV for a Volume Migration pair including LDEV other than that for Dynamic Provisioning, Dynamic Tiering, or active flash is specified.
- The host group/iSCSI target on the port with the LU security set to OFF is specified.

-server_operation delete_host_grp -port <port> [<host group name>]

Specifies the iSCSI virtual port ID when the specified port is iSCSI and the iSCSI virtual port mode is enabled. If this option is omitted, 0 is assumed.

From the server specified for `-server_id <id>` or `-server_name <name>`, delete the host groups or iSCSI targets specified with `-port <port> [<host group name>]`.

The LDEV is deleted from the server if an LU path is defined for the specified host groups or iSCSI targets, and if the LU path is the last one of the LDEV associated with the LU path on the specified server.

For `-port <port> [<host group name>]`, specify the port number and host groups (iSCSI target ID for iSCSI), or the host group name (target alias for iSCSI). If the host group name or the target alias is more than 64 characters, the host group name or the target alias cannot be specified. Use the host group ID or the target ID.

For example:

- CL1-A -g (g is from 1 to 254)
- CL1-A Linux_X86
- CL1-A Target00

To check if host groups or an iSCSI targets are registered on the server, execute the **raidcom get host_grp** command with the `-key server` option specified.

If any of the following conditions is met, the command execution fails:

- The specified server and host group/iSCSI target are not installed.
- The IDs of the specified host groups or iSCSI targets are 0.
- The host group/iSCSI target, LU path, or LDEV information on the specified server is being updated.

Examples

Adding CL1-A-1 to server ID: 0.

```
#raidcom modify server -server_id 0 -request_id auto -server_operation add_host_grp -
port CL1-A-1
REQID:0
```

Deleting CL1-A-1 from server ID: 0.

```
#raidcom modify server -server_id 0 -request_id auto -server_operation
delete_host_grp -port CL1-A-1
REQID:0
```

raidcom get server

Supported storage systems:

- VSP One B20
- VSP E series
- VSP G130, G/F350, G/F370, G/F700, G/F900

Displays a server managed by Storage Advisor Embedded or VSP One Block Administrator. To execute this command, you need to log in to the storage system with the user account belonging to the user group to which "all resource groups" are set. In addition, you must be able to perform operations with all implemented resource groups. Therefore, this command cannot be executed if any of the following conditions is met:

- The command device belonging to a resource group other than meta_resource is used in the in-band management.
- The virtual storage machine that contains only some of the implemented resource groups is specified as HORCM_VCMD in the configuration file.

Syntax

```
raidcom get server [-key <keyword> [-server_id <id>|-server_name <name>]]
```

Options and parameters

[-key <keyword> [-server_id <id>|-server_name <name>]]

Specifies the server information to be displayed.

- **opt:** Specifies the detailed information about the server. To specify this parameter, make sure to specify the server ID or nickname that displays the detailed information in `-server_id <id>` or `-server_name <name>`.

You need to specify the server ID (0 through 4095) in decimal or hexadecimal in `-server_id <id>`. To specify in hexadecimal, add "0x" at the beginning of the server ID.

For example:

- `-server_id 128`
- `-server_id 0x80`

To specify the server nickname, specify the server nickname in `-server_name <name>` with one character through 229 characters.

The following are the characters that can be used:

0 through 9, A through Z, a through z, -, ., /, @, \, _ single-byte space

Example

Displaying server information.

```
#raidcom get server
SRVID NAME
0 "server0"
1 "server1"
155 "server2"
```

Description of each column in the output example:

SRVID

Displays the server ID in decimal.

NAME

Displays the server nickname. A double quotation (") is added at the beginning and end of the server nickname if the `-key opt` option is not specified.

Specifying the server information to be displayed.

```
#raidcom get server -key opt -server_id 0
SRVID : 0
NAME : server0
STS : RCV DGG1
PROTOCOL : FIBRE
OS : WIN
OS_OPT : 1 3
```

Description of each column in the output example:

STS

Status of the server. The status of the server is not displayed if it is not in any of the following statuses:

- RSV: Protocol, OS type, and OS type option are not set.
When setting the host group/iSCSI target to the server with the **raidcom modify server** command, the OS type is set based on the host mode of the host group/iSCSI target.
- ILU: No LU path is defined between any of the LDEV registered on the server and any of the host group/iSCSI target registered on the server.
- IOT: The OS type of the server does not correspond to the host mode of any of the host group/iSCSI target registered on the server.
- IOP: The OS type of the server does not correspond to the host mode option of any of the host group/iSCSI target registered on the server.
- RCV: The host group/iSCSI target, LU path, or LDEV information on the server is being updated.
- DGG1: For VSP G130, G/F350, G/F370, G/F700, G/F900, the microcode version cannot be downgraded to 88-01-xx-xx/xx or earlier. For VSP One B20, VSP E series, ignore this item.

PROTOCOL

Displays the protocol of the server.

- FIBRE: The protocol is Fiber Channel.
- ISCSI: The protocol is iSCSI.
- - (hyphen): The protocol is not either Fiber Channel or iSCSI.

OS

Displays the OS type of the server. If other OS types are displayed, the OS type that is not supported by Storage Advisor Embedded or VSP One Block Administrator is set.

- LINUX/IRIX: The OS type is "Linux".
- VMWARE: The OS type is "(Deprecated) VMware".
- HP-UX: The OS type is "HP-UX".

- OVMS: The OS type is "OpenVMS".
- TRU 64: The OS type is "Tru64".
- SOLARIS: The OS type is "Solaris".
- NETWARE: The OS type is "NetWare".
- WIN: The OS type is "(Deprecated) Windows".
- AIX: The OS type is "AIX".
- VMWARE_EX: The OS type is "VMWare".
- WIN_EX: The OS type is "Windows".
- - (hyphen): No OS type is set.

OS_OPT

OS type options that are set on the server. A hyphen (-) is displayed if the OS type is not set.

raidcom get apn

Displays the AP numbers used in the storage system. These commands are used for troubleshooting. Use them only if you get any directions.

Syntax

```
raidcom get apn [-fx]
```

Options and parameters

-fx

Specifies the LDEV number of the command device in hexadecimal.

Displays the AP numbers used in the storage system.

```
#raidcom get apn
AP#    LDEV#    WWN
0001      1    5060000000000001
0002      2    5060000000000002
0eff     12    -
```

Description of each column in output example:

AP#

Displays the AP# being used in hexadecimal.

LDEV

Displays the LDEV number of the command device using AP#.

WWN

If the command device is used via Fibre Channel, the HBA WWN is displayed.

If a command device is used via iSCSI, a hyphen(-) is displayed.

raidcom add qos_grp

Supported storage systems:

- VSP 5000 series
- VSP One B20

Creates a QoS group or adds an LDEV to the specified QoS group.

This command is executed asynchronously with the command input. Check the completion of this processing using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add qos_grp -qos_grp_id <qos group#> [-ldev_id <ldev#>] -request_id auto
```

Options and parameters

-qos_grp_id <qos-group#>

Specifies a QoS group ID.

[-ldev_id <ldev#>]

Specifies an LDEV number (0-65279) to be set in the QoS group.

This option is available only when an LDEV is added to the QoS group.

Example:

- -ldev_id 200

If the command is being run continuously by specifying this option, run it within the allowable numbers of commands that the storage system can accept. For more information, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#). The command ends normally even if an LDEV specified with -ldev exists in a QoS group specified with **-qos_grp_id**.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Creating the QoS group whose QoS group ID is 1.

```
# raidcom add qos_grp -qos_grp_id 1 -request_id auto
```

Adding LDEV ID: 101 to the QoS group whose QoS group ID is 1.

```
# raidcom add qos_grp -qos_grp_id 1 -ldev_id 101 -request_id auto
```

Adding LDEV IDs: 1, 2, and 3 to the QoS group whose QoS group ID is 1.

```
# raidcom get ldev -ldev_id 1 -cnt 3 -key front_end | rmawk @L-ne:1
exe="raidcom add qos_grp -qos_grp_id 1 -ldev_id @2 -request_id auto"
REQID : 1
REQID : 2
REQID : 3
# raidcom get command_command status
# raidcom reset command_status -request_id all
# raidcom reset command_status
```

raidcom delete qos_grp

Supported storage systems:

- VSP 5000 series
- VSP One B20

Deletes an LDEV from the QoS group or deletes a QoS group.

This command is executed asynchronously with the command input. Check the completion of this processing using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete qos_grp -qos_grp_id <qos group#> [-ldev_id <ldev#>] -request_id auto
```

Options and parameters

-qos_grp_id <qos-group#>

Specifies a QoS group ID.

[-ldev_id <ldev#>]

Specifies an LDEV number (0-65279) to be deleted from the QoS group.

Example:

- -ldev_id 200

If the command is being run continuously by specifying this option, run it within the allowable numbers of commands that the storage system can accept. For more information, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#). The command ends normally even if an LDEV specified with -ldev does not exist in a QoS group specified with -qos_grp_id.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Deleting LDEV ID: 101 from the QoS group whose QoS group ID is 1.

```
# raidcom delete qos_grp -qos_grp_id 1 -ldev_id 101 -request_id auto
```

Deleting the QoS group whose QoS group ID is 1.

```
# raidcom delete qos_grp -qos_grp_id 1 -request_id auto
```

Deleting up to 2000 LDEVs from the QoS group whose QoS group ID is 1.

```
#raidcom get qos_grp -qos_grp_id 1 -key resource | rmawk @L-ne:1 -a @L-le:2001
exe="raidcom delete qos_grp -qos_grp_id 1 -ldev_id @2 -request_id auto"
REQID : 1
REQID : 2
REQID : 3
...
# raidcom get command_command status
# raidcom reset command_status -request_id all
# raidcom reset command_status
```

raidcom modify qos_grp

Supported storage systems:

- VSP 5000 series
- VSP One B20

Changes the QoS settings for each QoS group.

This command is executed asynchronously with the command input. Check the completion of this processing using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify qos_grp -qos_grp_id <qos group#> {-upper_throughput_io <upper throughput io> | -upper_data_trans_mb <upper data trans mb> | -upper_alert_time <upper alert time> | -response_priority <#priority> | -response_alert_time <response alert time>} -request_id auto
```

Options and parameters

-qos_grp_id <qos-group#>

Specifies a QoS group ID.

[-upper_throughput_io <upper throughput io>]

Specifies the upper limit of the throughput per second. The upper limit is the total throughput for all LDEVs in the group. To disable it, specify zero.

[-upper_data_trans_mb <upper data trans mb>]

Specifies the upper limit of the data transfer volume in megabytes (MB) per second. The upper limit is the total data transfer volumes for all LDEVs in the group. To disable it, specify zero.

[-upper_alert_time <upper alert time>]

Specifies the alert notification time (second) when the upper limit is exceeded. To disable it, specify zero.

An alert is issued if either of the following conditions continues for the alert notification time or longer:

- The number of received commands continuously exceeds the upper limit of the throughput per second.
- The transfer volume of received commands continuously exceeds the upper limit of the data transfer volume per second.

[-response_priority <#priority>]

Specifies the I/O processing priority with 1 to 3. To disable it, specify 0. 1 represents the lowest priority and a greater number represents a higher priority.

For details about the I/O processing priority, see the *Performance Guide*.

[-response_alert_time <response alert time>]

Specifies the alert notification time (second) when the target response time is not reached. If the target response time has not been continuously reached for more than the alert notification time, an alert is issued. To disable it, specify 0.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Changing the upper limit of the data transfer volume per second to 100 megabytes per second (MB/s) in the QoS group whose QoS group ID is 1.

```
# raidcom modify qos_grp -qos_grp_id 1 -upper_data_trans_mb 100 -request_id auto
```

Specifying the I/O processing priority to 1 in the QoS group whose QoS group ID is 1.

```
# raidcom modify qos_grp -qos_grp_id 1 -response_priority 1 -request_id auto
```

raidcom get qos_grp

Supported storage systems:

- VSP 5000 series
- VSP One B20

Displays the QoS group information.

Syntax

```
raidcom get qos_grp [-qos_grp_id <qos group#>] [-key <resource | monitor>][--time_zone <time zone>]
```

Options and parameters

-qos_grp_id <qos-group#>

Specifies a QoS group ID. This option displays only the specified QoS group information.

[-key <resource | monitor>]

Specifies the QoS group information to be displayed.

- resource: Specifies the information on LDEVs included in the QoS group. To specify this parameter, make sure to specify the QoS group ID.
- monitor: Specifies the performance monitor information in the QoS group. If the command is run immediately after creating the QoS group, the QoS group might not be displayed in the command output. In this case, rerun the command after a few seconds.

[-time_zone <time zone>]

Specifies the time zone to be displayed. If this option is not specified, the time in the time zone set for the storage system is displayed.

utc: displays the time in Coordinated Universal Time (UTC).

Example 1

Displaying the list of QoS groups. A hyphen (-) is displayed if the QoS information is not set.

```
# raidcom get qos_grp
QID UPPER_THROUGHPUT(I/Ops) UPPER_DATA_TRANS(MBps) UPPER_ALERT(s) UPPER_ALERT_NOTICE
PRIORITY RESPONSE_TARGET(ms) RESPONSE_ALERT(s) RESPONSE_ALERT_NOTICE
1 80 100 30 2019-07-31T10:15:20 1 80 100 2019-07-31T10:15:20
2 80 100 30 2019-07-31T10:15:20 1 80 100 2019-07-31T10:15:20
3 80 100 30 2019-07-31T10:15:20 1 80 100 2019-07-31T10:15:20
```

Description of each column in output example:**QID**

Displays the QoS group ID.

UPPER_THROUGHPUT(IOps)

Displays the upper limit of the throughput per second. The upper limit is the total throughput of all LDEVs in the group.

UPPER_DATA_TRANS(MBps)

Displays the upper limit of the data transfer volume in megabytes (MB) per second. The upper limit is the total data transfer volumes of all LDEVs in the group.

UPPER_ALERT(s)

Displays the alert notification time (second) when the upper limit is exceeded.

An alert is issued if either of the following conditions continues for the alert notification time or longer:

- The number of received commands continuously exceeds the upper limit of the throughput per second.
- The transfer volume of received commands continuously exceeds the upper limit of the data transfer volume per second.

UPPER_ALERT_NOTICE

Displays the last alert time of the following alerts:

- The alert issued when the number of received commands continuously exceeds the upper limit of the throughput per second.
- The alert issued when the transfer volume of received commands continuously exceeds the upper limit of the data transfer volume per second.

If the last alert time does not exist, a hyphen (-) is displayed.

PRIORITY

Displays the I/O processing priority. 1 represents the lowest priority and a greater number represents a higher priority.

If the DKCMAIN microcode version is not supported, a hyphen (-) is displayed.

RESPONSE_TARGET(ms)

Displays the target response time.

If the DKCMAIN microcode version is not supported, a hyphen (-) is displayed.

RESPONSE_ALERT(s)

Displays the alert notification time (second) when the target response time is not reached. If the target response time has not been continuously reached for more than the alert notification time, an alert is issued.

If the DKCMAIN microcode version is not supported, a hyphen (-) is displayed.

RESPONSE_ALERT_NOTICE

Displays the last alert time when the target response time has not been continuously reached.

If the last alert time does not exist, or the DKCMAIN microcode version is not supported, a hyphen (-) is displayed.

Example 2

Displaying the information on the QoS group whose QoS group ID is 1. A hyphen (-) is displayed if the QoS information is not set.

```
# raidcom get qos_grp -qos_grp_id 1
QID UPPER_THROUGHPUT(Iops) UPPER_DATA_TRANS(MBps) UPPER_ALERT(s) UPPER_ALERT_NOTICE
PRIORITY RESPONSE_TARGET(ms) RESPONSE_ALERT(s) RESPONSE_ALERT_NOTICE
1 80 100 30 2019-07-31T10:15:20 1 80 100 2019-07-31T10:15:20
```

Displaying the information on LDEVs included in the QoS group whose QoS group ID is 1.

```
# raidcom get qos_grp -qos_grp_id 1 -key resource
QID LDEV#
1 100
1 101
```

Description of each column in output example:

LDEV#

Displays the numbers of LDEVs included in the QoS group.

Example 3

Displaying the performance monitor information on all LDEVs in the QoS group. If all performance monitor information is not valid, nothing is displayed.

```
# raidcom get qos_grp -key monitor
QID RCV_CMD(Iops) RCV_CMD_TRANS(KBps) THROUGHPUT(Iops) DATA_TRANS(KBps) RESPONSE(us)
MONITOR_TIME
1 3840 15 2650 10 5 2019-07-31T10:15:20
```

Description of each column in output example:

RCV_CMD(Iops)

Displays the total numbers of received commands for all LDEVs in the QoS group.

RCV_CMD_TRANS(KBps)

Displays the total data transfer volumes of received commands for all LDEVs in the QoS group.

THROUGHPUT(Iops)

Displays the total throughput per second for all LDEVs in the QoS group.

DATA_TRANS(KBps)

Displays the total data transfer volumes per second for all LDEVs in the QoS group.

RESPONSE(us)

Displays the average response time for all LDEVs in the QoS group.

MONITOR_TIME

Displays the monitoring collection time.

raidcom monitor resource

Supported storage systems:

- VSP 5000 series
- VSP One B20

Displays the performance information on the storage system. To run this command, you need to log in to the storage system with the user account belonging to the user group to which all resource groups are set. In addition, you must be able to perform the operations with all implemented resource groups. Therefore, this command cannot be run if either of the following conditions is met:

- The command device belonging to a resource group other than meta_resource is used in the in-band management.
- The virtual storage machine that contains only some of the implemented resource groups is specified as HORCM_VCMD in the configuration file.

Syntax

```
raidcom monitor resource -monitor_type qos_ldev [-qos_grp_id <qos group id>] [-start
_id <start id#>] [-cnt <count>] -monitor_interval <interval> [<repetition count>] [-
time_zone <time zone>]
raidcom monitor resource -monitor_type qos_grp [-start _id <start id#>] [-cnt <count>]
-monitor_interval <interval> [<repetition count>] [-time_zone <time zone>]
```

Options and parameters**-monitor_type qos_ldev**

Displays the QoS target LDEV performance information that is collected every one second on the storage system. If the command is run immediately after setting the LDEV as a QoS target, the LDEV might not be displayed in the command output. In this case, rerun the command after a few seconds.

-monitor_type qos_grp

Displays the QoS group performance information that is collected every one second on the storage system. If the command is run immediately after creating the QoS group, the QoS group might not be displayed in the command output. In this case, rerun the command after a few seconds.

[-qos_grp_id <qos group id>]

Specifies a QoS group ID. This option displays the performance information on the LDEV in the QoS group specified with <qos group id>. This option is available only when `-monitor_type qos_ldev` is specified.

[-start_id <start id#>]

Specifies a minimum ID of the resource with which the display is started in decimal.

If `-monitor_type qos_ldev` is specified, specify a minimum LDEV ID with which the display is started.

If `-monitor_type qos_grp` is specified, specify a minimum QoS group ID with which the display is started.

[-cnt <count>]

Specifies a maximum number (1-65535) of the displayed resource in decimal. This option displays the resource information by the number specified with <count> in increasing order from the smallest ID of the resources that meet the display conditions. If the number of resources that meet the display conditions is less than specified with <count>, the resources that meet all requirements are displayed.

If `-monitor_type qos_ldev` is specified, a maximum number of the displayed LDEVs is specified.

If `-monitor_type qos_grp` is specified, a maximum number of the displayed QoS groups is specified.

-monitor_interval <interval> [<repetition count>]

Specifies the repetition interval for monitoring the performance information in 100 ms as a decimal number (1-36000). (0.1[s]-1[hour])

If <repetition count> (1-200000) is specified, the monitoring is performed the number of times specified with <repetition count>.

If <repetition count> is not specified, the monitoring is performed until the user terminates the command with CTL+C/SIGTERM.

<repetition count> specifies a decimal number.

If the number of displayed resources is large relative to the command execution rate, all information might not be collected within the repetition interval specified with <interval>. In that case, reduce the number of displayed resources or increase the value of <interval>.

[-time_zone <time zone>]

Specifies the time zone to be displayed in MONITOR_TIME. If this option is not specified, the time in the time zone set for the storage system is displayed.

Example 1

Of the LDEVs belonging to the QoS group ID: 1, displaying the performance information on at least two LDEVs whose LDEV number is 200 or more every one second. The number of times to display is 3 times.

```
# raidcom monitor resource -monitor_type qos_ldev -qos_grp_id 1 -start_id 200 -cnt 2 -
monitor_interval 10 3
```



```

LDEV# RCV_CMD (IOps) RCV_CMD_TRANS (KBps) THROUGHPUT (IOps) DATA_TRANS (KBps) RESPONSE (us)
MONITOR_TIME
200 3840 15 2650 10 5 2019-07-31T10:15:20
201 3840 15 2650 10 5 2019-07-31T10:15:20
200 3840 15 2650 10 5 2019-07-31T10:15:21
201 3840 15 2650 10 5 2019-07-31T10:15:21
200 3840 15 2650 10 5 2019-07-31T10:15:22
201 3840 15 2650 10 5 2019-07-31T10:15:22

```

Description of each column in output example:

LDEV#

Displays the numbers of LDEVs included in the QoS group.

RCV_CMD(IOps)

Displays the numbers of received commands.

RCV_CMD_TRANS(KBps)

Displays the data transfer volumes of received commands.

THROUGHPUT(IOps)

Displays the throughput per second.

DATA_TRANS(KBps)

Displays the data transfer volumes per second.

RESPONSE(us)

Displays the response time.

MONITOR_TIME

Displays the monitoring collection time.

Example 2

Displaying the performance information on at least two QoS groups whose QoS group ID is 200 or more every one second. The number of times to display is 3 times.

```

# raidcom monitor resource -monitor_type qos_grp -start_id 200 -cnt 2 -
monitor_interval 10 3
QID RCV_CMD (IOps) RCV_CMD_TRANS (KBps) THROUGHPUT (IOps) DATA_TRANS (KBps) RESPONSE (us)
MONITOR_TIME
200 3840 15 2650 10 5 2019-07-31T10:15:20
201 3840 15 2650 10 5 2019-07-31T10:15:20
200 3840 15 2650 10 5 2019-07-31T10:15:21
201 3840 15 2650 10 5 2019-07-31T10:15:21
200 3840 15 2650 10 5 2019-07-31T10:15:22
201 3840 15 2650 10 5 2019-07-31T10:15:22

```

Description of each column in output example:

QID

Displays the QoS group ID in decimal.

raidcom add nvm_subsystem

Supported storage systems:

- VSP 5000 series
- VSP E1090

Creates an NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom add nvm_subsystem -nvm_subsystem_id <nvm subsystem id> [-
nvm_subsystem_name <nvm subsystem name>] [-host_mode <host mode> [-
host_mode_opt <host mode option> ...]] [-namespace_security <enable |
disable>] [-virtual_nvm_subsystem_id <virtual nvm subsystem id>] -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

[-nvm_subsystem_name <nvm subsystem name>]

Specifies the NVM subsystem name.

If this option is omitted, the name is automatically set by the system.

[-host_mode <host mode>]

Specifies the host mode.

Specify the following character strings. For more information, see the *Provisioning Guide for Open Systems*.

- LINUX
- VMWARE
- VMWARE_EX
- AIX

[-host_mode_opt <host mode option> ...]

Specifies the host mode option. For more information, see the *Provisioning Guide for Open Systems*.

If this option is omitted, the host mode option is disabled.

[-namespace_security <enable | disable>]

Specifies the setting to disable or enable use of the namespace security.

- enable: Enable the namespace security.
- disable: Disable the namespace security.

If this option is omitted, the namespace security is enabled.

[-virtual_nvm_subsystem_id <virtual nvm subsystem id>]

Specifies the virtual NVM subsystem ID.

(VSP 5000 series and VSP One B20 only) If this option is omitted, the ID specified for the `-nvm_subsystem_id <nvm subsystem id>` option is used as the virtual NVM subsystem ID.

-request_id auto

`<request#>` is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Examples

Creating an NVM subsystem whose NVM subsystem ID is 1 and whose NVM subsystem name is `my_nvm_subsystem`.

```
# raidcom add nvm_subsystem -nvm_subsystem_id 1 -nvm_subsystem_name
my_nvm_subsystem -request_id auto
```

Creating an NVM subsystem whose NVM subsystem ID is 1 and for which the namespace security is disabled.

```
# raidcom add nvm_subsystem -nvm_subsystem_id 1 -namespace_security
disable -request_id auto
```

Creating an NVM subsystem whose NVM subsystem ID is 1, whose virtual NVM subsystem ID is 2, and whose NVM subsystem name is `my_nvm_subsystem`.

```
# raidcom add nvm_subsystem -nvm_subsystem_id 1 -nvm_subsystem_name
my_nvm_subsystem -virtual_nvm_subsystem_id 2 -request_id auto
```

raidcom modify nvm_subsystem

Supported storage systems:

- VSP 5000 series
- VSP E1090

Changes an NVM subsystem.

The NVM subsystem name, the host mode, the host mode option, and the namespace security can be changed.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the `raidcom get command_status` command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom modify nvmm_subsystem {-nvmm_subsystem_id <nvmm subsystem id>} {-
rename_nvmm_subsystem <nvmm subsystem nickname> | -host_mode <host mode>
| -set_host_mode_opt <host mode option> ... | -reset_host_mode_opt | -
namespace_security <enable | disable>} -request_id auto
```

Options and parameters

-nvmm_subsystem_id <nvmm subsystem id>

Specifies the NVM subsystem ID.

-rename_nvmm_subsystem <nvmm subsystem nickname>

Specifies the NVM subsystem name to which it is changed.

[-host_mode <host mode>]

Specifies the host mode.

Specify the following character strings. For more information, see the *Provisioning Guide for Open Systems*.

- LINUX
- VMWARE
- VMWARE_EX
- AIX

-set host_mode_opt <host mode option> ...

Specifies the host mode option.

Specify the host mode option for <host mode option>. If you specify this option, all host mode options that you do not specify are cleared. For more information, see the *Provisioning Guide for Open Systems*.

-reset_host_mode_opt

Clears all host mode options.

For more information about the host mode option, see the *Provisioning Guide for Open Systems*.

-namespace_security <enable | disable>

Specifies the setting to disable or enable use of the namespace security.

- enable: Enable the namespace security.
- disable: Disable the namespace security.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Examples

Changing the NVM subsystem name with the NVM subsystem ID 1 to my_subsystem.

```
# raidcom modify nvm_subsystem -nvm_subsystem_id 1 -rename_nvm_subsystem
my_subsystem -request_id auto
```

Changing the namespace security for the NVM subsystem ID 1 to disable.

```
# raidcom modify nvm_subsystem -nvm_subsystem_id 1 -namespace_security
disable -request_id auto
```

raidcom delete nvm_subsystem

Supported storage systems:

- VSP 5000 series
- VSP E1090

Deletes an NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom delete nvm_subsystem {-nvm_subsystem_id <nvm subsystem id>} -
request_id auto
```

Options and parameters**-nvm_subsystem_id <nvm subsystem id>**

Specifies the NVM subsystem ID.

-reset_host_mode_opt

Clears all host mode options.

For more information about the host mode option, see the *Provisioning Guide for Open Systems*.

Example

Deleting the NVM subsystem whose NVM subsystem ID is 1.

```
# raidcom delete nvm_subsystem -nvm_subsystem_id 1 -request_id auto
```

raidcom get nvm_subsystem

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays the NVM subsystem information.

Syntax

```
raidcom get nvm_subsystem [-nvm_subsystem_id <nvm subsystem id>] [-key  
<keyword>]
```

Options and parameters**-nvm_subsystem_id <nvm subsystem id>**

Specifies the NVM subsystem ID.

[-key <keyword>]

Specifies one of the following values:

- **opt**: Displays the information about the subsystem NQN of the specified NVM subsystem.
- **optp**: Displays the information about the subsystem NQN of the specified NVM subsystem. This option displays the NQN that is generated from the actual serial number of the storage system and the NVM subsystem ID. The subsystem NQN displayed by this option is usually unused.
- **namespace**: Displays the information about the namespace created on the specified NVM subsystem.
- **port**: Displays the information about the port registered on the specified NVM subsystem.
- **detail**: Displays the information about the implemented and the unimplemented NVM subsystems.
- **undefined**: Displays the unimplemented NVM subsystem ID.

Example

Displaying the NVM subsystem list.

```
# raidcom get nvm_subsystem
NVMSS_ID RGID NVMSS_NAME SECURITY T10PI
HMD HMO_BITS V_NVMSS_ID
1 0 nvms_id_00001(default_name) ENABLE DISABLE
VMWARE - 1
```

Description of each column in output example:**NVMSS_ID**

Displays the NVM subsystem ID.

RGID

Displays the resource group ID.

NVMSS_NAME

Displays the NVM subsystem name.

SECURITY

Displays the setting to disable or enable use of the namespace security.

- **ENABLE**: Enable the namespace security.
- **DISABLE**: Disable the namespace security.
- **– (hyphen)**: This information is not available.

T10PI

Displays the setting to disable or enable the T10 PI mode.

- ENABLE: Enable the T10 PI mode.
- DISABLE: Disable the T10 PI mode.
- – (hyphen): This information is not available.

HMD

Displays the data transfer volumes per second.

- – (hyphen): This information is not available.

HMO_BITS

Displays BIT to set the host mode option for the host group. If multiple host mode options are specified, list the specified host mode options separated by a colon (:). If no host mode options are specified, a hyphen (-) is displayed. For more information about BIT, see the *Provisioning Guide for Open Systems*.

V_NVMSS_ID

Displays the virtual NVM subsystem ID.

If no virtual NVM subsystem ID is specified, a hyphen (-) is displayed.

Examples

Displaying the information about a subsystem NQN of the NVM subsystem ID 1.

```
# raidcom get nvm_subsystem -nvm_subsystem_id 1 -key opt
NVMSS_ID NVMSS_NAME NVMSS_NQN
1 nvmss_id_00001(default_name) nqn.xxxxx
```

Displaying the information about a subsystem NQN of the NVM subsystem ID 1 (the NQN is generated from the actual serial number of the storage system and the NVM subsystem ID).

```
# raidcom get nvm_subsystem -nvm_subsystem_id 1 -key optp
NVMSS_ID NVMSS_NAME NVMSS_NQN
1 nvmss_id_00001(default_name) nqn.xxxxx
```

Description of each column in output example:**NVMSS_NQN**

Displays the subsystem NQN. When the `-key optp` option is specified, this column displays the information about the subsystem NQN that is generated from the actual serial number of the storage system and the NVM subsystem ID.

Example

Displaying the information about a namespace registered on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom get nvm_subsystem -nvm_subsystem_id 1 -key namespace
NVMSS_ID NVMSS_NAME NSID LDEVID
```



```
1 nvmss_id_00001(default_name) 1 1
1 nvmss_id_00001(default_name) 2 2
```

Description of each column in output example:

NSID

Displays the namespace ID.

LDEVID

Displays the actual LDEV ID in the storage system.

Example

Displaying the information about an NVM subsystem port registered on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom get nvm_subsystem -nvm_subsystem_id 1 -key port
PORT NVMSS_ID NVMSS_NAME
CL1-A 1 nvmss_id_00001(default_name)
CL1-B 1 nvmss_id_00001(default_name)
```

Description of each column in output example:

PORT

Displays the port number.

Examples

Displays the list of the implemented and the unimplemented NVM subsystems.

```
# raidcom get nvm_subsystem -key detail
NVMSS_ID RGID NVMSS_NAME SECURITY T10PI
HMD HMO_BITS V_NVMSS_ID
0 0 nvmss_id_00001(default_name) ENABLE DISABLE
VMWARE 64 0
1 0 - - -
- - -
. . .
```

Displays the list of the unimplemented NVM subsystem ID.

```
# raidcom get nvm_subsystem -key undefined
NVMSS_ID
2
3
5
. . .
```

raidcom add nvm_subsystem_port

Supported storage systems:

- VSP 5000 series
- VSP E1090

Adds an NVM subsystem port to the NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom add nvm_subsystem_port {-nvm_subsystem_id <nvm subsystem id>} -
port <port#> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

-port <port#>

Specifies the port number.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Adding the port: CL1-A to the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom add nvm_subsystem_port -nvm_subsystem_id 1 -port CL1-A -
request_id auto
```

raidcom delete nvm_subsystem_port

Supported storage systems:

- VSP 5000 series
- VSP E1090

Deletes an NVM subsystem port from the NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom delete nvm_subsystem_port {-nvm_subsystem_id <nvm subsystem id>}
-port <port#> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

-port <port#>

Specifies the port number.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Deleting the port: CL1-A from the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom delete nvm_subsystem_port -nvm_subsystem_id 1 -port CL1-A -
request_id auto
```

raidcom get nvm_subsystem_port

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays the NVM subsystem port information.

Syntax

```
raidcom get nvm_subsystem_port {-port <port> | -nvm_subsystem_id <nvm
subsystem id>}
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

-port <port>

Specifies the port number.

Example

Displaying the information about an NVM subsystem port registered on the NVM subsystem ID 1.

```
# raidcom get nvm_subsystem_port -nvm_subsystem_id 1
PORT NVMSS_ID NVMSS_NAME
CL1-A 1 my_nvm_subsystem
CL1-C 1 my_nvm_subsystem
```

Description of each column in output example:**PORT**

Displays the port number.

NVMSS_ID

Displays the NVM subsystem ID.

NVMSS_NAME

Displays the NVM subsystem name.

Example

Displaying the information about an NVM subsystem on which the port: CL1-A is registered as an NVM subsystem port.

```
# raidcom get nvm_subsystem_port -port CL1-A
PORT NVMSS_ID NVMSS_NAME
CL1-A 1 my_nvm_subsystem
CL1-A 2 NVMSS_2
```

raidcom add namespace_path

Supported storage systems:

- VSP 5000 series
- VSP E1090

Adds a host-namespace path by specifying a host NQN and a namespace registered on the NVM subsystem.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the `raidcom get command_status` command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom add namespace_path {-nvm_subsystem_id <nvm subsystem id>} -ns_id
<ns#> -host_nqn <host_nqn> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>
Specifies the NVM subsystem ID.

--ns_id <ns#>
Specifies the namespace ID.

-host_nqn <host_nqn>
Specifies the host NQN.

-request_id auto
<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Adding a host-namespace path between the host NQN: `nqn.xxx` and the namespace with the namespace ID 1 on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom add namespace_path -nvm_subsystem_id 1 -ns_id 1 -host_nqn
nqn.xxx -request_id auto
```

raidcom delete namespace_path

Supported storage systems:

- VSP 5000 series
- VSP E1090

Deletes a host-namespace path.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

When you run this command, the processing is run asynchronously with the command input. Verify that the processing is complete by using the **raidcom get command_status** command. There is a limit to the number of asynchronous commands that can be accepted simultaneously. For details, see [Maximum number of acceptable asynchronous commands \(on page 269\)](#).

Syntax

```
raidcom delete namespace_path {-nvm_subsystem_id <nvm subsystem id>} -
ns_id <ns#> -host_nqn <host_nqn> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

--ns_id <ns#>

Specifies the namespace ID.

-host_nqn <host_nqn>

Specifies the host NQN.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Deleting a host-namespace path between the host NQN: `nqn.xxx` and the namespace with the namespace ID 1 on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom delete namespace_path -nvm_subsystem_id 1 -ns_id 1 -host_nqn
nqn.xxxx -request_id auto
```

raidcom get namespace_path

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays a host-namespace path.

Syntax

```
raidcom get namespace_path {-nvm_subsystem_id <nvm subsystem id> [-ns_id
<ns#>]}
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

--ns_id <ns#>

Specifies the namespace ID.

Examples

Displaying the information about a host-namespace path for the namespace with the namespace ID 1 on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom get namespace_path -nvm_subsystem_id 1 -ns_id 1
NVMSS_ID NVMSS_NAME NSID LDEV# HOST_NQN
1 nvms_id_00001(default_name) 1 10 nqn.xxx
```

Displaying the information about a host-namespace path for the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom get namespace_path -nvm_subsystem_id 1
NVMSS_ID NVMSS_NAME NSID LDEV# HOST_NQN
1 nvms_id_00001(default_name) 1 10 nqn.xxx
1 nvms_id_00001(default_name) 2 11 nqn.xxx
```

Description of each column in output example:**NVMSS_ID**

Displays the NVM subsystem ID.

NVMSS_NAME

Displays the NVM subsystem name.

NSID

Displays the namespace ID.

LDEV#

Displays the actual LDEV ID in the storage system.

HOST_NQN

Displays the host NQN.

raidcom add namespace

Supported storage systems:

- VSP 5000 series
- VSP E1090

Assigns an LDEV to a namespace that is created on the specified NVM subsystem.

This command is run asynchronously with the command input. Verify that the operation is complete by using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom add namespace {-nvm_subsystem_id <nvm subsystem id>} {-ns_id
<ns#> | -ns_id auto} -ldev_id <ldev#> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

{-ns_id <ns#> | -ns_id auto}

To create a namespace, specifies whether to assign the namespace ID manually or automatically.

-ldev_id <ldev#>

Specifies the LDEV number (0 to 65279).

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Assigning the LDEV with the LDEV number 10 to the namespace with the namespace ID automatically assigned that is created on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom add namespace -nvm_subsystem_id 1 -ns_id auto -ldev_id 10 -
request_id auto
```

raidcom delete namespace

Supported storage systems:

- VSP 5000 series
- VSP E1090

Deletes a namespace.

This command is run asynchronously with the command input. Verify that the operation is complete by using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom delete namespace {-nvm_subsystem_id <nvm subsystem id>} -ns_id
<ns#> -request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

-ns_id <ns#>

Specifies the namespace ID.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Deleting the namespace with the namespace ID 1 on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom delete namespace -nvm_subsystem_id 1 -ns_id 1 -request_id auto
```

raidcom modify namespace

Supported storage systems:

- VSP 5000 series
- VSP E1090

Changes a namespace. The namespace nickname can be changed.

This command is run asynchronously with the command input. Verify that the operation is complete by using the **raidcom get command_status** command.



Note: If a failure occurs in a processor on an MP blade or a controller, this command cannot be run even on an MP blade or a controller on which no failure occurs.

Syntax

```
raidcom modify namespace {-nvm_subsystem_id <nvm subsystem id>} -ns_id
<ns#> {-set_namespace_name <namespace name> | -reset_namespace_name} -
request_id auto
```

Options and parameters

-nvm_subsystem_id <nvm subsystem id>

Specifies the NVM subsystem ID.

-ns_id <ns#>

Specifies the namespace ID.

-set_namespace_name <namespace name>

Specifies the namespace nickname to which it is changed.

-reset_namespace_name

Deletes the namespace nickname.

-request_id auto

<request#> is a request ID assigned each time the command is run. For more information, see [Request ID function \(on page 293\)](#).

Example

Adding a nickname: `my_namespace` to the namespace with the namespace ID 1 on the NVM subsystem with the NVM subsystem ID 1.

```
# raidcom modify namespace -nvm_subsystem_id 1 -ns_id 1 -
set_namespace_name my_namespace -request_id auto
```

raidcom get namespace

Supported storage systems:

- VSP 5000 series
- VSP E1090

Displays the namespace information.

Syntax

```
raidcom get namespace {-nvm_subsystem_id <nvm subsystem id>} [-ns_id
<ns#>] [-key <keyword>]
```

Options and parameters**-nvm_subsystem_id <nvm subsystem id>**

Specifies the NVM subsystem ID.

-ns_id <ns#>

Specifies the namespace ID.

[-key <keyword>]

Specifies a display keyword:

- opt: Displays the detailed information about the namespace.

Example

Displaying the information about the namespace with the NVM subsystem ID 1.

```
# raidcom get namespace -nvm_subsystem_id 1
NVMSS_ID NVMSS_NAME NSID LDEVID CAPACITY(BLK)
1 nvmss_id_00001(default_name) 1 10 16777216
1 nvmss_id_00001(default_name) 2 11 16777216
```

Description of each column in output example:**NVMSS_ID**

Displays the NVM subsystem ID.

NVMSS_NAME

Displays the NVM subsystem name.

NSID

Displays the namespace ID.

LDEVID

Displays the actual LDEV ID in the storage system.

CAPACITY(BLK)

Displays the namespace capacity in block sized increments.

Example

Displaying the detailed information about the namespace with the NVM subsystem ID 1.

```
# raidcom get namespace -nvm_subsystem_id 1 -key opt
NVMSS_ID NVMSS_NAME NSID LDEVID CAPACITY(BLK)
NS_NAME
1 nvmss_id_00001(default_name) 1 10 16777216
my_namespace
1 nvmss_id_00001(default_name) 2 11 16777216
namespace2
```

NS_NAME

Displays the namespace nickname.

If the namespace nickname is not assigned, a hyphen (-) is displayed.

raidcom get ctg

Displays the CTG information.

Syntax

```
#raidcom get ctg -ctg_type <ctg type>
```

Options and parameters

-ctg_type <ctg type>

Specifies the type of the displayed CTG.

`local_replication` (VSP One B20): Displays the CTG information in ShadowImage and Thin Image.

`remote_replication`: Displays the CTG information in TrueCopy, TrueCopy for Mainframe, Universal Replicator, Universal Replicator for Mainframe, and global-active device.

Example 1

Displays the CTG information in ShadowImage and Thin Image.

```
# raidcom get ctg -ctg_type local_replication
CTG# STS NUMOF_PAIR
0 FREE -
1 CUSE 1
:
:
2047 FREE -
```

Example 2

Displays only the CUSE CTG information of the CTG information in ShadowImage and Thin Image.

```
# raidcom get ctg -ctg_type local_replication -I100 | rmawk @L=eq:1 -o @2=eq:"CUSE"
CTG# STS NUMOF_PAIR
1 CUSE 2
2 CUSE 2
3 CUSE 2
4 CUSE 2
5 CUSE 2
6 CUSE 2
```

Description of each column in output example:**CTG#**

Displays the CTG ID in decimal.

STS

Displays the CTG status.

- CUSE: The consistency group used in ShadowImage.
- SUSE: The consistency group used in Thin Image.
- FREE: The consistency group that is not used or registered.
- CHG: The consistency group status is being changed.
- - (hyphen): This information is not valid.

NUMOF_PAIR

Displays the number of pairs in the CTG in decimal. A hyphen (-) is displayed if the CTG is not available.

Example 3

Displays the CTG information in TrueCopy, TrueCopy for Mainframe, Universal Replicator, Universal Replicator for Mainframe, and global-active device.

```
# raidcom get ctg -ctg_type remote_replication
CTG#  STS
      0   FREE
      1   USE
```

Description of each column in output example:**CTG#**

Displays the CTG ID in decimal.

STS

Displays the CTG status.

- USE: The consistency group used in TrueCopy, TrueCopy for Mainframe, Universal Replicator, Universal Replicator for Mainframe, and global-active device
- FREE: The consistency group that is not used

Appendix A: Valid parameter ranges for CCI options (VSP One Block)

This table lists the valid parameter ranges for CCI options for VSP One Block.

Parameter	Valid values
Serial number	800000 to 899999
LDEV ID	VSP One B24: 0 to 16383 VSP One B26: 0 to 32767 VSP One B28: 0 to 49151
Pool ID	0 to 63
Port number	VSP One B24: CLx-y (x: 1 to 8, y: A,D) VSP One B26 and VSP One B28: CLx-y (x: 1 to 8, y: A to D)
LU number	0 to 2,047
Resource group ID	1 to 1,023
Parity group ID	VSP One B20: gno: 1 to 11 VSP One B20 sgno: <ul style="list-style-type: none">▪ If the gno is 1 to 10, 1 to 32▪ If the gno is 11, 1 to 24
Drive location	00-00 to 00-23 or 50-00 to 53-11
iSCSI or NVMe/TCP virtual port ID	0 to 15
SSID	0004 to FFFE
Loop ID for port	0x01 to 0xEF
VLAN ID	1 to 4094
CU #	VSP One B24: 0 to 3F VSP One B26: 0 to 7F VSP One B28: 0 to BF
DB number	50 to 53

Parameter	Valid values
LDEV ID used in a virtual storage machine	0 to 65279
Virtual port number	Same as the port number
Host group ID or iSCSI target ID	0 to 254
SSID for LDEV in a virtual storage machine	Same as SSID
Port speed	0 (auto), 1, 4, 8, 10, 16, 25, 32, 64, 100
NVM subsystem ID	0 to 2047
Virtual NVM subsystem ID	0 to 2047
Namespace ID	1 to 2048

Appendix B: Valid parameter ranges for CCI options (VSP 5000 series)

This table lists the valid parameter ranges for CCI options for VSP 5000 series.

Parameter	Valid values
Serial number	500000 to 599999
LDEV ID	0 to 65279
Pool ID	0 to 127
Port number	CLx-y x: 1 to C y: A to R (except I and O) A to D, or J to M, when x is 9 to C.
LU number	0 to 4095
Resource group ID	1 to 1023
Parity group ID	gno: 1 to 24 sgno: 1 to 32
iSCSI virtual port ID	0 to 15
SSID	0004 to FEFF
Tiering policy ID	0 to 31
Loop ID for port	0x01 to 0xEF
VLAN ID	1 to 4094
CU #	00 to FE
LDEV ID used in a virtual storage machine	0 to 65279
Virtual port number	CLx-y x: 1 to C y: A to R (except I and O)

Parameter	Valid values
	A to D, or J to M, when x is 9 to C.
Host group ID or iSCSI target ID	0 to 254
SSID for LDEV in a virtual storage machine	0004 to FEFF
Port speed	0 (auto), 4, 8, 10, 16, 32
NVM subsystem ID	0 to 2047
Virtual NVM subsystem ID	0 to 2047
Namespace ID	1 to 4096

Appendix C: Valid parameter ranges for CCI options (VSP G130, G/F350, G/F370, G/F700, G/F900)

This table lists the valid parameter ranges for CCI options for VSP G130, G/F350, G/F370, G/F700, G/F900.

Parameter	Valid values
Serial number	VSP E590: 600000 to 699999 VSP E790: 600000 to 699999 VSP E990: 400000 to 499999 VSP E1090: 700000 to 799999
LDEV ID	VSP E590: 0 to 32767 VSP E790: 0 to 49151 VSP E990, VSP E1090: 0 to 65279
Pool ID	VSP E590: 0 to 63 VSP E790: 0 to 63 VSP E990, VSP E1090: 0 to 127
Port number	VSP E590: CLx-y (x: 1 to 8, y: A to C) VSP E790: CLx-y (x: 1 to 8, y: A to C) VSP E990, VSP E1090: CLx-y (x: 1 to 8, y: A to M, except I)
LU number	0 to 2,047
Resource group ID	1 to 1,023
Parity group ID	VSP E590: gno: 1 to 15 ¹ , sgno: 1 to 32 ² VSP E790: gno 1 to 15 ¹ , sgno: 1 to 32 ² VSP E990: gno: 1 to 1, sgno: 1 to 32 VSP E1090: gno 1 to 15, sgno: 1 to 32
Drive location	0-0 to 23-59 (Maximum values using DB60. The maximum values might vary depending on your operation environment.)

Parameter	Valid values
iSCSI virtual port ID	0 to 15
SSID	0000 to FFFE
Tiering policy ID	0 to 31
Loop ID for port	0x01 to 0xEF
VLAN ID	1 to 4094
CU #	VSP E590: 0 to 7F VSP E790: 0 to BF VSP E990, VSP E1090: 0 to FE
DB number	0 to 47
LDEV ID used in a virtual storage machine	0 to 65279
Virtual port number	VSP E590: CLx-y (x: 1 to 8, y: A to C) VSP E790: CLx-y (x: 1 to 8, y: A to C) VSP E990, VSP E1090: CLx-y (x: 1 to 8, y: A to M, except I)
Host group ID or iSCSI target ID	0 to 254
SSID for LDEV in a virtual storage machine	0000 to FFFE
Port speed	0 (auto), 1, 4, 8, 10, 16, 25, 32
NVM subsystem ID	0 to 2047
Namespace ID	1 to 2048
Notes: <ol style="list-style-type: none"> 1. The valid value is from 1 to 1 for a parity group consisting only of drives on the controller chassis. The valid value is from 2 to 15 for a parity group containing drives on the drive box. 2. The valid value is from 1 to 8 for a parity group consisting only of drives on the controller chassis. The valid value is from 1 to 32 for a parity group containing drives on the drive box. 	

Hitachi Vantara

Corporate Headquarters
2535 Augustine Drive
Santa Clara, CA 95054 USA



HitachiVantara.com/contact